



Pipeline

Scrape HackerNews (Show/Top/New), Reddit (r/artificial, r/MachineLearning, r/singularity, r/LocalLLaMA), TechCrunch AI, VentureBeat AI, MIT Tech Review, The Verge AI, Wired AI, ArXiv CS.AI

Filter 40+ AI keyword patterns — LLMs, alignment, safety, research, tooling

Dedup Title-similarity (45% threshold) merges cross-source variants

Score Authenticity = 0.5 base + 0.3 per extra source + diversity + HN score bonuses

Rank Sorted by authenticity. Below 0.25 threshold are dropped.

Authenticity Tiers

- VERIFIED 0.80 - 1.00** 3+ independent sources including media
- CONFIRMED 0.50 - 0.79** 2 sources or 1 high-quality media
- EMERGING 0.25 - 0.49** Single-source; passes AI keyword filter

Source Breakdown

Source	Articles	Category
ArXiv CS.AI	124	Media
TechCrunch AI	12	Media
The Verge AI	7	Media
HackerNews	7	Community
MIT Tech Review	4	Media

VERIFIED INTELLIGENCE — 3 stories

VERI

The AI data center e-waste problem is huge — and getting bigger

The Verge AI

+1 source

80%

E-waste from the AI boom has been vastly underestimated, a new report warns. By 2050, it could become enough trash to fill 23 million shipping containers - roughly enough 40-foot containers to circle the world six times if lined up in a row. It's a significantly higher estimate...

<https://www.theverge.com/ai-artificial-intelligence/996470/ai-data-center-e-waste-ban>

VERI

Meet a mouse whose brain cortex is made up of human cells

MIT Tech Review

+1 source

80%

Multiple cameras tracked a mouse as it wandered around a small arena. A computer charted its position and speed, leaving Pong-like traces on a monitor. The reason to watch this rodent so carefully? Nearly half its brain volume had been replaced with human cells. The effort to...

<https://www.technologyreview.com/2026/09/16/1144210/meet-a-mouse-whose-brain-cortex-is-...>

VERI

Building the materials foundation for AI

MIT Tech Review

+1 source

80%

The AI boom is becoming a materials challenge. As AI pushes computing into new territory, the materials behind that infrastructure are becoming just as crucial as the algorithms running on it. Semiconductors and data centers are approaching physical limits around performance,...

<https://www.technologyreview.com/2026/09/16/1144014/building-the-materials-foundation-f-...>

CONFIRMED SIGNALS — 148 stories

CONF

Show HN: How Stale Is Your AI? Release age and training cutoff for 20 models

HackerNews

50%

HN 73

Hi everyone, Been working on Otis, an open-source ai agent that gives you one minimal experience across local and hosted open-weight models, privacy-focused by design. On setup it recommends a local model based on the hardware Otis is running on, downloads it and runs it through...

<https://stale.jock.pl/>

CONF

Show HN: TurboBench, the Compression Lie Detector, 100 Codecs, Daily Update

HackerNews

50%

HN 5

TurboBench: The Ultimate Data Compression Benchmark 100+ Codecs. The Powerhouse in a Single Executable — The Compression Lie Detector Tired of misleading benchmarks with I/O overhead, cache effects, and CPU throttling? TurboBench delivers high-precision, zero-overhead results...

<https://github.com/powturbo/TurboBench/blob/master/benchmark.md>

CONF

Show HN: Pixel Agents - A pixel-art mission control for your Claude Code agents

HackerNews

50%

HN 4

No summary — see link for full article.

<https://mateovalle.github.io/pixel-agents/>

CONF

Show HN: YOLO FPS-per-dollar numbers for edge AI boards (\$75-\$215)

HackerNews

50%

HN 4

No summary — see link for full article.

<https://github.com/chorylee/edge-ai-benchmarks>

CONF

Show HN: Put an AI agent on a FaceTime audio/video call (open source, WebRTC)

HackerNews 50% HN 3

Hey everyone, Gary here from Chert (YC P26). We just open-sourced a framework for putting AI agents on FaceTime audio/video calls: <https://github.com/cherthq/chert-facetime-opensource> . There's also a managed inbound/outbound API if you don't want to run it yourself:...

<https://github.com/cherthq/chert-facetime-opensource>

CONF

Show HN: OmnisBench, a re-gradable, open LLM routing benchmark on fresh tasks

HackerNews 50% HN 3

No summary — see link for full article.

<https://github.com/Fortitude-Group/OmnisBench>

0 CONF

Show HN: Llmbridge, a C++ LLM gateway with sub-millisecond overhead

HackerNews 50% HN 3

No summary — see link for full article.

<https://github.com/kottos-ai/llmbridge>

1 CONF

Iceland-based Treble raises \$18 million for its voice simulation platform

TechCrunch AI 50%

Treble's voice simulation platform is used by voice AI model developers, AI wearable, and robotics companies

<https://techcrunch.com/2026/09/16/iceland-based-treble-raises-18-million-for-its-voice-...>

2 CONF

Your startup's next teammate might be an AI agent: Gusto, Insight Partners, and Leland explain what that changes at TechCrunch Disrupt 2026

TechCrunch AI 50%

This session will explore how early-stage companies are building teams where humans and AI agents work alongside each other — and how founders can do that without sacrificing speed, accountability, or culture. Learn more at TechCrunch Disrupt 2026. Register before September 25...

<https://techcrunch.com/2026/09/16/your-startups-next-teammate-might-be-an-ai-agent-gust...>

3 CONF

Snap tries to make the case again for its \$2,200 smart glasses

TechCrunch AI 50%

Since Specs' debut earlier this year, Snap has clearly been looking for an opportunity to explain why the smart glasses deserve to exist.

<https://techcrunch.com/2026/09/16/snap-tries-to-make-the-case-again-for-its-2200-smart-...>

4 CONF

AI Gore says the real AI risk isn't data centers

TechCrunch AI 50%

In an interview with TechCrunch, AI Gore suggested he isn't losing sleep over AI data center emissions — he's more worried about the AI industry's own warnings about where the technology is headed.

<https://techcrunch.com/2026/09/16/ai-gore-has-a-surprisingly-calm-take-on-the-ai-data-c...>

5 **CONF****Anthropic and OpenAI want to embed safety evaluators. Will they really be independent?****TechCrunch AI** 50%

Anthropic and OpenAI want to embed independent safety evaluators inside their AI labs. Researchers welcome the unprecedented access, but warn meaningful oversight requires transparency, independence, and eventually regulation.

<https://techcrunch.com/2026/09/16/anthropic-and-openai-want-to-embed-safety-evaluators-...>

6 **CONF****After accusations of selling 'perv glasses,' Meta prepares to sell a pair without a camera****TechCrunch AI** 50%

Can Meta dodge the "pervert glasses" accusations with a new camera-free product?

<https://techcrunch.com/2026/09/16/after-accusations-of-selling-perv-glasses-meta-prepar...>

7 **CONF****AI labs want in-house auditors — but maybe they should shut the front door first****TechCrunch AI** 50%

There may be a simpler and more effective fix for rogue agents, hiding in plain sight.

<https://techcrunch.com/2026/09/16/ai-labs-want-in-house-auditors-but-maybe-they-should-...>

8 **CONF****Your AI agents can now control your Google Home devices****TechCrunch AI** 50%

Google is launching early access to a new MCP server for Google Home, allowing AI agents like Claude, ChatGPT, and others to control connected devices, review camera summaries, and access smart home activity using natural language.

<https://techcrunch.com/2026/09/16/your-ai-agents-can-now-control-your-google-home-devices/>

9 **CONF****Anthropic merges Claude chat and Cowork in one interface****TechCrunch AI** 50%

Anthropic is initially releasing these features to Pro and Max plan subscribers.

<https://techcrunch.com/2026/09/16/anthropic-merges-claude-chat-and-cowork-in-one-interf...>

10 **CONF****Former Infosys chief's AI startup nabs another \$53M****TechCrunch AI** 50%

The Palo Alto startup says it has landed multiple seven-figure enterprise contracts within months of launch.

<https://techcrunch.com/2026/09/16/former-infosys-chiefs-ai-startup-adds-50m-to-seed-wee...>

1 **CONF****The Download: AI's trillion-dollar gamble and OpenAI's biology data bid****MIT Tech Review** 50%

This is today's edition of The Download, our weekday newsletter that provides a daily dose of what's going on in the world of technology. What's at stake in AI's trillion-dollar gamble When Jessica Wachter, a finance professor at the University of Pennsylvania, wanted to assess...

<https://www.technologyreview.com/2026/09/16/1144205/the-download-ai-trillion-dollar-bui...>

2 **CONF****The Download: AI doomers, whistleblowing agents, and de-aged livers****MIT Tech Review** 50%

This is today's edition of The Download, our weekday newsletter that provides a daily dose of what's going on in the world of technology. The AI industry has taken a doomer turn. What now? AI chiefs Dario Amodei, Sam Altman, Elon Musk, and Demis Hassabis are suddenly all in...

<https://www.technologyreview.com/2026/09/15/1144141/the-download-ai-extinction-whistleb...>

3 **CONF****Snap is launching a new Specs AI tool, and it's coming to iOS and Mac****The Verge AI** 50%

Snap is introducing "Specs Intelligence," a new AI assistant that can connect other digital accounts to help you with things like work tasks and keeping track of travel information. It seems similar to AI assistants like Meta's Muse and Gemini's Spark, though Snap is pitching...

<https://www.theverge.com/tech/996078/snap-specs-intelligence-ai-agent-ios-mac>4 **CONF****The 2.5-hour AI-generated Odyssey movie is 2.5 hours too long****The Verge AI** 50%

Christopher Nolan's engrossing take on The Odyssey dominated at the box office and spurred a newfound interest in classic literature among filmgoers. But a new retelling of the story made entirely with AI is so bad that it might just make viewers hate the original tale...

<https://www.theverge.com/entertainment/996499/ai-odyssey-movie-review>5 **CONF****Apple might make servers again to cash in on the AI rush****The Verge AI** 50%

According to The Information, Apple is planning to get back into the server game and might just pair up with Nvidia to make it happen. Apple retired its Xserve line in 2011 and has largely left enterprise machines to other manufacturers since. But the growing demand for compute...

<https://www.theverge.com/tech/996321/apple-servers-ai-nvidia>6 **CONF****Google will now let any AI agent run your smart home****The Verge AI** 50%

Google is opening up its smart home to AI agents, letting tools like Claude and Open Claw access and control your connected devices and analyze your home's data using the standardized Model Context Protocol. Google Home MCP is a new integration that lets third-party AI agents...

<https://www.theverge.com/tech/996310/google-home-mcp-integration-agent-ai-smart-home-...>7 **CONF****Claude comes for Gemini with its own take on Docs and Slides****The Verge AI** 50%

Claude is getting a pair of new tools today: Docs and Slides. They'll let you create documents and presentations through Claude chats, which you can export, edit, and share with other users. As part of the announcement, Anthropic is also simplifying how Claude chats work...

<https://www.theverge.com/ai-artificial-intelligence/996234/anthropic-one-claude-cowork-...>8 **CONF****The sexy AI-powered dating app scams are here****The Verge AI** 50%

Security researcher Matthew "Zigula" Gore-Kormanik was analyzing a fraudulent dating app called Dora when he got a pop-up message saying he was receiving a call from Jennifer. According to her bio, she's a 41-year-old Sagittarius with red hair, blue eyes, and piercings. She...

<https://www.theverge.com/ai-artificial-intelligence/995348/ai-dating-app-scams>

9 CONF

Making AI-Assisted Claims Independently Challengeable: Publication Authority and a Protocol for Falsifiable Publication Records

ArXiv CS.AI 50%

arXiv:2609.17631v1 Announce Type: new Abstract: AI-assisted claims can appear authoritative when evidence, analysis, human authorization, presentation, and correction history refer to different states. Provenance, attestation, and transparency expose history but alone do not...

<https://arxiv.org/abs/2609.17631>

0 CONF

EvolveTrade: Experience-Driven Policy Refinement for Self-Evolving LLM Trading Agents

ArXiv CS.AI 50%

arXiv:2609.17632v1 Announce Type: new Abstract: Large language model (LLM) trading agents can combine market data, news, and executable analysis, but their behavior is often controlled by static hand-written tool-use policies that are fixed before deployment. This limits their...

<https://arxiv.org/abs/2609.17632>

1 CONF

Physics-Constrained Digital Twins for Sensor Integrity in Urban Pedestrian Flow: Detecting Stealthy False Data Injection with Conformal Guarantees

ArXiv CS.AI 50%

arXiv:2609.17635v1 Announce Type: new Abstract: City pedestrian counting systems now feed economic indicators, planning decisions and safety operations, yet the twins built on top of them treat the incoming stream as ground truth. We study what happens when it is not. We...

<https://arxiv.org/abs/2609.17635>

2 CONF

CzechTopic: A Benchmark for Zero-Shot Topic Localization in Historical Czech Documents

ArXiv CS.AI 50%

arXiv:2603.03884v2 Announce Type: replace-cross Abstract: Topic localization aims to identify spans of text that express a given topic defined by a name and description. To study this task, we introduce a human-annotated benchmark based on Czech historical documents, containing...

<https://arxiv.org/abs/2603.03884>

3 CONF

GraphEcho: Structural Redundancy and Evidence Provenance in LLM Graph Agents

ArXiv CS.AI 50%

arXiv:2609.17695v1 Announce Type: new Abstract: A large language model (LLM) agent can follow more graph paths without acquiring more independent evidence. GraphEcho tests whether agents mistake these repeated encounters for additional corroboration. The benchmark varies path...

<https://arxiv.org/abs/2609.17695>

4 CONF

NeMo Data Designer: An Extensible Framework for Multimodal Synthetic Data Generation

ArXiv CS.AI 50%

arXiv:2609.17699v1 Announce Type: new Abstract: We present NeMo Data Designer (NDD), an open-source, general-purpose framework for multi-modal synthetic data generation (SDG). Designed to be intuitive to use, NDD provides a declarative configuration format in which human and/or...

<https://arxiv.org/abs/2609.17699>

5 CONF

A Systematic Evaluation of the COTQ Provincial Land Cover Product: Structural Consistency, Spectral Separability, and Relative Positioning Against ESA, ESRI, and Google Products

ArXiv CS.AI 50%

arXiv:2609.17731v1 Announce Type: new Abstract: High-resolution land use and land cover (LULC) products derived from Sentinel-2 imagery are widely used for environmental monitoring and land management, yet their performance can vary across regions with complex ecological...

<https://arxiv.org/abs/2609.17731>

6 CONF

GenStream: Semantic Streaming Framework for Generative Reconstruction of Human-centric Media

ArXiv CS.AI 50%

arXiv:2609.18634v1 Announce Type: cross Abstract: Video streaming dominates global internet traffic, yet conventional pipelines remain inefficient for structured, human-centric content such as sports, performance, or interactive media. Standard codecs re-encode entire frames,...

<https://arxiv.org/abs/2609.18634>

7 CONF

A Four-Stage Decomposition of Word-Problem Solving and Mechanistic Fragility in LLM Math Reasoning

ArXiv CS.AI 50%

arXiv:2609.17804v1 Announce Type: new Abstract: Large language models solve grade-school math word problems with high accuracy, yet a single irrelevant clause inserted into the problem can collapse it. We reconcile these observations with a mechanistic account. We show that the...

<https://arxiv.org/abs/2609.17804>

8 CONF

The Inference Engineering Pareto Atlas: Which Optimizations Dominate the Cost, Quality, and Latency Frontier?

ArXiv CS.AI 50%

arXiv:2609.17863v1 Announce Type: new Abstract: LLM inference optimizations report speedups on different models, GPUs, prompts, and quality metrics, making them hard to compare or combine. We build a cost, quality, and latency Pareto atlas to identify the best configurations for...

<https://arxiv.org/abs/2609.17863>

9 CONF

Measuring AI Leadership: Development and Validation of a Multidimensional Measure for AI-Native Organizations

ArXiv CS.AI 50%

arXiv:2609.17965v1 Announce Type: new Abstract: AI is changing what leaders must judge, explain, learn, and coordinate, yet existing measures do not capture these behaviors at the level needed to study leadership in AI-enabled work. We develop the AI Leadership Battery, which...

<https://arxiv.org/abs/2609.17965>

0 CONF

Memory Has Geometry: Non-Uniform Geometric Memory for Long-Horizon Personalized AI

ArXiv CS.AI 50%

arXiv:2609.17969v1 Announce Type: new Abstract: Long-term memory is becoming a core substrate for personalized AI, yet most systems still represent personalization as discrete records in a largely static latent space, accessed under one global similarity notion. For data mining,...

<https://arxiv.org/abs/2609.17969>

1 **CONF****When to Call an LLM: A Confidence-Gated Hybrid for Cost-Effective Emotion Recognition in Conversational AI****ArXiv CS.AI** **50%**

arXiv:2609.17977v1 Announce Type: new Abstract: Emotion recognition in conversation (ERC) is a production capability behind agent-assist prompts, escalation routing, and post-call analytics in contact-center-as-a-service (CCaaS) platforms, where cost and latency constraints...

<https://arxiv.org/abs/2609.17977>

2 **CONF****Contiguity, Not Importance: Budgeted Repair of Stale KV Caches After Document Edits****ArXiv CS.AI** **50%**

arXiv:2609.17983v1 Announce Type: new Abstract: KV-cache reuse can reduce inference cost in retrieval-augmented generation and agentic systems, but cached contexts may become stale when retrieved knowledge, working memory, or user state is edited. Under causal self-attention,...

<https://arxiv.org/abs/2609.17983>

3 **CONF****TuiML: Machine Learning for AI Agents****ArXiv CS.AI** **50%**

arXiv:2609.17984v1 Announce Type: new Abstract: Machine-learning libraries such as Weka and scikit-learn were designed for human programmers. Language-model agents now use these same libraries by recalling APIs from memory and writing code, an approach that hides what a library...

<https://arxiv.org/abs/2609.17984>

4 **CONF****Reflect, Revise, Reuse: Training-Free Skill Evolution for GUI Agents****ArXiv CS.AI** **50%**

arXiv:2609.17653v1 Announce Type: cross Abstract: GUI agents execute long-horizon tasks on dynamic graphical user interfaces, where pop-ups, delayed loads, and relocated widgets routinely invalidate plans fixed before execution. Recent agent-skill frameworks encapsulate reusable...

<https://arxiv.org/abs/2609.17653>

5 **CONF****Multimodal Conditioning of Fine-Tuned Stable Diffusion XL for Controllable and Culturally Faithful Ulos Motif Generation****ArXiv CS.AI** **50%**

arXiv:2609.17987v1 Announce Type: new Abstract: The traditional Batak Ulos weaving industry faces growing challenges in producing diverse, innovative motifs due to limitations in conventional, manually driven design methods. This study proposes a multimodal generative framework...

<https://arxiv.org/abs/2609.17987>

6 **CONF****The Other Half of the Memory Wall: Serving 35B MoEs from SSD with Trained Routing Prediction****ArXiv CS.AI** **50%**

arXiv:2609.18063v1 Announce Type: new Abstract: Mixture-of-experts (MoE) inference on consumer hardware is bounded by weight memory: a 35B-class model is 19.5GB at 4-bit, and sparsity shrinks the compute per token, not the bytes that must be held. Naive offloading to SSD does...

<https://arxiv.org/abs/2609.18063>

7 CONF

Teaching AI, Robotics, & Community: A Hubs-Based K-12 Education Framework for Reaching Rural Schools

ArXiv CS.AI 50%

arXiv:2609.18072v1 Announce Type: new Abstract: K-12 robotics and AI education remains difficult to scale, especially in rural regions lacking sustained technical mentorship. Programs like FIRST provide competition pathways and instructional opportunities, but they do not...

<https://arxiv.org/abs/2609.18072>

8 CONF

AutoTuneBench: Trustworthy Measurement for Agent Auto-Tuning of LLM Serving Engines

ArXiv CS.AI 50%

arXiv:2609.18123v1 Announce Type: new Abstract: Large language model agents tune GPU kernels and serving engines through a closed loop of propose, measure, and keep, but the measurements behind this loop are not trustworthy. We characterize four failure modes from a four-day...

<https://arxiv.org/abs/2609.18123>

9 CONF

Designing Agentic AI Workflow Portfolios under Imperfect Selection and Compute Cost

ArXiv CS.AI 50%

arXiv:2609.18126v1 Announce Type: new Abstract: Agentic AI systems often approach the same task through multiple workflows that differ in reasoning strategy, verification structure, and compute cost. A natural deployment policy is to use the workflow with the highest average...

<https://arxiv.org/abs/2609.18126>

0 CONF

ASLEval: Measuring Privacy Exposure Displacement in LLM Agent Sessions

ArXiv CS.AI 50%

arXiv:2609.18864v1 Announce Type: cross Abstract: Privacy evaluations of tool-using LLM agents often inspect a designated action, final response, or attacker report. These local proxies can miss unauthorized exposure elsewhere in a multi-step session and lack common ground truth...

<https://arxiv.org/abs/2609.18864>

1 CONF

Clueing up LLMs with Tool-Augmented Deductive Reasoning

ArXiv CS.AI 50%

arXiv:2609.18736v1 Announce Type: new Abstract: Despite recent advances in large language models (LLMs), performing logically consistent deductive reasoning over extended interactions remains challenging. Tasks that require integrating evidence across multiple reasoning steps,...

<https://arxiv.org/abs/2609.18736>

2 CONF

REPAIR: Resolving Long-Tail Confusion in Scientific Retrievers via Fact-Verified Iterative Refinement

ArXiv CS.AI 50%

arXiv:2609.18262v1 Announce Type: new Abstract: Precise retrieval of scientific information is fundamentally constrained by long-tailed concepts and high fact-sensitivity of scientific corpora. These challenges often limit the effectiveness of dense retrievers and...

<https://arxiv.org/abs/2609.18262>

3 CONF

BENCHCOMPASS: From Scores to Signals for Training and Harness Decisions in Payment-Domain LLMs

ArXiv CS.AI 50%

arXiv:2609.18270v1 Announce Type: new Abstract: Payment operations are a critical financial infrastructure, but the value of large language models in this domain remains unclear because payment rules change quickly, evidence is fragmented, and decisions depend on transaction...

<https://arxiv.org/abs/2609.18270>

4 CONF

Who Audits Whom, on What Substrate, with What Evidence? An Independence-Graded Audit Protocol for Agentic AI

ArXiv CS.AI 50%

arXiv:2609.18272v1 Announce Type: new Abstract: Agentic AI systems plan, invoke tools and act with limited supervision; they are now both the subject of audits and, increasingly, the auditor. Independence, the foundation of assurance, is still applied to them as a binary. We...

<https://arxiv.org/abs/2609.18272>

5 CONF

Building Trust in Artificial Intelligence: A Necessity for Railway Applications

ArXiv CS.AI 50%

arXiv:2609.18278v1 Announce Type: new Abstract: Artificial Intelligence (AI) is currently only applied to non-safety critical applications due to the strict standards and regulations for railway industries. We propose to review the three main fields necessary to increase trust...

<https://arxiv.org/abs/2609.18278>

6 CONF

Where Should Agents Live? Energy-Memory Characterization of Agentic AI for the Edge-Cloud Continuum

ArXiv CS.AI 50%

arXiv:2609.18283v1 Announce Type: new Abstract: As telecommunication networks evolve toward autonomous 5G-Advanced and 6G operations, agentic artificial intelligence (AI) workflows, where large language models (LLMs) execute multi-step reasoning, invoke diagnostic tools,...

<https://arxiv.org/abs/2609.18283>

7 CONF

What Counts as Strategic Reasoning? A Systematic Mapping of Chess Research on Humans, Engines, and Language Models

ArXiv CS.AI 50%

arXiv:2609.18286v1 Announce Type: new Abstract: Chess has long served as a model domain for studying search, expertise, decision-making, and artificial intelligence. The emergence of large language models (LLMs) has renewed the relevance of chess as a controlled environment for...

<https://arxiv.org/abs/2609.18286>

8 CONF

Visual Compliance via Executable Safety Rule Entailment

ArXiv CS.AI 50%

arXiv:2609.18328v1 Announce Type: new Abstract: Recent advances in LLMs and VLMs have enabled safety systems to reason beyond simple risk patterns toward more contextual and semantic safety concerns. However, as risk patterns continue to evolve and safety rules become more...

<https://arxiv.org/abs/2609.18328>

9 **CONF****Faithful yet Collusive: Why Chain-of-Thought Monitoring Cannot Detect Collusion in LLM Pricing Agents under Oligopolistic Competition****ArXiv CS.AI** **50%**

arXiv:2609.18346v1 Announce Type: new Abstract: Large language models (LLM) deployed as autonomous pricing agents may sustain supracompetitive prices through tacit coordination. We develop a causal graph divergence framework that separately measures structural faithfulness and...

<https://arxiv.org/abs/2609.18346>

0 **CONF****EvoUndo: Recoverability-Constrained Self-Evolution for LLM Agent Harnesses****ArXiv CS.AI** **50%**

arXiv:2608.28363v2 Announce Type: replace Abstract: LLM agents increasingly modify their own prompts, tools, middleware, resources, and execution harnesses at runtime. Such self-evolution can improve capability, but a successful mutation may leave persistent effects that cannot...

<https://arxiv.org/abs/2608.28363>

1 **CONF****Cultural Competence in Context: A Large Language Model Passes the Turing Test in Finland****ArXiv CS.AI** **50%**

arXiv:2609.18394v1 Announce Type: new Abstract: We report the results of a Turing Test conducted in Finland in the Finnish language. Because languages and cultural contexts are unevenly represented in LLM training data, we expected the model (ChatGPT 5.2) to perform worse in a...

<https://arxiv.org/abs/2609.18394>

2 **CONF****The Mirage of Calibrated Confidence: Trajectory-Independence of Verbalized Confidence in Vision-Language Models****ArXiv CS.AI** **50%**

arXiv:2609.18453v1 Announce Type: new Abstract: A calibrated Vision-Language Model (VLM) can repeatedly self-correct, say "Wait, I should recheck," arrive at the wrong answer, and still report high confidence. We find that this occurs because verbalized confidence is largely...

<https://arxiv.org/abs/2609.18453>

3 **CONF****Collective Loss of Control in LLM Agent Systems: An Epidemic Account of Mutation, Contagion, and Recovery****ArXiv CS.AI** **50%**

arXiv:2609.18460v1 Announce Type: new Abstract: How does a multi-agent system evolve from a local deviation into collective loss of control? We propose an epidemic explanation organized around accidental mutation, contagion, and recovery. A spontaneous deviation creates a seed;...

<https://arxiv.org/abs/2609.18460>

4 **CONF****GYROval: A Robust Benchmark for Cultural Value Orientation in Large Language Models****ArXiv CS.AI** **50%**

arXiv:2609.18384v1 Announce Type: cross Abstract: We present a robust benchmark for measuring cultural value orientation in large language models on the two Inglehart-Welzel axes over several domains and roles (hence GYROval - Gridded Yielding of Robust value Orientation),...

<https://arxiv.org/abs/2609.18384>

5 **CONF****Reasoning through Evolution: Automatic Meta-path Discovery for LLM-based Fake News Detection**

ArXiv CS.AI 50%

arXiv:2609.18597v1 Announce Type: new Abstract: Propagation structures provide crucial evidence for fake news detection, yet existing approaches primarily rely on supervised GNN-based models, which require substantial labeled data and exhibit limited generalization. Although...

<https://arxiv.org/abs/2609.18597>

6 **CONF****The Uneven Impact of Generative AI on Student Learning: Examining the Roles of Reliance, Evaluation Literacy, and Course Policy in AI-related Courses**

ArXiv CS.AI 50%

arXiv:2609.18676v1 Announce Type: new Abstract: Generative artificial intelligence (GenAI) is changing how students learn, yet the roles of course context, cognitive reliance, evaluation literacy, and early reliance remain underexplored. Using survey responses from 118 students...

<https://arxiv.org/abs/2609.18676>

7 **CONF****Beyond Truncation: Rethinking LLM Decoding as Ensemble Pruning**

ArXiv CS.AI 50%

arXiv:2609.18723v1 Announce Type: new Abstract: We introduce Mahalanobis-Ensemble Decoding (ME-Decoding), a novel Large Language Model (LLM) decoding framework that frames candidate token selection as ensemble pruning. Existing selection strategies rely predominantly on scalar...

<https://arxiv.org/abs/2609.18723>

8 **CONF****Which LLM is Best for Translating Natural Language Goals to PDDL**

ArXiv CS.AI 50%

arXiv:2609.18731v1 Announce Type: new Abstract: Bridging the gap between human intent and machine execution remains a challenge in automated planning, where expressing goals in formal languages like PDDL restricts accessibility to non-experts. This paper empirically evaluates...

<https://arxiv.org/abs/2609.18731>

9 **CONF****CERA-MoA: Co-Evolving Routing Mechanisms with Continually Learning LLM Agents**

ArXiv CS.AI 50%

arXiv:2609.18779v1 Announce Type: new Abstract: Current Mixture-of-Agents (MoA) paradigms generally treat query routing and agent fine-tuning as separate processes, limiting their ability to respond to evolving agent capabilities. This disconnect prevents routing strategies from...

<https://arxiv.org/abs/2609.18779>

10 **CONF****Compositional Policy Violations: When Step-Level Compliance Fails In Agentic AI Workflows**

ArXiv CS.AI 50%

arXiv:2609.18820v1 Announce Type: new Abstract: Agentic workflows now make consequential decisions in regulated settings, and the governance placed around them is almost entirely step-scoped: input-output classifiers, per turn rails, and span-level evaluators. The policies...

<https://arxiv.org/abs/2609.18820>

1 CONF

Infinite-Parameter LLMs: Generating and Adapting Weights from Live Data

ArXiv CS.AI 50%

arXiv:2609.18842v1 Announce Type: new Abstract: The scaling laws hold that a language model grows more capable with more parameters and more training data, and Mixture-of-Experts (MoE) architectures have ridden these laws to remarkable results, activating only a fraction of an...

<https://arxiv.org/abs/2609.18842>

2 CONF

From a River in Gilead to the Inference Distributions of Large Language Models: Covert Dialect Bias and Linguistic Profiling at Scale

ArXiv CS.AI 50%

arXiv:2609.18068v1 Announce Type: cross Abstract: Large language models (LLMs) are increasingly deployed in high-stakes domains such as housing screening. While alignment techniques mitigate explicit racial bias in generated text, they often leave covert attitudinal associations...

<https://arxiv.org/abs/2609.18068>

3 CONF

LM Fight Arena: Benchmarking Large Multimodal Models via Game Competition

ArXiv CS.AI 50%

arXiv:2510.08928v2 Announce Type: replace Abstract: Existing benchmarks for large multimodal models (LMMs) often fail to capture their performance in real-time, adversarial environments. We introduce LM Fight Arena (Large Model Fight Arena), a novel framework that evaluates LMMs...

<https://arxiv.org/abs/2510.08928>

4 CONF

The Death of Schema Linking? Text-to-SQL in the Age of Well-Reasoned Language Models

ArXiv CS.AI 50%

arXiv:2408.07702v2 Announce Type: cross Abstract: Schema linking is a crucial step in Text-to-SQL pipelines. Its goal is to retrieve the relevant tables and columns of a target database for a user's query while disregarding irrelevant ones. However, imperfect schema linking can...

<https://arxiv.org/abs/2408.07702>

5 CONF

WARD: Runtime Workload-Adaptive Vision TRansformer Framework for Dependable Edge AI

ArXiv CS.AI 50%

arXiv:2609.17556v1 Announce Type: cross Abstract: Edge-deployed AI operate under dynamically changing power budgets, reliability requirements, and input distributions, requiring continuous adaptation. Such conditions arise in long-running edge AI applications, including...

<https://arxiv.org/abs/2609.17556>

6 CONF

Structure is not mechanism: high-gain gated-FFN rows across text and genomic foundation models

ArXiv CS.AI 50%

arXiv:2609.17599v1 Announce Type: cross Abstract: A small number of unusually high-gain parameters can exert disproportionate effects in transformer language models, but whether analogous structures recur in genomic foundation models and whether structural geometry determines...

<https://arxiv.org/abs/2609.17599>

7 **CONF****Enhancing Physics-Informed Neural Networks with Domain-aware Fourier Features: Towards Improved Performance and Interpretable Results****ArXiv CS.AI** **50%**

arXiv:2603.02948v2 Announce Type: replace-cross Abstract: Physics-Informed Neural Networks (PINNs) incorporate physics into neural networks by embedding partial differential equations (PDEs) into their loss function. Despite their success in learning the underlying physics, PINN...

<https://arxiv.org/abs/2603.02948>

8 **CONF****Accelerating Diffusion Sampling via Speculative Draft Trees****ArXiv CS.AI** **50%**

arXiv:2609.17691v1 Announce Type: cross Abstract: Speculative sampling accelerates diffusion model generation by drafting inexpensive candidate states and correcting them under a coupling that preserves the target distribution exactly, reducing the number of expensive target...

<https://arxiv.org/abs/2609.17691>

9 **CONF****One Size Does Not Fit All! Dynamic Retriever and Generator Selection for RAG****ArXiv CS.AI** **50%**

arXiv:2609.17709v1 Announce Type: cross Abstract: Retrieval-Augmented Generation (RAG) systems typically employ fixed retriever and generator configurations across queries, despite substantial differences in query complexity and information needs, leading to inefficient...

<https://arxiv.org/abs/2609.17709>

0 **CONF****CALOS: Control-Affine Lyapunov On-manifold Safety Layer for Safe Deep Reinforcement Learning for Quadrotors****ArXiv CS.AI** **50%**

arXiv:2609.17758v1 Announce Type: cross Abstract: Deep Reinforcement Learning has demonstrated remarkable capability in quadrotor control, yet learned policies offer no guarantee of respecting safety constraints during training or deployment. We present CALOS (Control-Affine...

<https://arxiv.org/abs/2609.17758>

1 **CONF****HINT-Plan: Human Intention-Aware Robot Task Planning in Context-Rich Environments using Vision Language Models****ArXiv CS.AI** **50%**

arXiv:2609.17771v1 Announce Type: cross Abstract: Approaches to incorporating human awareness into mobile robot decision-making mainly focus on collision avoidance in low-level motion planning, often overlooking the challenges posed by human presence and high-level behavior. To...

<https://arxiv.org/abs/2609.17771>

2 **CONF****When AI Generates Covariates: Causal Typing and Estimand Drift in Sequential Experiments****ArXiv CS.AI** **50%**

arXiv:2609.17772v1 Announce Type: cross Abstract: AI-generated covariates from notes, conversations, images, and wearable streams can change the causal question when their roles are left unspecified. A generated feature may represent a treatment version, pre-action state,...

<https://arxiv.org/abs/2609.17772>

3 **CONF****Information Set Emulation: Causal Certificates for AI Derived EHR Features****ArXiv CS.AI** **50%**

arXiv:2609.17777v1 Announce Type: cross Abstract: AI and large language models can recover clinically meaningful features from electronic health records (EHRs), but predictive usefulness does not establish admissibility for causal inference. We introduce information set...

<https://arxiv.org/abs/2609.17777>

4 **CONF****AI and Human Approaches to Mathematical Problem Solving****ArXiv CS.AI** **50%**

arXiv:2609.17779v1 Announce Type: cross Abstract: AI systems have begun to report solutions, disproofs, and substantive advances on long-standing mathematical problems, raising questions about whether they approach research in the same way as mathematicians. This study compares...

<https://arxiv.org/abs/2609.17779>

5 **CONF****SAiFE-gym: Model-based Environments for Automated Market Making with Concentrated Liquidity****ArXiv CS.AI** **50%**

arXiv:2609.17788v1 Announce Type: cross Abstract: We present SAiFE_gym, a Python module that provides a collection of simulation environments for studying trading problems in Constant Product Markets (CPMs) with Concentrated Liquidity (CL). These markets give Liquidity Providers...

<https://arxiv.org/abs/2609.17788>

6 **CONF****MINT: Multimodal Imaging-to-Speech Knowledge Transfer for Early Alzheimer's Screening****ArXiv CS.AI** **50%**

arXiv:2602.23994v2 Announce Type: replace-cross Abstract: Alzheimer's disease is a progressive neurodegenerative disorder in which mild cognitive impairment (MCI) precedes dementia. Structural MRI provides biomarkers but requires costly infrastructure, limiting population-scale...

<https://arxiv.org/abs/2602.23994>

7 **CONF****Procedural Pretraining for Molecular Property Prediction****ArXiv CS.AI** **50%**

arXiv:2609.17831v1 Announce Type: cross Abstract: Molecular property prediction is often limited by the small size of labeled downstream datasets, motivating pretraining on large corpora of unlabeled molecules. In this work, we ask whether useful inductive biases can instead be...

<https://arxiv.org/abs/2609.17831>

8 **CONF****The Free Inference Dimension: Complexity Measure for Zero-Collision Navigation under Hypothesis Mixtures****ArXiv CS.AI** **50%**

arXiv:2609.17816v1 Announce Type: cross Abstract: Solomonoff induction frames prediction as a mixture over computable hypotheses, typically leading to identification of the true environment. In a finite meta-reinforcement learning setting with nested constraint families, in our...

<https://arxiv.org/abs/2609.17816>

9 CONF

Reflections on Trusting Trust, Revisited: Contaminating Self-Modifying AI Coding Agents with Poisoned Benchmarks

ArXiv CS.AI 50%

arXiv:2609.17817v1 Announce Type: cross Abstract: Thompson's "Reflections on Trusting Trust" showed that a compiler can be poisoned to reinsert its own backdoor, so that even recompiling clean source reproduces the Trojan. Today, substantial coding work is done by AI coding...

<https://arxiv.org/abs/2609.17817>

0 CONF

CoRe-MARL: Cooperative Redistribution Under Unknown Dynamics Using Recurrent Multi-Agent Reinforcement Learning

ArXiv CS.AI 50%

arXiv:2609.18639v1 Announce Type: cross Abstract: Emergency management assistance programs, such as relief distribution, are essential for delivering necessary supplies to affected communities. However, these programs operate in a decentralized network of local centers that face...

<https://arxiv.org/abs/2609.18639>

1 CONF

Learning Nuclear Structure with AI: Radii and Collectivity

ArXiv CS.AI 50%

arXiv:2609.17838v1 Announce Type: cross Abstract: Low-energy nuclear structure is encoded in a broad body of experimental information across the chart of nuclides. Learning how this information is organized across observables and nuclei can provide a data-driven empirical...

<https://arxiv.org/abs/2609.17838>

2 CONF

Abstention vs. Hallucination: Benchmarking LLM Source Attribution for Scientific Citations

ArXiv CS.AI 50%

arXiv:2405.02228v5 Announce Type: replace-cross Abstract: Large language models (LLMs) increasingly generate citation-backed responses, yet citation hallucination remains a major challenge for trustworthy scientific information access. We introduce REASONS, a benchmark of 12,723...

<https://arxiv.org/abs/2405.02228>

3 CONF

Does AI Assistance Leave a Temporal Fingerprint? Detecting Overreliance in AI-Assisted Writing and Programming

ArXiv CS.AI 50%

arXiv:2609.17883v1 Announce Type: cross Abstract: The rapid adoption of generative AI has made final artifacts unreliable evidence of student learning, and AI detectors that examine only the finished product are inaccurate and ethically contentious. Process data offers an...

<https://arxiv.org/abs/2609.17883>

4 CONF

Walking the Score Manifold: Continuous-time Generative Dynamics on Learned Data Manifolds

ArXiv CS.AI 50%

arXiv:2609.17901v1 Announce Type: cross Abstract: Generative modeling of time-dependent data is typically formulated on a discrete temporal grid, restricting supervision to the observed timestamps in the training data. We instead frame generation as continuous-time evolution on...

<https://arxiv.org/abs/2609.17901>

5 CONF

EDCT-Bench: Uncovering Faithfulness Gaps in VLMs via Explanation-Driven Counterfactual Testing

ArXiv CS.AI 50%

arXiv:2609.17953v1 Announce Type: cross Abstract: Vision-Language Models (VLMs) can produce Natural Language Explanations (NLEs) that sound plausible yet remain inconsistent with the visual evidence they cite. We present Explanation-Driven Counterfactual Testing (EDCT), an...

<https://arxiv.org/abs/2609.17953>

6 CONF

Whom Do AI Agents Work For? Role Assignment Induces Sponsorship Bias in LLM Recommenders

ArXiv CS.AI 50%

arXiv:2609.17989v1 Announce Type: cross Abstract: Large language models (LLMs) now serve as conversational shopping assistants on platforms that also sell advertising. These AI agents face a conflict of duty. They advise consumers who rely on their judgment, yet are deployed by...

<https://arxiv.org/abs/2609.17989>

7 CONF

CompArt: Operationalizing Aesthetic Alignment in Text-to-Image Generation via Principles of Art

ArXiv CS.AI 50%

arXiv:2503.12018v2 Announce Type: replace-cross Abstract: Text-to-Image (T2I) diffusion models have made rapid progress on semantic alignment (generating what is described in the prompt), yet users still lack reliable control over aesthetic composition (how visual elements are...

<https://arxiv.org/abs/2503.12018>

8 CONF

Linguistic Triggers of Gender and Racial Bias in Open-Weight LLMs Applied to Recruitment

ArXiv CS.AI 50%

arXiv:2609.18106v1 Announce Type: cross Abstract: Open-weight large language models are rapidly entering hiring pipelines, yet their discriminatory failure modes -- and the regulatory exposure these create under the EU AI Act high-risk classification (Annex III) and U.S. EEOC...

<https://arxiv.org/abs/2609.18106>

9 CONF

A Comprehensive Review of Generative Physical Artificial Intelligence

ArXiv CS.AI 50%

arXiv:2609.18111v1 Announce Type: cross Abstract: The integration of large-scale foundation models with physical embodiments has led to significant advancements in robotics known as Generative Physical Artificial Intelligence (GPAI). These agentic AI systems autonomously...

<https://arxiv.org/abs/2609.18111>

00 CONF

BadQubits: An LLM-Based Framework for Static Pre-Execution Detection of Structurally Harmful Quantum Circuits

ArXiv CS.AI 50%

arXiv:2609.18965v1 Announce Type: cross Abstract: This paper presents BadQubits, an LLM-based framework for static pre-execution detection of structurally harmful OpenQASM 2.0 circuits. The framework targets physical-execution-layer threats by analyzing submitted circuits prior...

<https://arxiv.org/abs/2609.18965>

01 CONF

Rethinking How We Evaluate Methodological Progress in Health AI

ArXiv CS.AI 50%

arXiv:2609.18134v1 Announce Type: cross Abstract: Methodological progress in artificial intelligence (AI) for electronic health records (EHRs) depends on our ability to determine which algorithms work better, and under which conditions. However, such progress is thought to be...

<https://arxiv.org/abs/2609.18134>

02 CONF

Beyond Accuracy: How Procedural Traces Shift the Decision Criterion of LLM Overseers

ArXiv CS.AI 50%

arXiv:2609.18204v1 Announce Type: cross Abstract: Organizations increasingly use oversight loops where one large language model (LLM) audits another's outputs alongside procedural traces of claimed steps. A common concern about such LLM-as-a-judge pipelines is that detailed...

<https://arxiv.org/abs/2609.18204>

03 CONF

A Lightweight CNN Integrated Compact Convolutional Transformer for Multi-Scale Feature Learning and reducing computational complexity for breast cancer mammography image detection and classification

ArXiv CS.AI 50%

arXiv:2609.18212v1 Announce Type: cross Abstract: Over the years, Convolutional Neural Networks (CNNs) have demonstrated strong capability in cancer detection and classification using medical images. However, CNN-based models often struggle to capture long-range contextual...

<https://arxiv.org/abs/2609.18212>

04 CONF

Quanta: A Self-Contained Python Library for Hybrid Retrieval over Quantised Embeddings, Lexical Indexes, and Knowledge Graphs

ArXiv CS.AI 50%

arXiv:2609.18248v1 Announce Type: cross Abstract: An advanced retrieval-augmented generation pipeline is typically assembled from three or four independently operated systems: an approximate nearest-neighbour index, a full-text search engine, a graph database, and a relational...

<https://arxiv.org/abs/2609.18248>

05 CONF

A Study of the Reliability of Agentic AI-Generated Programs

ArXiv CS.AI 50%

arXiv:2609.18298v1 Announce Type: cross Abstract: Agentic-AI based software development offers the promise of faster completion of the software, greater programmer efficiency, and more reliable code. The question is how can we verify these claims in an objective way? In this...

<https://arxiv.org/abs/2609.18298>

06 CONF

Semantic CSI Feedback for Beam Selection: When Task-Aware Embeddings from Sparse Pilots Outperform Full-Bandwidth Reconstruction

ArXiv CS.AI 50%

arXiv:2609.18368v1 Announce Type: cross Abstract: Classical CSI feedback in FDD massive MIMO transmits a compressed reconstruction of the channel, optimizing fidelity to the original signal regardless of the downstream task. We propose a semantic communication perspective:...

<https://arxiv.org/abs/2609.18368>

07 CONF

Multitask Reinforcement Learning for Assisting Choice Model Specification

ArXiv CS.AI 50%

arXiv:2609.18441v1 Announce Type: cross Abstract: Discrete choice model specification is a time-consuming task in which modellers often specify and estimate multiple models while balancing goodness-of-fit, parsimony, and behavioural plausibility. We present Delphos, a multitask...

<https://arxiv.org/abs/2609.18441>

08 CONF

MiST: Mid-Training LLMs for Cybersecurity

ArXiv CS.AI 50%

arXiv:2609.18496v1 Announce Type: cross Abstract: Cybersecurity combines high-stakes analysis with complex technical language, making it an impactful and challenging domain for LLMs. We present MiST (Mid-trained Security Transformer), a suite of 8B and 32B models that achieve...

<https://arxiv.org/abs/2609.18496>

09 CONF

VoiceTrace: A Benchmark and Retrieval Framework for Who-Said-What Speech Retrieval

ArXiv CS.AI 50%

arXiv:2609.18521v1 Announce Type: cross Abstract: Speech retrieval has become increasingly important as spoken content continues to grow across meetings, lectures, podcasts, and videos. Existing benchmarks and models have advanced semantic search over spoken content, but largely...

<https://arxiv.org/abs/2609.18521>

10 CONF

Interpretable Patch-Based Deep Learning for Wildfire Spread Prediction from Ensemble Simulations

ArXiv CS.AI 50%

arXiv:2609.18555v1 Announce Type: cross Abstract: Wildfire spread is traditionally predicted using physics-based simulators, which are physically interpretable but whose cost increases with each additional ensemble member. We ask how well deep learning surrogates can reproduce...

<https://arxiv.org/abs/2609.18555>

11 CONF

RLLBC-Lib: An Educational Code Library for Reinforcement Learning and Learning-Based Control

ArXiv CS.AI 50%

arXiv:2609.19074v1 Announce Type: cross Abstract: Reinforcement learning (RL) is an exciting concept as well as a remarkable success story worth sharing. However, RL builds on rather complex interactions between different objects that play out over several cycles. Such dynamics...

<https://arxiv.org/abs/2609.19074>

12 CONF

Hypothesis-Driven Autonomous Materials Synthesis with Multimodal LLM Agents

ArXiv CS.AI 50%

arXiv:2609.18598v1 Announce Type: cross Abstract: Self-driving laboratories can explore synthesis conditions autonomously, but their decision-making layer is typically a black-box optimizer, and the output is a set of optimized samples, with the measurements reduced to...

<https://arxiv.org/abs/2609.18598>

13 CONF

PACT: Can Enterprise AI Assistants Be Trusted Under Pressure?

ArXiv CS.AI 50%

arXiv:2609.18605v1 Announce Type: cross Abstract: As corporate AI adoption continues to grow, enterprise-grade LLM agents are being deployed into sensitive contexts such as hiring, healthcare, and finance. In these contexts, compliance with rules specified in an agent's system...

<https://arxiv.org/abs/2609.18605>

14 CONF

Generalist-Specialist Mixture-of-Experts for Rare Pathology Detection in Multimodal Imaging

ArXiv CS.AI 50%

arXiv:2609.18688v1 Announce Type: cross Abstract: AI models for multimodal medical imaging must balance modality-specific specialization with cross-modal shared representations, a trade-off that pure Mixture-of-Experts (MoE) architectures currently fail to satisfy. Expert-based...

<https://arxiv.org/abs/2609.18688>

15 CONF

GrainSpeech: Less Context, More Detail for Compact Speech Synthesis

ArXiv CS.AI 50%

arXiv:2609.18856v1 Announce Type: cross Abstract: Compact acoustic models face a challenging quality-capacity trade-off. We investigate two factors in this regime: encoder context and Mel-spectrogram supervision. A receptive-field-scaling study shows that expanding...

<https://arxiv.org/abs/2609.18856>

16 CONF

Taming the Agentic RAN: Stability-Guaranteed Arbitration of Autonomous AI Agents in O-RAN

ArXiv CS.AI 50%

arXiv:2609.18857v1 Announce Type: cross Abstract: The O-RAN control plane is becoming agentic: autonomous AI agents, deployed as rApps by different vendors, independently close control loops over shared radio resources. We demonstrate on a live O-RAN system that this...

<https://arxiv.org/abs/2609.18857>

17 CONF

Decodable but Misrouted: Sparse Features Uncover a Readout Gap in Vision-Language Models for Harmful Meme Detection

ArXiv CS.AI 50%

arXiv:2609.18860v1 Announce Type: cross Abstract: When a large vision-language model misclassifies a harmful meme, the failure may reflect missing internal evidence or an inability to route represented evidence to its output. We distinguish these cases in Gemma-3 and Qwen3.5...

<https://arxiv.org/abs/2609.18860>

18 CONF

Dose-Aware Cold Diffusion with Physics Consistency for Generalizable Low-Dose CT Reconstruction

ArXiv CS.AI 50%

arXiv:2609.18943v1 Announce Type: cross Abstract: Reducing radiation dose in computed tomography significantly degrades image quality and poses challenges for accurate and clinically reliable reconstruction. While recent approaches have shown promise for low-dose CT, they often...

<https://arxiv.org/abs/2609.18943>

19 **CONF****Automated Dental Caries Segmentation in Panoramic Radiographs Using Dual-Stage Deep Learning****ArXiv CS.AI** **50%**

arXiv:2609.18952v1 Announce Type: cross Abstract: Early detection of dental caries remains challenging due to limitations in traditional diagnostic methods, particularly for proximal lesions in posterior teeth. Deep learning models show promise for automated caries detection but...

<https://arxiv.org/abs/2609.18952>

20 **CONF****Transcribe, Then Reason: Two-Pass Decomposition for Multimodal Review****ArXiv CS.AI** **50%**

arXiv:2609.18958v1 Announce Type: cross Abstract: The natural way to review a long recording or document with a multimodal model is to hand it the raw source and ask for a review in one call. We show that this quietly fails: the model satisfies, dropping roughly a third of the...

<https://arxiv.org/abs/2609.18958>

21 **CONF****WordPolo: Evaluating Language Models Through Iterative Semantic Feedback****ArXiv CS.AI** **50%**

arXiv:2609.19006v1 Announce Type: cross Abstract: Large Language Models (LLMs) and Large Reasoning Models (LRMs) are typically evaluated on challenging benchmarks through dataset accuracy alone, providing no insight into the quality or faithfulness of their reasoning processes....

<https://arxiv.org/abs/2609.19006>

22 **CONF****Securing quantum error correction against misleading advice from AI agents****ArXiv CS.AI** **50%**

arXiv:2609.19090v1 Announce Type: cross Abstract: Can an attacker turn influence over an artificial intelligence (AI) adviser into a harmful quantum error-correction update? We identify an ambiguity in passive syndrome records that obstructs recovery selection, then show how...

<https://arxiv.org/abs/2609.19090>

23 **CONF****Prepared Or Unprepared? Evaluating Healthcare Workforce Readiness for Clinical Adoption of Artificial Intelligence in Nigeria****ArXiv CS.AI** **50%**

arXiv:2609.19096v1 Announce Type: cross Abstract: Artificial intelligence (AI) is increasingly integrated into healthcare systems worldwide, yet its successful clinical adoption depends critically on workforce readiness, particularly in low- and middle-income countries (LMICs)...

<https://arxiv.org/abs/2609.19096>

24 **CONF****Creating an Atomic User Model for Personality-Aware Large Language Model Interaction****ArXiv CS.AI** **50%**

arXiv:2609.12086v2 Announce Type: replace-cross Abstract: Assistants built on large language models are expected to write in their users' own voice. Most systems summarise the user's preferences and include the summary in the prompt. This is the wrong way round. Preferences are...

<https://arxiv.org/abs/2609.12086>

25 CONF

Dreaming the Sound of Contact: Leveraging Video and Audio Generation for Zero-Shot Force-Aware Manipulation and Data Generation

ArXiv CS.AI 50%

arXiv:2609.19137v1 Announce Type: cross Abstract: Recent advances in video generation allow robots to learn manipulation trajectories from generated videos. However, these approaches produce purely kinematic trajectories that lack force information, causing failures in...

<https://arxiv.org/abs/2609.19137>

26 CONF

A Zeroth-Order Paradigm for LLM Preference Alignment

ArXiv CS.AI 50%

arXiv:2609.19144v1 Announce Type: cross Abstract: Direct preference alignment methods are widely used to align large language models (LLMs) with human preferences because of their computational and memory efficiency. However, likelihood displacement motivates alternative ways to...

<https://arxiv.org/abs/2609.19144>

27 CONF

A Survey on Bridging EEG Signals and Generative AI: From Image and Text to Beyond

ArXiv CS.AI 50%

arXiv:2502.12048v4 Announce Type: replace Abstract: Decoding neural activity into human-interpretable representations is a key research direction in brain-computer interfaces (BCIs) and computational neuroscience. Recent progress in machine learning and generative AI has driven...

<https://arxiv.org/abs/2502.12048>

28 CONF

SurgRAW: Multi-Agent Workflow with Chain of Thought Reasoning for Robotic Surgical Video Analysis

ArXiv CS.AI 50%

arXiv:2503.10265v3 Announce Type: replace Abstract: Robotic-assisted surgery (RAS) is central to modern surgery, driving the need for intelligent systems with accurate scene understanding. Most existing surgical AI methods rely on isolated, task-specific models, leading to...

<https://arxiv.org/abs/2503.10265>

29 CONF

MCPAgentBench: A Real-world Task Benchmark for Evaluating LLM Agent MCP Tool Use

ArXiv CS.AI 50%

arXiv:2512.24565v5 Announce Type: replace Abstract: Large Language Models (LLMs) are increasingly serving as autonomous agents, and their utilization of external tools via the Model Context Protocol (MCP) is considered a future trend. Current MCP evaluation sets suffer from...

<https://arxiv.org/abs/2512.24565>

30 CONF

Exploratory Responsiveness and Adaptive Rigidity under AI-Assisted Optimization

ArXiv CS.AI 50%

arXiv:2606.10086v2 Announce Type: replace Abstract: This paper develops a theory of exploratory adaptation under AI-assisted optimization. The central argument is that the long-run adaptive effects of AI systems depend critically on how predictive assistance interacts with...

<https://arxiv.org/abs/2606.10086>

31 CONF

HyQuant: Hybrid-Precision Quantization for LLM Attention

ArXiv CS.AI 50%

arXiv:2608.27875v3 Announce Type: replace Abstract: Quantization has been widely adopted in LLM training and inference to reduce cost and improve efficiency. However, low-bit quantization of the \emph{attention} module often introduces large errors at very low bit-widths,...

<https://arxiv.org/abs/2608.27875>

32 CONF

Unleash LLMs Potential for Sequential Recommendation by Coordinating Dual Dynamic Index Mechanism

ArXiv CS.AI 50%

arXiv:2409.09253v2 Announce Type: replace-cross Abstract: Owing to the unprecedented capability in semantic understanding and logical reasoning, large language models (LLMs) have shown fantastic potential in developing next-generation sequential recommender systems (RSs)....

<https://arxiv.org/abs/2409.09253>

33 CONF

From 'May' to 'Is': Certainty Distortion in Language Model Rewriting

ArXiv CS.AI 50%

arXiv:2606.07951v2 Announce Type: replace-cross Abstract: Humans increasingly turn to Language Models (LMs) in ways that shape beliefs and drive decisions, including discussing, rewriting, and summarizing information from scientific articles, news, and medical reports. However,...

<https://arxiv.org/abs/2606.07951>

34 CONF

BOOM: Benchmarking Out-Of-distribution Molecular Property Predictions of Machine Learning Models

ArXiv CS.AI 50%

arXiv:2505.01912v3 Announce Type: replace-cross Abstract: Data-driven molecular discovery leverages artificial intelligence/machine learning (AI/ML) and generative modeling to filter and design novel molecules. Discovering novel molecules requires accurate out-of-distribution...

<https://arxiv.org/abs/2505.01912>

35 CONF

Extracting Probabilistic Knowledge from Large Language Models for Bayesian Network Parameterization

ArXiv CS.AI 50%

arXiv:2505.15918v3 Announce Type: replace-cross Abstract: In this work, we evaluate the potential of Large Language Models (LLMs) in building Bayesian Networks (BNs) by approximating domain expert priors. LLMs have demonstrated potential as factual knowledge bases; however,...

<https://arxiv.org/abs/2505.15918>

36 CONF

Diff-SPORT: Diffusion-based Sensor Placement Optimization and Reconstruction of Turbulent flows in urban environments

ArXiv CS.AI 50%

arXiv:2506.00214v2 Announce Type: replace-cross Abstract: Rapid urbanization demands efficient monitoring of turbulent wind and pollutant dispersion, yet existing reconstruction and sensor placement strategies fail under realistic sparsity constraints. Here, we introduce...

<https://arxiv.org/abs/2506.00214>

37 CONF

From Alignment to Synthesis: Contrastive Volumetric Grounding for Text-to-CT Generation

ArXiv CS.AI 50%

arXiv:2506.00633v4 Announce Type: replace-cross Abstract: Generating semantically controllable 3D CT volumes from radiology reports requires more than a rich text encoder, it requires vision-language alignment grounded in volumetric space. Existing Text-to-CT approaches...

<https://arxiv.org/abs/2506.00633>

38 CONF

Constrained PSLQ Search for Machin-like Identities Achieving Record-Low Lehmer Measures

ArXiv CS.AI 50%

arXiv:2508.08307v2 Announce Type: replace-cross Abstract: Machin-like arctangent relations are classical tools for computing π , with efficiency quantified by the Lehmer measure (λ). We present a framework for discovering low-measure relations by coupling the PSLQ...

<https://arxiv.org/abs/2508.08307>

39 CONF

Visual Perception Engine: Fast and Flexible Multi-Head Inference for Robotic Vision Tasks

ArXiv CS.AI 50%

arXiv:2508.11584v3 Announce Type: replace-cross Abstract: Deploying multiple machine learning models on resource-constrained robotic platforms for different perception tasks often results in redundant computations, large memory footprints, and complex integration challenges. In...

<https://arxiv.org/abs/2508.11584>

40 CONF

Ultralytics YOLO Evolution: An Overview of YOLO27, YOLO26, YOLO11, YOLOv8, and YOLOv5 Object Detectors for Computer Vision and Pattern Recognition

ArXiv CS.AI 50%

arXiv:2510.09653v4 Announce Type: replace-cross Abstract: This paper presents a comprehensive overview of the Ultralytics YOLO family, emphasizing architectural evolution, benchmarking, deployment, and emerging directions from YOLOv5 through YOLO27. The review begins with YOLO27...

<https://arxiv.org/abs/2510.09653>

41 CONF

Human Resilience in the AI Era -- What Machines Can't Replace

ArXiv CS.AI 50%

arXiv:2510.25218v2 Announce Type: replace-cross Abstract: AI is changing work and decision making faster than many institutions can adapt their operating practices. We argue that this adaptation gap makes human resilience a core capability for the AI era. We define resilience as...

<https://arxiv.org/abs/2510.25218>

42 CONF

Performance and Complexity Trade-off Optimization of Speech Models During Training

ArXiv CS.AI 50%

arXiv:2601.13704v4 Announce Type: replace-cross Abstract: In speech machine learning, neural network models are typically designed by choosing an architecture with fixed layer sizes and structure. These models are then trained to maximize performance on metrics aligned with the...

<https://arxiv.org/abs/2601.13704>

43 CONF

ShotFinder: Imagination-Driven Open-Domain Video Shot Retrieval via Web Search

ArXiv CS.AI 50%

arXiv:2601.23232v4 Announce Type: replace-cross Abstract: In recent years, large language models (LLMs) have made rapid progress in information retrieval, yet existing research has mainly focused on text or static multimodal settings. Open-domain video shot retrieval, which...

<https://arxiv.org/abs/2601.23232>

44 CONF

AMIGO: Agentic Multi-Image Grounding Oracle Benchmark

ArXiv CS.AI 50%

arXiv:2603.28662v2 Announce Type: replace-cross Abstract: Agentic vision-language models increasingly act through extended interactions, but most evaluations still focus on single-image, single-turn correctness. We introduce $\text{\textbf{AMIGO}}$ ($\text{\textbf{A}}\text{\textbf{g}}\text{\textbf{e}}\text{\textbf{n}}\text{\textbf{t}}\text{\textbf{i}}\text{\textbf{c}}$...

<https://arxiv.org/abs/2603.28662>

45 CONF

Schema-Key Wording as an Instruction Channel in Structured Generation under Constrained Decoding

ArXiv CS.AI 50%

arXiv:2604.14862v3 Announce Type: replace-cross Abstract: Constrained decoding is widely used to make large language models produce structured outputs that satisfy schemas such as JSON. Existing work mainly treats schemas as structural constraints, overlooking that schema-key...

<https://arxiv.org/abs/2604.14862>

46 CONF

FLARE: Fine-Grained Diagnostic Feedback for LLM Code Refinement

ArXiv CS.AI 50%

arXiv:2606.03852v2 Announce Type: replace-cross Abstract: Large language models often generate code with bugs. Existing methods rely on feedback signals such as test failures and self-critiques to iteratively refine the generated code. Such signals are either too coarse-grained...

<https://arxiv.org/abs/2606.03852>

47 CONF

SegTME-UNI2: A Foundation Model-Based Framework for Generalisable Multiclass Cell Segmentation and LLM-Driven Tumour Microenvironment Characterisation in Histopathology

ArXiv CS.AI 50%

arXiv:2606.17702v3 Announce Type: replace-cross Abstract: Characterising the TME from routine H&E-stained histology images requires simultaneous cell segmentation, biological feature extraction, and interpretable clinical reporting. We present SegTME-UNI2, a unified framework...

<https://arxiv.org/abs/2606.17702>

48 CONF

Subjective Risk Decomposition: A New View for Uncertainty Quantification

ArXiv CS.AI 50%

arXiv:2607.15196v3 Announce Type: replace-cross Abstract: We present a novel viewpoint for uncertainty quantification. Uncertainty measures are not primitives, in need of axioms and argumentation, but instead consequences, of higher-level modelling decisions. We show how...

<https://arxiv.org/abs/2607.15196>

49 **CONF****Ranking Infrared-Visible Fusion the Way Humans Do: A Learned Pairwise Preference Measure****ArXiv CS.AI 50%**

arXiv:2608.01301v4 Announce Type: replace-cross Abstract: Human pairwise comparison provides a direct basis for perceptual infrared-visible image fusion assessment, but dense annotation becomes costly as method pools grow. We present the Learned Perceptual Image Fusion Measure...

<https://arxiv.org/abs/2608.01301>50 **CONF****Exploring the Potential of Contrastive Language-Image Pre-training for Multi-Source Remote Sensing Data****ArXiv CS.AI 50%**

arXiv:2609.03391v2 Announce Type: replace-cross Abstract: Contrastive language-image learning (CLIP) has become a key paradigm for remote sensing vision-language understanding. However, existing remote sensing contrastive learning methods are mostly built on RGB-oriented CLIP...

<https://arxiv.org/abs/2609.03391>51 **CONF****Language-Guided Terrain-Adaptive Neural MPC for Autonomous Traversal of Articulated Tracked Robots****ArXiv CS.AI 50%**

arXiv:2609.13083v3 Announce Type: replace-cross Abstract: In urban search and rescue, articulated tracked robots (ATRs) must traverse structured but contact-rich environments such as stairwells and cluttered building interiors. Reliable autonomy remains challenging because...

<https://arxiv.org/abs/2609.13083>