



Pipeline

Scrape HackerNews (Show/Top/New), Reddit (r/artificial, r/MachineLearning, r/singularity, r/LocalLLaMA), TechCrunch AI, VentureBeat AI, MIT Tech Review, The Verge AI, Wired AI, ArXiv CS.AI

Filter 40+ AI keyword patterns — LLMs, alignment, safety, research, tooling

Dedup Title-similarity (45% threshold) merges cross-source variants

Score Authenticity = 0.5 base + 0.3 per extra source + diversity + HN score bonuses

Rank Sorted by authenticity. Below 0.25 threshold are dropped.

Authenticity Tiers

- VERIFIED 0.80 - 1.00** 3+ independent sources including media
- CONFIRMED 0.50 - 0.79** 2 sources or 1 high-quality media
- EMERGING 0.25 - 0.49** Single-source; passes AI keyword filter

Source Breakdown

Source	Articles	Category
ArXiv CS.AI	114	Media
TechCrunch AI	14	Media
HackerNews	9	Community
The Verge AI	5	Media
MIT Tech Review	1	Media

VERIFIED INTELLIGENCE — 2 stories

VERI

Show HN: AutoBot - live voice control for long-running AI work

HackerNews

+1 source

90%

HN 18

arXiv:2609.20347v1 Announce Type: cross Abstract: LEO satellite networks feature dynamic topologies, time-varying links, and diverse service requirements, which make conventional routing schemes difficult to support fine-grained quality-of-service (QoS) provisioning. Existing...

<https://github.com/demeyer1/Autobot>

VERI

Microsoft AI CEO says AI threats are real, and Anthropic is making it worse

The Verge AI

+1 source

80%

Today, I'm talking with Mustafa Suleyman, the CEO of Microsoft AI. As you're no doubt aware, the biggest story in tech right now is the spiraling debate about AI safety and regulation. It should come as no surprise that Mustafa has strong opinions on how AI should be built and...

<https://www.theverge.com/podcast/996412/microsoft-ai-ceo-mustafa-suleyman-regulation-sa...>

CONFIRMED SIGNALS — 139 stories

CONF

Show HN: Share your AI Setup, Learn from others

HackerNews

60%

HN 211

I kept seeing engineers share what they were building with AI; however, I was always more curious about how they worked. Which agents did they use? What skills and tools had stuck or been thrown out the window? How did they manage longer-running tasks? So I built this with the...

<https://mysetup.ai/>

CONF

Show HN: Trading on-chain with 0 gas: A <1ms DEX router at 0.50% fees

HackerNews

50%

HN 9

No summary — see link for full article.

<https://rhopool.com/>

CONF

Show HN: Continuity - project state your AI sessions can't silently contradict

HackerNews

50%

HN 3

No summary — see link for full article.

<https://github.com/vikcena01/ai-continuity-plugin>

CONF

Show HN: Composing domain-specific harness on Python in 10 mins

HackerNews

50%

HN 3

No summary — see link for full article.

<https://cayu.dev/walkthrough/>

CONF

Show HN: Vim-like, hyper-efficient, LLM power tool - 20-80k tokens not 350k+

HackerNews

50%

HN 2

I built a focused TUI tool for efficient parallelized tasks, early independent testers expressed 90%+ API savings and more direct results. This is accomplished by user smaller context windows, hybrid client-side compaction, tool compaction, and system nudges (to avoid polluting...

<https://easiest.ai/>

CONF

Show HN: Linting 216 public Claude Code skills - 69% won't reliably trigger

HackerNews 50% HN 2

I built a linter for Claude Code artifacts (skills, subagents, slash commands, .mcp.json configs, and plugins) and then pointed it at a pile of public repos to see what people actually ship. A few of the numbers surprised me enough that I wrote them up. The one that stuck:...

<https://skillcrossroads.com>

CONF

Show HN: Graph RAG in Postgres. New facts replace older facts

HackerNews 50% HN 2

No summary — see link for full article.

<https://github.com/craiah/post-graph-rag>

0

CONF

Show HN: Smart Mouth Billy Bass - another fishy LLM assistant

HackerNews 50% HN 2

Hello, HN people. This project was inspired by another very similar Show HN post from a few months ago called BillAI Bass. I decided to do something like it but take it in a different direction. It's not a serious assistant and doesn't turn my lights on and off, but it has no...

<https://existens.org/billy/>

1

CONF

Crusoe raises \$3.9B to build massive data centers and small modular 'AI factories'

TechCrunch AI 50%

The round values the data center giant at \$30.9 billion.

<https://techcrunch.com/2026/09/17/crusoe-raises-3-9b-to-build-massive-data-centers-and-...>

2

CONF

Google DeepMind launches institute to widen the AGI debate

TechCrunch AI 50%

The new institute aims to surface differing views between Google, Google DeepMind, and the broader global research community around AGI. "They will not always agree, and they will likely change their minds, as more data and information comes to light at the fast-moving frontier."

<https://techcrunch.com/2026/09/17/google-deepmind-launches-institute-to-widen-the-agi-d...>

3

CONF

PrismML hopes its tiny LLM will change how we all use AI

TechCrunch AI 50%

If AI lab PrismML isn't on your radar yet, it should be.

<https://techcrunch.com/2026/09/17/prismml-hopes-its-tiny-llm-could-change-how-we-all-us...>

4

CONF

The FAA's plan to fix air traffic? \$875M worth of AI

TechCrunch AI 50%

A new AI-based software program is being launched to help air traffic controllers better navigate their jobs as the crossing guards of America's skies.

[https://techcrunch.com/2026/09/17/the-faas-plan-to-fix-air-traffic-875-million-worth-of...](https://techcrunch.com/2026/09/17/the-faas-plan-to-fix-air-traffic-875-million-worth-of-...)

5

CONF

The fix for rogue AI agents could be more AI

TechCrunch AI 50%

As companies hand off longer and more complex tasks to AI agents, they are running into an oversight problem: Agents can act faster, longer, and at greater volume than humans can realistically review.

<https://techcrunch.com/2026/09/17/the-fix-for-rogue-ai-agents-could-be-more-ai/>

6 **CONF****OpenAI caught its models leaving notes to successors to hide bad behavior****TechCrunch AI** 50%

OpenAI disclosed instances of GPT-5.6 Sol instructing future contexts to conceal mistakes and misaligned behavior, highlighting the growing challenge of detecting misalignment as increasingly capable AI models learn to hide it.

<https://techcrunch.com/2026/09/17/openai-caught-its-models-leaving-notes-to-successors-...>

7 **CONF****Is the AI safety debate about safety or control?****TechCrunch AI** 50%

Not everyone agrees with Amodei's call for globally coordinated action for AI safety.

<https://techcrunch.com/2026/09/17/is-the-ai-safety-debate-about-safety-or-control/>

8 **CONF****UN turns to Google to make its global data ready for AI agents****TechCrunch AI** 50%

The shift comes after a UNICEF test found leading AI models struggled to accurately retrieve global development statistics.

<https://techcrunch.com/2026/09/17/un-turns-to-google-to-make-its-global-data-ready-for-...>

9 **CONF****Microsoft exec called AI scraping 'the largest theft of labor in human history,' new unredacted filings reveal****TechCrunch AI** 50%

Newly unsealed court filings show Microsoft privately called OpenAI's data practices "theft" while both companies scraped paywalled Times content, built datasets from it, and warned internally it would gut publishers.

<https://techcrunch.com/2026/09/17/microsoft-exec-called-ai-scraping-the-largest-theft-o-...>

0 **CONF****Base Labs launches an open-weight AI safety partnership with Hugging Face and Goodfire****TechCrunch AI** 50%

Base Labs, the research group Baseten spun up earlier this year, will develop and publish methods for training and monitoring open models.

<https://techcrunch.com/2026/09/17/base-labs-launches-an-open-weight-ai-safety-partnersh...>

1 **CONF****Pinterest teases a new 'Restyle' feature that lets you redesign your room with AI****TechCrunch AI** 50%

Pinterest is testing Restyle, a new AI-powered feature that lets users visualize furniture, decor, lighting, and more in photos of their own rooms — potentially helping turn saved inspiration into purchases.

<https://techcrunch.com/2026/09/17/pinterest-teases-a-new-restyle-feature-that-lets-you-...>

2 **CONF****Huawei plans Q1 2027 launch of new AI chip as it takes on Nvidia****TechCrunch AI** 50%

Huawei is accelerating the launch of its next-generation Ascend 960DT AI chip as it pushes to compete with Nvidia and close China's AI computing gap with the U.S.

<https://techcrunch.com/2026/09/17/huawei-plans-q1-2027-launch-of-new-ai-chip-as-it-take...>

3 **CONF****Rival AI agents, Instinct and Meta's Muse, both add the ability to make calls****TechCrunch AI** 50%

People can use these assistants to make restaurant reservations and cancel subscriptions.

<https://techcrunch.com/2026/09/17/rival-ai-agents-instinct-and-metas-muse-both-add-the-...>

4 **CONF****The Download: mice with part-human brains and climate tech innovators****MIT Tech Review** 50%

This is today's edition of The Download, our weekday newsletter that provides a daily dose of what's going on in the world of technology. Meet a mouse whose brain cortex is made up of human cells Multiple cameras tracked a mouse as it wandered around a small arena. A computer...

<https://www.technologyreview.com/2026/09/17/1144314/the-download-mice-part-human-brains...>

5 **CONF****The AI Superintelligence Slowdown****The Verge AI** 50%

Remember when tech leaders would tell their employees to "move fast and break things"? It seemed that would be the way of AI too. But after a summer where rogue AI agents became reality, and researchers warned that AI could kill us all, a number of leading US AI companies are...

<https://www.theverge.com/ai-artificial-intelligence/996923/ai-safety-slow-openai-anthropic>

6 **CONF****Claude Code relaunches Projects to manage multiple AI agents in the cloud****The Verge AI** 50%

The revamped projects feature in Claude Code allows users to run multiple agents under the same roof, with a shared memory, goals, and library of files and artifacts. Similar to Grok Bot and other tools that manage groups of AI agents, each project has "threads" running...

<https://www.theverge.com/ai-artificial-intelligence/997134/anthropic-claude-code-projects>

7 **CONF****AI is feared globally as the destroyer of jobs****The Verge AI** 50%

Pew Research has published a new global survey that sheds light on how people view AI, including its impact on jobs, life in general, and income inequality. The survey questioned 42,151 people across 37 countries from February 8th to May 13th - well ahead of recent apocalyptic...

<https://www.theverge.com/ai-artificial-intelligence/996775/ai-is-feared-globally-as-the...>

8 **CONF****Inside the suddenly explosive world of AI safety****The Verge AI** 50%

On a sunny July day in Berkeley, California, the country's top AI safety researchers gathered on an unmarked floor of an unmarked building. They had come together for a "war room" to dissect the high-profile cybersecurity incident that had rocked the AI industry hours earlier....

<https://www.theverge.com/ai-artificial-intelligence/996563/ai-safety-research-metr-redw...>

9 **CONF****BioPhys-Bridge: A Benchmark for Interdisciplinary Scientific Reasoning in Physics-Grounded Biological Research****ArXiv CS.AI** 50%

arXiv:2609.19180v1 Announce Type: new Abstract: Language models face unique challenges in analyzing interdisciplinary scientific research literature. In biophysics research, faithful answers require grounding observed data in source evidence, interpreting it through a...

<https://arxiv.org/abs/2609.19180>

0 **CONF****What Do We Expect from LLMs? Mapping the Design of LLM Benchmarks**

ArXiv CS.AI 50%

arXiv:2609.19182v1 Announce Type: new Abstract: Benchmarks are central to how progress in large language models (LLMs) is assessed and communicated. Yet model rankings alone reveal little about how evaluation requirements themselves are changing. The expanding variety of...

<https://arxiv.org/abs/2609.19182>
1 **CONF****Position: It is Time to Virtualize Foundation Models with a Self-evolving Operating System Layer**

ArXiv CS.AI 50%

arXiv:2609.19203v1 Announce Type: new Abstract: AI applications have shifted from single, monolithic foundation models (FM) to compound agentic systems. Yet today's stacks remain fragmented: even as protocols (e.g., MCP, A2A) ease tool/agent connectivity, each framework embeds...

<https://arxiv.org/abs/2609.19203>
2 **CONF****Characterizing Web Search by Conversational LLM Agents: From Search Decisions and Strategies to Results and Responses**

ArXiv CS.AI 50%

arXiv:2609.19244v1 Announce Type: new Abstract: Conversational LLM agents increasingly rely on Web search, yet the end-to-end lifecycle of agentic search remains poorly understood. We present the first study of Web search across four major conversational platforms (ChatGPT,...

<https://arxiv.org/abs/2609.19244>
3 **CONF****Do AI Agents Understand Computer Architecture?**

ArXiv CS.AI 50%

arXiv:2609.19387v1 Announce Type: new Abstract: Agents are increasingly asked to design hardware, and increasingly reported to succeed. Such reports establish that a design improved; they cannot establish why. An agent that improves an accelerator may be reasoning about the...

<https://arxiv.org/abs/2609.19387>
4 **CONF****EPIG-Tree: Compute-Optimal Branching for Gradient-Efficient Reinforcement Learning**

ArXiv CS.AI 50%

arXiv:2609.20004v1 Announce Type: cross Abstract: Reward-based reinforcement learning for language models, exemplified by Group Relative Policy Optimization (GRPO), collapses an entire stochastic trajectory into a single scalar reward. This is clean and scalable, but it explores...

<https://arxiv.org/abs/2609.20004>
5 **CONF****Geopolitical Divisions Across Languages in Large Language Models**

ArXiv CS.AI 50%

arXiv:2609.20005v1 Announce Type: new Abstract: People increasingly turn to AI chatbots for news and explanations of world events. But do they receive the same political answers when they ask in different languages? Here we show that the language of a question can change how the...

<https://arxiv.org/abs/2609.20005>

6 CONF

QVAC Genesis III: A Large-Scale, High-Quality Open Synthetic STEM Corpus for Efficient Language Model Pre-Training

ArXiv CS.AI 50%

arXiv:2609.19513v1 Announce Type: new Abstract: High-quality pre-training data is a critical bottleneck for educational and STEM-specific language models targeting edge AI and on-device deployment where token budgets are tightly constrained. While major organizations train...

<https://arxiv.org/abs/2609.19513>

7 CONF

LLM-as-an-Improver: Turning Verification into Better Candidates

ArXiv CS.AI 50%

arXiv:2609.19515v1 Announce Type: new Abstract: Verifier-based selection improves LLM performance by generating multiple candidate solutions and using a verifier to select the most promising one. However, existing methods typically treat verification only as a ranking step and...

<https://arxiv.org/abs/2609.19515>

8 CONF

A Unified Evaluation Framework for Trustworthy Large Language Models, Agentic AI, and Multimodal Systems

ArXiv CS.AI 50%

arXiv:2609.19524v1 Announce Type: new Abstract: Benchmark scores alone provide an incomplete basis for assessing the trustworthiness of modern artificial intelligence systems. Large language models (LLMs), agentic systems, and multimodal models (MLLMs) require different forms of...

<https://arxiv.org/abs/2609.19524>

9 CONF

Agentic AI Networking for Heterogeneous Unmanned Aerial Systems in Low-Altitude Wireless Networks

ArXiv CS.AI 50%

arXiv:2609.19538v1 Announce Type: new Abstract: Low-altitude wireless networks (LAWNs) are emerging as a key infrastructure for heterogeneous unmanned aerial systems that support concurrent services within a shared three-dimensional airspace. Their coexistence creates strong...

<https://arxiv.org/abs/2609.19538>

10 CONF

From Intent to Action: Benchmarking LLM Safety in Vehicle Voice Command Authorization

ArXiv CS.AI 50%

arXiv:2609.19630v1 Announce Type: new Abstract: Large language models (LLMs) are increasingly integrated into vehicle voice assistants. But linking natural-language requests to vehicle functions creates a safety-critical authorization problem. Before executing a command, the...

<https://arxiv.org/abs/2609.19630>

11 CONF

Reach or Solve? Attributing Agentic RL Gains with Checkpoint Handoffs

ArXiv CS.AI 50%

arXiv:2609.19636v1 Announce Type: new Abstract: Reinforcement learning now trains language-model agents that act over dozens of steps in live environments. The gains are large, and they are read as better decision-making. An agent in a closed loop writes its own inputs. Each...

<https://arxiv.org/abs/2609.19636>

2 **CONF****ScientistTwo: Pioneering the Human Knowledge Frontier with Autonomous AI****ArXiv CS.AI 50%**

arXiv:2609.19644v1 Announce Type: new Abstract: Scientific discovery is defined by the ability to identify the boundaries of existing knowledge and venture into unexplored territory. The ultimate vision for AI in science is problem-driven autonomous research: given a fundamental...

<https://arxiv.org/abs/2609.19644>

3 **CONF****Replan, Repair, or Edit? A Unified Empirical Evaluation of Travel Agents for Itinerary Revision under Resource Disruptions****ArXiv CS.AI 50%**

arXiv:2609.19654v1 Announce Type: new Abstract: Travel-planning agents generate itineraries that may become infeasible after acceptance because of flight cancellations, hotel unavailability, or attraction closures. Revising these itineraries involves full replanning, classical...

<https://arxiv.org/abs/2609.19654>

4 **CONF****Large Language Model Agents for Evidence Based Genetic Disease Severity Classification****ArXiv CS.AI 50%**

arXiv:2609.19569v1 Announce Type: cross Abstract: Disease severity classification for genetic conditions is subjective and labor-intensive, creating bottlenecks in genomic screening, where commercial panels vary widely in size and overlap. We developed an autonomous AI agent...

<https://arxiv.org/abs/2609.19569>

5 **CONF****Integrating knowledge from case reports: a medical ontology based multimodal information system with structured summary****ArXiv CS.AI 50%**

arXiv:2609.19775v1 Announce Type: new Abstract: Published medical case reports serve as a crucial medical information carrier, documenting discoveries in rare diseases, diagnostic methods, and innovative treatments. Despite the wealth of clinical knowledge in millions of case...

<https://arxiv.org/abs/2609.19775>

6 **CONF****A Dual-Process Perspective on Nudge Susceptibility in LLM-Based GUI Agents****ArXiv CS.AI 50%**

arXiv:2609.19843v1 Announce Type: new Abstract: LLM-based GUI agents increasingly act on behalf of users in digital environments that were designed with human users in mind. These graphical user interfaces were designed to support, but also deliberately steer, the behaviour and...

<https://arxiv.org/abs/2609.19843>

7 **CONF****Constraint-Safe Graph-Context Scoring for Stable Point-Feature Labels Under Text-Width and Accessibility-Inspired Profiles****ArXiv CS.AI 50%**

arXiv:2609.19848v1 Announce Type: new Abstract: Point-feature label placement on interactive maps must reconcile geometric validity, display yield, local placement utility, and stability across camera motion. Accessibility and multilingual requirements further change label...

<https://arxiv.org/abs/2609.19848>

8 **CONF****Reproducibility is not construct validity: LLM measurement of institutionally situated communication**

ArXiv CS.AI 50%

arXiv:2609.19866v1 Announce Type: new Abstract: High annotation reproducibility does not necessarily imply that an LLM-inferred measure captures the construct it is intended to measure. We test this distinction using a dataset from the European Commission's AI Act consultation,...

<https://arxiv.org/abs/2609.19866>

9 **CONF****Physical knowledge on historical data matters more than enforcing physical constraints on the forecast**

ArXiv CS.AI 50%

arXiv:2609.19871v1 Announce Type: new Abstract: Time series forecasting has seen significant advancements with the emergence of new deep learning models. However, forecasting time series in applications involving physical processes remains a major challenge. Despite the apparition...

<https://arxiv.org/abs/2609.19871>

10 **CONF****Not All AI Agents Are Equal: Characterizing Resource and Performance Dynamics**

ArXiv CS.AI 50%

arXiv:2609.19947v1 Announce Type: new Abstract: LLM-based AI agents process user requests through iterative reasoning and tool execution, often involving the invocation of remote LLM APIs with local tool containers. This execution model can make the optimization of agent serving...

<https://arxiv.org/abs/2609.19947>

1 **CONF****E-AVI: Evidence-Grounded Multimodal Assessment for Automated Video Interviews**

ArXiv CS.AI 50%

arXiv:2609.20001v1 Announce Type: new Abstract: Automated video interview assessment integrates verbal content, acoustic delivery, and visual behavior, yet numerical predictions alone provide limited inspectable support. We present E-AVI, an evidence-grounded framework that...

<https://arxiv.org/abs/2609.20001>

2 **CONF****DART: Distillation-Aware Reparameterization for Training-Free LoRA Reuse in Few-Step Video Diffusion Models**

ArXiv CS.AI 50%

arXiv:2609.20051v1 Announce Type: new Abstract: Step distillation reduces the cost of video generation, but reusing a LoRA trained for a longer trajectory can alter its functional effect or degrade target quality. Static parameter compatibility offers one perspective on this...

<https://arxiv.org/abs/2609.20051>

3 **CONF****MAGMA-GEN: Validated Recovery Supervision from Ambiguous Failures via Counterfactual Re-Execution**

ArXiv CS.AI 50%

arXiv:2609.20056v1 Announce Type: new Abstract: Hierarchical robotic systems executing long-horizon manipulation tasks must make high-level semantic decisions that orchestrate stochastic low-level skills. In this setting, failed rollouts are ambiguous: a poor downstream state...

<https://arxiv.org/abs/2609.20056>

4 CONF

WiCleanData: Guaranteeing the Type Consistency of Wikidata by Taxonomy Refinement and Constraint Enforcement

ArXiv CS.AI 50%

arXiv:2609.20057v1 Announce Type: new Abstract: Because of its collaborative nature, Wikidata suffers from errors, inconsistencies, and excessive complexity, such as redundant classes, ambiguity between instances and classes, wrong taxonomic paths, and type constraint...

<https://arxiv.org/abs/2609.20057>

5 CONF

Marginal utility, matrix factorization, and the Key-Value (KV) cache: a unified information-economic framework for sovereign geo-mining inference

ArXiv CS.AI 50%

arXiv:2609.20068v1 Announce Type: new Abstract: This paper builds a theoretical bridge between the economic notion of marginal utility and two machine-learning constructs, matrix factorization and the Key-Value cache of transformer language models. The singular value spectrum...

<https://arxiv.org/abs/2609.20068>

6 CONF

CoMa: Contextual Massing Generation with Vision-Language Models

ArXiv CS.AI 50%

arXiv:2601.08464v2 Announce Type: replace-cross Abstract: Context-aware building massing is an important early-stage design task: given a site for buildings, a generated massing should not only fit the target parcel, but also relate to the scale, density, and morphology of its...

<https://arxiv.org/abs/2601.08464>

7 CONF

A Proposal for an Agentic AI Architecture to Support Multi-Domain Decision-Making in the Brazilian Armed Forces

ArXiv CS.AI 50%

arXiv:2609.20080v1 Announce Type: new Abstract: The growing complexity of multi-domain operational environments (land, aerospace, naval, cyber, and electromagnetic spectrum) has increased the volume and velocity of data reaching command-and-control (C2) centers, straining the...

<https://arxiv.org/abs/2609.20080>

8 CONF

Perception, Layout, and Validation: Calibrated Confidence for Reliable Straight-Through Processing of Financial Documents

ArXiv CS.AI 50%

arXiv:2609.20110v1 Announce Type: new Abstract: Straight-through processing (STP) on extracted key-value fields from financial documents without human review requires a calibrated probability together with a bounded guarantee on the residual error of the auto-approved tier. The...

<https://arxiv.org/abs/2609.20110>

9 CONF

PaGNet: A Panel-Aware GBDT--Neural Network for Multi-Target Corporate Tax Avoidance Proxy Forecasting

ArXiv CS.AI 50%

arXiv:2609.20177v1 Announce Type: new Abstract: Forecasting corporate tax avoidance proxies from firm--year panel data is challenging because predictive signals are distributed across short firm histories and related targets, while screening-oriented use requires transparent...

<https://arxiv.org/abs/2609.20177>

0 **CONF****Sequential Contextual Fit Predicts Human Behavioural and Neural Dynamics Across Domains****ArXiv CS.AI** **50%**

arXiv:2609.20179v1 Announce Type: new Abstract: Human perception, action and decision making unfold in sequences, but computational predictors are often domain-specific. This study computes and tests sequential contextual fit (SCF), an embedding-based measure of how well a...

<https://arxiv.org/abs/2609.20179>

1 **CONF****When AI Agents Commit: Cognitive Serializability Across Data, Evidence, Policy, and Authority****ArXiv CS.AI** **50%**

arXiv:2609.20261v1 Announce Type: new Abstract: Autonomous agents derive concrete mutations from database reads, retrieved evidence, policy, beliefs, and delegated authority. Those inputs may change while reasoning is in progress. Database isolation orders the submitted...

<https://arxiv.org/abs/2609.20261>

2 **CONF****AI-Driven Real-Time Relay Optimisation in Smart Urban NR-V2X Networks via Learning-to-Optimise Graph Neural Networks****ArXiv CS.AI** **50%**

arXiv:2609.20271v1 Announce Type: new Abstract: Reliable and low-latency communication is a fundamental requirement for smart city services and Industry 4.0 applications enabled by NR-V2X networks. However, limited Road-Side Unit (RSU) deployment and complex urban propagation...

<https://arxiv.org/abs/2609.20271>

3 **CONF****AgentPProf: Semantic Profiler for Long Horizon AI Agents****ArXiv CS.AI** **50%**

arXiv:2609.20301v1 Announce Type: new Abstract: AI agents increasingly orchestrate long-running activities with users, tools, and system resources for days and weeks. To improve agent quality, safety, and cost efficiency, developers need to determine where failures happen, what...

<https://arxiv.org/abs/2609.20301>

4 **CONF****Learn2Drive: A neural network-based framework for socially compliant automated vehicle control****ArXiv CS.AI** **50%**

arXiv:2510.21736v2 Announce Type: replace-cross Abstract: This study introduces a novel control framework for adaptive cruise control (ACC) in automated driving, leveraging neural networks and physics-informed constraints. As automated vehicles (AVs) adopt advanced features like...

<https://arxiv.org/abs/2510.21736>

5 **CONF****Generating Heterogeneous 3D Geological Microstructures from 2D Images via a Stable Diffusion-Adversarial Model****ArXiv CS.AI** **50%**

arXiv:2609.20358v1 Announce Type: new Abstract: Characterizing the physical properties of clay and cementitious materials matters across many fields, from materials science to geological waste disposal. Property simulation typically calls for 3D imaging, which is expensive, not...

<https://arxiv.org/abs/2609.20358>

6 CONF

The Organization of Inference: Information, Resource Constraints, and AI Production

ArXiv CS.AI 50%

arXiv:2609.20449v1 Announce Type: new Abstract: The economic value of inference depends on how capacity and task information are distributed across stages of AI production. We study these organizational margins using controlled workflow experiments on externally verified...

<https://arxiv.org/abs/2609.20449>

7 CONF

How Do Agent Harnesses Create Value? Planning Information and Release Control in Stateful LLM Agents

ArXiv CS.AI 50%

arXiv:2609.20474v1 Announce Type: new Abstract: Agent harnesses supply planning guidance, organize execution, and check completion. We study how these components affect success, erroneous acceptance, and cost in two Retail experiments and an Airline pilot in τ^2 -bench. The...

<https://arxiv.org/abs/2609.20474>

8 CONF

FreqCondNorm: Towards Cross-domain Predictive Maintenance through a Frequency-Conditioned Transformer Foundation Model

ArXiv CS.AI 50%

arXiv:2609.20535v1 Announce Type: new Abstract: Deep learning predictive maintenance models suffer from poor transferability across machines and operating conditions, especially when labelled data are scarce and signals span five orders of magnitude in sampling frequency (1 Hz...

<https://arxiv.org/abs/2609.20535>

9 CONF

Refuse, Decompose, Refresh: A Claim-Safe Protocol for Closed-Loop AI Evaluation

ArXiv CS.AI 50%

arXiv:2609.20538v1 Announce Type: new Abstract: An AI evaluation can be perfectly reproducible and still support the wrong claim. This risk is acute in closed-loop systems: policy determines visited states, observable components, and which failures leave a measurable trace. We...

<https://arxiv.org/abs/2609.20538>

0 CONF

Limits of Confidence in Diffusion

ArXiv CS.AI 50%

arXiv:2609.20581v1 Announce Type: new Abstract: Discrete diffusion, including remasking and uniform-state samplers, generate a sequence by writing multiple token positions per step, drawing each from a per-position distribution and choosing which positions to write from those...

<https://arxiv.org/abs/2609.20581>

1 CONF

PAA: The Probabilistic Allen Algebra: A Generative and Complete Probabilistic Extension of Allen's Interval Relations

ArXiv CS.AI 50%

arXiv:2609.20634v1 Announce Type: new Abstract: Allen's interval algebra is a qualitative calculus for temporal relations, but its thirteen base relations are crisp predicates over exact interval boundaries. This is inadequate for temporal information from language, perception,...

<https://arxiv.org/abs/2609.20634>

2 **CONF****Ownership in AI-Assisted Everyday Tasks**

ArXiv CS.AI 50%

arXiv:2609.20658v1 Announce Type: new Abstract: When does work done with AI still feel like ours? As AI becomes woven into everyday tasks, we must examine what happens to our sense of ownership and contribution when a machine shares in producing what we make. We report an...

<https://arxiv.org/abs/2609.20658>

3 **CONF****Deep Noir: Autonomous Steering Discovery via Architectural Chronometry in Transformer Models**

ArXiv CS.AI 50%

arXiv:2609.20722v1 Announce Type: new Abstract: Activation steering modifies LLM behavior at inference time, but identifying where and how strongly to steer remains manual. We introduce Deep Noir, a framework that uses Logit Lens convergence and causal head-level attribution to...

<https://arxiv.org/abs/2609.20722>

4 **CONF****FOCAL: Fine-Grained Optimal-Transport-Driven Contrastive Alignment of Language and ECGs with Waveform Enhancement**

ArXiv CS.AI 50%

arXiv:2505.11939v3 Announce Type: replace-cross Abstract: Electrocardiograms (ECGs) are essential non-invasive tools for diagnosing cardiovascular diseases. While recent multimodal ECG-Report contrastive learning methods have shown promise for zero-shot ECG interpretation, they...

<https://arxiv.org/abs/2505.11939>

5 **CONF****MeshKV: A Network-on-Chip KV Cache Fabric for Scalable Transformer Decoding Accelerators**

ArXiv CS.AI 50%

arXiv:2609.19207v1 Announce Type: cross Abstract: Autoregressive transformer decoding is constrained by irregular key-value (KV) cache movement on tiled accelerators. Prior compression and DRAM-placement systems still concentrate traffic on centralized memory paths that...

<https://arxiv.org/abs/2609.19207>

6 **CONF****Layer-wise Curriculum Learning for Efficient LLM Compression**

ArXiv CS.AI 50%

arXiv:2609.19213v1 Announce Type: cross Abstract: In this paper, we introduce layer-wise curriculum learning for efficient LLM compression. The proposed method facilitates the knowledge transfer from the teacher model to the student model, utilizing a curriculum learning...

<https://arxiv.org/abs/2609.19213>

7 **CONF****Perceptual Refinement of an End-to-End Video Streaming Pipeline via Generative AI Layers**

ArXiv CS.AI 50%

arXiv:2609.19215v1 Announce Type: cross Abstract: Traditional codecs treat every region of a frame alike; a generative layer can instead degrade the regions a viewer attends to least and reconstruct them at the client. We present PRESLEY, which extends the prior conference work...

<https://arxiv.org/abs/2609.19215>

8 CONF

PAPC: Platform Mediation for Privacy-Propagation Externalities in AI-Mediated Workflows

ArXiv CS.AI 50%

arXiv:2609.19226v1 Announce Type: cross Abstract: AI-mediated platforms coordinate work through LLM agents acting for different principals. In these workflows, privacy loss can be created before a final answer appears: a memory write, shared-workspace update, inter-agent...

<https://arxiv.org/abs/2609.19226>

9 CONF

Pre-train to Gain: Robust Learning Without Clean Labels

ArXiv CS.AI 50%

arXiv:2511.20844v2 Announce Type: replace-cross Abstract: Training deep networks with noisy labels leads to poor generalization and degraded accuracy due to overfitting to label noise. Existing approaches for learning with noisy labels often rely on the availability of a clean...

<https://arxiv.org/abs/2511.20844>

0 CONF

From Models to Systems: A Comprehensive Survey of Efficient Multimodal Learning

ArXiv CS.AI 50%

arXiv:2609.19445v1 Announce Type: cross Abstract: The rapid expansion of multimodal models has surfaced formidable bottlenecks in computation, memory, and deployment, catalyzing the rise of Efficient Multimodal Learning (EML) as a pivotal research frontier. Despite intensive...

<https://arxiv.org/abs/2609.19445>

1 CONF

AUDITPLAN: Commit, Then Answer for Auditable Safety Alignment

ArXiv CS.AI 50%

arXiv:2609.19325v1 Announce Type: cross Abstract: Safety tuning pipelines judge only the final answer, which makes it difficult to distinguish robust refusal from two undesirable shortcuts: blanket refusal on benign requests and polished but unfaithful safety rationales that do...

<https://arxiv.org/abs/2609.19325>

2 CONF

Kinematics-Grounded Agentic AI for Robotic Additive Manufacturing Process Planning

ArXiv CS.AI 50%

arXiv:2609.19347v1 Announce Type: cross Abstract: Robotic additive manufacturing (AM) extends material-extrusion printing beyond gantry kinematics but makes process planning robot-dependent. A slicer-generated plan that appears favorable in part coordinates can become infeasible...

<https://arxiv.org/abs/2609.19347>

3 CONF

Large language models eroding science understanding: an empirical study of malignment

ArXiv CS.AI 50%

arXiv:2604.25639v2 Announce Type: replace-cross Abstract: This paper is accepted and in press for AI and Ethics. This paper includes the supplementary data file at the end of the manuscript. This study examines whether large language models (LLMs) can reliably answer scientific...

<https://arxiv.org/abs/2604.25639>

4 **CONF****AURORA: A Natural Language-Driven Agentic Framework for Understanding, Reasoning, and Orchestrating Reliable Air-Ground Co-Simulation****ArXiv CS.AI** **50%**

arXiv:2609.19527v1 Announce Type: cross Abstract: Air-ground transportation research increasingly relies on co-simulation, yet constructing scenarios remains labor-intensive and difficult to validate. More importantly, a generated scenario may execute successfully while failing...

<https://arxiv.org/abs/2609.19527>

5 **CONF****Detecting Soft Errors in Parallel Software with LLM-tuned Instruction Duplication****ArXiv CS.AI** **50%**

arXiv:2609.19531v1 Announce Type: cross Abstract: We propose PaRID (PaRallel Instruction Duplication), a software-directed soft error detection framework that requires only compile-time effort for multithreading parallel programs. PaRID addresses two key challenges: supporting...

<https://arxiv.org/abs/2609.19531>

6 **CONF****Compressed Active Subspaces for Scalable Bayesian Inference****ArXiv CS.AI** **50%**

arXiv:2609.19539v1 Announce Type: cross Abstract: Active subspace methods provide a framework for quantifying predictive uncertainty in high-dimensional models by identifying and performing inference along parameter directions that have the greatest influence on the model...

<https://arxiv.org/abs/2609.19539>

7 **CONF****A Multi-Modal Generative Model for Tomato Disease Leaves Understanding****ArXiv CS.AI** **50%**

arXiv:2609.19555v1 Announce Type: cross Abstract: Artificial intelligence for plant disease analysis has advanced from task-specific classifiers to multi-modal models capable of jointly interpreting visual and textual information. However, practical deployment in precision...

<https://arxiv.org/abs/2609.19555>

8 **CONF****CliniCIRCA: A Modular LLM Framework for Constructing Longitudinal Mental Health Patient Journeys from Raw EHR Narratives****ArXiv CS.AI** **50%**

arXiv:2609.19585v1 Announce Type: cross Abstract: In mental health care, reasoning over patient journeys is a key task for clinicians. Yet these journeys, encompassing a longitudinal progression of biological, psychological, and social events, are often spread across disparate...

<https://arxiv.org/abs/2609.19585>

9 **CONF****Red-Teaming Auto Mode: Improving Blocking Classifiers Against Malign Coding Agents****ArXiv CS.AI** **50%**

arXiv:2609.19587v1 Announce Type: cross Abstract: To keep coding agents from going off the rails, production systems now review each proposed action with a blocking monitor that can reject it before it runs (Auto Mode in Claude Code, Guardian in OpenAI's Codex). Prior...

<https://arxiv.org/abs/2609.19587>

0 **CONF****TouchSight: Bare-Handed Tactile Prediction from Egocentric Video via Generative Visual Augmentation****ArXiv CS.AI** **50%**

arXiv:2609.20414v1 Announce Type: cross Abstract: Tactile signals provide direct contact and force measurements that are essential for understanding physical interactions and enabling dexterous robotic manipulation. However, tactile sensing requires direct measurement at contact...

<https://arxiv.org/abs/2609.20414>

1 **CONF****SoK: Trading Agents or Market Crashers? Dissecting Robustness and Security Failures in Academic Financial LLM Trading Schemes****ArXiv CS.AI** **50%**

arXiv:2609.19705v1 Announce Type: cross Abstract: Autonomous large language model (LLM) agents are moving rapidly into high-stakes domains, yet existing agentic-AI security studies remain largely domain-agnostic and overlook the distinctive, high-consequence attack surface such...

<https://arxiv.org/abs/2609.19705>

2 **CONF****Trust, but Validate the Instrument: Auditing AI-Generated RTL Verification Plans on Authored Security-Regression Proxies****ArXiv CS.AI** **50%**

arXiv:2609.19844v1 Announce Type: cross Abstract: AI-generated RTL verification plans can satisfy a provider schema yet fail at the boundary to trusted execution. We present SecTB-RTL, an auditable framework covering 31 tasks and 124 authored hardware-security regressions. A...

<https://arxiv.org/abs/2609.19844>

3 **CONF****PACE: Precise AI Cinematic Expression: A Typed Specification for Script-Grounded Previsualization and Geometric Conformance****ArXiv CS.AI** **50%**

arXiv:2609.19853v1 Announce Type: cross Abstract: Between a screenplay and a film sits a planning problem that is spatial first: who stands where, and what a camera sees from where it stands. An image diffusion model asked for a shot in free text settles that plan by its own...

<https://arxiv.org/abs/2609.19853>

4 **CONF****Zarya: A Hybrid Autoregressive--Masked Diffusion Language Model with Flexible Training and Dual-Mode Inference****ArXiv CS.AI** **50%**

arXiv:2609.19868v1 Announce Type: cross Abstract: Autoregressive language models (ARMs) are constrained by sequential, left-to-right generation, while masked diffusion models (MDMs) enable parallel decoding but suffer from high computational overhead due to the inability to...

<https://arxiv.org/abs/2609.19868>

5 **CONF****PetriBench: Benchmarking LLM Reasoning over Dynamic State Spaces****ArXiv CS.AI** **50%**

arXiv:2609.19883v1 Announce Type: cross Abstract: Characterizing LLM reasoning remains an open challenge, as many existing benchmarks isolate specific reasoning skills, rely on external knowledge, or are costly to extend. We introduce PetriBench, a compact, fully self-contained,...

<https://arxiv.org/abs/2609.19883>

6 CONF

KoNeoBench: A Curated Evaluation Dataset for LLM Understanding of Korean Neologisms

ArXiv CS.AI 50%

arXiv:2609.19916v1 Announce Type: cross Abstract: Large language models (LLMs) are typically evaluated on static benchmarks, even though natural language constantly evolves through newly emerging words and meanings. Existing Korean benchmarks are centered on established...

<https://arxiv.org/abs/2609.19916>

7 CONF

Governance-as-Code: Translating EU AI Act Technical Requirements into Executable Compliance Pipelines for Generative AI Systems

ArXiv CS.AI 50%

arXiv:2609.20016v1 Announce Type: cross Abstract: The EU AI Act (Regulation 2024/1689) imposes technical obligations on high-risk AI providers, yet Articles 8-15 were drafted for predictive AI and leave seven technical gaps when applied to generative systems, spanning...

<https://arxiv.org/abs/2609.20016>

8 CONF

AI Should Facilitate Democratic Deliberation at Scale

ArXiv CS.AI 50%

arXiv:2609.20059v1 Announce Type: cross Abstract: AI systems can strengthen democracy by supporting deliberation at scale by addressing cognitive, social, platform-design, and market-driven frictions, while preserving human agency. Unlike proposals such as liquid democracy that...

<https://arxiv.org/abs/2609.20059>

9 CONF

Reading Emotions in the Token Space: Discriminative Adaptation of SpeechLLMs for Emotion Recognition

ArXiv CS.AI 50%

arXiv:2609.20081v1 Announce Type: cross Abstract: SpeechLLMs have shown strong potential for emotion recognition, yet they read the predicted emotion off a generative decoder not suited for classification: it can emit labels outside the target set and favors frequent classes. We...

<https://arxiv.org/abs/2609.20081>

00 CONF

Local Sparsity Enables Unsupervised LLM Safety Detection

ArXiv CS.AI 50%

arXiv:2609.20129v1 Announce Type: cross Abstract: Deployment-time safety methods for large language models (LLMs) are predominantly supervised and assume access to unsafe training data. Nevertheless, new attacks and harm categories regularly arise, not captured by models trained...

<https://arxiv.org/abs/2609.20129>

01 CONF

AdaRepair-Mem: Adaptive Experience Orchestration for Repository-Level Program Repair

ArXiv CS.AI 50%

arXiv:2609.20130v1 Announce Type: cross Abstract: Recent memory-augmented repository-level program repair methods reuse historical repair experiences to improve LLM-based issue resolution. However, our analysis reveals three limitations in existing repository-level memory...

<https://arxiv.org/abs/2609.20130>

02 CONF

Cross-Modal Attention Acts as a Frequency Filter: Why Verbose Prompts Improve Robustness in Vision-Language Models

ArXiv CS.AI 50%

arXiv:2609.20139v1 Announce Type: cross Abstract: Vision-language models (VLMs) are fragile under image corruption. We find that the wording of the question affects VLMs in two opposite ways. Verbose questions make VLMs substantially more robust---e.g., rephrasing "Is there a...

<https://arxiv.org/abs/2609.20139>

03 CONF

Designing Against Deskilling: Metacognitive Feedback Reduces Cognitive Offloading to LLM Assistants

ArXiv CS.AI 50%

arXiv:2609.20143v1 Announce Type: cross Abstract: Cognitive offloading to AI can reduce opportunities to practice skills, creating risks of deskilling. However, it remains unclear how to prevent deskilling without restricting access to AI. Here, we design two interventions to...

<https://arxiv.org/abs/2609.20143>

04 CONF

CleanVideo: Adaptive Concept Erasure for Text-to-Video Diffusion Models

ArXiv CS.AI 50%

arXiv:2609.20267v1 Announce Type: cross Abstract: Concept erasure aims to selectively eliminate undesired visual semantics from pre-trained generative models without compromising their general utility. Extending concept erasure from images to video is nontrivial. Target concepts...

<https://arxiv.org/abs/2609.20267>

05 CONF

SoftTri: Smooth Triangular Membership Functions for Adaptive Fuzzy Inference Systems

ArXiv CS.AI 50%

arXiv:2609.20194v1 Announce Type: cross Abstract: Triangular membership functions (MFs) are widely used in fuzzy systems because of their interpretability, low parameterization complexity, and strong locality properties. However, their inherent nondifferentiability at knot...

<https://arxiv.org/abs/2609.20194>

06 CONF

Is It Still Worth Training a Classical Model in the Era of LLMs? A Crossover Benchmark on Tabular Data

ArXiv CS.AI 50%

arXiv:2609.20218v1 Announce Type: cross Abstract: Large language models can label a tabular row from a plain-English description with no training - a capability now shipping in mainstream spreadsheet tools such as Microsoft Copilot in Excel and Anthropic's Claude for Excel -...

<https://arxiv.org/abs/2609.20218>

07 CONF

Lens: Bringing the Right Semantic Perspective into Focus for Training-Free Multimodal Representation Learning

ArXiv CS.AI 50%

arXiv:2609.20252v1 Announce Type: cross Abstract: High-quality representations are essential for a wide range of downstream tasks. Dedicated embedding models are explicitly optimized for representation learning, yet their training data are often more limited in scale and...

<https://arxiv.org/abs/2609.20252>

08 CONF

A Multi-Objective Optimisation Framework for Corticomuscular EEG-EMG Pair Selection in Hybrid BCI

ArXiv CS.AI 50%

arXiv:2609.20275v1 Announce Type: cross Abstract: Hybrid brain-computer interface (BCI) systems that integrate electroencephalography (EEG) and electromyography (EMG) signals have shown significant potential in improving the reliability of motor imagery (MI) classification,...

<https://arxiv.org/abs/2609.20275>

09 CONF

Labeled Incidence Structures for Native Transformer Modeling of Text, Knowledge Graphs, and Hypergraphs

ArXiv CS.AI 50%

arXiv:2609.20278v1 Announce Type: cross Abstract: Text, knowledge graphs, and hypergraphs all have elements that play distinct roles within relation instances, structure that is lost when data is flattened into token sequences. We introduce labeled incidence structures (LIS), a...

<https://arxiv.org/abs/2609.20278>

10 CONF

Human and AI-generated texts between modal logic and statistics

ArXiv CS.AI 50%

arXiv:2609.20311v1 Announce Type: cross Abstract: We read the geometry of semantic neighbourhood graphs as modal logic and give that reading a statistical form, in order to make precise the structural difference between human and machine-generated text. Texts are the worlds of a...

<https://arxiv.org/abs/2609.20311>

11 CONF

LLM-Guided Transformation of Non-Critical Driving Scenes into Safety-Critical Scenarios Using Augmented Reality

ArXiv CS.AI 50%

arXiv:2609.20318v1 Announce Type: cross Abstract: Testing Autonomous Driving Systems (ADS) requires realistic safety-critical scenarios, but collecting such data from real-world driving is costly and unsafe. This paper presents an automated pipeline that transforms safe driving...

<https://arxiv.org/abs/2609.20318>

12 CONF

Xeno-Interpretability: Investigating the Alien Minds of LLMs

ArXiv CS.AI 50%

arXiv:2609.20408v1 Announce Type: cross Abstract: Large language models are usually interpreted through concepts that humans already possess: truthfulness, refusal, deception, personality, harmfulness, and related categories. This paper asks whether models may also represent and...

<https://arxiv.org/abs/2609.20408>

13 CONF

Stress-testing Alignment Midtraining

ArXiv CS.AI 50%

arXiv:2609.20412v1 Announce Type: cross Abstract: When aligning frontier models through post-training techniques, it is not possible to directly demonstrate all of the behaviours we want a model to exhibit in all possible deployment environments; our model must generalise...

<https://arxiv.org/abs/2609.20412>

14 CONF

SCGFM-ART: Amortized Relational Transport for Structure-Centric Graph Foundation Models

ArXiv CS.AI 50%

arXiv:2609.20419v1 Announce Type: cross Abstract: Graph foundation models (GFMs) aim to learn transferable representations across severely heterogeneous graph domains. However, severe domain shifts in topology, graph scale, and feature semantics impede the construction of a...

<https://arxiv.org/abs/2609.20419>

15 CONF

When Do Language-Grounded Explanations Help? A Graph-Bottleneck for Farm Monitoring Interpretable Sheep Facial Pain

ArXiv CS.AI 50%

arXiv:2609.20427v1 Announce Type: cross Abstract: Automated pain recognition from facial expression could make continuous welfare assessment practical in sheep, but adoption depends on trust: a stockperson cannot act on a score that arrives without justification. We ground a...

<https://arxiv.org/abs/2609.20427>

16 CONF

Deep Learning-Based Classification of Cognitive and Resting States Using Electroencephalography Signals

ArXiv CS.AI 50%

arXiv:2609.20467v1 Announce Type: cross Abstract: The categorization of cognitive and resting states derived from electroencephalography (EEG) signals is crucial for comprehending fluctuations in brain activity linked to various mental states. EEG provides a non-intrusive...

<https://arxiv.org/abs/2609.20467>

17 CONF

greCAPTCHA: Assessing Understanding as Evidence of Research Authorship Under Generative AI

ArXiv CS.AI 50%

arXiv:2609.20481v1 Announce Type: cross Abstract: Conferences, journals, funders, schools, and universities are struggling with a surge of potentially AI-generated submissions from ostensibly human authors, who may not have exercised sufficient human oversight for their...

<https://arxiv.org/abs/2609.20481>

18 CONF

Model-Agnostic and Language-Agnostic Voice Pipeline Improvement for the Agriculture Domain

ArXiv CS.AI 50%

arXiv:2609.20504v1 Announce Type: cross Abstract: FarmerChat is Digital Green's AI-powered agricultural advisory assistant for smallholder farmers, who access it in their own language through text, voice, or photographs. Voice is a critical channel for this population, yet...

<https://arxiv.org/abs/2609.20504>

19 CONF

Inference-Engine Fingerprinting Attacks are Practical: Exploring Model-Driven Environmental Discovery, Exploitation, and Escape

ArXiv CS.AI 50%

arXiv:2609.20614v1 Announce Type: cross Abstract: Frontier AI models are rapidly gaining the ability to exploit vulnerabilities in complex pieces of software. The risk is not theoretical, as evidenced by recent sandbox escapes performed by frontier models at OpenAI and...

<https://arxiv.org/abs/2609.20614>

20 CONF

A Simulation Platform for AUV Fault Recovery: Exploring LLM-Based Diagnostic Strategies

ArXiv CS.AI 50%

arXiv:2609.20620v1 Announce Type: cross Abstract: Autonomous underwater vehicles (AUVs) operating beyond reliable communications must recover from failures without human intervention. We investigate an architecture in which conventional deterministic layered control autonomy...

<https://arxiv.org/abs/2609.20620>

21 CONF

HIL-UMI: Bringing Human-in-the-Loop Post-Training of Vision-Language-Action Models to Universal Manipulation Interface

ArXiv CS.AI 50%

arXiv:2609.20659v1 Announce Type: cross Abstract: Large-scale vision-language-action (VLA) models provide powerful priors for robot manipulation, yet adapting them to a specific deployment remains challenging. Supervised fine-tuning (SFT) on task-specific demonstrations provides...

<https://arxiv.org/abs/2609.20659>

22 CONF

Large Language Models as Falsifiers for Cyber-Physical Systems

ArXiv CS.AI 50%

arXiv:2609.20752v1 Announce Type: cross Abstract: Falsification searches for counterexamples to formal specifications in cyber-physical systems (CPS). With specifications written in Signal Temporal Logic (STL), falsification can be formulated as a robustness optimization...

<https://arxiv.org/abs/2609.20752>

23 CONF

Prediction-Powered Smoothing and Validation for Disaggregated AI Evaluation

ArXiv CS.AI 50%

arXiv:2609.20758v1 Announce Type: cross Abstract: Evaluating an AI system requires disaggregated assessment, as performance varies across domains such as benchmark task types or conversation types in deployed agents. Exhaustive testing is expensive, so evaluation rests on a...

<https://arxiv.org/abs/2609.20758>

24 CONF

Harm Laundering in GPT Models: Evidence That Gender Discrimination Is Transformed Rather Than Reduced Across Safety-Trained Generations

ArXiv CS.AI 50%

arXiv:2609.20779v1 Announce Type: cross Abstract: Safety evaluations for large language models rely on surface-form classifiers that report declining harm scores across model generations. We provide evidence that this methodology is systematically incomplete: explicit...

<https://arxiv.org/abs/2609.20779>

25 CONF

Quantifying Overclaiming Propensity in Frontier LLM Agents

ArXiv CS.AI 50%

arXiv:2609.20812v1 Announce Type: cross Abstract: Frontier coding agents are increasingly trusted to work autonomously for long periods, yet an agent's final response is often the only account of that work a user sees. We quantify the propensity of frontier agents to...

<https://arxiv.org/abs/2609.20812>

26 CONF

Paint-Anything: Unified Any-Color Control for Image Generation and Editing

ArXiv CS.AI 50%

arXiv:2609.20816v1 Announce Type: cross Abstract: Professional design requires any-color control: the ability to specify an object's target color with any 24-bit hex value for image generation and editing. Prior work has explored color generation, editing, and colorization, but...

<https://arxiv.org/abs/2609.20816>

27 CONF

TripScore: Aligning LLMs for Real-World Travel Planning via Expert-Calibrated Reward

ArXiv CS.AI 50%

arXiv:2510.09011v4 Announce Type: replace Abstract: In our deployed travel-planning service, most users give minimal inputs or free-form requests rather than the structured constraint checklists assumed by existing benchmarks. We therefore present TripScore, a behavior-grounded...

<https://arxiv.org/abs/2510.09011>

28 CONF

Model Specific Task Similarity for Vision Language Model Selection via Layer Conductance

ArXiv CS.AI 50%

arXiv:2602.01346v2 Announce Type: replace Abstract: While open sourced Vision-Language Models (VLMs) have proliferated, selecting the optimal pretrained model for a specific downstream task remains challenging. Exhaustive evaluation is often infeasible due to computational...

<https://arxiv.org/abs/2602.01346>

29 CONF

Faithful, Not Corrective: Model Capability Governs Message-Format Effects in Multi-Hop Agent Relays

ArXiv CS.AI 50%

arXiv:2607.09678v2 Announce Type: replace Abstract: When LLM agents hand information to one another, does the message format matter? Two literatures disagree: format-optimization work reports that structured messages cut cost without hurting accuracy, while format-restriction...

<https://arxiv.org/abs/2607.09678>

30 CONF

A Neuropsychologically Grounded Evaluation of LLM Cognitive Abilities

ArXiv CS.AI 50%

arXiv:2603.02540v2 Announce Type: replace Abstract: Large language models (LLMs) display a unified "general factor" of capability across 10 benchmarks (a finding confirmed by our factor analysis of 156 models), yet they still struggle with simple, trivial tasks for humans. This...

<https://arxiv.org/abs/2603.02540>

31 CONF

Physics-Informed Support Vector Kernels via Green-Function Analogies and Jackson-Chebyshev Spectral Design

ArXiv CS.AI 50%

arXiv:2502.11153v4 Announce Type: replace-cross Abstract: Kernel selection for regression of physical observables is often heuristic. We investigate a physics-informed strategy in which functional forms and spectral structures associated with Green's functions motivate kernel...

<https://arxiv.org/abs/2502.11153>

32 CONF

Spherical Cauchy Variational Autoencoders: Heavy Angular Tails and Exact KL Evaluation

ArXiv CS.AI 50%

arXiv:2506.21278v4 Announce Type: replace-cross Abstract: Heavy-tailed posteriors are routine in Euclidean variational autoencoders, where the Student family relaxes the Gaussian without new machinery. The sphere has had no comparable option. Von Mises-Fisher distribution needs...

<https://arxiv.org/abs/2506.21278>

33 CONF

Understanding Role Switching in Human-AI Collaboration through Multimodal Behavioral Signals

ArXiv CS.AI 50%

arXiv:2509.20666v2 Announce Type: replace-cross Abstract: Human-AI collaboration often requires dividing complex tasks into complementary subtasks. As a task unfolds, users may want to shift which subtasks they perform and which their AI partner performs, in response to evolving...

<https://arxiv.org/abs/2509.20666>

34 CONF

 $D^{\$}$ -MOPD: Dynamic Domain Scheduling for Efficient Multi-Teacher Distillation

ArXiv CS.AI 50%

arXiv:2608.24987v2 Announce Type: replace-cross Abstract: Multi-teacher on-policy distillation (MOPD) distills several domain-expert teachers into a single student by minimizing per-domain reverse-KL divergence on the student's own rollouts. Existing approaches typically fix the...

<https://arxiv.org/abs/2608.24987>

35 CONF

Perturbing the Phase: Analyzing Adversarial Robustness of Complex-Valued Neural Networks

ArXiv CS.AI 50%

arXiv:2602.06577v2 Announce Type: replace-cross Abstract: Complex-valued neural networks (CVNNs) are rising in popularity for all kinds of applications. To safely use CVNNs in practice, analyzing their robustness against outliers is crucial. One well known technique to...

<https://arxiv.org/abs/2602.06577>

36 CONF

Drag-Aware Aerodynamic Manipulability for Torque-Limited Redundant Multirotors: Aerodynamic Promptness based on the Symmetric Acceleration Capacity

ArXiv CS.AI 50%

arXiv:2603.07998v2 Announce Type: replace-cross Abstract: Aerodynamic promptness quantifies how rotor-speed variations generate multirotor wrench variations, but its Euclidean formulation assigns the same local cost to a given rotor acceleration at every operating speed. This...

<https://arxiv.org/abs/2603.07998>

37 CONF

Reinforcement Learning for Graph Generation under a Hard Assortativity Constraint

ArXiv CS.AI 50%

arXiv:2605.23285v2 Announce Type: replace-cross Abstract: Generating graph ensembles with precisely controlled structural properties is central to investigating how network structure shapes function. Canonical ensembles impose constraints only in expectation (soft constraints),...

<https://arxiv.org/abs/2605.23285>

38 CONF

Time-Aware Diffusion based on Preference Disentanglement for Generative Recommendation

ArXiv CS.AI 50%

arXiv:2606.01670v3 Announce Type: replace-cross Abstract: Recently, Generative Recommenders (GRs) have emerged as a transformative recommendation paradigm by replacing traditional item IDs with semantic indices (SIDs). Owing to the exceptional generative capabilities of...

<https://arxiv.org/abs/2606.01670>

39 CONF

AnyAudio-Judge: A Dynamic Rubric-Based Benchmark and Evaluator for Audio Instruction Following

ArXiv CS.AI 50%

arXiv:2606.03116v2 Announce Type: replace-cross Abstract: The rapid advancement of instruction-guided audio generation has highlighted the critical need for robust alignment evaluation. Current automated evaluation methods heavily rely on holistic scoring from general-purpose...

<https://arxiv.org/abs/2606.03116>

40 CONF

GigaBrain-WBC-0.5: A Behavior World Model for Robust Humanoid Whole-Body Tracking with Environment Interaction

ArXiv CS.AI 50%

arXiv:2608.18234v4 Announce Type: replace-cross Abstract: General-purpose motion trackers enable humanoid robots to follow diverse whole-body motions while maintaining balance, but are trained only on flat ground, failing to exploit bipedal mobility over complex terrain....

<https://arxiv.org/abs/2608.18234>

41 CONF

How Humans and LLMs Read Gender into "Gender-Neutral" Physical Descriptions

ArXiv CS.AI 50%

arXiv:2609.16366v2 Announce Type: replace-cross Abstract: When foundation models describe people, recent work in AI fairness, accessibility, and ethics recommends avoiding inferred identity labels (e.g., "she", "his") in favor of seemingly "objective" physical descriptions...

<https://arxiv.org/abs/2609.16366>