



Pipeline

Scrape HackerNews (Show/Top/New), Reddit (r/artificial, r/MachineLearning, r/singularity, r/LocalLLaMA), TechCrunch AI, VentureBeat AI, MIT Tech Review, The Verge AI, Wired AI, ArXiv CS.AI

Filter 40+ AI keyword patterns — LLMs, alignment, safety, research, tooling

Dedup Title-similarity (45% threshold) merges cross-source variants

Score Authenticity = 0.5 base + 0.3 per extra source + diversity + HN score bonuses

Rank Sorted by authenticity. Below 0.25 threshold are dropped.

Authenticity Tiers

- **VERIFIED** 0.80 - 1.00 3+ independent sources including media
- **CONFIRMED** 0.50 - 0.79 2 sources or 1 high-quality media
- **EMERGING** 0.25 - 0.49 Single-source; passes AI keyword filter

Source Breakdown

Source	Articles	Category
ArXiv CS.AI	77	Media
HackerNews	6	Community
The Verge AI	4	Media
TechCrunch AI	3	Media

VERIFIED INTELLIGENCE — 4 stories

VERI

Show HN: jevals - replacing LLM judges with typed Jev decisions

HackerNews

+1 source

90%

HN 19

arXiv:2608.18265v4 Announce Type: replace-cross Abstract: We introduce a general, easy-to-implement AI-based modeling technique for analyzing human behavior. A key feature of this approach, which contrasts with existing modeling techniques, is that it combines the flexibility...

<https://github.com/openlayer-ai/jevals>

VERI

The AI regulation smackdown isn't over

The Verge AI

+1 source

80%

At the start of this week, the who's-who of AI seemed - at least tentatively - on the side of AI regulation. Over the weekend, Anthropic CEO Dario Amodei had proposed a three-step plan for slowing AI development, including by embedding third-party evaluators in labs,...

<https://www.theverge.com/ai-artificial-intelligence/997706/the-ai-regulation-smackdown-...>

VERI

Trump now says he wants to form an 'AI Force'

The Verge AI

+1 source

80%

The president posted on Truth Social that he wanted to appoint an "AI czar" to lead a new "AI force." He made the announcement amid growing calls from across the political spectrum and even within the industry to pump the brakes on AI development. He posted that his...

<https://www.theverge.com/ai-artificial-intelligence/997867/trump-ai-force-ai-czar>

VERI

Visual Navigation Transformer with Pose Attention

ArXiv CS.AI

+1 source

80%

arXiv:2609.21212v1 Announce Type: cross Abstract: Learned navigation policies typically consume observations as a temporally ordered history, with positional encodings tying each observation to when it was seen, making it difficult to reuse experience from earlier traversals of...

<https://arxiv.org/abs/2609.21212>

CONFIRMED SIGNALS — 82 stories

CONF

Show HN: A competition for small neural networks that play strategy games

HackerNews

50%

HN 69

Hi HN - long-time lurker (since 2012!), first time poster. Pizza Bot is a self-hosted desktop app for Mac, Windows, and Linux that runs AI agents in the background and exposes them through an email-like UI. Finished work shows up in Unread, and anything waiting on your approval...

<https://tinybrains.dev>

CONF

Show HN: Mini-AGI - Dynamic continual learning model trained on 8GB VRAM

HackerNews

50%

HN 54

Sorry for the pretentious name, I know, I know.. It just contains all the pieces I would like to see a AGI model to have, and I can't stand the temptation. Before throwing rocks at me, please take a glance at the README, and I hope it will cover your mood a little bit. So, first...

<https://github.com/volotat/mini-agi/>

CONF

Show HN: Airmash - HTML5 Multiplayer Missile Warfare

HackerNews 50% HN 6

This is the official reboot of the 2017 hit game. Back then it was #1 post on HN for 3 consecutive days: <https://news.ycombinator.com/item?id=15892066> The game was rewritten from the ground up: New graphics, WebGL shader effects, Map obstacles, Completely new UI, Additional...

<https://airma.sh/>

CONF

Show HN: Architectural Guardrails for AI native development

HackerNews 50% HN 5

No summary — see link for full article.

<https://github.com/cxgrd/cli>

CONF

Show HN: Testing a non-generative decision model on 5,500 CLINC150 inputs

HackerNews 50% HN 3

No summary — see link for full article.

<https://github.com/chr-kelly/jev-cookbook>

0 CONF

No one is surprised that Nvidia's Jensen Huang thinks AI fears are overblown

The Verge AI 50%

The man who may stand to make the most money from the AI boom seems to think he knows better than anyone else, including researchers who have studied and worked on AI for decades. In an interview with CBS Sunday Morning, he claimed there was a "0% chance" of AI being the end of...

<https://www.theverge.com/ai-artificial-intelligence/997936/nvidia-jensen-huang-ai-fears...>

1 CONF

Humans, not rogue AI, are still the biggest cybersecurity risk to energy systems

The Verge AI 50%

Before recent high-profile hacks raised the specter of AI possibly "killing all humans," our energy systems were already disturbingly vulnerable to cyberattack - and the risk is growing. "We were always prey. We were just kind of surviving at the appetite of our predators,"...

<https://www.theverge.com/science/997834/ai-cyberattack-energy-critical-infrastructure>

2 CONF

RBS-Attention: Radius-Bounded Sparse Prefill for Long-Context Large Language Models

ArXiv CS.AI 50%

arXiv:2609.20971v1 Announce Type: new Abstract: Long-context large language model inference is increasingly limited by prefill, where dense self-attention processes the entire prompt before generation begins. Sparse block selection can reduce this cost, but a block centroid may...

<https://arxiv.org/abs/2609.20971>

3 CONF

CaLR: Causal Latent Revision for Robust Diffusion Reasoning

ArXiv CS.AI 50%

arXiv:2609.20981v1 Announce Type: new Abstract: Autoregressive (AR) models suffer from local greediness, while diffusion language models (DLMs) often lack the strict causal structure required for reasoning. To combine the advantages and overcome the drawbacks of the dual, we...

<https://arxiv.org/abs/2609.20981>

4 **CONF****LoRA Enhanced Contrastive Learning with SAS Vision Transformers****ArXiv CS.AI** **50%**

arXiv:2609.21061v1 Announce Type: new Abstract: Automatic target recognition (ATR) with synthetic aperture sonar (SAS) supports advanced naval capabilities, but deep learning is constrained by scarce target imagery, background clutter, and human-in-the-loop assessment. We adapt...

<https://arxiv.org/abs/2609.21061>

5 **CONF****Detecting Hallucination in LLMs: Tracing the Topological Signatures of Impaired Context Sharing****ArXiv CS.AI** **50%**

arXiv:2609.21096v1 Announce Type: new Abstract: In this work, we examine the topology of information flow patterns within attention graphs to effectively distinguish hallucinated from non-hallucinated responses. We analyze the Forman-Ricci curvature to identify structural...

<https://arxiv.org/abs/2609.21096>

6 **CONF****Decoupling Internal Representational Changes and Causal Importance in Fine-Tuned Large Language Models****ArXiv CS.AI** **50%**

arXiv:2609.21113v1 Announce Type: new Abstract: Fine-tuning has emerged as a widely adopted approach for adapting LLMs to a variety of downstream tasks. However, how it reshapes their internal mechanisms remains poorly understood. To address this, we investigate how fine-tuning...

<https://arxiv.org/abs/2609.21113>

7 **CONF****TinyCeNN-LM: Quality-Gated Conversion of Pretrained Attention with CeNN-Inspired Cellular-Recurrent Layers****ArXiv CS.AI** **50%**

arXiv:2609.21139v1 Announce Type: new Abstract: Replacing attention in a pretrained language model is a compatibility problem: a plausible substitute may alter representations expected by later layers. TinyCeNN-LM introduces a \emph{quality-gated post-training conversion}...

<https://arxiv.org/abs/2609.21139>

8 **CONF****Intent-Governed Tool Authorization for AI Agents****ArXiv CS.AI** **50%**

arXiv:2606.22916v4 Announce Type: replace Abstract: Tool-using AI agents commonly operate under integration credentials whose static permissions exceed a user's current request. We present Intent-Governed Access Control (IGAC), a server-side authorization layer that converts a...

<https://arxiv.org/abs/2606.22916>

9 **CONF****Efficient Benchmarking in Production: A Study of an Evolving LLM Agent****ArXiv CS.AI** **50%**

arXiv:2609.21267v1 Announce Type: new Abstract: Production LLM agents are evaluated repeatedly as they evolve, but full agent benchmarks are costly to rerun. We study efficient recurring evaluation for a production analytics agent serving tens of thousands of monthly active...

<https://arxiv.org/abs/2609.21267>

0 **CONF****AI-GRACE: A Use-Case Operationalization Framework for Agentic AI: From Organizational Objectives and Obligations to Deployment Capabilities and Architecture**

ArXiv CS.AI 50%

arXiv:2609.21192v1 Announce Type: new Abstract: Organizations deploying agentic artificial intelligence must determine more than whether a model is trustworthy; they must establish what to validate, control, and observe for a use case to deliver its intended outcome while...

<https://arxiv.org/abs/2609.21192>

1 **CONF****Information-Gain Rewards over Diversity-Pruned Tests: GT-Anchored Verifier Co-Training for Reliable Code Generation**

ArXiv CS.AI 50%

arXiv:2609.21208v1 Announce Type: new Abstract: Self-play methods that co-train a single language model as both coder and test author promise to move code-generation RL beyond fixed test suites, but they suffer from two coupled pathologies: permissiveness collapse, where...

<https://arxiv.org/abs/2609.21208>

2 **CONF****PlaceReasoner-Beta: Reasoning-Driven Macro Placement and Benchmarking**

ArXiv CS.AI 50%

arXiv:2609.21263v1 Announce Type: new Abstract: Automated macro placement remains a fundamental challenge in VLSI physical design. Despite decades of research, existing approaches predominantly optimize hand-crafted proxy objectives, such as estimated wirelength, and typically...

<https://arxiv.org/abs/2609.21263>

3 **CONF****GameASG-Bench: Benchmarking Autonomous Software Generation for Game Development**

ArXiv CS.AI 50%

arXiv:2609.21293v1 Announce Type: new Abstract: Autonomous software generation (ASG) aims to turn human requirements into executable applications, but delivering these applications does not necessarily establish that their interacting components satisfy the specified behavioral...

<https://arxiv.org/abs/2609.21293>

4 **CONF****LEGIT: Credentialing Protocol for Trustworthy AI Agent Marketplaces**

ArXiv CS.AI 50%

arXiv:2609.21325v1 Announce Type: new Abstract: Agentic marketplaces are emerging where AI agents with varying capabilities autonomously complete specialized tasks for buyers. A major challenge of such marketplaces is that buyers cannot easily determine which agent will perform...

<https://arxiv.org/abs/2609.21325>

5 **CONF****Self-Reference in Large Language Models: The Introspection Threshold for Recursive Self-Improvement**

ArXiv CS.AI 50%

arXiv:2607.04277v2 Announce Type: replace-cross Abstract: The pursuit of self-evolving AI raises a critical question: when is autonomous self-improvement sustainable rather than degenerative? Drawing an analogy to von Neumann's complexity threshold for self-reproducing automata,...

<https://arxiv.org/abs/2607.04277>

6 CONF

GVPO++: Group Variance Policy Optimization for LLM Post-Training and On-Policy Distillation

ArXiv CS.AI 50%

arXiv:2609.21432v1 Announce Type: new Abstract: Post-training plays a pivotal role in enhancing the reasoning capabilities and task-specific expertise of large language models (LLMs). Despite recent advances in post-training methods, such as Group Relative Policy Optimization...

<https://arxiv.org/abs/2609.21432>

7 CONF

The Communication Bottleneck: A Round-Trip Study of Tree-Structured Expression Serialization in Language Models

ArXiv CS.AI 50%

arXiv:2609.21509v1 Announce Type: new Abstract: When language models reason in chain-of-thought or exchange free-text intermediates, they serialize structured information into natural language. How much tree-structured compositional content survives this bottleneck? We propose a...

<https://arxiv.org/abs/2609.21509>

8 CONF

Learning-to-Optimize as the Missing Architectural Layer of AI-Native Networks

ArXiv CS.AI 50%

arXiv:2609.21519v1 Announce Type: new Abstract: Artificial Intelligence (AI) is becoming a fundamental design principle of future AI-native communication networks, enabling autonomous resource management, adaptive control, and zero-touch network operation. While current...

<https://arxiv.org/abs/2609.21519>

9 CONF

Reducing Barriers to Academic Support: Evaluating a Course-Specific RAG System for Addressing Help-Seeking Disparities in Higher Education

ArXiv CS.AI 50%

arXiv:2609.21600v1 Announce Type: new Abstract: Access to academic support is a key determinant of student success, yet students experience it unequally: some readily seek help from lecturers or tutors, while others hesitate due to anxiety, fear of judgement, uncertainty about...

<https://arxiv.org/abs/2609.21600>

0 CONF

Accelerating Dense LLMs via L0-regularized Mixture-of-Experts

ArXiv CS.AI 50%

arXiv:2609.21672v1 Announce Type: new Abstract: Large language models (LLMs) achieve strong performance but suffer from slow and costly inference. Existing acceleration methods often lead to noticeable performance degradation, while Mixture-of-Experts (MoE) models require...

<https://arxiv.org/abs/2609.21672>

1 CONF

ECG Mirage: Revealing and Mitigating the Underutilisation of ECGs in Vision-Language Models for Clinical Prediction

ArXiv CS.AI 50%

arXiv:2609.21755v1 Announce Type: new Abstract: Emergency department (ED) decision-making relies on heterogeneous clinical information, including patient history, vital signs, laboratory results, and electrocardiograms (ECGs). Vision-language models (VLMs) can jointly process...

<https://arxiv.org/abs/2609.21755>

2 **CONF****LLM-Generated Feature Pools for Time Series Anomaly Detection**

ArXiv CS.AI 50%

arXiv:2609.21801v1 Announce Type: new Abstract: We study how far a simple statistical pipeline can go on univariate time series anomaly detection under a strict selection protocol. The method extracts a small pool of statistics over sliding windows, scores each window with a...

<https://arxiv.org/abs/2609.21801>
3 **CONF****MIST: Multimodal Survival Prediction with Genomic-Guided Histology Attention**

ArXiv CS.AI 50%

arXiv:2609.21811v1 Announce Type: new Abstract: Multimodal survival models can combine complementary prognostic information from whole-slide images and genomic profiles, but effective fusion remains challenging amid external cohort shift and computational complexity. To address...

<https://arxiv.org/abs/2609.21811>
4 **CONF****EnterpriseVal: Quantifying the Efficacy, Reliability and Value of Generative AI in the Enterprise**

ArXiv CS.AI 50%

arXiv:2609.21841v1 Announce Type: new Abstract: Frontier language models now produce professional deliverables that expert graders judge to match human work on a substantial share of economically valuable tasks, yet most enterprise GenAI initiatives fail to show a measurable...

<https://arxiv.org/abs/2609.21841>
5 **CONF****dSTAR: Straggler Tolerant and Byzantine Resilient Distributed SGD**

ArXiv CS.AI 50%

arXiv:2412.07151v1 Announce Type: cross Abstract: Distributed model training needs to be adapted to challenges such as the straggler effect and Byzantine attacks. When coordinating the training process with multiple computing nodes, ensuring timely and reliable gradient...

<https://arxiv.org/abs/2412.07151>
6 **CONF****AgenticRL: Agentic Reinforcement Learning with Self-Refinement for Complex UAV Navigation**

ArXiv CS.AI 50%

arXiv:2606.03963v4 Announce Type: replace-cross Abstract: Deep reinforcement learning enables autonomous robots to learn complex navigation tasks, but still relies heavily on time consuming manual reward design and fine tuning. Existing automated reward generation and refinement...

<https://arxiv.org/abs/2606.03963>
7 **CONF****SpaceDiffusion: Over-the-Orbit Diffusion for Space Generate-and-Forward Communications**

ArXiv CS.AI 50%

arXiv:2609.20899v1 Announce Type: cross Abstract: Satellite communications are an essential component of sixth-generation (6G) mobile networks, which provide ubiquitous connectivity for global services. However, the satellite uplink remains a critical bottleneck for ground...

<https://arxiv.org/abs/2609.20899>

8 CONF

Bio-MF: Low-Latency and High-Fidelity EEG-to-fNIRS Cross-Modal Generation for Hybrid Motor-Imagery Brain-Computer Interfaces

ArXiv CS.AI 50%

arXiv:2609.20904v1 Announce Type: cross Abstract: Hybrid motor-imagery brain-computer interfaces (MI-BCIs) combining EEG and fNIRS can outperform EEG-only systems by exploiting complementary electrophysiological and hemodynamic information. To obtain such hybrid information when...

<https://arxiv.org/abs/2609.20904>

9 CONF

Trustworthy FinAInce: Unpacking How AI-Mediated Financial Advice is Judged

ArXiv CS.AI 50%

arXiv:2609.20989v1 Announce Type: cross Abstract: As generative AI is increasingly used as a source of personal financial guidance, understanding how people appraise such advice is important for supporting appropriate reliance. We conducted a randomized vignette experiment with...

<https://arxiv.org/abs/2609.20989>

10 CONF

How Much of a Real Workload Can LLM-Generated GPU Kernels Actually Reach?

ArXiv CS.AI 50%

arXiv:2609.21058v1 Announce Type: cross Abstract: Language models can now write GPU kernels that outperform PyTorch. We evaluate five model configurations on KernelBench level 1 and find that a frontier model produces correct kernels for 91.1% of problems and independently...

<https://arxiv.org/abs/2609.21058>

1 CONF

Geometry of Values: Task Vector Composition for Ethical Preference Alignment in Language Models

ArXiv CS.AI 50%

arXiv:2609.21094v1 Announce Type: cross Abstract: Large Language Models (LLMs) are increasingly deployed in applications that must weigh clashing moral values, yet even strong models exhibit hidden biases and brittle instruction-following across languages. We introduce a...

<https://arxiv.org/abs/2609.21094>

2 CONF

The Stochastic Shift: A New Evaluation Paradigm for Text-to-SQL with AI Operators

ArXiv CS.AI 50%

arXiv:2609.21133v1 Announce Type: cross Abstract: SQL has been augmented with AI operators, enabling modern data analytics platforms to derive insights from both structured and unstructured data. We observe that while current Text-to-SQL systems can successfully generate these...

<https://arxiv.org/abs/2609.21133>

3 CONF

EnSol: an environment-aware graph neural network for molecular solubility prediction

ArXiv CS.AI 50%

arXiv:2609.21151v1 Announce Type: cross Abstract: Molecular solubility directly affects key aspects of molecular development such as reaction feasibility, formulation performance, separation efficiency, and solvent selection. However, experimental measurement across solutes,...

<https://arxiv.org/abs/2609.21151>

4 **CONF****SWE-Proof: Can Language Models Resolve Real-World Issues with Machine-Checked Proofs?**

ArXiv CS.AI 50%

arXiv:2609.21190v1 Announce Type: cross Abstract: Ensuring the correctness of LLM-generated code is a core challenge for modern software engineering. Benchmarks for agentic code generation check correctness with held-out test suites, which are inherently incomplete and...

<https://arxiv.org/abs/2609.21190>
5 **CONF****Fewer Steps, Better Actions: Rethinking Flow-Matching Inference for VLA Policies**

ArXiv CS.AI 50%

arXiv:2609.21216v1 Announce Type: cross Abstract: Vision-language-action (VLA) policies based on flow matching generate action chunks through repeated evaluations of an action expert. Increasing the number of integration steps raises inference cost, but does not necessarily...

<https://arxiv.org/abs/2609.21216>
6 **CONF****Beyond Exact Match: Task-Aware GRPO for Cross-Domain PCBA Visual Question Answering**

ArXiv CS.AI 50%

arXiv:2609.21276v1 Announce Type: cross Abstract: In automated Printed Circuit Board Assembly (PCBA) inspection, standards-guided decisions require systems to jointly reason over fine-grained visual cues, component semantics, and manufacturing knowledge. Although large...

<https://arxiv.org/abs/2609.21276>
7 **CONF****Authorization Revocation for Long-Running AI Agents: Root-Scoped Quiescence under Delegation and Asynchronous Execution**

ArXiv CS.AI 50%

arXiv:2609.21284v1 Announce Type: cross Abstract: Long-running AI agents outlive initiating processes through credentials, delegated tasks, queues, callbacks, reservations, and provider-side operations. Cancellation, process exit, and credential revocation neither close every...

<https://arxiv.org/abs/2609.21284>
8 **CONF****CESBench: Benchmarking Large Language Models on Cryptographic Engineering Security for IoT Devices**

ArXiv CS.AI 50%

arXiv:2609.21344v1 Announce Type: cross Abstract: For Internet of Things (IoT) devices, a secure algorithm alone is not enough: an attacker with physical access can attack the implementation directly, and its flaws are hard to fix once deployed. Large language models (LLMs) are...

<https://arxiv.org/abs/2609.21344>
9 **CONF****AgentVidBench: A Multi-Hop Video Question Answering Benchmark for Evaluating MLLM Agents**

ArXiv CS.AI 50%

arXiv:2609.21386v1 Announce Type: cross Abstract: Comprehensive video understanding is crucial for advancing artificial intelligence toward the intricate dynamics of the physical world. While recent advances in Multimodal Large Language Models (MLLMs) have demonstrated...

<https://arxiv.org/abs/2609.21386>

0 **CONF****Talking Past the Machine: Morality, Politeness, and Alignment in Human-AI Dialogue****ArXiv CS.AI 50%**

arXiv:2609.21401v1 Announce Type: cross Abstract: Conversational AI systems produce fluent, socially appropriate responses, yet whether they participate in cooperative communication or merely simulate its surface forms remains unclear - a question central to how these systems...

<https://arxiv.org/abs/2609.21401>

1 **CONF****AtomEgo: Exploring Ego-Robot Integration for Embodied Foundation Model Pretraining****ArXiv CS.AI 50%**

arXiv:2609.21461v1 Announce Type: cross Abstract: Embodied foundation models are constrained by the limited scale and diversity of robot demonstrations, motivating the use of large-scale egocentric human interaction data. However, how to effectively incorporate such data into...

<https://arxiv.org/abs/2609.21461>

2 **CONF****HE-Guardrail: A Homomorphic Guardrail Against Jailbreak Attacks for Encrypted Large Language Model Inference****ArXiv CS.AI 50%**

arXiv:2609.21484v1 Announce Type: cross Abstract: Homomorphic encryption (HE) has emerged as a promising approach to privacy-preserving machine learning (PPML), enabling computation directly over encrypted data. In HE-based PPML, a client submits an encrypted input to the...

<https://arxiv.org/abs/2609.21484>

3 **CONF****VidOmni-Bench: A Benchmark for Fine-Grained Video Understanding via Spatio-Temporal Event Verification across Complexity and Duration****ArXiv CS.AI 50%**

arXiv:2609.21521v1 Announce Type: cross Abstract: While Video Large Language Models (Video-LLMs) have recently demonstrated strong performance, reliably evaluating their fine-grained video understanding remains challenging. Existing benchmarks often rely on question answering or...

<https://arxiv.org/abs/2609.21521>

4 **CONF****Micro-Collaborative Poisoning: A Distributed Attack on RAG Systems****ArXiv CS.AI 50%**

arXiv:2609.21573v1 Announce Type: cross Abstract: Retrieval-Augmented Generation (RAG) improves large language models by grounding outputs in external knowledge sources, but this dependency also creates a surface for poisoning attacks. This paper introduces Micro-Collaborative...

<https://arxiv.org/abs/2609.21573>

5 **CONF****Prompt-Driven Exploration: Language as an Exploration Space for VLA Reinforcement Learning****ArXiv CS.AI 50%**

arXiv:2607.08837v4 Announce Type: replace-cross Abstract: Exploration is essential to RL since a policy cannot improve by repeatedly sampling the behaviors it already prefers. Standard methods inject stochasticity in the action space, but such jitter only yields rollouts close...

<https://arxiv.org/abs/2607.08837>

6 CONF

Steering LLMs Responses Towards Moral Foundations on the Norwegian MFQ-30

ArXiv CS.AI 50%

arXiv:2609.21636v1 Announce Type: cross Abstract: Recent work applies human psychometric questionnaires to large language models to elicit moral and value profiles, but it is not clear whether these instruments measure anything stable in models or whether the resulting profiles...

<https://arxiv.org/abs/2609.21636>

7 CONF

Chinese Competitive Debating Dataset and Benchmark

ArXiv CS.AI 50%

arXiv:2609.21637v1 Announce Type: cross Abstract: Debate adjudication requires tracking how arguments develop through interaction, yet existing datasets rarely combine fine-grained debate transcripts with professional judgments collected during real competitions under a shared...

<https://arxiv.org/abs/2609.21637>

8 CONF

SynthDemo-RL: Breaking the Zero-Reward Barrier in VLA Adaptation with LLM-Guided Synthetic Demonstrations

ArXiv CS.AI 50%

arXiv:2609.21650v1 Announce Type: cross Abstract: Fine-tuning Vision-Language-Action (VLA) models commonly relies on human teleoperation demonstrations, while reinforcement learning (RL) with sparse binary rewards faces an exploration challenge when successful trajectories are...

<https://arxiv.org/abs/2609.21650>

9 CONF

When Steering Fails in Latent Reasoning: A Latent-to-Language Transition Gap

ArXiv CS.AI 50%

arXiv:2609.21662v1 Announce Type: cross Abstract: Activation steering has become a widely used approach for controlling language models during explicit chain-of-thought (CoT) reasoning, motivating its extension to latent CoT. However, we find that steering continuous thoughts...

<https://arxiv.org/abs/2609.21662>

10 CONF

TERMon: Detecting Persistent Behavioral Threats in Edge AI via Hardware-Native Ternary Runtime Monitor

ArXiv CS.AI 50%

arXiv:2609.21713v1 Announce Type: cross Abstract: Edge AI accelerators are increasingly deployed in safety-critical environments, where model outputs may control physical actuators, make access-control decisions, or trigger alarms. In these settings, runtime failures often...

<https://arxiv.org/abs/2609.21713>

11 CONF

Balanced Prompt Adaptation against Entropy-Induced Collapse for Test-Time Binary Segmentation

ArXiv CS.AI 50%

arXiv:2609.21743v1 Announce Type: cross Abstract: Entropy minimization is a standard objective for test-time adaptation (TTA), but it can fail in imbalanced binary segmentation. Unlike image classification, dense segmentation aggregates thousands of pixel predictions, allowing...

<https://arxiv.org/abs/2609.21743>

2 **CONF****Do Personality-Tuned LLMs Make Better Social Agents?**

ArXiv CS.AI 50%

arXiv:2609.21857v1 Announce Type: cross Abstract: LLMs are increasingly used in social simulations for socially interactive agents and robots, offering more flexibility than rule-based systems. However, even though they mimic human behaviour very well, there is a persistent...

<https://arxiv.org/abs/2609.21857>

3 **CONF****When Should a Failing Robot Ask? Initiating Corrective Human-Robot Dialogue from Audited Sensor Evidence**

ArXiv CS.AI 50%

arXiv:2609.21942v1 Announce Type: cross Abstract: A robot that fails at a task faces the first decision in corrective dialogue: act on its own diagnosis, consult another onboard sensor, or interrupt a person. Choosing well requires knowing how much the robot's sensors reveal...

<https://arxiv.org/abs/2609.21942>

4 **CONF****Bayesian Belief Layer for Controllable Opinion Dynamics in LLM Agents**

ArXiv CS.AI 50%

arXiv:2609.21997v1 Announce Type: cross Abstract: LLM agents in social simulation revise their opinions implicitly, in context: how open an agent is to persuasion can neither be specified nor verified, and collective outcomes inherit the model's training prior. We introduce...

<https://arxiv.org/abs/2609.21997>

5 **CONF****Gricea: An Open Science Platform for Conversational AI Research**

ArXiv CS.AI 50%

arXiv:2609.22039v1 Announce Type: cross Abstract: We need studies on conversational AI (CAI) at scale to understand human behavior and shape CAI design. However, fragmented reporting of systems and study configurations hinders replication, extension, and knowledge accumulation....

<https://arxiv.org/abs/2609.22039>

6 **CONF****Value-Sensitive Delegation in Everyday AI Agent Use: Evidence from OpenClaw**

ArXiv CS.AI 50%

arXiv:2609.22067v1 Announce Type: cross Abstract: Users increasingly delegate work to autonomous AI agents, yet evaluations typically measure task completion rather than the values users prioritize. Using Value Sensitive Design, we analyzed, with LLM assistance, 73,093...

<https://arxiv.org/abs/2609.22067>

7 **CONF****Representation Before Training: A Practical Benchmark for Generative Medical Event Model Tokenization**

ArXiv CS.AI 50%

arXiv:2604.16775v3 Announce Type: replace-cross Abstract: Generative medical event models use tokenized sequences of patient timelines as input, but practical guidance on the many decisions around tokenization is limited. We benchmark quantization granularity, reference-range...

<https://arxiv.org/abs/2604.16775>

8 **CONF****Planetary Prediction Engine: Autonomous Geospatial Prediction via Intelligent Data Selection and Foundation Model Embeddings****ArXiv CS.AI** **50%**

arXiv:2608.26088v2 Announce Type: replace Abstract: Addressing critical global challenges, from food security and disaster risk to disease outbreaks and socio-economic vulnerability, demands high-fidelity geospatial modeling. However, building predictive planetary models remains...

<https://arxiv.org/abs/2608.26088>

9 **CONF****Balance of Benchmarks: Semantic Density Reweighting for Task-Conditioned Model Comparison****ArXiv CS.AI** **50%**

arXiv:2608.30044v3 Announce Type: replace Abstract: Model comparison increasingly relies on large collections of publicly reported benchmark scores, yet common aggregation strategies trade off evidence coverage against control over capability weighting. Manually curated suites...

<https://arxiv.org/abs/2608.30044>

0 **CONF****MOSCOPT: Mixture-of-Skills Collective Optimization for LLM Agents****ArXiv CS.AI** **50%**

arXiv:2609.14399v2 Announce Type: replace Abstract: Natural language prompts and skills serve as the strategic backbone of LLM-based agents. Recent advances in prompt and skill optimization have achieved notable gains, yet all existing methods optimize a \emph{single} text...

<https://arxiv.org/abs/2609.14399>

1 **CONF****NeuSOGA3D: A Neuro-Symbolic Framework for Explainable 3D Geometric Reconstruction****ArXiv CS.AI** **50%**

arXiv:2609.20323v2 Announce Type: replace Abstract: Three-dimensional reconstruction from unorganized point clouds remains a challenging problem in computer vision, geometric modeling, and computer-aided design. While neural implicit methods achieve impressive reconstruction...

<https://arxiv.org/abs/2609.20323>

2 **CONF****Continuous Spiking Graph Neural Networks****ArXiv CS.AI** **50%**

arXiv:2404.01897v3 Announce Type: replace-cross Abstract: Continuous graph neural networks (CGNNs) have garnered significant attention due to their ability to generalize existing discrete graph neural networks (GNNs) by introducing continuous dynamics. They typically draw...

<https://arxiv.org/abs/2404.01897>

3 **CONF****Benchmarking Autonomous Driving Planners Across Leaderboards: A Unified CARLA-Based Evaluation****ArXiv CS.AI** **50%**

arXiv:2509.22754v2 Announce Type: replace-cross Abstract: Autonomous driving remains a highly active research domain that seeks to enable vehicles to perceive dynamic environments, predict the future trajectories of traffic agents such as vehicles, pedestrians, and cyclists and...

<https://arxiv.org/abs/2509.22754>

4 **CONF****Auditing a KB Elicitation of Frontier LLM Knowledge: A Multi-dimensional Analysis of GPTKB v1.5****ArXiv CS.AI 50%**

arXiv:2510.07024v3 Announce Type: replace-cross Abstract: LLMs are remarkable artifacts that have revolutionized a range of knowledge-intensive tasks. A significant contributor is their factual knowledge, which, to date, remains poorly understood, and is usually analyzed from...

<https://arxiv.org/abs/2510.07024>

5 **CONF****The Impact of Semantic Pairs on Self-Supervised Representation Learning****ArXiv CS.AI 50%**

arXiv:2510.08722v4 Announce Type: replace-cross Abstract: Instance discrimination learns visual representations by treating different augmented views of the same image as positive pairs. While this encourages invariance to handcrafted transformations, same-image positives can...

<https://arxiv.org/abs/2510.08722>

6 **CONF****Deep Learning-Enhanced Real-Time Wi-Fi Sensing Through Single Transceiver Pair****ArXiv CS.AI 50%**

arXiv:2511.02845v3 Announce Type: replace-cross Abstract: The advancement of next-generation Wi-Fi technology heavily relies on sensing capabilities, which play a pivotal role in enabling sophisticated applications. In response to the growing demand for large-scale deployments,...

<https://arxiv.org/abs/2511.02845>

7 **CONF****Large Language Models As Shannon Lossy Compressors Not Solomonoff Induction Estimators: The Singularity Is Not Near Without Symbolic Model Synthesis****ArXiv CS.AI 50%**

arXiv:2601.05280v5 Announce Type: replace-cross Abstract: On the one hand, the question of whether Large Language Models (LLMs) are Solomonoff induction estimators has become an explicit question at the intersection of Algorithmic Information Theory (AIT) and Machine Learning...

<https://arxiv.org/abs/2601.05280>

8 **CONF****MENASpeechBank: A Reference Voice Bank with Persona-Conditioned Multi-Turn Conversations for AudioLLMs****ArXiv CS.AI 50%**

arXiv:2602.07036v2 Announce Type: replace-cross Abstract: Audio large language models (AudioLLMs) enable instruction following over speech and general audio, but progress is limited by the scarcity of diverse, conversational, and instruction-aligned speech-text data. This gap...

<https://arxiv.org/abs/2602.07036>

9 **CONF****Taming the Adversary: A Cost-to-Disturbance Ratio Approach to Adversarial Reinforcement Learning****ArXiv CS.AI 50%**

arXiv:2603.12110v2 Announce Type: replace-cross Abstract: Reinforcement learning (RL) policies trained in simulation often degrade once deployed on real systems, where the controller must reject external disturbances that were never encountered in simulation. Robust RL addresses...

<https://arxiv.org/abs/2603.12110>

0 **CONF****Constraint Decay: The Fragility of LLM Agents in Backend Code Generation****ArXiv CS.AI 50%**

arXiv:2605.06445v2 Announce Type: replace-cross Abstract: Large Language Model (LLM) agents demonstrate strong performance in autonomous code generation under loose specifications. However, production-grade software requires strict adherence to structural constraints, such as...

<https://arxiv.org/abs/2605.06445>1 **CONF****GeoSelect: Spatial-Program Execution for Training-Free Referring Remote Sensing Image Segmentation****ArXiv CS.AI 50%**

arXiv:2607.03869v2 Announce Type: replace-cross Abstract: Referring remote sensing image segmentation segments the object named by a natural-language expression in an aerial image. Existing training-free methods resolve the expression through implicit vision-language activations...

<https://arxiv.org/abs/2607.03869>2 **CONF****Teaching LLMs to Self-Evolve: Cultivating Core Meta-Skills with Reinforcement Learning****ArXiv CS.AI 50%**

arXiv:2607.21971v2 Announce Type: replace-cross Abstract: Test-time scaling through iterative self-evolution with environment feedback, as demonstrated by AlphaEvolve, shows remarkable performance gains. We hypothesize that the success of such evolution frameworks hinges on...

<https://arxiv.org/abs/2607.21971>3 **CONF****A Training-Free Proactive Defense Against Partial Speech Manipulation via Self-Embedding Steganography****ArXiv CS.AI 50%**

arXiv:2608.25285v2 Announce Type: replace-cross Abstract: Partial deepfake speech, where only limited segments of an utterance are synthesized or manipulated, poses a significant challenge to existing deepfake detection systems. As the proportion of spoofed regions decreases,...

<https://arxiv.org/abs/2608.25285>4 **CONF****PAVE: Predictive Alignment and Value-Guided Evolution for World-Action Policies****ArXiv CS.AI 50%**

arXiv:2608.30378v3 Announce Type: replace-cross Abstract: Direct vision-language-action policies generate continuous robot actions efficiently, but standard behavior cloning leaves two complementary gaps: their representations are not explicitly required to describe how the...

<https://arxiv.org/abs/2608.30378>5 **CONF****PentestChain: A Cost-Aware, MCP-Orchestrated Framework for Automated Penetration Testing with Free-Tier LLMs****ArXiv CS.AI 50%**

arXiv:2609.18120v2 Announce Type: replace-cross Abstract: AI-driven penetration testing has been demonstrated with premium frontier models such as GPT-4, but the per-engagement token cost makes continuous, automated testing unaffordable for the smaller organisations that need it...

<https://arxiv.org/abs/2609.18120>

6 CONF

PACE: Precise AI Cinematic Expression

ArXiv CS.AI 50%

arXiv:2609.19853v2 Announce Type: replace-cross Abstract: Between a screenplay and a film sits a planning problem that is spatial first: who stands where, and what a camera sees from where it stands. An image diffusion model asked for a shot in free text settles that plan by its...
<https://arxiv.org/abs/2609.19853>