



Pipeline

Scrape HackerNews (Show/Top/New), Reddit (r/artificial, r/MachineLearning, r/singularity, r/LocalLLaMA), TechCrunch AI, VentureBeat AI, MIT Tech Review, The Verge AI, Wired AI, ArXiv CS.AI

Filter 40+ AI keyword patterns — LLMs, alignment, safety, research, tooling

Dedup Title-similarity (45% threshold) merges cross-source variants

Score Authenticity = 0.5 base + 0.3 per extra source + diversity + HN score bonuses

Rank Sorted by authenticity. Below 0.25 threshold are dropped.

Authenticity Tiers

- VERIFIED 0.80 - 1.00** 3+ independent sources including media
- CONFIRMED 0.50 - 0.79** 2 sources or 1 high-quality media
- EMERGING 0.25 - 0.49** Single-source; passes AI keyword filter

Source Breakdown

Source	Articles	Category
TechCrunch AI	8	Media
HackerNews	7	Community
ArXiv CS.AI	4	Media
The Verge AI	4	Media
MIT Tech Review	3	Media

VERIFIED INTELLIGENCE — 2 stories

VERI

Show HN: Less Prompts, More Guardrails

HackerNews

+1 source

90%

HN 3

arXiv:2603.17839v4 Announce Type: replace-cross Abstract: Verbal confidence -- prompting LLMs to state their confidence as a number or category -- is widely used to extract uncertainty estimates from black-box models. However, how LLMs internally generate such scores remains...

<https://yasyf.com/writing/less-prompts-more-guardrails/>

VERI

Runtime Authorization for Resources Acquired by AI Agents

ArXiv CS.AI

+1 source

80%

arXiv:2609.14744v2 Announce Type: replace Abstract: By acquiring compute, credentials, accounts, services, and other agents, autonomous AI agents can introduce new authority into a task. Payment, budget, OAuth, mandate, and fulfillment checks can validate transaction conditions...

<https://arxiv.org/abs/2609.14744>

CONFIRMED SIGNALS — 22 stories

CONF

Show HN: Agent Chaperone - Screen AI agent tool calls and results with Jev

HackerNews

50%

HN 4

Hi everyone. TypeSafe released Jev last week, and it pushed me to build agent-chaperone. It is a screening layer for AI agent tool calls and tool results. It screens risky, off-task, or secret-leaking actions before they run, and prompt injections or other suspicious...

<https://github.com/agent-chaperone/agent-chaperone>

CONF

Show HN: AI Changed How We Write Code. Why Hasn't Code Review Changed?

HackerNews

50%

HN 3

No summary — see link for full article.

<https://github.com/otobongfp/code-graph-view>

CONF

Show HN: PokerTools Arena - Local AI vs. AI Poker LLM Benchmark Table

HackerNews

50%

HN 3

No summary — see link for full article.

<https://github.com/pokertools-arena/pokertools-arena.github.io>

CONF

Show HN: Nanosamurai - Ollama for Voice Models (STT)

HackerNews

50%

HN 3

nanosamurai is a speech-to-text platform that supports different models for realtime, semi-realtime, and batch transcription and provides a unified stack for robust speech processing, agentic workflows, and webhooks. You can run it locally on your computer via electron app +...

<https://github.com/nanosamurai/nanosamurai>

CONF

Show HN: Primescript - Academic Markdown editor with vector search and LaTeX

HackerNews

50%

HN 2

No summary — see link for full article.

<https://primescript.com>

CONF

Show HN: Maki, an open-source multi-agent LLM framework (local or hosted)

HackerNews 50% HN 2

No summary — see link for full article.

<https://github.com/BowlOfData/maki>

CONF

The man who built Apple's stores doesn't buy Silicon Valley's bet on AI shopping

TechCrunch AI 50%

Apple Store architect Ron Johnson says Apple's secret sauce has always been its people.

<https://techcrunch.com/2026/09/21/the-man-who-built-apples-stores-doesnt-buy-silicon-va...>

0 CONF

OpenAI forms math advisory group as its AI resolves more than 100 open problems

TechCrunch AI 50%

The group won't be given leeway to slow down or redirect OpenAI's ongoing mathematical research.

<https://techcrunch.com/2026/09/21/openai-forms-math-advisory-group-as-its-ai-resolves-m...>

1 CONF

Meta's Muse is outpacing ChatGPT's early mobile launch

TechCrunch AI 50%

Meta's new AI agent Muse has racked up more downloads and daily active users in the U.S. and Canada than ChatGPT did over the same period after its mobile debut, according to new estimates from Appfigures.

<https://techcrunch.com/2026/09/21/metas-muse-is-outpacing-chatgpts-early-mobile-launch/>

2 CONF

Meta's AI agent has been blocked from using Amazon.com

TechCrunch AI 50%

Amazon has its own cohort of foundation models, along with one of the most popular inference platforms on the internet. As long as they're under no legal obligation to open the doors to Muse, why would they?

<https://techcrunch.com/2026/09/21/metas-ai-agent-has-been-blocked-from-using-amazon-com/>

3 CONF

With Tabby, a former accountant is using AI to make accountants obsolete

TechCrunch AI 50%

Tabby is designed to be a real-time bookkeeping interface, handling clients' paperwork as it gives them up-to-the-minute data on their business's profit and loss.

<https://techcrunch.com/2026/09/21/with-tabby-a-former-accountant-is-using-ai-to-make-ac...>

4 CONF

Where will the next breakout startup come from? Benchmark's full partnership weighs in at TechCrunch Disrupt 2026

TechCrunch AI 50%

Where will the next breakout startup come from? Benchmark's full partnership weighs in on the main stage at TechCrunch Disrupt 2026. Save up to \$200 before September 25 at 11:59 p.m. PT. Register now.

<https://techcrunch.com/2026/09/21/where-will-the-next-breakout-startup-come-from-benchm...>

5 CONF

Google's \$899 Googlebook is a bet that you'll buy a new laptop for Gemini

TechCrunch AI 50%

Google's AI-native Googlebook ties Gemini to the cursor, dictation, widgets, and other parts of the desktop experience.

<https://techcrunch.com/2026/09/21/googles-899-googlebook-is-a-bet-that-youll-buy-a-new-...>

6 CONF

Is the AI industry really ready to slow down?

TechCrunch AI 50%

On Equity, we debated whether AI executives are serious about wanting to slow down.

<https://techcrunch.com/2026/09/20/is-the-ai-industry-really-ready-to-slow-down/>

7 CONF

4 ways to address the failures we found along the US border's "virtual wall"

MIT Tech Review 50%

MIT Technology Review today published our investigation into how many people have died near the "virtual wall" of surveillance towers that the US government has installed along the US-Mexico border. We found cases of people who walked undetected through areas surveilled by...

<https://www.technologyreview.com/2026/09/21/1144164/border-towers-surveillance-policy-r...>

8 CONF

She died at the San Diego border. A surveillance camera was in plain sight

MIT Tech Review 50%

She had only walked for a couple of hours, and already she was lost. It was early afternoon on Sept. 14, 2025, when 30-year-old Graciela Gómez Hernández crossed the border from the eastern edge of Tijuana into Southern California, sending voice messages to her mother and sister...

<https://www.technologyreview.com/2026/09/21/1144170/border-towers-surveillance-feature/>

9 CONF

Could AI really kill us all? Your questions, answered.

MIT Tech Review 50%

On Wednesday, MIT Technology Review hosted a live Roundtables event for subscribers that asked the question everyone's asking right now: Could AI really kill us all? But attendees had so many more questions than we had time to answer in the 30 minute session. So we asked our...

<https://www.technologyreview.com/2026/09/18/1144435/could-ai-really-kill-us-all-your-qu...>

0 CONF

California tightens rules on AI data center energy and water use

The Verge AI 50%

California Gov. Gavin Newsom has signed seven bills designed to prevent AI data centers from passing utility costs onto residents, as reported earlier by the Los Angeles Times. The package of laws requires the California Public Utilities Commission to introduce a new rate...

<https://www.theverge.com/ai-artificial-intelligence/998453/california-ai-data-center-bills>

1 CONF

iPhone owners can now submit claims in Apple's \$250 million Siri AI settlement

The Verge AI 50%

Apple is paying \$250 million to settle claims that it failed to deliver an AI-upgraded Siri - and now, eligible iPhone owners can submit a claim for a payout. If you live in the US and purchased an iPhone 15 Pro, iPhone 15 Pro Max, or any iPhone 16 model between June 10th, 2024,...

<https://www.theverge.com/tech/998191/apple-siri-ai-iphone-16-class-action-lawsuit-settl...>

2 CONF

No one is surprised that Nvidia's Jensen Huang thinks AI fears are overblown

The Verge AI 50%

The man who may stand to make the most money from the AI boom seems to think he knows better than anyone else, including researchers who have studied and worked on AI for decades. In an interview with CBS Sunday Morning, he claimed there was a "0% chance" of AI being the end of...

<https://www.theverge.com/ai-artificial-intelligence/997936/nvidia-jensen-huang-ai-fears...>

3 CONF

World Modeling in Transformers

ArXiv CS.AI 50%

arXiv:2609.21748v1 Announce Type: new Abstract: Behavioral failures can make a transformer appear to lack a world model even when it has learned faithful representations of its environment. We demonstrate this in TaxiGPT, a transformer trained on random walks through Manhattan...
<https://arxiv.org/abs/2609.21748>

4 CONF

A Lie Detector Test for Language Models: Reading Knowledge a Model Won't Reveal

ArXiv CS.AI 50%

arXiv:2609.21996v1 Announce Type: new Abstract: Large language models can hold knowledge they do not report. A model may sandbag on a capability evaluation, or answer against what it internally knows, and its outputs alone cannot tell whether it is hiding an answer or simply...
<https://arxiv.org/abs/2609.21996>