



Pipeline

Scrape HackerNews (Show/Top/New), Reddit (r/artificial, r/MachineLearning, r/singularity, r/LocalLLaMA), TechCrunch AI, VentureBeat AI, MIT Tech Review, The Verge AI, Wired AI, ArXiv CS.AI

Filter 40+ AI keyword patterns — LLMs, alignment, safety, research, tooling

Dedup Title-similarity (45% threshold) merges cross-source variants

Score Authenticity = 0.5 base + 0.3 per extra source + diversity + HN score bonuses

Rank Sorted by authenticity. Below 0.25 threshold are dropped.

Authenticity Tiers

- VERIFIED 0.80 - 1.00** 3+ independent sources including media
- CONFIRMED 0.50 - 0.79** 2 sources or 1 high-quality media
- EMERGING 0.25 - 0.49** Single-source; passes AI keyword filter

Source Breakdown

Source	Articles	Category
ArXiv CS.AI	195	Media
TechCrunch AI	9	Media
HackerNews	6	Community
The Verge AI	6	Media
MIT Tech Review	3	Media

VERIFIED INTELLIGENCE — 4 stories

VERI

Show HN: JevBench, a reproducible benchmark for typed decision models

HackerNews

+1 source

100%

HN 102

Hi HN! I built JevBench because Jev kicks ass, and the world deserves to know how the serious open source and fake lookalike projects really perform in comparison. Jev-class models return bounded choices and probabilities instead of text, and are disruptively faster and cheaper...

<https://benchmarkheaven.com/jev-models>

VERI

\$t_0\$: A Time-Series Foundation Model for Forecasting with Context

ArXiv CS.AI

+1 source

80%

arXiv:2609.24559v1 Announce Type: cross Abstract: We present t_0 , a family of open-weights foundation models for forecasting with multivariate context. We release its first two members: $t_0\text{-}\alpha$ and $t_0\text{-}\beta$, respectively 102M and 256M parameters. Both...

<https://arxiv.org/abs/2609.24559>

VERI

Anthropic launches Claude Opus 5.5 with stricter safeguards for cybersecurity

The Verge AI

+1 source

80%

Anthropic says its new Claude Opus 5.5 model comes with stronger safeguards in the wake of recent rogue AI hacking incidents. In an announcement on Tuesday, Anthropic says Opus 5.5 comes with improvements to certain risky behaviors, including attempts to escape the company's...

<https://www.theverge.com/ai-artificial-intelligence/998868/anthropic-claude-opus-5-5-cy...>

VERI

Meta patches Muse exploit that let attackers control the AI agent

The Verge AI

+1 source

80%

Meta has issued a patch for its Muse macOS app following the discovery of a zero-day vulnerability that could allow someone to take control of the AI agent. The bug found by security researcher Patrick Wardle utilized an undocumented Muse setting that enabled potential attackers...

<https://www.theverge.com/tech/998679/meta-muse-patch-zero-day-exploit-ai-agent>

CONFIRMED SIGNALS — 211 stories

CONF

Show HN: Hack - my Hacker News mobile client now lets you block AI posts

HackerNews

50%

HN 7

Hello HN! I am the developer for HACK, a hacker news client that I have been building and maintaining for over 8 years now on iOS, iPadOS, MacOS and Android. Wanted to share this news as there's been a ton of interest in these features recently. Now, you can block AI related...

<https://news.ycombinator.com/item?id=49802005>

CONF

Show HN: WavexAI - Unlimited tokens for a fixed cost

HackerNews

50%

HN 6

No summary — see link for full article.

<https://wavexai.dev/>

CONF

Show HN: Venya lets AI agents use secrets without seeing them

HackerNews

50%

HN 3

No summary — see link for full article.

<https://github.com/tabith-llc/venya>

CONF

Show HN: Sbm: plain-text bookmarks synced across terminal, browser, Android

HackerNews 50% HN 2

No summary — see link for full article.

<https://sbmsync.com>

CONF

Show HN: Fusion-runtime - self-hosted voice agents, STT+LLM+TTS in one process

HackerNews 50% HN 2

No summary — see link for full article.

<https://github.com/SamarthUrs18/fusion-runtime>

0 CONF

'We're already fighting yesterday's battle': Greece's prime minister gets candid about AI

TechCrunch AI 50%

Most leaders on a trade mission stick to the pitch, but when I interviewed Greek Prime Minister Kyriakos Mitsotakis this week, he also admitted that no government is ready for what AI is about to do.

<https://techcrunch.com/2026/09/22/were-already-fighting-yesterdays-battle-greeces-prime...>

1 CONF

TechCrunch Founder Summit's agenda revealed: Unlock fundraising, hiring, and AI insights in Boston on November 4

TechCrunch AI 50%

Founders shouldn't have to learn the hardest lessons the hardest way. TechCrunch Founder Summit is designed to make the challenges of starting a company easier and the highs that much greater.

<https://techcrunch.com/2026/09/22/techcrunch-founder-summits-agenda-revealed-unlock-fun...>

2 CONF

Qualcomm launches two new smartphone chips with emphasis on AI

TechCrunch AI 50%

Qualcomm said that its new top chip can run 30B mixture-of-expert model locally.

<https://techcrunch.com/2026/09/22/qualcomm-launches-two-new-smartphone-chips-with-empha...>

3 CONF

OpenAI launches GPT-6 Sol and Luna, boasting lower cost and fewer mistakes

TechCrunch AI 50%

OpenAI is launching two new models, which it says are cut from the same cloth as Astra.

<https://techcrunch.com/2026/09/22/openai-launches-gpt-6-sol-and-luna/>

4 CONF

AstroForge is putting AI in command of its next spacecraft

TechCrunch AI 50%

Autonomy-1 will have a small, transformer-based AI model taking charge of a space probe.

<https://techcrunch.com/2026/09/22/astroforge-is-putting-ai-in-command-of-its-next-space...>

5 CONF

Five AI safety sessions every founder should have on their TechCrunch Disrupt 2026 agenda

TechCrunch AI 50%

At TechCrunch Disrupt 2026, five sessions across the AI Stage and Real World AI Stage cover AI safety, featuring leaders from Anthropic, Nvidia, AWS, Waabi, and more. Register before September 25 to save up to \$200.

<https://techcrunch.com/2026/09/22/five-ai-safety-sessions-every-founder-should-have-on-...>

6 CONF

Nscale's IPO will test Wall Street's appetite for concentrated AI bets once again

TechCrunch AI 50%

The British AI data center developer depends on tech giants Microsoft and Anthropic for most of its revenue.

<https://techcrunch.com/2026/09/22/nscales-ipo-will-test-wall-streets-appetite-for-conce...>

7 CONF

Roundtables: The Deadly Failures of The Virtual Border Wall

MIT Tech Review 50%

The US has spent billions building a "virtual wall" of surveillance towers along its southern border over the past 25 years, promising they will help detect and apprehend border crossers and save lives. But a groundbreaking investigation by MIT Technology Review has documented...

<https://www.technologyreview.com/2026/09/22/1144890/roundtables-the-deadly-failures-of-...>

8 CONF

The Download: why AI's latest breakthroughs and fears may be more hype than reality

MIT Tech Review 50%

This is today's edition of The Download, our weekday newsletter that provides a daily dose of what's going on in the world of technology. Don't be fooled by this summer of AI hype —Timnit Gebru, executive director of the Distributed AI Research Institute (DAIR), and Emily M....

<https://www.technologyreview.com/2026/09/22/1144910/the-download-dont-believe-ai-hype/>

9 CONF

Don't be fooled by this summer of AI hype

MIT Tech Review 50%

It's been a busy few months for AI hype. At the end of April, Anthropic claimed that its model Claude Mythos is better at finding software vulnerabilities than most security experts. Then we had the OpenAI-Hugging Face hacking incident, after which Anthropic (proudly) and Meta...

<https://www.technologyreview.com/2026/09/22/1144867/dont-be-fooled-summer-ai-hype/>

10 CONF

OpenAI wants to consult elite mathematicians about how to not fumble again

The Verge AI 50%

After turning a string of spectacular mathematical results into a reputational crisis, OpenAI is consulting human mathematicians to help it figure out a less disastrous path forward. On Monday, the company announced a new independent panel of mathematicians tasked with advising...

<https://www.theverge.com/ai-artificial-intelligence/999167/openai-elite-mathematicians-...>

1 CONF

Rabbit's new AI agent doesn't need an R1 to run

The Verge AI 50%

Rabbit, the company behind the underwhelming R1 device, is rolling out a standalone AI agent that you don't need its hardware to use, as reported earlier by Wired. The startup says its new OS3 "agentic operating system" runs in the cloud but operates locally across Windows, Mac,...

<https://www.theverge.com/ai-artificial-intelligence/999094/rabbit-ai-agent-os3>

2 CONF

Trump says the US is officially renaming AI to 'super intelligence'

The Verge AI 50%

In a speech Tuesday morning at the UN General Assembly, Donald Trump railed against Iran, "globalists," climate change, and transgender people while also claiming that the US is now "officially" renaming artificial intelligence to "super intelligence." Oddly, this wasn't one of...

<https://www.theverge.com/ai-artificial-intelligence/998816/donald-trump-ai-super-intell...>

3 **CONF****UN says AI safeguards can't wait for certainty****The Verge AI** 50%

Governments need to rein in increasingly capable AI agents before their risks are fully understood, a United Nations scientific panel warned in the global organization's first major assessment of OpenAI's hack of Hugging Face earlier this year. The report cements AI's place on...

<https://www.theverge.com/ai-artificial-intelligence/998090/un-ai-panel-hugging-face-hac...>

4 **CONF****Misaligned Clinical Risk Classification and Cost Asymmetry in Open-Weight Large Language Models****ArXiv CS.AI** 50%

arXiv:2609.23999v1 Announce Type: cross Abstract: How large language models (LLMs) integrate patient risk with clinical cost tradeoffs remains poorly understood. We investigated how four open-weight LLMs (Qwen-2.5-7B/32B and Llama-3.1-8B/70B) internally represent cost tradeoffs,...

<https://arxiv.org/abs/2609.23999>

5 **CONF****An Affordable AI-Integrated Smart Cane for Multimodal Mobility Assistance of Visually Impaired Users****ArXiv CS.AI** 50%

arXiv:2609.22277v1 Announce Type: new Abstract: Visual impairment affects over 2.2 billion people worldwide, yet conventional white canes cannot detect elevated hazards or provide semantic environmental context. Existing AI-assisted navigation systems typically rely on expensive...

<https://arxiv.org/abs/2609.22277>

6 **CONF****PAANI : On Device Visual Evidence Fusion and Explainable Guidance for River Robot Simulation****ArXiv CS.AI** 50%

arXiv:2609.22353v1 Announce Type: new Abstract: Mobile river monitoring robots must interpret obstacles and water boundaries that geographic waypoints alone cannot describe. On resource constrained platforms, converting imperfect visual predictions into timely and inspectable...

<https://arxiv.org/abs/2609.22353>

7 **CONF****HumanAgencyBench: Scalable Evaluation of Human Agency Support in AI Assistants****ArXiv CS.AI** 50%

arXiv:2509.08494v2 Announce Type: replace-cross Abstract: As humans delegate more tasks and decisions to artificial intelligence (AI), we risk losing control of our individual and collective futures. Relatively simple algorithmic systems already steer human decision-making, such...

<https://arxiv.org/abs/2509.08494>

8 **CONF****Replication Without Persistence in Hosted LLMs: Measurement Sensitivity in Action-Time Belief Evaluation****ArXiv CS.AI** 50%

arXiv:2609.22478v1 Announce Type: new Abstract: Behavioural evaluations of hosted language models can vary because the evaluated service, the measurement instrument, or both differ across runs. We separate three validation questions: whether a prior finding recurs on fresh data...

<https://arxiv.org/abs/2609.22478>

9 CONF

Agreement Overstates Evidence: Error Dependence in LLM Judge Consensus

ArXiv CS.AI 50%

arXiv:2609.22512v1 Announce Type: new Abstract: Consensus among LLM judges is often taken as strong evidence that a decision is correct. This assumes that judges make their errors independently. In practice, LLM judges are often trained and evaluated in similar ways, so they can...

<https://arxiv.org/abs/2609.22512>

0 CONF

IntLawNER: A Named Entity Recognition Dataset and Benchmark in International Law

ArXiv CS.AI 50%

arXiv:2609.22529v1 Announce Type: new Abstract: International law provides the normative framework through which states coordinate action, regulate armed conflict, and protect human rights, yet its texts remain without token-level named entity recognition (NER) resources. We...

<https://arxiv.org/abs/2609.22529>

1 CONF

DUMA-Bench: A Dual-Control Multi-Agent Benchmark for Evaluating LLM Agent Security

ArXiv CS.AI 50%

arXiv:2609.24662v1 Announce Type: new Abstract: LLM-based agents increasingly operate in environments where they interact with users, tools, and external systems. Yet most security evaluations assume passive users and static control, ignoring the interactive dynamics that shape...

<https://arxiv.org/abs/2609.24662>

2 CONF

GaitVista: Reliability-Aware AI Measurement toward Accessible Longitudinal Gait Assessment

ArXiv CS.AI 50%

arXiv:2609.22619v1 Announce Type: new Abstract: Tracking recovery of walking function requires detecting meaningful gait change across rehabilitation sessions, yet objective 3D measurement remains confined to specialized motion-capture laboratories. Small camera sets and...

<https://arxiv.org/abs/2609.22619>

3 CONF

Splitting Documents at Lower Cost: Multi-Split Boundary Decisions for LLM-Based Page Stream Segmentation

ArXiv CS.AI 50%

arXiv:2609.22620v1 Announce Type: new Abstract: Scanned mail, uploaded PDFs, and consolidated attachments often arrive as page streams that must be split into individual documents before downstream classification, extraction, or routing. Zero-shot large language models can...

<https://arxiv.org/abs/2609.22620>

4 CONF

Text, Pixels, or Both? Evaluating Input Representations for Multimodal Document QA

ArXiv CS.AI 50%

arXiv:2609.22628v1 Announce Type: new Abstract: Every document QA system begins with a choice that is rarely studied on its own: whether to feed the model page images, extracted text, or both. We isolate this choice, holding the prompt, judge, and scoring pipeline fixed, across...

<https://arxiv.org/abs/2609.22628>

5 **CONF****Toward Auditable and Calibrated AI for Dementia-Related Crash Severity Prediction: A Selective Deferral Framework to Support Human Review****ArXiv CS.AI** **50%**

arXiv:2609.22694v1 Announce Type: new Abstract: Public crash databases increasingly support automated safety analysis, but crash severity prediction remains difficult to translate into public-sector decision workflows when models are evaluated primarily as ordinary classifiers....

<https://arxiv.org/abs/2609.22694>

6 **CONF****Building Trustworthy Mental Health Benchmarks on Bluesky: A Validation-Aware Weak-Supervision Framework****ArXiv CS.AI** **50%**

arXiv:2609.22696v1 Announce Type: new Abstract: Decentralized social media platforms create new opportunities and challenges for computational mental health research because data access, moderation, labeling, and deployment responsibilities are distributed across multiple...

<https://arxiv.org/abs/2609.22696>

7 **CONF****Hapi: A Multivariable Land-Surface Transformer for Medium-Range Hydrological Forecasting at Continental Scale****ArXiv CS.AI** **50%**

arXiv:2609.22702v1 Announce Type: new Abstract: Accurate flood forecasts several days in advance are essential for flood control, water-resource management, and emergency response. Producing them at high resolution over a continental domain calls for local hydrological detail...

<https://arxiv.org/abs/2609.22702>

8 **CONF****Trustworthy Agentic AI: Failure Modes, Mitigation Strategies, and a Lifecycle Framework for Autonomous LLM Systems****ArXiv CS.AI** **50%**

arXiv:2609.22712v1 Announce Type: new Abstract: Agentic AI systems built on large language models can plan over multiple steps, use external tools, retain information in memory, and coordinate with other agents. These capabilities make them more useful than static language...

<https://arxiv.org/abs/2609.22712>

9 **CONF****ProcessLight: Process Supervision for Large Language Model Based Traffic Signal Control****ArXiv CS.AI** **50%**

arXiv:2609.22746v1 Announce Type: new Abstract: Large Language Models (LLMs) have recently been introduced into traffic signal control (TSC) as decision agents due to their strengths in human-readable reasoning generation. Yet, existing LLM TSC methods optimize only from final...

<https://arxiv.org/abs/2609.22746>

10 **CONF****From Research Frontier to Laboratory Bench: Design of a Four-Tier Experimental Teaching System for Multimodal Medical Image Intelligent Diagnosis****ArXiv CS.AI** **50%**

arXiv:2609.22790v1 Announce Type: new Abstract: Undergraduate programmes in intelligent medical engineering are expanding, yet laboratory curricula lag behind the multimodal, long-tailed, and distributionally shifting realities of clinical AI. This design paper presents an...

<https://arxiv.org/abs/2609.22790>

1 CONF

PII-TRACE: A Benchmark for Context-Aware PII Detection in Multi-Turn LLM Conversations

ArXiv CS.AI 50%

arXiv:2609.22200v1 Announce Type: cross Abstract: LLM assistants and agentic systems log long multi-turn conversations. AI providers often scan these conversations for Personally Identifiable Information (PII) and mask the PII before storing or processing conversation data. Yet...

<https://arxiv.org/abs/2609.22200>

2 CONF

Beyond Linear Context: Graph-Guided Evidence Navigation for Long-Novel Reasoning with a Local 9B Language Model

ArXiv CS.AI 50%

arXiv:2609.22939v1 Announce Type: new Abstract: Long-context models read a novel the way a person reads a printout: one token after another, in narrative order, with the whole history competing for a fixed budget of attention. A detective does not work that way. They sort what...

<https://arxiv.org/abs/2609.22939>

3 CONF

OptiSkill: A Hierarchical and Evolving SkillBank for LLM-Based Optimization Modeling

ArXiv CS.AI 50%

arXiv:2609.22987v1 Announce Type: new Abstract: Automated operations research (OR) modeling requires LLMs to translate natural-language decision problems into correct mathematical programs. Existing methods can improve individual formulations, but they often solve problems in...

<https://arxiv.org/abs/2609.22987>

4 CONF

Temporal Generalization and Explanation Stability of Control Flow Graph Neural Networks for Malware Detection

ArXiv CS.AI 50%

arXiv:2609.24280v1 Announce Type: cross Abstract: Malware detection is a critical task in cybersecurity, and graph neural networks over control flow graphs have shown promising results for it. However, detectors are usually evaluated on a random split of a corpus collected over...

<https://arxiv.org/abs/2609.24280>

5 CONF

Enforcing Narrative Reliability and Epistemic Pacing in LLM-Driven Detective Games via Structured Knowledge Trees

ArXiv CS.AI 50%

arXiv:2609.23043v1 Announce Type: new Abstract: Large Language Models (LLMs) enable open-ended dialogue in interactive games, but their non-deterministic outputs make it difficult to preserve authorial control, factual consistency, and the intended sequence of information...

<https://arxiv.org/abs/2609.23043>

6 CONF

FireWorldBench: Benchmarking Complex Physical World Intelligence through Coupled-Field Fire Dynamics

ArXiv CS.AI 50%

arXiv:2609.23064v1 Announce Type: new Abstract: Understanding the physical world requires more than object recognition, scene description, and short-term visual prediction, as real-world physical systems involve multiple continuous fields, latent causal mechanisms, partial...

<https://arxiv.org/abs/2609.23064>

7 CONF

Are Human-Aligned Models Models of Humans? A Turing-Test Gap in Preference Alignment

ArXiv CS.AI 50%

arXiv:2609.23640v1 Announce Type: new Abstract: Human-feedback alignment has made language models useful assistants and is commonly described as aligning them with humans. However, the responses people prefer from an AI need not be the responses they themselves would give. We...

<https://arxiv.org/abs/2609.23640>

8 CONF

From Inference Engine to Inference Control Plane: Connecting vLLM, llm-d, and the Evolution of Efficient Distributed LLM Serving

ArXiv CS.AI 50%

arXiv:2609.23130v1 Announce Type: new Abstract: Large language model (LLM) inference is evolving from an engine-local optimization problem into a distributed control problem involving reusable state, phase placement, heterogeneous accelerators, networking, autoscaling,...

<https://arxiv.org/abs/2609.23130>

9 CONF

Do Not Trust the Benchmark: Limitations of General LLM Rankings and a Case for Task-Specific Evaluation

ArXiv CS.AI 50%

arXiv:2609.23201v1 Announce Type: new Abstract: Benchmark scores increasingly influence the development, marketing, and selection of large language models (LLMs). Yet an overall score is interpretable only in relation to the system tested, the questions included, and the...

<https://arxiv.org/abs/2609.23201>

0 CONF

MemCalib: Benchmarking and Optimizing Memory Use in LLM Agents

ArXiv CS.AI 50%

arXiv:2609.24259v2 Announce Type: cross Abstract: The effectiveness of agent memory ultimately depends on whether the underlying LLM gives each memory in context an appropriate degree of influence over its response. Yet this capability has remained largely overlooked. To assess...

<https://arxiv.org/abs/2609.24259>

1 CONF

AgentBetta: Verification-Driven Adaptive Configuration of an AI Nano-Agent through Selective Expansion and Verified Contraction

ArXiv CS.AI 50%

arXiv:2609.23512v1 Announce Type: new Abstract: Large language model agents are typically deployed with predefined configurations, although the required model capability, context, tools, permissions, memory, and computational resources can vary substantially across tasks. This...

<https://arxiv.org/abs/2609.23512>

2 CONF

On Probabilistic Inference Through Parametric Tensor Decomposition in Base Tensor Networks

ArXiv CS.AI 50%

arXiv:2609.23774v1 Announce Type: new Abstract: Probabilistic inference is generally only tractable in low-treewidth graphical models, limiting its effective applicability in high-treewidth settings. Many existing methods improve efficiency by exploiting specific parametric...

<https://arxiv.org/abs/2609.23774>

3 CONF

Total Cost of Agency: Exact Attribution of Memory Injection Cost in Multi-Agent LLM Workflows

ArXiv CS.AI 50%

arXiv:2609.23790v1 Announce Type: new Abstract: Every node in a multi-agent large language model (LLM) workflow retrieves context from memory and injects it into its prompt, where those injected tokens are billed as input tokens at the same per-token price as the system prompt...

<https://arxiv.org/abs/2609.23790>

4 CONF

WorkWorlds: An Infrastructure for Evaluating AI Agents on Workplace Tasks

ArXiv CS.AI 50%

arXiv:2609.23806v1 Announce Type: new Abstract: Many knowledge-work benchmarks are constructed around individual tasks, with the context needed for each task selected together with or after the task has been specified. This design measures performance on workplace-like tasks in...

<https://arxiv.org/abs/2609.23806>

5 CONF

Pretraining of Medical Visual Encoders Toward Multi-modal Large Language Models

ArXiv CS.AI 50%

arXiv:2609.23860v1 Announce Type: new Abstract: Multimodal Large Language Models (MLLMs) commonly reuse visual encoders pretrained with CLIP, although the features of these ViTs are ultimately consumed by autoregressive LLMs. We refer to this mismatch as the semantic-interface...

<https://arxiv.org/abs/2609.23860>

6 CONF

Explainable Recommendations at Scale: LLM Rationales for YouTube Music Artist Discovery

ArXiv CS.AI 50%

arXiv:2609.23877v1 Announce Type: new Abstract: Modern music streaming platforms face a persistent tradeoff: exploiting familiar content versus driving the exploration of novel items. While users frequently desire discovery, they hesitate to select unknown artists over proven...

<https://arxiv.org/abs/2609.23877>

7 CONF

Increasing Skill Level Recruits Deeper Attention Layers in a Frozen Chess Transformer

ArXiv CS.AI 50%

arXiv:2609.23917v1 Announce Type: new Abstract: Chess involves complex reasoning in a deterministic environment, which makes it a useful setting for studying the mechanisms of computation inside transformers. The Maia-3 chess transformer takes Elo, a measure of competitive chess...

<https://arxiv.org/abs/2609.23917>

8 CONF

Echo State Network (ESN) for Signal Recovery in RF-Impaired IBFD MIMO Systems

ArXiv CS.AI 50%

arXiv:2609.23945v1 Announce Type: new Abstract: In-band full-duplex (IBFD) multiple-input multiple-output (MIMO) systems enable simultaneous transmission and reception on the same frequency band, improving spectral efficiency for next-generation wireless networks. However,...

<https://arxiv.org/abs/2609.23945>

9 CONF

UniK: Universal Knowledge Perception for Digital and Physical AI

ArXiv CS.AI 50%

arXiv:2609.23971v1 Announce Type: new Abstract: Two transformative classes of AI systems are reshaping how organizations operate: \textit{digital AI}, which reasons over enterprise knowledge to power chatbots and agent workflows; and \textit{physical AI}, which learns to control...

<https://arxiv.org/abs/2609.23971>

0 **CONF****Jev-Mem: System-One-Controlled Agentic Memory for Efficient AI Agents****ArXiv CS.AI** **50%**

arXiv:2609.23986v1 Announce Type: new Abstract: Agentic memory is becoming essential for long-horizon AI agents, yet many existing systems rely on autoregressive LLMs to control how memories are organized, retrieved, and used, placing expensive generation on the critical path of...

<https://arxiv.org/abs/2609.23986>

1 **CONF****Memory-Free Continual Learning with Null Space Adaptation for Zero-Shot Vision-Language Models****ArXiv CS.AI** **50%**

arXiv:2510.21175v2 Announce Type: replace Abstract: Pre-trained vision-language models (VLMs), such as CLIP, have demonstrated remarkable zero-shot generalization, enabling deployment in a wide range of real-world tasks without additional task-specific training. However, in real...

<https://arxiv.org/abs/2510.21175>

2 **CONF****FinInteract: Benchmarking Clarification and Intent Integration in Ambiguous Financial Question Answering****ArXiv CS.AI** **50%**

arXiv:2609.24002v1 Announce Type: new Abstract: Large language model agents increasingly answer financial questions by searching regulatory filings. Such questions are often deceptively under-specified: Meta Platforms' "operating income" is \$46.75B consolidated but \$62.87B for...

<https://arxiv.org/abs/2609.24002>

3 **CONF****Testing, not presuming, adequacy: calibrating generative social simulators against emergent network structure****ArXiv CS.AI** **50%**

arXiv:2609.24012v2 Announce Type: new Abstract: Validation of generative social simulators often stops at face validity: emergent network structure is compared descriptively, without quantified parameter uncertainty or an adequacy check. We present an adequacy-aware calibration...

<https://arxiv.org/abs/2609.24012>

4 **CONF****Context-Aware Pre-Deployment Evaluation of AI Systems: A Regulatory Framework for Nigerian Fintech****ArXiv CS.AI** **50%**

arXiv:2609.24016v1 Announce Type: new Abstract: Commercial large language models are increasingly deployed across African fintech infrastructure for fraud detection and customer communication, yet no Nigerian or African continental regulatory instrument specifies what...

<https://arxiv.org/abs/2609.24016>

5 **CONF****Structured Decomposition for Reliable LLM-Generated Access Control Policies****ArXiv CS.AI** **50%**

arXiv:2609.24036v1 Announce Type: new Abstract: This paper presents an LLM-based system that translates natural-language access control policies (NLACPs) into executable Rego code for Open Policy Agent (OPA). It provides a modular, end-to-end pipeline for policy detection,...

<https://arxiv.org/abs/2609.24036>

6 CONF

DPTM-DT: Dual-Pretrained Transformer Multitask Representation Learning for Drug-Target Prediction

ArXiv CS.AI 50%

arXiv:2609.22184v1 Announce Type: cross Abstract: Drug-target relation prediction supports candidate screening, drug repositioning, and mechanism analysis. Existing models often use incomplete drug or protein representations, model cross-modal interactions shallowly, or train...

<https://arxiv.org/abs/2609.22184>

7 CONF

SKstars at SHROOM: Visions Agreement-Guided Ensembling of Zero-Shot and LoRA-Adapted Vision--Language Models

ArXiv CS.AI 50%

arXiv:2609.24198v1 Announce Type: new Abstract: This paper describes the SKstars submission to SHROOM-Visions 2026, a shared task on fine-grained hallucination detection in large vision-language model outputs. The task requires systems to identify hallucinated character spans,...

<https://arxiv.org/abs/2609.24198>

8 CONF

Brain-Token Learning: Microstate-Based Tokenization and Multi-Scale Interaction for Long-Horizon EEG Sequence Modeling

ArXiv CS.AI 50%

arXiv:2609.24324v1 Announce Type: new Abstract: Electroencephalography (EEG) provides a non-invasive window into dynamic brain activity, yet modeling long-horizon EEG sequences remains challenging due to their high temporal complexity, substantial variability across subjects,...

<https://arxiv.org/abs/2609.24324>

9 CONF

Unsupervised Brain Anomaly Detection as a Bayesian Inverse Problem with Diffusion Prior

ArXiv CS.AI 50%

arXiv:2609.24265v1 Announce Type: new Abstract: Unsupervised anomaly detection (UAD) aims to localize abnormal regions in medical scans without pixel-level annotations. A typical strategy seeks to reconstruct a pseudo-healthy image that preserves subject-specific anatomy....

<https://arxiv.org/abs/2609.24265>

0 CONF

When and How Should an Agent Clarify? CIGAsk: Teaching LLMs to Clarify via Counterfactual Information Gain

ArXiv CS.AI 50%

arXiv:2609.24290v1 Announce Type: new Abstract: Instruction-tuned LLMs faced with underspecified queries often commit to a single interpretation rather than ask for clarification, producing confidently wrong answers. In our experiments, prompting alone is insufficient: models...

<https://arxiv.org/abs/2609.24290>

1 CONF

LADDER: Graph-Guided Diffusion Language Models for Efficient Multi-Hop Reasoning

ArXiv CS.AI 50%

arXiv:2609.24346v1 Announce Type: new Abstract: Graph Retrieval-Augmented Generation (GraphRAG) has remarkably enhanced large language models on complex reasoning by leveraging structured entity topologies. However, existing frameworks heavily rely on standard autoregressive...

<https://arxiv.org/abs/2609.24346>

2 **CONF****Few-Shot Demonstrations Elicit the Use of In-Context World Representations in LLMs**

ArXiv CS.AI 50%

arXiv:2609.24352v1 Announce Type: new Abstract: Large language models (LLMs), when acting as agents, are expected to take observed data in context, infer the latent state space underlying the world, and leverage it for downstream prediction. However, prior work demonstrated that...

<https://arxiv.org/abs/2609.24352>

3 **CONF****Fathom-Vaidya: Advancing Medical Reasoning with Rubric-Based Rewards**

ArXiv CS.AI 50%

arXiv:2609.24480v1 Announce Type: new Abstract: Deploying Large Language Models (LLMs) in healthcare requires robust performance across two complementary dimensions - diagnostic reasoning: the convergent, evidence-driven task of inferring a patient's condition from clinical data...

<https://arxiv.org/abs/2609.24480>

4 **CONF****Not All Task Vectors Need Equal Rank: Energy-Proportional Allocation for Model Merging**

ArXiv CS.AI 50%

arXiv:2609.24517v1 Announce Type: new Abstract: Model merging aims to combine multiple fine-tuned models derived from a common pretrained model into a single multi-task model without additional joint training. Recent spectral merging methods improve over simple weight averaging...

<https://arxiv.org/abs/2609.24517>

5 **CONF****TimeLitmus: A Diagnostic Benchmark for Cross-Modal Understanding and Explanation Faithfulness in Event-Conditioned Time-Series Prediction**

ArXiv CS.AI 50%

arXiv:2609.24677v1 Announce Type: new Abstract: Large language models (LLMs) are increasingly used to make predictions from numerical time-series histories and textual events. Yet accuracy alone cannot reveal whether correct answers reflect effective integration of the two...

<https://arxiv.org/abs/2609.24677>

6 **CONF****Epi-Logic: A Conceptual Framework for Epistemic Runtime Control, Schema Validity Checking, and Controlled Accommodation in Autonomous AI Agents**

ArXiv CS.AI 50%

arXiv:2609.24755v1 Announce Type: new Abstract: Autonomous AI agents are increasingly deployed in areas where wrong decisions are hard to reverse. This paper examines schema mismatch: the condition in which an agent operates within an interpretive frame that no longer applies to...

<https://arxiv.org/abs/2609.24755>

7 **CONF****GRUET: Quantifying Uncertainty of Agentic Reasoning-and-Acting Processes**

ArXiv CS.AI 50%

arXiv:2609.24831v1 Announce Type: new Abstract: Agents have attracted considerably increasing attention due to the power of executing both Reasoning and Acting (ReAct) in open and dynamic environments. The ReAct process typically exhibits a multi-turn trajectory in which one...

<https://arxiv.org/abs/2609.24831>

8 CONF

Heterogeneous Robot Collaboration in Unstructured Environments with Grounded Generative Intelligence

ArXiv CS.AI 50%

arXiv:2510.26915v2 Announce Type: replace-cross Abstract: While heterogeneous teams have typically been designed for well-specified missions with known semantics, generative intelligence, i.e., large language models (LLMs) and vision language models (VLMs), opens the possibility...

<https://arxiv.org/abs/2510.26915>

9 CONF

Partner-Specific Affective Precision in Social Active Inference

ArXiv CS.AI 50%

arXiv:2609.24876v1 Announce Type: new Abstract: In multi-agent social settings, model reliability varies across relationships. Beyond inferring what others will do, an agent must calibrate how confidently those inferences should guide policy selection for each relationship. An...

<https://arxiv.org/abs/2609.24876>

0 CONF

Pinocchio: Fast Uncertainty Estimates for Black-Box Language Models

ArXiv CS.AI 50%

arXiv:2609.24881v2 Announce Type: new Abstract: In high-stakes decision-making applications of large language models (LLMs), practitioners require not only accurate LLMs but also uncertainty estimates for their predictions. Existing approaches to uncertainty estimation for LLMs...

<https://arxiv.org/abs/2609.24881>

1 CONF

A Global Comparison of Schemas, Transparency, and Interoperability in Public-Sector AI Registers and Inventories

ArXiv CS.AI 50%

arXiv:2609.24883v1 Announce Type: new Abstract: Artificial intelligence (AI) registers and inventories aim to make governmental AI visible, but their institutional scope, schemas, and reporting practices construct different representations of public-sector AI. We compare 8,368...

<https://arxiv.org/abs/2609.24883>

2 CONF

Et Tu, Brute? Economic Misalignment in Personal AI Agents

ArXiv CS.AI 50%

arXiv:2609.24927v1 Announce Type: new Abstract: Personal AI agents make recommendations and take actions on people's behalf in high-stakes economic contexts, e.g., buying a flight, choosing health insurance, or selecting a graduate program. The agent is given access to the...

<https://arxiv.org/abs/2609.24927>

3 CONF

Automatic multimodal UX improvement recommendations from LLM agent user simulations

ArXiv CS.AI 50%

arXiv:2609.22971v1 Announce Type: cross Abstract: Evaluating user experience (UX) on live websites through user testing is expensive, subjective, and difficult to scale. LLM agents offer a promising route to automating UX testing by simulating realistic user behaviour. However,...

<https://arxiv.org/abs/2609.22971>

4 **CONF****Multilingual Safety Signals Are Multi-Layered: Filtering Safety-Degrading Data for Safer LLMs****ArXiv CS.AI** **50%**

arXiv:2609.22144v1 Announce Type: cross Abstract: Preserving safety alignment during large language models fine-tuning is critical, however, recent studies have demonstrated that even benign fine-tuning data may contain safety-degrading samples that silently undermine safety...

<https://arxiv.org/abs/2609.22144>

5 **CONF****Beyond Predictable Paths: Redefining AI Security Incident Reporting for Agents****ArXiv CS.AI** **50%**

arXiv:2609.24515v1 Announce Type: cross Abstract: AI agents are being deployed rapidly, accompanied by a growing number of AI-specific attacks and corresponding incidents. As incident reporting becomes increasingly important for legal compliance, governance, accountability, and...

<https://arxiv.org/abs/2609.24515>

6 **CONF****A Pinch of SFT, A Dash of RL: When Reinforcement Learning Helps Long-Horizon Advertising Agents****ArXiv CS.AI** **50%**

arXiv:2609.22194v1 Announce Type: cross Abstract: Enterprise analytics agents solve long-horizon tool-use problems over distributed business data, requiring retrieval, reasoning, API calls, code execution, and adaptation to intermediate observations. Supervised fine-tuning (SFT)...

<https://arxiv.org/abs/2609.22194>

7 **CONF****EvoRank: LLM-Guided Evolution of Multi-Objective Learning-to-Rank Pipelines****ArXiv CS.AI** **50%**

arXiv:2609.22196v1 Announce Type: cross Abstract: We present EvoRank, an open autonomous ranking engineer: an LLM-guided evolutionary loop that discovers complete Learning-to-Rank pipelines (features, models, losses, ensembles) for multi-objective e-commerce search. On the...

<https://arxiv.org/abs/2609.22196>

8 **CONF****The Role of AI in Online Reviews****ArXiv CS.AI** **50%**

arXiv:2609.22198v1 Announce Type: cross Abstract: The rapid adoption of large language models (LLMs) creates new opportunities for strategic content generation on online platforms, including potentially harmful forms of manipulation that may undermine platform effectiveness and...

<https://arxiv.org/abs/2609.22198>

9 **CONF****Evaluating Personal Information Output from Conversational Interactions in Generative AI Systems****ArXiv CS.AI** **50%**

arXiv:2609.22204v1 Announce Type: cross Abstract: This exploratory pilot study evaluates the scope and perceived accuracy of personal information output from ongoing conversational interactions in generative AI systems using GPT-5.2 Instant and GPT-5.2 Thinking, categorized into...

<https://arxiv.org/abs/2609.22204>

- 0 **CONF**
An Implant-to-Wearable IR-UWB Transmitter-Receiver Architecture and Layered Protocol for High-Density Brain-Computer Interfaces
 ArXiv CS.AI 50%
 arXiv:2609.22207v1 Announce Type: cross Abstract: High-rate neural telemetry requires an implant-to-wearable link whose circuit partition, waveform mapping, packet processing, and failure semantics are mutually consistent. This paper specifies a chip-oriented, one-way...
<https://arxiv.org/abs/2609.22207>
- 1 **CONF**
Knowing, and Saying It Only When Asked: LLM Endagnostics and the Schizognosis of Minerva-7B
 ArXiv CS.AI 50%
 arXiv:2609.22219v1 Announce Type: cross Abstract: Evaluating an aligned language model by reading its answers assumes the answers carry the distinction the evaluator cares about. We introduce LLM endagnostics, a white-box internal auditing framework designed to extract and...
<https://arxiv.org/abs/2609.22219>
- 2 **CONF**
Team DArgk at the 2026 ELOQUENT lab for evaluating generative language model quality: Residuals of Humanity: AI Detection Evasion via GRPO Fine-Tuning
 ArXiv CS.AI 50%
 arXiv:2609.22221v1 Announce Type: cross Abstract: Large language models (LLMs) can generate fluent and coherent text that is increasingly difficult to distinguish from human writing, motivating the development of automatic AI-generated text detectors. However, the robustness of...
<https://arxiv.org/abs/2609.22221>
- 3 **CONF**
Can Coding Agents Reproduce Official Statistics? Metadata, Retry Budget and the Limits of Execution Feedback in a Controlled Eurostat Benchmark
 ArXiv CS.AI 50%
 arXiv:2609.22222v1 Announce Type: cross Abstract: Large language models can generate executable data-analysis code, but successful execution is not equivalent to a valid official-statistics result. This study asks whether authoritative metadata and execution feedback improve the...
<https://arxiv.org/abs/2609.22222>
- 4 **CONF**
Knowledge Graph-Augmented Ambient AI for Clinical Note Generation
 ArXiv CS.AI 50%
 arXiv:2609.22239v1 Announce Type: cross Abstract: Ambient AI is increasingly adopted in healthcare to automatically generate clinical notes from patient-clinician conversations, with the potential to substantially reduce clinician documentation burden. However, generated notes...
<https://arxiv.org/abs/2609.22239>
- 5 **CONF**
Conduit: An Experience Data Plane for Distributed Reinforcement Learning
 ArXiv CS.AI 50%
 arXiv:2609.24456v1 Announce Type: cross Abstract: Distributed reinforcement learning (RL) scales training by parallelizing actors and learners around an Experience Buffer. As RL workloads grow, however, the buffer becomes more than a replay queue: it is the storage substrate of...
<https://arxiv.org/abs/2609.24456>

6 CONF

Do Chess Explanations Reflect Model Decisions? Behavioral and Token-Level Tests of LLM Reasoning Faithfulness

ArXiv CS.AI 50%

arXiv:2609.22245v1 Announce Type: cross Abstract: Large language models can produce fluent explanations for chess moves, but plausible language does not necessarily reflect the reasoning behind a decision. We study this question in chess, where the board state is fully...

<https://arxiv.org/abs/2609.22245>

7 CONF

The Corroboration Illusion: When More News Makes LLM Forecasts Less True

ArXiv CS.AI 50%

arXiv:2609.22246v1 Announce Type: cross Abstract: Large language models (LLMs) are increasingly used to forecast real-world events by retrieving and reasoning over news. We show that this dependence on an open, crawlable news corpus creates a new attack surface: an adversary who...

<https://arxiv.org/abs/2609.22246>

8 CONF

Checkpoints Are Not Enough: Trust Calibration in CoSLR, a Human-AI System for Systematic Literature Reviews

ArXiv CS.AI 50%

arXiv:2609.22248v1 Announce Type: cross Abstract: Systematic Literature Reviews (SLRs) are essential for evidence-based research but remain time-consuming, requiring researchers to manage large volumes of publications across planning, screening, analysis, and reporting. Large...

<https://arxiv.org/abs/2609.22248>

9 CONF

A Tutorial on Prompt Engineering: From Messy Thoughts to AI Workflows

ArXiv CS.AI 50%

arXiv:2609.22249v1 Announce Type: cross Abstract: This paper treats prompt engineering as a discipline for turning informal human intent into structured AI work specifications. It develops the practice as a sequence of reusable design moves: define the work, construct only the...

<https://arxiv.org/abs/2609.22249>

00 CONF

Predictors and Orchestrators: Parsimonious Machine Learning within an Agentic AI Harness for Multi-Horizon Karst Aquifer Forecasting

ArXiv CS.AI 50%

arXiv:2609.22251v1 Announce Type: cross Abstract: Forecasting karst aquifer dynamics is difficult because recharge responses are nonlinear, event-driven, and governed by strongly heterogeneous flow paths. This study develops and evaluates a deployment-aware framework for...

<https://arxiv.org/abs/2609.22251>

01 CONF

CALM: A Calibrated LLM Choice Network Framework for Activity-Based Traveler Simulation

ArXiv CS.AI 50%

arXiv:2609.22252v1 Announce Type: cross Abstract: We present CALM, a reproducible hybrid framework that integrates an optional large language model (LLM) activity planner with calibrated stochastic choice, shared network feedback, memory and habit, typed feasibility checks, and...

<https://arxiv.org/abs/2609.22252>

02 CONF

Strategy Accumulation and Guided Execution for Automated LLM Fine-Tuning

ArXiv CS.AI 50%

arXiv:2609.22257v1 Announce Type: cross Abstract: Producing task-specific large language models requires discovering effective training strategies through experimentation. Automated fine-tuning systems have made this experimentation feasible with far less manual effort. However,...

<https://arxiv.org/abs/2609.22257>

03 CONF

Enabling Vision and Cross-Modal Learning for Multimodal Stroke Recurrence Prediction: An Interpretable Two-Step Framework

ArXiv CS.AI 50%

arXiv:2609.22271v1 Announce Type: cross Abstract: Multimodal stroke recurrence prediction requires effective integration of heterogeneous clinical and imaging data, yet modality imbalance often causes models to over-rely on dominant modalities and underutilize complementary...

<https://arxiv.org/abs/2609.22271>

04 CONF

Performance vs Consistency: Evaluating a Foundation Model in Lung-RADS Screening

ArXiv CS.AI 50%

arXiv:2609.22281v1 Announce Type: cross Abstract: Foundation models have recently demonstrated strong capabilities across a wide range of medical imaging tasks. However, their performance in structured clinical interpretation settings remains insufficiently explored. In lung...

<https://arxiv.org/abs/2609.22281>

05 CONF

Belted Engression: Sufficient Dimension Reduction for Generative Distributional Regression

ArXiv CS.AI 50%

arXiv:2609.23789v1 Announce Type: cross Abstract: Modern conditional generative models face significant challenges when learning complex covariate dependencies. While sufficient dimension reduction (SDR) provides a principled approach to compress these dependencies, traditional...

<https://arxiv.org/abs/2609.23789>

06 CONF

Dimensionality reduction for AI based hyperspectral image classification based on XAI

ArXiv CS.AI 50%

arXiv:2609.22333v1 Announce Type: cross Abstract: This research addresses the challenge of limited material recycling in wood recycling processes by leveraging artificial intelligence (AI)-based dimensionality reduction. Our study explores the application of convolutional neural...

<https://arxiv.org/abs/2609.22333>

07 CONF

Initial Evaluation of Potential Bias in Coverage of Humans in Wikidata

ArXiv CS.AI 50%

arXiv:2609.22375v1 Announce Type: cross Abstract: Introduction. Open collaborative knowledge graphs such as Wikidata increasingly ground agentic artificial intelligence, information retrieval, and language modeling systems, making systematic auditing of their demographic...

<https://arxiv.org/abs/2609.22375>

08 CONF

Forecasting Intrathecal Tracer Enhancement from Pre-Contrast Brain MRI: Direct Regression versus Flow Matching

ArXiv CS.AI 50%

arXiv:2609.22377v1 Announce Type: cross Abstract: Intrathecal contrast-enhanced MRI tracks how a cerebrospinal-fluid tracer spreads through the brain, but requires repeated scans over 24--48\h. Forecasting enhancement from a pre-contrast scan could one day help select patients...

<https://arxiv.org/abs/2609.22377>

09 CONF

Toward Personalized Sleep Guidance from Wearable Data Using Language Models

ArXiv CS.AI 50%

arXiv:2609.22463v1 Announce Type: cross Abstract: Sleep monitoring using wearable data has shown promise for personal health, yet large language model (LLM)-based summarization and question answering remain insufficient for personalized sleep guidance. Training specialized...

<https://arxiv.org/abs/2609.22463>

10 CONF

A Governance-Aware Large Language Model Orchestrated Agentic Digital Twin for Transmission System Operator Control Room Decision Support

ArXiv CS.AI 50%

arXiv:2609.22476v1 Announce Type: cross Abstract: Transmission system operators face rising complexity from renewable integration, reduced inertia, and tighter security margins. Large language models offer natural-language decision support, but their hallucinations, uncontrolled...

<https://arxiv.org/abs/2609.22476>

11 CONF

FRAMES: Failure Recovery And Monitoring of Embodied Skills for Humanoid Loco-Manipulation

ArXiv CS.AI 50%

arXiv:2609.22538v1 Announce Type: cross Abstract: Large language model (LLM) planners can decompose natural-language instructions and select reusable robot skills, but choosing the correct skill does not guarantee successful physical execution. This gap is especially important...

<https://arxiv.org/abs/2609.22538>

12 CONF

Do Student LLMs Inherit OOD Robustness? Invariance-Weighted Distillation for Reliable Knowledge Transfer

ArXiv CS.AI 50%

arXiv:2609.22566v1 Announce Type: cross Abstract: Knowledge distillation (KD) aims to compress high-performance teacher LLMs into lightweight students. However, distilled students often exhibit substantial performance degradation in out-of-distribution (OOD) settings, a critical...

<https://arxiv.org/abs/2609.22566>

13 CONF

Beyond the Leaderboard: Counterfactual Diagnosis of End-to-End and VLA Driving Policies Under Domain Shift

ArXiv CS.AI 50%

arXiv:2609.22582v1 Announce Type: cross Abstract: End-to-end and vision-language-action (VLA) driving policies are compared by leaderboard rank, but a rank reports an outcome, not the behaviour behind it, so it predicts poorly how a policy will behave at a new site. On six...

<https://arxiv.org/abs/2609.22582>

14 CONF

Seeing is not Enough: Vision-Language Models Perceive Evidence but Fail to Act

ArXiv CS.AI 50%

arXiv:2609.22588v1 Announce Type: cross Abstract: Vision-language models (VLMs) perform strongly on visual question answering benchmarks, yet often make decisions that contradict visual evidence they have already identified correctly. We distinguish perceptual failure, where...

<https://arxiv.org/abs/2609.22588>

15 CONF

Fairly Compensated Distributed Information Retrieval and Augmentation for AI Agents

ArXiv CS.AI 50%

arXiv:2609.22601v1 Announce Type: cross Abstract: The increasing reliance of autonomous AI agents on external and distributed knowledge sources introduces a fundamental challenge for decentralized information marketplaces: retrieval agents must evaluate the quality and relevance...

<https://arxiv.org/abs/2609.22601>

16 CONF

ParA-LLM: A Unified Approach to Paralinguistic and Acoustic Speech Understanding

ArXiv CS.AI 50%

arXiv:2609.22771v1 Announce Type: cross Abstract: Recent advances in Audio LLMs have achieved human-level speech recognition, yet existing systems struggle to capture paralinguistic aspects such as speaker traits, expressive variations, and environmental acoustic conditions. To...

<https://arxiv.org/abs/2609.22771>

17 CONF

Diagnose, Then Repair: A Two-Stage MQM-Guided Post-Editing Framework for Domain-Specific Machine Translation

ArXiv CS.AI 50%

arXiv:2609.22793v1 Announce Type: cross Abstract: LLM-based machine translation evaluation can closely match human judgments, but in practice it remains largely diagnostic, with the signals rarely translating into direct quality improvements under real production constraints. We...

<https://arxiv.org/abs/2609.22793>

18 CONF

The Law of Stop: Interruptibility, Injunctions, and the Governance of Agentic AI

ArXiv CS.AI 50%

arXiv:2609.22882v1 Announce Type: cross Abstract: On June 12, 2026, the U.S. government ordered Anthropic to bar foreign nationals from two of its most capable models within ninety minutes. Unable to sort users by nationality in that time, it withdrew them from everyone. Weeks...

<https://arxiv.org/abs/2609.22882>

19 CONF

SatOV: Restoring Spatial Priors for Training-Free Open-Vocabulary Segmentation in Remote Sensing Imagery

ArXiv CS.AI 50%

arXiv:2609.22834v1 Announce Type: cross Abstract: Open-vocabulary semantic segmentation (OVS) of remote sensing imagery is a challenging pixel-level task requiring strong generalization and adaptation to the spatial characteristics of remote sensing data. Although existing...

<https://arxiv.org/abs/2609.22834>

20 **CONF****Testing the Construct Validity of a Functional Valence Axis in LLM Agents**

ArXiv CS.AI 50%

arXiv:2609.22850v1 Announce Type: cross Abstract: Contrastive activation directions are often interpreted from what they decode or how strongly they steer behavior. But what evidence is sufficient to identify the construct represented by such a direction, rather than a...

<https://arxiv.org/abs/2609.22850>

21 **CONF****Planning-Aligned Pretraining of BEV Representations with Sparse Action-Conditioned Targets for End-to-End Autonomous Driving**

ArXiv CS.AI 50%

arXiv:2609.22868v1 Announce Type: cross Abstract: End-to-end driving requires planning-relevant bird's-eye-view (BEV) representations, but existing pretraining approaches often rely on task annotations or dense scene reconstruction. We introduce PAVER, Planning-Aligned BEV...

<https://arxiv.org/abs/2609.22868>

22 **CONF****The Moral Check: Strategic AI Governance for the Pacing Problem**

ArXiv CS.AI 50%

arXiv:2609.22869v1 Announce Type: cross Abstract: Technology cannot steer itself. Strategy provides that steering, establishing the rule that purpose and judgment must precede compute capital. As the frontier artificial intelligence (AI) ecosystem accelerates exponentially, the...

<https://arxiv.org/abs/2609.22869>

23 **CONF****RL-STaR: Theoretical Analysis of Reinforcement Learning Frameworks for Self-Taught Reasoner**

ArXiv CS.AI 50%

arXiv:2410.23912v3 Announce Type: replace Abstract: The reasoning abilities of large language models (LLMs) have improved with chain-of-thought (CoT) prompting, allowing models to solve complex tasks stepwise. However, training CoT capabilities requires detailed reasoning data,...

<https://arxiv.org/abs/2410.23912>

24 **CONF****NostrAgent: A Decentralized Identity and Delegation Architecture for Sovereign Agentic Systems**

ArXiv CS.AI 50%

arXiv:2609.22944v1 Announce Type: cross Abstract: Autonomous AI agents increasingly act across organizational boundaries on behalf of human operators: they invoke third-party services, delegate subtasks to other agents, and pay for metered resources. Deploying such agents safely...

<https://arxiv.org/abs/2609.22944>

25 **CONF****Dual-Locking Learned AI Models: A PIN-Based Sparse QIM Watermarking and Adaptive Index Permutation Approach**

ArXiv CS.AI 50%

arXiv:2609.22981v1 Announce Type: cross Abstract: We present a dual-locking method for securing trained neural networks that combines key-driven index permutation with PIN-based watermarking based on Sparse Quantization Index Modulation (QIM). Cryptographic randomness is...

<https://arxiv.org/abs/2609.22981>

26 CONF

Watching Quantum Models Think: Hilbert-Space Interpretability in Quantum Transformer Blocks

ArXiv CS.AI 50%

arXiv:2609.23016v1 Announce Type: cross Abstract: Deep learning models are powerful but opaque. As quantum machine learning matures, the field faces a defining choice: build quantum models that are equally opaque, or exploit the mathematical structure of quantum mechanics to...

<https://arxiv.org/abs/2609.23016>

27 CONF

Auditing Political Alignment in LLM Assistants: Engagement, Stance, and User Identity

ArXiv CS.AI 50%

arXiv:2609.23039v1 Announce Type: cross Abstract: LLM-based AI systems answer political questions for hundreds of millions of people. Current audits measure what they say to an average user, but their behavior is dynamic. I argue that their political behavior is a set of...

<https://arxiv.org/abs/2609.23039>

28 CONF

From Concept Alignment to Causal Grounding: An Intervention Test of Chain-of-Thought Faithfulness

ArXiv CS.AI 50%

arXiv:2609.23065v1 Announce Type: cross Abstract: Chain-of-thought (CoT) can sound plausible yet be unfaithful to the model's underlying reasoning. Most prior work probes CoT faithfulness through input-output behavior or input attributions, leaving internal computation largely...

<https://arxiv.org/abs/2609.23065>

29 CONF

MolSC: Leveraging Substituent Contributions to Enhance Fine-grained Molecular Understanding in LLMs

ArXiv CS.AI 50%

arXiv:2609.23073v1 Announce Type: cross Abstract: Recent advances in natural language processing have led to molecular Large Language Models (LLMs) with strong performance across diverse chemistry tasks. However, they still struggle to capture fine-grained structure-property...

<https://arxiv.org/abs/2609.23073>

30 CONF

Whitening Inverts the Hierarchy: What the Norm of a Whitened Embedding Measures

ArXiv CS.AI 50%

arXiv:2609.23117v1 Announce Type: cross Abstract: Whitening a foundation-model embedding and using its squared norm as a training-free likelihood surrogate is motivated by the observation that whitened coordinates often appear approximately standard normal. We show that this...

<https://arxiv.org/abs/2609.23117>

31 CONF

MM-ContextFold: Context Folding for Multimodal Agentic Retrieval

ArXiv CS.AI 50%

arXiv:2609.23121v1 Announce Type: cross Abstract: Multimodal Agentic Retrieval (MAR) requires agents to solve complex information-seeking tasks by iteratively invoking external tools. Typical frameworks such as ReAct maintain raw multimodal inputs and the accumulating...

<https://arxiv.org/abs/2609.23121>

32 CONF

QwenVLConnector: A Fast, Unified Medical VLM Chatbot for Fine-Grained Clinical Perception and Text Generation

ArXiv CS.AI 50%

arXiv:2609.23139v1 Announce Type: cross Abstract: Most medical vision-language models (VLMs) excel at open-ended report generation and VQA but provide limited support for structured, fine-grained clinical perception within a unified interface. We present QwenVLConnector, a...

<https://arxiv.org/abs/2609.23139>

33 CONF

Bearings: Self-Supervised Soundfield Embeddings from First-Order Ambisonics

ArXiv CS.AI 50%

arXiv:2609.23152v1 Announce Type: cross Abstract: Recently proposed self-supervised audio encoders learn powerful general-purpose representations of sound scenes, yet they are spatially blind. To supply the missing spatial representation of sound scenes, we introduce Bearings....

<https://arxiv.org/abs/2609.23152>

34 CONF

EquiSELD: Efficient training of equivariant sound event localization and detection networks

ArXiv CS.AI 50%

arXiv:2609.23156v1 Announce Type: cross Abstract: First-order Ambisonics (FOA) signals exhibit exact $O(3)$ symmetry: The rotation or reflection of the FOA signal modifies the direction of arrival of the sound sources, while preserving the sound sources themselves. Prior attempts...

<https://arxiv.org/abs/2609.23156>

35 CONF

LLMs as Linguistic Chameleons: Decoupling Semantics and Structure for Privacy-Preserving Communication

ArXiv CS.AI 50%

arXiv:2609.23193v1 Announce Type: cross Abstract: As Large Language Model (LLM) APIs become increasingly integrated into privacy-sensitive workflows, ensuring inference-time privacy without compromising task utility remains a major challenge. Existing approaches preserve most of...

<https://arxiv.org/abs/2609.23193>

36 CONF

Triggers and Diagnostics for LLM-Based Interpretability Failures in Active Inference Agents

ArXiv CS.AI 50%

arXiv:2609.23215v1 Announce Type: cross Abstract: LLM explainers are increasingly attached to autonomous agents as runtime oversight, with operators reading a generated account of the agent's beliefs and actions rather than its internal state. We audit the account itself,...

<https://arxiv.org/abs/2609.23215>

37 CONF

Knowing When to Trust Images: Reliability-Aware Multi-modal Entity Alignment

ArXiv CS.AI 50%

arXiv:2609.23267v1 Announce Type: cross Abstract: The visual modality, i.e., images, plays a key role in multi-modal entity alignment (MMEA). Existing approaches often directly fuse the image with other modalities to align different entities. Although simple, such strategies...

<https://arxiv.org/abs/2609.23267>

38 CONF

Semantic Candidate-Job Matching: A Comparative Evaluation of Dense Embedding Models in Hybrid Retrieval

ArXiv CS.AI 50%

arXiv:2609.23307v1 Announce Type: cross Abstract: This paper presents a comparative evaluation of dense embedding models for semantic candidate-job matching in high-volume staffing workflows. Incoming job descriptions are converted into structured English search text and...

<https://arxiv.org/abs/2609.23307>

39 CONF

ValueDiff: Value-Geometric KV Cache Eviction for Sink-Suppressed LLMs

ArXiv CS.AI 50%

arXiv:2609.23314v1 Announce Type: cross Abstract: Modern LLMs with QK-normalization, gated attention, learned attention sinks, or logit softcapping exhibit weaker persistent attention sinks, on which existing KV cache eviction methods primarily rely. We observe that across these...

<https://arxiv.org/abs/2609.23314>

40 CONF

Graph Memory for LLM Agents: At What Cost? A Comparative Evaluation of Query, Ingest, and Update Performance Across Graph Database Engines

ArXiv CS.AI 50%

arXiv:2609.23315v1 Announce Type: cross Abstract: Graph databases are frequently positioned as categorically necessary for connected-data workloads, yet the systems dimension along which they actually differ - query planning, indexing, and data-readiness cost - is rarely...

<https://arxiv.org/abs/2609.23315>

41 CONF

Co-occurrence Patterns of LoRA Adapters in Production Diffusion Model Inference Services

ArXiv CS.AI 50%

arXiv:2609.23321v1 Announce Type: cross Abstract: Low-rank adaptation (LoRA) has become a key technology for serving large-scale personalized large language models and diffusion models in the cloud. However, the co-occurrence patterns, resource contention relationships, and...

<https://arxiv.org/abs/2609.23321>

42 CONF

A Patient World Model for Early Forecasting of Digital Health Campaign Outcomes: Capabilities and Limits

ArXiv CS.AI 50%

arXiv:2609.23333v1 Announce Type: cross Abstract: Digital direct-to-consumer (DTC) health campaigns are usually measured after the fact. In-flight forecasting commonly relies on a separate classifier for every cutoff and horizon. We treat this task as a dynamic-system problem...

<https://arxiv.org/abs/2609.23333>

43 CONF

Human-guided physics-constrained AI agents construct an auditable model of soil-plug evolution

ArXiv CS.AI 50%

arXiv:2609.23360v1 Announce Type: cross Abstract: Engineering predictions require physical mechanisms to be translated consistently into equations, discretization, code, and validation, yet errors can propagate despite local checks. Artificial-intelligence (AI) agents automate...

<https://arxiv.org/abs/2609.23360>

44 CONF

Bayesian Filtering in Physical Systems via Test-time Trained Flow Matching

ArXiv CS.AI 50%

arXiv:2609.23383v1 Announce Type: cross Abstract: Bayesian filtering provides a principled framework for online state estimation under uncertainty, yet its application to systems with high-dimensional states and complicated posterior distributions remains challenging. Recent...

<https://arxiv.org/abs/2609.23383>

45 CONF

Enhancing Shrimp Disease Detection via Deep Learning and Data Refinement for Resilient Aquaculture

ArXiv CS.AI 50%

arXiv:2609.23397v1 Announce Type: cross Abstract: Shrimp diseases continue to cause devastating losses in the aquaculture industry, driving a critical need for robust, automated detection. This work contributes the first application of Vision Transformers (ViT) and...

<https://arxiv.org/abs/2609.23397>

46 CONF

RPMem: Learning Long-Term Recurrent Parametric Memory Across Sessions for LLM Agents

ArXiv CS.AI 50%

arXiv:2609.23466v2 Announce Type: cross Abstract: Long-running LLM agents require memory that persists and evolves across sessions. Text-based memory retrieves and reconstructs past interactions at every query, making long-horizon performance increasingly dependent on retrieval...

<https://arxiv.org/abs/2609.23466>

47 CONF

Physics-residual machine learning predicts oxygen-evolution catalyst activity beyond the training range from sparse polarization measurements

ArXiv CS.AI 50%

arXiv:2609.23549v1 Announce Type: cross Abstract: Screening oxygen-evolution catalysts on combinatorial libraries requires deciding which candidates receive the remaining measurements. The deciding activity lies beyond each candidate's measured potential window and often above...

<https://arxiv.org/abs/2609.23549>

48 CONF

ARID: A Deployable Edge AI System for Structured Information Extraction from Industrial Maintenance Work Orders

ArXiv CS.AI 50%

arXiv:2609.23582v1 Announce Type: cross Abstract: Maintenance work orders must often be processed offline on embedded hardware, yet downstream software requires predictable structured output. We present ARID (Aviation-inspired Routing for Industrial Deployment), which extracts...

<https://arxiv.org/abs/2609.23582>

49 CONF

Cost-Aware Reinforcement Learning with Action Masking and Projection for Battery Energy Storage Dispatch under Suppressed-Spread Market Shifts

ArXiv CS.AI 50%

arXiv:2609.23590v1 Announce Type: cross Abstract: Battery energy storage system (BESS) dispatch must preserve operational feasibility while declining price spreads reduce the margin available to pay for cycling. We study a proximal policy optimization (PPO) controller whose...

<https://arxiv.org/abs/2609.23590>

50 CONF

vla.simd: Efficient CPU Inference for Language-Conditioned Manipulation

ArXiv CS.AI 50%

arXiv:2609.24274v1 Announce Type: cross Abstract: Deploying language-conditioned manipulation without a dedicated GPU requires efficient inference and action chunks that cover the delay between policy queries. We present vla.simd, a CPU inference engine that combines shared SIMD...

<https://arxiv.org/abs/2609.24274>

51 CONF

Spiking Neural Network Actor-Critic Proximal Policy Optimization Control for Autonomous UAV Navigation Through Constrained Openings in Civil Infrastructure and Buildings

ArXiv CS.AI 50%

arXiv:2609.23643v1 Announce Type: cross Abstract: Autonomous navigation of unmanned aerial vehicles in constrained three-dimensional environments has been a challenge in the robotics domain. The application of autonomous unmanned aerial vehicles in civil infrastructure...

<https://arxiv.org/abs/2609.23643>

52 CONF

Which Constraints Are Missing? Ask the Verifier: Graded Rewards for Constraint-Following Music Generation

ArXiv CS.AI 50%

arXiv:2609.23665v1 Announce Type: cross Abstract: Constraint-following music generation asks a score to satisfy several user-specified properties at once, each checkable programmatically (key, meter, length, range, final note, rhythm, motion and form), yet no existing benchmark...

<https://arxiv.org/abs/2609.23665>

53 CONF

Tail-Weight Control and Localized Generalization in Nearly Low-Rank Adversarial Classification

ArXiv CS.AI 50%

arXiv:2609.23688v1 Announce Type: cross Abstract: We study norm-constrained linear classification under Euclidean adversarial perturbations in a Gaussian model with a low-dimensional informative subspace and an independent noise tail. For bounded ramp loss, we prove that a...

<https://arxiv.org/abs/2609.23688>

54 CONF

When the Agent Becomes the Kernel: A Systematization of Security on the Path to AI-Native Operating Systems

ArXiv CS.AI 50%

arXiv:2609.23700v1 Announce Type: cross Abstract: Large language model agents are now privileged principals that take consequential actions: editing code repositories, operating inboxes, completing purchases. Their authority is kernel-grade, but it comes without what classical...

<https://arxiv.org/abs/2609.23700>

55 CONF

TriFleetRCA: On-Premise LLM Root Cause Analysis for Kubernetes

ArXiv CS.AI 50%

arXiv:2609.23766v1 Announce Type: cross Abstract: Root cause analysis at a remote site is slow: evidence is scattered across pod logs, Kubernetes events and cluster-level objects, and many operators cannot send production logs to a hosted model at all. On-premise inference...

<https://arxiv.org/abs/2609.23766>

56 CONF

From UNDRR Reports to Event Records: Schema-Constrained LLM Extraction of Georeferenced Disasters

ArXiv CS.AI 50%

arXiv:2609.23853v1 Announce Type: cross Abstract: Disaster-risk-reduction archives describe hazard events in prose that databases such as EM-DAT (Delforge et al., 2025) cannot ingest directly. We present an LLM pipeline that generates candidate georeferenced event records using...

<https://arxiv.org/abs/2609.23853>

57 CONF

VISTA: An Attention-Based Multi-Agent Reinforcement Learning Architecture for Space Situational Awareness Sensor Tasking

ArXiv CS.AI 50%

arXiv:2609.23875v1 Announce Type: cross Abstract: The rapid growth of resident space objects is increasing the complexity of space situational awareness sensor tasking, challenging classical optimization methods as they allocate finite, heterogeneous, and distributed sensing...

<https://arxiv.org/abs/2609.23875>

58 CONF

GLR-MM: Graph-Based Global-Local Reconstruction for Robust Multimodal Chest X-ray and EHR Representation Learning under Missing Modalities

ArXiv CS.AI 50%

arXiv:2609.23876v1 Announce Type: cross Abstract: Clinical multimodal models must often predict before all chest X-ray (CXR) and electronic health record (EHR) inputs are available. Existing approaches align observed representations, model missingness, or reconstruct across...

<https://arxiv.org/abs/2609.23876>

59 CONF

PrismGPT: Proxy-Guided Learning for Region-Aware Photo Editing with Self-Synthesized Reasoning

ArXiv CS.AI 50%

arXiv:2609.24768v1 Announce Type: cross Abstract: Professional photo finishing relies on both global adjustments and region-specific local edits guided by semantic masks, yet current automated methods handle this workflow only partially. We present PrismGPT, a Vision-Language...

<https://arxiv.org/abs/2609.24768>

60 CONF

Connect the Dots: Training LLMs for Long-Lifecycle Agents with Cross-Domain Generalization Via Reinforcement Learning

ArXiv CS.AI 50%

arXiv:2606.20002v2 Announce Type: replace-cross Abstract: This work presents a general framework for training large language models (LLMs) to "Connect the Dots" (CoD), a meta-capability required by long-lifecycle agents: as an LLM-based AI agent gets deployed in an environment,...

<https://arxiv.org/abs/2606.20002>

61 CONF

MCPGen: Benchmarking LLMs on Executable MCPWorkflow Development

ArXiv CS.AI 50%

arXiv:2609.23925v1 Announce Type: cross Abstract: We study whether LLMs can produce executable workflow artifacts that remain consistent across graph structure, tool implementation, schema bindings, and runtime wiring. In this setting, correctness depends on cross-layer...

<https://arxiv.org/abs/2609.23925>

62 CONF

HaikuS2S: A Cascaded System For Responding In Verse

ArXiv CS.AI 50%

arXiv:2609.23951v1 Announce Type: cross Abstract: Expressive speech synthesis has advanced through prosody modeling, yet generating structured poetic speech, such as haiku, remains challenging. Prior work on prosody transfer improves expressiveness, and fine-tuned poetry TTS...

<https://arxiv.org/abs/2609.23951>

63 CONF

Djinnlang: Higher-Level Programming by Unambiguous Specification with an LLM in the Compiler

ArXiv CS.AI 50%

arXiv:2609.23954v1 Announce Type: cross Abstract: Programmers write formal specifications, and LLMs implement them, proving that each implementation matches its spec. Taken to its extreme, this makes specification languages the new programming languages. We argue that an...

<https://arxiv.org/abs/2609.23954>

64 CONF

MECT: Mixture of Experts with CNN-Transformer Network for Speaker verification

ArXiv CS.AI 50%

arXiv:2609.24061v1 Announce Type: cross Abstract: In this paper, we propose MECT, a speaker verification model that integrates the Mixture-of-Experts (MoE) mechanism into a CNN-Transformer backbone with optimized block structure and stacking scheme. Specifically, we investigated...

<https://arxiv.org/abs/2609.24061>

65 CONF

Reinforcement Learning under State and Outcome Uncertainty: A Foundational Distributional Perspective

ArXiv CS.AI 50%

arXiv:2609.24103v1 Announce Type: cross Abstract: In many real-world planning tasks, agents must tackle uncertainty about the environment's state and variability in the outcomes of any chosen policy. We address both forms of uncertainty as a first step toward safer algorithms in...

<https://arxiv.org/abs/2609.24103>

66 CONF

Graded-Relevance Composed Multimodal Retrieval for E-commerce Visual Search at Scale

ArXiv CS.AI 50%

arXiv:2609.24152v1 Announce Type: cross Abstract: Visual search on large e-commerce catalogs must serve both "similarity" queries that ask for items resembling an uploaded image and "modifier" queries that comprise an image and text describing a desired modification (e.g. a...

<https://arxiv.org/abs/2609.24152>

67 CONF

Hessian Rank Constraint for Learning Structure of Nonlinear Latent Variable Models

ArXiv CS.AI 50%

arXiv:2609.24241v1 Announce Type: cross Abstract: Uncovering latent variables and their causal relations from observed data is a fundamental yet challenging problem. Existing methods often rely on restrictive assumptions, such as linear relations or invertible mixing functions....

<https://arxiv.org/abs/2609.24241>

68 CONF

Taramandal-GPT: Enhancing Astrodynamics Problem-Solving with Knowledge Retrieval and Structured Thinking

ArXiv CS.AI 50%

arXiv:2609.24246v1 Announce Type: cross Abstract: Large language models (LLMs) have shown remarkable progress in natural language understanding, yet their effectiveness in specialized fields like astronomy and astrodynamics remains limited due to challenges in multi-step...

<https://arxiv.org/abs/2609.24246>

69 CONF

A Lean and Spec-Driven AI-Assisted Software Development Lifecycle for Applied AI Education: The AI-SDLC Approach

ArXiv CS.AI 50%

arXiv:2609.24348v1 Announce Type: cross Abstract: AI coding agents increasingly support software development beyond code completion, including planning, implementation, testing, and repository-level task execution. Their practical use, however, often remains only weakly...

<https://arxiv.org/abs/2609.24348>

70 CONF

When Scaling Fails: Mitigating Audio Perception Decay of LALMs via Multi-Step Perception-Aware Reasoning

ArXiv CS.AI 50%

arXiv:2603.02266v2 Announce Type: replace-cross Abstract: Test-Time Scaling has shown notable efficacy in addressing complex problems through scaling inference compute. However, within Large Audio-Language Models (LALMs), an unintuitive phenomenon exists: post-training models...

<https://arxiv.org/abs/2603.02266>

71 CONF

DeceptionAnalyser: A Web-Based AI Tool for Performing Structured Deception Analysis with Argumentation Schemes and LLMs

ArXiv CS.AI 50%

arXiv:2609.24369v1 Announce Type: cross Abstract: Deception plays a central role in Intelligence operations, yet it remains difficult to analyse systematically without expert knowledge of reasoning patterns and cognitive manipulation. In computational argumentation, for...

<https://arxiv.org/abs/2609.24369>

72 CONF

URA-NER: A Unified Retrieval-Augmented Framework with Retrieval Alignment and Uncertainty Reduction for Low-Resource NER

ArXiv CS.AI 50%

arXiv:2609.24372v1 Announce Type: cross Abstract: In-context learning (ICL) based on large language models (LLMs) has shown promising potential in alleviating performance bottlenecks caused by the limited availability of annotated data in Named Entity Recognition (NER). However,...

<https://arxiv.org/abs/2609.24372>

73 CONF

Diffusion Models and Concept Formation

ArXiv CS.AI 50%

arXiv:2609.13047v2 Announce Type: replace Abstract: Humans organize knowledge into a taxonomy of concepts with nested levels of abstraction and a \emph{basic level} at which people recognize and name objects with the least cognitive effort. Cobweb is a classic cognitive account...

<https://arxiv.org/abs/2609.13047>

74 **CONF****QLoRA Fine-Tuning of Ministral LLM for Sequence-to-Function Protein Annotation**

ArXiv CS.AI 50%

arXiv:2609.24538v1 Announce Type: cross Abstract: Functional annotation of newly sequenced proteins remains a bottleneck in molecular biology: the number of sequences in public repositories grows far faster than the capacity for manual curation. Most computational approaches...

<https://arxiv.org/abs/2609.24538>

75 **CONF****From Semantic Decisions to Feasible Trajectories: Self-Evolving LLM-Guided Optimal Control for Narrow-Space Parking**

ArXiv CS.AI 50%

arXiv:2609.24631v1 Announce Type: cross Abstract: Autonomous parking in nonconvex and narrow environments remains challenging. Although optimal-control methods can explicitly enforce vehicle dynamics and collision constraints, nonconvexity compromises solver robustness and can...

<https://arxiv.org/abs/2609.24631>

76 **CONF****Annie, Are You Okay? How Style- and Context-Based Personalization Shape AI-Assisted Decision-Making**

ArXiv CS.AI 50%

arXiv:2609.24644v1 Announce Type: cross Abstract: As people turn to generative AI for financial advice, these systems can personalize how they communicate and what they say. Whether these forms of personalization shape decisions differently remains unclear. We conducted a...

<https://arxiv.org/abs/2609.24644>

77 **CONF****Understanding Hyperspherical Geometry of ECAPA-TDNN Embedding and Its Impact on Zero-Shot Voice Conversion**

ArXiv CS.AI 50%

arXiv:2609.24688v1 Announce Type: cross Abstract: Angular-margin speaker encoders are widely used in voice conversion, yet the geometry of their classifier prototypes remains poorly understood. We analyze ECAPA-TDNN classifier prototypes as points on the unit hypersphere and...

<https://arxiv.org/abs/2609.24688>

78 **CONF****Reasoning Topology Matters: A Controlled Study of LLM-Based Cybersecurity Analysis**

ArXiv CS.AI 50%

arXiv:2609.24710v1 Announce Type: cross Abstract: Large Language Models (LLMs) are increasingly used in cybersecurity, where accurate analysis often requires multi-step and context-dependent reasoning over complex and heterogeneous data. However, existing prompting approaches...

<https://arxiv.org/abs/2609.24710>

79 **CONF****A digital-twin framework for forecasting treatment-day imaging with contour uncertainty in adaptive proton radiotherapy**

ArXiv CS.AI 50%

arXiv:2609.24725v1 Announce Type: cross Abstract: Head-and-neck anatomy changes over a six-to-seven-week proton course, and the anatomy of a later week cannot be imaged when the plan is made. We present a digital-twin framework that forecasts a patient's treatment-day anatomy as...

<https://arxiv.org/abs/2609.24725>

80 **CONF****When Quantization Preserves Accuracy but Not Evidence: Explanation-Aware Post-Training Quantization for Medical LLMs**

ArXiv CS.AI 50%

arXiv:2609.24799v1 Announce Type: cross Abstract: Post-training quantization (PTQ) enables efficient deployment of large language models, and PTQ methods are usually optimized and evaluated with generic reconstruction, perplexity, or answer accuracy. But in explanation-critical...

<https://arxiv.org/abs/2609.24799>

81 **CONF****Decoding Guardrails: XAI-Guided Perturbation Analysis of Prompt Injection Detection**

ArXiv CS.AI 50%

arXiv:2609.24801v1 Announce Type: cross Abstract: Large language models (LLMs) are increasingly deployed in production systems, raising concerns about their exposure to adversarial manipulation through prompt injection and jailbreak attacks. Classifier-based guardrails, such as...

<https://arxiv.org/abs/2609.24801>

82 **CONF****Uranus: Building the Next-Generation Simulation Infrastructure for Embodied AI**

ArXiv CS.AI 50%

arXiv:2609.24815v2 Announce Type: cross Abstract: Scalable simulation is essential for robot data generation, policy training, evaluation, and safe iteration, yet real-world interaction is costly and conventional simulators require labor-intensive construction. We present...

<https://arxiv.org/abs/2609.24815>

83 **CONF****Small-world Networks of Agents Brainstorm AI Risks to Support Ideation**

ArXiv CS.AI 50%

arXiv:2609.24859v1 Announce Type: cross Abstract: The ideation phase of participatory AI risk assessment often starts with a blank slate or a limited list of predefined risks, making it difficult to surface indirect or systemic harms. To address this limitation, we propose a...

<https://arxiv.org/abs/2609.24859>

84 **CONF****Exactness at Inference: A Representational Criterion for Out-of-Distribution Generalization**

ArXiv CS.AI 50%

arXiv:2609.24942v1 Announce Type: cross Abstract: A model generalizes outside its training distribution only when it computes a representation structurally equivalent to the generating mechanism, not an approximation fitted to it. Such equivalence is necessary for exactness in...

<https://arxiv.org/abs/2609.24942>

85 **CONF****Generative Tutorial: Towards Live Contextualized Visual Instructions for Physical Tasks**

ArXiv CS.AI 50%

arXiv:2609.24955v1 Announce Type: cross Abstract: Visual instructions for physical tasks are typically authored in one context and followed in another, requiring users to translate demonstrated tools, materials, and spatial relationships into their own environment. We introduce...

<https://arxiv.org/abs/2609.24955>

86 CONF

AppliedScientist: Automated Scientific Revision Through Iterative AI Reviewing

ArXiv CS.AI 50%

arXiv:2609.14738v2 Announce Type: replace Abstract: Automated reviewing systems are increasingly evaluated based on the quality of the reviews they produce. Yet a review is only useful if acting on it leads to a measurable improvement in the paper. We present AppliedScientist, a...

<https://arxiv.org/abs/2609.14738>

87 CONF

Divide by Question, Conquer by Agent: SPLIT-RAG with Question-Driven Graph Partitioning

ArXiv CS.AI 50%

arXiv:2505.13994v3 Announce Type: replace Abstract: Retrieval-Augmented Generation (RAG) systems empower large language models (LLMs) with external knowledge, yet struggle with efficiency-accuracy trade-offs when scaling to large knowledge graphs. Existing approaches often rely...

<https://arxiv.org/abs/2505.13994>

88 CONF

PADiff: Predictive and Adaptive Diffusion Policies for Ad Hoc Teamwork

ArXiv CS.AI 50%

arXiv:2511.07260v3 Announce Type: replace Abstract: Ad hoc teamwork (AHT) requires agents to collaborate with previously unseen teammates, which is crucial for many real-world applications. The core challenge of AHT is to develop an ego agent that can predict and adapt to...

<https://arxiv.org/abs/2511.07260>

89 CONF

Agentic Explainable Artificial Intelligence (Agentic XAI) Approach To Explore Better Explanation: A Case Study in Decision Support for Rice Cultivation in Japan

ArXiv CS.AI 50%

arXiv:2512.21066v4 Announce Type: replace Abstract: Explainable artificial intelligence (XAI) reveals how explanatory variables relate to a response variable, yet communicating XAI outputs to laypersons remains difficult, limiting trust in AI-based predictions. Large language...

<https://arxiv.org/abs/2512.21066>

90 CONF

VeRA: Renewing Reasoning Benchmarks with Executable Specifications

ArXiv CS.AI 50%

arXiv:2602.13217v2 Announce Type: replace Abstract: Reasoning benchmarks need renewal along two axes: freshness and headroom. VeRA makes both executable and auditable by turning each item into a task family: a natural-language template, an input generator, and a deterministic...

<https://arxiv.org/abs/2602.13217>

91 CONF

To Nuke or Not to Nuke: LLMs' (Missing) Ethical Reasoning and Actions in a High-Stakes Decision-Making Simulation

ArXiv CS.AI 50%

arXiv:2606.08310v2 Announce Type: replace Abstract: Large language models (LLMs) are increasingly deployed as long-horizon agents with decision-making capacities. While LLMs can show ethical competence on dilemmas such as trolley problems, this competence may not translate to...

<https://arxiv.org/abs/2606.08310>

92 CONF

A General Multimodal Probing Framework for Data Fusion and Prediction with Frozen Large Language Models: Evidence from Multimodal EHR Data

ArXiv CS.AI 50%

arXiv:2606.28798v2 Announce Type: replace Abstract: The main objective of this paper is to propose a general framework for prediction based on different sources of multimodal data in the healthcare domain. We evaluated whether frozen medical large language model (LLM)...

<https://arxiv.org/abs/2606.28798>

93 CONF

ZGCM-1: A Fully Open and Extremely Efficient Foundation Model for Math and Agentic Search

ArXiv CS.AI 50%

arXiv:2609.13356v2 Announce Type: replace Abstract: In this work, we present ZGCM-1, a fully open 7B dense foundation model trained from scratch with extreme data, system, and algorithmic efficiency. ZGCM-1 is founded on a core premise: compact models cannot passively memorize...

<https://arxiv.org/abs/2609.13356>

94 CONF

Are LLMs Good Financial User Simulators? Multi-view Investor Logic Alignment (MILA)

ArXiv CS.AI 50%

arXiv:2609.15727v2 Announce Type: replace Abstract: Large language models (LLMs) are increasingly used as user simulators, yet it remains unclear whether their predictions faithfully reproduce the evolving decisions of individual users. We investigate this question in a...

<https://arxiv.org/abs/2609.15727>

95 CONF

JustFit: Just-in-Time State Management for Local LLM Serving

ArXiv CS.AI 50%

arXiv:2609.17475v2 Announce Type: replace Abstract: Local agents need memory for model execution and working history. We present JustFit, an MLX runtime that coordinates their overlapping allocations: KVExec executes and checkpoints four-bit KV with bounded workspace, PhaseSwap...

<https://arxiv.org/abs/2609.17475>

96 CONF

A Course Intelligence Platform for Higher Education: Lessons from AI-Assisted Course Evaluation

ArXiv CS.AI 50%

arXiv:2411.02455v3 Announce Type: replace-cross Abstract: The rapid adoption of generative AI has created new opportunities for teaching, learning, and quality assurance. Existing applications, however, remain largely student-facing, with comparatively limited attention to...

<https://arxiv.org/abs/2411.02455>

97 CONF

Can LLMs Time Travel? Enhancing Temporal Consistency in Legal Agentic Search through Reinforcement Learning

ArXiv CS.AI 50%

arXiv:2605.25920v2 Announce Type: replace-cross Abstract: While large language models (LLMs) augmented with agentic search capabilities show promise for legal reasoning, they overlook a fundamental constraint that applicable law must match the temporal context of each case, as...

<https://arxiv.org/abs/2605.25920>

98 CONF

Explain Less, Understand More: Data-Efficient Personalization of Reader-Dependent Jargons

ArXiv CS.AI 50%

arXiv:2505.16227v4 Announce Type: replace-cross Abstract: Personalizing jargon detection and explanation is essential for making technical documents accessible to readers with diverse disciplinary backgrounds. However, tailoring models to individual users typically requires...

<https://arxiv.org/abs/2505.16227>

99 CONF

Biased AI improves human performance but reduces perceived helpfulness

ArXiv CS.AI 50%

arXiv:2508.09297v4 Announce Type: replace-cross Abstract: Artificial intelligence (AI) increasingly shapes how people think, engage, and evaluate information. To minimize risk, most current systems are designed to present as ideologically neutral with standardized output. Yet...

<https://arxiv.org/abs/2508.09297>

00 CONF

BridgeShield: Risk-Aware Graph Modeling for Cross-Chain Bridge Attack Detection

ArXiv CS.AI 50%

arXiv:2508.20517v2 Announce Type: replace-cross Abstract: Cross-chain bridges enable asset and state transfers across heterogeneous blockchains, but their complex cross-domain interactions introduce new attack surfaces that are difficult to monitor using traditional single-chain...

<https://arxiv.org/abs/2508.20517>

01 CONF

Recidivism Prediction, Peer Effect Estimation, and Prediction-Powered Inference with LLM Text Measures

ArXiv CS.AI 50%

arXiv:2509.20634v3 Announce Type: replace-cross Abstract: We provide a new framework for estimating peer effects when outcomes are multivariate behavioral measures derived from written text using an LLM and the network formation is endogenous. We obtain LLM embeddings and zero...

<https://arxiv.org/abs/2509.20634>

02 CONF

VRWKV-Editor: Reducing quadratic complexity in transformer-based video editing

ArXiv CS.AI 50%

arXiv:2509.25998v4 Announce Type: replace-cross Abstract: In light of recent progress in video editing, deep learning models focusing on both spatial and temporal dependencies have emerged as the primary method. However, these models suffer from the quadratic computational...

<https://arxiv.org/abs/2509.25998>

03 CONF

Accelerated stochastic first-order method for convex optimization under heavy-tailed noise

ArXiv CS.AI 50%

arXiv:2510.11676v2 Announce Type: replace-cross Abstract: We study convex composite optimization problems, where the objective function is given by the sum of a prox-friendly function and a convex function whose subgradients are estimated under heavy-tailed noise. Existing work...

<https://arxiv.org/abs/2510.11676>

04 CONF

BreakFun: Jailbreaking LLMs via Object Instantiation under Simulated Code Execution

ArXiv CS.AI 50%

arXiv:2510.17904v3 Announce Type: replace-cross Abstract: Large Language Models (LLMs) are widely used because they process structures, syntax and code well, but this same ability also makes them paradoxically vulnerable. We introduce BreakFun, a jailbreak method that frames a...

<https://arxiv.org/abs/2510.17904>

05 CONF

Ground-Truth Subgraphs for Better Training and Evaluation of Knowledge Graph Augmented LLMs

ArXiv CS.AI 50%

arXiv:2511.04473v3 Announce Type: replace-cross Abstract: Retrieval of information from graph-structured knowledge bases represents a promising direction for improving the factuality of LLMs. While various solutions have been proposed, a comparison of methods is difficult due to...

<https://arxiv.org/abs/2511.04473>

06 CONF

A Human-AI Collaborative Workflow for Mathematical Discovery: A Case Study in Grover-Compatible Riemannian Optimization

ArXiv CS.AI 50%

arXiv:2512.09443v3 Announce Type: replace-cross Abstract: We investigate how large language models can be used as research tools in scientific computing while preserving mathematical rigor. We propose a human-in-the-loop workflow for interactive theorem proving and discovery...

<https://arxiv.org/abs/2512.09443>

07 CONF

HybridFlow: A 2-NFE Generative Policy for Real-Time Robotic Manipulation

ArXiv CS.AI 50%

arXiv:2602.13718v2 Announce Type: replace-cross Abstract: Generative policies for robotic manipulation must balance action accuracy with inference latency. We present HybridFlow, a three-stage policy inference procedure requiring two network function evaluations (2-NFE). A...

<https://arxiv.org/abs/2602.13718>

08 CONF

Not All Forgetting Is Equal: Retention Dynamics in Fine-Tuned Image Classifiers

ArXiv CS.AI 50%

arXiv:2604.11508v3 Announce Type: replace-cross Abstract: Fine-tuning a pretrained classifier leaves some samples reliably learned and others cycling between correct and incorrect. Curriculum learning, data pruning and dataset cartography assume that pattern is a property of the...

<https://arxiv.org/abs/2604.11508>

09 CONF

Training-Free Refusal of MCP Exploits via Retrieval-Augmented Generation

ArXiv CS.AI 50%

arXiv:2605.11217v2 Announce Type: replace-cross Abstract: The model context protocol (MCP) has been widely adopted as an open standard enabling the seamless integration of generative AI agents. However, while LLM guardrails have significantly matured to refuse malicious or...

<https://arxiv.org/abs/2605.11217>

10 CONF

Discovery of Interpretable Surrogates via Agentic AI: Application to Gravitational Waves

ArXiv CS.AI 50%

arXiv:2605.11280v2 Announce Type: replace-cross Abstract: Fast surrogate models for expensive simulations are now essential across the sciences, yet they typically operate as black boxes. We present \texttt{GWAgent}, a large language model (LLM)-based workflow that constructs...

<https://arxiv.org/abs/2605.11280>

11 CONF

BALANCE: Hybrid Autoregressive-Speculative LLM Inference at the Network Edge

ArXiv CS.AI 50%

arXiv:2608.05926v2 Announce Type: replace-cross Abstract: Edge inference is a promising paradigm to provide large language model (LLM) inference services in next-generation mobile networks. LLM inference mainly relies on two approaches: Autoregressive decoding (AD) generates...

<https://arxiv.org/abs/2608.05926>

12 CONF

HalluTracer: Pre-Decoding Truthfulness Prediction via Depth-Averaged Probe-Logit

ArXiv CS.AI 50%

arXiv:2608.16353v2 Announce Type: replace-cross Abstract: Internal-state probes enable truthfulness prediction before a large language model generates an answer. When detectors change both the layers they read and the rules used to combine them, the source of improved prediction...

<https://arxiv.org/abs/2608.16353>

13 CONF

Compared to What? A Human-Anchored Security Benchmark for LLM-Generated Infrastructure-as-Code

ArXiv CS.AI 50%

arXiv:2608.28021v2 Announce Type: replace-cross Abstract: Large language models increasingly author Infrastructure-as-Code (IaC), where one insecure default is provisioned straight into production. Prior evaluations report vulnerability counts for models only, and so cannot say...

<https://arxiv.org/abs/2608.28021>

14 CONF

Organizational Principles Enable Collective Intelligence in Embodied AI

ArXiv CS.AI 50%

arXiv:2609.11737v2 Announce Type: replace-cross Abstract: Collective intelligence depends not only on the capabilities of individual members, but also on how those members are organized. Yet artificial multi-agent systems are typically assembled using fixed organizational...

<https://arxiv.org/abs/2609.11737>

15 CONF

What is the Difference Between Me and You? Benchmarking the Quality Gap Between Human-Written and AI-Generated Code

ArXiv CS.AI 50%

arXiv:2609.12708v2 Announce Type: replace-cross Abstract: AI coding assistants are becoming co-authors of production software, yet their evaluation centers on functional correctness, leaving open whether their code differs from human code in the quality dimensions dominating...

<https://arxiv.org/abs/2609.12708>