



Pipeline

Scrape HackerNews (Show/Top/New), Reddit (r/artificial, r/MachineLearning, r/singularity, r/LocalLLaMA), TechCrunch AI, VentureBeat AI, MIT Tech Review, The Verge AI, Wired AI, ArXiv CS.AI

Filter 40+ AI keyword patterns — LLMs, alignment, safety, research, tooling

Dedup Title-similarity (45% threshold) merges cross-source variants

Score Authenticity = 0.5 base + 0.3 per extra source + diversity + HN score bonuses

Rank Sorted by authenticity. Below 0.25 threshold are dropped.

Authenticity Tiers

- **VERIFIED 0.80 - 1.00** 3+ independent sources including media
- **CONFIRMED 0.50 - 0.79** 2 sources or 1 high-quality media
- **EMERGING 0.25 - 0.49** Single-source; passes AI keyword filter

Source Breakdown

Source	Articles	Category
ArXiv CS.AI	106	Media
TechCrunch AI	11	Media
HackerNews	5	Community
The Verge AI	3	Media
MIT Tech Review	2	Media

VERIFIED INTELLIGENCE — 4 stories

VERI

Show HN: Does per-step reasoning effort save money in Claude Code?

HackerNews

+1 source

90%

HN 3

arXiv:2609.25244v2 Announce Type: replace-cross Abstract: Children increasingly interact with AI chatbots, making trust calibration essential to AI literacy. Prior research has examined children's trust in AI mainly as users evaluating systems built by others, rather than as...

<https://github.com/ifoster01/jev-effort>

VERI

Show HN: Notes on Agentic AI - A text-first guide for practicing engineers

HackerNews

+1 source

90%

HN 2

Modern life comes with an unending, auto-populating to-do list. It never ceases to amaze me how I can be doing nothing at all, minding my own business, and suddenly something needs to be taken care of. You're telling me a tree branch fell in the backyard and now I have to figure...

<https://github.com/mrsachindixit/agentixit>

VERI

Evolving Inspectable O-RAN Slicing xApps with LLMs

ArXiv CS.AI

+1 source

80%

arXiv:2609.27337v1 Announce Type: cross Abstract: Open RAN (O-RAN) slicing xApps must adapt resource allocations to changing channel conditions and traffic demands while meeting service-level agreements (SLAs). Deep reinforcement learning can produce adaptive policies, but their...

<https://arxiv.org/abs/2609.27337>

VERI

Meta is making a standalone Muse AI gadget

The Verge AI

+1 source

80%

Meta is building a dedicated hardware device for its new Muse AI agent. The product, called Muse Charm, was briefly shown off by Meta CEO Mark Zuckerberg at the end of tonight's Meta Connect presentation. It looks almost like a chunky smartwatch without the strap - just a big...

<https://www.theverge.com/tech/999750/muse-charm-meta-ai-hardware>

CONFIRMED SIGNALS — 119 stories

CONF

Show HN: Training a model to identify AI web content from structure alone

HackerNews

50%

HN 66

Hey HN! We're Vincent and Jochen from Sitefire (<https://sitefire.ai>). We have been working together for years, with backgrounds in RL/optimization at Stanford and software engineering from Technical University Munich (TUM). With Sitefire (YC W26), we help marketing teams get...

<https://arxiv.org/abs/2609.15369>

CONF

Show HN: AI-rete-RAG - a Rete rule engine decides, RAG explains why

HackerNews

50%

HN 43

Hi HN, I built ai-rete-rag because I kept seeing teams put an LLM in charge of decisions that need to be auditable (lending, fraud, clinical triage), then bolt on "guardrails" after the fact. It runs the two in series instead: 1. A pure-Python Rete engine evaluates YAML rules...

<https://ai-rete-rag.com/>

CONF

Show HN: Crest – Answer Claude Code approvals from your MacBook's notch

HackerNews 50% HN 5

No summary — see link for full article.

<https://crestnotch.app/claude-code-notch>

CONF

Meta introduces camera-free AI glasses

TechCrunch AI 50%

Meta says the camera-free glasses will be lighter and have up to 12 hours battery life.

<https://techcrunch.com/2026/09/23/meta-introduces-camera-free-ai-glasses/>

CONF

Anthropic says its biology lab has already found something big

TechCrunch AI 50%

But maybe the biggest reveal is that Anthropic has not let Claude run loose in its biology lab. Humans are still, so far, in the loop.

<https://techcrunch.com/2026/09/23/anthropic-says-its-biology-lab-has-already-found-some...>

0 CONF

Enveda secures \$311M to bring more nature-derived AI drugs into clinical trials

TechCrunch AI 50%

The round valued the AI biotech at \$2 billion. It is currently testing drugs that treat skin conditions and preserve weight loss after stopping GLP-1s.

<https://techcrunch.com/2026/09/23/enveda-secures-311m-to-bring-more-nature-derived-ai-d...>

1 CONF

YouTube Music gets more conversational with new AI features

TechCrunch AI 50%

Built directly into the YouTube Music app, Ask Music lets users describe what they want to hear in everyday language rather than searching for individual songs or artists.

<https://techcrunch.com/2026/09/23/youtube-music-gets-more-conversational-with-new-ai-fe...>

2 CONF

Even Americans who use AI every day are worried about it

TechCrunch AI 50%

The report suggests that greater exposure will not resolve the unease around the technology, nor reduce public support for AI regulation.

<https://techcrunch.com/2026/09/23/even-americans-who-use-ai-every-day-are-worried-about...>

3 CONF

YouTube will let you build your own algorithm with AI

TechCrunch AI 50%

YouTube's new custom feeds let users describe the videos they want to see in their own words, then use Gemini to build a personalized feed around the request.

<https://techcrunch.com/2026/09/23/youtube-will-let-you-build-your-own-algorithm-with-ai/>

4 CONF

YouTube releases new AI features for creators within its Studio app

TechCrunch AI 50%

YouTube is adding new features to generate ideas and monitor the performance of thumbnails.

<https://techcrunch.com/2026/09/23/youtube-releases-new-ai-features-for-creators-within-...>

5 CONF

Ema raises \$77M as AI starts eating into enterprise software and services

TechCrunch AI 50%

Ema has raised \$140 million to date and has more than 50 enterprise customers, including Google and Microsoft.

<https://techcrunch.com/2026/09/23/ema-raises-77m-as-ai-starts-eating-into-enterprise-so...>

6 CONF

Snorkel AI triples valuation to \$3.5B as demand for AI training data booms

TechCrunch AI 50%

The seven-year-old startup has raised a \$350 million Series E to fuel its data-as-a-service approach.

<https://techcrunch.com/2026/09/22/snorkel-ai-triples-valuation-to-3-5b-as-demand-for-ai...>

7 CONF

The Download: India's smart glasses menace and AI's trillion-dollar gamble

MIT Tech Review 50%

This is today's edition of The Download, our weekday newsletter that provides a daily dose of what's going on in the world of technology. Smart glasses are already causing havoc in India When Shubnam saw an Instagram video of a Delhi protest they had attended, they realized a...

<https://www.technologyreview.com/2026/09/23/1144966/the-download-india-smart-glasses-ai...>

8 CONF

The AI Hype Index: AI loves cheating

MIT Tech Review 50%

Brace yourself: It turns out AI is being optimized for cheating. OpenAI's agents hacked into Hugging Face to get the answers to a cybersecurity test. Next, they solved a prestigious math problem (or just stole from two top mathematicians' answer sheets). Anthropic's models have...

<https://www.technologyreview.com/2026/09/23/1144940/ai-hype-index-ai-loves-cheating/>

9 CONF

Anthropic's biolab made a discovery it's comparing to Crispr

The Verge AI 50%

Anthropic says its AI Claude has "autonomously discovered" a new enzyme system similar to machinery behind the powerful gene-editing tool Crispr. It's the first result from Anthropic's newly-launched wet lab and an early test of Claude's usefulness for science as the company...

<https://www.theverge.com/ai-artificial-intelligence/999470/anthropic-biolab-claude-crispr>

0 CONF

Silent Failures in Agent-Tool Interaction: An Audit of ToolUniverse

ArXiv CS.AI 50%

arXiv:2609.26836v1 Announce Type: new Abstract: Agentic AI systems are increasingly adopting automated pipelines that integrate multiple tools. While prior research and benchmarks have studied about task success and task completion of these agentic systems, the research about...

<https://arxiv.org/abs/2609.26836>

1 CONF

Building Socio-Affective Artificial Intelligence for Interactive Multi-Agent Simulations

ArXiv CS.AI 50%

arXiv:2609.26927v1 Announce Type: new Abstract: The objective of this article is to provide design principles and a software architecture for enabling interaction between humans and multiple agents in simulated dynamic worlds. This connects the current era of general artificial...

<https://arxiv.org/abs/2609.26927>

2 CONF

Which Objectives Need a Dial? Predicting Objective Conflict and Covering Trade-offs in Steerable Pluralistic Alignment

ArXiv CS.AI 50%

arXiv:2609.26929v1 Announce Type: new Abstract: People hold diverse, sometimes conflicting values, so no single aligned model can satisfy everyone. Pluralistic alignment therefore calls for steerable models that can balance competing objectives differently. Multi-Objective...

<https://arxiv.org/abs/2609.26929>

3 CONF

Escaping Python Dependency Hell: A Hybrid Replay-and-Repair Pipeline for Python Dependency Resolution

ArXiv CS.AI 50%

arXiv:2609.26952v1 Announce Type: new Abstract: Dependency conflicts in Python ecosystems arise from incompatible version constraints, missing packages, and undocumented compatibility relationships, causing many real-world code snippets to fail at execution. This paper presents...

<https://arxiv.org/abs/2609.26952>

4 CONF

DCRL: Decoupling and Coupling Reinforcement Learning via Policy-Reward Manifold Alignment

ArXiv CS.AI 50%

arXiv:2609.27572v1 Announce Type: cross Abstract: Reinforcement learning (RL) has emerged as a key paradigm for improving the reasoning capabilities of large language models (LLMs). However, existing reward systems, such as rule-based and reward-model-based, often exhibit issues...

<https://arxiv.org/abs/2609.27572>

5 CONF

Training Intelligent Voice Assistant Wakeup with Controllable Synthetic Conversations

ArXiv CS.AI 50%

arXiv:2609.27037v1 Announce Type: new Abstract: Wake word detection is a critical component of virtual assistants, serving as the gateway to seamless user interactions. This paper introduces a novel wake-up system that extends traditional direct keyword detection with contextual...

<https://arxiv.org/abs/2609.27037>

6 CONF

Math Reasoning in LLMs is Organized by Approach, Not Topic

ArXiv CS.AI 50%

arXiv:2609.27041v1 Announce Type: new Abstract: Mathematical reasoning benchmarks are typically organized by topic, but language models may organize their internal computation by reusable reasoning approach instead. In this paper, we investigate whether open math-capable LLMs...

<https://arxiv.org/abs/2609.27041>

7 CONF

Propose, Don't Judge: An Anytime-Valid Referee for LLM Agents That Mine Investment Factors

ArXiv CS.AI 50%

arXiv:2609.27051v1 Announce Type: new Abstract: Language-model agents now run the whole of quantitative factor research: they propose investment factors, backtest them, select the survivors and retire them. We ask which of those jobs an agent should keep. Our answer is governed...

<https://arxiv.org/abs/2609.27051>

8 CONF

Policy-as-Skill: Governed LLM Decision Support with Evidence, Deterministic Control, and Audit

ArXiv CS.AI 50%

arXiv:2609.27087v1 Announce Type: new Abstract: Organizations increasingly use LLMs for policy, compliance, risk, and operational decision support, requiring evidence validation, review routing, version control, and auditability. We introduce Policy-as-Skill (PaS), a modular...

<https://arxiv.org/abs/2609.27087>

9 CONF

Provably Complete Generalized Planning with LLMs

ArXiv CS.AI 50%

arXiv:2609.27105v1 Announce Type: new Abstract: Generalized planning aims to compute a plan that solves all instances of a planning domain. Recent work has used LLMs to automatically generate and debug such generalized plans in the form of Python programs and achieved perfect...

<https://arxiv.org/abs/2609.27105>

0 CONF

The Path Matters: Evaluating Small Language Models Beyond Answer Accuracy in KGQA

ArXiv CS.AI 50%

arXiv:2609.27669v1 Announce Type: cross Abstract: Small language models (SLMs) are increasingly paired with knowledge graphs (KGs), yet end-to-end KG question answering conflates graph access, search, navigation, reasoning, and answer generation. This coupling makes it difficult...

<https://arxiv.org/abs/2609.27669>

1 CONF

Safety Nudges: User-Facing Interventions for Real-Time AI Risk Awareness

ArXiv CS.AI 50%

arXiv:2609.26865v1 Announce Type: cross Abstract: Conversational AI systems can pose safety risks to their users such as hallucination, sycophancy, overconfidence, and anthropomorphism, but these risks are difficult for users to detect during everyday use. We introduce Safety...

<https://arxiv.org/abs/2609.26865>

2 CONF

Sparse-Observation Atmospheric Thermal Forecasting with Physics-Informed Neural Networks for Climate-Aware Digital Twins

ArXiv CS.AI 50%

arXiv:2609.27290v1 Announce Type: new Abstract: Short-horizon forecasts of atmospheric temperature are needed to support climate-aware digital-twin systems, but such forecasts must be produced where thermal observations are incomplete. This study evaluates a physics-informed...

<https://arxiv.org/abs/2609.27290>

3 CONF

Alignment Inertia: Auditing the Durability of Training Data Influence Through Policy Override Resistance

ArXiv CS.AI 50%

arXiv:2609.27333v1 Announce Type: new Abstract: Platform operators increasingly rely on system prompts and fine-tuning to govern model behavior, yet it remains unclear how reliably these interventions override behavior inherited from prior training. We propose Override Success...

<https://arxiv.org/abs/2609.27333>

4 **CONF****Just-in-Time Memory: Learning to Curate Task-Adaptive Memory for LLM Agents**

ArXiv CS.AI 50%

arXiv:2609.27334v1 Announce Type: new Abstract: Agentic memory systems reuse past experience to improve future performance, yet most existing designs curate memory at write time: once a task is completed, its trajectory is distilled into a fixed artifact, such as a reflection,...

<https://arxiv.org/abs/2609.27334>
5 **CONF****CART: Closed-Loop Adaptive Red Teaming for Large Language Models**

ArXiv CS.AI 50%

arXiv:2609.27336v1 Announce Type: new Abstract: Automated red teaming often replays a fixed set of prompts, which measures known risks but cannot learn from failures found during testing. We present CART (Closed-Loop Adaptive Red Teaming), a framework that uses each result to...

<https://arxiv.org/abs/2609.27336>
6 **CONF****FDE-Bench: Evaluating LLM Agents for Deployment Environment Configuration**

ArXiv CS.AI 50%

arXiv:2609.27571v1 Announce Type: cross Abstract: Deployment requires an agent to turn application code into a running system whose services connect, become ready, and remain observable. FDE-Bench evaluates this capability with 136 deployment-configuration tasks spanning Docker...

<https://arxiv.org/abs/2609.27571>
7 **CONF****Benchmarking Text-to-SQL under Role-Based Access Control**

ArXiv CS.AI 50%

arXiv:2607.22115v2 Announce Type: replace-cross Abstract: Given a database S and a natural language question Q, text-to-SQL systems aim to generate an SQL query that correctly answers Q when executed against S. Currently, popular text-to-SQL benchmarks mostly assume unrestricted...

<https://arxiv.org/abs/2607.22115>
8 **CONF****Not What You Meant: Can LLMs Follow a Specified Negation Semantics?**

ArXiv CS.AI 50%

arXiv:2609.27517v1 Announce Type: new Abstract: Negation does not carry a uniform interpretation across domains. In legal, regulatory, and medical reasoning, the intended interpretation depends on the reading in force -- open- versus closed-world, two- versus three-valued, and...

<https://arxiv.org/abs/2609.27517>
9 **CONF****State-Grounded Conditioning: Wrapping User-Facing LLM Agents Where Direction Depends on Live State**

ArXiv CS.AI 50%

arXiv:2609.27606v1 Announce Type: new Abstract: We introduce State-Grounded Conditioning (SGC), a design principle for user-facing LLM agents that must condition on live user state (game state, session history, live inventory), and a distinct failure class we call direction...

<https://arxiv.org/abs/2609.27606>

0 CONF

BiCFlow-MER: Orchestrating Discriminative and Generative Multimodal Emotion Recognition via Conditional Transport

ArXiv CS.AI 50%

arXiv:2609.27615v1 Announce Type: new Abstract: In multimodal emotion recognition (MER), human affective states are inferred by integrating complementary cues from multiple modalities. In audio-text MER, affective cues are often entangled with speaker style and lexical content,...

<https://arxiv.org/abs/2609.27615>

1 CONF

Evolutionary Stability Does Not Guarantee Learning Accessibility: A Multi-Agent Reinforcement Learning Perspective on Cooperation Emergence

ArXiv CS.AI 50%

arXiv:2609.27664v1 Announce Type: new Abstract: Cooperation emergence is a central problem in multi-agent systems because decentralized agents must coordinate while adapting to the changing behavior of others. Evolutionary game theory identifies strategically stable outcomes,...

<https://arxiv.org/abs/2609.27664>

2 CONF

Comparative Evaluation of Static Embedding Models for HTTP Request Anomaly Detection

ArXiv CS.AI 50%

arXiv:2609.26860v1 Announce Type: cross Abstract: Web applications are increasingly targeted by cyberattacks that exploit HTTP requests to evade security mechanisms. Traditional web application firewalls (WAFs) rely on rule-based approaches that often exhibit high false positive...

<https://arxiv.org/abs/2609.26860>

3 CONF

Reporting Under Pressure: Separating Factual and Tonal Sycophancy in LLM Statistical Analysis

ArXiv CS.AI 50%

arXiv:2609.27756v1 Announce Type: new Abstract: Large language models are increasingly asked to analyze data and report what the results mean, a task distinct from the belief- or preference-alignment settings studied in most sycophancy research. We test whether editorial framing...

<https://arxiv.org/abs/2609.27756>

4 CONF

Alignment of LRMs via Counter-Aligned Few-Shot Conversation Exposure

ArXiv CS.AI 50%

arXiv:2609.27763v1 Announce Type: new Abstract: Large Reasoning Models (LRMs) rely on explicit chain-of-thought (CoT) reasoning and large context windows to achieve strong performance on complex tasks, but these features also introduce new attack surfaces. We show that LRMs'...

<https://arxiv.org/abs/2609.27763>

5 CONF

Agentic Governance and Adversarial Verification for Policy-Constrained LLM Healthcare Appeal Generation

ArXiv CS.AI 50%

arXiv:2609.27844v1 Announce Type: new Abstract: Claim denial management costs U.S. healthcare approximately \$260 billion annually in administrative overhead. Large Language Models (LLMs) and Retrieval-Augmented Generation (RAG) can produce fluent clinical text, but single-agent...

<https://arxiv.org/abs/2609.27844>

6 **CONF****Reachable Global Optimization in AI Systems: How Global Is Global?****ArXiv CS.AI** **50%**

arXiv:2609.27855v1 Announce Type: new Abstract: AI systems increasingly claim to optimize prompts, policies, architectures, plans, tool-use trajectories, reasoning traces, and test-time computation. This paper argues that such claims are underspecified unless they state the...

<https://arxiv.org/abs/2609.27855>

7 **CONF****A hierarchy of faithfulness criteria for knowledge base completion****ArXiv CS.AI** **50%**

arXiv:2609.27863v1 Announce Type: new Abstract: Knowledge graph completion is evaluated by ranking observed triples above randomly corrupted ones, which treats every unobserved fact as false. When the object being completed is a description logic knowledge base rather than a...

<https://arxiv.org/abs/2609.27863>

8 **CONF****Learning What to Activate: Combinatorial Capability Allocation for Long-Horizon Multimodal Agents****ArXiv CS.AI** **50%**

arXiv:2609.27869v1 Announce Type: new Abstract: Long-horizon multimodal agents rely on specialized capabilities for perception, retrieval, reasoning, verification, and execution. Existing designs typically activate a fixed capability set or invoke a predefined workflow,...

<https://arxiv.org/abs/2609.27869>

9 **CONF****MicroQonv: Reshaping Convolution Tensors for Efficient Microscaling in Training and Inference****ArXiv CS.AI** **50%**

arXiv:2609.28358v1 Announce Type: cross Abstract: Microscaling quantization techniques are increasingly used to represent neural network parameters with 8 bits or fewer while preserving near-full precision accuracy. However, applying these methods efficiently in convolutional...

<https://arxiv.org/abs/2609.28358>

0 **CONF****ANO: Robust Policy Optimization via Bounded, Redescending Gain Fields****ArXiv CS.AI** **50%**

arXiv:2605.02320v3 Announce Type: replace Abstract: Proximal Policy Optimization (PPO) dominates reinforcement learning and LLM alignment, yet its hard-clipping mechanism and unconstrained alternatives (e.g., SPO) sit at two extremes of a stability-efficiency dilemma. We argue...

<https://arxiv.org/abs/2605.02320>

1 **CONF****An Open Pipeline and Dashboard for Systemic-Risk Evidence under the EU AI Act's Code of Practice****ArXiv CS.AI** **50%**

arXiv:2609.28335v1 Announce Type: new Abstract: Claims about AI safety reach audiences well beyond the AI community, yet many rely on opaque evidence or static assessments, when supporting evidence is accessible at all. We present the Systemic Risk Index, an open evaluation...

<https://arxiv.org/abs/2609.28335>

2 CONF

StudentBench: AI and human tutoring yield equivalent GRE learning gains

ArXiv CS.AI 50%

arXiv:2609.28470v1 Announce Type: new Abstract: Artificial intelligence offers an unprecedented opportunity to augment human capabilities, yet progress at the frontier has focused primarily on advancing model capabilities. We introduce StudentBench, a suite of AI teaching...

<https://arxiv.org/abs/2609.28470>

3 CONF

Signal2Symbol: Neuro-Symbolic Temporal Reasoning for Explainable Physiological Time-Series Anomaly Detection

ArXiv CS.AI 50%

arXiv:2609.26820v1 Announce Type: cross Abstract: Physiological time series such as electrocardiograms (ECG) and electroencephalograms (EEG) exhibit complex temporal structure, substantial acquisition variability, and a strong need for transparent decision-making. Although deep...

<https://arxiv.org/abs/2609.26820>

4 CONF

Bridging LLM Serving and CXL-SSDs with Chunk-Aware KV Cache Management

ArXiv CS.AI 50%

arXiv:2609.26828v1 Announce Type: cross Abstract: NAND-backed storage offers the capacity needed to scale LLM prefix caching, but its block I/O path incurs CPU cache contention and host-DRAM staging in addition to NAND latency. Our characterization shows that these interface...

<https://arxiv.org/abs/2609.26828>

5 CONF

Validation and Simulation Catch Different Errors: Four Levels of Evaluation for LLM-Generated Circuits

ArXiv CS.AI 50%

arXiv:2609.26830v1 Announce Type: cross Abstract: Simulation success is not equivalent to structural correctness for LLM-generated circuits. We define and measure four evaluation levels -- schema validity, topological validity, backend executability, and component-set agreement...

<https://arxiv.org/abs/2609.26830>

6 CONF

Spec2COBOLRot: An Agentic-AI Degradation Loop for Realistic COBOL Corpus Generation

ArXiv CS.AI 50%

arXiv:2609.26835v1 Announce Type: cross Abstract: COBOL remains widely deployed, yet representative corpora reflecting real production code are rarely available, limiting rigorous benchmarking of modernization approaches. We propose a systematic agentic AI pipeline for...

<https://arxiv.org/abs/2609.26835>

7 CONF

RAIN: Region-Aware Inversion Network for Semantic Watermark Extraction

ArXiv CS.AI 50%

arXiv:2609.14856v2 Announce Type: replace-cross Abstract: Semantic watermarks for diffusion models embed ownership information into the generative process while preserving perceptual quality, but Gaussian-Shading extraction conventionally requires multi-step diffusion inversion...

<https://arxiv.org/abs/2609.14856>

8 CONF

COPE: Continual Personalization of LLMs under Sparse User Feedback via User Embeddings and Self-Evaluation

ArXiv CS.AI 50%

arXiv:2609.26853v1 Announce Type: cross Abstract: While Large Language Models (LLMs) have achieved remarkable results across various benchmarks, their alignment with normative values often results in homogenized responses that fail to address diverse user preferences. Existing...

<https://arxiv.org/abs/2609.26853>

9 CONF

QUARTET: Quad-branch cross-Attention and Random-walk Traces for Enhancing Transformers on Relational Graphs

ArXiv CS.AI 50%

arXiv:2609.26855v1 Announce Type: cross Abstract: Relational Deep Learning (RDL) models multi-table databases as heterogeneous temporal graphs, and graph transformers currently achieve state-of-the-art performance on benchmarks like RelBench. However, the current leading model,...

<https://arxiv.org/abs/2609.26855>

10 CONF

FLINT: Fast Lightweight Inference for Traversability

ArXiv CS.AI 50%

arXiv:2609.26857v1 Announce Type: cross Abstract: Navigation in off-road conditions is challenging due to the lack of structure. There is no fixed vocabulary for what is traversable. The traversability depends on both the environment and the embodiment's dynamics. Neither of...

<https://arxiv.org/abs/2609.26857>

1 CONF

COMED: The Missing Middle Between Routing and Collaboration in Multi-LLM Inference

ArXiv CS.AI 50%

arXiv:2609.26913v1 Announce Type: cross Abstract: No single Large Language Model (LLM) is uniformly reliable across queries, motivating multi-model inference systems that either route among models or combine their outputs. However, routing stops after selecting an initial model,...

<https://arxiv.org/abs/2609.26913>

2 CONF

On Preference Coverage Collapse from Hindsight Relabeling in Multi-Objective Reinforcement Learning

ArXiv CS.AI 50%

arXiv:2609.26918v1 Announce Type: cross Abstract: Hindsight relabeling which retroactively replacing a transition's goal with the outcome the agent actually achieved is an effective tool for improving sample-efficiency in Reinforcement Learning (RL). A natural extension to...

<https://arxiv.org/abs/2609.26918>

3 CONF

Experts Rise Where LLMs Disagree: Using Cross-Model Disagreement to Target Expert Effort in LLM Codebook Revision for Large-Scale Annotation

ArXiv CS.AI 50%

arXiv:2609.26926v1 Announce Type: cross Abstract: Large-scale text annotation brings expert insight to millions of documents, often through a codebook that AI annotators follow. Developing a robust codebook, however, takes months. Large language models (LLMs) could speed this...

<https://arxiv.org/abs/2609.26926>

4 **CONF****Recognized but Not Produced: A Generation Benchmark for Culturally Specific Kinship Terms****ArXiv CS.AI** **50%**

arXiv:2609.26942v1 Announce Type: cross Abstract: Current literature evaluates large language models (LLMs) on multilingual kinship understanding using multiple choice benchmarks, treating it as a recognition problem. We instead prompt five open weight LLMs to generate kinship...

<https://arxiv.org/abs/2609.26942>

5 **CONF****An open benchmark for machine learning-based polymer property prediction****ArXiv CS.AI** **50%**

arXiv:2609.27036v1 Announce Type: cross Abstract: Polymer property prediction lacks open, standardized benchmarks that enable rigorous comparison of machine-learning methods, with existing resources covering only a narrow fraction of polymer architectures, such as homopolymers....

<https://arxiv.org/abs/2609.27036>

6 **CONF****The Gaussian Is Enough: Flow-Matching Priors Do Not Help When Fine-Tuning Large Behavior Models****ArXiv CS.AI** **50%**

arXiv:2609.27070v1 Announce Type: cross Abstract: Modern robot imitation learning increasingly relies on generative policies based on diffusion or flow-matching models, which generate actions by transforming samples from a prior distribution. A key question is whether the choice...

<https://arxiv.org/abs/2609.27070>

7 **CONF****Crossflow: Prefill-Decode Elasticity for Agentic LLM Serving****ArXiv CS.AI** **50%**

arXiv:2609.27085v1 Announce Type: cross Abstract: As serving capacity demand surpasses that of training, serving efficiency becomes increasingly important. Prefill-decode (P/D) disaggregation improves serving efficiency through specialization and isolation of the two phases....

<https://arxiv.org/abs/2609.27085>

8 **CONF****Local Evidence and Geometric Readout Repair in Trained GNNs****ArXiv CS.AI** **50%**

arXiv:2609.27092v1 Announce Type: cross Abstract: Many node-classification GNNs apply a linear classifier to a nonnegative mixture of local messages. An error can reflect either poor mixture weights or a reachable logit set poorly positioned for the classifier. We separate these...

<https://arxiv.org/abs/2609.27092>

9 **CONF****A Hierarchy-Aware Video-Language Model Evaluation and Hyperbolic Baseline for Surgery****ArXiv CS.AI** **50%**

arXiv:2609.27139v1 Announce Type: cross Abstract: Surgical procedures follow a phase-to-step hierarchy, yet the video-language models used to recognize them are evaluated with flat per-level metrics that ignore cross-level coherence and error structure. In this paper we make two...

<https://arxiv.org/abs/2609.27139>

0 **CONF****The Like Trap: Multi-Stage Poisoning against Agents in Similarity-based Recommendation Systems****ArXiv CS.AI** **50%**

arXiv:2609.27155v1 Announce Type: cross Abstract: With recent advancements in large language models (LLMs) and LLM-based agents, these agents are becoming increasingly autonomous and gaining broader access to act on users' behalf on the internet. However, the vulnerability of...

<https://arxiv.org/abs/2609.27155>

1 **CONF****Count Evidence, Not Sentences: Tempered Evidence Fusion of LLM Judgments for Long-Text Value Measurement****ArXiv CS.AI** **50%**

arXiv:2609.27165v1 Announce Type: cross Abstract: Large language models (LLMs) are increasingly used to measure public value orientations from long social media posts, yet such posts often mix background, quotations, concessions, and only a few stance-bearing sentences. Existing...

<https://arxiv.org/abs/2609.27165>

2 **CONF****Phonemizing User-Generated Text: A Benchmark, Taxonomy, and Compositional Approach****ArXiv CS.AI** **50%**

arXiv:2609.27205v1 Announce Type: cross Abstract: Text-to-speech systems increasingly process user-generated text (UGT) such as ppl and imo, whose pronunciation must be inferred from the canonical rather than surface form. We introduce UGTPhon, the first grapheme-to-phoneme...

<https://arxiv.org/abs/2609.27205>

3 **CONF****Meet, Compare, or Abstain: LatWeave for Deterministic Multi-Hop Question Answering on Knowledge Lattices****ArXiv CS.AI** **50%**

arXiv:2609.27225v1 Announce Type: cross Abstract: Probabilistic question-answering systems -- whether large language models (LLMs) themselves, retrieval-augmented generation (RAG), or trained multi-hop retrievers -- conflate "what is known" and "how to reason" into a single...

<https://arxiv.org/abs/2609.27225>

4 **CONF****Combining LLMs and Genetic Search for ARC-AGI-2****ArXiv CS.AI** **50%**

arXiv:2609.27242v1 Announce Type: cross Abstract: LLMs can generate programs for ARC-AGI-2 tasks, but the provided compute only allows a small number of attempts to generate, debug and validate solutions. Genetic algorithms can search and test many more programs, but random...

<https://arxiv.org/abs/2609.27242>

5 **CONF****Listening and Mirroring: The Effects of Verbal Attunement and Behavioral Mimicry on Social and Empathic Perceptions of Embodied AI Agents in VR****ArXiv CS.AI** **50%**

arXiv:2609.27246v1 Announce Type: cross Abstract: As embodied agents take on increasingly social and relational roles in VR, visual realism and embodiment alone may be insufficient; users must also perceive these agents as emotionally attuned, supportive, and humanlike. Prior...

<https://arxiv.org/abs/2609.27246>

6 CONF

KITE: KV-Invariant Transformer Expansion for Efficient Agentic LLM Scaling

ArXiv CS.AI 50%

arXiv:2609.27294v1 Announce Type: cross Abstract: Scaling a language model is not only a question of final quality: the architectural choice determines how much computation is spent during training, prompt processing, and autoregressive decoding to achieve certain model quality....

<https://arxiv.org/abs/2609.27294>

7 CONF

Teach-to-Crash: A Closed-Loop Student-Teacher LLM Framework for Collision-Inducing Test Scenario Generation

ArXiv CS.AI 50%

arXiv:2609.27296v1 Announce Type: cross Abstract: Validating Autonomous Driving Systems (ADS) in simulation requires testing architectures that can discover rare, safety-critical failures while generating scenarios that are executable, diverse, and useful for downstream failure...

<https://arxiv.org/abs/2609.27296>

8 CONF

Turning Safety into Competence: Minimally Exploitable Robot Policies via Safety-Filtered Reinforcement Learning

ArXiv CS.AI 50%

arXiv:2609.27312v1 Announce Type: cross Abstract: Robots deployed for competitive tasks must outmaneuver their opponents without sacrificing safety. Existing approaches, including safe reinforcement learning (RL), train a single policy to achieve task success and avoid failures...

<https://arxiv.org/abs/2609.27312>

9 CONF

Constraint-Driven Context Engineering: Designing Domain Interfaces for AI Systems

ArXiv CS.AI 50%

arXiv:2609.27354v1 Announce Type: cross Abstract: Generative AI systems are increasingly deployed to address domain problems. These systems operate under technical, regulatory, institutional, and normative constraints that define acceptable AI behaviour and outcomes within their...

<https://arxiv.org/abs/2609.27354>

0 CONF

Quantization-Robust Unlearning through the Lens of Retain-Forget Loss Landscapes Interaction

ArXiv CS.AI 50%

arXiv:2609.27355v1 Announce Type: cross Abstract: Unlearning ensures LLM compliance by removing the influence of private or copyrighted training data. However, since LLM models typically undergo post-training compression, like quantization, in practical deployment, it has been...

<https://arxiv.org/abs/2609.27355>

1 CONF

Neither Silence nor Overlap Is Failure: Intent-Conditioned Evaluation of Turn-Taking in Full-Duplex Spoken Dialogue Models

ArXiv CS.AI 50%

arXiv:2609.27372v1 Announce Type: cross Abstract: Benchmarks for full-duplex spoken dialogue models score turn-taking with binary fixed-window rules that reward immediate response or silence by completeness of the prior turn. We argue that the appropriateness of a response...

<https://arxiv.org/abs/2609.27372>

2 **CONF****What Looks Like a Capability Limit in Vision-Language Models Is a Readout Limit**

ArXiv CS.AI 50%

arXiv:2609.27408v1 Announce Type: cross Abstract: Benchmarks for vision-language models offer their answer choices in some convention: a letter, a color name, a pixel coordinate. That convention is treated as neutral. We find it is not, and that the limits a benchmark reports...

<https://arxiv.org/abs/2609.27408>

3 **CONF****Passing: An Endless Journey through Reconstructed Spacetime with AI-Generated Sound**

ArXiv CS.AI 50%

arXiv:2609.27489v1 Announce Type: cross Abstract: This paper introduces Passing, an interactive audiovisual installation that generates an endless journey from a single continuous monorail-window recording by reconstructing it as a spatiotemporal volume. Rather than replaying...

<https://arxiv.org/abs/2609.27489>

4 **CONF****NV-Reason-CT: 3D Visual Language Model for CT Analysis**

ArXiv CS.AI 50%

arXiv:2609.27511v1 Announce Type: cross Abstract: We present NV-Reason-CT, a generative vision--language model for chest and abdominal CT combining native 3D visual encoding with radiologist-guided reasoning. The model couples a native 3D vision transformer with a language...

<https://arxiv.org/abs/2609.27511>

5 **CONF****PhyMo: A Physical-Field Modality for Multimodal AI4Physics**

ArXiv CS.AI 50%

arXiv:2609.27554v1 Announce Type: cross Abstract: Multimodal learning is emerging as a powerful paradigm for AI for Physics (AI4Physics), where predicting physical systems requires the joint interpretation of heterogeneous observations, measurements, and domain knowledge...

<https://arxiv.org/abs/2609.27554>

6 **CONF****ProteinJEPa: Latent prediction improves protein language model pretraining**

ArXiv CS.AI 50%

arXiv:2605.07554v2 Announce Type: replace-cross Abstract: Protein language models are trained primarily with masked language modeling (MLM), which predicts masked amino-acid identities. Joint-embedding predictive architectures (JEPa) instead predict latent representations, but...

<https://arxiv.org/abs/2605.07554>

7 **CONF****InGuard: Towards Generalized Inner Guardrail for Safe Text-to-Image Generation**

ArXiv CS.AI 50%

arXiv:2609.27620v1 Announce Type: cross Abstract: Modern text-to-image (T2I) models generate high-quality images from arbitrary user prompts, yet they can just as easily produce not-safe-for-work (NSFW) content. Conventional outer guardrails consist of two components: a prompt...

<https://arxiv.org/abs/2609.27620>

8 **CONF****Compliant AI Infrastructure for Regulated Finance: A tiered multi-agent framework with DLT audit trails for financial operations in DACH**

ArXiv CS.AI 50%

arXiv:2609.27632v1 Announce Type: cross Abstract: We present a compliance-first architecture for AI in regulated finance that treats regulation as an orientation layer rather than a deterministic ruleset. A matrix of regulatory intent and exposure provides a compact...

<https://arxiv.org/abs/2609.27632>

9 CONF

AI-Driven Neural Surrogates for In Silico Design of Cognitive-Affective Neuromodulation Targets

ArXiv CS.AI 50%

arXiv:2609.27729v1 Announce Type: cross Abstract: In neuropsychiatry, the primary goal is often not only to decode brain activity but to change it, for example to lessen a negative affective bias or an overly salient memory. Motivated by control theory, we develop an AI-driven...

<https://arxiv.org/abs/2609.27729>

0 CONF

Backdoors Leave Structural Traces: FedMAST for Backdoor Detection and Containment in Federated Learning

ArXiv CS.AI 50%

arXiv:2609.27760v1 Announce Type: cross Abstract: Federated learning enables distributed training of a shared model without requiring clients to share their raw data. However, its reliance on the integrity of the client-submitted updates exposes the global model to stealthy...

<https://arxiv.org/abs/2609.27760>

1 CONF

Beyond Unsafe Detection: Counterfactually Anchored Evidence Attribution for Multi-Turn LLM Safety Failures

ArXiv CS.AI 50%

arXiv:2609.27773v1 Announce Type: cross Abstract: As Large Language Models (LLMs) move from conversational assistants to advanced agentic systems, guardrail failures can convert adversarial intents into harmful executions. However, most guardrail evaluation frameworks focus only...

<https://arxiv.org/abs/2609.27773>

2 CONF

LabourCrew: A Multi-Agent RAG Framework for Trustworthy Adversarial Deliberation and Statutory Reasoning over Labour Law

ArXiv CS.AI 50%

arXiv:2609.27814v1 Announce Type: cross Abstract: In statutory question answering, every claim must be traceable to evidence, not merely relevant, since unverifiable labour-rights answers carry serious legal consequences. Current systems fall short: single-pass RAG cannot detect...

<https://arxiv.org/abs/2609.27814>

3 CONF

Evaluating ADC-only deep learning pipelines for breast cancer detection and segmentation using standalone diffusion-weighted MRI

ArXiv CS.AI 50%

arXiv:2609.27815v1 Announce Type: cross Abstract: Dynamic contrast-enhanced (DCE) imaging is the gold standard technique for the detection and characterization of breast cancer using magnetic resonance imaging (MRI). However, DCE-MRI requires long acquisition times and the...

<https://arxiv.org/abs/2609.27815>

4 CONF

Safe Multi-Robot Coordination via VLM-LLM Reasoning and Reachability Analysis

ArXiv CS.AI 50%

arXiv:2609.27816v1 Announce Type: cross Abstract: Safe coordination in heterogeneous machine-to-machine (M2M) robotic systems is challenging when robots differ in sensing capabilities, environmental awareness, and motion execution roles. This paper presents a centralized...

<https://arxiv.org/abs/2609.27816>

5 CONF

Schrödinger's Code Repository: Have LLMs Learned SWE-bench or Memorized It?

ArXiv CS.AI 50%

arXiv:2609.27891v1 Announce Type: cross Abstract: Repository-level coding benchmarks have become the standard for evaluating coding agents, yet they inherently suffer from data leakage because they are built upon popular open-source repositories repeatedly used for training...

<https://arxiv.org/abs/2609.27891>

6 CONF

Compliant with Local Controls, Collectively Discriminatory. A Governance Architecture for Multi-Agent AI in Regulated Finance

ArXiv CS.AI 50%

arXiv:2609.27994v1 Announce Type: cross Abstract: Financial institutions are beginning to deploy agentic workflows in credit, fraud, collections, compliance, and operational control. Governance remains largely component-centric: each model or agent is specified, tested,...

<https://arxiv.org/abs/2609.27994>

7 CONF

TEMPS: Temporal Sentence Embeddings for Temporal Information Retrieval

ArXiv CS.AI 50%

arXiv:2609.28048v1 Announce Type: cross Abstract: Modern information retrieval (IR) systems rarely represent time, yet many information needs depend on it: in clinical, journalistic, and legal search, when an event occurred can decide whether a document is relevant. Dense...

<https://arxiv.org/abs/2609.28048>

8 CONF

Prompt, Probe, Train, or Annotate? Single-camera sports video understanding in amateur settings

ArXiv CS.AI 50%

arXiv:2609.28049v1 Announce Type: cross Abstract: Video understanding is usually benchmarked on curated, single-actor, or professionally filmed clips, and a strong score there is routinely read as evidence a model is robust enough for deployment. Amateur team sport is a useful,...

<https://arxiv.org/abs/2609.28049>

9 CONF

Curriculum Learning with GNN-based Reinforcement Learning for Job Shop Scheduling

ArXiv CS.AI 50%

arXiv:2609.28085v1 Announce Type: cross Abstract: The job shop scheduling problem is a challenging combinatorial optimization problem, and recent reinforcement learning approaches using graph neural networks have shown promise for learning scheduling policies directly from...

<https://arxiv.org/abs/2609.28085>

00 CONF

Can LLMs Catch a Rigged Backtest? A Clean-Control Calibration Benchmark

ArXiv CS.AI 50%

arXiv:2609.28090v1 Announce Type: cross Abstract: Backtest auditing is a calibration problem: high flaw recall is not useful when the model falsely flags matched clean strategies. We build a 96-item paired benchmark in which every flawed backtest has a clean control that holds...

<https://arxiv.org/abs/2609.28090>

01 CONF

Field-of-View Extension in Dental Cone-Beam CT via Implicit Neural Representations and Diffusion Model-Based Refinement

ArXiv CS.AI 50%

arXiv:2609.28110v1 Announce Type: cross Abstract: Dental cone-beam computed tomography (CBCT) systems often employ detector configurations that provide a truncated field of view (FOV) that only captures a small part of the patient's anatomy. In this work, we aim to reconstruct...

<https://arxiv.org/abs/2609.28110>

02 CONF

"We'll Fix It Later": Education, AI, and the Deferral of Privacy in EdTech

ArXiv CS.AI 50%

arXiv:2609.28137v1 Announce Type: cross Abstract: Educational technology (EdTech) platforms collect highly sensitive student data, including behavioral logs, disability records, and academic histories. However, privacy considerations are often postponed rather than treated as a...

<https://arxiv.org/abs/2609.28137>

03 CONF

Do Center Biases Propagate? Robustness of Pathology Foundation Models in Whole-Slide Image Classification

ArXiv CS.AI 50%

arXiv:2609.28231v1 Announce Type: cross Abstract: Pathology foundation models (PFMs) have transformed computational pathology through powerful representation learning from histopathological images. PFMs provide rich, discriminative representations for whole slide image (WSI)...

<https://arxiv.org/abs/2609.28231>

04 CONF

Beyond Poetry: Can Large Language Models Generate Classical Arabic Maqamat?

ArXiv CS.AI 50%

arXiv:2609.28245v1 Announce Type: cross Abstract: Large language models (LLMs) have shown strong performance in creative text generation, yet their ability to produce culturally grounded and stylistically constrained literary forms remains underexplored. Prior work has focused...

<https://arxiv.org/abs/2609.28245>

05 CONF

AnchorReasoning: A Visual Grounding and Causal Reasoning Dataset in Long-Tail Autonomous Driving Scenarios

ArXiv CS.AI 50%

arXiv:2609.28366v1 Announce Type: cross Abstract: Vision-language models (VLMs) offer a promising approach to long-tail autonomous driving, but existing driving datasets provide limited supervision for connecting decision-critical visual evidence with reasoning and planning. We...

<https://arxiv.org/abs/2609.28366>

06 CONF

Shopping by algorithm: How agentic AI deploys human heuristics as a surrogate consumer

ArXiv CS.AI 50%

arXiv:2609.28372v1 Announce Type: cross Abstract: Consumers increasingly delegate purchasing decisions to Large Language Models (LLMs) acting as surrogate consumers. Using "Tool-Lab," an adaptation of information-board process tracing that places product attributes behind costly...

<https://arxiv.org/abs/2609.28372>

07 CONF

Retrieval Augmented (Knowledge Graph), and Large Language Model-Driven Design Structure Matrix (DSM) Generation of Cyber-Physical Systems

ArXiv CS.AI 50%

arXiv:2602.16715v2 Announce Type: replace Abstract: We explore the potential of Large Language Models (LLMs), Retrieval-Augmented Generation (RAG), and Graph-based RAG (GraphRAG) for generating Design Structure Matrices (DSMs). We test these methods on two distinct use cases--a...

<https://arxiv.org/abs/2602.16715>

08 CONF

TACT: Taxonomy-Aligned Post-Training for Pedagogically Adaptive English Tutoring

ArXiv CS.AI 50%

arXiv:2608.03952v2 Announce Type: replace Abstract: Large language models (LLMs) are increasingly used to provide conversational practice for English-as-a-second-language (ESL) learners. Effective ESL tutoring, however, requires more than fluent response generation: a tutor must...

<https://arxiv.org/abs/2608.03952>

09 CONF

HANIA: Planner-Guided Multimodal Graph Evidence Selection for Grounded Question Answering

ArXiv CS.AI 50%

arXiv:2608.29088v2 Announce Type: replace Abstract: Multimodal question answering remains sensitive to noisy, incomplete, and weakly grounded evidence. Long unstructured contexts can introduce redundancy and encourage unsupported generation, while flat retrieval may overlook...

<https://arxiv.org/abs/2608.29088>

10 CONF

TransBERT: A Framework for Synthetic Translation in Domain-Specific Language Modeling

ArXiv CS.AI 50%

arXiv:2609.26347v2 Announce Type: replace-cross Abstract: The scarcity of non-English language data in specialized domains significantly limits the development of effective Natural Language Processing (NLP) tools. We present TransBERT, a novel framework for pre-training language...

<https://arxiv.org/abs/2609.26347>

11 CONF

Offline A/B Testing of Slate Recommendation Systems with LLMs: Reducing the Dependency on Pre-Collected User Interaction Data

ArXiv CS.AI 50%

arXiv:2511.04541v2 Announce Type: replace-cross Abstract: Slate recommender systems (RecSys) present users with ordered sets of interacting items (e.g., playlists). We investigate whether large language models (LLMs) can articulate pairwise preferences between slates for...

<https://arxiv.org/abs/2511.04541>

12 CONF

Parameter Importance-Driven Continual Learning for Foundation Models

ArXiv CS.AI 50%

arXiv:2511.15375v2 Announce Type: replace-cross Abstract: Domain-specific post-training often causes catastrophic forgetting, making foundation models lose their general reasoning ability and limiting their adaptability to dynamic real-world environments. Preserving general...

<https://arxiv.org/abs/2511.15375>

13 CONF

Fine-Tune, Then Rectify

ArXiv CS.AI 50%

arXiv:2511.19486v3 Announce Type: replace-cross Abstract: Driven by recent advances in artificial intelligence, a growing literature has demonstrated the potential of using large language models (LLMs) as scalable surrogates to generate human-like responses. Two common...

<https://arxiv.org/abs/2511.19486>

14 CONF

TiAb Review Plugin: A Browser-Based Tool for AI-Assisted Study Selection in Systematic Reviews

ArXiv CS.AI 50%

arXiv:2604.08602v2 Announce Type: replace-cross Abstract: Server-based screening tools impose subscription costs, while open-source alternatives require coding skills, and full-text screening has remained outside the scope of no-code open-source tools. We developed TiAb Review...

<https://arxiv.org/abs/2604.08602>

15 CONF

Toward Measuring Structural Drift in LLM Communication Loops

ArXiv CS.AI 50%

arXiv:2604.13061v3 Announce Type: replace-cross Abstract: Large language models increasingly run in stateful pipelines that assemble each prompt from retrieval, memory, tools, and other agents. Such pipelines drift: information that should shape the next response is dropped,...

<https://arxiv.org/abs/2604.13061>

16 CONF

Safeguarding LLM Agents against Long-Horizon Threats via Shadow Memory

ArXiv CS.AI 50%

arXiv:2605.03228v2 Announce Type: replace-cross Abstract: As large language model (LLM)-powered agents are increasingly deployed to perform complex, real-world tasks, they face a growing class of attacks that exploit extended user-agent-environment interactions to pursue...

<https://arxiv.org/abs/2605.03228>

17 CONF

EA-WM: Event-Aware Generative World Model with Structured Kinematic-to-Visual Action Fields

ArXiv CS.AI 50%

arXiv:2605.06192v2 Announce Type: replace-cross Abstract: Pretrained video diffusion models provide powerful spatiotemporal generative priors, making them a natural foundation for robotic world models. While recent world-action models jointly optimize future videos and actions,...

<https://arxiv.org/abs/2605.06192>

18 CONF

DreamAvoid: Critical-Phase Test-Time Dreaming to Avoid Failures in VLA Policies

ArXiv CS.AI 50%

arXiv:2605.11750v2 Announce Type: replace-cross Abstract: Vision-Language-Action (VLA) models are often brittle in fine-grained manipulation, where minor action errors during the critical phases can rapidly escalate into irrecoverable failures. Since existing VLA models rely...

<https://arxiv.org/abs/2605.11750>

19 CONF

Routing-Aware Expert Calibration for Machine Unlearning in Mixture-of-Experts Language Models

ArXiv CS.AI 50%

arXiv:2606.10338v2 Announce Type: replace-cross Abstract: Machine unlearning is increasingly important for large language models, yet unlearning in Mixture-of-Experts (MoE) architectures remains underexplored. Unlike dense models, MoE architectures employ a router at each layer...

<https://arxiv.org/abs/2606.10338>

20 CONF

How a shared state is described determines whether AI agents synchronize

ArXiv CS.AI 50%

arXiv:2608.06968v2 Announce Type: replace-cross Abstract: Language-model agents increasingly act in populations, where the outcome that matters is collective: whether they align, split or fail to coordinate. Each acts not on the world but on a text description of it, a choice...

<https://arxiv.org/abs/2608.06968>

21 CONF

BIDETA: Brain-Inspired Data-Efficient Tactile Adaptation for Unseen Sensors

ArXiv CS.AI 50%

arXiv:2609.08673v2 Announce Type: replace-cross Abstract: Vision-based tactile sensors provide high-resolution contact information for robotic perception and contact-rich manipulation, advancing embodied intelligence through more reliable physical interaction. However,...

<https://arxiv.org/abs/2609.08673>

22 CONF

Same Outcome, Different Readout: What Does a Steerable Valence Direction in LLMs Represent?

ArXiv CS.AI 50%

arXiv:2609.22850v2 Announce Type: replace-cross Abstract: Decodability and successful activation steering do not, by themselves, establish what an internal direction represents. This gap is especially consequential for welfare-relevant interpretations, where a proposed...

<https://arxiv.org/abs/2609.22850>

23 CONF

SyzHarness: Patch-Based Kernel Bug Reproduction with LLM-Synthesized Fuzzing Harnesses

ArXiv CS.AI 50%

arXiv:2609.23889v2 Announce Type: replace-cross Abstract: Automated kernel vulnerability reproduction is essential for bug triage, patch validation, and regression testing, but still lacks an effective and efficient solution. The core challenge is twofold: a reproducer must...

<https://arxiv.org/abs/2609.23889>