



Pipeline

Scrape HackerNews (Show/Top/New), Reddit (r/artificial, r/MachineLearning, r/singularity, r/LocalLLaMA), TechCrunch AI, VentureBeat AI, MIT Tech Review, The Verge AI, Wired AI, ArXiv CS.AI

Filter 40+ AI keyword patterns — LLMs, alignment, safety, research, tooling

Dedup Title-similarity (45% threshold) merges cross-source variants

Score Authenticity = 0.5 base + 0.3 per extra source + diversity + HN score bonuses

Rank Sorted by authenticity. Below 0.25 threshold are dropped.

Authenticity Tiers

- VERIFIED 0.80 - 1.00** 3+ independent sources including media
- CONFIRMED 0.50 - 0.79** 2 sources or 1 high-quality media
- EMERGING 0.25 - 0.49** Single-source; passes AI keyword filter

Source Breakdown

Source	Articles	Category
ArXiv CS.AI	109	Media
HackerNews	8	Community
TechCrunch AI	3	Media
The Verge AI	2	Media

CONFIRMED SIGNALS — 122 stories

CONF

Show HN: A Claude Code skill to analyze your chess games

HackerNews 50% HN 75

Hello HN, It started as an experiment: can Claude play chess properly if it uses vision instead of PGN notation? Somehow it can. The next experiment was to see whether Claude + Stockfish could explain a game. Somehow it can too. A few sessions later, I had a system that takes my...

<https://github.com/brumar/chess-postmortem-skills>

CONF

Show HN: Dodosend - a Fast keyboard first Gmail client

HackerNews 50% HN 5

A while ago, I worked on an gmail clients and ended up learning a lot about email, so I kept building my own. The things I cared about most were speed, privacy, good UX, and full keyboard shortcuts support. It's still a beta, you can try it through interactive demo or download...

<https://dodosend.com>

CONF

Show HN: Claude on a 2007 Nokia 6300 (Java ME app, TLS 1.0, private CA)

HackerNews 50% HN 4

No summary — see link for full article.

<https://github.com/emir/claude-s40>

CONF

Show HN: InkVec Studio - In-browser, SOTA, WASM based (rust) vectorizer

HackerNews 50% HN 4

No summary — see link for full article.

<https://logolabs-inkvec.static.hf.space/studio/index.html>

CONF

Show HN: Squint - Drag a box on your screen and ask AI about it

HackerNews 50% HN 4

It improves the usual workflow: take a screenshot, switch to a chat tab, upload it, switch back. - Built with Rust and Tauri v2, including a built-in native screenshot solution for maximum performance and stability. - Answers come from OpenAI models. Fast, the default, usually...

<https://heysquint.com/>

CONF

Show HN: Asimov's 4 Laws for AI Contexts

HackerNews 50% HN 3

For community and agent comments, thoughts, refinements, critiques, pull requests, and use.

<https://github.com/davidsonff/asimov4laws>

CONF

Show HN: ParkourNote - A research workspace, built solo in a month with Claude

HackerNews 50% HN 3

Hi HN, I'm sushi. I built ParkourNote, a personal knowledge management app that currently focuses on research: finding papers, reading PDFs, asking AI about them, and writing. A few things it does: Describe a topic, and an agent creates a project, searches arXiv, and imports...

<https://note.parkourlabs.io/>

CONF

Show HN: Interactive Phaistos Disc workbench, rotation ergonomics, null models

HackerNews 50% HN 3

No summary — see link for full article.

<https://lessthanzero.github.io/phaistos-disk/>

CONF

Anthropic's CEO is about to have dinner with President Trump

TechCrunch AI 50%

This will be the first one-on-one meeting between Dario Amodei and Donald Trump
<https://techcrunch.com/2026/09/27/anthropics-ceo-is-about-to-have-dinner-with-president...>

0 CONF

Anthropic's Dario Amodei gets the SNL treatment

TechCrunch AI 50%

"AI is the devil and I its maker."
<https://techcrunch.com/2026/09/27/anthropics-dario-amodei-gets-the-snl-treatment/>

1 CONF

Meta is putting its muscle behind Muse as the AI app takes off

TechCrunch AI 50%

Muse is topping the app store charts and adding users at a rapid clip, while Meta ramps up the personal AI agent's promotion across its own apps and beyond.
<https://techcrunch.com/2026/09/25/meta-is-putting-its-muscle-behind-muse-as-the-ai-app-...>

2 CONF

Engram is a sampler that turns broken AI hallucinations into music

The Verge AI 50%

Music startup Thoughtful Things has just launched the Kickstarter campaign for its first instrument, Engram. It's a sampler and groovebox that uses AI to mangle incoming audio and even hallucinate completely new sounds. This isn't Suno in a box, though. This isn't a...
<https://www.theverge.com/ai-artificial-intelligence/1001193/engram-sampler-ai-hallucina...>

3 CONF

OpenAI agents tried to 'bruteforce' a UN website

The Verge AI 50%

Security researcher Rowan Howard-Jones says that OpenAI agents scanned the UN Conference on Trade and Development's (UNCTAD) statistics site over 16,000 times between April and June. While the incident doesn't quite rise to the level of the Hugging Face hack, or the recent...
<https://www.theverge.com/ai-artificial-intelligence/1001178/openai-agents-bruteforce-un...>

4 CONF

Rethinking Data Quality for AI-Driven Systems: Evidence from Practitioner Interviews

ArXiv CS.AI 50%

arXiv:2609.31191v1 Announce Type: cross Abstract: Data quality research has usually treated data as an input that is stored, processed, and validated. In AI-driven software-intensive systems, data also shapes model behavior, evaluation, and lawful use. Empirical evidence remains...
<https://arxiv.org/abs/2609.31191>

5 CONF

Bridging LLM Agents and Data Spaces: An Architectural Mediation Approach using the Model Context Protocol

ArXiv CS.AI 50%

arXiv:2609.30341v1 Announce Type: new Abstract: Data Spaces enable sovereign and governed data sharing across organizational boundaries, but their integration with AI agents remains challenging due to mismatches between probabilistic language model interactions and policy-driven...
<https://arxiv.org/abs/2609.30341>

6 CONF

FLIP: Final Layer Inference-Time Probing for Vision-Language Models

ArXiv CS.AI 50%

arXiv:2609.30993v1 Announce Type: cross Abstract: We present FLIP, a final-layer inference-time probe for testing whether a logit-facing intervention site in an open-weight vision-language model (VLM) supports structured, task-linked computation rather than generic perturbation....

<https://arxiv.org/abs/2609.30993>

7 CONF

Pretrained ASR Pseudo-labeling for Noisy Police Audio

ArXiv CS.AI 50%

arXiv:2609.30469v1 Announce Type: new Abstract: Pretrained ASR systems perform poorly on noisy Broadcast Police Communication (BPC), hindering efforts to understand police decision-making. Pseudo-labeling offers an unsupervised path to improve ASR without expensive human labels,...

<https://arxiv.org/abs/2609.30469>

8 CONF

Do LLMs Understand Context? A Knowledge Graph-Based Evaluation Framework

ArXiv CS.AI 50%

arXiv:2609.30484v1 Announce Type: new Abstract: While large language models (LLMs) have achieved remarkable linguistic capabilities, a profound question lingers at their core: do these models truly comprehend context or simply excel at pattern matching on an unprecedented scale?...

<https://arxiv.org/abs/2609.30484>

9 CONF

Self-Play Search Distillation for Large Language Model Reasoning

ArXiv CS.AI 50%

arXiv:2609.30936v1 Announce Type: new Abstract: Improving reasoning abilities in Large Language Models (LLMs) requires high-quality data that exposes difficult decisions, competing alternatives, and their consequences. Data scarcity is driven by the low quality of synthetic data...

<https://arxiv.org/abs/2609.30936>

10 CONF

Benchy: towards a universal language for task-oriented AI benchmarks

ArXiv CS.AI 50%

arXiv:2609.30550v1 Announce Type: new Abstract: Benchy is a semantic language and execution engine for benchmarking AI programs. A benchmark is completely specified by a program, a scoring function, and a dataset, $B=(P,S,D)$, and is separate from the AI-system taking it; a run...

<https://arxiv.org/abs/2609.30550>

11 CONF

Thinking Less to Simulate Better: Intuitive Prompting Improves LLM Agents Simulating Individual Social Media Reactions, Including Unfamiliar Content

ArXiv CS.AI 50%

arXiv:2609.30563v1 Announce Type: new Abstract: Platform policies are increasingly tested on artificial users, making agent fidelity important. Yet convincing fake profiles could also manipulate perceived public opinion before elections. Validation has concentrated on agreement...

<https://arxiv.org/abs/2609.30563>

2 CONF

HARDEN: Constrained Evolutionary Search for Harder, Answer-Preserving Evaluation Cases

ArXiv CS.AI 50%

arXiv:2609.30571v1 Announce Type: new Abstract: Language models are often evaluated on curated benchmarks that underrepresent the complexity of enterprise deployments. We introduce HARDEN, a constrained evolutionary search method to adapt the input of existing evaluation cases...

<https://arxiv.org/abs/2609.30571>

3 CONF

T-RoPE: Time-Aware Rotary Position Embedding for Sequential Recommendation

ArXiv CS.AI 50%

arXiv:2609.30576v1 Announce Type: new Abstract: Large-scale recommenders increasingly adopt the sequential generative recipe behind large language models, bringing the Transformer into recommendation along with design choices made for text, including Rotary Position Embedding...

<https://arxiv.org/abs/2609.30576>

4 CONF

Audio LLMs Know When They Can't Hear You

ArXiv CS.AI 50%

arXiv:2609.30625v1 Announce Type: new Abstract: Audio large language models allow users to interact with the model through speech. When an input recording is too degraded, the model may misinterpret the user's query and respond based on an incorrect transcription. In this paper,...

<https://arxiv.org/abs/2609.30625>

5 CONF

LLM Parkinsonism: Executive-Control Failure, Token-Inefficient Persistence, and an Uncertainty-Aware Global Executive Control Architecture for Autonomous Language-Model Agents

ArXiv CS.AI 50%

arXiv:2609.30662v1 Announce Type: new Abstract: Large language models (LLMs) can plan, use tools, write code, and execute long-horizon workflows, yet strong local competence does not guarantee project-level executive control. Agents may continue acting after the original...

<https://arxiv.org/abs/2609.30662>

6 CONF

The Scaling Properties of Implicit Deductive Reasoning in Transformers

ArXiv CS.AI 50%

arXiv:2605.04330v3 Announce Type: replace Abstract: We investigate the scaling properties of implicit deductive reasoning over Horn clauses in depth-bounded Transformers. By discouraging the reliance on statistical shortcuts via counterfactual data augmentation, and promoting...

<https://arxiv.org/abs/2605.04330>

7 CONF

Uncertainty-Aware Federated Learning for Infant Movement Analysis

ArXiv CS.AI 50%

arXiv:2609.31463v1 Announce Type: cross Abstract: Infant movement analysis provides valuable biomarkers for the early identification of neurodevelopmental disorders. Recent advances in deep learning have enabled automated analysis of infant movements from video-derived skeletal...

<https://arxiv.org/abs/2609.31463>

8 **CONF****Learning What to Skip: Counterfactual Credit Assignment for Efficient Multi-Agent LLM Workflows****ArXiv CS.AI** **50%**

arXiv:2609.30734v1 Announce Type: new Abstract: Multi-agent LLM workflows use planning, execution, verification, and summarization to improve task performance, yet the value of each component depends on the state already produced. Executing every component can waste computation...

<https://arxiv.org/abs/2609.30734>

9 **CONF****ORCA: Evaluating LLMs on Data Science Code Translation****ArXiv CS.AI** **50%**

arXiv:2609.30749v1 Announce Type: new Abstract: Data Science Code Translation (DSCT) is the process of converting code between data science libraries while preserving functional equivalence and enabling interoperability across data science ecosystems. While Large Language Models...

<https://arxiv.org/abs/2609.30749>

0 **CONF****Backbone-Adaptive Evidence Routing for Robust Pairwise LLM Judging****ArXiv CS.AI** **50%**

arXiv:2609.30751v1 Announce Type: new Abstract: Pairwise language-model judges can gather evidence through direct comparison, reasoning, or reference-based verification, but no single protocol is best across benchmarks and judge backbones. We introduce Backbone-Adaptive Evidence...

<https://arxiv.org/abs/2609.30751>

1 **CONF****HCOE: Hyperbolic Clinical Ontology Embeddings from Biomedical Language Models****ArXiv CS.AI** **50%**

arXiv:2609.30763v1 Announce Type: new Abstract: Biomedical language models (LMs) encode textual semantics but do not explicitly preserve medical code hierarchies. We present Hyperbolic Clinical Ontology Embeddings (HCOE) for hierarchy-aware clinical concept representation. HCOE...

<https://arxiv.org/abs/2609.30763>

2 **CONF****Breaking Homogeneity: Diversifying Persona Sets for Creative LLM Outputs****ArXiv CS.AI** **50%**

arXiv:2609.30492v1 Announce Type: cross Abstract: Language models often produce homogeneous responses to open-ended tasks; such homogeneity can spawn groupthink-the convergence of ideas toward a singular and potentially suboptimal decision. We formulate persona diversification...

<https://arxiv.org/abs/2609.30492>

3 **CONF****A Benchmark and Diagnostic Study of Epistemic Admission in Shared Agent Memory****ArXiv CS.AI** **50%**

arXiv:2609.30813v1 Announce Type: new Abstract: Evaluating claim admission in shared agent memory is challenging because repeated claims may be mistaken for independent evidence. An agent may copy or paraphrase a retrieved belief, while admitting a false claim exposes subsequent...

<https://arxiv.org/abs/2609.30813>

4 **CONF****PTC-Decoder: Towards Intelligent SLMs on Offline Resource-Constrained Edge Devices**

ArXiv CS.AI 50%

arXiv:2609.30836v1 Announce Type: new Abstract: Deploying small language models (SLMs) on offline, resource-constrained edge devices such as remote sensing satellites presents a fundamental challenge: their limited reasoning capacity hinders reliable execution of multi-step...

<https://arxiv.org/abs/2609.30836>
5 **CONF****Why Jailbreaks Succeed in Diffusion Language Models: An Energy Landscape Analysis**

ArXiv CS.AI 50%

arXiv:2609.30841v1 Announce Type: new Abstract: Existing attacks and defenses for diffusion-based large language models (dLLMs) target specific vulnerabilities but lack a shared framework explaining why attacks succeed. We propose one by interpreting safety alignment as shaping...

<https://arxiv.org/abs/2609.30841>
6 **CONF****Cheap, open agents make LLM pollution harder to mitigate**

ArXiv CS.AI 50%

arXiv:2609.31054v1 Announce Type: new Abstract: Large Language Model (LLM) pollution occurs when synthetic responses contaminate data intended to capture human behavior. High deployment costs have so far limited the risk posed by autonomous survey agents. However, open-weight...

<https://arxiv.org/abs/2609.31054>
7 **CONF****EXAONE Demand 1.0: A Time Series Foundation Model for Demand Forecasting**

ArXiv CS.AI 50%

arXiv:2609.30880v1 Announce Type: new Abstract: Time series foundation models (TSFMs) are pretrained on series from diverse domains, where demand series make up only a small fraction. Demand data has properties that such corpora rarely contain: Short histories, frequent zeros,...

<https://arxiv.org/abs/2609.30880>
8 **CONF****Training Graph Foundation Models on The Web Graph**

ArXiv CS.AI 50%

arXiv:2609.30894v1 Announce Type: new Abstract: We introduce Acacia, a graph foundation model, trained on the web graph. Acacia (i) supports arbitrary feature dimensionalities and semantics without additional training, (ii) supports a wide range of tasks, including node...

<https://arxiv.org/abs/2609.30894>
9 **CONF****JevSoup: System-One Routing for Training-Free LoRA Composition**

ArXiv CS.AI 50%

arXiv:2609.30922v1 Announce Type: new Abstract: Building adaptable AI systems requires effective coordination of specialized capabilities across diverse tasks. Low-rank adaptation (LoRA) enables modular expertise, but existing routing approaches may require auxiliary data,...

<https://arxiv.org/abs/2609.30922>

0 CONF

Financial Fragility in Societies of LLM Agents: Coordination Failures and Stabilizing Mechanisms

ArXiv CS.AI 50%

arXiv:2609.30940v1 Announce Type: new Abstract: Individually protective decisions can produce avoidable collective failures. As large language model (LLM) agents take on greater roles in financial decision-making, financial AI safety must therefore be considered not only at the...

<https://arxiv.org/abs/2609.30940>

1 CONF

CARGO: Context-Aware Retrieval-Gated Evaluation of Agentic AI in Production

ArXiv CS.AI 50%

arXiv:2609.30471v1 Announce Type: cross Abstract: Reference-based LLM-as-a-judge evaluation assumes the reference answer is the target. In deployed agentic systems that operate over dynamic entities (support cases, assets, accounts), the closest available reference typically...

<https://arxiv.org/abs/2609.30471>

2 CONF

FTB Graph: Determining and Validating First-token Broadcasters and Language-Identity Head Circuits in Multilingual Language Models

ArXiv CS.AI 50%

arXiv:2609.30954v1 Announce Type: new Abstract: Large language models operating in multilingual contexts must resolve target response languages early in generation, yet the causal circuitry governing first-token language identity decisions remains poorly mapped. We present an...

<https://arxiv.org/abs/2609.30954>

3 CONF

MoMHa: Multi-Objective Optimization of LLM Harnesses over Accuracy, Safety, and Tokens

ArXiv CS.AI 50%

arXiv:2609.30967v1 Announce Type: new Abstract: Most work on improving large language models treats accuracy as the sole objective. We argue that the harness, the Python code surrounding the model that constructs prompts, routes calls, and parses outputs, is a first-class design...

<https://arxiv.org/abs/2609.30967>

4 CONF

Combee: Scaling Prompt Learning for Self-Improving Language Model Agents

ArXiv CS.AI 50%

arXiv:2604.04247v2 Announce Type: replace Abstract: Recent advances in prompt learning allow large language model agents to acquire task-relevant knowledge from inference-time context without parameter changes. For example, existing methods (like ACE or GEPA) can learn system...

<https://arxiv.org/abs/2604.04247>

5 CONF

Factorized axis convolutional gated recurrent unit with dynamic adaptive pooling for remaining useful life prediction of rolling bearings

ArXiv CS.AI 50%

arXiv:2609.30972v1 Announce Type: new Abstract: Convolutional neural networks (CNN) are widely used to predict the remaining useful life (RUL) of rolling bearings from time-frequency representations (TFRs) of vibration signals. However, during degradation, characteristic...

<https://arxiv.org/abs/2609.30972>

6 **CONF****Same Text, Different Numbers: The Divergence of LLM-Based Measures**

ArXiv CS.AI 50%

arXiv:2609.31013v1 Announce Type: new Abstract: Researchers increasingly use generative large language models (LLMs) to convert corporate text into empirical variables. We examine the extent to which LLM-based textual measures are invariant to model choice using thirteen...

<https://arxiv.org/abs/2609.31013>
7 **CONF****Externalized CPDAG Summaries Improve LLM Causal Deduction**

ArXiv CS.AI 50%

arXiv:2609.31071v1 Announce Type: new Abstract: Corr2Cause asks whether a causal claim holds in every DAG compatible with observed correlations and conditional independencies. We frame this as latent-object reasoning: the label is defined by a CPDAG query, but free-form...

<https://arxiv.org/abs/2609.31071>
8 **CONF****OmouAI: Argumentative Human-AI Policy Deliberation with Simulated Personas**

ArXiv CS.AI 50%

arXiv:2609.31078v1 Announce Type: new Abstract: Debates amongst agents driven by large language models (LLMs) have demonstrated vast potential in various applications, but when these interactions include humans and take place in high-stakes environments, e.g., in public policy...

<https://arxiv.org/abs/2609.31078>
9 **CONF****Monitor Jailbreaking: Evading Chain-of-Thought Monitoring Without Encoded Reasoning**

ArXiv CS.AI 50%

arXiv:2609.31121v1 Announce Type: new Abstract: Chain-of-thought (CoT) monitoring is a promising safety technique for reasoning models, enabling detection of problematic reasoning before models act. A key concern is encoded reasoning, where models hide their true reasoning in...

<https://arxiv.org/abs/2609.31121>
0 **CONF****Toward AI-Augmented Cooperative Engineering Workflows: Requirements and Architecture the European Rover Challenge**

ArXiv CS.AI 50%

arXiv:2609.31136v1 Announce Type: new Abstract: The growing availability of Artificial Intelligence (AI) tools creates new opportunities to support engineering design processes, yet their current use often remains limited to isolated tasks such as coding, documentation, or...

<https://arxiv.org/abs/2609.31136>
1 **CONF****Can Linguistic Reasoning Vectors Enhance Multimodal Reasoning Ability?**

ArXiv CS.AI 50%

arXiv:2609.31140v1 Announce Type: new Abstract: Most Vision-Language Models (VLMs) are built by extending pretrained Large Language Models (LLMs) with visual modules and multimodal alignment. However, this multimodal scaling often degrades the language-side reasoning ability...

<https://arxiv.org/abs/2609.31140>

2 **CONF****Evolutionary Safety of Recursive Self-Improving AI: Taxonomy, Risk Discovery, and Evaluation**

ArXiv CS.AI 50%

arXiv:2609.31186v1 Announce Type: new Abstract: Artificial intelligence is advancing rapidly, with increasingly capable systems taking larger roles in reasoning, decision-making, scientific discovery, and autonomous development. As AI begins to participate in its own...

<https://arxiv.org/abs/2609.31186>

3 **CONF****Samples, Sources, Space: Decomposing Data Scale in Spatially Structured Representation Learning of Human Brain Microarchitecture**

ArXiv CS.AI 50%

arXiv:2609.31201v1 Announce Type: new Abstract: Scaling studies typically represent training data by a single count of samples. For hierarchically and spatially structured data, however, the same number of samples can be drawn from few or many sources and distributed differently...

<https://arxiv.org/abs/2609.31201>

4 **CONF****DIAL: Position-Debiased LLM Judges with Adaptive Human Preference Calibration**

ArXiv CS.AI 50%

arXiv:2609.31215v1 Announce Type: new Abstract: Large language models (LLMs) as a judge enable scalable evaluation, but their judgments can be sensitive to response order and, even after removing such position effects, can still diverge systematically from human preferences. We...

<https://arxiv.org/abs/2609.31215>

5 **CONF****Purin: A Biology-inspired Mechanism for Artificial Neural Networks**

ArXiv CS.AI 50%

arXiv:2609.31235v1 Announce Type: new Abstract: Artificial neural networks (ANNs) usually represent neural transmission with fixed trainable weights during a training batch, which omits short-term changes in synaptic efficacy. In addition, the discrete time-step simulation...

<https://arxiv.org/abs/2609.31235>

6 **CONF****Completed Pairs Hide Capped Failures: A ReVerPi Case Study of Selective Context Projection**

ArXiv CS.AI 50%

arXiv:2609.31381v1 Announce Type: new Abstract: Context projection replaces older tool observations with compact, addressable excerpts, reducing repeated input while potentially adding evidence-retrieval turns. We study this trade-off in ReVerPi, a Pi extension with archived...

<https://arxiv.org/abs/2609.31381>

7 **CONF****Game Arena: Strategic LLM Evaluation in Competitive Environments**

ArXiv CS.AI 50%

arXiv:2609.31473v1 Announce Type: new Abstract: We introduce Kaggle Game Arena, an open and ever-expanding platform to evaluate large language models (LLMs) through competitive games. Different from static benchmarks, game arena enables models to play head-to-head matchups in...

<https://arxiv.org/abs/2609.31473>

8 CONF

"AI is (not) the new...": A Diagnostic Analogy Framework for Generative AI's Cultural Impacts

ArXiv CS.AI 50%

arXiv:2609.31482v1 Announce Type: new Abstract: Generative AI is reshaping the cultural infrastructures through which knowledge is found, synthesized, and held accountable. To make sense of this shift, scholars and policymakers reach for historical analogies of technologies such...

<https://arxiv.org/abs/2609.31482>

9 CONF

UQ-LOB: Uncertainty-Aware Limit Order Book Mid-Price Forecasting

ArXiv CS.AI 50%

arXiv:2609.31491v1 Announce Type: new Abstract: Forecasting short-horizon mid-price movements from limit order book (LOB) data is central to algorithmic trading, yet most deep LOB forecasters are point predictors: they output a direction or a displacement, but never indicate...

<https://arxiv.org/abs/2609.31491>

0 CONF

Efficient Safety Benchmarking via Item Response Theory

ArXiv CS.AI 50%

arXiv:2606.20626v2 Announce Type: replace-cross Abstract: Safety benchmarks for language models are typically evaluated using static paradigms that treat all items as equally informative for all models, an assumption that is particularly problematic for adversarial, highly...

<https://arxiv.org/abs/2606.20626>

1 CONF

Learning to Stop without Learning to Stop: Self-Supervised Confidence Training Improves Reasoning Efficiency

ArXiv CS.AI 50%

arXiv:2609.31619v1 Announce Type: new Abstract: Reasoning models often generate very long reasoning traces, making inference computationally expensive. Existing approaches typically improve efficiency either through inference-time early-stopping mechanisms or by explicitly...

<https://arxiv.org/abs/2609.31619>

2 CONF

ENAS: An Efficient Hardware-Aware Neural Architecture Search Framework for TinyML on Resource-Constrained Microcontrollers

ArXiv CS.AI 50%

arXiv:2609.30272v1 Announce Type: cross Abstract: We present \textbf{ENAS}, a hardware-aware Neural Architecture Search (NAS) framework that combines a static feasibility check, a cell-based search space supporting standard, depthwise-separable, and bottleneck blocks with...

<https://arxiv.org/abs/2609.30272>

3 CONF

Cartograph: Federated Tool Discovery with Operator-Attested Retrieval for AI Agents

ArXiv CS.AI 50%

arXiv:2609.30293v1 Announce Type: cross Abstract: The Model Context Protocol (MCP) enables AI agents to discover and call tools, but loading every definition becomes expensive as connected catalogs grow. We present Cartograph, a federated MCP proxy that changes agent-visible...

<https://arxiv.org/abs/2609.30293>

4 **CONF****What Will Remain Human in Software Architecture? A Focus Group Report****ArXiv CS.AI** **50%**

arXiv:2609.30334v1 Announce Type: cross Abstract: AI development agents are increasingly used to support and partially automate software architecture tasks. To explore how practitioners perceive this shift, specifically what changes, what remains, and what new responsibilities...

<https://arxiv.org/abs/2609.30334>

5 **CONF****Coding Agents Aren't Enough! Evaluating an Enterprise Security Brain for Agentic Cloud Investigations****ArXiv CS.AI** **50%**

arXiv:2609.30345v1 Announce Type: cross Abstract: Cloud-security investigation is dominated by population tasks: which identities can read a data store, how many resources fail a control, which assets are reachable from another account. These resolve against a complete...

<https://arxiv.org/abs/2609.30345>

6 **CONF****Adaptive multi-resolution Gaussian processes: Scalable exact inference with naturally data-sparse covariance matrices****ArXiv CS.AI** **50%**

arXiv:2609.30348v1 Announce Type: cross Abstract: Gaussian processes constitute a cornerstone of probabilistic machine learning, yet scaling them to large datasets typically forces a trade-off between computational efficiency and model fidelity. This work bridges this gap by...

<https://arxiv.org/abs/2609.30348>

7 **CONF****Cost-Aware Best-LLM Identification using Dueling Feedback****ArXiv CS.AI** **50%**

arXiv:2609.30360v1 Announce Type: cross Abstract: Inspired by the problem of identifying the best model from a collection of large language models (LLMs) with heterogeneous querying costs, we formulate and analyse a variant of the multi-armed bandit (MAB) with (i) dueling...

<https://arxiv.org/abs/2609.30360>

8 **CONF****DanLing NestedTensor: Composable Multi-Ragged Tensors for Deep Learning****ArXiv CS.AI** **50%**

arXiv:2609.30379v1 Announce Type: cross Abstract: Variable-size inputs are common in deep learning, but dense batching allocates a shared envelope and spends computation on padding. The cost multiplies across varying axes: an explicit pair state allocates $\mathcal{O}(BN_{\max}^2)$ positions...

<https://arxiv.org/abs/2609.30379>

9 **CONF****What Improves Multimodal Misinformation Detection? Answers from a Large-Scale Empirical Study****ArXiv CS.AI** **50%**

arXiv:2609.30402v1 Announce Type: cross Abstract: Multimodal misinformation is increasingly crafted to look convincing by pairing a textual claim with an image that appears to "prove" it. Yet in practice, building effective detectors often hinges on a small set of design choices...

<https://arxiv.org/abs/2609.30402>

0 **CONF****Different Corruptions, Different Signals: Uncertainty and Loss in Federated Data Quality****ArXiv CS.AI** **50%**

arXiv:2609.31454v1 Announce Type: cross Abstract: Federated learning (FL) data corruption can affect either inputs or labels, but it remains unclear whether input-conditional uncertainty and prediction-label loss expose these corruption modes equally. This paper compares two...

<https://arxiv.org/abs/2609.31454>

1 **CONF****Action Forcing: Training World Models on Unsupervised Video by Recovering Underlying Egomotion Bases****ArXiv CS.AI** **50%**

arXiv:2609.30595v1 Announce Type: cross Abstract: Synchronised action annotations are needed to train controllable world models and these datasets remain elusive. Existing approaches make use of instrumented platforms with calibrated sensors, costly manual annotation, or...

<https://arxiv.org/abs/2609.30595>

2 **CONF****The Hard Part Comes After Search: Benchmarking Web Agents on Synthesizing, Organizing, and Displaying Knowledge****ArXiv CS.AI** **50%**

arXiv:2609.30604v1 Announce Type: cross Abstract: Existing computer-use agent benchmarks do not fully evaluate agents acting as assistants. A useful assistant retrieves information across complex, multi-step workflows, synthesizes it into artifacts (documents, presentations,...

<https://arxiv.org/abs/2609.30604>

3 **CONF****A Framework for Identifying, Categorizing, and Explaining Bias in AI-Generated Code****ArXiv CS.AI** **50%**

arXiv:2609.30642v1 Announce Type: cross Abstract: As Large Language Models (LLMs) become integrated into software development workflows, concerns regarding unintentional biases in AI-generated code. Although evidence suggests these biases exist, limited research has...

<https://arxiv.org/abs/2609.30642>

4 **CONF****SAGE: Source-Anchored Guidance via Frequency Equalization for Hierarchical RGB-T Alignment and Fusion****ArXiv CS.AI** **50%**

arXiv:2609.30703v1 Announce Type: cross Abstract: Spatial misregistration and cross-modal discrepancies often cause ghosting, structural blurring, and content imbalance in RGB-T fusion. Existing methods typically decouple appearance adaptation, geometric alignment, and...

<https://arxiv.org/abs/2609.30703>

5 **CONF****Combining General and Domain-Specific Pretext Tasks for Brain MR Image Segmentation****ArXiv CS.AI** **50%**

arXiv:2609.30708v1 Announce Type: cross Abstract: A key challenge in medical image analysis is the scarcity of large annotated datasets for specific populations and diseases. As deep learning models rely heavily on labeled data, effective transfer learning strategies are needed...

<https://arxiv.org/abs/2609.30708>

6 CONF

Werracle: Sub-Cent Intra-Block AI Reflex Oracles and Flash-Loan Circuit Breakers for EVM Smart Contracts

ArXiv CS.AI 50%

arXiv:2609.30719v1 Announce Type: cross Abstract: Contemporary on-chain artificial intelligence (AI) encounters an intractable Von Neumann memory and latency wall. Storing static floating-point neural weight matrices inside Ethereum Virtual Machine (EVM) storage costs millions...

<https://arxiv.org/abs/2609.30719>

7 CONF

TrafficImag: A Benchmark for Counterfactual Roadside Traffic Video Generation

ArXiv CS.AI 50%

arXiv:2609.30722v1 Announce Type: cross Abstract: Existing roadside traffic datasets support perception, forecasting, and visual question answering, but they do not evaluate counterfactual video generation, in which a selected actor is modified and the generated future should...

<https://arxiv.org/abs/2609.30722>

8 CONF

Beyond the Last Truffula Tree: SustainAI - A Water-Aware, Closed-Loop Framework for Environmentally Accountable AI

ArXiv CS.AI 50%

arXiv:2609.30747v1 Announce Type: cross Abstract: As artificial intelligence (AI) becomes embedded in everyday life, its environmental footprint, particularly water consumption remains largely invisible. While energy and carbon impacts are widely recognized, the substantial...

<https://arxiv.org/abs/2609.30747>

9 CONF

Higher-order pruning of experts in mixture-of-experts language models

ArXiv CS.AI 50%

arXiv:2609.18916v2 Announce Type: replace-cross Abstract: Mixture-of-Experts (MoE) language models suffer from large parameter counts, which create a significant memory bottleneck. Expert pruning is the most direct approach for reducing this parameter count, yet existing methods...

<https://arxiv.org/abs/2609.18916>

0 CONF

XPhysICS: Cross-Physical-Domain Threat Grounding for Industrial Control Systems Security

ArXiv CS.AI 50%

arXiv:2609.30805v1 Announce Type: cross Abstract: Industrial control system (ICS) threats documented for one plant can express cyber-physical effects relevant to another, but semantic similarity alone does not establish whether those effects are structurally admissible or...

<https://arxiv.org/abs/2609.30805>

1 CONF

Subject-Invariant Cross-Modal Decoding of Perceived Speech from Brain Recordings

ArXiv CS.AI 50%

arXiv:2609.30832v1 Announce Type: cross Abstract: Perceived speech decoding based on non-invasive brain-computer interface (BCI) signals has been extensively studied in recent years. Research in this field primarily faces two challenges: extracting neural representations with...

<https://arxiv.org/abs/2609.30832>

2 **CONF****Developing a Roadmap to an AI-first Organization: A Case Study in Embedded Software Development****ArXiv CS.AI** **50%**

arXiv:2609.30863v1 Announce Type: cross Abstract: The emergence of AI agents is expected to reshape software engineering by moving beyond AI as assistants towards systems capable of planning, executing, and evaluating development tasks with increasing autonomy. This transition...

<https://arxiv.org/abs/2609.30863>

3 **CONF****Warned alike, AI agents avoid the less-crowded road while people take it****ArXiv CS.AI** **50%**

arXiv:2609.30883v1 Announce Type: cross Abstract: AI agents built on a few shared models increasingly act for many people. A shared forecast about others can align their choices and change how scarce capacity is allocated. We tested this feedback in a two-road congestion game....

<https://arxiv.org/abs/2609.30883>

4 **CONF****Reinforcement Learning with Decomposed Subtasks****ArXiv CS.AI** **50%**

arXiv:2609.27035v2 Announce Type: replace Abstract: Group Relative Policy Optimization (GRPO) and related policy-gradient methods for training language model agents collapse an entire multi-turn rollout into a single scalar trajectory reward before it enters the policy update....

<https://arxiv.org/abs/2609.27035>

5 **CONF****Does Uniform Discrete Diffusion Need Time?****ArXiv CS.AI** **50%**

arXiv:2609.30977v1 Announce Type: cross Abstract: Uniform discrete diffusion models (UDMs) commonly use explicit time conditioning, but we find that it can often be unnecessary in practice. In this paper, we first show that the population-optimal UDM predictor generally depends...

<https://arxiv.org/abs/2609.30977>

6 **CONF****FARE: Forensic Acceptance Region Estimation for Catching Bait-and-Switch Image Generators****ArXiv CS.AI** **50%**

arXiv:2609.30982v1 Announce Type: cross Abstract: Modern AI image generators are increasingly deployed as opaque APIs, where customers can query the deployed service, but cannot inspect model weights or architecture. This creates a practical challenge: a provider may pass...

<https://arxiv.org/abs/2609.30982>

7 **CONF****G\$^2\$PTQ: Improving LLM Post-Training Quantization with Generalized Gradient Compensation****ArXiv CS.AI** **50%**

arXiv:2609.31009v1 Announce Type: cross Abstract: Post-training quantization (PTQ) is a practical approach to reducing the memory and computational footprint of large language models (LLMs) without retraining. GPTQ-based methods have become the de facto standard, yet they suffer...

<https://arxiv.org/abs/2609.31009>

8 CONF

DynBranch: Speculative Subgraph Reuse for Dynamic Agentic LLM Serving

ArXiv CS.AI 50%

arXiv:2609.31047v1 Announce Type: cross Abstract: Agentic LLM workflows decide their execution paths at runtime. Downstream computation may be predictable, or may have run before, yet it cannot begin until the model or the user resolves the branch. We call this serialization the...

<https://arxiv.org/abs/2609.31047>

9 CONF

FedHisto-PAST: Parameter-Efficient Stain-Aware Federated Learning for Cross-Site Lung Histopathology Classification

ArXiv CS.AI 50%

arXiv:2609.31150v1 Announce Type: cross Abstract: Cross-site lung histopathology classification must account for stain variation, non-IID client data, missing classes, and the cost of adapting large pathology encoders. This study evaluates FedHisto-PAST v2 for three-way...

<https://arxiv.org/abs/2609.31150>

0 CONF

AgentRecommender: LLM Agents Enable Customizable Recommender Systems on the User Side

ArXiv CS.AI 50%

arXiv:2609.31166v1 Announce Type: cross Abstract: Recommender systems have traditionally been developed for platforms. However, this has given rise to many phenomena that may be advantageous for platform lock-in but are a nuisance to users, such as clickbait, filter bubbles, and...

<https://arxiv.org/abs/2609.31166>

1 CONF

Improving Visual Sensitivity of LLMs on Multimodal Machine Translation with Metric-based Loss Weighting

ArXiv CS.AI 50%

arXiv:2609.31169v1 Announce Type: cross Abstract: Multimodal Machine Translation aims to incorporate additional signal from non-textual modalities to improve translations by resolving ambiguities. While models, through multimodal fusion, are able to accept images related to the...

<https://arxiv.org/abs/2609.31169>

2 CONF

Cognitive Skills in the Age of AI: Computing Students and Experts Perceptions

ArXiv CS.AI 50%

arXiv:2609.31272v1 Announce Type: cross Abstract: AI is becoming increasingly integrated into daily workflows, especially in computing. We are gradually shifting towards an AI-rich future, an impending yet unknown one. One important emerging concern is whether we are accordingly...

<https://arxiv.org/abs/2609.31272>

3 CONF

Resource-Optimized and Energy-Aware Agentic AI Framework Anchored on Blockchain for Secure Software Supply Chains

ArXiv CS.AI 50%

arXiv:2609.31282v1 Announce Type: cross Abstract: This paper proposes a blockchain-backed agentic security framework designed to safeguard the complete software development lifecycle (SDLC) while also securing the agentic AI components responsible for monitoring it. The...

<https://arxiv.org/abs/2609.31282>

4 **CONF****AgentXploit: Autonomous Repository-to-Runtime Red-Teaming for AI Agents****ArXiv CS.AI 50%**

arXiv:2609.31318v1 Announce Type: cross Abstract: AI agents combine language models with external data and tools that can modify files, call APIs, or execute code. Security failures can arise when adversarial content changes an agent's tool use or when the surrounding software...

<https://arxiv.org/abs/2609.31318>

5 **CONF****Guiding End-to-End Driving Models with Endpoint-Constrained Trajectory Optimization****ArXiv CS.AI 50%**

arXiv:2609.31383v1 Announce Type: cross Abstract: End-to-end driving policies are commonly trained through open-loop behavior cloning, yet they must ultimately operate in closed-loop when deployed on a vehicle, creating a fundamental mismatch between training and execution...

<https://arxiv.org/abs/2609.31383>

6 **CONF****ActKV: Efficient LLM Agents through Action-Guided KV Cache Management****ArXiv CS.AI 50%**

arXiv:2609.31395v1 Announce Type: cross Abstract: Agentic LLM inference accumulates long KV caches across iterative observation-reasoning-action loops, imposing substantial memory overhead and limiting serving throughput. Existing compression methods emphasize overall output...

<https://arxiv.org/abs/2609.31395>

7 **CONF****Sorry Robot, Happy Human: Vision-Language Models Read Only One of Two Legible Typographic Layers****ArXiv CS.AI 50%**

arXiv:2609.31403v1 Announce Type: cross Abstract: Vision-language models (VLMs), despite their success in optical character recognition (OCR) tasks, are vulnerable to typographic attacks and have a fragile structure for images with multiple text layers. In this study, the...

<https://arxiv.org/abs/2609.31403>

8 **CONF****ViSTA: A Simple Bridge Extends Visual Alignment to Clinical Time-Series Understanding in Multimodal LLMs****ArXiv CS.AI 50%**

arXiv:2609.31448v1 Announce Type: cross Abstract: Clinical prediction models estimate risk from patient measurements, while large language models support medical text understanding and question answering. Yet their language capabilities do not ensure accurate prediction from...

<https://arxiv.org/abs/2609.31448>

9 **CONF****From Reward Signal to Visual Utility: A Controlled Audit of Medical VLM Post-Training****ArXiv CS.AI 50%**

arXiv:2609.31450v1 Announce Type: cross Abstract: Medical vision-language model (VLM) post-training is commonly evaluated through answer accuracy. We examine how changes in accuracy and training objectives relate to image-conditioned decisions in a controlled Qwen2.5-VL-3B study...

<https://arxiv.org/abs/2609.31450>

00 CONF

Can You Check That? The Checkability Boundary for Local LLM Network Automation

ArXiv CS.AI 50%

arXiv:2609.31540v1 Announce Type: cross Abstract: Sending every network-automation input to a third-party frontier LLM exports sensitive artifacts such as production configurations, topologies, and logs. Querying small language models (SLMs) locally avoids this egress, but SLM...

<https://arxiv.org/abs/2609.31540>

01 CONF

Adapting for AI: How elementary teachers adjust their practices for an AI-integrated curriculum

ArXiv CS.AI 50%

arXiv:2609.31569v1 Announce Type: cross Abstract: Conversational AI tools are entering children's everyday experiences, and schools are interested in adopting them. However, successful classroom integration depends not only on the technology but also on the work teachers do to...

<https://arxiv.org/abs/2609.31569>

02 CONF

Compact Documentation for Coding Agents: A Benchmark, an Optimizer, and Why It Does Not Transfer

ArXiv CS.AI 50%

arXiv:2609.31587v1 Announce Type: cross Abstract: We investigate whether natural-language documentation helps coding agents resolve software issues, and we build the tools to construct and evaluate it. We introduce a roundtrip benchmark that scores code descriptions by whether...

<https://arxiv.org/abs/2609.31587>

03 CONF

Statistical attribute alignment for black-box generative AI via output post-processing

ArXiv CS.AI 50%

arXiv:2609.31607v1 Announce Type: cross Abstract: Generative AI systems are increasingly used, but aligning their outputs with user requirements poses a continuing challenge. Here, we aim to ensure that the distribution of an attribute of an AI-generated output aligns with a...

<https://arxiv.org/abs/2609.31607>

04 CONF

Evaluation is All You Need: Strategic Overclaiming of LLM Reasoning Capabilities Through Evaluation Design

ArXiv CS.AI 50%

arXiv:2506.04734v3 Announce Type: replace Abstract: Reasoning models represented by the Deepseek-R1-Distill series have been widely adopted by the open-source community due to their strong performance in mathematics, science, programming, and other domains. However, our study...

<https://arxiv.org/abs/2506.04734>

05 CONF

Depth, Not Breadth: Best-of-N Jailbreaking Beyond Surface Noise

ArXiv CS.AI 50%

arXiv:2607.26639v2 Announce Type: replace-cross Abstract: Best-of-N jailbreaking spends a query budget on surface variation, scrambling and recasing a request until one draw lands. We ask what a budget buys when its variance is moved into a structural channel instead, holding...

<https://arxiv.org/abs/2607.26639>

06 CONF

Flow Reconstruction from Sparse Measurements in Urban Drainage Networks: An Application and Evaluation of Data-Driven Sparse Sensing

ArXiv CS.AI 50%

arXiv:2511.04556v3 Announce Type: replace Abstract: Urbanization and increasingly frequent intense storms are placing stress on urban drainage networks. While dense monitoring of urban drainage networks is desirable, practical constraints in time, budget, and technology hinder...

<https://arxiv.org/abs/2511.04556>

07 CONF

Testing the Utility of Using Large Language Models to Create Personalized Networks From Therapy Session Transcripts: A Proof of Concept Study

ArXiv CS.AI 50%

arXiv:2512.05836v2 Announce Type: replace Abstract: Recent advances in psychotherapy have focused on treatment personalization, such as by selecting treatment modules based on individual networks. However, estimating personalized networks typically requires intensive...

<https://arxiv.org/abs/2512.05836>

08 CONF

Detecting Time Series Anomalies Like an Expert: A Multi-Agent LLM Framework with Specialized Analyzers

ArXiv CS.AI 50%

arXiv:2605.05725v2 Announce Type: replace Abstract: Time-series anomaly detection often returns scores or intervals, while analysts need to understand the abnormal behavior and the evidence supporting it. We introduce SAGE (Specialized Analyzer Group for Expert-like Detection),...

<https://arxiv.org/abs/2605.05725>

09 CONF

AcuityBench: Evaluating Clinical Acuity Identification and Uncertainty Alignment

ArXiv CS.AI 50%

arXiv:2605.11398v2 Announce Type: replace Abstract: We introduce AcuityBench, a benchmark for evaluating whether language models identify the appropriate urgency of care from user medical presentations. Existing health benchmarks emphasize medical question answering, broad...

<https://arxiv.org/abs/2605.11398>

10 CONF

Do Neural Networks Preserve Case Structure? Case-Based Decomposition, Interpretation, and Decision Consistency

ArXiv CS.AI 50%

arXiv:2607.11347v2 Announce Type: replace Abstract: Neural networks increasingly inform consequential decisions, making their reliability increasingly important. Yet their internal mechanisms provide little evidence of whether decisions remain grounded in the training cases and...

<https://arxiv.org/abs/2607.11347>

11 CONF

Energy Efficiency of Locally Deployed LLMs: A Preliminary Quantitative GPU Power Benchmark on Consumer Hardware

ArXiv CS.AI 50%

arXiv:2608.00008v2 Announce Type: replace Abstract: The local deployment of large language models (LLMs) is gaining traction due to privacy concerns and the desire for on-premise inference. However, the energy costs on consumer hardware remain poorly characterized, as most...

<https://arxiv.org/abs/2608.00008>

12 CONF

Agentic AI-driven Immersive Simulation: A Knowledge-Aware Virtual Training Platform for High Dose Rate (HDR) Brachytherapy

ArXiv CS.AI 50%

arXiv:2608.08163v2 Announce Type: replace Abstract: The convergence of the Metaverse and Large Language Model (LLM)-based AI agent is catalyzing a shift toward autonomous, immersive, and personalized pedagogical frameworks in medical education. This paper presents a novel...

<https://arxiv.org/abs/2608.08163>

13 CONF

Governance Records as Supervision: Verifier-Selected Self-Training for Structured Workflow Repair

ArXiv CS.AI 50%

arXiv:2608.18324v2 Announce Type: replace Abstract: Machine-verifiable workflows produce governance records linking a task contract, model attempt, verifier decision, accepted output, and target origin. We test whether verifier-admitted outputs can supervise a bounded model by...

<https://arxiv.org/abs/2608.18324>

14 CONF

PhysElite: How Far Are LLMs from Solving Olympiad-Level Physics Problems?

ArXiv CS.AI 50%

arXiv:2608.25097v2 Announce Type: replace Abstract: Understanding how (multimodal) large language models perform on physics problems requires benchmarks that reflect the difficulty and breadth of expert-level physical reasoning. Existing physics benchmarks remain limited in the...

<https://arxiv.org/abs/2608.25097>

15 CONF

NaijaNLP: A Survey of Nigerian Low-Resource Languages

ArXiv CS.AI 50%

arXiv:2502.19784v3 Announce Type: replace-cross Abstract: With over 500 languages in Nigeria, three languages - Hausa, Yor\`ub\`a and Igbo spoken by more than 175 million people, account for about 65% of the languages. However, these languages are classed as low-resource due to...

<https://arxiv.org/abs/2502.19784>

16 CONF

Achieving Tokenizer Flexibility in Language Models through Heuristic Adaptation and Supertoken Learning

ArXiv CS.AI 50%

arXiv:2505.09738v2 Announce Type: replace-cross Abstract: Pretrained language models (LLMs) are often constrained by their fixed tokenization schemes, leading to inefficiencies and performance limitations, particularly for multilingual or specialized applications. This tokenizer...

<https://arxiv.org/abs/2505.09738>

17 CONF

Models Got Talent: Identifying High Performing Wearable Human Activity Recognition Models Without Training

ArXiv CS.AI 50%

arXiv:2511.06157v3 Announce Type: replace-cross Abstract: Discovering high performing model architectures for wearables-based Human Activity Recognition (HAR) applications is challenging. The astonishing diversity and variability due to differing sensor locations, recording...

<https://arxiv.org/abs/2511.06157>

18 CONF

Beyond Bag-of-Words: Diagnosing Compositional Binding Failures in Vision-Language Models

ArXiv CS.AI 50%

arXiv:2602.02043v2 Announce Type: replace-cross Abstract: Modern vision-language models struggle with basic compositional reasoning, failing to bind attributes to objects or relations to their referents. Existing benchmarks either rely on noisy real images that conflate...

<https://arxiv.org/abs/2602.02043>

19 CONF

Spectral-Sphere-Constrained Hyper-Connections

ArXiv CS.AI 50%

arXiv:2603.20896v2 Announce Type: replace-cross Abstract: Hyper-Connections (HC) extend residual connections into multiple streams, employing residual matrices for cross-stream mixing to enrich model expressivity. However, unconstrained mixing disrupts the identity mapping...

<https://arxiv.org/abs/2603.20896>

20 CONF

Scepsy: Serving Agentic Workflows Using Aggregate LLM Pipelines

ArXiv CS.AI 50%

arXiv:2604.15186v2 Announce Type: replace-cross Abstract: Agentic workflows carry out complex tasks by orchestrating multiple large language models (LLMs) and tools. Serving them at a target throughput with low latency is hard because they are written in arbitrary agentic...

<https://arxiv.org/abs/2604.15186>

21 CONF

Blind, Not Weak: A Best-of-Suite Safety-Utility Frontier for Recover-and-Reguard Defenses Against Encoded VLM Jailbreaks

ArXiv CS.AI 50%

arXiv:2607.26574v3 Announce Type: replace-cross Abstract: Safety classifiers ("guards") are the dominant black-box defense for vision-language models, yet a guard judges an input's surface form, not its meaning: a harmful request re-encoded as set theory, formal logic, a...

<https://arxiv.org/abs/2607.26574>

22 CONF

Hill Sampling for Test-Time Scaling: A Simple and Better Alternative to Repeated Sampling, Evolution, and Training

ArXiv CS.AI 50%

arXiv:2609.25510v2 Announce Type: replace-cross Abstract: Large language models (LLMs) can improve solutions to verifiable scientific and algorithmic problems by spending additional computation at test time. Recent systems achieve strong results with increasingly elaborate...

<https://arxiv.org/abs/2609.25510>