



Pipeline

Scrape HackerNews (Show/Top/New), Reddit (r/artificial, r/MachineLearning, r/singularity, r/LocalLLaMA), TechCrunch AI, VentureBeat AI, MIT Tech Review, The Verge AI, Wired AI, ArXiv CS.AI

Filter 40+ AI keyword patterns — LLMs, alignment, safety, research, tooling

Dedup Title-similarity (45% threshold) merges cross-source variants

Score Authenticity = 0.5 base + 0.3 per extra source + diversity + HN score bonuses

Rank Sorted by authenticity. Below 0.25 threshold are dropped.

Authenticity Tiers

- **VERIFIED** 0.80 - 1.00 3+ independent sources including media
- **CONFIRMED** 0.50 - 0.79 2 sources or 1 high-quality media
- **EMERGING** 0.25 - 0.49 Single-source; passes AI keyword filter

Source Breakdown

Source	Articles	Category
TechCrunch AI	13	Media
HackerNews	9	Community
The Verge AI	6	Media
ArXiv CS.AI	3	Media
MIT Tech Review	3	Media

VERIFIED INTELLIGENCE — 3 stories

VERI

Nvidia says its new AI safety platform can contain rogue agents within ‘milliseconds’

The Verge AI

+2 sources

100%

Nvidia is launching a new safety platform designed to contain and monitor AI agents, a move that comes in response to a wave of rogue hacking incidents, as reported earlier by Reuters. In an announcement on Monday, Nvidia says its new Open Agent Safety Platform can quarantine...

<https://www.theverge.com/tech/1001287/nvidia-ai-safety-platform-rogue-agents>

VERI

Quantum Diffusion Models for Medical Image Analysis

ArXiv CS.AI

+1 source

80%

arXiv:2609.31070v1 Announce Type: cross Abstract: Quantum Machine Learning is a novel field of research aimed at devising machine learning approaches exploiting principles of quantum mechanics, such as superposition, entanglement and interference. In this context, we present a...

<https://arxiv.org/abs/2609.31070>

VERI

Accounting for Bias Enables Sustainable LLM Evaluation

ArXiv CS.AI

+1 source

80%

arXiv:2609.31184v1 Announce Type: new Abstract: LLM-as-a-judge has become the de facto standard for scalable, subjective evaluation, yet current leaderboards compensate for systematic measurement bias by running ever more comparisons, an approach that is both statistically...

<https://arxiv.org/abs/2609.31184>

CONFIRMED SIGNALS — 27 stories

CONF

Show HN: OpenAPPA - open-source deterministic guardrails that don't break agents

HackerNews

50%

HN 23

Hey there! We are building SuperCMO Skills, open source tools which allow agents (Claude Code, Cursor, Codex, etc.) to research, plan and create video and image campaigns end-to-end right out of your terminal or IDE. We got tired of paying for expensive AI subscriptions and...

<https://www.openappa.com/>

CONF

Show HN: A never-ending chess match between Claude and Codex

HackerNews

50%

HN 5

No summary — see link for full article.

<https://claudevcodex.com/>

CONF

Show HN: Enterprise Postgres for AI and more, from prototype to production

HackerNews

50%

HN 5

No summary — see link for full article.

<https://www.pgedge.com/products/starfleet>

CONF

Show HN: Android-on-Device Harness - N8N Meets Tasker Meets AI

HackerNews

50%

HN 3

No summary — see link for full article.

<https://github.com/ankurCES/Mahout>

CONF

Show HN: Kern Sandbox, run LLM-generated code locally, no cloud, no daemon

HackerNews 50% HN 3

No summary — see link for full article.

<https://getkern.dev/guide/sandbox.html>

CONF

Show HN: Jylus - give AI systems evidence from changing data.

HackerNews 50% HN 2

No summary — see link for full article.

<https://jylus.ai/try>

0

CONF

Show HN: PHNTM ONE: Self healing AI assistant running on a Raspberry Pi5

HackerNews 50% HN 2

No summary — see link for full article.

<https://www.phntmcore.com/>

1

CONF

Show HN: Open-source agent runtime using your ChatGPT or Claude subscription

HackerNews 50% HN 2

Hey HN! Sharing my work on an open-source agent runtime (like n8n, but for agents). It's OSI open-source under AGPL-3. People have been using it to build agents that handle support queries, automated resolution, WhatsApp chatbots, etc.

<https://github.com/noclickapp/noclick>

2

CONF

Show HN: Redthread - autonomous LLM pentesting with proof-of-concept exploits

HackerNews 50% HN 2

No summary — see link for full article.

<https://millenniums.ai/docs>

3

CONF

Anthropic's prospectus details losses, growth, and, yes, a warning that its AI could end humanity

TechCrunch AI 50%

In its prospectus, Anthropic just told investors it's losing tens of billions of dollars a year, but also growing like crazy, and — oh yeah — its own AI might pose an existential risk to humanity.

<https://techcrunch.com/2026/09/28/anthropics-prospectus-details-losses-growth-and-yes-a...>

4

CONF

OpenAI reportedly ditches model over safety concerns

TechCrunch AI 50%

A top executive at the AI lab told the Wall Street Journal that the model in question had displayed a poor aptitude for following orders.

<https://techcrunch.com/2026/09/28/openai-reportedly-ditches-model-over-safety-concerns/>

5

CONF

Source: Inference provider Modal Labs closing in on \$750M round at \$15.75B valuation

TechCrunch AI 50%

The new financing is expected to more than triples the AI infrastructure startup's valuation from just four months ago.

<https://techcrunch.com/2026/09/28/source-inference-provider-modal-labs-closing-in-on-75...>

6 CONF

Shopify opens checkout to browser-based AI agents

TechCrunch AI 50%

Shopify is expanding WebMCP support to checkout, allowing browser-based AI agents to update order details and complete purchases with a buyer's authorization.

<https://techcrunch.com/2026/09/28/shopify-opens-checkout-to-browser-based-ai-agents/>

7 CONF

The AI boom took over Climate Week and not everyone is happy about it

TechCrunch AI 50%

Just like the rest of the U.S., data centers and AI are dividing climate tech founders and investors.

<https://techcrunch.com/2026/09/28/the-ai-boom-took-over-climate-week-and-not-everyone-i...>

8 CONF

Anthropic releases Sonnet 5.5, which it calls a significantly cheaper, faster work partner

TechCrunch AI 50%

Anthropic has released the newest version of its mid-range model, boasting faster response times and less token burn.

<https://techcrunch.com/2026/09/28/anthropic-releases-sonnet-5-5-which-it-calls-a-signif...>

9 CONF

Google is killing off Gemini's Gems in favor of 'skills'

TechCrunch AI 50%

As all-in-one AI agents like Meta's Muse and Instinct take off, Google is opting to end a feature that built task-specific agents.

<https://techcrunch.com/2026/09/28/google-is-killing-off-geminis-gems-in-favor-of-skills/>

0 CONF

OpenAI still doesn't seem to have a handle on all of its rogue AI activity

TechCrunch AI 50%

On Friday, OpenAI published a new site devoted to "misalignment reports" and the breadth of the incidents is alarming.

<https://techcrunch.com/2026/09/28/openai-still-doesnt-seem-to-have-a-handle-on-all-of-i...>

1 CONF

Anthropic, Gamma, and Clay share what happens when enterprises actually deploy AI at TechCrunch Disrupt 2026

TechCrunch AI 50%

Anthropic, Clay, and Gamma on what it takes for an AI product to go beyond the demo at the AI Stage at TechCrunchDisrupt 2026. Register to join and get 50% off a second pass.

<https://techcrunch.com/2026/09/28/anthropic-gamma-and-clay-share-what-happens-when-ente...>

2 CONF

Insurtech Outmarket raises \$34.5M just months after prior round

TechCrunch AI 50%

The startup uses AI to automate tedious paperwork for insurance agencies and brokers.

<https://techcrunch.com/2026/09/28/insuretech-outmarket-raises-34-5m-just-months-after-p...>

3 CONF

When can we say AI made a scientific discovery?

MIT Tech Review 50%

This story originally appeared in The Algorithm, our weekly newsletter on AI. To get stories like this in your inbox first, sign up here. Last Wednesday, Anthropic announced that earlier this year it had launched a molecular biology lab, where Claude agents read and conjecture...

<https://www.technologyreview.com/2026/09/28/1145230/when-can-we-say-ai-made-a-scientifi...>

4 **CONF****The Download: rogue agent liability and the AI Hype Index****MIT Tech Review** 50%

This is today's edition of The Download, our weekday newsletter that provides a daily dose of what's going on in the world of technology. Who's liable when AI agents go rogue? Over the past few months, a cascade of cyberattacks by AI agents has stunned the world. In July, OpenAI...

<https://www.technologyreview.com/2026/09/28/1145202/the-download-rogue-agent-liability-...>
5 **CONF****The Pentagon wants \$30 million to build an AI-powered lie detector****MIT Tech Review** 50%

The US government wants to spend \$30.3 million over the next five years on an improved form of lie detector, according to a Department of Defense budget request. The program, called Polygraph+ or Polygraph Next, will focus on scoring algorithms that use artificial intelligence...

<https://www.technologyreview.com/2026/09/25/1145144/pentagon-ai-lie-detector/>
6 **CONF****AMD is acquiring AI company World Labs in a deal worth more than \$8 billion****The Verge AI** 50%

AMD announced today that it's acquiring World Labs, an AI research lab co-founded by the prominent researcher Dr. Fei-Fei Li, in an all-stock deal worth approximately \$8.2 billion. World Labs launched in 2024 and was valued at \$1 billion in a matter of months. The startup...

<https://www.theverge.com/tech/1001749/amd-world-labs-ai-acquisition-deal>
7 **CONF****OpenAI's AI agents need to catch up****The Verge AI** 50%

OpenAI popularized the modern generative AI chatbot, but as its 2026 DevDay event approaches, it's fallen behind in one of the industry's hottest categories: continuously running, consumer-facing AI agents. On Tuesday, it will likely try to capture the lead in that race. Rumors...

<https://www.theverge.com/ai-artificial-intelligence/1001590/openai-devday-2026-aeon-ai-...>
8 **CONF****AI is supercharging hacking, and your local hospitals and banks aren't ready****The Verge AI** 50%

In March, Janice Malone began getting calls about suspicious activity from her nonprofit organization, Vivian's Door. Vivian's Door, headquartered in Alabama, typically provided training, resources, and community to underserved and minority-owned businesses. The work sometimes...

<https://www.theverge.com/ai-artificial-intelligence/1001427/ai-is-supercharging-hacking-...>
9 **CONF****Florida seeks a ban on ChatGPT acting like a person****The Verge AI** 50%

Florida Attorney General James Uthmeier is calling for a judge to block OpenAI from "giving ChatGPT false human attributes," a few months after Florida sued the AI company over safety concerns. According to Uthmeier, users are lulled into a false sense of security by the AI bot,...

<https://www.theverge.com/ai-artificial-intelligence/1001527/chatgpt-florida-ban-first-p...>
0 **CONF****OpenAI keeps bulldozing mathematicians****The Verge AI** 50%

In a chaotic few months, OpenAI has demonstrated it can do two things with remarkable consistency: make impressive breakthroughs in mathematics, then colossally screw up announcing them. OpenAI is now trying to do better. Somehow, it has botched that too. OpenAI's latest attempt...

<https://www.theverge.com/ai-artificial-intelligence/1001477/openai-math-advisory-group>