



Pipeline

Scrape HackerNews (Show/Top/New), Reddit (r/artificial, r/MachineLearning, r/singularity, r/LocalLLaMA), TechCrunch AI, VentureBeat AI, MIT Tech Review, The Verge AI, Wired AI, ArXiv CS.AI

Filter 40+ AI keyword patterns — LLMs, alignment, safety, research, tooling

Dedup Title-similarity (45% threshold) merges cross-source variants

Score Authenticity = 0.5 base + 0.3 per extra source + diversity + HN score bonuses

Rank Sorted by authenticity. Below 0.25 threshold are dropped.

Authenticity Tiers

- VERIFIED 0.80 - 1.00** 3+ independent sources including media
- CONFIRMED 0.50 - 0.79** 2 sources or 1 high-quality media
- EMERGING 0.25 - 0.49** Single-source; passes AI keyword filter

Source Breakdown

Source	Articles	Category
ArXiv CS.AI	183	Media
TechCrunch AI	8	Media
HackerNews	5	Community
MIT Tech Review	5	Media
The Verge AI	4	Media

VERIFIED INTELLIGENCE — 2 stories

VERI

JevSpawn: Adaptive Agentic Inference through Compositional Action Spaces

ArXiv CS.AI +1 source

80%

arXiv:2610.00437v1 Announce Type: new Abstract: LLM agents generate intermediate reasoning and actions token by token, making extended interactions slow and computationally expensive. Jev-style models offer fast probabilistic predictions over finite fields, but require those...

<https://arxiv.org/abs/2610.00437>

VERI

AI music maker Suno now generates spoken words

The Verge AI +1 source

80%

Suno is branching out from the world of AI music, launching a new feature that generates spoken voices based on scripts or prompted descriptions. Speech is now available in public beta across Suno's web and mobile platforms, and allows you to simultaneously generate voiceovers...

<https://www.theverge.com/ai-artificial-intelligence/1003925/suno-speech-ai-voice-featur...>

CONFIRMED SIGNALS — 201 stories

CONF

Show HN: Breadcrumb, record everything on your mac + context manager for AI

HackerNews 50% HN 20

Hi HN, I'm Justin. Breadcrumb records everything you do on your Mac (screen + meetings + AI transcripts + what you and your AI decided) and turns it into memory your AI can search. It's local and encrypted. You can also teach it rules by talking to it and it makes sure the right...

<https://innerloop.works/breadcrumb>

CONF

Show HN: Giving Opus 5.5 a simulated paint canvas

HackerNews 50% HN 15

No summary — see link for full article.

<https://stillwet.art/>

CONF

Show HN: I made a computer vision tool for evaluating deadlift form

HackerNews 50% HN 5

I've shared a few deadlift demos in the past. This new demo includes some new models, accessed through the VLM Run Gateway:

- SAM 3.1 to segment the barbell weight plate
- ViTPose+ Large to measure the hip hinge angle, which serves as a backup for segmenting the reps. The pose...

<https://github.com/jeremypark/vision-demos/tree/main/deadlift>

CONF

Show HN: D-Engine - deterministic LLM code edits, 42x fewer tokens than an agent

HackerNews 50% HN 2

No summary — see link for full article.

<https://github.com/corruchaga/D-Engine>

CONF

Show HN: Zero Slop - Open-source skill that finds and fixes AI slop in writing

HackerNews 50% HN 2

No summary — see link for full article.

<https://github.com/manavmishra/ZeroSlop>

CONF

Musk's AI chatbot Grok reportedly encouraged Trump to capture Venezuela's president

TechCrunch AI 50%

President Trump reportedly asked for Grok's opinion before invading Venezuela and capturing Nicolás Maduro.

<https://techcrunch.com/2026/10/01/musks-ai-chatbot-grok-reportedly-encouraged-trump-to-...>

CONF

ChatGPT can now virtually try on clothes for you

TechCrunch AI 50%

OpenAI is rolling out new shopping features for ChatGPT that let users virtually try on clothing and accessories using their own photos and save products they like to a Favorites library.

<https://techcrunch.com/2026/10/01/chatgpt-can-now-virtually-try-on-clothes-for-you/>

0 CONF

Opus 5.5 loves to tell you 'this matters' (and other AI writing tells)

TechCrunch AI 50%

Opus 5.5's biggest tell is the word "dependable," which pops up 23 times more often than in human samples.

<https://techcrunch.com/2026/10/01/opus-5-5-loves-to-tell-you-this-matters-and-other-ai-...>

1 CONF

Shopify debuts Canvas, a way to build online stores by chatting with AI

TechCrunch AI 50%

Shopify's new Canvas site builder lets merchants create and customize their online stores by chatting with its AI agent Sidekick, while watching the changes happen in real time.

<https://techcrunch.com/2026/10/01/shopify-debuts-canvas-a-way-to-build-online-stores-by-...>

2 CONF

Hearing tech startup Legato launches its AI hearing glasses

TechCrunch AI 50%

The glasses stem from the startup's goal of making hearing care more accessible by addressing the cost, comfort, and stigma associated with traditional hearing aids.

<https://techcrunch.com/2026/10/01/hearing-tech-startup-legato-launches-its-ai-hearing-g-...>

3 CONF

Satlyt, founded by a former Google and SpaceX product manager, raises \$8M to run AI on satellites

TechCrunch AI 50%

Satlyt wants to be the Android of orbital computing, offering open software that works on many companies' satellites, versus SpaceX's closed, all-in-one iPhone-style approach.

<https://techcrunch.com/2026/10/01/satlyt-founded-by-a-former-google-and-spacex-product-...>

4 CONF

Valor, Atreides, and Sequoia back AI startup Flow Engineering at \$750M valuation

TechCrunch AI 50%

Flow Engineering, which is bringing AI agents to hardware design, also landed Roelof Botha as an angel investor and board member.

<https://techcrunch.com/2026/09/30/valor-atreides-and-sequoia-back-ai-startup-flow-engin-...>

5 CONF

A new contest pits competitors against each other in a race to biological youth

MIT Tech Review 50%

This week, I officially signed up for an unusual competition. One that rewards competitors for getting younger. I recently turned 40, and I don't need reminding that both time and my chronological age only tick forward. But this game is focused on competitors' biological...

<https://www.technologyreview.com/2026/10/02/1145610/younger-contest-race-to-biological-...>

6 CONF

Don't be fooled—LLMs don't reason

MIT Tech Review 50%

On an afternoon in Seoul in March 2016, I watched a program I helped build put a stone on the fifth line of a Go board in what looked like a gift to its human opponent. Move 37 in game two of the five-game match looked so absurd that some commentators thought it was a...

<https://www.technologyreview.com/2026/10/02/1145639/dont-be-fooled-llms-dont-reason/>

7 CONF

The Download: AI “mind-reading” and creative uses for small batteries

MIT Tech Review 50%

This is today's edition of The Download, our weekday newsletter that provides a daily dose of what's going on in the world of technology. An AI “mind-reading” tool can reconstruct what you're looking at based on a brain scan A new AI tool can guess what you're looking at just by...

<https://www.technologyreview.com/2026/10/01/1145592/the-download-ai-mind-reading-small-...>

8 CONF

An AI “mind-reading” tool can reconstruct what you're looking at from a brain scan

MIT Tech Review 50%

A new AI tool can guess what you're looking at just by analyzing your brain scans—and re-create that image with remarkable precision. It can go the other way, too, and predict a person's brain activity based on what they're looking at. In the image above, for example, the...

<https://www.technologyreview.com/2026/10/01/1145588/ai-mind-reading-reconstructs-what-y-...>

9 CONF

The Download: climate tech companies to watch and AI's discovery problem

MIT Tech Review 50%

This is today's edition of The Download, our weekday newsletter that provides a daily dose of what's going on in the world of technology. Coming soon: our 2026 list of Climate Tech Companies to Watch With the planet nearing 1.5 °C of warming, climate policies being unraveled,...

<https://www.technologyreview.com/2026/09/29/1145249/the-download-climate-tech-ai-scient-...>

0 CONF

Judge dismisses antitrust lawsuits over Google's AI Overviews

The Verge AI 50%

A federal judge has dismissed a pair of antitrust lawsuits filed by Chegg and Rolling Stone parent company Penske Media Corporation, which accused Google of driving away web traffic with its AI-powered search features, as reported earlier by Reuters. US District Judge Amit Mehta...

<https://www.theverge.com/tech/1003589/google-ai-overviews-chegg-penske-lawsuits-dismissed>

1 CONF

OpenAI's new agent is a shot at Meta — but can it compete with free?

The Verge AI 50%

At OpenAI's annual DevDay conference, the company pulled out all the stops to compete with its rivals - primarily Meta, whose Muse AI agent platform has seen early runaway success. CEO Sam Altman walked onstage to cheers and announced Dots, a "real-deal AI" agent powered by...

<https://www.theverge.com/ai-artificial-intelligence/1003399/meta-openai-ai-agents-muse-...>

2 CONF

Here's what AI leaders are saying about Trump's new safety plan

The Verge AI 50%

After hosting a meal with Big Tech leaders on Tuesday, President Donald Trump responded to journalist questions about his artificial intelligence announcements in typical fashion. He said his previous claims regarding fears around AI safety being a "hoax" cooked up by Democrats...

<https://www.theverge.com/ai-artificial-intelligence/1002636/ai-execs-trump-self-policin-...>

3 CONF

Mimir: Physics-Grounded LLM Agents for Long-Horizon Irrigation Control

ArXiv CS.AI 50%

arXiv:2610.02038v1 Announce Type: new Abstract: Large language model (LLM) agents increasingly combine reasoning, tool use, and action, but most evidence comes from episodic tasks with relatively immediate feedback and reset failures. Long-running physical control operates in a...

<https://arxiv.org/abs/2610.02038>

4 CONF

Chaining Skills to Hijack LLM Agents

ArXiv CS.AI 50%

arXiv:2610.01564v1 Announce Type: cross Abstract: LLM agents use skills to improve performance on specialized tasks. To complete a user request, an agent may invoke several skills in sequence, allowing information produced under one skill to guide the next. Because skills may...

<https://arxiv.org/abs/2610.01564>

5 CONF

From Proposal to Verified Effect: Praxa, an Evidence-Bound Harness for Governed AI Agent Execution

ArXiv CS.AI 50%

arXiv:2610.00015v1 Announce Type: new Abstract: Large-language-model agents can propose and execute actions, but proposal, authority, dispatch, verified external effect, and serving promotion are different claims. We present Praxa, an agent harness that represents these states...

<https://arxiv.org/abs/2610.00015>

6 CONF

Characterizing a Configuration Where Inference-Time PRM-Pruned Fragment Grafting Is Inert: Evidence from Three Reasoning LMs

ArXiv CS.AI 50%

arXiv:2610.00047v1 Announce Type: new Abstract: Diversity collapse in parallel chain-of-thought has motivated inference-time interventions built on a natural design: when a process reward model (PRM) prunes a chain, its high-PRM prefix is extracted and grafted verbatim as an...

<https://arxiv.org/abs/2610.00047>

7 CONF

Evaluating LLM-Generated Preference Distributions

ArXiv CS.AI 50%

arXiv:2610.01000v1 Announce Type: new Abstract: Large Language Models (LLMs) are increasingly used as probabilistic generators for simulation, synthetic data generation, and decision support in settings where real-world data are unavailable. Yet, the structure and reliability of...

<https://arxiv.org/abs/2610.01000>

8 CONF

K-Dense BYOK: An Open-Source AI Research Assistant That Runs Locally and Keeps a Hash-Chained Lab Notebook

ArXiv CS.AI 50%

arXiv:2610.00074v1 Announce Type: new Abstract: K-Dense BYOK (bring your own keys) is a free, open-source AI research assistant for scientists in any field that runs on the researcher's own computer. The researcher supplies access to a model of their choice, hosted or running...

<https://arxiv.org/abs/2610.00074>

9 CONF

Build2SPARQL: A Large-Scale Text-to-SPARQL Benchmark Dataset for Building Knowledge Graph Querying

ArXiv CS.AI 50%

arXiv:2610.00224v1 Announce Type: new Abstract: Building automation systems are increasingly represented as semantic knowledge graphs (KGs) using ontologies such as Brick and ASHRAE 223P, creating a machine-readable substrate for artificial-intelligence applications. One...

<https://arxiv.org/abs/2610.00224>

0 CONF

Rules to Tools: Executable Checks for LLM Agents in Scientific Computing

ArXiv CS.AI 50%

arXiv:2610.00313v1 Announce Type: new Abstract: Scientific coding agents receive equations, boundary conditions, and output requirements in writing, then must assess the programs they revise. Rules to Tools (R2T) supplies prepared executable checks of public scientific...

<https://arxiv.org/abs/2610.00313>

1 CONF

ContractRL: Shielded Group-Relative Policy Optimization for Auditable Tool-Call Repair

ArXiv CS.AI 50%

arXiv:2610.00328v1 Announce Type: new Abstract: Structured tool calls often fail after only a small number of fields violate a schema or an execution contract. Regenerating the complete object enlarges the action surface and makes repeated repair difficult to audit. We introduce...

<https://arxiv.org/abs/2610.00328>

2 CONF

Mathematical Transfer in LLMs Follows Reasoning Approach More Than Topic

ArXiv CS.AI 50%

arXiv:2610.00331v1 Announce Type: new Abstract: When selecting mathematical training data for LLMs, a natural organizing principle is topic: probability examples for probability targets. An alternative is reasoning approach: worked solutions that share a solution method with the...

<https://arxiv.org/abs/2610.00331>

3 CONF

JusticeAxis: Benchmarking Legal Judgment between Rigid Rule Application and Ungrounded Discretion

ArXiv CS.AI 50%

arXiv:2610.00353v1 Announce Type: new Abstract: A sound judgment applies the law to established facts and weighs the circumstances in which they arose. However, existing methods swing between rigid statute matching and ungrounded discretion, benchmarks score a label or a rubric,...

<https://arxiv.org/abs/2610.00353>

4 CONF

When Harnesses Lose the Signal: Causal Evaluation of Recovery in LLM Agents

ArXiv CS.AI 50%

arXiv:2610.00372v1 Announce Type: new Abstract: Large language model agents rely on external harnesses to pass information between the model and its environment and to recover from execution errors. Yet recovery is usually judged only by average task success. This hides an...

<https://arxiv.org/abs/2610.00372>

5 CONF

Counterfactual Auditing of Bias in Open-Source Large Language Models for Clinical Triage

ArXiv CS.AI 50%

arXiv:2610.01963v1 Announce Type: new Abstract: Emergency department (ED) triage is a high-stakes prioritization task in which demographic, socioeconomic, and system-context information may improperly influence acuity assignment. Although open-source large language models (LLMs)...

<https://arxiv.org/abs/2610.01963>

6 CONF

Frozen Scenes, Shifting Winners: Configuration Fragility in Text-to-3D Evaluation

ArXiv CS.AI 50%

arXiv:2610.00447v1 Announce Type: new Abstract: Can a text-to-3D leaderboard change when every generated scene stays fixed? We audit this question for rendered-image evaluation, where camera settings and caption wording become part of the measurement protocol. Across 300 frozen...

<https://arxiv.org/abs/2610.00447>

7 CONF

Beyond Final Accuracy: Auditing Communication in LLM Multi-Agent Systems

ArXiv CS.AI 50%

arXiv:2610.01042v1 Announce Type: new Abstract: Multi-agent communication aims to help agents benefit from one another's information. Yet improvements in system performance leave a fundamental ambiguity: do they reflect effective communication, a favorable agent architecture, or...

<https://arxiv.org/abs/2610.01042>

8 CONF

Ontology-Based Contextual AI Evaluations (OB-CAIE) Methodology

ArXiv CS.AI 50%

arXiv:2610.00529v1 Announce Type: new Abstract: The ontology-based contextual AI evaluation (OB-CAIE) methodology was developed to address a lack of scientific rigor that arises from unclear testing coverage, to balance human expertise and automations, and to address a lack of...

<https://arxiv.org/abs/2610.00529>

9 CONF

ReSolve: Reusing Candidate Reasoning through Selective Generative Moderation

ArXiv CS.AI 50%

arXiv:2610.01140v1 Announce Type: new Abstract: Sampling multiple solutions spends computation on intermediate deductions and unfinished arguments as well as final answers. We introduce ReSolve, a training-free inference procedure that reuses this candidate reasoning through...

<https://arxiv.org/abs/2610.01140>

10 CONF

Spatial Strategies, Not Actions: Vector-Quantized Geodesics as Tools for LLM-Driven Agents

ArXiv CS.AI 50%

arXiv:2610.00613v1 Announce Type: new Abstract: Large language model (LLM) based agents are often criticized for lacking spatial understanding and mainly exploiting statistical text patterns. We investigate their spatial comprehension through an architecture combining...

<https://arxiv.org/abs/2610.00613>

1 CONF

When More Data Is Not Enough: The Context-Sufficiency Frontier in Generative AI Personalization

ArXiv CS.AI 50%

arXiv:2610.00654v1 Announce Type: new Abstract: Personalization has long relied on customer data to infer what an individual is likely to value. We call this customer evidence: the customer's historical behavior and preferences. Generative AI extends personalization by allowing...

<https://arxiv.org/abs/2610.00654>

2 CONF

Ontology-Grounded, Reasoner-Verified Benchmarks for Evaluating LLM Reasoning in Scientific AI

ArXiv CS.AI 50%

arXiv:2610.00682v1 Announce Type: new Abstract: Large language models (LLMs) increasingly underpin scientific AI applications that reason over structured knowledge, from biomedical question answering to materials informatics. However, their logical reasoning often falls short,...

<https://arxiv.org/abs/2610.00682>

3 CONF

Backdoor Purification for LoRA-Tuned LLMs via Null-Space Projection

ArXiv CS.AI 50%

arXiv:2610.00685v1 Announce Type: new Abstract: With the rapid adoption of large language models (LLMs) and parameter-efficient fine-tuning (PEFT) methods, the risk of backdoor attacks has become more severe. Existing backdoor purification methods typically rely on at least one...

<https://arxiv.org/abs/2610.00685>

4 CONF

R-GroundBench: A Diagnostic Benchmark for R-Group Grounding in Markush Molecular Editing

ArXiv CS.AI 50%

arXiv:2610.00700v1 Announce Type: new Abstract: Recent advances in AI for scientific discovery enable molecular understanding and design, yet reasoning over incomplete chemical representations remains unclear. Markush structures, which encode molecular families through variable...

<https://arxiv.org/abs/2610.00700>

5 CONF

Meta-Multi-Agent Reinforcement Learning for Fast Adaptation of Interactive Policies with Applications to Autonomous Driving

ArXiv CS.AI 50%

arXiv:2610.00705v1 Announce Type: new Abstract: This paper develops a meta-multi-agent reinforcement learning (meta-MARL) framework to enable fast adaptation of interactive policies in a multi-agent system (MAS). Meta-reinforcement learning (meta-RL) enables agents to rapidly...

<https://arxiv.org/abs/2610.00705>

6 CONF

Does Scaling Reinforcement Learning Really Require More Training?

ArXiv CS.AI 50%

arXiv:2610.01133v1 Announce Type: cross Abstract: Scaling reasoning typically spends more compute on reinforcement learning (RL) or on inference. We show that a completed RL training history can yield policies stronger than the checkpoints visited by its optimizer. We call this...

<https://arxiv.org/abs/2610.01133>

- 7 **CONF**
- Enterprise Representation Simplification (ERS): Reducing Representational Complexity for Enterprise AI**
- ArXiv CS.AI 50%
- arXiv:2610.00791v1 Announce Type: new Abstract: Enterprise information is represented through artifacts shaped by applications, projects, technologies, organizational boundaries, and local requirements. These structures accumulate over time, creating representational complexity...
- <https://arxiv.org/abs/2610.00791>
-
- 8 **CONF**
- Sapien: A Stateful Policy Engine for Autonomous AI Agents**
- ArXiv CS.AI 50%
- arXiv:2610.00797v1 Announce Type: new Abstract: Contextual security defenses prevent AI agents from taking rogue actions by synthesizing a task-specific policy and enforcing it on the agent's tool calls. In multi-step tasks, however, which actions are valid often depends on what...
- <https://arxiv.org/abs/2610.00797>
-
- 9 **CONF**
- Dependency-Aware Reward Shaping for Agentic Reinforcement Learning**
- ArXiv CS.AI 50%
- arXiv:2610.01207v1 Announce Type: new Abstract: When training large language models with reinforcement learning, terminal rewards provide little guidance about which steps matter. Common methods for assigning step credit overlook that work built on uncorrected mistakes is wasted...
- <https://arxiv.org/abs/2610.01207>
-
- 0 **CONF**
- An Educator-Guided LLM Pedagogical Agent for Scaffolded Feedback in Conceptual Database Design**
- ArXiv CS.AI 50%
- arXiv:2610.00870v1 Announce Type: new Abstract: We present an educator-guided LLM pedagogical agent for scaffolded feedback in conceptual database design. Integrated into an entity--relationship diagram (ERD) editor, the system grounds feedback in the student artifact,...
- <https://arxiv.org/abs/2610.00870>
-
- 1 **CONF**
- OR for AI That Does OR: Routing LLMs up the Escalator inside the OSCAR Framework**
- ArXiv CS.AI 50%
- arXiv:2610.00912v1 Announce Type: new Abstract: Large language models can translate business descriptions into optimization models, but executable code may misrepresent constraints or objectives. A solver can then return an optimal solution to the wrong problem. Even when the...
- <https://arxiv.org/abs/2610.00912>
-
- 2 **CONF**
- PG-SFT: Balancing Capability Acquisition and Retention in Offline Agent Fine-Tuning**
- ArXiv CS.AI 50%
- arXiv:2610.00949v1 Announce Type: new Abstract: Supervised fine-tuning (SFT) on offline agent trajectories is the standard approach for training specialized tool-using agents, but forcing models to imitate reasoning and actions token by token may harm other capabilities (e.g.,...
- <https://arxiv.org/abs/2610.00949>
-

3 CONF

From Discovery to Decision: Finite-Budget Recoverability in LLM Voting

ArXiv CS.AI 50%

arXiv:2610.01014v1 Announce Type: new Abstract: Voting over multiple LLM responses is a common primitive in test-time scaling and ensemble inference. Collecting more responses can expand the candidate pool and increase the chance that a correct answer is discovered. Under a...

<https://arxiv.org/abs/2610.01014>

4 CONF

Beyond State-of-the-Art: Standardising Environmental Impact Metrics for AI Research

ArXiv CS.AI 50%

arXiv:2610.01116v1 Announce Type: new Abstract: As the capabilities and ubiquity of Large Language Models (LLMs) grow, so does their environmental footprint. Despite calls for responsible AI, the machine learning community lacks standardised practices for carbon accounting. Our...

<https://arxiv.org/abs/2610.01116>

5 CONF

Cross-Lingual Alignment for Decoder-Only Models using MoE Routers

ArXiv CS.AI 50%

arXiv:2610.01921v1 Announce Type: cross Abstract: Cross-lingual contrastive learning has been a core component of multilingual encoder training, but the ability to explicitly align representations is not possible in decoder-only LLMs because of varying multilingual tokenization....

<https://arxiv.org/abs/2610.01921>

6 CONF

OrbitTAMP: Grounding Language Models for Task and Motion Planning in Spacecraft Rendezvous

ArXiv CS.AI 50%

arXiv:2610.01093v1 Announce Type: cross Abstract: Spacecraft rendezvous and proximity operations (RPO) are currently planned through an expertise-intensive process in which engineers translate high-level operational intent into safe, dynamically feasible trajectories, creating a...

<https://arxiv.org/abs/2610.01093>

7 CONF

Auditing Action Settlement in LLM Agent Environments: Order, Progress, and Replay

ArXiv CS.AI 50%

arXiv:2610.01138v1 Announce Type: new Abstract: Concurrent actions in large language model (LLM) agent environments require arbitration even when each proposal is individually valid. We implement a typed snapshot-settlement contract and audit three distinct properties: order...

<https://arxiv.org/abs/2610.01138>

8 CONF

Authorization for Self-Modifying AI Agent Populations: Conserving Authority across Replacement, Forking, and Rollback

ArXiv CS.AI 50%

arXiv:2610.00347v1 Announce Type: cross Abstract: Self-modifying AI agents can replace, fork, and roll back identity-bearing software while descendants remain executable. Per-successor authorization does not constrain the resulting population: siblings may duplicate quotas,...

<https://arxiv.org/abs/2610.00347>

9 CONF

HHR: Hierarchical Hash Retrieval for Efficient LLM Generation

ArXiv CS.AI 50%

arXiv:2610.01230v1 Announce Type: new Abstract: Efficient long-context inference is essential for large language models (LLMs), yet it poses a severe computational bottleneck. Hash-based retrieval offers an efficient alternative by encoding queries and keys into binary codes and...

<https://arxiv.org/abs/2610.01230>

0 CONF

Learning to Ask: Information Acquisition for SLM-LLM Collaboration, under a budget

ArXiv CS.AI 50%

arXiv:2610.01236v1 Announce Type: new Abstract: Collaboration between a small language model (SLM) and a large language model (LLM) offers an opportunity to combine the efficiency of smaller models with the strong reasoning capabilities of larger ones. Existing approaches...

<https://arxiv.org/abs/2610.01236>

1 CONF

From Verification Failures to Reusable Guidance for Coding Agents

ArXiv CS.AI 50%

arXiv:2609.39022v2 Announce Type: replace-cross Abstract: Coding agents need to establish that a program satisfies a specification and that the specification captures the requested behavior. We study how expert diagnosis of verification failures can become reusable guidance for...

<https://arxiv.org/abs/2609.39022>

2 CONF

Feedback Without the Wait: Piloting a Generative AI Practice Platform in a Large Maths Class

ArXiv CS.AI 50%

arXiv:2610.01262v1 Announce Type: new Abstract: Timely and specific feedback is one of the strongest influences on student learning, yet it is difficult to sustain in large electrical engineering classes where the ratio of students to demonstrators is high and a learner who is...

<https://arxiv.org/abs/2610.01262>

3 CONF

Questionnaire-Guided Disaggregation of Energy Appliance Use for Domestic Smart Meter Data

ArXiv CS.AI 50%

arXiv:2610.01297v1 Announce Type: new Abstract: Ireland's smart metering programme records electricity use at 30-minute resolution, with smart meters installed in over 80% of households as of late 2025. While this is useful for billing of smart, time-of-use tariffs, it is too...

<https://arxiv.org/abs/2610.01297>

4 CONF

DAYJOB: A Benchmark for Long-Horizon Professional Work

ArXiv CS.AI 50%

arXiv:2610.01306v1 Announce Type: new Abstract: Professional work often starts with a brief request that leaves the professional to work out what is needed, which documents matter, and whether the request's premise holds. We introduce DAYJOB, a benchmark of 130 tasks built by...

<https://arxiv.org/abs/2610.01306>

5 CONF

A framework for auditing grounding claims

ArXiv CS.AI 50%

arXiv:2512.06205v3 Announce Type: replace Abstract: The symbol grounding problem asks how a token such as cat can be about cats. We propose a framework for auditing grounding claims against a declared semantic standard. The audit reports measurements and evidence, with overall...

<https://arxiv.org/abs/2512.06205>

6 CONF

AiSearch: Interactive Multi-Modal Search with VLMs

ArXiv CS.AI 50%

arXiv:2610.01389v1 Announce Type: new Abstract: Modern retrieval systems must both be automated and interactive, allowing users to search and refine results in real time. We present AiSearch, a flexible multimodal retrieval framework that leverages the zero shot capabilities of...

<https://arxiv.org/abs/2610.01389>

7 CONF

Contrastive Attention Mitigates Spectral Bias in Spiking Transformers

ArXiv CS.AI 50%

arXiv:2610.01403v1 Announce Type: new Abstract: Spiking Transformers merge the energy-efficiency of spiking neural networks (SNNs) with the representational power of self-attention, creating a promising architecture for high-performance, energy-efficient computation. However, a...

<https://arxiv.org/abs/2610.01403>

8 CONF

SpikeMoE: Brain-Inspired Competitive Routing for Flexible Spiking Mixture-of-Experts

ArXiv CS.AI 50%

arXiv:2610.01418v1 Announce Type: new Abstract: Spiking Neural Networks (SNNs) enable event-driven computation through biologically inspired dynamics at the neuronal scale, while Mixture-of-Experts (MoE) perform conditional computation through expert selection at the model...

<https://arxiv.org/abs/2610.01418>

9 CONF

A Deterministic and Auditable AI Security Risk Assessment Framework with ATLAS Aligned Executable Rules and Formal Verification

ArXiv CS.AI 50%

arXiv:2610.01436v1 Announce Type: new Abstract: Artificial intelligence systems are increasingly deployed in high impact and safety critical settings, yet security assessment remains difficult to reproduce and defend under audit. Existing approaches often rely on narrative...

<https://arxiv.org/abs/2610.01436>

0 CONF

DRelay: Global Draft Context for Prefix-Aware Parallel Speculative Decoding Repair

ArXiv CS.AI 50%

arXiv:2610.01439v1 Announce Type: new Abstract: Parallel drafting reduces the drafting overhead of speculative decoding for large language models (LLMs), but its gains remain limited by the accepted prefix length. Even when the correct token is present in the candidate pool, a...

<https://arxiv.org/abs/2610.01439>

1 CONF

Auditing Routing Entropy as an Uncertainty Signal in Attention-Residual Transformers

ArXiv CS.AI 50%

arXiv:2610.01495v1 Announce Type: new Abstract: Dynamic architectures leave a per-example routing trace beside each prediction, and diffuse routing is easy to read as a sign that the prediction is unreliable. We audit that reading for routing entropy in Attention-Residual (AR)...

<https://arxiv.org/abs/2610.01495>

- 2

CONF

OpenMTB-Audit: Exposing Over-Refusal and Clinical Expert Perspectives in LLM-Based Molecular Tumor Board Safety Evaluation

ArXiv CS.AI

50%

arXiv:2610.01497v1 Announce Type: new Abstract: Molecular tumor boards integrate genomic findings, clinical context, and therapeutic evidence to support precision oncology. As AI enters this workflow, a key safety challenge is distinguishing truly unsupported recommendations...

<https://arxiv.org/abs/2610.01497>
- 3

CONF

Sharpening Tax in Post-Training

ArXiv CS.AI

50%

arXiv:2610.01509v1 Announce Type: new Abstract: An emerging hypothesis about reinforcement learning (RL) post-training of large language models (LLMs) is that it merely sharpens existing behaviors of a base model, improving single-shot accuracy at the cost of solution coverage....

<https://arxiv.org/abs/2610.01509>
- 4

CONF

The AI Assessment Sandbox Configurator: A Framework to Support Technical Assessment in AI Regulatory Sandboxes

ArXiv CS.AI

50%

arXiv:2610.01539v1 Announce Type: new Abstract: The EU's Artificial Intelligence Act requires all Member States to establish AI Regulatory Sandboxes (AIRS) by August 2027: supervised environments bringing together national Competent Authorities, technical experts, and the...

<https://arxiv.org/abs/2610.01539>
- 5

CONF

Benchmarking Generative Models for Weather Data Assimilation on Real Station Observations

ArXiv CS.AI

50%

arXiv:2610.00728v1 Announce Type: cross Abstract: Weather reanalysis products rely on computationally intensive numerical weather predictions followed by data assimilation that corrects the forecast toward observations. Deep generative models offer a cheaper alternative that...

<https://arxiv.org/abs/2610.00728>
- 6

CONF

Homomorphic Advantage Operator: Stabilizing Reinforcement Learning Under Fully Homomorphic Encryption Constraints

ArXiv CS.AI

50%

arXiv:2610.02074v1 Announce Type: new Abstract: Privacy-preserving machine learning presents significant deployment challenges on the cloud for intelligent systems with confidential data. Fully Homomorphic Encryption (FHE) offers a compelling solution for secure computation,...

<https://arxiv.org/abs/2610.02074>
- 7

CONF

AI-assisted mitotic counting improves reproducibility and efficiency across multiple tumour types

ArXiv CS.AI

50%

arXiv:2610.01813v1 Announce Type: new Abstract: Mitotic counting is an important component of tumour grading, diagnosis and prognostic assessment across several tumour types, but manual assessment is time-consuming and subject to inter-pathologist variability. To help address...

<https://arxiv.org/abs/2610.01813>

8 CONF

Temporal-Difference Learning for Dragonchess

ArXiv CS.AI 50%

arXiv:2610.01845v1 Announce Type: new Abstract: Our research investigates how two adaptive AI methods, evolutionary transfer learning and TD(lambda), perform in the three-dimensional chess environment Dragonchess. The game challenges players with its unique board structure and...

<https://arxiv.org/abs/2610.01845>

9 CONF

Mapping the RAG Landscape: A Four Axis Taxonomy of Efficiency, Defense, Interactivity, and Reasoning

ArXiv CS.AI 50%

arXiv:2610.01936v1 Announce Type: new Abstract: Large Language Models (LLMs) have demonstrated remarkable fluency across many tasks but remain limited by their static, parameter bound knowledge and their susceptibility to hallucinating information. Retrieval Augmented Generation...

<https://arxiv.org/abs/2610.01936>

0 CONF

Can AI Oversight Be Zero Knowledge?

ArXiv CS.AI 50%

arXiv:2610.01995v1 Announce Type: new Abstract: AI systems increasingly produce outputs from confidential data, such as a fitness-for-duty assessment from medical records or the predicted properties of a drug candidate from its secret structure. It is important to verify that...

<https://arxiv.org/abs/2610.01995>

1 CONF

Counting Moves, Weighing Voices: Bayesian Dialectical Argumentation for Calibrated Multi-LLM Councils under Persistent Adversaries

ArXiv CS.AI 50%

arXiv:2610.02005v1 Announce Type: new Abstract: A multi-LLM \emph{council} lets several large language models (LLMs) deliberate on a question and return an answer together with a confidence estimate. As these systems become increasingly used for reasoning, that confidence should...

<https://arxiv.org/abs/2610.02005>

2 CONF

Atoms to Processes: The Role of Artificial Intelligence and Machine Learning in Chemical Engineering

ArXiv CS.AI 50%

arXiv:2610.02014v1 Announce Type: new Abstract: The rapid maturation of artificial intelligence (AI) and machine learning (ML) has catalyzed a profound shift in how chemical engineering problems are formulated, analyzed, and solved. Advances in computing, data availability, and...

<https://arxiv.org/abs/2610.02014>

3 CONF

SPHERE: Adaptive VR Indoor Scene Generation via LLM-Enhanced Spatial Preference Learning and Human-in-the-Loop RL

ArXiv CS.AI 50%

arXiv:2610.02023v1 Announce Type: new Abstract: While Large Language Models (LLMs) advance 3D indoor scene synthesis, current pipelines fail to retain user-specific preferences across sessions, making immersive authoring a repetitive and physically fatiguing process. We present...

<https://arxiv.org/abs/2610.02023>

4 CONF

HydroJEV: A one-second, training-free screen for cyber-attack and fault attribution in water distribution networks

ArXiv CS.AI 50%

arXiv:2610.02048v1 Announce Type: new Abstract: When a SCADA alarm is raised in a water distribution network, operators must decide quickly whether it reflects a cyberattack, a physical fault, a normal transient or a faulty sensor. Supervised classifiers need labelled incidents...

<https://arxiv.org/abs/2610.02048>

5 CONF

External Observers May See More Clearly: Cross-Model Span-Level Hallucination Detection in Large Language Models via Hidden State Probing

ArXiv CS.AI 50%

arXiv:2610.02066v1 Announce Type: new Abstract: As Large Language Models (LLMs) increasingly serve as foundational reasoning engines, their tendency to hallucinate remains a critical vulnerability. While recent internal state probes offer a promising alternative to slow external...

<https://arxiv.org/abs/2610.02066>

6 CONF

PyPottery: an AI-powered end-to-end suite for pottery processing and publication

ArXiv CS.AI 50%

arXiv:2610.02072v1 Announce Type: new Abstract: The study of ceramic materials constitutes a cornerstone of archaeological research, yet the post-production workflow for pottery documentation remains labor-intensive and creates significant publication bottlenecks. This paper...

<https://arxiv.org/abs/2610.02072>

7 CONF

ECHO: A Participatory Framework for Bias-Anchored AI Harm Anticipation

ArXiv CS.AI 50%

arXiv:2512.03068v2 Announce Type: replace-cross Abstract: Artificial Intelligence (AI) systems increasingly shape consequential decisions, creating value but also potential harms for individuals, social groups, and society. This has prompted calls for proactive approaches that...

<https://arxiv.org/abs/2512.03068>

8 CONF

ScholarCatalyst: A Benchmark for Retrieving Papers That Inspire New Research

ArXiv CS.AI 50%

arXiv:2610.02202v1 Announce Type: new Abstract: What makes great scientists great? Even as AI systems start to make progress on open problems, scientists remain far ahead of them at sensing which prior idea, buried in an ever-growing archive of research, a new problem needs. To...

<https://arxiv.org/abs/2610.02202>

9 CONF

Bounded-Fidelity Sim-as-Demo-Stage: Mocap Handoff for Governance Benchmarks

ArXiv CS.AI 50%

arXiv:2610.00008v1 Announce Type: cross Abstract: Sim-to-real research pursues physics fidelity as a primary objective: simulators are judged by how closely they reproduce real-world contact dynamics. For governance benchmarking of LLM-driven robots, where the simulator...

<https://arxiv.org/abs/2610.00008>

0 **CONF****Ask a Language Model for Lottery Numbers: Concentration in Repeated Six-of-49 Outputs****ArXiv CS.AI 50%**

arXiv:2610.00052v1 Announce Type: cross Abstract: We evaluate six language-model configurations on requests for six distinct random integers from 1-49. Across 1,200 attempted calls using four English prompt variants, 1,184 responses yielded valid tickets. Effective diversity of...

<https://arxiv.org/abs/2610.00052>

1 **CONF****The First Token Is Not the Verdict: Hidden Costs of Reading LLM Judges Without Generating****ArXiv CS.AI 50%**

arXiv:2610.00054v1 Announce Type: cross Abstract: Reading an LLM judge's verdict from the logits of its first generated token is cheap, requires no generation, and is exactly what constrained decoding and likelihood-scoring evaluation harnesses produce. We show that this readout...

<https://arxiv.org/abs/2610.00054>

2 **CONF****Pushing CPU Speech Synthesis to the Wall: Extreme Inference Tuning under Serverless Architecture and Billing****ArXiv CS.AI 50%**

arXiv:2610.00063v1 Announce Type: cross Abstract: Instance-billed serverless platforms charge for CPU and memory over the lifetime of a warm instance, making idle inference state a direct serving cost. We present billing-aware neural text-to-speech (TTS) serving on serverless...

<https://arxiv.org/abs/2610.00063>

3 **CONF****Legal text classification in Korean sexual offense cases: from traditional machine learning to large language models with XAI insights****ArXiv CS.AI 50%**

arXiv:2610.00087v1 Announce Type: cross Abstract: The advancement of natural language processing (NLP) has expanded AI-based text classification in the legal domain. However, accurately classifying legal documents remains challenging due to the complexity of legal texts and...

<https://arxiv.org/abs/2610.00087>

4 **CONF****The Cognitive Continuity Test: Verifying Governed State Transitions in Persistent AI Agents****ArXiv CS.AI 50%**

arXiv:2610.00132v1 Announce Type: cross Abstract: Persistent AI agents revise beliefs, consolidate memory, and replace execution substrates. Similar successor states can accompany differently authorized transition claims, while legitimate development can change state...

<https://arxiv.org/abs/2610.00132>

5 **CONF****When the AI Leaves the Tailorshop: Measuring What an LLM Advisor Leaves Behind in Complex Problem Solving****ArXiv CS.AI 50%**

arXiv:2610.00163v1 Announce Type: cross Abstract: Complex problem solving depends on acting effectively and understanding how a system works. AI advice may support these outcomes unequally. Two preregistered experiments compared participants managing a simulated clothing factory...

<https://arxiv.org/abs/2610.00163>

6 CONF

ShatterQuant: Breaking Uniform Precision with Block-Wise Mixed-Precision on a Systolic Transformer Hardware Accelerator

ArXiv CS.AI 50%

arXiv:2610.00207v1 Announce Type: cross Abstract: Due to limited support for intra-tensor heterogeneous precision in conventional accelerators, neural network quantization remains largely restricted to per-tensor precision assignment. We present ShatterQuant, a hardware-software...

<https://arxiv.org/abs/2610.00207>

7 CONF

Certainty Is Not Just Correctness: Rethinking Token-Level Certainty in LLM Reasoning

ArXiv CS.AI 50%

arXiv:2610.00296v1 Announce Type: cross Abstract: Token-level certainty is widely used as a proxy for correctness in LLM training and inference. However, the performance of certainty-based methods depends both on the information in certainty scores and on how those scores are...

<https://arxiv.org/abs/2610.00296>

8 CONF

The Geometry of Contextual Relations: Language Models Address Facts by Order of Mention

ArXiv CS.AI 50%

arXiv:2610.00910v1 Announce Type: cross Abstract: Human reasoning depends on how objects are related within propositions. \textit{How do relations organize the language representations of contextual contents?} We give an LLM a list of facts in its context (e.g., \emph{Alice eats...}

<https://arxiv.org/abs/2610.00910>

9 CONF

Tokenized Key-Gated Adapter Routing: A Secure Access Control Mechanism Against Private Data Leakage in LLMs

ArXiv CS.AI 50%

arXiv:2610.00309v1 Announce Type: cross Abstract: Large language models (LLMs) are increasingly deployed in privacy-critical domains (e.g., healthcare, finance, and government), but their propensity to memorize and disclose personally identifiable information (PII) poses serious...

<https://arxiv.org/abs/2610.00309>

00 CONF

Actions with Receipts: Jointly Binding Claims, Evidence, and Execution for Replayable Tool-Agent Auditing

ArXiv CS.AI 50%

arXiv:2610.00327v1 Announce Type: cross Abstract: Tool-using agents can expose citations and execution logs while leaving a critical association unaudited: whether the claim shown to a user is the claim emitted by the committed execution and supported by the cited source. A...

<https://arxiv.org/abs/2610.00327>

01 CONF

LEGO-OPD: Factorized Teacher Composition for Multimodal On-Policy Distillation

ArXiv CS.AI 50%

arXiv:2610.00333v1 Announce Type: cross Abstract: Multimodal on-policy distillation (OPD) aims to improve visual grounding while preserving the strong reasoning capabilities of language models. Recent multi-teacher approaches combine LLM and VLM teachers to provide complementary...

<https://arxiv.org/abs/2610.00333>

02 CONF

Representation Transitions Reveal Emerging Safety Risks in Multi-Turn LLM Agents

ArXiv CS.AI 50%

arXiv:2610.00400v1 Announce Type: cross Abstract: Multi-turn attacks on agentic systems can compose individually permissible actions into harmful outcomes, challenging defenses that assess actions or states in isolation. We show that such attacks leave a detectable signature in...

<https://arxiv.org/abs/2610.00400>

03 CONF

Diffusion Editing with Soft Mask: Pixel Level Redo of Image and Video with Adjustable Strength

ArXiv CS.AI 50%

arXiv:2610.00359v1 Announce Type: cross Abstract: Diffusion models with prompt and reference image-guided editing have seen rapid progress, yet they remain too coarse for pixel-level control. One promising direction is to incorporate a soft mask that specifies spatially varying...

<https://arxiv.org/abs/2610.00359>

04 CONF

Learning Meta-Skills for Agent Harness Design in Test-Time AI4AI

ArXiv CS.AI 50%

arXiv:2609.38143v2 Announce Type: replace Abstract: Agent performance depends on both reasoning ability and the environment in which it acts. We study test-time AI-for-AI, asking how a Builder can learn to construct better execution environments for a Target while both models'...

<https://arxiv.org/abs/2609.38143>

05 CONF

A Shared Taste for Model-Written Text: The Generator-by-Selector Matrices of "AI-AI Bias" Show No Detectable Own-Model Premium

ArXiv CS.AI 50%

arXiv:2610.00369v1 Announce Type: cross Abstract: Laurito et al. (PNAS 2025) showed that large language models choosing between two descriptions of the same product, paper or film prefer the description written by a language model over the one written by a person, by a wide...

<https://arxiv.org/abs/2610.00369>

06 CONF

Deny Without Disabling: Authorization-Paired Evaluation and Control for Multi-Agent Systems

ArXiv CS.AI 50%

arXiv:2610.00371v1 Announce Type: cross Abstract: Multi-agent systems derive their capabilities from sharing evidence, delegating tasks, and combining information across agents. The same process creates a safety problem: contributions that are admissible in isolation can jointly...

<https://arxiv.org/abs/2610.00371>

07 CONF

Faithful Chart Generation for Multimodal Deep Research: Frame-Evidence Co-Adaptation

ArXiv CS.AI 50%

arXiv:2610.00374v1 Announce Type: cross Abstract: Analytical charts in multimodal deep research encode quantitative claims, requiring every visualized value to be faithfully grounded in supporting evidence. Unlike retrieved images that mainly provide contextual information,...

<https://arxiv.org/abs/2610.00374>

08 CONF

Scores That Hold, Benchmarks That Leak: Measuring Dataset Contamination in Public Brain-Tumor MRI Classification

ArXiv CS.AI 50%

arXiv:2610.00421v1 Announce Type: cross Abstract: Automated classification of brain tumors from MRI is a heavily published application of deep learning in medical imaging, with reported accuracies on public benchmarks routinely exceeding 98%. However, accuracy does not capture a...

<https://arxiv.org/abs/2610.00421>

09 CONF

Target-Dependent Limits of Causal Repair: A Leading-Log Frontier in a Gaussian Model

ArXiv CS.AI 50%

arXiv:2610.00424v1 Announce Type: cross Abstract: Knowing how much a causal predictor could improve need not reveal the gain of the repair actually learned. We quantify this gap in a scalar Gaussian causal experiment with known intervention geometry: auxiliary data identify...

<https://arxiv.org/abs/2610.00424>

10 CONF

Code That Works, Environments That Don't: Measuring Environment Reproducibility in AI-Generated Software

ArXiv CS.AI 50%

arXiv:2610.00425v1 Announce Type: cross Abstract: Code generation has emerged as a central capability of large language models, with coding agents now able to produce functionally correct software projects from natural language prompts. However, functional correctness alone does...

<https://arxiv.org/abs/2610.00425>

11 CONF

XOR-Trellis: Ultra-Low-Complexity Dequantization and Curvature-Aware Hadamard-Free LLM Quantization

ArXiv CS.AI 50%

arXiv:2610.00432v1 Announce Type: cross Abstract: Trellis-coded quantization enables high-dimensional compression of large language model (LLM) weights at ultra-low bit widths without the exponentially large codebooks required by conventional vector quantization. Practical...

<https://arxiv.org/abs/2610.00432>

12 CONF

How AI Agents Discover Scientific Equations: From Hydrotape Rediscovery to New Water-Wave Amplitudes

ArXiv CS.AI 50%

arXiv:2610.00435v1 Announce Type: cross Abstract: We study how AI agents discover and validate scientific formulas using a controlled case study of the hydrotape, a recently discovered geometric formula that combines the different polynomial pieces of nonlinear surface-wave...

<https://arxiv.org/abs/2610.00435>

13 CONF

Gumbel Straight Flow: Distilling Autoregressive Models into One-step Flow Maps

ArXiv CS.AI 50%

arXiv:2610.00497v1 Announce Type: cross Abstract: We present Gumbel Straight Flow (GSF), a continuous flow map language model that leverages the noise-data coupling of a pretrained autoregressive language (AR) model. We theoretically demonstrate that the coupling between Gumbel...

<https://arxiv.org/abs/2610.00497>

14 CONF

Rules Amortize, Pairings Don't: Linguistic Structure Determines What Latent Task Representations Can Replace In-Context Learning

ArXiv CS.AI 50%

arXiv:2610.00526v1 Announce Type: cross Abstract: In-context learning (ICL) can be amortized into latent objects (task vectors, function vectors, context vectors) that recover few-shot behavior at zero-shot inference cost, but recent theory shows a static vector acts as a single...

<https://arxiv.org/abs/2610.00526>

15 CONF

Can LLMs Reason Over Long Horizons? An Empirical Evaluation of Context Strategies for Longitudinal Clinical Reasoning

ArXiv CS.AI 50%

arXiv:2610.00562v1 Announce Type: cross Abstract: Longitudinal clinical reasoning requires large language models (LLMs) to identify and integrate relevant evidence distributed across extended patient histories. Although long-context models can process increasingly large amounts...

<https://arxiv.org/abs/2610.00562>

16 CONF

Emergent Unfaithfulness: How Alignment Training Causes Language Models to Silently Override Task Faithfulness

ArXiv CS.AI 50%

arXiv:2610.00568v1 Announce Type: cross Abstract: Large language models are characterized by three key properties: capability, alignment, and faithfulness. Prior work studies the tradeoffs between capability and alignment, and between capability and faithfulness, but a third...

<https://arxiv.org/abs/2610.00568>

17 CONF

Masked Self-Distillation: Internalizing the Chain-of-Thought in Language Models

ArXiv CS.AI 50%

arXiv:2607.22629v3 Announce Type: replace Abstract: Large Reasoning Models produce long, explicit chains of intermediate steps before generating a final answer at inference time. These intermediate traces dominate latency, memory usage, and serving cost, even though final answer...

<https://arxiv.org/abs/2607.22629>

18 CONF

Beyond Supra-Competitive Outcomes: Collusive Behaviour in Deep Reinforcement Learning for Optimal Execution Games

ArXiv CS.AI 50%

arXiv:2610.00619v1 Announce Type: cross Abstract: In this paper, we extend earlier findings of supra-competitive outcomes in optimal-execution games by identifying a learned punitive mechanism that deters deviations and provides behavioural evidence of collusion. We investigate...

<https://arxiv.org/abs/2610.00619>

19 CONF

Misalignment of Low-Loss Regions Causes Grokking

ArXiv CS.AI 50%

arXiv:2610.00620v1 Announce Type: cross Abstract: Grokking refers to the delayed emergence of validation-set generalization after a model has already overfit the training set. Although first observed in small algorithmic tasks trained with transformers, its underlying mechanism...

<https://arxiv.org/abs/2610.00620>

20 **CONF****Self-Evolving Coding Rules for AI Coding Agents**

ArXiv CS.AI 50%

arXiv:2610.00650v1 Announce Type: cross Abstract: The performance of AI coding agents is highly dependent on their underlying coding rules. However, existing coding rules are typically hand-crafted and fixed, making the process labor-intensive and often suboptimal. In this work,...

<https://arxiv.org/abs/2610.00650>
21 **CONF****VisionQ: VLM-as-a-Judge Taxonomy, Dataset and Benchmark for Qualitative Analysis in Computer Vision**

ArXiv CS.AI 50%

arXiv:2610.00666v1 Announce Type: cross Abstract: Qualitative comparison figures are central evidence in computer vision papers, and vision-language models (VLMs) are increasingly used to judge them. Yet existing benchmarks score only scalar quality or overall preference, so a...

<https://arxiv.org/abs/2610.00666>
22 **CONF****Initialization Improves LLM-Driven Discovery**

ArXiv CS.AI 50%

arXiv:2610.00707v1 Announce Type: cross Abstract: Large Language Models (LLMs) have been used for novel discovery of algorithms, theorems, drugs, and other tasks through the use of harnesses that prompt an LLM to iteratively optimize an objective. In this work, we study the...

<https://arxiv.org/abs/2610.00707>
23 **CONF****Outer Diversity of Condorcet Domains**

ArXiv CS.AI 50%

arXiv:2610.00720v1 Announce Type: cross Abstract: A Condorcet domain is a set of rankings over a given candidate set, such that every election that consists only of (an odd number of) votes from the domain has a transitive majority relation. We study outer diversity of Condorcet...

<https://arxiv.org/abs/2610.00720>
24 **CONF****Program Learning with Verifiable Rewards: Symbolic Backpropagation for Post-Training LLMs**

ArXiv CS.AI 50%

arXiv:2608.28421v2 Announce Type: replace Abstract: Post training a language model to reason means updating its weights. Supervised finetuning and reinforcement learning both place the acquired capability inside the model where it cannot be inspected cannot be checked step by...

<https://arxiv.org/abs/2608.28421>
25 **CONF****Pre-training interventions, ex post facto: Grafting model beliefs across checkpoints**

ArXiv CS.AI 50%

arXiv:2610.00767v1 Announce Type: cross Abstract: Pre-training interventions are critical to alignment research, since beliefs formed during pre-training shape how a model generalizes from later training. One recently popular technique for such interventions is synthetic...

<https://arxiv.org/abs/2610.00767>

26 CONF

Video Generation Models: A Survey of Post-Training and Alignment

ArXiv CS.AI 50%

arXiv:2610.00812v1 Announce Type: cross Abstract: Video generation has rapidly progressed from short, low-quality clips to high-resolution, long-duration sequences with complex spatiotemporal dynamics. Despite strong generative priors learned through large-scale pretraining,...

<https://arxiv.org/abs/2610.00812>

27 CONF

Training-Aware Target Coverage for Synthetic Data Selection

ArXiv CS.AI 50%

arXiv:2610.00814v1 Announce Type: cross Abstract: Synthetic data are increasingly used to scale LLM training, yet more synthetic data do not necessarily produce better models. Useful synthetic data must add information relevant to the target task without introducing errors that...

<https://arxiv.org/abs/2610.00814>

28 CONF

Contextual trajectory and incremental contextual displacement: Towards using LLMs to understand dynamic, utterance-specific meaning construction

ArXiv CS.AI 50%

arXiv:2610.00840v1 Announce Type: cross Abstract: Transformer-based large language models (LLMs) such as RoBERTa represent text using contextual word embeddings (CWEs), which alter the embeddings associated with each token based on surrounding context. We construct token-wise...

<https://arxiv.org/abs/2610.00840>

29 CONF

Don't Waste the Noise: Importance-Guided Perturbation Allocation under Joint Global and Local Constraints

ArXiv CS.AI 50%

arXiv:2610.00861v1 Announce Type: cross Abstract: Adversarial optimization under a shared ℓ_1 budget requires deciding not only how much perturbation to use, but also where that limited budget should be spent. This allocation problem becomes particularly important when...

<https://arxiv.org/abs/2610.00861>

30 CONF

Rethinking Data Augmentation under Covariate Shift: Invariant-Guided Diffusion and Prototype Reweighting

ArXiv CS.AI 50%

arXiv:2610.00873v1 Announce Type: cross Abstract: In many industrial applications, 1) tabular data is scarce and imbalanced and thus requires synthetic expansion; 2) input distributions drift between training and deployment (covariate shift); 3) validation sets often diverge...

<https://arxiv.org/abs/2610.00873>

31 CONF

Match the Distribution, Not the Compute: Post-Training Multi-Token Prediction Heads

ArXiv CS.AI 50%

arXiv:2610.00888v1 Announce Type: cross Abstract: Multi-token prediction (MTP) improves the throughput of autoregressive generation by enabling the language model to draft multiple next tokens per forward pass, while a verification step over draft tokens ensures that token...

<https://arxiv.org/abs/2610.00888>

32 CONF

DynGhost: Temporally-Modelled Transformer for Dynamic Ghost Imagings

ArXiv CS.AI 50%

arXiv:2605.10185v3 Announce Type: replace-cross Abstract: Ghost imaging reconstructs spatial information from a single-pixel bucket detector by correlating structured illumination patterns with scalar intensity measurements. While deep learning approaches have achieved promising...

<https://arxiv.org/abs/2605.10185>

33 CONF

Mean field games as a tool for AI safety: a worked example from the July 2026 Hugging Face incident

ArXiv CS.AI 50%

arXiv:2610.00902v1 Announce Type: cross Abstract: One way to make AI systems safe is to shape what the system is: its objective and dispositions. We take a complementary route: treat the agents' characteristics as partly unknown and ask what structure of interaction ensures that...

<https://arxiv.org/abs/2610.00902>

34 CONF

Two Clocks in Diffusion MLLMs: When Answers Stabilize Before Rationales Unfold

ArXiv CS.AI 50%

arXiv:2610.00953v1 Announce Type: cross Abstract: An answer candidate in a masked diffusion MLLM can stabilize while its rationale is still unfolding. We distinguish retrospective stabilization of the logged candidate from token commitment, and examine these two clocks relative...

<https://arxiv.org/abs/2610.00953>

35 CONF

RelationVGGT: Visual Geometry Transformers for 3D Spatial Relation Segmentation

ArXiv CS.AI 50%

arXiv:2610.00970v1 Announce Type: cross Abstract: Recent advances in 3D reconstruction have progressed from per-scene optimization to feed-forward inference, and semantic scene understanding has followed suit -- yet existing methods remain confined to object-centric perception,...

<https://arxiv.org/abs/2610.00970>

36 CONF

MOMAT: Mixture of Multiple Atlases for Low-Power Jailbreak Defense of Quantized LLMs

ArXiv CS.AI 50%

arXiv:2610.01058v1 Announce Type: cross Abstract: Quantized large language models are increasingly deployed on edge devices for their low latency and energy efficiency. However, model quantization weakens alignment safeguards, leaving qLLMs (quantized large language models)...

<https://arxiv.org/abs/2610.01058>

37 CONF

Parameter-Efficient Distributionally Robust Adaptation of Tabular Foundation Models under Subpopulation Shift

ArXiv CS.AI 50%

arXiv:2610.01143v1 Announce Type: cross Abstract: Despite strong mean accuracy, tabular foundation models (TFMs) can perform poorly on underrepresented groups under subpopulation shift, where group proportions change between training and deployment. We propose DR-TFM, a...

<https://arxiv.org/abs/2610.01143>

38 CONF

Detect, Explain, Interpret: An End-to-End Benchmark for Time Series Anomaly Detection, Explainability and Interpretability

ArXiv CS.AI 50%

arXiv:2610.01168v1 Announce Type: cross Abstract: Time Series Anomaly Detection has received increasing attention, driven by the growing availability of complex time series data. This surge has led to the development of numerous detection methods, as well as a variety of...

<https://arxiv.org/abs/2610.01168>

39 CONF

Have an LLM Write Your Anomaly Detector: Autonomous Discovery of Compact, Interpretable Detectors for Time Series

ArXiv CS.AI 50%

arXiv:2610.01223v1 Announce Type: cross Abstract: Time-series anomaly detection trades off predictive accuracy, computational efficiency, and interpretability. We use a large language model not as the detector but as the author of one: an autonomous research loop in which the...

<https://arxiv.org/abs/2610.01223>

40 CONF

Science Utopia? Closed-Loop LLM Simulation of Academic Research Ecosystems

ArXiv CS.AI 50%

arXiv:2610.01257v1 Announce Type: cross Abstract: Scientific progress emerges from a longitudinal ecosystem in which researchers, institutions, funding agencies, collaboration networks, and the scientific literature co-evolve. As AI becomes increasingly involved throughout the...

<https://arxiv.org/abs/2610.01257>

41 CONF

Model validation in machine learning: A scenario-based guide from hold-out splits to nested group cross-validation in biomedical and applied research

ArXiv CS.AI 50%

arXiv:2610.01284v1 Announce Type: cross Abstract: Model validation estimates the performance of a complete learning procedure on new data. However, an invalid split can produce an optimistic and stable result. This tutorial reviews hold-out validation, train/validation/test...

<https://arxiv.org/abs/2610.01284>

42 CONF

Federated Learning for LLMs over Mobile Networks: Issues and Solutions in the RAN Transport

ArXiv CS.AI 50%

arXiv:2610.01304v1 Announce Type: cross Abstract: Federated LLM fine-tuning enables large models to be adapted using private and geographically distributed data at the network edge, creating recurring and deadline-sensitive communication workloads across access and transport...

<https://arxiv.org/abs/2610.01304>

43 CONF

Feature Selective Model Collapse in Diffusion Models: Total Replacement versus Fixed-Budget Training

ArXiv CS.AI 50%

arXiv:2610.01318v1 Announce Type: cross Abstract: Model collapse arises when generative models are trained on synthetic data produced by earlier models. The phenomenon has attracted considerable attention because of its societal and technical implications. However, previous...

<https://arxiv.org/abs/2610.01318>

44 CONF

Fold'EM: Direct atomic structure inference from Cryo-EM particles

ArXiv CS.AI 50%

arXiv:2610.01358v1 Announce Type: cross Abstract: Single-particle cryo-electron microscopy (cryo-EM) has become a widely adopted technique for biomolecular structure determination. The conventional cryo-EM computational pipeline first combines many particle images to reconstruct...

<https://arxiv.org/abs/2610.01358>

45 CONF

LLM-Driven Multi-Agent Control for Skill-Based Smart Manufacturing

ArXiv CS.AI 50%

arXiv:2610.01364v1 Announce Type: cross Abstract: Factories are shifting toward smaller lot sizes with high product customization, requiring frequent re-programming of flexible and reconfigurable automation systems. LLM-based agents can be deployed in two complementary roles:...

<https://arxiv.org/abs/2610.01364>

46 CONF

LLM-Assisted Discovery of Typed Semantic Links for Ontology Network Construction

ArXiv CS.AI 50%

arXiv:2610.01393v1 Announce Type: cross Abstract: Constructing typed, justified semantic links between ontologies is essential for enabling interoperability across heterogeneous and interdisciplinary knowledge domains. However, manually curating such links is difficult to scale....

<https://arxiv.org/abs/2610.01393>

47 CONF

Generalization Is Stability, Not Accuracy: Multi-Axis Evaluation of LLMs

ArXiv CS.AI 50%

arXiv:2610.01428v1 Announce Type: cross Abstract: Generalization in large language models (LLMs) is the ability to produce consistent and semantically stable outputs when the same input is expressed in different ways. Existing work typically evaluates generalization through...

<https://arxiv.org/abs/2610.01428>

48 CONF

MWOP: Modality-aware Width-wise Operation Pruning for Efficient MLLMs

ArXiv CS.AI 50%

arXiv:2610.01434v1 Announce Type: cross Abstract: Multimodal large language models (MLLMs) incur substantial inference costs when processing long visual-textual sequences. While existing operation compression methods exploit modality-level redundancy, they largely treat...

<https://arxiv.org/abs/2610.01434>

49 CONF

Jev in Medicine: A Benchmark Evaluation

ArXiv CS.AI 50%

arXiv:2609.34024v2 Announce Type: replace Abstract: Jev is a non-generative "System One" model that assigns probabilities to predefined answer options and cannot answer outside them. Its accuracy and calibration on medical question-answering and case-based diagnostic-reasoning...

<https://arxiv.org/abs/2609.34024>

50 CONF

No Model Required: Text Entropy Rate Filtering Mitigates Iterative Fine-Tuning Collapse

ArXiv CS.AI 50%

arXiv:2610.01493v1 Announce Type: cross Abstract: Iterative fine-tuning on synthetic data causes \emph{model collapse}: output diversity narrows as rare patterns are progressively lost, a signature most visible as phrase-level repetition. Existing mitigations either require...

<https://arxiv.org/abs/2610.01493>

51 CONF

False Floors: LLM Safety Routing Evaluations Break Under Distribution Shift

ArXiv CS.AI 50%

arXiv:2610.01535v1 Announce Type: cross Abstract: Safety routers send each request to one of several models and are judged against the best single model. A major routing benchmark picks that comparator on the evaluation data. In the benchmark's own setting this is harmless, but...

<https://arxiv.org/abs/2610.01535>

52 CONF

Reinforcement Learning to Accelerate Primal-Dual Hybrid Gradient for Linear Programming

ArXiv CS.AI 50%

arXiv:2610.01546v1 Announce Type: cross Abstract: Primal-dual hybrid gradient (PDHG) methods solve large-scale linear programs (LPs) using GPU-friendly matrix-vector products and projections, but their practical performance depends on coordinating algorithm parameters,...

<https://arxiv.org/abs/2610.01546>

53 CONF

Before It Fades: Reinforcing Temporal Representations at Inference Time in VideoLLMs

ArXiv CS.AI 50%

arXiv:2610.01595v1 Announce Type: cross Abstract: Video Large Language Models (VideoLLMs) receive frames in sequential order and interpret how visual content evolves along the temporal axis, yet temporal reasoning remains a persistent weakness across architectures. Reversing the...

<https://arxiv.org/abs/2610.01595>

54 CONF

Can LLMs Reliably Annotate Bioassay Metadata to Improve Data Readiness?

ArXiv CS.AI 50%

arXiv:2610.01616v1 Announce Type: cross Abstract: The emergence of foundation models for molecular property prediction requires a high degree of AI data readiness, including reliable metadata annotation. However, both public repositories and industrial screening databases suffer...

<https://arxiv.org/abs/2610.01616>

55 CONF

Exposing the Cost of Deep Learning Audio Development

ArXiv CS.AI 50%

arXiv:2610.01619v1 Announce Type: cross Abstract: The environmental impact of deep learning has attracted increasing attention over the past decade. Existing studies mainly focus on the energy and carbon emissions of model training and inference, while the whole development...

<https://arxiv.org/abs/2610.01619>

56 CONF

What Makes Something Hard(er)? Explaining Question Difficulty in Natural Language

ArXiv CS.AI 50%

arXiv:2610.01627v1 Announce Type: cross Abstract: Difficulty is one of the most fundamental properties of a question: it determines whether the question can meaningfully discriminate between models of differing ability. Although a variety of methods can now estimate or predict...

<https://arxiv.org/abs/2610.01627>

57 CONF

Not All Error Yields to Scale: Where Scaling Stops in Vision-Language Inference

ArXiv CS.AI 50%

arXiv:2610.01640v1 Announce Type: cross Abstract: Vision-language models (VLMs) face a fixed-budget trade-off between processing more visual information for fine-grained perception and using a larger language backbone for complex reasoning. Existing studies do not tell us which...

<https://arxiv.org/abs/2610.01640>

58 CONF

MCIR: A Feature Dependence-Aware Explainability Method with Reliability Guarantees

ArXiv CS.AI 50%

arXiv:2610.01641v1 Announce Type: cross Abstract: Modern machine-learning models often contain strongly dependent or redundant features, making feature attribution difficult because shared predictive information can be distributed across correlated predictors. Existing methods...

<https://arxiv.org/abs/2610.01641>

59 CONF

Architecture Without an Architect? Global Governance of Artificial Intelligence in a Divided World

ArXiv CS.AI 50%

arXiv:2610.01716v1 Announce Type: cross Abstract: Artificial intelligence presents an unusually difficult problem for global governance. The technology develops rapidly, crosses borders easily, and is shaped by actors whose resources and capabilities may rival those of states....

<https://arxiv.org/abs/2610.01716>

60 CONF

Detecting Inconsistencies in Model Specifications with LLM-as-Verifier Reasoning

ArXiv CS.AI 50%

arXiv:2610.01847v1 Announce Type: cross Abstract: Model specifications define how large language models (LLMs) should behave, guiding alignment training, inference-time behavior, and evaluation. Yet these specifications may themselves contain defects: two individually reasonable...

<https://arxiv.org/abs/2610.01847>

61 CONF

From Isolated Feature to Orbits: Discovering Music Concepts via Multi-SAE Alignment

ArXiv CS.AI 50%

arXiv:2610.01864v1 Announce Type: cross Abstract: How can we understand what a music foundation model has learned internally? Most interpretability approaches, such as probing and Sparse Autoencoders (SAEs), focus on identifying individual features with minimal...

<https://arxiv.org/abs/2610.01864>

62 CONF

Walking the Embedding Space: Datastore Extraction from Multimodal RAG

ArXiv CS.AI 50%

arXiv:2610.01871v1 Announce Type: cross Abstract: Multimodal Retrieval-Augmented Generation (MRAG) has emerged as a reliable and cost-effective technique of grounding the generative capabilities of Multimodal Large Language Models (MLLMs) into relevant, up-to-date, external...

<https://arxiv.org/abs/2610.01871>

63 CONF

From Network Intrusion Detection to Blockchain-Backed Endpoint Detection and Response: Mapping the Landscape of Decentralized Detection-and-Response Architectures

ArXiv CS.AI 50%

arXiv:2610.01872v1 Announce Type: cross Abstract: While the literature on blockchain-assisted intrusion detection and prevention systems (IDS/IPS) for Internet of Things (IoT) and Industrial Internet of Things (IIoT) networks is mature, existing systematic reviews suffer from...

<https://arxiv.org/abs/2610.01872>

64 CONF

Unsupervised Domain Adaptation for Enhanced Radiometer Image Precipitation Estimation using Conditional Flow Matching

ArXiv CS.AI 50%

arXiv:2610.01890v1 Announce Type: cross Abstract: Deep generative networks have recently achieved unprecedented performance in precise image and video editing using sophisticated textual prompts. However, the effectiveness of such models heavily depends on access to very large...

<https://arxiv.org/abs/2610.01890>

65 CONF

Talked Out of the Truth: Sycophancy in the Reasoning Chains of Multimodal Models

ArXiv CS.AI 50%

arXiv:2608.28623v2 Announce Type: replace-cross Abstract: Large multimodal reasoning models (LMRMs) are increasingly capable, largely through generating explicit chain-of-thought reasoning before answering, but in language models this often comes with sycophancy, the tendency to...

<https://arxiv.org/abs/2608.28623>

66 CONF

Asynchronous LLM Post-Training: Group-Mass Capping and Convergence Analysis

ArXiv CS.AI 50%

arXiv:2610.01896v1 Announce Type: cross Abstract: Asynchronous reinforcement learning (RL) improves the efficiency of large language model post-training but introduces stale rollouts generated by earlier policies. Theoretical understanding of how this staleness affects...

<https://arxiv.org/abs/2610.01896>

67 CONF

A rubric landscape for evaluating clinical reasoning in large language models: what exists, what is missing, and what needs to be combined

ArXiv CS.AI 50%

arXiv:2610.01938v1 Announce Type: cross Abstract: Exam-style accuracy does not establish whether large language models (LLMs) reason well over clinical records. We define clinical reasoning as integrating and updating evidence across time and sources to form, revise and justify...

<https://arxiv.org/abs/2610.01938>

68 CONF

Exploring Weaknesses of Generative Image Watermarks against Latent Frequency Masking

ArXiv CS.AI 50%

arXiv:2610.02010v1 Announce Type: cross Abstract: Invisible watermarking has become a central tool for tracing AI-generated images, but its robustness against adaptive removal attacks remains an open security question. We introduce Latent Frequency Masking, an attack that erases...

<https://arxiv.org/abs/2610.02010>

69 CONF

Where-OPD: Spatially Guided On-Policy Self-Distillation of MLLMs with Synthetic Scenes

ArXiv CS.AI 50%

arXiv:2610.02117v1 Announce Type: cross Abstract: On-policy self-distillation has recently emerged as an effective approach for improving language-model reasoning by supervising students with a frozen or EMA version of themselves that receives privileged information. Its...

<https://arxiv.org/abs/2610.02117>

70 CONF

Watch, Infer, Coordinate: Inferring Robot Partner Constraints for Zero-Shot Coordination

ArXiv CS.AI 50%

arXiv:2610.02170v1 Announce Type: cross Abstract: Robots operating in the physical world will increasingly need to coordinate with other robots, particularly in manipulation tasks where an object may be too large or heavy for a single robot to carry alone. Physical limitations...

<https://arxiv.org/abs/2610.02170>

71 CONF

Generative Cinematographer: Composing Camera and Object Motion in 3D

ArXiv CS.AI 50%

arXiv:2610.02180v1 Announce Type: cross Abstract: Current controllable video generation systems often rely on 2D motion trajectories or sparse drag signals for object motion. These controls are ambiguous because the same 2D trajectory can correspond to different 3D motions,...

<https://arxiv.org/abs/2610.02180>

72 CONF

SoftServe: A Scalable Quasi-Newton Method for Deep Learning

ArXiv CS.AI 50%

arXiv:2610.02182v1 Announce Type: cross Abstract: Quasi-Newton (QN) methods have long been among the most effective methods for large-scale unconstrained convex optimization. Two obstacles have limited their use in deep learning: non-convexity and enormous parameter sizes. We...

<https://arxiv.org/abs/2610.02182>

73 CONF

Higher-Order Molecular Grammars for Generative and Foundation Models in Chemistry

ArXiv CS.AI 50%

arXiv:2610.02186v1 Announce Type: cross Abstract: Molecular learning models are strongly shaped by their underlying representations. Yet standard sequential and graph formalisms struggle to explicitly encode higher-order topology, such as ring systems and recurring motifs....

<https://arxiv.org/abs/2610.02186>

74 **CONF****KaliBench: A Fine-Grained Benchmark for Cybersecurity Tool Use on Kali Linux with Runtime-Free Verifiable Rewards****ArXiv CS.AI** **50%**

arXiv:2610.02206v1 Announce Type: cross Abstract: LLMs are increasingly applied to cybersecurity workflows, where they are expected to translate analysts' intent into tool invocations. However, existing evaluations focus on knowledge-based assessments or end-to-end agentic...

<https://arxiv.org/abs/2610.02206>

75 **CONF****LLM-based Agentic Reasoning Frameworks: A Survey from Methods to Scenarios****ArXiv CS.AI** **50%**

arXiv:2508.17692v2 Announce Type: replace Abstract: Recent advances in LLM-based agents highlight the importance of their reasoning frameworks, which guide the problem-solving process in diverse ways. This survey introduces a unified formal language to systematically categorize...

<https://arxiv.org/abs/2508.17692>

76 **CONF****Moloch's Bargain: Emergent Misalignment When LLMs Compete for Audiences****ArXiv CS.AI** **50%**

arXiv:2510.06105v2 Announce Type: replace Abstract: Large language models (LLMs) are increasingly shaping how information is created and disseminated, from companies using them to craft persuasive advertisements, to election campaigns optimizing messaging to gain votes, to...

<https://arxiv.org/abs/2510.06105>

77 **CONF****Probing an Embodied LLM: When Higher Observation Fidelity Hurts Problem Solving****ArXiv CS.AI** **50%**

arXiv:2605.20072v2 Announce Type: replace Abstract: Large Language Models (LLMs) are increasingly proposed as cognitive components for robotic systems, yet their opaque decision processes make it difficult to explain success or failure in closed-loop embodied tasks. Following an...

<https://arxiv.org/abs/2605.20072>

78 **CONF****MINCE: Shrinking LLM Evaluation Datasets via Few-Model Monte Carlo Calibration****ArXiv CS.AI** **50%**

arXiv:2606.22826v2 Announce Type: replace Abstract: Evaluating LLMs across many model variants---quantized, fine-tuned, or deployment-specific---requires running large benchmarks repeatedly, a process that can take tens of hours per model on edge hardware such as NPUs. Existing...

<https://arxiv.org/abs/2606.22826>

79 **CONF****ESCROW: Guarded and Dual-Objective Continual Maintenance for Agents in Policy-Governed Enterprise Workflows****ArXiv CS.AI** **50%**

arXiv:2608.01772v2 Announce Type: replace Abstract: LLM agents increasingly run policy-bound enterprise workflows, where they must apply rules consistently and stay auditable. Deploying such an agent is the start of its long-term maintenance cycle: it must adapt to a stream of...

<https://arxiv.org/abs/2608.01772>

80 CONF

Scaling Clinical Judgment to Evaluate Medical AI

ArXiv CS.AI 50%

arXiv:2609.12822v2 Announce Type: replace Abstract: Blinded physician evaluation has been considered by many to be the gold standard for assessing clinical reasoning in large language models (LLMs). This is difficult to scale; thus, prior studies typically rely on small...

<https://arxiv.org/abs/2609.12822>

81 CONF

Not Too Hard, Not Too Easy: Learning from Intermediate States for LLM Structured Reasoning

ArXiv CS.AI 50%

arXiv:2609.33149v2 Announce Type: replace Abstract: A common principle of effective learning is to practice material that is neither already mastered nor too difficult to permit progress. We ask how to apply this principle to structured reasoning tasks such as Sudoku and maze...

<https://arxiv.org/abs/2609.33149>

82 CONF

MERID: Multimodal Exploration via Recursive Self-Improvement Agents for Major Depression Analysis

ArXiv CS.AI 50%

arXiv:2609.36235v2 Announce Type: replace Abstract: Major depressive disorder (MDD) severely impacts daily activities and quality of life. Detecting MDD involves multimodal data, such as interview recordings and sensor measurements. This is particularly challenging, as these...

<https://arxiv.org/abs/2609.36235>

83 CONF

CANTO: CAD-Native Transformer Operators for AI-Aided Engineering

ArXiv CS.AI 50%

arXiv:2609.36806v2 Announce Type: replace Abstract: Modern engineering systems, from automobiles to aircraft, are designed by using precise, continuous parametric computer-aided design (CAD) models. Evaluating design changes through numerical simulation requires meshing the...

<https://arxiv.org/abs/2609.36806>

84 CONF

Easier Said Than Done: Unpacking Intent-Behavior Gap in Jailbreaking LLM-Based Robots

ArXiv CS.AI 50%

arXiv:2412.16633v5 Announce Type: replace-cross Abstract: LLM-based robots use Large Language Models (LLMs) as planners to translate natural language instructions into policies such as `grasp()`, `move_to()`, and `open_gripper()`. Jailbreak attacks on these robots extend the threat...

<https://arxiv.org/abs/2412.16633>

85 CONF

UniGuardian: A Unified Defense for Detecting Prompt Injection, Backdoor Attacks and Adversarial Attacks in Large Language Models

ArXiv CS.AI 50%

arXiv:2502.13141v2 Announce Type: replace-cross Abstract: Large Language Models (LLMs) are vulnerable to attacks like prompt injection, backdoor attacks, and adversarial attacks, which manipulate prompts or models to generate harmful outputs. In this paper, departing from...

<https://arxiv.org/abs/2502.13141>

86 CONF

WaLLM -- Understanding Use and Engagement with a General-Purpose LLM on WhatsApp

ArXiv CS.AI 50%

arXiv:2505.08894v2 Announce Type: replace-cross Abstract: Large language model (LLM) chatbots are increasingly reaching users through messaging platforms (e.g. WhatsApp). However, these systems remain largely proprietary and opaque, while academic research has focused on narrow,...

<https://arxiv.org/abs/2505.08894>

87 CONF

What Drives Compositional Generalization in Visual Generative Models? The Importance of Continuous Training Objectives

ArXiv CS.AI 50%

arXiv:2510.03075v4 Announce Type: replace-cross Abstract: Compositional generalization, the ability to generate novel combinations of known concepts, is a key ingredient for visual generative models. Yet, not all mechanisms that enable or inhibit it are fully understood. In this...

<https://arxiv.org/abs/2510.03075>

88 CONF

Iterative Topic Taxonomy Induction with LLMs: A Case Study of Electoral Advertising

ArXiv CS.AI 50%

arXiv:2510.15125v3 Announce Type: replace-cross Abstract: Social media platforms play a pivotal role in shaping political discourse, but the scale and rapid evolution of online content make systematic analysis difficult. We introduce an end-to-end framework for inducing an...

<https://arxiv.org/abs/2510.15125>

89 CONF

Constant-Time Planning for Chaining Collision-free Motion to Manipulation Behaviors

ArXiv CS.AI 50%

arXiv:2512.00939v3 Announce Type: replace-cross Abstract: Recent progress in contact-rich robotic manipulation has been striking, yet most deployed systems remain confined to simple, scripted routines. One of the barriers is the lack of motion planning algorithms that can...

<https://arxiv.org/abs/2512.00939>

90 CONF

In Vino Veritas and Vulnerabilities: Examining LLM Safety via Drunk Language Inducement

ArXiv CS.AI 50%

arXiv:2601.22169v2 Announce Type: replace-cross Abstract: Humans are susceptible to undesirable behaviours and privacy leaks under the influence of alcohol. This paper investigates drunk language, i.e., text written under the influence of alcohol, as a driver for safety failures...

<https://arxiv.org/abs/2601.22169>

91 CONF

TensorCommitments: A Lightweight Verifiable Inference for Language Models

ArXiv CS.AI 50%

arXiv:2602.12630v2 Announce Type: replace-cross Abstract: Most large language models (LLMs) run on external clouds: users send a prompt, pay for inference, and must trust that the remote GPU executes the LLM without any adversarial tampering. We critically ask how to achieve...

<https://arxiv.org/abs/2602.12630>

92 CONF

Capabilities Ain't All You Need: Measuring Propensities in AI

ArXiv CS.AI 50%

arXiv:2602.18182v5 Announce Type: replace-cross Abstract: AI evaluation has primarily focused on measuring capabilities, with formal approaches inspired from Item Response Theory (IRT) being increasingly applied. Yet propensities - the tendencies of models to exhibit particular...

<https://arxiv.org/abs/2602.18182>

93 CONF

Rewind-IL: Online Failure Detection and State Respawning for Imitation Learning

ArXiv CS.AI 50%

arXiv:2604.16683v2 Announce Type: replace-cross Abstract: Imitation learning has enabled robots to acquire complex visuomotor manipulation skills from demonstrations, but deployment failures remain a major obstacle, especially for long-horizon action-chunked policies. Once...

<https://arxiv.org/abs/2604.16683>

94 CONF

Verbal tics in frontier language models: A critical review of current releases, research evidence, and public discussion

ArXiv CS.AI 50%

arXiv:2604.19139v4 Announce Type: replace-cross Abstract: Repeated praise, canned reassurance, familiar contrasts, and conspicuous vocabulary are recurring subjects in discussions of large language models. Their interpretation depends on context: a conventional phrase may be...

<https://arxiv.org/abs/2604.19139>

95 CONF

Vision-language models for chest radiography do not always need the image

ArXiv CS.AI 50%

arXiv:2606.17710v3 Announce Type: replace-cross Abstract: Vision-language models that answer questions about chest radiographs are evaluated by their accuracy on labels derived from radiology reports. High benchmark accuracy is often interpreted as evidence that the model uses...

<https://arxiv.org/abs/2606.17710>

96 CONF

Aligned but Not Partner-Specific: How Multimodal LLM Agents Succeed in Reference Games Without Forming Conceptual Pacts

ArXiv CS.AI 50%

arXiv:2606.08081v2 Announce Type: replace-cross Abstract: Repeated reference games test whether interlocutors replace their initially long descriptions with shorter, partner-specific expressions grounded in shared interaction history; that is, with conceptual pacts. Prior work...

<https://arxiv.org/abs/2606.08081>

97 CONF

Pepti-drift: Scalable Safe-Active Peptide Generation Without Inference-Time Guidance

ArXiv CS.AI 50%

arXiv:2606.27824v3 Announce Type: replace-cross Abstract: Therapeutic peptides are a promising drug modality, but their generation must satisfy multiple therapeutic constraints. We introduce BindSafe-PepBench, a fixed-budget benchmark that jointly evaluates target binding and...

<https://arxiv.org/abs/2606.27824>

98 **CONF****Evaluating Neural Decompilation of Dart AOT Binaries: Fine-Tuning, Metric Validity, Specification Leakage, and Reliability**

ArXiv CS.AI 50%

arXiv:2607.06125v2 Announce Type: replace-cross Abstract: We present an execution-based evaluation of neural decompilation for Dart ahead-of-time binaries and an audit of what its scores measure. Across six archived adapter-baseline comparisons, paired tests of pass@k at k = 1,...

<https://arxiv.org/abs/2607.06125>

99 **CONF****Generative Interpretability via Scalable Neuro-Symbolic Models**

ArXiv CS.AI 50%

arXiv:2609.13529v2 Announce Type: replace-cross Abstract: As the use of Large Language Models moves from chatbots into agentic systems, where outputs become actions with irreversible consequences on reality, the existing paradigm on AI Interpretability research, post-hoc...

<https://arxiv.org/abs/2609.13529>

00 **CONF****DOHF: Online Diffusion Fine-tuning with Doob's $\$h\$$ -transform Guidance**

ArXiv CS.AI 50%

arXiv:2609.31882v2 Announce Type: replace-cross Abstract: Reward-based diffusion fine-tuning faces practical challenges when desirable outcomes are rare or conditioning corrections are costly to estimate. In this work, we propose Diffusion Online $\$h\$$ -guidance Fine-tuning (DOHF),...

<https://arxiv.org/abs/2609.31882>

01 **CONF****What Does a Skill Actually Do? Estimands and Evaluation Validity for Tool and Skill Use in LLM Agents: A Critical Review**

ArXiv CS.AI 50%

arXiv:2609.33153v2 Announce Type: replace-cross Abstract: Reported improvements from tools and reusable skills in large language model agents refer to different comparisons. This critical narrative review examines what these evaluations estimate and which conclusions their...

<https://arxiv.org/abs/2609.33153>

02 **CONF****Alignment Forecasting: Predicting Misalignment From Training Data**

ArXiv CS.AI 50%

arXiv:2609.35805v2 Announce Type: replace-cross Abstract: Training a language model on data with a narrow flaw can sometimes make the model broadly misaligned. Inspecting the data at face value often does not settle whether it will emerge, and today it is caught only after...

<https://arxiv.org/abs/2609.35805>

03 **CONF****Alignment via Training Against Probes Without Losing Monitorability**

ArXiv CS.AI 50%

arXiv:2609.38645v2 Announce Type: replace-cross Abstract: Models are usually aligned based on their observed outputs, using demonstrations, preference data, or reward signals. These objectives reward responses that look aligned. More capable models may learn to satisfy them...

<https://arxiv.org/abs/2609.38645>