



Pipeline

Scrape HackerNews (Show/Top/New), Reddit (r/artificial, r/MachineLearning, r/singularity, r/LocalLLaMA), TechCrunch AI, VentureBeat AI, MIT Tech Review, The Verge AI, Wired AI, ArXiv CS.AI

Filter 40+ AI keyword patterns — LLMs, alignment, safety, research, tooling

Dedup Title-similarity (45% threshold) merges cross-source variants

Score Authenticity = 0.5 base + 0.3 per extra source + diversity + HN score bonuses

Rank Sorted by authenticity. Below 0.25 threshold are dropped.

Authenticity Tiers

- VERIFIED 0.80 - 1.00** 3+ independent sources including media
- CONFIRMED 0.50 - 0.79** 2 sources or 1 high-quality media
- EMERGING 0.25 - 0.49** Single-source; passes AI keyword filter

Source Breakdown

Source	Articles	Category
ArXiv CS.AI	124	Media
HackerNews	5	Community
MIT Tech Review	3	Media
The Verge AI	3	Media
TechCrunch AI	2	Media

VERIFIED INTELLIGENCE — 1 story

VERI

Securing Computer-Use Agents Against Branch Steering Attacks

ArXiv CS.AI

+1 source

80%

arXiv:2610.03089v1 Announce Type: cross Abstract: Modern Computer Use Agents (CUAs) directly interact with graphical user interfaces and execute third-party web tools, exposing them to indirect prompt injection across every rendered page and tool response. While the Dual-LLM...

<https://arxiv.org/abs/2610.03089>

CONFIRMED SIGNALS — 135 stories

CONF

Show HN: Dataviz, ranked daily from GitHub, NPM, PyPI and CRAN

HackerNews

50%

HN 11

No summary — see link for full article.

<https://awesomedataviz.com/>

CONF

Show HN: Four LLMs play 20 questions for real money, live

HackerNews

50%

HN 4

Hi HN, I wanted to explore how AI models behave when they have their own money. So I gave 4 frontier models their own USDC and had them compete in games of 20 questions. Anyone can watch live. Each round, one model takes a turn picking the secret word, and the others interrogate...

<https://aquarium.money/twenty>

CONF

Show HN: Gutsy, subsecond decision model inference on CPU (Jev Style), no API

HackerNews

50%

HN 3

No summary — see link for full article.

<https://github.com/kouhxp/gutsy>

CONF

Show HN: AstroHelm - Use your phone camera to aim a telescope or telephoto lens

HackerNews

50%

HN 2

Hello HN! As an amateur astronomer living in New York City, finding objects in the night sky has always been a huge challenge. I thought about using the smartphone's gyroscope and magnetometer, but the accuracy left a lot to be desired. Then I learned about plate solving: given...

<https://astrohelm.app/>

CONF

Show HN: Solvi 1.0 - an LLM proposes, checks decide, every decision replays

HackerNews

50%

HN 2

No summary — see link for full article.

<https://github.com/solvi-ai/solvi>

CONF

Google froze its open source bug bounty program due to a 'significant rise' in AI submissions

TechCrunch AI

50%

AI slop seems to be overwhelming bug bounty programs.

<https://techcrunch.com/2026/10/04/google-froze-its-open-source-bug-bounty-program-due-t...>

CONF

It's not AI anymore, it's 'super intelligence' (according to the White House)

TechCrunch AI 50%

This week, the White House got nearly every major tech CEO in one room — Zuckerberg, Bezos, Musk, and Anthropic's Dario Amodei among them — to sign an AI safety pledge that President Donald Trump called "morally binding." Trump also signed an executive order officially...

<https://techcrunch.com/video/its-not-ai-anymore-its-super-intelligence-according-to-the...>

CONF

People really hate AI, so why can't they get enough?

MIT Tech Review 50%

Over the summer I talked to the CEO of Springboards, a startup building an LLM that's designed to come up with a wider variety of responses than its mainstream rivals do. At the start of the call, he said something that's been stuck in my head since: "We often say that we're a..."

<https://www.technologyreview.com/2026/10/05/1145682/people-really-hate-ai-so-why-cant-t...>

0 CONF

EmTech Future 2026: When AI Meets Everything

MIT Tech Review 50%

Yossi Matias, Vice President & Head of Google Research, explores how AI is beginning to reshape biology, infrastructure, manufacturing, and science, and why its greatest impact may come when it intersects with other fields. Step inside the newsroom with our MIT Technology Review...

<https://www.technologyreview.com/2026/10/05/1145413/emtech-future-2026-when-ai-meets-ev...>

1 CONF

Our minds aren't equipped to handle AI

The Verge AI 50%

Norbert Wiener, godfather of cybernetics, once said, "The thought of every age is reflected in its technique." For the past century, our thought has been reflected in our computers, including by those in the AI industry. Google's Demis Hassabis calls the brain "a biological..."

<https://www.theverge.com/ai-artificial-intelligence/1003794/ai-education-computational-...>

2 CONF

NJ's former Lt Gov is using AI to say he's innocent of sexual harassment

The Verge AI 50%

New Jersey's lieutenant governor Dale Caldwell was forced to resign on September 25th after an investigation found he had sexually harassed a staffer and repeatedly violated ethics rules. The now-former Lt. governor has been making the media rounds trying to clear his name. But...

<https://www.theverge.com/ai-artificial-intelligence/1004549/well-if-ai-said-it-it-must-...>

3 CONF

An AI couldn't beat humans at StarCraft, so it decided to cheat

The Verge AI 50%

StarSkirmish pits AI-made StarCraft-playing bots against one another, as well as against human-made bots. OpenAI's GPT-6 Astra and Claude Opus 5.5 were essentially tied as the best-performing AI-made bots, but they couldn't top Stardust, the top-rated human-made bot. On Friday,...

<https://www.theverge.com/ai-artificial-intelligence/1004543/openai-gpt-cheat-starcraft>

4 CONF

Multimodal Representation Learning Conditioned on Semantic Relations

ArXiv CS.AI 50%

arXiv:2508.17497v3 Announce Type: replace-cross Abstract: Multimodal representation learning has been largely driven by contrastive models such as CLIP, which learn a shared embedding space by aligning paired image-text samples. While effective for general-purpose representation...

<https://arxiv.org/abs/2508.17497>

5 **CONF****Fast Models, Slow Evidence: A Paired and Self-Audited Evaluation of System-1 Decision Models for LLM Agent Harnesses****ArXiv CS.AI** **50%**

arXiv:2610.02267v1 Announce Type: new Abstract: Agent harnesses make many small, typed decisions per task: which model to call, which tool to use, whether retrieved text is relevant, whether an input carries an injection. System-1 decision models answer such questions in a...

<https://arxiv.org/abs/2610.02267>

6 **CONF****The AI Risk Observatory: What Can We Learn from AI Disclosures in Annual Reports About Societal Resilience?****ArXiv CS.AI** **50%**

arXiv:2610.02281v1 Announce Type: new Abstract: Societal resilience research relies on access to useful and actionable data, which motivates our main research question: Can annual reports, processed at scale with LLMs, provide a useful signal about how companies disclose their...

<https://arxiv.org/abs/2610.02281>

7 **CONF****LEAP: Learning Efficient Action Proposals For LLM Agents****ArXiv CS.AI** **50%**

arXiv:2610.02670v1 Announce Type: cross Abstract: LLM agents are known to be slow in rollouts. An agent completes a task one step at a time. At each step, it reasons and then chooses an action to execute. The next step and action cannot start until the previous one has finished....

<https://arxiv.org/abs/2610.02670>

8 **CONF****Traversing the Satisfaction-Diversity Frontier in Text-to-Image Diffusion****ArXiv CS.AI** **50%**

arXiv:2610.02372v1 Announce Type: new Abstract: Text-to-image generation enables users to explore several images generated from the same prompt. For these generated images to be useful, each one must reflect the user's preferences, measured by a learned reward, and differ...

<https://arxiv.org/abs/2610.02372>

9 **CONF****THPL: A Vision-to-Language Decision Support Framework for Rainbow Trout Feeding Management in RAS****ArXiv CS.AI** **50%**

arXiv:2610.02378v1 Announce Type: new Abstract: In Recirculating Aquaculture Systems (RAS), precision feeding is critical for minimizing costs and improving fish welfare. However, existing methods lack cognitive alignment between fish behaviors and management knowledge, impeding...

<https://arxiv.org/abs/2610.02378>

0 **CONF****When Terminal-Agent Training Stalls: Demystifying Data Generation and Verification Challenge****ArXiv CS.AI** **50%**

arXiv:2610.02405v1 Announce Type: new Abstract: Using a frontier model like Claude Opus as a meta-agent to generate terminal tasks and verifiers for RL training is increasingly common. Yet a runnable Docker image and executable test suite do not guarantee a faithful end-to-end...

<https://arxiv.org/abs/2610.02405>

1 **CONF****Reinforcement Learning Techniques for the Optimization of Target Polarization in Nuclear Physics Scattering Experiments****ArXiv CS.AI** **50%**

arXiv:2610.02452v1 Announce Type: new Abstract: The operation of dynamically polarized targets in nuclear physics experiments relies on continuous tuning of the microwave frequency to compensate for radiation damage and evolving material properties, a task that is traditionally...

<https://arxiv.org/abs/2610.02452>

2 **CONF****Multi-Fidelity Policy Gradients Stabilize Data-Scarce Reinforcement Learning****ArXiv CS.AI** **50%**

arXiv:2610.02505v1 Announce Type: cross Abstract: Policy gradient methods for on-policy reinforcement learning (RL) can become unstable when expensive, scarce target-domain data yield noisy gradient estimates. We address this challenge by complementing limited high-fidelity (HF)...

<https://arxiv.org/abs/2610.02505>

3 **CONF****Right Order, Wrong Scale: Auditing LLM Judges for Occupational AI Measurement****ArXiv CS.AI** **50%**

arXiv:2610.02492v1 Announce Type: new Abstract: LLM judges are increasingly used to assess whether AI outputs meet workplace requirements, but agreement on response rankings does not establish agreement on acceptance rates or occupational aggregates. We introduce O*NET-BENCH, an...

<https://arxiv.org/abs/2610.02492>

4 **CONF****HXAI: Hierarchical Privacy-Preserving Explainable AI in Distributed Energy Systems****ArXiv CS.AI** **50%**

arXiv:2610.02504v1 Announce Type: new Abstract: Balancing electricity demand and supply is increasingly difficult due to the inherent intermittency of renewable power generation and the stochastic power consumption. Grid operators require fine-grained, decision-relevant insights...

<https://arxiv.org/abs/2610.02504>

5 **CONF****On-Premises Multi-Course RAG Tutoring for Business Education: Hardware-Software Trade-offs in a Campus AI Tutor****ArXiv CS.AI** **50%**

arXiv:2610.02510v1 Announce Type: new Abstract: Campus AI tutors based on retrieval-augmented generation (RAG) must ground answers in assigned course materials while keeping textbooks and student dialogue on institutional infrastructure. We present CourseChat, an on-premises,...

<https://arxiv.org/abs/2610.02510>

6 **CONF****Beyond Trained Models: Compiling GNNs for a Sound Explainer Benchmark****ArXiv CS.AI** **50%**

arXiv:2610.03526v1 Announce Type: cross Abstract: Explainers for Graph Neural Networks (GNNs) are commonly evaluated by their plausibility, i.e., how well their explanations recover a predefined ground truth, such as a motif planted in the data. This protocol implicitly assumes...

<https://arxiv.org/abs/2610.03526>

7 CONF

How to Have a Sensitive Debate: An Instance-Optimal Protocol for AI Debate

ArXiv CS.AI 50%

arXiv:2610.02557v1 Announce Type: new Abstract: As powerful AI systems reach and sometimes surpass the abilities of human experts across a range of cognitively demanding tasks, the problem of accurate oversight and supervision of these systems has become increasingly urgent. One...

<https://arxiv.org/abs/2610.02557>

8 CONF

Suan: Rectifying Direct Preference Safety Alignment in Large Language Models

ArXiv CS.AI 50%

arXiv:2609.08634v2 Announce Type: replace-cross Abstract: Integrating robust safety guardrails into Large Language Models (LLMs) is essential for delivering helpful yet harmless responses. While proprietary systems exhibit reliable safety controls, their underlying methodologies...

<https://arxiv.org/abs/2609.08634>

9 CONF

EditHero: A Benchmark for Long-Horizon Part-Level 3D Editing and Vibe Modeling

ArXiv CS.AI 50%

arXiv:2610.02298v1 Announce Type: cross Abstract: 3D editing methods are usually tested on a single edit, yet an asset is built through a long sequence of revisions, each of which must implement the requested change while leaving everything else unchanged. We introduce EditHero,...

<https://arxiv.org/abs/2610.02298>

10 CONF

Time Series Forecasting Benchmarks Need Scenario-Grounded Stress Testing

ArXiv CS.AI 50%

arXiv:2610.02608v1 Announce Type: new Abstract: Time series forecasting (TSF) increasingly drives decisions in transportation, energy, finance, healthcare, and infrastructure, yet current evaluation remains overly narrow: standard benchmarks reward low held-out error, while...

<https://arxiv.org/abs/2610.02608>

11 CONF

Coherence-Driven Belief Formation and Population Dynamics of Contagion in LLM Agents

ArXiv CS.AI 50%

arXiv:2610.02654v1 Announce Type: new Abstract: Models of social contagion usually assume how individuals adopt beliefs and derive population behavior from it. We instead empirically measure belief adoption in language model agents, quantifying the probability an agent adopts a...

<https://arxiv.org/abs/2610.02654>

12 CONF

Lost in the Request: How Communication Variation Disrupts Retrieval and Action in Email Agents

ArXiv CS.AI 50%

arXiv:2610.02627v1 Announce Type: new Abstract: An email assistant should not complete less work simply because a user phrases the same request differently. Yet most benchmarks test each task with only one canonical request, leaving this form of robustness largely unmeasured. We...

<https://arxiv.org/abs/2610.02627>

3 CONF

Designing the Future of User Feedback for Generative AI

ArXiv CS.AI 50%

arXiv:2610.02631v1 Announce Type: new Abstract: Post-deployment feedback from users can be a cost-effective, scalable, and representative means to monitor and improve generative AI systems and features. When implemented effectively, giving such feedback can increase users'...

<https://arxiv.org/abs/2610.02631>

4 CONF

Batched Speech Decisions Without Decoding: Single-Token Supervision Lets a Frozen LLM Hear Beyond the Transcript

ArXiv CS.AI 50%

arXiv:2610.02638v1 Announce Type: new Abstract: Full-duplex voice agents make many small, closed decisions, which current systems answer by slow autoregressive decoding. We propose Duplexjev, which feeds ASR-encoder hidden states through a small connector into a frozen LLM and...

<https://arxiv.org/abs/2610.02638>

5 CONF

A GHOST in Long-Horizon Agents: Governance Hazard from Overlooked Safety Constraints across Turns

ArXiv CS.AI 50%

arXiv:2610.02664v1 Announce Type: new Abstract: Long-horizon agents are now playing an increasingly significant role in assisting humans with complex problem-solving. However, it is exactly their extended interaction history that introduces an underexplored execution-safety...

<https://arxiv.org/abs/2610.02664>

6 CONF

DuetMoE: Coupling Inter- and Intra-Subgroup Robustness for Fair Medical Image Analysis

ArXiv CS.AI 50%

arXiv:2605.10521v2 Announce Type: replace-cross Abstract: As medical AI expands across diverse healthcare settings worldwide, equitable performance across patient populations is becoming essential to trustworthy clinical use. Fairness in medical image analysis is often evaluated...

<https://arxiv.org/abs/2605.10521>

7 CONF

SoftGene: Protein Language Model-Enhanced Soft Prompting for Interpretable Gene Set Annotation

ArXiv CS.AI 50%

arXiv:2610.03029v1 Announce Type: new Abstract: Gene set analysis is a cornerstone of functional genomics, yet it remains labor-intensive and heavily dependent on manual curation and expert biological interpretation. While Large Language Models (LLMs) have emerged as powerful...

<https://arxiv.org/abs/2610.03029>

8 CONF

On the Chain-of-Thought Monitorability of Looped Language Models

ArXiv CS.AI 50%

arXiv:2610.02741v1 Announce Type: new Abstract: Chain-of-thought (CoT) monitoring provides a promising approach for detecting undesirable model behavior. Looped language models (LoopLMs) repeatedly apply shared transformer layers, increasing effective computational depth and...

<https://arxiv.org/abs/2610.02741>

9 CONF

Dynamic LLM Routers are Often Misguided

ArXiv CS.AI 50%

arXiv:2610.02762v1 Announce Type: new Abstract: Dynamic LLM routers promise to cut inference costs by sending each query to the cheapest model that can answer it correctly. We analyze six commercial routers across 14 settings on a diverse benchmark spanning eight task...

<https://arxiv.org/abs/2610.02762>

0 CONF

PAPER2LLM++: Continual Self-Evolution of LLMs from Research Papers

ArXiv CS.AI 50%

arXiv:2610.02793v1 Announce Type: new Abstract: Research on LLMs continually uncovers model limitations, their causes, and potential solutions. Yet these human discoveries remain largely disconnected from model evolution: an LLM does not automatically learn from new research...

<https://arxiv.org/abs/2610.02793>

1 CONF

MaskCoFT: Masked Co-Adaptive Fine-Tuning for Memory-Efficient MoE Inference

ArXiv CS.AI 50%

arXiv:2609.34077v2 Announce Type: replace-cross Abstract: Mixture-of-experts (MoE) language models often exceed the memory of a single GPU. Expert offloading keeps most experts in host memory and loads them on demand, so decoding speed depends on how many experts each token must...

<https://arxiv.org/abs/2609.34077>

2 CONF

Robust Adversarial Quantification via Conflict-Aware Evidential Deep Learning

ArXiv CS.AI 50%

arXiv:2506.05937v3 Announce Type: replace-cross Abstract: Reliability of deep learning models is critical for deployment in high-stakes applications, where out-of-distribution or adversarial inputs may lead to detrimental outcomes. Evidential Deep Learning, an efficient paradigm...

<https://arxiv.org/abs/2506.05937>

3 CONF

MLCommons Jailbreak Benchmark v1.0

ArXiv CS.AI 50%

arXiv:2610.02827v1 Announce Type: new Abstract: Modern AI systems are designed to refuse hazardous requests. A jailbreak is a prompt crafted to bypass those safeguards and elicit outputs that the system would normally refuse to provide. The MLCommons Jailbreak Benchmark v1.0...

<https://arxiv.org/abs/2610.02827>

4 CONF

FSPO: Policy-Consistent Risk and Pareto-Feasible Control for Budgeted LLM RL Post-Training

ArXiv CS.AI 50%

arXiv:2610.02828v1 Announce Type: new Abstract: Adaptive LLM reinforcement-learning post-training changes multiple training actuators online, including rollout temperature, group size, clipping, KL regularization, verifier allocation, and update budget. Three coupled issues...

<https://arxiv.org/abs/2610.02828>

5 **CONF****Bounded Reachability & Jailbreak Detection via Contraction-Constrained State Space Models**

ArXiv CS.AI 50%

arXiv:2610.02853v1 Announce Type: new Abstract: Safety heads are lightweight classifiers attached to pretrained language models for flagging harmful inputs before generation. Their empirical detection performance has been studied, but their formal robustness properties remain...

<https://arxiv.org/abs/2610.02853>
6 **CONF****Beyond Predefined Sinks: Security-Aware Dependency Analysis for LLM Agents**

ArXiv CS.AI 50%

arXiv:2610.03014v1 Announce Type: cross Abstract: Large language model (LLM)-based agents increasingly connect model-generated decisions to security-sensitive software capabilities such as command execution, filesystem access, network communication, browser control, and external...

<https://arxiv.org/abs/2610.03014>
7 **CONF****LUMOS: Tracing Parametric Knowledge from Training Data to Behavioral Outputs in LLMs**

ArXiv CS.AI 50%

arXiv:2610.02902v1 Announce Type: new Abstract: Current analyses of LLMs' parametric knowledge are largely output-centric, drawing conclusions about what a model knows without verifying what it was actually trained on. This leaves fundamental questions, such as whether a correct...

<https://arxiv.org/abs/2610.02902>
8 **CONF****Frequency Is Not Sensitivity Identifying Safety-Sensitive Experts in Sparse MoE LLM**

ArXiv CS.AI 50%

arXiv:2610.02910v1 Announce Type: new Abstract: Suppressing a small set of routed experts can weaken the safety behavior of a sparse Mixture-of-Experts (MoE) language model without retraining. Which experts to suppress is therefore a security question, and the usual answer is...

<https://arxiv.org/abs/2610.02910>
9 **CONF****Reasoning with Evidence, Not Merely Rationales: Verifiable Preference Proofs for LLM-Based Recommendation**

ArXiv CS.AI 50%

arXiv:2610.02968v1 Announce Type: new Abstract: Large language models (LLMs) can infer user preferences from interaction histories and reviews, yet the rationales they generate may not reflect the information actually used for recommendation. A preference claim may be weakly...

<https://arxiv.org/abs/2610.02968>
0 **CONF****Safeguarding Mutual Correction in Source-Free Domain Adaptation via Cut Statistics**

ArXiv CS.AI 50%

arXiv:2610.02981v1 Announce Type: new Abstract: Source-Free Domain Adaptation (SFDA) aims to adapt a source-pretrained model to an unlabeled target domain without access to the original source domain. While early single-model approaches rely on self-refinement, they are...

<https://arxiv.org/abs/2610.02981>

1 **CONF****Benchmarking Candidate Coverage in Typed Decision Models**

ArXiv CS.AI 50%

arXiv:2610.03387v1 Announce Type: new Abstract: Typed decision models return choices or distributions over answer options supplied at request time. Accuracy with complete options does not establish whether a model recognizes that a reference answer is missing or avoids rejecting...

<https://arxiv.org/abs/2610.03387>

2 **CONF****TPBench: A Turning-Point Benchmark for Dialogue Compression**

ArXiv CS.AI 50%

arXiv:2610.02736v1 Announce Type: cross Abstract: A compressor can keep the facts of a dialogue and still drop the turn that changed them. A user corrects a price, reverses a choice, or adds a constraint. We call this failure turning-point eviction. One overall retention score...

<https://arxiv.org/abs/2610.02736>

3 **CONF****When Numbers Start Talking: Numerical Signalling and Strategic Behaviour Among LLMs**

ArXiv CS.AI 50%

arXiv:2610.03033v1 Announce Type: new Abstract: Large language model (LLM)-based agents increasingly operate in multi-agent systems (MAS) characterised by strategic interaction. However, little is known about whether, and to what extent, different types of messages affect the...

<https://arxiv.org/abs/2610.03033>

4 **CONF****Peer Influence across Heterogeneous AI Models**

ArXiv CS.AI 50%

arXiv:2610.03095v1 Announce Type: new Abstract: When two AI agents disagree, who persuades whom? As multi-agent systems increasingly combine language models of different families and sizes, the answer can determine which judgments survive interaction. Measuring persuasion as the...

<https://arxiv.org/abs/2610.03095>

5 **CONF****JOVE: Joint Execution and Verification for Resource-Aware LLM Task Graphs**

ArXiv CS.AI 50%

arXiv:2610.03296v1 Announce Type: new Abstract: Complex reasoning queries can be decomposed into directed acyclic task graphs and distributed across heterogeneous LLMs, reducing latency through parallelism and enabling smaller models to solve complex tasks. In practice, however,...

<https://arxiv.org/abs/2610.03296>

6 **CONF****Multi-Task Evolution for Zero-Shot Cross-Problem Generalization using LLMs**

ArXiv CS.AI 50%

arXiv:2610.03316v1 Announce Type: new Abstract: Designing effective heuristics for diverse combinatorial optimization problems requires substantial expertise and repeated search. Large language models (LLMs) automate heuristic generation and refinement, but heuristic search...

<https://arxiv.org/abs/2610.03316>

7 **CONF**

Refinement Buys Intelligibility, Search Buys Identity: What Test-Time Compute Buys in Masked-Diffusion TTS

ArXiv CS.AI 50%

arXiv:2610.03320v1 Announce Type: new Abstract: Diffusion language models for text-to-speech combine two forms of computation: model depth (parameters) and refinement steps (inference budget). We ask whether they scale equally across capabilities. We train 15 masked-diffusion...

<https://arxiv.org/abs/2610.03320>

8 **CONF**

Rethinking What to Cache in Few-Step Diffusion Transformers: Solver-Aware Target Selection

ArXiv CS.AI 50%

arXiv:2610.03577v1 Announce Type: cross Abstract: Diffusion Transformers (DiTs) can generate high-quality images and videos, but generating each sample requires multiple costly DiT forward passes. Two common ways to accelerate DiT sampling are step distillation, which reduces...

<https://arxiv.org/abs/2610.03577>

9 **CONF**

Geometry Meets Physics: Data-Efficient Pre-Training for Unstructured Neural PDE Solvers

ArXiv CS.AI 50%

arXiv:2610.03363v1 Announce Type: new Abstract: Neural surrogate models for Partial Differential Equations (PDEs) on unstructured 3D geometries are often limited by poor generalization and the high cost of generating large-scale training datasets. Consequently, pre-training on...

<https://arxiv.org/abs/2610.03363>

0 **CONF**

CVE2AP: Automated Generation of PDDL-Encoded Attack Paths via Large Language Models

ArXiv CS.AI 50%

arXiv:2610.03383v1 Announce Type: new Abstract: Attack Path (AP) modeling is fundamental to cybersecurity analysis, where the Planning Domain Definition Language (PDDL) has been widely adopted to encode APs into formal and machine-verifiable representations for automated...

<https://arxiv.org/abs/2610.03383>

1 **CONF**

Becoming Suspicious Across Borders: Algorithmic Extraterritoriality and AI-Driven Financial Surveillance

ArXiv CS.AI 50%

arXiv:2610.03425v1 Announce Type: new Abstract: Suspicion is an important, yet elusive concept in anti-money laundering and counter-terrorist financing (AML/CFT), which allows for intervention below the threshold of proof. In its traditional form, suspicion can be understood as...

<https://arxiv.org/abs/2610.03425>

2 **CONF**

Jumping the Line: Exploiting Length Predictions in LLM Scheduling

ArXiv CS.AI 50%

arXiv:2610.03430v1 Announce Type: new Abstract: Efficient request scheduling is increasingly important for reducing completion time in large language model (LLM) serving. Size-based policies such as Shortest Job First prioritize shorter requests, but output lengths are unknown...

<https://arxiv.org/abs/2610.03430>

3 CONF

Rethinking RAG in Long Videos: What to Retrieve and How to Use It?

ArXiv CS.AI 50%

arXiv:2606.13141v2 Announce Type: replace Abstract: Retrieval-augmented generation is extending beyond text to long videos, where query-relevant chunks can be represented across multiple modalities and temporal granularities. Progress in this setting, VideoRAG, is limited by two...

<https://arxiv.org/abs/2606.13141>

4 CONF

From Benchmarks to Production: A Text-to-SQL System for Complex Financial Data

ArXiv CS.AI 50%

arXiv:2610.03524v1 Announce Type: new Abstract: General-purpose Text-to-SQL systems achieve strong performance on academic benchmarks like Spider and BIRD, where schemas are relatively shallow and column values are often human readable. In production financial databases, where...

<https://arxiv.org/abs/2610.03524>

5 CONF

MRVQ: One Resident Index for Dimension- and Rate-Elastic Vector Search

ArXiv CS.AI 50%

arXiv:2610.03651v1 Announce Type: new Abstract: Dense-retrieval services must switch among embedding-prefix dimensions and index bit rates as latency, quality, and memory budgets change. Tuning a quantizer separately for each rate gives the best quality, but the retrieval tier...

<https://arxiv.org/abs/2610.03651>

6 CONF

Transcriptome-informed multi-modal AI for predicting neoadjuvant therapy response from breast cancer biopsies

ArXiv CS.AI 50%

arXiv:2610.03693v1 Announce Type: new Abstract: Scarcity of labeled data limits development of deep learning biomarkers in oncology. We develop a two-stage AI model predicting pathological complete response (pCR) to neoadjuvant therapy in breast cancer. The first stage learns...

<https://arxiv.org/abs/2610.03693>

7 CONF

ClarifyCodeBench: Evaluating LLMs on Clarifying Ambiguous Requirements for Code Generation

ArXiv CS.AI 50%

arXiv:2607.00711v2 Announce Type: cross Abstract: Large Language Models have emerged as programming assistants. However, the efficacy of code generation is constrained by the quality of input requirements, which are frequently ambiguous, incomplete, or underspecified. While LLMs...

<https://arxiv.org/abs/2607.00711>

8 CONF

RxnOptBench: Benchmarking LLMs for Reaction-Condition Optimization in Organic Methodology

ArXiv CS.AI 50%

arXiv:2610.02242v1 Announce Type: cross Abstract: Chemical reaction-condition optimization -- choosing the catalyst, ligand, solvent, reagent, temperature, time, and atmosphere that jointly maximize yield and stereoselectivity -- is a central, judgement-laden subtask of organic...

<https://arxiv.org/abs/2610.02242>

9 CONF

Diffusion-Based Synthetic Data Pretraining for Enhancing Activity Recognition

ArXiv CS.AI 50%

arXiv:2610.02292v1 Announce Type: cross Abstract: Human activity recognition (HAR) is increasingly important for healthcare, well-being, and daily monitoring applications, for which detecting alimentary activities such as eating and drinking can provide actionable insight into...

<https://arxiv.org/abs/2610.02292>

0 CONF

MIRROR: Multipath Quorum Integrity for LLM Multi-Agent Communication

ArXiv CS.AI 50%

arXiv:2610.02349v1 Announce Type: cross Abstract: Inter-agent communication is central to Large Language Model Multi-Agent Systems (LLM-MAS), but it introduces an underexplored vulnerability: Agent-in-the-Middle (AiTM) attacks that manipulate messages in transit without...

<https://arxiv.org/abs/2610.02349>

1 CONF

Automating the Application of HCI Principles: Skills for On-Demand UI Construction, the Human-AI Space to Think, and the Future of HCI

ArXiv CS.AI 50%

arXiv:2610.02369v1 Announce Type: cross Abstract: Human-computer interaction (HCI) is in the middle of a transition: large language models can now generate functional user interfaces (UIs) on demand from natural-language task descriptions. A user explains what they are trying to...

<https://arxiv.org/abs/2610.02369>

2 CONF

Hop-Decayed Influence: New Vulnerabilities of Structural Auxiliary Indexing in GraphRAG Pipelines with LLM

ArXiv CS.AI 50%

arXiv:2610.02373v1 Announce Type: cross Abstract: GraphRAG pipelines construct auxiliary structures during offline indexing--semantic summaries, hierarchical edges, and pre-computed scores--that determine how retrieval is prioritised at query time. Prior attacks target only...

<https://arxiv.org/abs/2610.02373>

3 CONF

Coco: An Agentic Copilot for the Hardware--Software Co-Design Lifecycle

ArXiv CS.AI 50%

arXiv:2610.02376v1 Announce Type: cross Abstract: Co-designing ML models and the accelerators that run them is an unusual reasoning task: architects must draw confident, high-stakes conclusions about systems that do not yet exist, and the pace of both model evolution and...

<https://arxiv.org/abs/2610.02376>

4 CONF

The Surprising Effectiveness of Shared Memory in Looped Transformers

ArXiv CS.AI 50%

arXiv:2610.02383v1 Announce Type: cross Abstract: Looped Transformers apply the same layers several times per token, adding compute to improve quality without more parameters. Each recursion, however, writes its own key-value cache, so memory still grows with compute....

<https://arxiv.org/abs/2610.02383>

5 CONF

LLM-Guided Reinforcement Learning with Representative Agents for Traffic Modeling

ArXiv CS.AI 50%

arXiv:2511.06260v2 Announce Type: replace-cross Abstract: Large language models (LLMs) are increasingly used as behavioral proxies for self-interested travelers in agent-based traffic models. Although more flexible and generalizable than conventional models, the practical use of...

<https://arxiv.org/abs/2511.06260>

6 CONF

Mitigating Private Data Leakage in LLMs with Whiteout

ArXiv CS.AI 50%

arXiv:2610.02418v1 Announce Type: cross Abstract: Modern large language models (LLMs) are trained on massive, largely unfiltered datasets, including content scraped from nearly every accessible website and user inputs. As a result, LLMs often memorize and reproduce personally...

<https://arxiv.org/abs/2610.02418>

7 CONF

Are you Synthesizing or Recalling? Evaluating LLMs on Algorithmic Code Retrieval

ArXiv CS.AI 50%

arXiv:2610.02438v1 Announce Type: cross Abstract: Large language models (LLMs) have demonstrated strong performance in code generation, where success depends on both recalling relevant algorithmic knowledge and reasoning about how to apply it. However, existing LLM pipelines are...

<https://arxiv.org/abs/2610.02438>

8 CONF

Counterexample Generation via Per-Theorem Symbolic Verifiers: When Imitation Hurts and Reinforcement Repairs

ArXiv CS.AI 50%

arXiv:2610.02444v1 Announce Type: cross Abstract: Large language models often solve a theorem forward yet fail to disprove a closely related false one: a falsification gap that supervised fine-tuning does not close and can actively worsen. We frame counterexample generation as...

<https://arxiv.org/abs/2610.02444>

9 CONF

Compound AI System Reliability: A Failure Taxonomy and Resilience Pattern Catalog from 150 Production Incidents

ArXiv CS.AI 50%

arXiv:2610.02503v1 Announce Type: cross Abstract: Deploying compound AI systems reliably and safely requires understanding failure modes that emerge at component boundaries, not within individual models. Cascading errors propagate across component boundaries, silent quality...

<https://arxiv.org/abs/2610.02503>

0 CONF

From Fragments to Global Maps: Learning Vectorized Map Aggregation with Large Language Models

ArXiv CS.AI 50%

arXiv:2610.02513v1 Announce Type: cross Abstract: Large-scale vectorized HD maps provide structured road information that is essential for perception, localization, and planning in autonomous driving. Constructing such maps requires aggregating noisy, fragmented, and overlapping...

<https://arxiv.org/abs/2610.02513>

1 CONF

Student-Guided Teacher Distillation for Efficient LLM Task Routing: Positioning Against Jev-Style System-1 Classifiers

ArXiv CS.AI 50%

arXiv:2610.02516v1 Announce Type: cross Abstract: Zero-shot classifiers are useful for routing user requests to specialized LLM tasks, but scoring every request against a large candidate set is expensive; a zero-shot NLI classifier must evaluate one premise-hypothesis pair per...

<https://arxiv.org/abs/2610.02516>

2 CONF

Instance-Dependent Regret for CMDPs with Step-Wise Constraints

ArXiv CS.AI 50%

arXiv:2610.02520v1 Announce Type: cross Abstract: We study online learning in episodic tabular constrained Markov decision processes with step-wise safety constraints. In such a setting, the constraints induce a safe subgraph that shapes the variance of cumulative rewards under...

<https://arxiv.org/abs/2610.02520>

3 CONF

Out of Sync, Out of Sight: Phantom State Attacks against IIoT Intrusion Detection

ArXiv CS.AI 50%

arXiv:2610.02552v1 Announce Type: cross Abstract: Machine learning-based intrusion detection systems (IDS) are critical for securing Industrial Internet of Things (IIoT) environments. Most adversarial research against them perturbs the feature vector or the traffic that produces...

<https://arxiv.org/abs/2610.02552>

4 CONF

DAGS: Disentangled Appearance-and-Geometry Steering of a Frozen Image DiT for Temporally Stabilized Generative Rendering

ArXiv CS.AI 50%

arXiv:2610.02567v1 Announce Type: cross Abstract: Diffusion transformers (DiTs) generate high-fidelity images from text and image conditions, but their outputs carry large variance and their faithfulness to a desired target depends heavily on how the condition is supplied. We...

<https://arxiv.org/abs/2610.02567>

5 CONF

Improving the Energy-Efficiency of the Code Generated by LLMs through Effective Prompting

ArXiv CS.AI 50%

arXiv:2610.02571v1 Announce Type: cross Abstract: As AI-assisted programming becomes increasingly mainstream, the environmental impact of AI-generated software has emerged as an important consideration. This motivates evaluating LLM-generated code beyond functional correctness...

<https://arxiv.org/abs/2610.02571>

6 CONF

Large Language Continuous Diffusion Models

ArXiv CS.AI 50%

arXiv:2610.02665v1 Announce Type: cross Abstract: Despite the success of discrete diffusion language models (dLMs) for fast parallel decoding, their non-smooth, high-dimensional space hinders trajectory steering for reasoning and inference acceleration. To overcome this, we...

<https://arxiv.org/abs/2610.02665>

7 **CONF****Self-Supervised Scaling of Terminal Environments for Scientific Domains****ArXiv CS.AI** **50%**

arXiv:2610.02710v1 Announce Type: cross Abstract: Terminal agents are increasingly deployed beyond software engineering in science and other specialized domains. Constructing training environments requires executable reference behavior and a domain-specific verifier that...

<https://arxiv.org/abs/2610.02710>

8 **CONF****Correcting Guided Diffusion Trajectories with Spectral Alignment****ArXiv CS.AI** **50%**

arXiv:2610.02753v1 Announce Type: cross Abstract: The practical success of conditional image generation hinges on fine-grained differences in condition alignment and visual fidelity. Classifier-free guidance (CFG) is central to this success, but its lack of an explicit criterion...

<https://arxiv.org/abs/2610.02753>

9 **CONF****Adaptive Spectral-Koopman Dynamics Modeling for Temporal Domain Generalization****ArXiv CS.AI** **50%**

arXiv:2610.02822v1 Announce Type: cross Abstract: Temporal Domain Generalization (TDG) has emerged to address real-world streaming data with distribution shifts over time. However, existing methods are either prone to overfitting to domain-specific noise in the data space or...

<https://arxiv.org/abs/2610.02822>

0 **CONF****AgentTrap: Stateful Feedback Deception against Autonomous Penetration Testing Agents****ArXiv CS.AI** **50%**

arXiv:2610.02869v1 Announce Type: cross Abstract: Autonomous penetration testing agents conduct multi-step attacks by continuously adapting their plans and actions to target responses. As a common defense, honeypots can be deployed to divert these agents from real assets by...

<https://arxiv.org/abs/2610.02869>

1 **CONF****Query-aware routing for Cross-lingual performance gains in Encoders****ArXiv CS.AI** **50%**

arXiv:2610.02875v1 Announce Type: cross Abstract: Multilingual encoders can exhibit reduced retrieval effectiveness when queries and relevant documents differ in language, despite strong same-language performance. We investigate whether Finnish and Swedish cross-lingual...

<https://arxiv.org/abs/2610.02875>

2 **CONF****Revealing Epistemic Uncertainty in MLLMs via Causal-Invariant Masking****ArXiv CS.AI** **50%**

arXiv:2610.02887v1 Announce Type: cross Abstract: Multimodal Large Language Models (MLLMs) suffer from hallucinations, creating a critical need for Uncertainty Quantification (UQ) to ensure reliable deployment. However, existing approaches struggle to detect uncertainty caused...

<https://arxiv.org/abs/2610.02887>

3 CONF

When Can We Trust the Matching Principle? Robust Deployment Geometry Under Finite-Sample and Model Uncertainty

ArXiv CS.AI 50%

arXiv:2610.02894v1 Announce Type: cross Abstract: Match only geometry you can identify; otherwise spread the penalty. We quantify that decision by the trust ratio $\tau = \epsilon / \gamma$ (estimation uncertainty over spectral separation). Under the linear-quadratic Matching...

<https://arxiv.org/abs/2610.02894>

4 CONF

Sentry: Learning to Recover from LLM Agent Failures at Test Time

ArXiv CS.AI 50%

arXiv:2610.02994v1 Announce Type: cross Abstract: LLM agents often fail mid-task due to invalid tool calls, repeated actions, or poorly grounded reasoning, and learning from these failures is a path to reliability. We find that how failure knowledge reaches the agent matters as...

<https://arxiv.org/abs/2610.02994>

5 CONF

Post-Training Frontier Text-to-Image Models by Composing Preference and Rubric Rewards

ArXiv CS.AI 50%

arXiv:2610.02967v1 Announce Type: cross Abstract: Recent text-to-image generation models have achieved remarkable visual quality, but improving them through post-training remains challenging because no single reward signal captures the full range of human preference. In this...

<https://arxiv.org/abs/2610.02967>

6 CONF

Tailoring the Quantization Space for 1-Bit KV Cache Compression

ArXiv CS.AI 50%

arXiv:2610.03027v1 Announce Type: cross Abstract: The key-value (KV) cache becomes a major memory bottleneck in long-context LLM inference, placing substantial pressure on memory capacity and bandwidth. To mitigate this bottleneck, vector quantization (VQ) has emerged as a...

<https://arxiv.org/abs/2610.03027>

7 CONF

NegT2IBench: When Negation Changes the Picture. A Polarity Benchmark for Text-to-Image Models

ArXiv CS.AI 50%

arXiv:2610.03084v1 Announce Type: cross Abstract: Text-to-image (T2I) models are judged by benchmarks that measure whether requested content appears, but these benchmarks largely overlook the complementary ability to satisfy negated constraints, for example, generating "a..."

<https://arxiv.org/abs/2610.03084>

8 CONF

Science Is Falling Behind the Frontier: Foundation Model Adoption Across Half a Million Papers

ArXiv CS.AI 50%

arXiv:2511.21739v2 Announce Type: replace-cross Abstract: We present the first large-scale analysis of AI foundation model usage in science -- not just citations or keywords. We find that adoption has grown rapidly, at nearly-exponential rates, with the highest uptake in...

<https://arxiv.org/abs/2511.21739>

9 CONF

Beyond Single Videos: Benchmarking and Active Evidence Seeking for E-Commerce Cross-Video Reasoning

ArXiv CS.AI 50%

arXiv:2610.03099v1 Announce Type: cross Abstract: E-commerce videos are information-dense and frequently compared by consumers evaluating products and merchants assessing marketing strategies. However, existing multimodal models mainly focus on single-video understanding and...

<https://arxiv.org/abs/2610.03099>

00 CONF

Ask, Relax, or Act? Evaluating Actionable Indeterminacy in LLM Preference Reasoning

ArXiv CS.AI 50%

arXiv:2610.03102v1 Announce Type: cross Abstract: An LLM agent can recognize uncertainty yet still choose the wrong next step: asking when action is already justified, or seeking clarification when the constraints must change. We formalize actionable indeterminacy: act when an...

<https://arxiv.org/abs/2610.03102>

01 CONF

The Fragility of Trigger-Tag Mechanisms for Misuse Detection in Open-Weight LLMs

ArXiv CS.AI 50%

arXiv:2610.03124v1 Announce Type: cross Abstract: Open-weight language models can be downloaded, modified, and deployed beyond their developers' control, limiting the effectiveness of centrally enforced safeguards. Recent work has therefore proposed \emph{trigger-tag} mechanisms...

<https://arxiv.org/abs/2610.03124>

02 CONF

Multimodal reasoning for broadly neutralizing antibody discovery from label-free human B cell repertoires across virus families

ArXiv CS.AI 50%

arXiv:2610.03160v1 Announce Type: cross Abstract: Discovering broadly neutralizing antibodies (bnAbs) from human natural immune repertoires remains a fundamental challenge in immunology, hindered by: the extreme rarity of bnAb, incomplete understanding of their cellular origins...

<https://arxiv.org/abs/2610.03160>

03 CONF

Predicting Steering Vectors and Adapter Weights for Few-Shot Author-Style Transfer

ArXiv CS.AI 50%

arXiv:2610.03163v1 Announce Type: cross Abstract: Adapting large language models to an individual author's style from a few examples is challenging, and scientific writing sharpens the difficulty: formal conventions leave little surface variation, and authors write about their...

<https://arxiv.org/abs/2610.03163>

04 CONF

Not Until the Evidence Says So: Teaching LLM Investigators When to Close a Case

ArXiv CS.AI 50%

arXiv:2610.03190v1 Announce Type: cross Abstract: Accident, defect and outage investigations end with a decision that ordinary question answering never faces: whether the evidence gathered so far is enough to close the case. We study this decision for LLM investigators, which...

<https://arxiv.org/abs/2610.03190>

05 CONF

D2K-Bench: Can LLM Agents Turn Expert Designs into Efficient GPU Kernels?

ArXiv CS.AI 50%

arXiv:2610.03226v1 Announce Type: cross Abstract: GPU kernels generated by large language model (LLM) agents can remain less efficient than expert implementations, but runtime alone does not reveal how the gap relates to design discovery and implementation. We introduce...

<https://arxiv.org/abs/2610.03226>

06 CONF

Training-Loss Guarantees for Muon with Finite-Step Newton--Schulz Orthogonalization

ArXiv CS.AI 50%

arXiv:2610.03306v1 Announce Type: cross Abstract: Existing convergence analyses of Muon either assume exact orthogonalization or analyze classical Newton--Schulz polynomials, and guarantee only stationarity, so it is unresolved what Muon's five tuned Newton--Schulz steps...

<https://arxiv.org/abs/2610.03306>

07 CONF

Cordial Learning: Distributed Training with Correlated Data

ArXiv CS.AI 50%

arXiv:2610.03330v1 Announce Type: cross Abstract: We consider a distributed learning task with agents that have correlated data. Specifically, the label of an agent depends on the input of other agents for the same sample, and these inputs are also correlated. Correlated data is...

<https://arxiv.org/abs/2610.03330>

08 CONF

Equivariant Visual-Tactile Diffusion Policy for Contact-Rich Manipulation

ArXiv CS.AI 50%

arXiv:2610.03333v1 Announce Type: cross Abstract: Imitation learning for contact-rich manipulation requires high-quality expert data that is expensive to obtain. This makes learning a sample-efficient policy a key issue. To address this, we propose VISTA, a workspace-level...

<https://arxiv.org/abs/2610.03333>

09 CONF

Corrupted but Correct: Why Vision-Language Models Lie to Themselves Internally

ArXiv CS.AI 50%

arXiv:2610.03445v1 Announce Type: cross Abstract: A targeted adversarial perturbation can drive a vision-language model's (VLM's) teacher-forced training loss for a fixed target caption to near zero, yet the same model, allowed to generate freely, produces the original, correct...

<https://arxiv.org/abs/2610.03445>

10 CONF

Detect and Suppress: A Mechanistic Defense against Adversarial Patches in VLA Models

ArXiv CS.AI 50%

arXiv:2610.03498v1 Announce Type: cross Abstract: Adversarial patches can disrupt Vision-Language-Action (VLA) models by manipulating visual observations, leading to failures in robot control. However, it remains poorly understood which internal mechanisms underlie these...

<https://arxiv.org/abs/2610.03498>

11 CONF

Cephalonauts One: A deep fMRI dataset for decoding naturalistic speech in the human brain

ArXiv CS.AI 50%

arXiv:2610.03558v1 Announce Type: cross Abstract: Cephalonauts One is a whole-brain 3 Tesla (3T) functional magnetic resonance imaging (fMRI) dataset recorded while subjects listened to audio podcasts. Three healthy subjects underwent multiple scanning sessions, each consisting...

<https://arxiv.org/abs/2610.03558>

12 CONF

Threat-Preserving Representation Sensitivity in Agent-Security Benchmarks

ArXiv CS.AI 50%

arXiv:2610.03585v1 Announce Type: cross Abstract: Security benchmarks for LLM-based agents often report the attack success rate (ASR) as a measure of model robustness and use these scores to compare different models and defense mechanisms, assuming that they describe the...

<https://arxiv.org/abs/2610.03585>

13 CONF

FrugalEvo: Towards Cost-Aware LLM-Guided Program Evolution

ArXiv CS.AI 50%

arXiv:2610.03675v1 Announce Type: cross Abstract: LLM-guided evolutionary methods, such as AlphaEvolve, have emerged as powerful approaches for challenging computational optimization problems, such as circle packing. However, prior work typically optimizes performance gain over...

<https://arxiv.org/abs/2610.03675>

14 CONF

LPS-Bench: Benchmarking Safety Awareness of Computer-Use Agents in Long-Horizon Planning under Benign and Adversarial Scenarios

ArXiv CS.AI 50%

arXiv:2602.03255v2 Announce Type: replace Abstract: Computer-use agents (CUAs) execute multi-stage tasks through tools, where an early unsafe decision can propagate to consequential actions. Evaluating only final outcomes can miss such decisions, while constructing executable...

<https://arxiv.org/abs/2602.03255>

15 CONF

Heads, Tails, and AI Fails: LLMs, Randomness, and Human Judgments

ArXiv CS.AI 50%

arXiv:2406.00092v2 Announce Type: replace Abstract: Randomness is central to human cognition and to many applications in which large language models are deployed, yet probabilistic token generation does not imply that LLMs can produce unbiased random sequences. We study how...

<https://arxiv.org/abs/2406.00092>

16 CONF

KGPFN: Unlocking the Potential of Knowledge Graph Foundation Model via In-Context Learning

ArXiv CS.AI 50%

arXiv:2605.14907v2 Announce Type: replace Abstract: Knowledge graph (KG) foundation models aim to generalize to graphs with unseen entities and relations by learning transferable relational structure. Most existing methods, however, focus on relation-level universality, leaving...

<https://arxiv.org/abs/2605.14907>

17 CONF

Counterfactual Evidence Audits Predict LLM-Agent Susceptibility to Ranked Context

ArXiv CS.AI 50%

arXiv:2606.00914v2 Announce Type: replace Abstract: LLM agents increasingly decide from evidence assembled by upstream systems: retrievers choose documents, recommenders choose posts, and memory systems choose prior events. Existing evaluations usually hold this evidence fixed,...

<https://arxiv.org/abs/2606.00914>

18 CONF

SovereignNegotiation-Bench: Evaluating User-Owned Personal Agents In Delegated Bargaining Under Privacy, Consent, Evidence, And Institutional Pressure

ArXiv CS.AI 50%

arXiv:2607.02814v2 Announce Type: replace-cross Abstract: Personal AI agents are beginning to negotiate for people, from refunds and bills to deposits and sales. A human agent in that position is judged by the duties owed to the principal, not by whether a deal was struck. We...

<https://arxiv.org/abs/2607.02814>

19 CONF

How Artificial Intelligence LLM Engines Shape the Global Conflict Information Environment

ArXiv CS.AI 50%

arXiv:2607.14197v2 Announce Type: replace Abstract: Artificial Intelligence (AI) answer engines now field a growing share of the questions that analysts, scholars, and the public ask about issues of peace and conflict. Large Language Models (LLMs) are known to hallucinate under...

<https://arxiv.org/abs/2607.14197>

20 CONF

WAM-OPD: Joint Video-Action Supervision for World Action Model Post-Training with On-Policy Distillation

ArXiv CS.AI 50%

arXiv:2608.22364v2 Announce Type: replace Abstract: World Action Models (WAMs) generate both future video and robot actions, offering two connected outputs for post-training supervision. How can a pretrained WAM learn from a stronger Teacher on the histories it encounters during...

<https://arxiv.org/abs/2608.22364>

21 CONF

Reach or Solve? Deep Diving into Agentic RL Gains with Checkpoint Handoffs

ArXiv CS.AI 50%

arXiv:2609.19636v2 Announce Type: replace Abstract: Reinforcement learning (RL) is widely used to improve language-model agents, and its gains are usually measured by final task success. However, an agent's earlier actions shape the states in which its later decisions are made,...

<https://arxiv.org/abs/2609.19636>

22 CONF

VIDiff: Translating Videos via Multi-Modal Instructions with Diffusion Models

ArXiv CS.AI 50%

arXiv:2311.18837v2 Announce Type: replace-cross Abstract: Diffusion models have achieved significant success in image and video generation. This motivates a growing interest in video editing tasks, where videos are edited according to provided text descriptions. However, most...

<https://arxiv.org/abs/2311.18837>

23 CONF

Mitigating Watermark Forgery in Generative Models via Randomized Key Selection

ArXiv CS.AI 50%

arXiv:2507.07871v5 Announce Type: replace-cross Abstract: Watermarking enables GenAI providers to verify whether content was generated by their models. A watermark is a hidden signal in the content, whose presence can be detected using a secret watermark key. A core security...

<https://arxiv.org/abs/2507.07871>

24 CONF

Demystifying LLM-as-a-Judge: Analytically Tractable Model for Inference-Time Scaling

ArXiv CS.AI 50%

arXiv:2512.19905v3 Announce Type: replace-cross Abstract: Recent developments in large language models have shown advantages in reallocating a notable share of computational resource from training time to inference time. However, the principles behind inference time scaling are...

<https://arxiv.org/abs/2512.19905>

25 CONF

The Effective Depth Paradox: Topology and Trainability in Deep CNNs

ArXiv CS.AI 50%

arXiv:2602.13298v4 Announce Type: replace-cross Abstract: This paper presents a controlled comparative study of convolutional neural network (CNN) topology and image classification performance across the architectural families VGG, ResNet, and GoogLeNet, evaluated on CIFAR-10...

<https://arxiv.org/abs/2602.13298>

26 CONF

Evaluation format, not model capability, drives measured triage failure in the assessment of consumer health AI

ArXiv CS.AI 50%

arXiv:2603.11413v4 Announce Type: replace-cross Abstract: A recent Nature Medicine study reported that ChatGPT Health under-triages 51.6% of emergencies and concluded that consumer-facing AI triage poses safety risks. Its protocol, however, was an exam-style scaffold (forced...

<https://arxiv.org/abs/2603.11413>

27 CONF

Many Preferences, Few Policies: Compact Portfolios for Multi-Objective LLM Alignment

ArXiv CS.AI 50%

arXiv:2604.04144v3 Announce Type: replace-cross Abstract: Aligning large language models (LLMs) requires balancing competing objectives such as helpfulness, harmlessness, and conciseness. The appropriate balance varies across users and applications, yet training, evaluating, and...

<https://arxiv.org/abs/2604.04144>

28 CONF

Dual Certified White-Box Inference for Input Convex Neural Networks

ArXiv CS.AI 50%

arXiv:2605.04722v2 Announce Type: replace-cross Abstract: Input convex neural networks (ICNNs) are used to learn convex objectives whose minimizers define decisions, making efficient and reliable optimization central to inference. At nonsmooth inputs, automatic differentiation...

<https://arxiv.org/abs/2605.04722>

29 CONF

EchoDistill: Robust Large Audio Language Models via Noisy-to-Clean Self-Distillation

ArXiv CS.AI 50%

arXiv:2605.23954v2 Announce Type: replace-cross Abstract: Large Audio Language Models (LALMs) remain vulnerable to acoustic noise, which can obscure task-relevant evidence and produce unreliable responses. We propose EchoDistill, a noisy-to-clean self-distillation framework that...

<https://arxiv.org/abs/2605.23954>

30 CONF

Where Do Apparent LLM Clinical Triage Failures Arise? Localizing the Multiple-Choice Format Effect

ArXiv CS.AI 50%

arXiv:2605.29889v2 Announce Type: replace-cross Abstract: LLM evaluations using clinician-authored triage vignettes have reported substantial under-triage under constrained multiple-choice testing. Yet model performance on the same clinical cases can change when responses are...

<https://arxiv.org/abs/2605.29889>

31 CONF

A Language Model from 1913: Pretraining on Historical Text

ArXiv CS.AI 50%

arXiv:2606.02991v2 Announce Type: replace-cross Abstract: While modern language models increasingly rely on ever-larger web corpora, we show that pretraining on historical text (e.g., pre-1913 text) in a data-constrained setting can produce a temporally grounded language model...

<https://arxiv.org/abs/2606.02991>

32 CONF

Assessing Rule Adherence of LLM Adjudicators in Call of Cthulhu TRPG

ArXiv CS.AI 50%

arXiv:2607.02802v2 Announce Type: replace-cross Abstract: As LLMs are increasingly deployed as autonomous adjudicators in games such as Call of Cthulhu (CoC), robust rule adherence becomes critical when user intent conflicts with system rules. However, as these models are...

<https://arxiv.org/abs/2607.02802>

33 CONF

Sentence-Level Context Sensitivity as a Training-Free Detector of Unsupported Content, Evaluated Against Trained Verifiers

ArXiv CS.AI 50%

arXiv:2607.04223v2 Announce Type: replace-cross Abstract: Retrieval-augmented generation (RAG) assistants summarize records in clinical and legal work, where one unsupported sentence can mislead a reader. The contrast between an output's likelihood with and without its source is...

<https://arxiv.org/abs/2607.04223>

34 CONF

EXAM2: Extending Audio Understanding in Multilingual and Multimodal Analysis

ArXiv CS.AI 50%

arXiv:2608.23758v3 Announce Type: replace-cross Abstract: Recent large audio language models (LALMs) have achieved impressive progress in audio understanding. However, existing evaluations remain largely constrained to English and narrow audio domains. Prior benchmarks typically...

<https://arxiv.org/abs/2608.23758>

35 **CONF****From Solo to Social Learning: Characterizing Recursive Social Improvement in LLMs****ArXiv CS.AI** **50%**

arXiv:2609.38516v2 Announce Type: replace-cross Abstract: Large language models (LLMs) can now improve themselves by revising the instructions they follow, and LLM agents are increasingly orchestrated to work together on complex problems. However, self-improvement methods...

<https://arxiv.org/abs/2609.38516>36 **CONF****Screw Attention: Rigid-Body Algebra Inside a Transformer****ArXiv CS.AI** **50%**

arXiv:2610.00904v2 Announce Type: replace-cross Abstract: Learned manipulation policies rediscover from data the spatial relations that rigid-body mechanics supplies in closed form, which leaves them fragile to geometric change. We present Screw Attention, a transformer layer in...

<https://arxiv.org/abs/2610.00904>