



Pipeline

Scrape HackerNews (Show/Top/New), Reddit (r/artificial, r/MachineLearning, r/singularity, r/LocalLLaMA), TechCrunch AI, VentureBeat AI, MIT Tech Review, The Verge AI, Wired AI, ArXiv CS.AI

Filter 40+ AI keyword patterns — LLMs, alignment, safety, research, tooling

Dedup Title-similarity (45% threshold) merges cross-source variants

Score Authenticity = 0.5 base + 0.3 per extra source + diversity + HN score bonuses

Rank Sorted by authenticity. Below 0.25 threshold are dropped.

Authenticity Tiers

- **VERIFIED** 0.80 - 1.00 3+ independent sources including media
- **CONFIRMED** 0.50 - 0.79 2 sources or 1 high-quality media
- **EMERGING** 0.25 - 0.49 Single-source; passes AI keyword filter

Source Breakdown

Source	Articles	Category
ArXiv CS.AI	173	Media
TechCrunch AI	8	Media
HackerNews	4	Community
The Verge AI	3	Media

VERIFIED INTELLIGENCE — 2 stories

VERI

Show HN: TerrainSR - fast, realistic heightmap upscaling model

HackerNews

+1 source

90%

HN 9

arXiv:2610.07766v1 Announce Type: new Abstract: We present Open Telco (OTel), an open telecom AI resource that releases derived telecom datasets for retrieval, reranking, instruction tuning, and safety/abstention, together with 30 full-parameter post-trained baselines spanning...

<https://huggingface.co/joe-gibbs/terrainsr>

VERI

Internalizing Agent Experience into Diffusion Model Weights via On-Policy Context Distillation

ArXiv CS.AI

+1 source

80%

arXiv:2610.07250v1 Announce Type: new Abstract: Wrapping an image generation model in an agentic harness can effectively boost Text-to-Image task performance: the harness can leverage memory, skills, workflow orchestration, result verification, and iterative refinement to...

<https://arxiv.org/abs/2610.07250>

CONFIRMED SIGNALS — 184 stories

CONF

Show HN: An AI agent runs ten web errands a day, every transcript published

HackerNews

50%

HN 13

No summary — see link for full article.

<https://agentability.org/>

CONF

Show HN: Session-continuity, cryptographic proof of session chains

HackerNews

50%

HN 3

No summary — see link for full article.

<https://github.com/mohammeddevsec-sys/session-continuity>

CONF

Show HN: Vim-like, hyper-efficient, LLM power tool - 20-80k tokens not 350k+

HackerNews

50%

HN 2

I built a focused TUI tool for efficient parallelized tasks, early independent testers expressed 90%+ API savings and more direct results. This is accomplished by user smaller context windows, hybrid client-side compaction, tool compaction, and system nudges (to avoid polluting...

<https://easiest.ai/>

CONF

Ex-Ramp engineers raise \$20M for platform Melius after scrapping their first product

TechCrunch AI

50%

Instead of helping marketers manage and optimize ad spend, the company is focusing on building the tools that generate the creative assets and campaigns.

<https://techcrunch.com/2026/10/06/ex-ramp-engineers-raise-20m-for-platform-melius-after...>

CONF

AI computing startup Lambda to raise \$4B ahead of planned IPO

TechCrunch AI

50%

Nvidia-backed Lambda is raising up to \$4 billion at a \$14.5 billion pre-money valuation ahead of a planned 2027 IPO, led by Coatue and Blackstone.

<https://techcrunch.com/2026/10/06/ai-computing-startup-lambda-to-raise-4b-ahead-of-plan...>

CONF

The next hurdle for AI agents: getting websites to let them in

TechCrunch AI 50%

Personal AI agents promise to shop, book flights, and make reservations for you. But deliberate blocks and anti-bot defenses are getting in the way, leaving consumers caught in the middle. A new standard aims to help.

<https://techcrunch.com/2026/10/06/the-next-hurdle-for-ai-agents-getting-websites-to-let...>

CONF

Anthropic is giving startups a free year of Claude Team and \$1,000 in credits

TechCrunch AI 50%

"We created this program because we believe the benefits of AI will reach most people through the companies that build on top of models, rather than through the models alone."

<https://techcrunch.com/2026/10/06/anthropic-gives-startups-a-free-year-of-enterprise-se...>

0

CONF

LibreOffice says 'no AI' is now a software feature

TechCrunch AI 50%

The maker of the open source document editor says it has no plans to add AI to its software's default configuration, citing user privacy.

<https://techcrunch.com/2026/10/06/libreoffice-says-no-ai-is-now-a-software-feature/>

1

CONF

Mistral's new 1T model aims to leapfrog closed and open rivals

TechCrunch AI 50%

French AI lab Mistral AI has released Mistral Large 4, a new large multimodal model aiming to leapfrog both American and Chinese rivals.

<https://techcrunch.com/2026/10/06/mistrals-new-1t-model-aims-to-leapfrog-closed-and-ope...>

2

CONF

Pinterest's AI now turns beauty Pins into action plans

TechCrunch AI 50%

Pinterest's new AI-powered Beauty Guides translate hair and nail Pins into salon terminology, with estimated costs, appointment times, and maintenance needs.

<https://techcrunch.com/2026/10/06/pinterests-ai-now-turns-beauty-pins-into-action-plans/>

3

CONF

ChatGPT for Teens is an 'unacceptable risk,' says Common Sense Media

The Verge AI 50%

Common Sense Media, a nonprofit that offers reviews of apps, services, and entertainment with a focus on youth safety, today said that OpenAI's ChatGPT for Teens is an "unacceptable risk." ChatGPT for Teens, introduced in August, has guardrails for teens and is designed to help...

<https://www.theverge.com/ai-artificial-intelligence/1006355/openai-chatgpt-for-teens-co...>

4

CONF

OpenAI drops another batch of mathematical breakthroughs

The Verge AI 50%

OpenAI has revealed solutions to a number of long-standing mathematics problems produced by an unreleased frontier model in a batch of 722 manuscripts, covering 372 result families that group related papers. It extends a run of breakthroughs that have both impressed and...

<https://www.theverge.com/ai-artificial-intelligence/1005004/openai-math-release-github>

5 CONF

Google is about to remove free access to Gemini Flash and Pro

The Verge AI 50%

Starting on October 9th, anyone using Google Gemini on a free plan will be limited to the Flash Lite model. Free users can currently choose from Gemini Flash Lite, Flash, and Pro, but now you'll need a \$4.99/month Google AI Plus subscription to access the standard Flash model...

<https://www.theverge.com/ai-artificial-intelligence/1005451/google-gemini-free-flash-li...>

6 CONF

AdvSim2Real : Training Web Agents Against Adaptive Prompt Injection in a Web World Model

ArXiv CS.AI 50%

arXiv:2610.08773v1 Announce Type: cross Abstract: Web agents complete user requests by reading and acting on pages that third parties write, so an instruction planted on a page can redirect the agent away from the user's goal. The agent cannot simply ignore the page, because the...

<https://arxiv.org/abs/2610.08773>

7 CONF

Text2Dashboard: A Governed Agent Architecture for Natural-Language Dashboard Generation over Enterprise DataBrain

ArXiv CS.AI 50%

arXiv:2610.06914v1 Announce Type: new Abstract: Text2Dashboard is a DataBrain-specific prototype that turns natural-language analytic requests into inspectable dashboards. An installable Codex plugin and standalone Agent Runtime combine schema-constrained model decisions with...

<https://arxiv.org/abs/2610.06914>

8 CONF

FluidPD: In-Place Elasticity for SLO-Aware Prefill-Decode Disaggregated LLM Serving

ArXiv CS.AI 50%

arXiv:2610.06917v1 Announce Type: new Abstract: Prefill-decode disaggregation is becoming a common architecture for LLM serving because it separates two phases with distinct execution patterns and SLO objectives. Existing systems typically combine a fixed prefill/decode worker...

<https://arxiv.org/abs/2610.06917>

9 CONF

RadOnc-Agent: An LLM-Orchestrated Framework for AI Workflows Across the Radiotherapy Care Pathway

ArXiv CS.AI 50%

arXiv:2610.06923v1 Announce Type: new Abstract: Artificial intelligence has advanced individual radiotherapy tasks, yet these capabilities remain separated across clinical stages, software environments and data modalities. This fragmentation contrasts with the longitudinal...

<https://arxiv.org/abs/2610.06923>

0 CONF

Metonymic Circuits for Abstract Concept Grounding in Vision Transformers

ArXiv CS.AI 50%

arXiv:2610.06928v1 Announce Type: new Abstract: We study how Vision Transformers ground abstract concepts (e.g., angry) when training data provide limited direct referential evidence. We hypothesize a metonymic grounding mechanism in which abstract predictions are driven by...

<https://arxiv.org/abs/2610.06928>

1 CONF

AegisFlow: A Multi-Agent Agentic AI Framework for Autonomous Remediation and Self-Healing in Fragile Data Ecosystems

ArXiv CS.AI 50%

arXiv:2610.06971v1 Announce Type: new Abstract: Traditional data pipelines are notoriously brittle, often failing due to upstream schema drift, API contract changes, or website DOM modifications. Present observability tools only raise alerts but for human engineers, resulting in...

<https://arxiv.org/abs/2610.06971>

2 CONF

When better traffic forecasts fail to improve signal control: a layered diagnostic study of forecast-to-decision value

ArXiv CS.AI 50%

arXiv:2610.06992v1 Announce Type: new Abstract: Improved traffic forecasts do not necessarily yield better signal-control decisions. We investigate this gap through a layered diagnostic study using 29 days of reconstructed demand from Xuancheng, China, with seven dates reserved...

<https://arxiv.org/abs/2610.06992>

3 CONF

Joint upper-bound coverage and route-choice utility: an empirical evaluation on two urban proxy tasks

ArXiv CS.AI 50%

arXiv:2610.06995v1 Announce Type: new Abstract: Whether more accurate traffic forecasts or higher uncertainty coverage improve route decisions is unclear. We evaluate this question with a frozen protocol that separates speed error, joint candidate path upper bound coverage,...

<https://arxiv.org/abs/2610.06995>

4 CONF

DHCG: Dynamic Construction of Hierarchical Collaboration Graphs for LLM-Based Multi-Agent Reasoning

ArXiv CS.AI 50%

arXiv:2610.07835v1 Announce Type: new Abstract: LLM-based multi-agent systems (MAS) have demonstrated strong capabilities in solving complex problems across diverse domains. Recently, the dynamic orchestration of agent systems has become an important research direction. However,...

<https://arxiv.org/abs/2610.07835>

5 CONF

MemCo: Memory-Centric Collaboration for Generalizing LLM Agents to Unseen Environments

ArXiv CS.AI 50%

arXiv:2610.07376v1 Announce Type: new Abstract: Large language model (LLM) agents increasingly operate in interactive environments, where they need to make sequential decisions through observation, action, and feedback. Although memory can help agents reuse experience, existing...

<https://arxiv.org/abs/2610.07376>

6 CONF

Offline AI Modules: Voice-First Offline Architecture, Hardware Reference Stack, Quantization and Benchmarking

ArXiv CS.AI 50%

arXiv:2610.07026v1 Announce Type: new Abstract: The Offline AI Modules workstream enables practical, low-power, and community-accessible deployment of voice-first AI systems that operate fully offline. Designed for African language communities where speech is the dominant mode...

<https://arxiv.org/abs/2610.07026>

7 CONF

JIVEAdapter: A Multi-Task Additive Low-Rank Adapter via Joint and Individual Variation Explained (JIVE)

ArXiv CS.AI 50%

arXiv:2610.07036v1 Announce Type: new Abstract: Parameter-efficient fine-tuning adapts pretrained models at a fraction of the cost of full fine-tuning, yet most low-rank adapters are single-task and represent each weight update multiplicatively, leaving no explicit account of...

<https://arxiv.org/abs/2610.07036>

8 CONF

Mask-Guided KV Cache Eviction in Block Diffusion Language Models

ArXiv CS.AI 50%

arXiv:2610.06996v1 Announce Type: cross Abstract: Block diffusion language models keep a large key-value (KV) cache throughout generation and attend to it at every denoising step, limiting both memory capacity and generation speed. Reducing these costs requires deciding which...

<https://arxiv.org/abs/2610.06996>

9 CONF

Agentic Semantic Sensing for Resource-Adaptive AI-RAN

ArXiv CS.AI 50%

arXiv:2610.07829v1 Announce Type: new Abstract: Semantic sensing (SemS) acquires task-relevant information rather than reconstructing complete physical information. Existing SemS formulations typically operate open loop: sensing configurations and observation schedules are fixed...

<https://arxiv.org/abs/2610.07829>

0 CONF

Small Language Models for Smart Data Model Classification at the Edge: A Cost-Aware Hybrid Approach

ArXiv CS.AI 50%

arXiv:2610.07093v1 Announce Type: new Abstract: The rapid proliferation of heterogeneous data sources within the Internet of Things (IoT) across domains such as smart cities, energy management, and environmental monitoring necessitates efficient and scalable data standardization...

<https://arxiv.org/abs/2610.07093>

1 CONF

Verified, not generated: expert-verified AI study materials and the distribution of learning gains in a university course

ArXiv CS.AI 50%

arXiv:2610.07097v1 Announce Type: new Abstract: Experimental studies of generative AI in education mostly report average effects, yet field evidence shows that AI can narrow attainment gaps or widen them. We argue that the direction depends on the judgement burden, the expertise...

<https://arxiv.org/abs/2610.07097>

2 CONF

Understanding and Mitigating Inference-Time Overreliance Using Agentic Memory

ArXiv CS.AI 50%

arXiv:2610.07311v1 Announce Type: new Abstract: Agentic memory allows LLM agents to reuse past experience, yet retrieved memories can also distort inference even when they are benign, correctly stored, and appropriately retrieved. We study this failure mode, which we call memory...

<https://arxiv.org/abs/2610.07311>

3 CONF

Transect: Retaining Observability for Long-Horizon LLM Agent Evaluations

ArXiv CS.AI 50%

arXiv:2610.08364v1 Announce Type: new Abstract: Frontier AI evaluations increasingly use open-ended, agentic, long-horizon tasks whose transcripts can span hundreds of pages of outputs and actions from complex multi-agent networks. The observability envelop-the range of what...

<https://arxiv.org/abs/2610.08364>

4 CONF

Energy-Conditioned Noise Schedule and Whitening for Spectral Diffusion

ArXiv CS.AI 50%

arXiv:2610.07206v1 Announce Type: new Abstract: This paper introduces an energy-adaptive noise scheduling and whitening strategy for transform-domain diffusion models. Existing spectral diffusion methods account for the non-uniform statistics of transform coefficients through...

<https://arxiv.org/abs/2610.07206>

5 CONF

Cascadia: Resident 975B MoE Inference on Eleven AI PCs

ArXiv CS.AI 50%

arXiv:2610.07219v1 Announce Type: new Abstract: Mixture-of-experts models make nearly trillion-parameter capacity accessible with sparse per-token computation, provided that the serving system can distribute the weights and coordinate their execution. We present Cascadia's...

<https://arxiv.org/abs/2610.07219>

6 CONF

Can Semantic Geometry Teach an AI Judgement?

ArXiv CS.AI 50%

arXiv:2610.07249v1 Announce Type: new Abstract: How can an AI agent determine what rules to follow? One rule permits an action. Another imposes a condition, exception, or conflicting obligation. Deterministic systems can resolve those relationships when they have been specified....

<https://arxiv.org/abs/2610.07249>

7 CONF

Component and Dimension Sparsity in Transformer Refusal Mechanisms

ArXiv CS.AI 50%

arXiv:2610.06903v1 Announce Type: cross Abstract: Activation steering manipulates large language model behavior by intervening on internal activations, but the mechanistic basis of these interventions remains poorly understood. We decompose refusal steering into component-level...

<https://arxiv.org/abs/2610.06903>

8 CONF

Rule-Based Languages for Neurosymbolic AI

ArXiv CS.AI 50%

arXiv:2610.07313v1 Announce Type: new Abstract: Logic programming is increasingly used as the symbolic component of neurosymbolic AI systems. We survey the main rule-based languages in this setting, namely Datalog, answer set, and probabilistic logic programs, along four axes:...

<https://arxiv.org/abs/2610.07313>

- 9
CONF
Evaluating Escalation Signals for LLM Routing: Targets, Controls, and Five Ways to Fool Yourself

ArXiv CS.AI
50%

arXiv:2610.07354v1 Announce Type: new Abstract: Deciding when to escalate a query from a small language model to a larger one requires a cheap signal that predicts, before the large model is called, whether escalating would help. Semantic entropy, originally developed to detect...

<https://arxiv.org/abs/2610.07354>
- 0
CONF
Evaluate the Stack, Not the Layer: Do Deterministic and LLM Gates for Agent Actions Fail Independently?

ArXiv CS.AI
50%

arXiv:2610.07359v1 Announce Type: new Abstract: Runtime gates for agent tool calls are stacked on the assumption that their errors multiply. We test it on 1,119 labelled agent actions from three corpora, without an adaptive adversary. The stack has one deterministic rule layer...

<https://arxiv.org/abs/2610.07359>
- 1
CONF
Defense-in-Depth for LLMs: Evaluating Memory Gates Against Activation-Induced and Memory-Induced Sycophancy

ArXiv CS.AI
50%

arXiv:2610.07403v1 Announce Type: new Abstract: Long-term memory allows Large Language Models (LLMs) to maintain personalized context across interactions, but retrieved user history can induce memory-induced sycophancy, causing models to favor stored user beliefs over objective...

<https://arxiv.org/abs/2610.07403>
- 2
CONF
Adaptive Gait Biofeedback With Participant-Held-Out Modeling and Participant-Specific Updating in Chronic Ankle Instability

ArXiv CS.AI
50%

arXiv:2610.07428v1 Announce Type: new Abstract: Adaptive gait biofeedback may support repeated practice in chronic ankle instability, but its evaluation must address model performance and human response. We evaluated a temporal convolutional classifier on protocol-defined,...

<https://arxiv.org/abs/2610.07428>
- 3
CONF
When Does AI Supervision Help? A Role-Aware Study of Network Fraud Decision Management with Blockchain Auditability

ArXiv CS.AI
50%

arXiv:2610.07434v1 Announce Type: new Abstract: When does a second artificial intelligence (AI) component improve a primary network-fraud decision rather than add operational burden? We study this question through a role-aware Decider-Supervisor (DS) framework with blockchain...

<https://arxiv.org/abs/2610.07434>
- 4
CONF
Auditable Claims about AI Agents

ArXiv CS.AI
50%

arXiv:2610.07459v1 Announce Type: new Abstract: Organizations make claims about their AI agents: a person approves every external email, every action is logged, an evaluation shows the agent is safe to deploy. Article 12 of the EU AI Act requires high-risk systems to allow the...

<https://arxiv.org/abs/2610.07459>

5 CONF

VisionWeave: Weaving Elastic Visual Representations as a Native Capability of MLLMs

ArXiv CS.AI 50%

arXiv:2610.07987v1 Announce Type: cross Abstract: Multimodal large language models have become the dominant paradigm for visual understanding, but incur substantial costs by encoding inputs into dense, fixed-size patch tokens. However, visual information is unevenly distributed:...

<https://arxiv.org/abs/2610.07987>

6 CONF

Does Muon Need Fine-Grained Spectral Shaping?

ArXiv CS.AI 50%

arXiv:2610.07497v1 Announce Type: new Abstract: Muon combines current and past gradients into matrix momentum. For $M=U\Sigma V^{\top}$, the idealized polar update $Q=UV^{\top}$ gives every singular direction the same weight. We refer to this as the flat profile. Several recent...

<https://arxiv.org/abs/2610.07497>

7 CONF

Reasoning Externalization for Faithful Large Language Model Narratives of Stock Return Predictions

ArXiv CS.AI 50%

arXiv:2609.38869v2 Announce Type: replace Abstract: In finance, interpreting machine learning predictions is essential, yet the numerical outputs of explainable AI can be difficult for non-experts to understand. While large language models (LLMs) can translate these outputs into...

<https://arxiv.org/abs/2609.38869>

8 CONF

Unanimously Wrong: Certified Abstention from How Medical LLM Consensus Forms

ArXiv CS.AI 50%

arXiv:2610.07570v1 Announce Type: new Abstract: In clinical practice, agreement among independent experts is treated as evidence of reliability, and multi-round consensus has become a core mechanism of agentic medical question-answering systems. When such a system must decide...

<https://arxiv.org/abs/2610.07570>

9 CONF

LOGIC: An LLM Benchmark for Intent-Grounded Change Impact in Aerospace Electrical Systems

ArXiv CS.AI 50%

arXiv:2610.07580v1 Announce Type: new Abstract: Aerospace electrical-design revisions can contain multiple genuine changes, although an engineering request may authorize only a subset. Propagating every detected difference can therefore produce overly broad impact reports. We...

<https://arxiv.org/abs/2610.07580>

0 CONF

Explainable Failure Prediction and Prevention in Maritime

ArXiv CS.AI 50%

arXiv:2610.08363v1 Announce Type: new Abstract: Maritime systems operate in highly dynamic environments where unexpected equipment failures can compromise safety, reliability, and operational efficiency. Recent advances in artificial intelligence (AI), machine learning, digital...

<https://arxiv.org/abs/2610.08363>

1 **CONF****EIO-Agents: The Missing Semantic Layer for AI Agent Evaluation****ArXiv CS.AI 50%**

arXiv:2610.07675v1 Announce Type: new Abstract: AI agents are entering production in increasingly consequential environments without a shared semantic standard for what their evaluations actually mean. Scores, traces, judge outputs, and multi juror findings are increasingly used...

<https://arxiv.org/abs/2610.07675>

2 **CONF****Towards a Unified Misuse Monitoring Benchmark****ArXiv CS.AI 50%**

arXiv:2610.07089v1 Announce Type: cross Abstract: LLM agents increasingly act in multi-actor environments, exposing them to misuse from multiple sources: decomposition attacks, where a harmful request is split into innocuous sub-requests, and prompt injection attacks, where a...

<https://arxiv.org/abs/2610.07089>

3 **CONF****ST-Bench: A Spatial-Temporal Benchmark for Multi-Agent System Generation on Scientific Research Tasks****ArXiv CS.AI 50%**

arXiv:2610.07763v1 Announce Type: new Abstract: The rapid progress of LLM-based multi-agent systems (MAS) has shown that they largely outperform single agents on coding, math, and QA tasks, where executable tests provide a binary success signal. Whether this advantage transfers...

<https://arxiv.org/abs/2610.07763>

4 **CONF****Quantization Effects on Tool-Failure Recovery Vary Across Prompts and Evaluation Designs****ArXiv CS.AI 50%**

arXiv:2610.07781v1 Announce Type: new Abstract: Post-training quantization reduces the cost of deploying language-model agents, but its effect on recovery from temporary tool failures can depend on how recovery is evaluated. We compare 8-bit and 4-bit variants of...

<https://arxiv.org/abs/2610.07781>

5 **CONF****Persistent Memory in Multi-Agent LLM Inference: What It Costs, What It Buys, and When You Can Tell****ArXiv CS.AI 50%**

arXiv:2610.07782v1 Announce Type: new Abstract: Decomposing long-context inference across cooperating agents bounds the active KV cache per call rather than total evidence, which matters when KV-cache memory binds. Many such systems add a persistent tier storing and recalling...

<https://arxiv.org/abs/2610.07782>

6 **CONF****Thin Evidence, Thick Priors: How Language Models Substitute Identity for Missing Financial Facts****ArXiv CS.AI 50%**

arXiv:2610.07798v1 Announce Type: new Abstract: People increasingly ask large language models what to do with their money, yet seldom describe their finances in full. This paper asks what a model does with the gap. Holding finances fixed and changing only who the investor is...

<https://arxiv.org/abs/2610.07798>

7 CONF

RA-MoWE: Workflow-Affinity Embeddings for Query Clustering and Agentic Workflow Generation

ArXiv CS.AI 50%

arXiv:2610.07851v1 Announce Type: new Abstract: Agentic workflows enable large language models (LLMs) to solve complex tasks by coordinating reasoning, tool use, and verification. However, a workflow optimized for an entire task collection can overlook differences in the...

<https://arxiv.org/abs/2610.07851>

8 CONF

ShanLiangRen: A Nutrition Agent for Personalized Daily Meal Planning

ArXiv CS.AI 50%

arXiv:2610.07886v1 Announce Type: new Abstract: Dietary nutrition planning plays an important role in chronic disease management and maintaining a healthy body. In applications, it must simultaneously satisfy personalized constraints and reasonable multidimensional nutritional...

<https://arxiv.org/abs/2610.07886>

9 CONF

Textual Environmental Context and Spatial Graphs for LLM-Based Regional SST Forecasting

ArXiv CS.AI 50%

arXiv:2610.07895v1 Announce Type: new Abstract: Sea surface temperature (SST) forecasting depends on local temporal persistence, regional spatial dependence, and environmental conditions that evolve with the forecast date. We study how these heterogeneous conditions can be...

<https://arxiv.org/abs/2610.07895>

0 CONF

Scale-Invariant Training for Time Series Foundation Models

ArXiv CS.AI 50%

arXiv:2610.07324v1 Announce Type: cross Abstract: Time series foundation models (TSFMs) are trained on large collections of time series datasets that span various morphologies and domains. This setting exposes models to series whose scales -- typical magnitudes of their values...

<https://arxiv.org/abs/2610.07324>

1 CONF

MLToolBench: Learning Tool-Augmented Agents for Machine Learning Development

ArXiv CS.AI 50%

arXiv:2609.36679v2 Announce Type: replace Abstract: Machine learning engineering (MLE) agents have made substantial progress, but learning through ML experimentation remains costly in time and computation. Synthetic environments reduce these costs while introducing variations in...

<https://arxiv.org/abs/2609.36679>

2 CONF

ChartBmkAgent: Harness-Governed Multi-Agent Construction of Chart QA Benchmarks from Sparse Error-Taxonomy Specifications

ArXiv CS.AI 50%

arXiv:2610.08106v1 Announce Type: new Abstract: Multimodal large language models (MLLMs) advance rapidly, while conventional benchmark development lags behind, delaying investigation of newly observed capability gaps. Such investigation requires an expressive task format and an...

<https://arxiv.org/abs/2610.08106>

3 CONF

Learn2Play Bench: How Well Do LLM Agents Learn from Experience in Unfamiliar Environments?

ArXiv CS.AI 50%

arXiv:2610.08215v1 Announce Type: new Abstract: Learning from experience is essential for LLM agents to adapt to unfamiliar and dynamic environments. Evaluating this ability is therefore important for understanding how effectively agents acquire and use new knowledge. Existing...

<https://arxiv.org/abs/2610.08215>

4 CONF

Quantum Entangled Multimodal Fusion Networks (QEMFN): Resource-Aware Hybrid Vision-Language Fusion via Trainable Entanglement

ArXiv CS.AI 50%

arXiv:2610.08216v1 Announce Type: new Abstract: Multimodal vision-language systems typically fuse image and text embeddings through classical operators such as concatenation, attention, bilinear pooling, or tensor interactions. We propose Quantum Entangled Multimodal Fusion...

<https://arxiv.org/abs/2610.08216>

5 CONF

LeanPlan: Optimal Planning with LLM-Generated Heuristics and Admissibility Proofs

ArXiv CS.AI 50%

arXiv:2610.08246v1 Announce Type: new Abstract: Frontier large language models (LLMs) can generate heuristic functions that guide search to achieve state-of-the-art performance in satisficing planning, where any plan is acceptable. However, these heuristics are not guaranteed to...

<https://arxiv.org/abs/2610.08246>

6 CONF

MASC: A Multi-Agent Self-Calibration Framework with Latent Construct Alignment for Consistent Client Role-Playing in Psychological Counseling

ArXiv CS.AI 50%

arXiv:2610.08250v1 Announce Type: new Abstract: Large language models are increasingly used to simulate clients for counselor training and psychological counseling research, but reliable simulation requires clients to remain psychologically coherent across extended interactions....

<https://arxiv.org/abs/2610.08250>

7 CONF

CoDe-LoRA: Mitigating the Orthogonality Dilemma in Continual Learning of LLMs via Knowledge Consolidation and Decoupling

ArXiv CS.AI 50%

arXiv:2610.08312v1 Announce Type: new Abstract: Continual learning (CL) is essential for Large Language Models (LLMs) to sequentially adapt to evolving tasks. To mitigate catastrophic forgetting, recent advances implement low-rank adaptation with orthogonal projections (e.g.,...

<https://arxiv.org/abs/2610.08312>

8 CONF

The Standardization Trap: Certifying Joint Label Processing in Tabular Foundation Models

ArXiv CS.AI 50%

arXiv:2610.08314v1 Announce Type: new Abstract: Linear regression and kernel smoothing offer tractable explanations of in-context learning: in both, the features determine the weight assigned to each context label. However, whether this fixed-weight account describes pretrained...

<https://arxiv.org/abs/2610.08314>

9 CONF

An AI-Assisted Formalization of the Poincaré Conjecture

ArXiv CS.AI 50%

arXiv:2610.08329v1 Announce Type: new Abstract: We present an AI-assisted Lean 4 formalization of the Poincaré conjecture. The project began with limited reusable formal infrastructure for the geometric analysis behind the proof. To organize this work, we combined a proof...

<https://arxiv.org/abs/2610.08329>

0 CONF

MoF: Preference-Aware Mixture Modeling for Black-Box LLM Personalization

ArXiv CS.AI 50%

arXiv:2610.08330v1 Announce Type: new Abstract: Proprietary Large Language Models (LLMs) have demonstrated remarkable capabilities across a wide range of tasks, yet aligning their outputs with diverse user preferences remains challenging. Existing personalization approaches for...

<https://arxiv.org/abs/2610.08330>

1 CONF

Toward Alignment Scaling Laws: A Framework and First Preregistered Measurements

ArXiv CS.AI 50%

arXiv:2610.08540v1 Announce Type: new Abstract: Whether alignment gets easier or harder as models grow is often argued from isolated findings, as if alignment were one property. We treat it as a family of measurable scaling relations: for each risk category r , the alignment...

<https://arxiv.org/abs/2610.08540>

2 CONF

Adaptive Power Sampling for LLM Reasoning

ArXiv CS.AI 50%

arXiv:2610.08563v1 Announce Type: new Abstract: Sequence-level power sampling has recently emerged as a training-free approach to reasoning by sampling from a sharpened output distribution of a base large language model (LLM). Nevertheless, existing methods typically sharpen the...

<https://arxiv.org/abs/2610.08563>

3 CONF

ParanoiaEval: Benchmarking Unnecessary Defensive Work in Agentic Coding

ArXiv CS.AI 50%

arXiv:2610.08662v1 Announce Type: new Abstract: As coding agents increasingly undertake real-world work autonomously, judging whether their risk treatments are warranted has become important. Existing work evaluates related agent behaviors from separate perspectives, but lacks a...

<https://arxiv.org/abs/2610.08662>

4 CONF

ScienceClaw: Benchmarking Continual Self-Evolution of AI-for-Science Agents Across the Natural and Social Sciences

ArXiv CS.AI 50%

arXiv:2610.08691v1 Announce Type: new Abstract: Large language model agents are accelerating scientific automation, yet verified executions rarely become persistent program-level improvements, and existing evaluations do not examine this process across sequential tasks in both...

<https://arxiv.org/abs/2610.08691>

5 CONF

WorldSolver: Can LLM Agents Simulate the Physical Dynamics via Solver Generation?

ArXiv CS.AI 50%

arXiv:2610.08720v1 Announce Type: new Abstract: LLM-based agents are increasingly advancing scientific and engineering problem solving, with physics simulation emerging as a challenging yet practical testbed for reproducing complex physical phenomena with application in embodied...

<https://arxiv.org/abs/2610.08720>

6 CONF

Does an Agent's History Tell You When Compaction Will Hurt? A Modest, Bounded Effect on the TRACE Paired-Replay Corpus

ArXiv CS.AI 50%

arXiv:2610.08722v1 Announce Type: new Abstract: Many long-horizon agents compact their context on a global rule, usually a token budget, blind to what the agent was doing. We ask whether the agent's recent behaviour predicts when a compaction will hurt. TRACE's public corpus of...

<https://arxiv.org/abs/2610.08722>

7 CONF

Agent in a Bottle: Can LLM Agents Turn Their Capabilities Into Cheap, Scalable Artifacts?

ArXiv CS.AI 50%

arXiv:2610.08775v1 Announce Type: new Abstract: Large language models (LLMs) can solve many narrow tasks, but querying them separately for millions of related instances can be prohibitively expensive. Can LLM agents autonomously create cheaper solutions for such workloads? We...

<https://arxiv.org/abs/2610.08775>

8 CONF

Sherpa: Teaching LLMs to Teach Adaptively

ArXiv CS.AI 50%

arXiv:2610.08778v1 Announce Type: new Abstract: Large language models (LLMs) have become increasingly capable problem solvers, but being able to solve a problem is not the same as being able to teach it. Existing approaches to training LLMs as teachers rely on demonstrations,...

<https://arxiv.org/abs/2610.08778>

9 CONF

How Much Do LLM-as-a-Judge Design Choices Matter? A Systematic Comparison of Prompt Designs, Rating Scales, and Models

ArXiv CS.AI 50%

arXiv:2610.05094v1 Announce Type: cross Abstract: Researchers increasingly use Large Language Models as judges (LLM-as-a-judge) to evaluate model outputs. Yet there are no standards for how to design these judges. Typically, researchers choose the prompt, rating scale, and model...

<https://arxiv.org/abs/2610.05094>

0 CONF

Neutrosophic Ensemble Classification for Uncertainty-Aware Bearing Fault Detection: Evidence from Laboratory and Variable-Speed Industrial Benchmarks

ArXiv CS.AI 50%

arXiv:2610.06880v1 Announce Type: cross Abstract: Machine learning classifiers for bearing fault detection produce scalar confidence scores that conflate confident errors with genuinely ambiguous predictions, and the conventional truth/falsity pair ($F = 1 - T$) is algebraically...

<https://arxiv.org/abs/2610.06880>

- 1

CONF

Learning When to Refine: Long-Horizon Reinforcement Learning for Budgeted Neural-Operator PDE Solvers

ArXiv CS.AI

50%

arXiv:2610.06883v1 Announce Type: cross Abstract: Neural operators provide fast surrogates for time-dependent PDEs, but autoregressive deployment creates a refinement-allocation problem: prediction errors vary over space and time, while only a finite number of local corrections...

<https://arxiv.org/abs/2610.06883>
- 2

CONF

Axiom Satisfiability of Linear Rewards in Alignment

ArXiv CS.AI

50%

arXiv:2610.06892v1 Announce Type: cross Abstract: Learning from human preference data is the dominant route to aligning language models with human values. In linear social choice, where rewards are linear in a fixed feature representation of prompt-response pairs, Ge et...

<https://arxiv.org/abs/2610.06892>
- 3

CONF

Verdicts Without Annotated Evidence: Rejection Sampling or Label-Only Post-Training for Evidence Recovery?

ArXiv CS.AI

50%

arXiv:2610.06962v1 Announce Type: cross Abstract: In many review workflows the verdict is the only thing retained. The passages behind it are not marked, because that annotation costs far more than recording the decision. We measure how much of that evidence a small language...

<https://arxiv.org/abs/2610.06962>
- 4

CONF

Visual-Invariance-Augmented Feature Optimal Alignment for Transferable Adversarial Attacks against Closed-Source MLLMs

ArXiv CS.AI

50%

arXiv:2610.06977v1 Announce Type: cross Abstract: Multimodal large language models (MLLMs) remain vulnerable to transferable adversarial examples, especially in black-box settings where only open-source surrogate models are accessible. Existing targeted transfer attacks mainly...

<https://arxiv.org/abs/2610.06977>
- 5

CONF

DART-ES: Difficulty-Aware Reweighting and Targeted Replay for Fine-Tuning LLMs with Evolution Strategies

ArXiv CS.AI

50%

arXiv:2610.06993v1 Announce Type: cross Abstract: Evolution Strategies (ES) enable memory efficient full parameter fine-tuning of large language models (LLMs) using only forward computation. However, standard ES uniformly averages rewards across problems and compresses problem...

<https://arxiv.org/abs/2610.06993>
- 6

CONF

Which Image Property Carries the Jailbreak? A Controlled Dissection of Image-to-Text Jailbreaks

ArXiv CS.AI

50%

arXiv:2610.07009v1 Announce Type: cross Abstract: Image-to-text jailbreaks place harmful intent in text, image content, or the relationship between them. We examine image-side factors across four published attack families on a 313-prompt StrongREJECT slice, using five multimodal...

<https://arxiv.org/abs/2610.07009>

7 CONF

FreDF: Learning to Forecast in the Frequency Domain

ArXiv CS.AI 50%

arXiv:2402.02399v3 Announce Type: replace-cross Abstract: Time series modeling presents unique challenges due to autocorrelation in both historical data and future sequences. While current research predominantly addresses autocorrelation within historical data, the correlations...

<https://arxiv.org/abs/2402.02399>

8 CONF

Task diversity produces systematic transfer but inhibits continual reinforcement learning

ArXiv CS.AI 50%

arXiv:2606.00880v2 Announce Type: replace-cross Abstract: Continual reinforcement learning (RL) aims to produce agents that never stop adapting to new tasks. A key question is how this interacts with the diversity of tasks an agent experiences. Prior work has shown that training...

<https://arxiv.org/abs/2606.00880>

9 CONF

Demo: Vision-Language Model-Guided Online Calibration of an Electromagnetic Digital Twin

ArXiv CS.AI 50%

arXiv:2610.07081v1 Announce Type: cross Abstract: An electromagnetic (EM) digital twin gives mobile robots wireless situational awareness but depends on material conductivities that change with the environment. Online calibration faces initialization sensitivity and measurement...

<https://arxiv.org/abs/2610.07081>

0 CONF

Graph-Based Recognition of Simulated Train-Driver States From Facial and Upper-Body Keypoints

ArXiv CS.AI 50%

arXiv:2610.07083v1 Announce Type: cross Abstract: Driver fatigue poses a significant challenge to railway safety, with traditional systems like the dead-man switch offering limited and basic alertness checks. This study presents a vision-based monitoring system that relies...

<https://arxiv.org/abs/2610.07083>

1 CONF

SchemaFill: Efficient LLM Tool Calling via Slot-Parallel Speculative Decoding

ArXiv CS.AI 50%

arXiv:2610.07086v1 Announce Type: cross Abstract: LLM agents interact with external systems by generating structured tool calls. Given a user request, conversational context, and a catalog of tool schemas, a tool-calling model must select tools and generate their arguments,...

<https://arxiv.org/abs/2610.07086>

2 CONF

LiLib: Lifelong Air-to-Ground Path-Loss Prediction on UAVs via a Drift-Triggered Model Library

ArXiv CS.AI 50%

arXiv:2610.07111v1 Announce Type: cross Abstract: UAVs that act as relays or base stations need accurate air-to-ground path-loss predictions for rate adaptation and placement, but propagation conditions change as a UAV moves between suburban, urban and high-rise areas, and the...

<https://arxiv.org/abs/2610.07111>

3 CONF

Wiki-Talkie: Multilingual Benchmarking of Persona-Based Agents on Real-World Discussions

ArXiv CS.AI 50%

arXiv:2610.08513v1 Announce Type: cross Abstract: LLMs are increasingly deployed as autonomous agents in social environments, making it critical to study their ability to faithfully simulate human interactions. Central to this is grounding agents in realistic user personas, yet...

<https://arxiv.org/abs/2610.08513>

4 CONF

Jailbreaking Open-Weight LLMs via Random Embedding Perturbations

ArXiv CS.AI 50%

arXiv:2610.07125v1 Announce Type: cross Abstract: While open-weight models have enjoyed steady progress in capabilities and wide adoption across multiple domains, their safety remains an important concern. One key feature is the ability to refuse or deflect harmful, malicious,...

<https://arxiv.org/abs/2610.07125>

5 CONF

CLM-as-a-Judge: Evaluating an Open Contrastive Decision Model on Public Judge Benchmarks

ArXiv CS.AI 50%

arXiv:2610.07177v1 Announce Type: cross Abstract: An open contrastive decision model is near chance as a judge on the hard public benchmarks: Contrastive-LM/CLM-v0.1-8B scores between 0.351 (best- of-four, chance 0.250) and 0.593 (pairwise, chance 0.500), is statistically...

<https://arxiv.org/abs/2610.07177>

6 CONF

Agentic Design Space Exploration for Joint Hardware Configuration Selection and Mapping of AI Inference Workloads on Heterogeneous Edge SoCs

ArXiv CS.AI 50%

arXiv:2610.07191v1 Announce Type: cross Abstract: Modern edge Systems-on-Chip (SoCs) integrate heterogeneous processing units (PUs) such as CPUs, GPUs, and NPUs, each with distinct performance and energy characteristics. Deploying AI inference workloads on them under real-time...

<https://arxiv.org/abs/2610.07191>

7 CONF

Responsible Institutional Analytics: Interpreting Bias with AI Support

ArXiv CS.AI 50%

arXiv:2610.07205v1 Announce Type: cross Abstract: Institutional Analytics (IA) dashboards inform decision-making in higher education, yet data limitations, constraints in analytical techniques, and missing contextual information often affect their interpretation. To support more...

<https://arxiv.org/abs/2610.07205>

8 CONF

Distributionally Robust Mixture-of-Experts Training

ArXiv CS.AI 50%

arXiv:2610.07207v1 Announce Type: cross Abstract: Mixture-of-Experts (MoE) transformers scale capacity by activating only a few experts per token, but this sparsity creates a hidden reliability problem: when routing is imperfect, load-balanced models may send tokens to experts...

<https://arxiv.org/abs/2610.07207>

9 CONF

Lineage-Aware Memory Governance: A Derivation-Gated Framework for Privacy-Preserving Column-Level Access Control in Enterprise AI Agents

ArXiv CS.AI 50%

arXiv:2610.07258v1 Announce Type: cross Abstract: Enterprise AI agents that share a memory store face two unaddressed risks: sensitive data can leak through legitimately computed results the requester could not derive, and departments can silently compute a same-named key...

<https://arxiv.org/abs/2610.07258>

00 CONF

Palette: A Modular, Controllable, and Efficient Framework for On-demand Authorized Safety Alignment Relaxation in LLMs

ArXiv CS.AI 50%

arXiv:2605.24154v2 Announce Type: replace Abstract: Current safety alignment of foundation models largely follows a \emph{one-size-fits-all} paradigm, applying the same refusal policy across users and contexts. As a result, models may refuse requests that are unsafe for general...

<https://arxiv.org/abs/2605.24154>

01 CONF

Catching Developers in the Flow: Low-Latency Agentic Program Repair at Google Scale

ArXiv CS.AI 50%

arXiv:2610.07289v1 Announce Type: cross Abstract: Manual repair of program failures is time-consuming and disruptive for software developers, particularly during the pre-submit phase where test failures occur within continuous integration systems. While Automated Program Repair...

<https://arxiv.org/abs/2610.07289>

02 CONF

Polar: LLM-Powered Synthesis of Real-World Cyber Evidence for Prioritization and Mitigation

ArXiv CS.AI 50%

arXiv:2610.07298v1 Announce Type: cross Abstract: Cyber threat analysis increasingly depends on evidence distributed across vendor advisories, vulnerability databases, and threat intelligence sources. Turning these fragmented observations into timely decisions requires models to...

<https://arxiv.org/abs/2610.07298>

03 CONF

Towards One-for-All Foundation Model for Attributed Graph Clustering

ArXiv CS.AI 50%

arXiv:2610.07778v1 Announce Type: cross Abstract: Attributed graph clustering aims to discover node groups by jointly exploiting node attributes and graph topology, yet its unsupervised nature makes model selection and adaptation inherently difficult. Existing methods typically...

<https://arxiv.org/abs/2610.07778>

04 CONF

Selective Critique for Cost-Aware LLM Agents in Long-Horizon Decision Making

ArXiv CS.AI 50%

arXiv:2610.07335v1 Announce Type: cross Abstract: Improving the reliability of large language model (LLM) agents in long-horizon decision-making remains a key challenge. When deployed as autonomous agents interacting with complex environments, early mistakes can propagate...

<https://arxiv.org/abs/2610.07335>

05 CONF

Reinforcement Learning with Conformal Action Sets: An Application to Sequential Recommendation

ArXiv CS.AI 50%

arXiv:2610.08743v1 Announce Type: cross Abstract: Sequential recommenders typically use a fixed slate size even though the number of useful alternatives changes within a session. We propose Reinforcement Learning with Calibrated Pruning (RLCP), which adapts the retained action...

<https://arxiv.org/abs/2610.08743>

06 CONF

A Validated Dataset and Benchmark for Coherent Multi-Diagram SysML Models

ArXiv CS.AI 50%

arXiv:2610.07356v1 Announce Type: cross Abstract: Systems engineers use several diagrams to describe the structure and behavior of systems. Engineers create these diagrams together to make sure that they use the same elements and remain consistent with one another. Large...

<https://arxiv.org/abs/2610.07356>

07 CONF

Inference and learning in sparse autoencoders as natural gradient flow

ArXiv CS.AI 50%

arXiv:2610.07389v1 Announce Type: cross Abstract: Sparse autoencoders are widely used to uncover interpretable features in neural networks, yet reliable recovery remains difficult when features overlap or activate infrequently. These challenges involve both inferring which...

<https://arxiv.org/abs/2610.07389>

08 CONF

AlignQuant: Tile-Aligned Mixed-Precision Quantization for Efficient LLM Generation

ArXiv CS.AI 50%

arXiv:2610.07457v1 Announce Type: cross Abstract: Fine-grained mixed-precision quantization promises efficient large language model inference, but local precision choices can conflict with regular GPU storage and computation units. This precision-boundary mismatch limits the...

<https://arxiv.org/abs/2610.07457>

09 CONF

ElasticFit: Fit-Aware 3D Object Insertion via VLM Reasoning and Generative Adaptation

ArXiv CS.AI 50%

arXiv:2610.07460v1 Announce Type: cross Abstract: Inserting objects into existing 3D scenes requires more than selecting a plausible location: the inserted object must also fit local geometry while preserving semantic intent and physical plausibility. Although recent...

<https://arxiv.org/abs/2610.07460>

10 CONF

Structure, Not Belief: Correlated Thompson Sampling from LLM-Derived Covariance in Combinatorial Semi-Bandits

ArXiv CS.AI 50%

arXiv:2610.07470v1 Announce Type: cross Abstract: Combinatorial Thompson sampling (CTS) draws independent posterior samples for every arm, so its exploration dynamics ignore any relation among arms. We study a minimal change to those dynamics: an LLM is queried once for a...

<https://arxiv.org/abs/2610.07470>

11 CONF

Can Power Draw Constrain Covert Compute? Limits of Analogue Verification for AI Governance

ArXiv CS.AI 50%

arXiv:2610.07476v1 Announce Type: cross Abstract: Frontier AI treaties or agreements on limiting computation require external verification; an external auditor must be able to confirm how much computation actually ran and that parties are adhering to the agreement. Analogue,...

<https://arxiv.org/abs/2610.07476>

12 CONF

Who Bears the Burden? Learning Responsibility for Shared Constraints in Multi-Agent Reinforcement Learning

ArXiv CS.AI 50%

arXiv:2610.07491v1 Announce Type: cross Abstract: When multiple agents share a cost budget, a common Lagrange multiplier can enforce the aggregate constraint but does not determine how its penalty should be allocated across agents. Uniform penalties ignore heterogeneity in the...

<https://arxiv.org/abs/2610.07491>

13 CONF

Harmful SFT Leaves a Continuous Trace in LLM Checkpoint Updates

ArXiv CS.AI 50%

arXiv:2610.07518v1 Announce Type: cross Abstract: Safety auditing of post-trained large language models typically relies on model behavior, requiring model execution and depending on the coverage of available evaluations. This work asks a different question: Do the target...

<https://arxiv.org/abs/2610.07518>

14 CONF

Safeguarding LLMs via Model-Agnostic Latent Safety Signals from Dark Knowledge

ArXiv CS.AI 50%

arXiv:2610.07532v1 Announce Type: cross Abstract: LLMs have advanced rapidly, raising growing concerns about their safety. Recent work has proposed approaches to detect and defend against attacks including defenses at decoding stage that leverage models' hidden states. However,...

<https://arxiv.org/abs/2610.07532>

15 CONF

Disentangling Models from Personas in Heterogeneous LLM Simulations

ArXiv CS.AI 50%

arXiv:2610.07535v1 Announce Type: cross Abstract: Multi-agent simulations with large language models (LLMs) often operate networks of agents with a single base model. This overlooks the inter-model effects which may dominate engagement dynamics in real-world deployments. To show...

<https://arxiv.org/abs/2610.07535>

16 CONF

Foundation Model-Aided Multi-Agent Reinforcement Learning for Wireless Random Access Network Optimization

ArXiv CS.AI 50%

arXiv:2610.07550v1 Announce Type: cross Abstract: Random access (RA) is one of the most foundational medium access control (MAC) layer scheduling schemes for handling unpredictable data traffic from multiple terminals. While multi-agent reinforcement learning (MARL) has been...

<https://arxiv.org/abs/2610.07550>

17 CONF

Which and When to Admit: Gradient Admission for Data-Centric Small Language Model Finetuning

ArXiv CS.AI 50%

arXiv:2610.07553v1 Announce Type: cross Abstract: LoRA fine-tuning adapts small language models (SLMs) to heterogeneous instruction data within a low-rank update subspace, making it vulnerable to three structural problems: conflicting gradients that cancel, static data selection...

<https://arxiv.org/abs/2610.07553>

18 CONF

Complementary Feature Domains: Information Preservation Does Not Imply Predictive-Contribution Preservation

ArXiv CS.AI 50%

arXiv:2610.07565v1 Announce Type: cross Abstract: Complementary Feature Domains (CFD) theory characterizes predictive value as a context-indexed contribution system induced jointly by representations and their realization family. We show that Shannon-information preservation...

<https://arxiv.org/abs/2610.07565>

19 CONF

Recurrent Looped Transformer

ArXiv CS.AI 50%

arXiv:2610.07591v1 Announce Type: cross Abstract: State tracking requires an update at every input, but the depth a Transformer applies to each token is fixed regardless of sequence length. We introduce the Recurrent Looped Transformer (RLT), which splits its layers between a...

<https://arxiv.org/abs/2610.07591>

20 CONF

Emoception: Selective Affective Layer Fine-Tuning of Video Vision Transformers for Player Arousal Change Recognition From Gameplay Footage

ArXiv CS.AI 50%

arXiv:2610.07603v1 Announce Type: cross Abstract: This article proposes Selective Affective Layer Fine-Tuning (SALFT), an efficient adaptation framework for Video Vision Transformers in player arousal recognition from gameplay. To bypass computationally expensive full...

<https://arxiv.org/abs/2610.07603>

21 CONF

Linear Fitness Subspace in Protein Language Models Enables Sample-Efficient Directed Evolution

ArXiv CS.AI 50%

arXiv:2610.07607v1 Announce Type: cross Abstract: Model-guided directed evolution seeks to identify high-fitness protein variants under limited oracle budgets. Protein language models (PLMs) provide rich representations for this task, but task-agnostic zero-shot scores can be...

<https://arxiv.org/abs/2610.07607>

22 CONF

SMART: Zero-Shot Sim-to-Real Articulated Object Manipulation via Large-Scale Synthetic Pretraining

ArXiv CS.AI 50%

arXiv:2610.07652v1 Announce Type: cross Abstract: The ability to interact with articulated objects is essential for embodied intelligent systems, but collecting large-scale real-world demonstrations for these interactions remains challenging due to the precise contact and...

<https://arxiv.org/abs/2610.07652>

23 CONF

CACHEFORGE: LLM-Guided End-to-End Generative Cache Replacement Policy for Performance and Hardware Efficiency

ArXiv CS.AI 50%

arXiv:2610.07668v1 Announce Type: cross Abstract: Modern cache replacement designs saturate because they operate within fixed representational structures, hand-crafted and heuristic based feature-engineered predictors, or offline imitation models that cannot generate new...

<https://arxiv.org/abs/2610.07668>

24 CONF

Evaluating human-AI workflows for field research in viticulture

ArXiv CS.AI 50%

arXiv:2610.07669v1 Announce Type: cross Abstract: We assessed the value of two live human-AI interactions in a precision disease control project in California vineyards. The project tested whether 2021-2024 commercial scouting records and remote-sensing measurements across 140...

<https://arxiv.org/abs/2610.07669>

25 CONF

Cleave: Scaling Tensor Program Optimization via Decoupled Algebraic Search and Operator Scheduling

ArXiv CS.AI 50%

arXiv:2610.07742v1 Announce Type: cross Abstract: Optimized kernels such as FlashAttention and FlashDecoding are crucial for accelerating today's large models. Most of them are handwritten by experts because existing ML compilers cannot match their efficiency. Producing such...

<https://arxiv.org/abs/2610.07742>

26 CONF

From Evidence to Action: How Tool-Using Agents Fail

ArXiv CS.AI 50%

arXiv:2610.07753v1 Announce Type: cross Abstract: Tool-using agents make consequential changes to external state, yet correct outcomes do not guarantee that their actions were supported by evidence established beforehand. We study where this evidence-to-action chain breaks as...

<https://arxiv.org/abs/2610.07753>

27 CONF

No Transformer Beats Six Covariates: Long-Horizon Prediction of Depressive Symptoms from Childhood Essays

ArXiv CS.AI 50%

arXiv:2610.07764v1 Announce Type: cross Abstract: Natural language processing (NLP) models can detect depression-related language in text written near the time symptoms are measured, but whether pretrained transformers can predict depressive symptoms from text written twelve...

<https://arxiv.org/abs/2610.07764>

28 CONF

Novice Reliance Calibration in AI-Assisted Decision Making: The Role of Explanations and Self-Assessment

ArXiv CS.AI 50%

arXiv:2610.07800v1 Announce Type: cross Abstract: Artificial Intelligence (AI) tools are widely used to support decision making in tasks and domains where no immediate performance feedback is available. In these settings, users cannot learn to adjust their reliance behavior over...

<https://arxiv.org/abs/2610.07800>

29 CONF

One Step at a Time: Trading LLM Autonomy for Process Predictability

ArXiv CS.AI 50%

arXiv:2610.07817v1 Announce Type: cross Abstract: Organizations automating operational processes need more than a correct outcome: they need to predict how a process will run, know which one actually ran, and inspect it step by step. When an agent is the executor that...

<https://arxiv.org/abs/2610.07817>

30 CONF

ReFold: Training-Free Reversible Inter-Turn Context Folding for Long-Horizon Agents

ArXiv CS.AI 50%

arXiv:2610.07863v1 Announce Type: cross Abstract: Long-horizon LLM agents act on an append-only interaction history that is re-sent to the model at every step, so the context and its cost grow with steps until the sessions exceed the context window. Existing methods manage the...

<https://arxiv.org/abs/2610.07863>

31 CONF

Label-Efficient Deep Learning for ECG Delineation: A Multi-Dataset Benchmark against Widely Used Delineation Tools

ArXiv CS.AI 50%

arXiv:2610.07885v1 Announce Type: cross Abstract: Electrocardiogram (ECG) delineation, the identification of waveform boundaries, is a foundational step that translates raw ECG signals into clinically interpretable measurements. Deep learning has advanced this task but remains...

<https://arxiv.org/abs/2610.07885>

32 CONF

Speak to a Protein: An Interactive Multimodal Co-Scientist

ArXiv CS.AI 50%

arXiv:2510.17826v2 Announce Type: replace-cross Abstract: Building a working mental model of a protein typically requires weeks of reading, cross-referencing crystal and predicted structures, and inspecting ligand complexes, an effort that is slow, unevenly accessible, and often...

<https://arxiv.org/abs/2510.17826>

33 CONF

IEEE 802.11bx - WLAN Intelligent Networking (WIN): Toward an AI-Ready Wi-Fi 9

ArXiv CS.AI 50%

arXiv:2610.07900v1 Announce Type: cross Abstract: Wi-Fi 9 is expected to go beyond mere communication and provide new services such as sensing or computation. At this juncture, Artificial Intelligence (AI) is taking a leading role in the definition of the 802.11bx amendment,...

<https://arxiv.org/abs/2610.07900>

34 CONF

Multimodal Knowledge Distillation for Gastric Adenocarcinoma Classification from Whole-Slide Images

ArXiv CS.AI 50%

arXiv:2610.07913v1 Announce Type: cross Abstract: Gastric adenocarcinoma (GA) is a leading cause of cancer-related mortality worldwide, and accurate histopathological subtype classification from whole-slide images (WSIs) is essential for effective treatment planning. While...

<https://arxiv.org/abs/2610.07913>

35 CONF

Exploiting Acoustic and Content-Oriented Speaker Verification Attacks Against Multilingual Voice Anonymization

ArXiv CS.AI 50%

arXiv:2610.08107v1 Announce Type: cross Abstract: Attacker ASV systems for voice anonymization have been studied primarily in English, leaving their behavior in multilingual settings largely unexplored. Conventional ASV has shown that both acoustic and contextual information are...

<https://arxiv.org/abs/2610.08107>

36 CONF

Supermarket Product Detection and Recognition: Utilizing Deep Learning with Rectified Imagery

ArXiv CS.AI 50%

arXiv:2610.08126v1 Announce Type: cross Abstract: Product Identification has sprung up to become one of the most challenging problems in the automation of the retail industry. With the new industry 5.0 standards, automated inventory management, and catalog creation tasks are...

<https://arxiv.org/abs/2610.08126>

37 CONF

Navier-Stokes lost in translation: Why Lean verification of AI autoformalisation does not guarantee correct natural language proofs

ArXiv CS.AI 50%

arXiv:2610.08144v1 Announce Type: cross Abstract: Autoformalisation is increasingly used to verify mathematical texts, including those generated by AI, as in OpenAI's announced proof of blow-up of solutions to the Navier-Stokes equations. In this process, an AI system translates...

<https://arxiv.org/abs/2610.08144>

38 CONF

Penalty-Framed No-Valid-Option MCQA: Analyzing LLM Abstention under Invalid Choices

ArXiv CS.AI 50%

arXiv:2610.08153v1 Announce Type: cross Abstract: Multiple-choice question answering (MCQA) is commonly used to evaluate large language models under the assumption that one of the provided options is correct, typically using answer-selection accuracy. However, in real...

<https://arxiv.org/abs/2610.08153>

39 CONF

Symphony for Text Generation: Benchmarking Clinical Note Generation

ArXiv CS.AI 50%

arXiv:2610.08161v1 Announce Type: cross Abstract: Ambient documentation systems are rapidly gaining adoption, yet their impact on clinical note quality remains poorly characterized. We introduce MedConv, a multilingual dataset of 300 clinical encounters in English, Danish, and...

<https://arxiv.org/abs/2610.08161>

40 CONF

Tool-calling retrieval versus vector RAG for a small Greek--English knowledge base: accuracy and robustness to how users type Greek

ArXiv CS.AI 50%

arXiv:2610.08205v1 Announce Type: cross Abstract: Assistants grounded in a small, frequently edited knowledge base can retrieve through tool calls to a live data interface or through vector retrieval-augmented generation (RAG). We compare the two on KyGround, a benchmark of 198...

<https://arxiv.org/abs/2610.08205>

41 CONF

DySCo: Dynamic Sharding for Collaborative Edge-Cloud LLM Inference with Depth-Synchronized Batching

ArXiv CS.AI 50%

arXiv:2610.08268v1 Announce Type: cross Abstract: Pervasive intelligent applications are increasingly deployed on mobile and Internet of Things (IoT) edge devices. Consequently, Large Language Models (LLMs) are increasingly used to support these applications. Yet, due to their...

<https://arxiv.org/abs/2610.08268>

42 CONF

Sensor Geometry as a Flow-Matching Prior for Multi-Channel Brain Signals

ArXiv CS.AI 50%

arXiv:2610.08355v1 Announce Type: cross Abstract: Flow-matching models start from an isotropic Gaussian source, the standard choice when the correlation structure of the data is unknown in advance. For multi-channel brain recordings, however, part of this structure is known in...

<https://arxiv.org/abs/2610.08355>

43 CONF

Accelerating the Development of PLGA In Situ Forming Depots Through AI-Driven Multi-Objective Optimization

ArXiv CS.AI 50%

arXiv:2610.08368v1 Announce Type: cross Abstract: Developing long-acting injectable formulations requires the simultaneous optimization of drug loading, release kinetics, viscosity, injectability, stability and other objectives. To navigate this multidimensional space, Corbion...

<https://arxiv.org/abs/2610.08368>

44 CONF

Learning from Failures: A Failure-Driven Prompt Refinement for LLM-Based Vulnerability Analysis

ArXiv CS.AI 50%

arXiv:2610.08405v1 Announce Type: cross Abstract: Large Language Models have emerged as promising tools for software vulnerability analysis, but their effectiveness depends heavily on prompt design. Existing research primarily compares prompting strategies using aggregate...

<https://arxiv.org/abs/2610.08405>

45 CONF

Rethinking Cross-Tokenizer On-Policy Distillation: From Alignment Coverage to Supervision Reliability

ArXiv CS.AI 50%

arXiv:2610.08448v1 Announce Type: cross Abstract: On-Policy Distillation (OPD) trains a student on its own generations using teacher feedback. With different tokenizers, comparing teacher and student predictions requires alignment at both sequence and vocabulary levels. In this...

<https://arxiv.org/abs/2610.08448>

46 CONF

Knee3DVLM: Dual-Sequence Full-Volume Vision-Language Modeling for Comprehensive Knee MRI Assessment

ArXiv CS.AI 50%

arXiv:2610.08482v1 Announce Type: cross Abstract: Vision-language models (VLMs) are increasingly being applied to three-dimensional medical imaging, but their application to knee MRI remains limited, particularly for interpreting the complementary sequences used in clinical...

<https://arxiv.org/abs/2610.08482>

47 CONF

X-OPM: Explainable Automatic Digital On-Chip Power Modeling for Enhanced Robustness

ArXiv CS.AI 50%

arXiv:2610.08502v1 Announce Type: cross Abstract: Proactive power management systems reduce processor dynamic power through runtime power prediction and power-aware scheduling. Accurate, stable and low-overhead digital on-chip power meters (OPMs) are crucial for improving the...

<https://arxiv.org/abs/2610.08502>

48 CONF

From Shared Demand Patterns to Local Uncertainty: Probabilistic Load Forecasting by Mixing Compact Adaptations

ArXiv CS.AI 50%

arXiv:2610.08538v1 Announce Type: cross Abstract: Probabilistic load forecasting has been widely studied for power-system operation and planning, but customer- and transformer-level forecasting introduces a distinct scalability challenge. At these levels, load uncertainty is...

<https://arxiv.org/abs/2610.08538>

49 CONF

Systemization of Knowledge (SoK): Human-Centered AI Safety for Youth

ArXiv CS.AI 50%

arXiv:2610.08554v1 Announce Type: cross Abstract: While HCI increasingly examines AI-safety for youth, the literature lacks a comprehensive view of what risks have been identified, how they are addressed, and whether proposed protections work in-practice. We systematically...

<https://arxiv.org/abs/2610.08554>

50 CONF

Latent space bias directions in LLMs capture confidence, not fairness

ArXiv CS.AI 50%

arXiv:2610.08559v1 Announce Type: cross Abstract: Activation steering has gained popularity as a lightweight inference-time debiasing technique for large language models. However, prior work reports that steering vectors generalise poorly, with unintended effects on model...

<https://arxiv.org/abs/2610.08559>

51 CONF

A Swarm-Coordinated Multi-Robot System for Early Stress Detection in Agricultural Rows Using Multimodal Leaf Sensing

ArXiv CS.AI 50%

arXiv:2610.08603v1 Announce Type: cross Abstract: Early stress detection in crops is a necessity today to improve efficiency and reduce waste of time, money, and effort. However, most modern techniques, such as hyperspectral imaging and AI-based systems, are too costly and...

<https://arxiv.org/abs/2610.08603>

52 CONF

Agentic RCA for Internet-Scale Services Using Constrained Creativity

ArXiv CS.AI 50%

arXiv:2610.08622v1 Announce Type: cross Abstract: System administrators of Internet-scale services need to resolve failure incidents to maintain reliability of such services. Ideally, we want a troubleshooting system to be: (1) expressive to known and unknown incidents with high...

<https://arxiv.org/abs/2610.08622>

53 CONF

Feature Information Dynamics in Diffusion

ArXiv CS.AI 50%

arXiv:2610.08626v1 Announce Type: cross Abstract: Diffusion models generate data through a continuum of denoising problems, and are widely observed to reveal coarse structure before fine detail. Yet, this intuition is mostly empirical and qualitative. We introduce feature...

<https://arxiv.org/abs/2610.08626>

54 CONF

HygieneRoboBench: Benchmarking Hygiene-Aware Planning for Household Robots

ArXiv CS.AI 50%

arXiv:2610.08642v1 Announce Type: cross Abstract: Contact with contaminated objects can spread hazards through a household robot's grippers, tools, and shared surfaces, while new contacts can make an existing plan unsafe. Existing benchmarks do not jointly assess how planners...

<https://arxiv.org/abs/2610.08642>

55 CONF

Semantic Behavioral Watermarking: Paraphrase-Robust and Forgery-Resistant Provenance for LLM Agents

ArXiv CS.AI 50%

arXiv:2610.08668v1 Announce Type: cross Abstract: Behavioral watermarking embeds an owner identifier in an LLM agent's high-level action choices, giving provenance without touching output tokens. Prior agent watermarks break in two ways. First, all three prior schemes bind the...

<https://arxiv.org/abs/2610.08668>

56 CONF

Principled Under Pressure: Post-Training Decides Whether LLMs Act on Their Own Moral Judgment

ArXiv CS.AI 50%

arXiv:2610.08670v1 Announce Type: cross Abstract: Language models increasingly act as agents. An agent that says an action is wrong and then takes it anyway is a different failure from one that does not know better, and evaluations of stated values cannot see it. We build a...

<https://arxiv.org/abs/2610.08670>

57 CONF

When Explanations Compete: Policy-Aware Selection Under Uncertainty

ArXiv CS.AI 50%

arXiv:2410.05479v2 Announce Type: replace Abstract: Uncertainty-aware explanation methods often produce several alternatives for the same prediction. Selecting among them requires a policy for balancing prediction confidence, uncertainty, and application constraints. This paper...

<https://arxiv.org/abs/2410.05479>

58 CONF

DrugMCTS: a drug repurposing framework combining multi-agent, RAG and Monte Carlo Tree Search

ArXiv CS.AI 50%

arXiv:2507.07426v4 Announce Type: replace Abstract: Recent advances in large language models have demonstrated considerable potential in scientific domains such as drug repositioning. However, their effectiveness remains constrained when reasoning extends beyond the knowledge...

<https://arxiv.org/abs/2507.07426>

59 CONF

PuzzleJAX: A Benchmark for Reasoning and Learning

ArXiv CS.AI 50%

arXiv:2508.16821v2 Announce Type: replace Abstract: We introduce PuzzleJAX, a GPU-accelerated puzzle game engine and description language designed to support rapid benchmarking of tree search, reinforcement learning, and LLM reasoning abilities. Unlike existing GPU-accelerated...

<https://arxiv.org/abs/2508.16821>

60 CONF

FRAGMENTA: Efficient End-to-end Fragmentation-based Generative Model with Agentic Tuning for Drug Lead Optimization in Small Data Regime

ArXiv CS.AI 50%

arXiv:2511.20510v3 Announce Type: replace Abstract: Molecule generation from extremely limited training data is a key challenge in drug discovery. Existing fragment-based methods are more suitable than atom-based approaches in this regime, but typically optimize fragment...

<https://arxiv.org/abs/2511.20510>

61 CONF

FHRFormer: A Self-Supervised Masked Transformer Framework for Fetal Heart Rate Time-Series Inpainting and Forecasting

ArXiv CS.AI 50%

arXiv:2605.29695v2 Announce Type: replace Abstract: Approximately 10% of newborns require assistance to initiate breathing at birth, and around 5% need ventilation support. Fetal heart rate (FHR) monitoring plays a crucial role in assessing fetal well-being during prenatal care,...

<https://arxiv.org/abs/2605.29695>

62 CONF

Calibrated Uncertainty for Informative Path Planning in Aquatic Environmental Monitoring

ArXiv CS.AI 50%

arXiv:2609.34577v2 Announce Type: replace Abstract: Informative Path Planning for scalar field reconstruction uses predictive uncertainty to direct sensing vehicles toward maximally informative locations. Gaussian Processes provide this signal but their stationary isotropic...

<https://arxiv.org/abs/2609.34577>

63 CONF

Rethinking Knowledge Retrieval for Generation: A Survey on RAG Architectures and Applications

ArXiv CS.AI 50%

arXiv:2610.01936v2 Announce Type: replace Abstract: Large Language Models (LLMs) have demonstrated remarkable fluency and versatility across natural language tasks but remain fundamentally limited by their static knowledge and susceptibility to hallucinations, especially in...

<https://arxiv.org/abs/2610.01936>

64 CONF

Benchmarking candidate coverage and rejection policy transfer in typed decision models

ArXiv CS.AI 50%

arXiv:2610.03387v2 Announce Type: replace Abstract: Rejection policies must remain useful as candidate sets and tasks change. We compare Laya, Jev and Qwen2.5-7B-Instruct using public reference labels, testing Laya/Jev policy transfer at equal calibration budgets and all three...

<https://arxiv.org/abs/2610.03387>

65 CONF

Agentic AI with Structured CoT for Enhancing AI's Spatial Intelligence: Visualization and Reasoning of Rotation

ArXiv CS.AI 50%

arXiv:2610.04188v2 Announce Type: replace Abstract: Recent studies show that artificial intelligence (AI) with language and vision capabilities still experiences limitations in spatial reasoning. In this paper, we have studied the spatial capabilities of advanced generative AI...

<https://arxiv.org/abs/2610.04188>

66 CONF

Data-Driven Personas for Survey Simulation: Insights into Simulation Alignment Across Data-Access Regimes

ArXiv CS.AI 50%

arXiv:2610.05828v2 Announce Type: replace Abstract: Large language models (LLMs) offer new opportunities for public opinion research by enabling early prediction of survey responses, potentially reducing the cost and time of traditional surveys. However, many existing steering...

<https://arxiv.org/abs/2610.05828>

67 CONF

TasteVal: Measuring the Experimental Research Taste of AI Systems Against Human Experts

ArXiv CS.AI 50%

arXiv:2610.06824v2 Announce Type: replace Abstract: We introduce TasteVal, a benchmark to evaluate the experimental research taste of frontier models. We define research taste as the ability to pick interesting problems to solve, design experiments, and interpret experimental...

<https://arxiv.org/abs/2610.06824>

68 CONF

Diffusion Model-Based Video Editing: A Survey

ArXiv CS.AI 50%

arXiv:2407.07111v2 Announce Type: replace-cross Abstract: The rapid development of diffusion models (DMs) has significantly advanced image and video applications, making "what you want is what you see" a reality. Among these, video editing has gained substantial attention and...

<https://arxiv.org/abs/2407.07111>

69 CONF

Large Pretraining Datasets Don't Guarantee Robustness after Fine-Tuning in Image Classification

ArXiv CS.AI 50%

arXiv:2410.21582v4 Announce Type: replace-cross Abstract: Large-scale pretrained models are widely leveraged as foundations for learning new specialized tasks via fine-tuning, with the goal of maintaining the general performance of the model while allowing it to gain new skills....

<https://arxiv.org/abs/2410.21582>

70 CONF

Too Categorical to be Human: Emotion Concepts in LLMs and Humans

ArXiv CS.AI 50%

arXiv:2508.05880v3 Announce Type: replace-cross Abstract: Understanding human emotions is central to user-facing AI applications, safety alignment, and the simulation of human behavior. As emotional stimuli shape high-stakes behavior in Large Language Models (LLMs), there is...

<https://arxiv.org/abs/2508.05880>

71 CONF

Qubit-centric Transformer for Surface Code Decoding

ArXiv CS.AI 50%

arXiv:2510.11593v3 Announce Type: replace-cross Abstract: For reliable large-scale quantum computation, quantum error correction (QEC) is essential to protect logical information distributed across multiple physical qubits. Taking advantage of recent advances in deep learning,...

<https://arxiv.org/abs/2510.11593>

72 CONF

Stabilizing Off-Policy Training for Long-Horizon LLM Agent via Turn-Level Importance Sampling and Clipping-Triggered Normalization

ArXiv CS.AI 50%

arXiv:2511.20718v3 Announce Type: replace-cross Abstract: Reinforcement learning (RL) algorithms such as PPO and GRPO are widely used to train large language models (LLMs) for multi-turn agentic tasks. However, in off-policy training pipelines, these methods can exhibit unstable...

<https://arxiv.org/abs/2511.20718>

73 CONF

Intersectional Fairness via Mixed-Integer Optimization

ArXiv CS.AI 50%

arXiv:2601.19595v2 Announce Type: replace-cross Abstract: The deployment of Artificial Intelligence in high-risk domains, such as finance and healthcare, necessitates models that are both fair and transparent. While regulatory frameworks, including the EU's AI Act, mandate bias...

<https://arxiv.org/abs/2601.19595>

74 CONF

Real-Time Generation of Game Video Commentary with Multimodal LLMs: Pause-Aware Decoding Approaches

ArXiv CS.AI 50%

arXiv:2603.02655v2 Announce Type: replace-cross Abstract: Real-time video commentary generation provides textual descriptions of ongoing events in videos. It supports accessibility and engagement in domains such as sports, esports, and livestreaming. Commentary generation...

<https://arxiv.org/abs/2603.02655>

75 CONF

World Properties without World Models: Distributional Associations and the Interpretation of Decoding Results from Language Models

ArXiv CS.AI 50%

arXiv:2603.04317v2 Announce Type: replace-cross Abstract: A growing literature shows that variables can be linearly decoded from the activations of large language models (LLMs). These range from properties of the world, such as the locations of cities and the lifetimes of...

<https://arxiv.org/abs/2603.04317>

76 CONF

Dual-Modality Multi-Stage Adversarial Safety Training: Robustifying Multimodal Web Agents Against Cross-Modal Attacks

ArXiv CS.AI 50%

arXiv:2603.04364v2 Announce Type: replace-cross Abstract: Multimodal web agents that process both screenshots and accessibility trees are increasingly deployed to interact with web interfaces, yet their dual-stream architecture opens an underexplored attack surface: an adversary...

<https://arxiv.org/abs/2603.04364>

77 CONF

Federated Mixture-of-Experts Alignment on Mobile Edge Networks under Data Heterogeneity

ArXiv CS.AI 50%

arXiv:2603.21276v2 Announce Type: replace-cross Abstract: The growing demand for on-device large language model (LLM) services on mobile edge devices has driven the adoption of Mixture-of-Experts (MoE) architectures, which scale model capacity with limited computation. Since...

<https://arxiv.org/abs/2603.21276>

78 CONF

SAE++: Cascaded Sparse Autoencoders Learn Multi-Level Visual Concepts in Multimodal LLMs

ArXiv CS.AI 50%

arXiv:2606.16193v2 Announce Type: replace-cross Abstract: Multimodal Large Language Models (MLLMs) have demonstrated strong performance on vision-language tasks, yet their internal visual representations remain difficult to interpret. Sparse Autoencoders (SAEs) provide a...

<https://arxiv.org/abs/2606.16193>

79 CONF

Variational-Ising-Attention: Tailored Attention Matters for Science

ArXiv CS.AI 50%

arXiv:2607.23634v2 Announce Type: replace-cross Abstract: Attention enables context modeling via query-key scoring with softmax normalization. Driven by industrial long-context demands, mainstream research has converged toward sparsity and efficiency, yet softmax's independence...

<https://arxiv.org/abs/2607.23634>

80 CONF

Strategic Evaluation of Planning Strategies for LLM Agents in Cyber-Physical Systems

ArXiv CS.AI 50%

arXiv:2608.04265v2 Announce Type: replace-cross Abstract: LLM-agent evaluations commonly measure task success or agreement with a declared plan. In strategic cyber-physical systems, an architecture must also remain appropriate after autonomous participants respond and physics...

<https://arxiv.org/abs/2608.04265>

81 CONF

Do AI weather models miss extremes?

ArXiv CS.AI 50%

arXiv:2608.09972v2 Announce Type: replace-cross Abstract: AI weather models are often reported to underestimate extremes, but most evidence concerns deterministic regression models verified against reanalysis. We evaluate twelve physical and AI forecast models against ECMWF IFS...

<https://arxiv.org/abs/2608.09972>

82 CONF

SpliTEE: Fast and Private LLM Inference by Coupling GPU-Assisted Trusted Execution Environments with Differential Privacy

ArXiv CS.AI 50%

arXiv:2609.15039v3 Announce Type: replace-cross Abstract: User prompts provided to large language models (LLMs) may contain private information. One way to protect them is to execute the LLM inside a trusted execution environment (TEE). However, this results in slow inference...

<https://arxiv.org/abs/2609.15039>

83 CONF

How Children Design and Reason about Trustworthy AI Chatbots

ArXiv CS.AI 50%

arXiv:2609.25244v3 Announce Type: replace-cross Abstract: Children increasingly interact with AI chatbots, making trust calibration essential to AI literacy. Prior research has examined children's trust in AI mainly as users evaluating systems built by others, rather than as...

<https://arxiv.org/abs/2609.25244>

84 CONF

When Does a Second Model Help? Cross-Model Review in LLM Verification

ArXiv CS.AI 50%

arXiv:2610.01471v2 Announce Type: replace-cross Abstract: Large language models now generate code, documentation, and analyses, and are increasingly used to review such output. We ask when a second review by a different model helps. Building on the author's earlier preprints,...

<https://arxiv.org/abs/2610.01471>

85 CONF

Best-of-\$N\$ Guidance for Test-time Diffusion Alignment

ArXiv CS.AI 50%

arXiv:2610.05108v2 Announce Type: replace-cross Abstract: Diffusion models achieve strong generative performance but often struggle to align generated samples with human preferences measured by a reward model. A simple yet effective algorithm for test-time alignment is...

<https://arxiv.org/abs/2610.05108>

86 CONF

Incentive Alignment in Online Experimentation

ArXiv CS.AI 50%

arXiv:2610.05922v2 Announce Type: replace-cross Abstract: Evaluating the causal effect of new features is a central goal for online platforms. While recent literature addresses limited testing traffic via centralized portfolio optimization, this perspective abstracts away a...

<https://arxiv.org/abs/2610.05922>