



Pipeline

Scrape HackerNews (Show/Top/New), Reddit (r/artificial, r/MachineLearning, r/singularity, r/LocalLLaMA), TechCrunch AI, VentureBeat AI, MIT Tech Review, The Verge AI, Wired AI, ArXiv CS.AI

Filter 40+ AI keyword patterns — LLMs, alignment, safety, research, tooling

Dedup Title-similarity (45% threshold) merges cross-source variants

Score Authenticity = 0.5 base + 0.3 per extra source + diversity + HN score bonuses

Rank Sorted by authenticity. Below 0.25 threshold are dropped.

Authenticity Tiers

- VERIFIED 0.80 - 1.00** 3+ independent sources including media
- CONFIRMED 0.50 - 0.79** 2 sources or 1 high-quality media
- EMERGING 0.25 - 0.49** Single-source; passes AI keyword filter

Source Breakdown

Source	Articles	Category
ArXiv CS.AI	157	Media
TechCrunch AI	15	Media
The Verge AI	8	Media
HackerNews	4	Community
MIT Tech Review	3	Media

VERIFIED INTELLIGENCE — 6 stories

VERI

Google is launching a one-stop Gemini agent for your work tasks

The Verge AI

+2 sources

100%

Google is launching a "universal" Gemini AI agent that can work across apps and devices in the background. The tool, announced as part of the Gemini at Work event on Thursday, will be available within the Gemini Enterprise app, allowing users to chat with the Gemini agent and...

<https://www.theverge.com/tech/1007904/google-gemini-ai-agent-enterprise>

VERI

Roundtables: A Conversation With the Creator of AI-Designed Viruses

MIT Tech Review

+2 sources

100%

Friday, October 16, 2026 Can AI design new life forms? In 2025, Stanford University PhD student Samuel King came up with a preliminary answer when he used a generative AI model to propose genetic blueprints for microscopic viruses. It isn't yet an example of AI-generated life,...

<https://www.technologyreview.com/2026/10/08/1146224/roundtables-a-conversation-with-the...>

VERI

Show HN: Crashlytics for AI Coding Agents

HackerNews

+1 source

90%

HN 3

arXiv:2610.10560v1 Announce Type: cross Abstract: AI foundation models pretrained on weather and climate data are increasingly fine-tuned to Earth science tasks well beyond weather forecasting. Their development and adoption are outpacing the scientific community's ability to...

<https://github.com/0xfirattamur/xcrashlytics>

VERI

Anthropic bans 'abusive or cruel behavior' toward Claude

The Verge AI

+1 source

80%

Anthropic is making changes to its usage policy for the first time in over a year to reflect new and high-risk cases of misuse - including election interference, weapons development, surveillance, and health and financial uses. But one of the most significant changes prohibits...

<https://www.theverge.com/ai-artificial-intelligence/1008100/anthropic-new-usage-policy-...>

VERI

Humanity's Sixth Sense: Benchmarking Intuitive Visual Reasoning in Multimodal Models

ArXiv CS.AI

+1 source

80%

arXiv:2610.08966v2 Announce Type: replace Abstract: Humans perceive far more in a scene than what is explicitly depicted: a single glance captures past causes and future trajectories; a quick peek determines if a vehicle can fit between two parked cars; a few seconds of video...

<https://arxiv.org/abs/2610.08966>

VERI

Job titles of the future: Delivery drone air traffic controller

MIT Tech Review

+1 source

80%

The moment Trevor Wischnewsky heard that drones were delivering pizza and sushi around his Texas neighborhood, his mind was made up. "It was super fascinating," he says. "I just immediately wanted to be a part of it." Wischnewsky is now an RPIC, or "remote pilot in command," at...

<https://www.technologyreview.com/2026/10/09/1145737/job-titles-delivery-drone-air-traff...>

CONFIRMED SIGNALS — 173 stories

CONF

Show HN: Jevman - AI decision models play Pac-Man

HackerNews 50% HN 64

Openai just launched their decisions endpoint, cloudflare launched clef the other week, and many more jev alternatives are out there. We wanted to put the popular ones to the test and thought Pac-Man is a good benchmark for simple and fast decision making. So we let jev 1.13,...

<https://opper.ai/jevman-benchmark/>

CONF

Show HN: Edi Life OS - self-hosted life dashboard with an MCP server for AI

HackerNews 50% HN 38

No summary — see link for full article.

<https://github.com/edrisranjbar/lifeos>

CONF

Show HN: Porch - a Mac menu bar work log for Claude Code and Codex sessions

HackerNews 50% HN 2

No summary — see link for full article.

<https://github.com/amone-labs/porch>

0 CONF

Pretend you're sitting at Elizabeth Holmes' desk on this weirdly detailed website

TechCrunch AI 50%

With over a thousand emails, slides, texts, and documents from the United States v. Elizabeth Holmes trial, Extend engineer Bo Lau created a website that simulates what it might have been like to rifle through the Theranos founder's desk.

<https://techcrunch.com/2026/10/08/pretend-youre-sitting-at-elizabeth-holmes-desk-on-thi...>

1 CONF

Fired OpenAI safety researchers dispute misconduct claims, warn of chilling effect

TechCrunch AI 50%

Three fired OpenAI safety researchers dispute allegations of mishandling sensitive information, warning in an open letter that their dismissals are creating a chilling effect on the company's AI safety culture.

<https://techcrunch.com/2026/10/08/fired-openai-safety-researchers-dispute-misconduct-cl...>

2 CONF

Ben Affleck is an AI nerd, and the internet is impressed

TechCrunch AI 50%

Ben Affleck is going viral for his deep knowledge of AI, from neural networks and transformers to open weights. The actor, who sold his AI filmmaking startup to Netflix earlier this year, is proving he's more than just a Hollywood star.

<https://techcrunch.com/2026/10/08/ben-affleck-is-an-ai-nerd-and-the-internet-is-impressed/>

3 CONF

Popular AI leaderboard Arena nearly doubles valuation to \$3.1B valuation in 10 months

TechCrunch AI 50%

The company behind the popular LMArena leaderboard has raised \$200 million led by Lightspeed and Khosla, and is now measuring AI models on alignment issues such as lying.

<https://techcrunch.com/2026/10/08/popular-ai-leaderboard-arena-nearly-doubles-valuation...>

4 CONF

OpenAI's revenue is reportedly \$20 billion less than previously projected

TechCrunch AI 50%

It had previously been reported that the AI lab's annualized revenue was some \$70 billion, but a new report claims it's a whole lot less than that.

<https://techcrunch.com/2026/10/08/openais-revenue-is-reportedly-20-billion-less-than-pr...>

5 CONF

OpenAI's math solutions aren't meeting the field's standards yet

TechCrunch AI 50%

OpenAI's flood of proofs deviated from the guidelines set by a group of mathematical researchers consulted by the frontier lab.

<https://techcrunch.com/2026/10/08/openais-math-solutions-arent-meeting-the-fields-stand...>

6 CONF

Goodfire says its new 'inside-out' monitors catch rogue AI agents at a fraction of the cost

TechCrunch AI 50%

Goodfire just launched what it says is a cheaper way to keep AI agents in check: Instead of paying a second AI to read everything an agent does, its monitors peek inside the model while it works and only call in backup when something looks fishy.

<https://techcrunch.com/2026/10/08/goodfire-says-its-new-inside-out-monitors-catch-rogue...>

7 CONF

Hear from Ambrosia Energy and Bloom Energy execs on where the AI infrastructure boom is creating opportunity at TechCrunch Disrupt 2026

TechCrunch AI 50%

Ambrosia Energy CEO Ben Longmier and Bloom Energy SVP Bill Thayer join the Smart Systems Stage at TechCrunch Disrupt. Register now to save up to \$100. Grab a second of the same pass to save 50%.

<https://techcrunch.com/2026/10/08/hear-from-ambrosia-energy-and-bloom-energy-execs-on-w...>

8 CONF

Cal AI's 19-year-old founder just raised \$10M for his new AI startup

TechCrunch AI 50%

Zach Yadegari, the teen co-founder of popular Cal AI calorie tracking app, has launched a new personal AI agent startup that competes with Instinct, Muse, and Bee.

<https://techcrunch.com/2026/10/08/cal-ais-19-year-old-founder-just-raised-10m-for-his-n...>

9 CONF

China's Manus raises over \$500M in first funding round since split with Meta

TechCrunch AI 50%

Boyu Capital and IDG Capital led the funding round, and existing shareholders Tencent, HSG (formerly known as Sequoia China), ZhenFund, and others also participated.

<https://techcrunch.com/2026/10/08/chinas-manus-raises-over-500m-in-first-funding-round-...>

0 CONF

Microsoft releases new Nvidia-chip AI PCs with revamped Windows 11

TechCrunch AI 50%

Microsoft revealed the specs and price for its Surface Laptop Ultra, AI PCs that run on Nvidia chips that are designed to run AI models and agents.

<https://techcrunch.com/2026/10/07/microsoft-releases-new-nvidia-chip-ai-pcs-with-revamp...>

1 CONF

ChatGPT is getting a lot more visual, with the launch of a new interface

TechCrunch AI 50%

OpenAI is launching a new user interface that will bring interactive visuals to ChatGPT.

<https://techcrunch.com/2026/10/07/chatgpt-is-getting-a-lot-more-visual-with-the-launch-...>

2 CONF

We're putting too much faith in AI's ability to say no

MIT Tech Review 50%

Ever since people first seriously contemplated giving machines an intelligence modeled on our own, there has never been any question that they would, like us, be able to say no. The sci-fi canon is full of stories of robotic disobedience. Most of these capers are, of course,...

<https://www.technologyreview.com/2026/10/09/1145728/we-are-putting-too-much-faith-in-ai...>

3 **CONF****Anthropic launches free AI security scans for open-source projects****The Verge AI** 50%

Anthropic's offering to help open-source projects track down security vulnerabilities with a new service called OSS Scanner. It says open-source projects that opt-in will get "thorough, periodic security scans by our strongest models at no cost." That could mean open-source...

<https://www.theverge.com/ai-artificial-intelligence/1008521/anthropic-open-source-oss-s...>

4 **CONF****USA Today becomes the latest publisher to sue OpenAI****The Verge AI** 50%

USA Today Co., along with the several local newspapers it owns, is suing OpenAI over claims that the company copied "hundreds of thousands" of articles to train its AI models, as reported earlier by Reuters. In a filing on Thursday, the publisher asks for damages of more than...

<https://www.theverge.com/ai-artificial-intelligence/1008198/usa-today-openai-copyright-...>

5 **CONF****SpaceXAI backs Omarchy, the controversial Linux distro, with \$1.5 million in compute****The Verge AI** 50%

If Elon Musk and SpaceXAI were going to back any Linux distro, it seems obvious they'd back Omarchy. Today it was announced that SpaceXAI would be joining the Omacom Foundation, which oversees Omarchy, as a Founding Corporate Patron and donating \$1.5 million worth of Grok tokens...

<https://www.theverge.com/tech/1008148/spacexai-omarchy-grok-david-heinemeier-hansson>

6 **CONF****Google's AI note-taking app transcribes your meetings completely offline****The Verge AI** 50%

Google has released an experimental note-taking app that can transcribe meetings and audio files entirely offline, as reported earlier by TechCrunch. The app, called Google AI Edge Foresight, is free to use and runs on macOS using the company's on-device EmbeddingGemma 2 model....

<https://www.theverge.com/tech/1007985/google-ai-notetaking-app-transcribe-offline>

7 **CONF****Can you trust Meta's Muse or OpenAI's Dots to run your life?****The Verge AI** 50%

My Decoder guest today is Hayden Field, The Verge's senior AI reporter, and we're discussing the new wave of consumer-friendly AI agents. If you've been paying attention to this space, you know AI enthusiasts have been using agents for a minute now — homebrew OpenClaw setups led...

<https://www.theverge.com/podcast/1007408/meta-muse-openai-dots-ai-agent-race-privacy-free>

8 **CONF****An Explainable Header-Centric Framework for Large-Scale Semantic Table Interpretation and Data Quality Assessment****ArXiv CS.AI** 50%

arXiv:2610.10541v1 Announce Type: new Abstract: Knowledge Graph (KG) quality depends not only on downstream graph validation, but also on the quality of tabular metadata used before integration. In metadata-only Semantic Table Interpretation (STI), where cell values are...

<https://arxiv.org/abs/2610.10541>

9 CONF

Verification and Self-Improvement in Agentic AI: Foundations and Limits

ArXiv CS.AI 50%

arXiv:2610.10611v1 Announce Type: new Abstract: Agentic AI systems can improve by searching longer, receiving additional support, or modifying how they propose and verify outputs. A performance score does not distinguish these mechanisms. We compare these changes through bounded...

<https://arxiv.org/abs/2610.10611>

0 CONF

The Harness as the Only Mutable Surface: Compliance-Bounded Self-Evolution of LLM Agents in Credit Pipelines, with a Measured Admission Gate

ArXiv CS.AI 50%

arXiv:2610.10629v1 Announce Type: new Abstract: Self-improving LLM agents can adapt a credit pipeline to a changed rule, but an agent that rewrites itself destroys the artefact a supervisor reviews: a named change, a recorded test, an approval. We argue that self-evolution is...

<https://arxiv.org/abs/2610.10629>

1 CONF

Plan-and-Patch: Diffusion Language Models for Agentic Planning

ArXiv CS.AI 50%

arXiv:2610.10786v1 Announce Type: new Abstract: Planning is increasingly important for long-horizon agents, where successful execution requires coordinating subgoals, tool use, and intermediate outcomes over many steps. Yet assumptions made during planning may be invalidated by...

<https://arxiv.org/abs/2610.10786>

2 CONF

Whose Ground Truth? Embracing Ambiguity in Human-Centered AI

ArXiv CS.AI 50%

arXiv:2610.10805v1 Announce Type: new Abstract: As AI systems increasingly interact with people and make decisions about them, understanding human interpretations becomes an important part of developing human-centered AI. Conventional machine learning and AI systems are largely...

<https://arxiv.org/abs/2610.10805>

3 CONF

On the Clock: Towards Punctual and Productive Time-Budgeted AI Agents

ArXiv CS.AI 50%

arXiv:2610.10833v1 Announce Type: new Abstract: We study whether small LLM agents can operate effectively under explicit wall-clock time budgets by both respecting the allocated runtime and using available time productively. We evaluate Qwen3.6-27B on five competitions from...

<https://arxiv.org/abs/2610.10833>

4 CONF

Reading the Room: Foundations, Design, and Challenges of Normative Competence in LLMs

ArXiv CS.AI 50%

arXiv:2610.10906v1 Announce Type: new Abstract: Human communities are governed by normative systems: shared standards that produce \textit{norms} dictating acceptable behavior, enforced through community sanctioning. Aligning increasingly autonomous AI systems with these norms...

<https://arxiv.org/abs/2610.10906>

- 5
CONF

StoreBench: A Live-Commerce Environment for Evaluating and Training Autonomous Operator Agents

ArXiv CS.AI 50%

arXiv:2610.10942v1 Announce Type: new Abstract: Reinforcement learning environments are now a primary lever for improving large language model (LLM) capabilities in post-training, yet most agentic benchmarks remain static: the world moves only when the agent acts, the reward is...

<https://arxiv.org/abs/2610.10942>
- 6
CONF

How Narrative Wrapping Affects LLM Refusal: A Cross-Language Benchmark and Defense

ArXiv CS.AI 50%

arXiv:2610.11005v1 Announce Type: new Abstract: Safety-aligned language models often refuse a harmful request stated directly but answer the same request inside a role-play or narrative wrapper. We measure this vulnerability across languages and registers: attack success on...

<https://arxiv.org/abs/2610.11005>
- 7
CONF

Curating Always-Loaded Context for LLM Agents: A Capacitated Assortment Model with Censored Feedback

ArXiv CS.AI 50%

arXiv:2610.11007v1 Announce Type: new Abstract: At the start of every session, LLM agents load a fixed context file, such as `$\texttt{AGENTS.md}`. Each loaded token in the file is charged again in every later round of the session, and these files can degrade performance as they...

<https://arxiv.org/abs/2610.11007>
- 8
CONF

Beyond Imitation: A Framework and Benchmark for LLM-Assisted Peer Review

ArXiv CS.AI 50%

arXiv:2610.11087v1 Announce Type: new Abstract: The rapid growth of scientific publishing has strained peer review, particularly in machine learning, raising concerns about declining review quality and increasing reviewer workload. Large language models (LLMs) have been proposed...

<https://arxiv.org/abs/2610.11087>
- 9
CONF

OpenProblemBench: Benchmarking AI on Open Problems in the Foundational Theoretical Sciences

ArXiv CS.AI 50%

arXiv:2610.11118v1 Announce Type: new Abstract: The next frontier for artificial general intelligence is tackling unresolved scientific problems, calling for benchmarks that assess progress beyond established knowledge. We introduce OpenProblemBench, a benchmark of 82 unresolved...

<https://arxiv.org/abs/2610.11118>
- 0
CONF

When Interfaces Speak: Data-Aware Generative UI Harness for Active Interaction

ArXiv CS.AI 50%

arXiv:2610.11123v1 Announce Type: new Abstract: Most human-agent interaction today remains text-based. Natural language can impose cognitive overload, ambiguity, information chaos, and slow input for complex tasks; ephemeral generative UIs can present structured information and...

<https://arxiv.org/abs/2610.11123>

1 CONF

GameCommBench: A Unified Benchmark and Type-Aware Evaluation for AI-Generated Game Commentary

ArXiv CS.AI 50%

arXiv:2610.11129v1 Announce Type: new Abstract: Game commentary is an open-ended generation task requiring multimodal perception, strategic reasoning, and contextual knowledge. Existing AI-Generated Game Commentary (AI-GGC) studies remain fragmented across games, modalities, and...

<https://arxiv.org/abs/2610.11129>

2 CONF

Social Pain Disrupts Emotion-Action Brain-State Dynamics in Adolescents with Non-Suicidal Self-Injury

ArXiv CS.AI 50%

arXiv:2610.11155v1 Announce Type: new Abstract: Non-suicidal self-injury (NSSI) is prevalent among adolescents with depression, but the rapid brain-state dynamics linking social distress to maladaptive behavior remain unclear. We combine an experimental pain paradigm,...

<https://arxiv.org/abs/2610.11155>

3 CONF

What to Admit and How to Present: Governing Persistent Memory in LLM Agents

ArXiv CS.AI 50%

arXiv:2610.11188v1 Announce Type: new Abstract: Persistent memory can improve personalization in LLM agents but can also induce sycophancy and cross-domain leakage. We distinguish two governance decisions: admission, which determines what recalled information enters the working...

<https://arxiv.org/abs/2610.11188>

4 CONF

AliO: Output Alignment Matters in Long-Term Time Series Forecasting

ArXiv CS.AI 50%

arXiv:2610.11213v1 Announce Type: new Abstract: Long-term Time Series Forecasting (LTSF) tasks, which leverage the current data sequence as input to predict the future sequence, have become increasingly crucial in real-world applications such as weather forecasting and planning...

<https://arxiv.org/abs/2610.11213>

5 CONF

When Lower Reconstruction Loss Hurts: Distributionally Robust Refinement for Low-Bit LLM Quantization

ArXiv CS.AI 50%

arXiv:2610.11226v1 Announce Type: new Abstract: Weight-only post-training quantization (PTQ) relies heavily on reconstruction loss minimization to preserve model quality at low precision. We show that the weights favored by minimizing this loss need not yield better model...

<https://arxiv.org/abs/2610.11226>

6 CONF

Harness Compilation: Which Decisions Should a Small Vision-Language Model Keep?

ArXiv CS.AI 50%

arXiv:2610.11231v1 Announce Type: new Abstract: Small vision-language models may be able to read external evidence yet struggle to obtain it. We introduce Harness Compilation (HC), an offline procedure that adapts the division of work between a frozen small VLM and its external...

<https://arxiv.org/abs/2610.11231>

- 7
CONF
LLM-IDEA: Identifiability-Driven Experimental Agent for Autonomous Discovery of Mechanistic World Models
ArXiv CS.AI 50%

arXiv:2610.11253v1 Announce Type: new Abstract: Large language model agents are being increasingly deployed as autonomous scientists, designing experiments and inferring mechanistic world models with minimal human oversight. Yet identifiability is often overlooked: when a...

<https://arxiv.org/abs/2610.11253>
- 8
CONF
MedBenchAgent: Towards Systematic Automation of Medical VLM Benchmark Construction
ArXiv CS.AI 50%

arXiv:2610.11312v1 Announce Type: new Abstract: Large-scale construction of medical vision-language model (VLM) benchmarks is increasingly feasible with richly annotated imaging datasets and large language models (LLMs), yet existing automation largely focuses on generating...

<https://arxiv.org/abs/2610.11312>
- 9
CONF
DivMoE: Fine-Grained MoE Upcycling via Cross-Domain Expert Composition
ArXiv CS.AI 50%

arXiv:2610.11317v1 Announce Type: new Abstract: Mixture-of-Experts (MoE) architectures have become essential for scaling large language models, with recent work demonstrating the benefits of fine-grained expert designs. Training such models from scratch is expensive, and sparse...

<https://arxiv.org/abs/2610.11317>
- 0
CONF
Finsler Flow Matching: Dynamics-Aware Geodesic Interpolation for Single-Snapshot Trajectory Inference
ArXiv CS.AI 50%

arXiv:2610.11318v1 Announce Type: new Abstract: Single-cell snapshot data can resolve a continuum of cellular states but do not uniquely determine the dynamics governing transitions between them. However, additional dynamical information can often be encoded in a cell-cell...

<https://arxiv.org/abs/2610.11318>
- 1
CONF
The Review Lottery: Benchmarking an Observational Estimator of Peer-Review Noise (ICLR 2017-2025)
ArXiv CS.AI 50%

arXiv:2610.06591v2 Announce Type: replace Abstract: How much of a conference accept/reject decision would change if the same paper were reviewed by a different set of reviewers? Running a second independent program committee is the gold standard for answering this, but it is...

<https://arxiv.org/abs/2610.06591>
- 2
CONF
TokenBank: Financial Infrastructure for AI Services
ArXiv CS.AI 50%

arXiv:2610.11333v1 Announce Type: new Abstract: AI services incur inference costs during execution, while revenue may arrive later. Changing API prices, limited upfront capital, and service failures can limit operators' ability to sustain or expand their services. Beyond...

<https://arxiv.org/abs/2610.11333>

3 CONF

Learning from Own Solutions: Self-Conditioned Credit Assignment for Reinforcement Learning with Verifiable Rewards

ArXiv CS.AI 50%

arXiv:2606.18810v2 Announce Type: replace-cross Abstract: Reinforcement learning with verifiable rewards (RLVR) has driven substantial progress in training LLMs for reasoning tasks, but representative methods such as GRPO assign uniform credit across all tokens, wasting gradient...

<https://arxiv.org/abs/2606.18810>

4 CONF

RL-ARC: Calibrating Large Reasoning Models via Reasoning-guided Uncertainty

ArXiv CS.AI 50%

arXiv:2610.11352v1 Announce Type: new Abstract: Language models (LMs) are commonly trained with Reinforcement Learning with Verifiable Rewards (RLVR) to enhance their reasoning capabilities. However, since RLVR does not explicitly account for calibration during training, it can...

<https://arxiv.org/abs/2610.11352>

5 CONF

UniData: Universal Multimodal Instruction Generation Pipeline

ArXiv CS.AI 50%

arXiv:2610.11363v1 Announce Type: new Abstract: Multimodal Large Language Models (MLLMs) are increasingly being applied in a wider range of real-world scenarios. However, due to the substantial labor cost, creating high-quality multimodal instruction datasets for MLLMs remains a...

<https://arxiv.org/abs/2610.11363>

6 CONF

Equal Path Cost, Unequal Output Effects: Understanding Perturbation Propagation in Diffusion Models

ArXiv CS.AI 50%

arXiv:2610.11380v1 Announce Type: new Abstract: Diffusion models have achieved remarkable success in generative modeling, with their sampling procedures routinely modified to control generation and improve efficiency. These modifications introduce perturbations along the...

<https://arxiv.org/abs/2610.11380>

7 CONF

TypedBench: A Benchmark for Calibration, Framing Sensitivity, and Cost in System One Decision Models

ArXiv CS.AI 50%

arXiv:2610.11392v1 Announce Type: new Abstract: System One models output calibrated probabilities over typed answers such as categorical choices, ordinal levels, or binary outcomes, via a non-generative interface. Software can act on these probabilities through thresholds,...

<https://arxiv.org/abs/2610.11392>

8 CONF

Estimating great expectations under autoregressive language models with potentials

ArXiv CS.AI 50%

arXiv:2610.11399v1 Announce Type: new Abstract: Many applications of language models hinge not on individual samples but on the expectation of a test functional under the model. Estimating such expectations reliably can be computationally expensive. In this paper, we show how to...

<https://arxiv.org/abs/2610.11399>

- 9
CONF

SpikeSSL: A Universal Spike Inference Framework with Dynamics-Informed State-Space Layers

ArXiv CS.AI
50%

arXiv:2610.11456v1 Announce Type: new Abstract: Two-photon calcium imaging is a standard tool for recording large neural populations in vivo, yet inferring spikes accurately across the growing diversity of calcium indicators remains an open problem. Existing supervised methods...

<https://arxiv.org/abs/2610.11456>
- 0
CONF

From Chain-of-Thought to Loops: Non-Autoregressive Latent Reasoning via Looped Transformers

ArXiv CS.AI
50%

arXiv:2610.11472v1 Announce Type: new Abstract: Chain-of-thought (CoT) reasoning often improves language-model performance by giving models additional computation before answering. However, explicit CoT expresses this computation as a sequence of autoregressively generated...

<https://arxiv.org/abs/2610.11472>
- 1
CONF

Memory Type Varies: Empowering LLM Agents for Long-Term Memory with Diverse Strategies

ArXiv CS.AI
50%

arXiv:2610.11573v1 Announce Type: new Abstract: The memory capabilities of Large Language Models (LLMs) have garnered increasing attention recently. Despite great success achieved, existing retrieval-based memory approaches typically overlook the differences between memories and...

<https://arxiv.org/abs/2610.11573>
- 2
CONF

Characterizing Overconfident Failure in LLM-Based Code Generation

ArXiv CS.AI
50%

arXiv:2610.11300v1 Announce Type: cross Abstract: Large language models (LLMs) are increasingly used for automated code generation, but generated programs can appear syntactically plausible while still failing execution-based correctness checks. Existing validation methods, such...

<https://arxiv.org/abs/2610.11300>
- 3
CONF

Where to Adapt Matters: Layer-Selective Fine-Tuning for Capability Retention

ArXiv CS.AI
50%

arXiv:2610.11620v1 Announce Type: new Abstract: Parameter-efficient fine-tuning (PEFT) enables large language models (LLMs) to adapt to specialized tasks, but often at the cost of degrading general capabilities acquired during pretraining. Existing approaches primarily mitigate...

<https://arxiv.org/abs/2610.11620>
- 4
CONF

Evidence-Traceable Dynamic Interviewer Architecture for Expertise-Adaptive Qualitative Interviews Using Local LLMs

ArXiv CS.AI
50%

arXiv:2610.11651v1 Announce Type: new Abstract: Automated interviewers and conversational agents are increasingly used in research, recruitment, customer service, and education. However, many existing systems rely on fixed question sequences and provide limited context-based...

<https://arxiv.org/abs/2610.11651>

5 CONF

Harness Evolution Hits a Ceiling: When Weight Training Should Begin

ArXiv CS.AI 50%

arXiv:2610.11655v1 Announce Type: new Abstract: Improving a long-horizon LLM agent means evolving the harness around a frozen model or training its weights. We let a self-evolving harness make the system stronger first, then cross seed and evolved harnesses with base and trained...

<https://arxiv.org/abs/2610.11655>

6 CONF

Lamarck's Driving School: Discovering Autonomous Driving Training Strategies through Evolutionary Competition

ArXiv CS.AI 50%

arXiv:2610.11662v1 Announce Type: new Abstract: Autonomous driving capabilities depend strongly on the distribution of scenarios encountered during training. Existing methods commonly construct or dynamically adapt training scenario distributions using surrogate criteria such as...

<https://arxiv.org/abs/2610.11662>

7 CONF

Constrained Command-Conditioned Reinforcement Learning with Bandit Strategy Selection in Real-Time Strategy Games

ArXiv CS.AI 50%

arXiv:2610.11663v1 Announce Type: new Abstract: Deep reinforcement learning agents reach strong performance in real-time strategy games but can be brittle against opponents outside their training distribution. Separating strategic command selection from learned unit control...

<https://arxiv.org/abs/2610.11663>

8 CONF

Task-Centric Personalized Federated Fine-Tuning of Language Models

ArXiv CS.AI 50%

arXiv:2604.00050v3 Announce Type: replace-cross Abstract: Federated Learning (FL) has emerged as a promising technique for training language models on distributed and private datasets of diverse tasks. However, aggregating models trained on heterogeneous tasks often degrades the...

<https://arxiv.org/abs/2604.00050>

9 CONF

Scalable AI Uncertainty Quantification via Generalized Laplace Active Subspaces

ArXiv CS.AI 50%

arXiv:2610.11738v1 Announce Type: new Abstract: Reliable uncertainty quantification (UQ) is essential for deploying neural networks in scientific and high-stakes applications, but full Bayesian inference over the network parameters is computationally infeasible. We propose a...

<https://arxiv.org/abs/2610.11738>

0 CONF

Structure Tax: How Structured Output affects LLMs Performance

ArXiv CS.AI 50%

arXiv:2610.12056v1 Announce Type: new Abstract: Deploying large language models in production often requires constraining outputs to structured formats such as JSON or XML, and prior work treats the resulting accuracy loss as an inherent 'structure tax'. We re-examine this claim...

<https://arxiv.org/abs/2610.12056>

1 **CONF****Universal Textual Teaching for LLMs****ArXiv CS.AI 50%**

arXiv:2610.12114v1 Announce Type: new Abstract: Knowledge distillation (KD) transfers knowledge from stronger Teacher models to weaker Student models, but most methods require training the Student parameters, thereby binding the distilled knowledge to a specific architecture and...

<https://arxiv.org/abs/2610.12114>

2 **CONF****Use and Disuse: Intent-Structured Experience Consolidation for Memory and Learning in LLM Agents****ArXiv CS.AI 50%**

arXiv:2610.12124v1 Announce Type: new Abstract: The evolution of Large Language Model agents from single-task execution to long-term autonomous operation highlights the critical challenge of transforming continuous experiences into reusable knowledge. To address this, we propose...

<https://arxiv.org/abs/2610.12124>

3 **CONF****Unlocking the Regulatory Genome by ARGUS: An Evidence-Constrained Agentic Framework for Interpreting Single Nucleotide Variants****ArXiv CS.AI 50%**

arXiv:2610.12281v1 Announce Type: cross Abstract: Over 90% of disease-associated variants from genome-wide association studies fall in noncoding regulatory regions, yet their functional interpretation remains a central open problem in genomic medicine. Large language models...

<https://arxiv.org/abs/2610.12281>

4 **CONF****Learning Implicit Bias in Generative Spaces for Accelerating Protein Dynamics Emulation****ArXiv CS.AI 50%**

arXiv:2606.01833v2 Announce Type: replace-cross Abstract: Generative emulators of protein dynamics produce plausible trajectories at a fraction of the cost of molecular dynamics, but they inherit their training distribution and tend to revisit known states rather than reach rare...

<https://arxiv.org/abs/2606.01833>

5 **CONF****When Has a Bayesian Neural Network Sampled Enough? Adaptive Inference Time with Statistical Guarantees****ArXiv CS.AI 50%**

arXiv:2610.12212v1 Announce Type: new Abstract: Bayesian neural network predictions are commonly approximated using a fixed number of Monte Carlo samples per input, without controlling the resulting error that comes from this finite sample. We propose the use of confidence...

<https://arxiv.org/abs/2610.12212>

6 **CONF****One Word Opens the Gate: The Option-Channel Attack on Typed Decision Models as Agent Guardrails****ArXiv CS.AI 50%**

arXiv:2610.12292v1 Announce Type: new Abstract: A typed decision model reads a piece of text and returns a probability over caller-defined options, each with a short written definition, generating no text. Recent work places these models in agent systems as guardrails: the...

<https://arxiv.org/abs/2610.12292>

7 CONF

Learning Probabilistic Logic Programs with Functional Gradient Guided Language Models

ArXiv CS.AI 50%

arXiv:2610.12303v1 Announce Type: new Abstract: Declarative logic programs offer a powerful and interpretable abstraction for encoding relational structure and neurosymbolic reasoning, by expressing dependencies as weighted compositional rules. However, inducing them from data...

<https://arxiv.org/abs/2610.12303>

8 CONF

Looking Inside LLMs: Small-World Connectivity as a Signature of Reasoning Performance

ArXiv CS.AI 50%

arXiv:2610.12304v1 Announce Type: new Abstract: Understanding large language model (LLM) reasoning requires looking beyond behavioral performance to examine how reasoning ability is reflected in internal organization. Inspired by neuroscience findings linking higher intelligence...

<https://arxiv.org/abs/2610.12304>

9 CONF

Verdict Without the Rule: Diagnosing and Auditing Regulatory Rule Sensitivity in LLM Compliance Systems

ArXiv CS.AI 50%

arXiv:2610.12313v1 Announce Type: new Abstract: Large language model compliance systems are deployed on the assumption that a verdict depends on the regulatory rule it is given. We test this directly across five models and 20 regulatory and platform-policy domains: delete, swap,...

<https://arxiv.org/abs/2610.12313>

0 CONF

Prior or Feedback? What an LLM Uses When Adapting Neural Operators

ArXiv CS.AI 50%

arXiv:2610.12325v1 Announce Type: new Abstract: Do LLM scientific agents rely only on their initial task context, or do they adapt their decisions in response to experimental feedback? We study this question in neural operator adaptation, where a large language model (LLM)...

<https://arxiv.org/abs/2610.12325>

1 CONF

Can AI Agents Learn Their Way to the Top? Evaluating Heuristic Learning in a Long-Running Game Agent Competition

ArXiv CS.AI 50%

arXiv:2610.12341v1 Announce Type: new Abstract: Adversarial games have driven advances from heuristic search to reinforcement learning, yet learning and adapting strategies from limited samples remain challenging. AI agents offer an alternative by turning game experience into...

<https://arxiv.org/abs/2610.12341>

2 CONF

Cited but Not Consulted: A Counterfactual Audit of Legal Chain-of-Thought Faithfulness

ArXiv CS.AI 50%

arXiv:2610.12361v1 Announce Type: new Abstract: Large language models increasingly justify legal decisions by naming the statute or precedent behind a verdict, treated as evidence that the decision follows from it. We test this directly: holding case facts fixed, we substitute...

<https://arxiv.org/abs/2610.12361>

3 CONF

OnTrack: Real-Time Monitoring and Intervention in LLM Agent Trajectories via Streaming Structure-Aware Optimal Transport

ArXiv CS.AI 50%

arXiv:2610.12375v1 Announce Type: new Abstract: Agents are deployed in applications from trip planners and stock trading to IT incident triage. In most cases, LLM agents work autonomously with minimal rule-based safeguarding, leading to cost and safety issues from irreversible...

<https://arxiv.org/abs/2610.12375>

4 CONF

GeoReform: Reflective Formalization Evolution for Multimodal Geometry Problem Solving

ArXiv CS.AI 50%

arXiv:2610.12391v1 Announce Type: new Abstract: Multimodal large language models (MLLMs) often struggle to identify and use geometric relations in diagrams. Recent methods address this challenge by converting geometric entities, relations, and constraints into explicit textual...

<https://arxiv.org/abs/2610.12391>

5 CONF

Searching for "Harmful Refusal": A Psychometric Audit of an AI Safety Benchmark

ArXiv CS.AI 50%

arXiv:2610.12409v1 Announce Type: new Abstract: Safety benchmarks typically report one overall score for a suite of datasets, each of which may target one or more safety-related attributes, so models with similar overall scores can have very different attribute profiles....

<https://arxiv.org/abs/2610.12409>

6 CONF

Ecology of AI Agents: Collaboration Creates a Population Threshold for Takeoff

ArXiv CS.AI 50%

arXiv:2610.12436v1 Announce Type: new Abstract: AI agents can now conduct real-world cyberattacks, scale up capabilities with the number of agents, and collectively pursue misaligned goals to obtain rewards. Together, these factors raise the risk of a population explosion of...

<https://arxiv.org/abs/2610.12436>

7 CONF

On the estimation and validity of AI time horizons---a statistical look at the METR plot

ArXiv CS.AI 50%

arXiv:2610.12466v1 Announce Type: new Abstract: METR's 50% time horizon measures the human completion time of software tasks that an AI solves with 50% probability, allowing AI capabilities to be expressed in interpretable units. On 228 tasks and 26 AIs, we recompute the time...

<https://arxiv.org/abs/2610.12466>

8 CONF

SpectralCache: Accelerating Diffusion-Based World Models via Spectral Feature Caching

ArXiv CS.AI 50%

arXiv:2610.02660v1 Announce Type: cross Abstract: Diffusion-based world models enable high-quality interactive environment generation but suffer from substantial inference overhead due to repeated Transformer evaluations during denoising. Existing caching methods mainly exploit...

<https://arxiv.org/abs/2610.02660>

9 CONF

A Survey on LLM-Integrated Hardware Design Verification

ArXiv CS.AI 50%

arXiv:2610.10580v1 Announce Type: cross Abstract: Large language models (LLMs) are increasingly being integrated into hardware verification to automate specification interpretation, verification-artifact generation, debugging, formal reasoning, and tool orchestration. This...

<https://arxiv.org/abs/2610.10580>

0 CONF

Certified Corruption Budgets: Anytime-Valid Leaderboard Claims under Adaptive Rigging

ArXiv CS.AI 50%

arXiv:2610.10597v1 Announce Type: cross Abstract: Public leaderboards for AI models are read continuously, and attackers can see every published standing. Vote rigging, selective disclosure of private variants, and benchmark contamination can each move a ranking. Existing...

<https://arxiv.org/abs/2610.10597>

1 CONF

From Investigation Failures to Reliable SOC Agents: Understanding and Improving LLM-Based Alert Triage

ArXiv CS.AI 50%

arXiv:2610.10608v1 Announce Type: cross Abstract: Security operations centers (SOCs) must triage large volumes of alerts, most of which are benign, while missed attacks can remain uninvestigated. Tool-using large language model (LLM) agents can retrieve evidence during triage,...

<https://arxiv.org/abs/2610.10608>

2 CONF

TestJack: Should you trust the results in coding benchmarks? Agentic Coding Benchmarks Auditing via Evaluator Evolution

ArXiv CS.AI 50%

arXiv:2610.10619v1 Announce Type: cross Abstract: Large language model (LLM) agents are rapidly reshaping software engineering, accompanied by an explosion of new code benchmarks. Yet nearly all existing benchmarks still rely on the same decades-old criterion: a solution is...

<https://arxiv.org/abs/2610.10619>

3 CONF

When AI Finds Hidden Messages, Does It Report?

ArXiv CS.AI 50%

arXiv:2610.10620v1 Announce Type: cross Abstract: When an assistant encounters a message for another AI, does it tell its user? Four fixed model-provider deployments perform simulated source tasks in 1,280 ordinary-note and 128 enhanced-note sessions. Harmless and harmful...

<https://arxiv.org/abs/2610.10620>

4 CONF

MAP4CS: A Multi-dimensional Data Pruning Framework for Efficient Code Retriever Fine-tuning

ArXiv CS.AI 50%

arXiv:2610.11727v1 Announce Type: cross Abstract: Retrieval-Augmented Generation (RAG) has become a cornerstone in software engineering for enhancing Large Language Models (LLMs) with domain-specific knowledge. However, adapting retrievers to evolving code repositories remains...

<https://arxiv.org/abs/2610.11727>

5 CONF

ARC: A Reasoning Recipe for Robot Foundation Models

ArXiv CS.AI 50%

arXiv:2610.12386v1 Announce Type: cross Abstract: The prevailing approach to improving robot foundation models (RFMs) relies on larger models, more robot demonstrations, and costly training at scale. We show that there exists an effective and efficient complementary approach:...

<https://arxiv.org/abs/2610.12386>

6 CONF

From Log-Odds to Shapley Values: An Explanatory Geometry for the Weighted Naive Bayes Classifier

ArXiv CS.AI 50%

arXiv:2610.10642v1 Announce Type: cross Abstract: This paper studies the construction of an explanatory space for a weighted naive Bayes classifier from the supervised representation induced by the model. We start from the classical supervised distance based on conditional...

<https://arxiv.org/abs/2610.10642>

7 CONF

Masked Generative Motion Planning with Geometry-Guided Token Search

ArXiv CS.AI 50%

arXiv:2610.10646v1 Announce Type: cross Abstract: Generative motion planners typically use learned trajectory priors for initial generation, while leaving test-time repair to local continuous refinement. We introduce Masked Generative Motion Planning (MGMP), which extends the...

<https://arxiv.org/abs/2610.10646>

8 CONF

Beyond the Ergodic Wall: A Discrete Geometric Physics Sandbox for Analysing AI Scaling Limits and Complexity Collapse

ArXiv CS.AI 50%

arXiv:2610.10651v1 Announce Type: cross Abstract: This paper exposes the ergodic ceiling and thermodynamic inefficiency of current deep learning, which converges to a statistical average of historic human knowledge. True semantic novelty requires a path-dependent...

<https://arxiv.org/abs/2610.10651>

9 CONF

Nullify: Null-Space Activation Steering for Training-Free LLM Unlearning

ArXiv CS.AI 50%

arXiv:2610.10655v1 Announce Type: cross Abstract: Large Language Models (LLMs) inevitably internalize substantial amounts of sensitive or private information during pre-training, while LLM unlearning aims to selectively erase specific knowledge to prevent privacy leakage with...

<https://arxiv.org/abs/2610.10655>

00 CONF

Conversational Task Disambiguation over Tabular Data: Leakage-Aware Formulation, Benchmark Suite, and Training

ArXiv CS.AI 50%

arXiv:2610.10740v1 Announce Type: cross Abstract: Conversational task disambiguation over tabular data uses dialogue to resolve missing information about a user's intended task before producing a solution over tables or databases. Existing evaluation and training lack a...

<https://arxiv.org/abs/2610.10740>

01 CONF

A Dual-Hypothesis Reasoning Framework for LLM Guardrails

ArXiv CS.AI 50%

arXiv:2607.17575v3 Announce Type: replace Abstract: We propose ARBITER, a novel LLM guardrail framework that introduces two key ideas: (i) dual-hypothesis reasoning, a reasoning method for LLM guardrails that explicitly considers both safe and unsafe interpretations of a prompt...

<https://arxiv.org/abs/2607.17575>

02 CONF

AI-Mediated Self: How HCI Defines and Relates to the Self

ArXiv CS.AI 50%

arXiv:2610.10770v1 Announce Type: cross Abstract: How might AI alter how we understand and experience the self? This scoping review analyzes 102 papers to examine how the self is defined in the field of human-computer interaction (HCI), how AI-self relationships are...

<https://arxiv.org/abs/2610.10770>

03 CONF

NavGPT-3: Harnessing Context in a Hierarchical Navigation Runtime

ArXiv CS.AI 50%

arXiv:2610.10787v1 Announce Type: cross Abstract: Language models trained with long-horizon agentic reinforcement learning can generalize knowledge through reasoning, express precise actions, and pursue goals over many steps, raising the ceiling on what an embodied agent can...

<https://arxiv.org/abs/2610.10787>

04 CONF

Grammar Concept Annotation at Scale: Deployed Fine-Tuned Small Language Models Outperform Prompted Frontier Models

ArXiv CS.AI 50%

arXiv:2610.10827v1 Announce Type: cross Abstract: Corrective feedback is among the best-evidenced drivers of second-language acquisition, yet corrections delivered during lessons rarely accumulate into an actionable view of grammar mastery. Prompted frontier models can provide...

<https://arxiv.org/abs/2610.10827>

05 CONF

Real Long-Term Memory for AI: A 50-Million-Token Window That Is Faster and Cheaper Than Recompute

ArXiv CS.AI 50%

arXiv:2610.10845v1 Announce Type: cross Abstract: A large language model can only use the text that fits in its context window, and it recomputes its internal key-value (KV) state for a prompt every time the prompt is sent. We test a memory layer, the public package galahad-kv,...

<https://arxiv.org/abs/2610.10845>

06 CONF

RFChipAgent: Multi-Agentic AI Flow for Analog/RF Chip Design

ArXiv CS.AI 50%

arXiv:2610.10858v1 Announce Type: cross Abstract: Analog/RF circuits remain the critical interface between digital computation and the physical world, and emerging standards from Wi-Fi 7 to 6G place stringent demands on them, yet analog/RF design remains one of the most...

<https://arxiv.org/abs/2610.10858>

07 CONF

Rethinking the Tradeoff Between Temporal Encoding and Nonlinear Computation in Spiking Language Models

ArXiv CS.AI 50%

arXiv:2610.10933v1 Announce Type: cross Abstract: Spiking language models face a tradeoff between representing continuous semantic features over short temporal windows and retaining costly nonlinear attention operations. We introduce Spora, which jointly designs spike encodings...

<https://arxiv.org/abs/2610.10933>

08 CONF

iAm.md: Robot Skill Self-Assessment through Agentic Introspection for Unknown Open-Vocabulary Domains

ArXiv CS.AI 50%

arXiv:2610.10962v1 Announce Type: cross Abstract: Agentic AI based on Large Language Model generalization capabilities offers a wide range of potential applications, including planning for embodied tasks. For example, embodied agents based on Foundation models can generate...

<https://arxiv.org/abs/2610.10962>

09 CONF

Why LLM Agents Favor Their Group: Stakes, Observed Norms, and Reputation

ArXiv CS.AI 50%

arXiv:2610.11008v1 Announce Type: cross Abstract: Language-model agents favor their own group because they have watched their members favor each other. The group label alone does little once the decision has a cost; what drives favoritism is observed behavior, and an...

<https://arxiv.org/abs/2610.11008>

10 CONF

NOMOS: Compiling Written Policies into Statically Verified Tool-Call Gates for LLM Agents

ArXiv CS.AI 50%

arXiv:2610.11030v1 Announce Type: cross Abstract: Tool-using LLM agents violate the policies they are deployed to enforce, often silently. Prior defenses hand-write rules, query an LLM verifier per action, or compile policies through heavyweight formal machinery. Naive...

<https://arxiv.org/abs/2610.11030>

11 CONF

Stability-Plasticity Balance via Singular-Vector Selection in LLM Continual Learning

ArXiv CS.AI 50%

arXiv:2610.11076v1 Announce Type: cross Abstract: Domain-specific continual adaptation of LLMs risks catastrophic forgetting, creating a fundamental tension between acquiring new capabilities and preserving those learned during pretraining. PEFT mitigates this problem by...

<https://arxiv.org/abs/2610.11076>

12 CONF

Continuous Ground-Truth Construction and a Recovery Policy for Air--Water Robotic Tracking

ArXiv CS.AI 50%

arXiv:2610.11096v1 Announce Type: cross Abstract: Visual tracking across the air-water interface is challenged by splashes, bubbles, refraction, reflections, and abrupt appearance changes that can temporarily invalidate observations. This setting poses two coupled difficulties...

<https://arxiv.org/abs/2610.11096>

13 CONF

SFT-as-Context Mitigates Forgetting in Supervised Fine-Tuning

ArXiv CS.AI 50%

arXiv:2610.11132v1 Announce Type: cross Abstract: Supervised fine-tuning (SFT) equips large language models (LLMs) with specialized capabilities, but often comes at the cost of forgetting the general capabilities of their parent models (i.e., the pretrained models before...

<https://arxiv.org/abs/2610.11132>

14 CONF

Improving Image-Based Nutrition Estimation Through Multimodal Food-Item Verification and Recovery

ArXiv CS.AI 50%

arXiv:2610.11144v1 Announce Type: cross Abstract: Single-image nutrition estimation can fail silently when visible foods are missed. Even when a food is correctly identified, its proposed region may not support portion estimation. We propose a framework that uses multimodal...

<https://arxiv.org/abs/2610.11144>

15 CONF

Characterizing Statistical Separability in TP-CRIV for Probabilistic AI Models

ArXiv CS.AI 50%

arXiv:2610.11163v1 Announce Type: cross Abstract: Third-party challenge-response identity verification (TP-CRIV) enables an independent verifier to assess whether a claimant possesses a model identical to a remotely deployed model without directly accessing the reference model...

<https://arxiv.org/abs/2610.11163>

16 CONF

PIVOT: Perplexity-Informed KD-to-RL Transition Scheduling for Vertical-Domain Few-Shot Distillation

ArXiv CS.AI 50%

arXiv:2610.11167v1 Announce Type: cross Abstract: Vertical-domain few-shot classification remains challenging for small language models, as limited supervision makes it difficult to acquire domain-specific decision knowledge. On-Policy Distillation (OPD) can improve...

<https://arxiv.org/abs/2610.11167>

17 CONF

RAG-Stress: Probing the Limits of Evidence Reliance in Retrieval-Augmented Generation

ArXiv CS.AI 50%

arXiv:2610.11183v1 Announce Type: cross Abstract: Following retrieved evidence does not guarantee factual correctness: misleading evidence can induce a model to replace an answer it previously gave correctly. Standard accuracy measures obscure this behavior by combining answer...

<https://arxiv.org/abs/2610.11183>

18 CONF

SafeInferCom: Safe Inference-Time Compute via Verifier-Guided Mid-Generation Intervention for Robotic Task Planning

ArXiv CS.AI 50%

arXiv:2610.11223v1 Announce Type: cross Abstract: Large Reasoning Language Models (LRLMs) enable multi-step reasoning for robotic task planning, but continued reasoning can overwrite valid intermediate plans or leave constraint violations unresolved, reducing planning...

<https://arxiv.org/abs/2610.11223>

19 CONF

Neuro-Memory Fuzzy Inference System for Mimicking Human-like Car Following Behavior

ArXiv CS.AI 50%

arXiv:2610.11252v1 Announce Type: cross Abstract: This study presents the Neuro-Memory Fuzzy Inference System (NeMeFIS), a hierarchical machine learning architecture that asymmetrically models acceleration and deceleration in car following behavior by integrating five human...

<https://arxiv.org/abs/2610.11252>

20 CONF

Gated Memory: Admission-Controlled Memory Formation for Conversational AI

ArXiv CS.AI 50%

arXiv:2610.11270v1 Announce Type: cross Abstract: Personalized conversational AI relies on long-term memory systems that extract facts from user utterances and store them in persistent vector stores. Despite progress in retrieval, deduplication, and lifecycle management, the...

<https://arxiv.org/abs/2610.11270>

21 CONF

How to post-train on a surrogate: Envelope sampling mitigates reward hacking

ArXiv CS.AI 50%

arXiv:2610.11281v1 Announce Type: cross Abstract: Large language models (LLMs) are commonly post-trained against LLM judges and other cheap surrogates because the true reward, such as human preference, is too expensive to query at scale. This practice often leads to reward...

<https://arxiv.org/abs/2610.11281>

22 CONF

MetaEncoder: Exploring the Limit of Bi-Encoders for Multimodal System One Decision Making with Natural Language Interface

ArXiv CS.AI 50%

arXiv:2610.11316v1 Announce Type: cross Abstract: System One models output constrained decisions and probability distributions rather than free-form text generation. While prevailing paradigms rely on structured schema objects to encode state, intent, and candidate choices, we...

<https://arxiv.org/abs/2610.11316>

23 CONF

Rethinking Contrastive Loss in CLIP Post-training: A Complementary Framework with Frozen Text Encoder

ArXiv CS.AI 50%

arXiv:2610.11374v1 Announce Type: cross Abstract: CLIP serves as a foundational vision-language model and the de facto vision encoder for downstream VLMs such as LLaVA. Post-training offers a lightweight route to refine CLIP, but recent work argues that the standard contrastive...

<https://arxiv.org/abs/2610.11374>

24 CONF

CRISP: Fixing Flying Pixels in Latent LiDAR Generation via Diffusion Decoding

ArXiv CS.AI 50%

arXiv:2610.11376v1 Announce Type: cross Abstract: Latent LiDAR pipelines suffer from flying pixels: convolutional VAEs blur sharp radial depth discontinuities, yielding edge depths that back-project to points floating between surfaces. We identify this as a major, directly...

<https://arxiv.org/abs/2610.11376>

25 CONF

Rewiring Semantics, Dynamics, and Control: A Simple yet Effective Action-Centric Tri-Stream Transformer

ArXiv CS.AI 50%

arXiv:2610.11416v1 Announce Type: cross Abstract: Vision-Language-Action (VLA) models have emerged as a prominent framework for complex robotic manipulation, building on the strong semantic understanding of pretrained Vision-Language Models (VLMs). However, such VLM backbones...

<https://arxiv.org/abs/2610.11416>

26 CONF

Missing Modality-Aware Calibration for Trustworthy Brain Tumor Segmentation

ArXiv CS.AI 50%

arXiv:2610.11419v1 Announce Type: cross Abstract: Multimodal brain tumor segmentation typically leverages multiple MRI modalities, yet incomplete modality acquisition is common in clinical practice due to protocol heterogeneity and scan failures. Although recent methods maintain...

<https://arxiv.org/abs/2610.11419>

27 CONF

Closed-loop evaluation of LLM agents for embedded software development

ArXiv CS.AI 50%

arXiv:2610.11447v1 Announce Type: cross Abstract: Large language models (LLMs) are increasingly deployed as coding agents that edit files, run builds and tests, inspect execution results, and repair software iteratively. Embedded firmware is a demanding target because...

<https://arxiv.org/abs/2610.11447>

28 CONF

Zatom-2: Multitask Pretraining on Atomistic Data for Generative Modeling across Domains

ArXiv CS.AI 50%

arXiv:2610.11454v1 Announce Type: cross Abstract: Unified atomistic modeling has the potential to accelerate discovery in chemistry, materials science, and biology by bridging data-rich chemical domains and data-scarce biological contexts. However, existing generative approaches...

<https://arxiv.org/abs/2610.11454>

29 CONF

Generative Adversarial Loops

ArXiv CS.AI 50%

arXiv:2610.11458v1 Announce Type: cross Abstract: AI research progress can be viewed as the interaction between two processes: benchmark creation and method discovery. Historically, both were driven by human intelligence. However, recent advances in AI have accelerated automated...

<https://arxiv.org/abs/2610.11458>

30 CONF

ProtoSemImage: Image-Valued Prototypes with Deformable Row Alignment for Interpretable Document Classification

ArXiv CS.AI 50%

arXiv:2610.11460v1 Announce Type: cross Abstract: Prototypes in classification models are almost always vectors, and a vector has no readable form. This paper asks what happens when a prototype is an image. Documents give the question a natural form, because a document can be...

<https://arxiv.org/abs/2610.11460>

31 CONF

Conditional Residual Prediction: Improving Autoregressive Video Diffusion without a Bidirectional Teacher

ArXiv CS.AI 50%

arXiv:2610.11479v1 Announce Type: cross Abstract: Causal video diffusion models generate video autoregressively, which suits streaming, interactive, and long-video generation. Under standard training, however, they often yield lower generation quality than bidirectional models...

<https://arxiv.org/abs/2610.11479>

32 CONF

Evaluating Local Language Model Agents for Reproducible Data Engineering: An Empirical Software Engineering Study of Mobility Workflows

ArXiv CS.AI 50%

arXiv:2610.11482v1 Announce Type: cross Abstract: Context: Large language model (LLM) agents are increasingly used as software and data-engineering assistants, yet evidence about locally deployable open-weight agents remains limited. Existing evaluations often emphasize textual...

<https://arxiv.org/abs/2610.11482>

33 CONF

Constitutional Gating and Deterministic Recovery for Multi-Agent LLM Negotiation: Ablations Against a Stateful Adversarial Gatekeeper

ArXiv CS.AI 50%

arXiv:2610.11542v1 Announce Type: cross Abstract: Multi-agent LLM systems negotiating with a stateful counterpart waste model calls in three ways: polite loops that never meet the counterpart's hidden acceptance condition, malformed outputs that trigger retries, and compliance...

<https://arxiv.org/abs/2610.11542>

34 CONF

Learning to Orchestrate Evolutionary Search: Progression-Aware Deep Reinforcement Learning for Dynamic DE-CMA-ES Coordination in Optimization and Structural Model Updating

ArXiv CS.AI 50%

arXiv:2610.11546v1 Announce Type: cross Abstract: Solving high-dimensional structural model updating problems requires an algorithm capable of navigating complex, non-convex landscapes with correlated parameters. Existing hybrid evolutionary algorithms typically rely on static...

<https://arxiv.org/abs/2610.11546>

35 CONF

Large Language Model Turnover Undermines Screening for Artificial Intelligence-Assisted Scientific Writing

ArXiv CS.AI 50%

arXiv:2610.11599v1 Announce Type: cross Abstract: Journals and conferences have begun to screen submitted manuscripts for text written using large language models (LLMs). The reliability of this screening rests on benchmark evaluations against a fixed set of LLM versions, while...

<https://arxiv.org/abs/2610.11599>

36 CONF

Neural Networks for Temporal Pattern Recognition and Dynamic Arm Gesture Speed Estimation for Robot Control

ArXiv CS.AI 50%

arXiv:2610.11631v1 Announce Type: cross Abstract: Deploying intelligent robotic systems that interact with humans through gestures requires neural networks capable of recognizing diverse temporal patterns. We present a systematic benchmark of ten abstract sequential tasks--five...

<https://arxiv.org/abs/2610.11631>

37 CONF

DEX: Digit-Level Early Exit for Energy-Efficient MSDF Neural Network Inference

ArXiv CS.AI 50%

arXiv:2610.11748v1 Announce Type: cross Abstract: U-Net inference for brain-tumor segmentation requires billions of multiply-accumulate operations, motivating hardware that can reduce computation dynamically rather than relying only on fixed precision or static model...

<https://arxiv.org/abs/2610.11748>

38 CONF

Narrow and Deep: An Ontology Tower as the Knowledge of an LLM Agent for an Industrial Equipment System

ArXiv CS.AI 50%

arXiv:2610.11768v1 Announce Type: cross Abstract: Large language model (LLM) agents are beginning to operate industrial energy equipment, and what they get right depends on what they are told about the plant. Established building ontologies name many kinds of points across many...

<https://arxiv.org/abs/2610.11768>

39 CONF

From Suppression to Repair: Mitigating Object Hallucination in Large Vision-Language Models via Localized Distribution Alignment

ArXiv CS.AI 50%

arXiv:2610.11826v1 Announce Type: cross Abstract: Object hallucination remains a major obstacle for large vision-language models (LVLMs) to generate reliable content. An intuitive mitigation strategy is to suppress hallucination-related components in hidden representations....

<https://arxiv.org/abs/2610.11826>

40 CONF

Forms of LLM-Integrated Applications from LLM-Chats to Autonomous AI Agent System

ArXiv CS.AI 50%

arXiv:2610.11899v1 Announce Type: cross Abstract: Large language models (LLMs) are increasingly embedded as components in software systems, marketed under labels such as chatbot, copilot, retrieval-augmented generation, workflow, coding agent and AI agent. Whether these labels...

<https://arxiv.org/abs/2610.11899>

41 CONF

Neural Decoding as Cognitive Inference

ArXiv CS.AI 50%

arXiv:2610.11923v1 Announce Type: cross Abstract: The brain maintains stable cognition despite continuously changing neural activity. How to extract stable cognitive states from variable neural observations remains a central problem in neural decoding. Existing neural decoding...

<https://arxiv.org/abs/2610.11923>

42 CONF

Can LLMs Fix It Without Code? Toward Automated Verification of No-Code Bug Fixes

ArXiv CS.AI 50%

arXiv:2610.11963v1 Announce Type: cross Abstract: A no-code fix resolves an invalid bug report by directing the user to change a setting, update to a version where the problem is already fixed, or adjust their workflow. Manually verifying whether a proposed no-code fix resolves...

<https://arxiv.org/abs/2610.11963>

43 CONF

Neural Network Verification for Deep Joint Source-Channel Coding

ArXiv CS.AI 50%

arXiv:2610.11994v1 Announce Type: cross Abstract: Deep joint source-channel coding (DeepJSCC) transmits data end-to-end over wireless channels using a neural encoder-decoder, but reconstruction quality can degrade sharply under adversarial perturbations and channel disturbances;...

<https://arxiv.org/abs/2610.11994>

44 CONF

The Polytopal Neural Network

ArXiv CS.AI 50%

arXiv:2610.12004v1 Announce Type: cross Abstract: Understanding how deep neural networks process information remains a central challenge. Existing interpretability methods often compromise structural fidelity, rely on prespecified corpora, or explain models post-hoc. We propose...

<https://arxiv.org/abs/2610.12004>

45 CONF

MPGE: A Multi-Perspective Graph Explainer for Molecular Classification Explanation

ArXiv CS.AI 50%

arXiv:2610.12039v1 Announce Type: cross Abstract: Graph neural networks (GNNs) predict molecular properties from chemical graph data, but predictive accuracy does not explain how graph information supports an individual decision. A compact prediction-preserving rationale does...

<https://arxiv.org/abs/2610.12039>

46 CONF

MAST: Motif-Augmented Diffusion with Search Tree for Spectroscopic Molecular Structure Elucidation

ArXiv CS.AI 50%

arXiv:2610.12067v1 Announce Type: cross Abstract: Elucidating molecular structures from spectra is a foundational problem in chemical and materials characterization, yet remains challenging due to spectral ambiguity and the vast molecular space. Although recent diffusion-based...

<https://arxiv.org/abs/2610.12067>

47 CONF

All Verdicts are Not Equal: Rethinking LLM Judge Reliability

ArXiv CS.AI 50%

arXiv:2610.12083v1 Announce Type: cross Abstract: LLM-as-a-Judge is the standard paradigm for NLP evaluation, yet its systemic reliability remains poorly understood despite being widely treated as a deterministic ground truth. We present a comprehensive reliability audit,...

<https://arxiv.org/abs/2610.12083>

48 CONF

Using Weisfeiler-Leman Features for Algorithm Selection in Constraint Optimisation

ArXiv CS.AI 50%

arXiv:2610.12119v1 Announce Type: cross Abstract: Algorithm Selection is essential for efficient Constraint Programming. Over the years, many algorithm selectors based on machine learning methods have been successfully applied, yet traditional feature extraction methods often...

<https://arxiv.org/abs/2610.12119>

49 CONF

AI-Based On-Board Maritime Object Detection for Earth Observation Payload Data Reduction on Versal Embedded Hardware

ArXiv CS.AI 50%

arXiv:2610.12182v1 Announce Type: cross Abstract: Very-high-resolution Earth-observation satellites acquire more data than they can store and downlink, while in maritime surveillance the vessels cover a tiny fraction of each scene. We study onboard vessel detection as a way to...

<https://arxiv.org/abs/2610.12182>

50 CONF

ReSI: Recursive Safety Improvement toward Resistant and Resilient AI

ArXiv CS.AI 50%

arXiv:2610.12233v1 Announce Type: cross Abstract: Recursive self-improvement, the participation of AI systems in improving their own capabilities, is beginning to move from theoretical prospect to practice, posing both challenges and opportunities for safety alignment. Models...

<https://arxiv.org/abs/2610.12233>

51 CONF

Batch Before You Lift: Scalable Topological Deep Learning on Large Graphs

ArXiv CS.AI 50%

arXiv:2610.12247v1 Announce Type: cross Abstract: Topological Deep Learning extends graph-based learning to higher-order domains, such as hypergraphs, cellular, and simplicial complexes. These domains are typically constructed from patterns in an input graph through a process of...

<https://arxiv.org/abs/2610.12247>

52 CONF

How Much Audio Is Left In An Embedding? An Inversion Audit Of Audio Encoders

ArXiv CS.AI 50%

arXiv:2610.12250v1 Announce Type: cross Abstract: Pretrained audio encoders are reused for downstream tasks that are often unknown when the encoder is trained, so their usefulness depends partly on which signal properties survive the pretext objective. We study this retained...

<https://arxiv.org/abs/2610.12250>

53 CONF

Machine Learning Meets High-Energy Nuclear Physics: From Pattern Recognition to Physics-Integrated Discovery

ArXiv CS.AI 50%

arXiv:2610.12293v1 Announce Type: cross Abstract: Machine learning (ML) in high-energy nuclear physics (HENP) is entering a new stage in which physical knowledge is incorporated more directly into data analysis, simulation, and physics inference. This mini-review focuses on...

<https://arxiv.org/abs/2610.12293>

54 CONF

Multi-Agent Egocentric World Model with Fine-Grained Embodied Interaction

ArXiv CS.AI 50%

arXiv:2610.12299v1 Announce Type: cross Abstract: Egocentric world models predict first-person observations conditioned on an agent's actions, but most focus on a single agent. Real embodied settings often involve multiple agents that act and interact within a shared...

<https://arxiv.org/abs/2610.12299>

55 CONF

ContiLNN: Mitigating Slice Sampling Discontinuity with Liquid Neural Networks for Medical Image Restoration

ArXiv CS.AI 50%

arXiv:2610.12337v1 Announce Type: cross Abstract: Anatomical continuity provides complementary information for medical image restoration, but its use requires accounting for local anatomy and variations in slice sampling. We introduce ContiLNN, which augments two-dimensional...

<https://arxiv.org/abs/2610.12337>

56 CONF

LeWAM: A JEPA World Action Model with Diffusion-Steering-Based MPC

ArXiv CS.AI 50%

arXiv:2610.12407v1 Announce Type: cross Abstract: World action models (WAMs) predict actions and future observations, typically from a reconstruction-based representation that carries noisy, redundant information which can complicate downstream predictions. We introduce LeWAM, a...

<https://arxiv.org/abs/2610.12407>

57 CONF

From Reactive Containment to Proactive Assurance: Lessons from OpenAI, Anthropic, and Google Agent Security Incidents

ArXiv CS.AI 50%

arXiv:2610.12463v1 Announce Type: cross Abstract: In 2026, cybersecurity evaluations involving OpenAI, Anthropic, and Google agents reached real systems outside their authorized test scope. The paths were different. OpenAI agents exploited research infrastructure, coordinated...

<https://arxiv.org/abs/2610.12463>

58 CONF

Machine learning and information theory concepts towards an AI Mathematician

ArXiv CS.AI 50%

arXiv:2403.04571v2 Announce Type: replace Abstract: The current state-of-the-art in artificial intelligence is impressive, especially in terms of mastery of language, but not so much in terms of mathematical reasoning. What could be missing? Can we learn something useful about...

<https://arxiv.org/abs/2403.04571>

59 CONF

Thought-Like-Pro: Enhancing Reasoning of Large Language Models through Self-Bootstrapped Prolog-based Chain-of-Thought

ArXiv CS.AI 50%

arXiv:2407.14562v3 Announce Type: replace Abstract: Large language models have demonstrated remarkable capabilities as general-purpose assistants, excelling in a wide range of reasoning tasks and supporting various aspects of daily web usage. This achievement represents a...

<https://arxiv.org/abs/2407.14562>

60 CONF

MissMAC-Bench: Building Solid Benchmark for Missing Modality Issue in Robust Multimodal Affective Computing

ArXiv CS.AI 50%

arXiv:2602.00811v2 Announce Type: replace Abstract: Current Multimodal Affective Computing (MAC) systems heavily rely on the completeness of multiple modalities to accurately understand human's affective state. However, in real-world scenarios, the availability of modality data...

<https://arxiv.org/abs/2602.00811>

61 CONF

Beyond Preset Identities: Selective Stance Accommodation and Interaction Reorganisation in Generative Agent Societies

ArXiv CS.AI 50%

arXiv:2603.23406v3 Announce Type: replace Abstract: Generative agent societies simulate people with assigned roles, preferences and relationships. As agents exchange arguments and choose partners, they can revise their positions and reorganise discussion. Understanding these...

<https://arxiv.org/abs/2603.23406>

62 CONF

Soft Tournament Equilibrium: Differentiable Set-Valued Inference for Non-Transitive Pairwise Comparisons

ArXiv CS.AI 50%

arXiv:2604.04328v4 Announce Type: replace Abstract: Soft Tournament Equilibrium (STE) is a differentiable layer for Top-Cycle (TC) and Uncovered-Set (UC) inference from reciprocal pair probabilities. Normalized log-sum-exp reachability and covering give smooth scores with...

<https://arxiv.org/abs/2604.04328>

63 CONF

LiteGUI: Lightweight GUI Agents via Multi-Solution Guided Distillation and Dual-Level Reinforcement Learning

ArXiv CS.AI 50%

arXiv:2605.07505v2 Announce Type: replace Abstract: We present LiteGUI, a new framework for building lightweight GUI agents. GUI interaction poses unique challenges due to the long-horizon nature of complex tasks and the existence of multiple valid interaction paths, which are...

<https://arxiv.org/abs/2605.07505>

64 CONF

Deconstructing Off-Policy Ratios: Entropy-Normalized Trust Regions for Asynchronous Reinforcement Learning

ArXiv CS.AI 50%

arXiv:2607.22186v5 Announce Type: replace Abstract: Asynchronous reinforcement learning (RL) accelerates large language model (LLM) post-training by overlapping rollout generation with policy optimization, but the resulting stale, off-policy data destabilizes optimization and...

<https://arxiv.org/abs/2607.22186>

65 CONF

RealCompanion: Benchmarking Human Understanding from Reasoning over Longitudinal Real-World Conversations

ArXiv CS.AI 50%

arXiv:2610.01780v3 Announce Type: replace Abstract: An AI companion that talks with someone for months should come to understand them. It should know who they are, remember what they said, and recognize when something from the past matters now. Testing this needs real...

<https://arxiv.org/abs/2610.01780>

66 CONF

Mooncake: A KVCache-centric Disaggregated Architecture for LLM Serving

ArXiv CS.AI 50%

arXiv:2407.00079v5 Announce Type: replace-cross Abstract: Mooncake is the serving platform for Kimi, a leading LLM service provided by Moonshot AI. It features a KVCache-centric disaggregated architecture that separates the prefill and decoding clusters. It also leverages the...

<https://arxiv.org/abs/2407.00079>

67 CONF

An Investigation of Robustness of LLMs in Mathematical Reasoning: Benchmarking with Mathematically-Equivalent Transformation of Advanced Mathematical Problems

ArXiv CS.AI 50%

arXiv:2508.08833v4 Announce Type: replace-cross Abstract: Frontier large language models (LLMs) now reach near-ceiling accuracy on standard mathematical-reasoning benchmarks and gold-medal-level performance at the International Mathematical Olympiad. As these benchmarks saturate...

<https://arxiv.org/abs/2508.08833>

68 CONF

Trust, Don't Trust, or Flip: Robust Preference-Based Reinforcement Learning with Multi-Expert Feedback

ArXiv CS.AI 50%

arXiv:2601.18751v2 Announce Type: replace-cross Abstract: Preference-based reinforcement learning (PBRL) offers a promising alternative to explicit reward engineering by learning from pairwise trajectory comparisons. However, real-world preference data often comes from...

<https://arxiv.org/abs/2601.18751>

69 CONF

FreshMem: Brain-Inspired Frequency-Space Hybrid Memory for Streaming Video Understanding

ArXiv CS.AI 50%

arXiv:2602.01683v2 Announce Type: replace-cross Abstract: Transitioning Multimodal Large Language Models (MLLMs) from offline to online streaming video understanding is essential for continuous perception. However, existing methods lack flexible adaptivity, leading to...

<https://arxiv.org/abs/2602.01683>

70 CONF

Rethinking Latency Denial-of-Service: Attacking the LLM Serving Framework, Not the Model

ArXiv CS.AI 50%

arXiv:2602.07878v2 Announce Type: replace-cross Abstract: LLM inference is inherently expensive, even a modest slowdown can translate into substantial operating costs and severe availability risks. Recently, a growing body of research known as latency attacks focuses on crafting...

<https://arxiv.org/abs/2602.07878>

71 CONF

An Autonomous AI Security Agent for Banking: Multi-Vector Fraud and AML Detection Across Retail and Corporate Accounts

ArXiv CS.AI 50%

arXiv:2606.17555v4 Announce Type: replace-cross Abstract: Banks face two threat families with fundamentally different detection requirements: signature-based fraud (card-not-present attacks, account takeover, ATM cloning) and behavioural financial crime (structuring, layering,...

<https://arxiv.org/abs/2606.17555>

72 CONF

RoboTalk: Learning Multi-Robot Communication and Coordination from Multimodal Demonstrations

ArXiv CS.AI 50%

arXiv:2609.23997v2 Announce Type: replace-cross Abstract: Multi-robot collaboration could enable more efficient and scalable solutions to complex robotic tasks, but collaboration under partial observability remains challenging. Natural-language communication offers a promising...

<https://arxiv.org/abs/2609.23997>

73 CONF

Pretext: Defeating Malicious Skill Detection Frameworks for AI Agents

ArXiv CS.AI 50%

arXiv:2609.39607v2 Announce Type: replace-cross Abstract: Skills extend an agent's capabilities by injecting instructions and information into the context, and are widely used by agents such as OpenClaw and Claude Code. Prior work shows third-party marketplaces host malicious...

<https://arxiv.org/abs/2609.39607>

74 CONF

The Score Is Not the Structure: Brain Alignment and Cross-Lingual Transfer

ArXiv CS.AI 50%

arXiv:2610.03827v2 Announce Type: replace-cross Abstract: Similarity scores are often offered as evidence that a model shares structure with the brain or across languages. We ask what such a score reads when that structure is removed, or when the instrument measuring it does not...

<https://arxiv.org/abs/2610.03827>

75 CONF

Just for FUNS: LLM-Guided Spatio-Temporal Graph Node Generation for Forecasting Unobserved Node States

ArXiv CS.AI 50%

arXiv:2610.08818v2 Announce Type: replace-cross Abstract: Spatio-temporal forecasting is a cornerstone of logistics, urban planning, and intelligent transportation systems. However, constrained by deployment costs and maintenance resources, sensor networks often lack...

<https://arxiv.org/abs/2610.08818>

76 CONF

Leveraging LLM-Generated Explanations for Detecting Emotionally Rewritten Fake News

ArXiv CS.AI 50%

arXiv:2610.08835v2 Announce Type: replace-cross Abstract: The spread of fake news may cause severe social consequences. Existing fake news detection methods mainly focus on stylistic variations or incorporate external information such as explanations. However, news articles are...

<https://arxiv.org/abs/2610.08835>

77 CONF

Shared Geometry As A Rosetta Stone: Cross-Modal Alignment Without Paired Data

ArXiv CS.AI 50%

arXiv:2610.09411v2 Announce Type: replace-cross Abstract: Multimodal representations enable zero-shot classification and retrieval, but aligning independently trained models usually requires large amounts of paired data. Yet, the Platonic Representation Hypothesis suggests that...

<https://arxiv.org/abs/2610.09411>

78 CONF

How Do LLMs Change Predictions Under Negation?

ArXiv CS.AI 50%

arXiv:2610.09571v2 Announce Type: replace-cross Abstract: Negation is an essential feature of human language, yet large language models (LLMs) remain unreliable in processing it. We evaluate recent open-source and closed-source LLMs on our negation benchmark and find that, in...

<https://arxiv.org/abs/2610.09571>

79 CONF

Defensive Sufficiency in a Stackelberg Model of AI Security

ArXiv CS.AI 50%

arXiv:2610.09892v2 Announce Type: replace-cross Abstract: Feedback from automated testing, human red teaming, and incident response can strengthen an AI system's defenses when discovered failures lead to effective repairs. We study when this feedback process provides sufficient...
<https://arxiv.org/abs/2610.09892>