NeoSignal

AI Intelligence Daily · Multi-Source · Cross-Verified

27 October 2026

245 verified · 464 raw · 5 sources

245	3	242	0
TOTAL	VERIFIED	CONFIRMED	EMERGING

Pipeline

Scrape HackerNews (Show/Top/New), Reddit (r/artificial, r/MachineLearning, r/singularity, r/LocalLLaMA), TechCrunch AI, VentureBeat AI, MIT Tech Review, The Verge AI, Wired AI, ArXiv CS.AI

Filter 40+ AI keyword patterns — LLMs, alignment, safety, research, tooling

Dedup Title-similarity (45% threshold) merges cross-source variants

Score Authenticity = 0.5 base + 0.3 per extra source + diversity + HN score bonuses

Rank Sorted by authenticity. Below 0.25 threshold are dropped.

Authenticity Tiers

VERIFIED	**0.80 - 1.00**	3+ independent sources including media
CONFIRMED	**0.50 - 0.79**	2 sources or 1 high-quality media
EMERGING	**0.25 - 0.49**	Single-source; passes AI keyword filter

Source Breakdown

Source	Articles	Category
ArXiv CS.AI	226	Media
HackerNews	7	Community
The Verge AI	6	Media
TechCrunch AI	5	Media
MIT Tech Review	4	Media

VERIFIED INTELLIGENCE — 3 stories

VERI

Show HN: Autorouter Benchmark on Various PCBs

HackerNews +1 source 90% HN 2

arXiv:2610.29968v2 Announce Type: replace-cross Abstract: As LLMs expand their language coverage, evaluating how well they understand and represent diverse cultures becomes increasingly important. However, many benchmarks remain English-centric, with limited cultural coverage...

https://github.com/Koppar-AI/autorouter-bench

VERI

Robust and Fine-Grained Detection of AI Generated Texts

ArXiv CS.AI +1 source 80%

arXiv:2504.11952v4 Announce Type: replace-cross Abstract: An ideal detection system for machine generated content is supposed to work well on any generator as many more advanced LLMs come into existence day by day. Existing systems often struggle with accurately identifying...

https://arxiv.org/abs/2504.11952

VERI

Anthropic is burning gas to feed the AI boom

The Verge AI +1 source 80%

As it prepares to go public, Anthropic's S-1 filing makes generative AI's methane dependency glaringly obvious. "Competition for access to electrical supply and data center locations with adequate power availability is intense." In documents filed with the US Securities and...

https://www.theverge.com/ai-artificial-intelligence/1018994/anthropic-ipo-ai-data-cente...

CONFIRMED SIGNALS — 242 stories

CONF

Show HN: Hob - Open-source, on device alternative to Claude Cowork and ChatGPT

HackerNews 60% HN 145

Hey HN, I'm Christian and this is a new OSS project I've been noodling on for the last couple of weeks. Hob is a free, MIT-licensed desktop app for Mac, Windows and Linux that works on your files the way the Claude and ChatGPT desktop apps do, except the model runs on your own...

https://github.com/FrigadeHQ/hob

CONF

Show HN: agent-usage-guard – a Claude Code mod to protect your usage window

HackerNews 50% HN 28

No summary — see link for full article.

https://github.com/rbartoli/agent-usage-guard

CONF

Show HN: AgentTrust – Zero-trust security and human approvals for AI agents

HackerNews 50% HN 14

No summary — see link for full article.

https://github.com/SRN0129/Agenttrust

CONF

Show HN: UnicornFork a Chess AI That Is Fact Checked by Stockfish

HackerNews 50% HN 3

No summary — see link for full article.

https://www.unicornfork.com/

CONF

Show HN: Claude Plays the Piano from Braille

HackerNews 50% HN 2

For the longest time, I've wanted MIDI files created from scans of sheet music to listen to. I can't play the piano myself, but I've been around music since I was a child. The existing software almost never sounded good, never expressive. This is an attempt to use an LLM to...

https://github.com/chris-mccoy/gottschalk-tarantelle

CONF

Show HN: Arc-decide – an MIT 0.4B decision model for RAG checks that runs on CPU

HackerNews 50% HN 2

No summary — see link for full article.

https://docs.docsgpt.cloud/Models/decision-model

0 CONF

Anthropic’s IPO filing reveals its strategy for keeping enterprise customers hooked

TechCrunch AI 50%

Anthropic’s S-1 shows more than half of its $100,000-plus customers now use both the Claude Platform and a first-party app, a cross-selling playbook reminiscent of Amazon that’s making enterprise customers stickier.

https://techcrunch.com/2026/10/26/anthropics-ipo-filing-reveals-its-strategy-for-keepin...

1 CONF

Anthropic officially files for its public debut

TechCrunch AI 50%

Anthropic has officially filed for an IPO, revealing soaring revenue, mounting losses, major enterprise adoption, and an unusual governance structure designed to preserve founder control and its AI safety mission.

https://techcrunch.com/2026/10/26/anthropic-officially-files-for-its-public-debut/

2 CONF

Will AI expose the rickety infrastructure of the internet?

TechCrunch AI 50%

The moment you realize the internet is "held together by glue and dreams."

https://techcrunch.com/2026/10/24/will-ai-expose-the-rickety-infrastructure-of-the-inte...

3 CONF

OpenAI says it will no longer train its AI models on the live internet, citing ‘significant risks’

TechCrunch AI 50%

OpenAI will cut frontier models off from the internet weeks after one of its engineers said that solution wouldn't work.

https://techcrunch.com/2026/10/23/openai-says-it-will-no-longer-train-its-ai-models-on-...

4 CONF

How solid-state transformers could help power the AI boom

MIT Tech Review 50%

Transformers are a crucial but nearly invisible part of our grid. The devices help efficiently move power across long distances, but they’re complicated and in short supply. And given that their basic design hasn’t changed much for about a century, they could use an upgrade. As...

https://www.technologyreview.com/2026/10/27/1147684/solid-state-transformers/

5 CONF

Will the real Anthropic please stand up

MIT Tech Review 50%

This story originally appeared in The Algorithm, our weekly newsletter on AI. To get stories like this in your inbox first, sign up here. After weeks of frenzied speculation, Anthropic has filed its IPO prospectus. The 335-page document (100 pages or so are focused on risks)...

https://www.technologyreview.com/2026/10/26/1147672/will-the-real-anthropic-please-stan...

6 CONF

The Download: everyday AI harms and the race to nuclear spaceflight

MIT Tech Review 50%

This is today’s edition of The Download, our weekday newsletter that provides a daily dose of what’s going on in the world of technology. Opinion: The focus on AI safety disasters ignores everyday harms —Sam Hiner, cofounder of the Young People’s Alliance, and Emma Leiken,...

https://www.technologyreview.com/2026/10/26/1147650/the-download-everyday-ai-harms-nucl...

7 CONF

The focus on AI safety disasters ignores everyday harms

MIT Tech Review 50%

Consider what passes for success in technology policy. An AI company promises that its chatbot will not encourage suicide. A social media platform makes it more difficult for children to buy fentanyl. These are important protections, and they should have existed all along. But...

https://www.technologyreview.com/2026/10/26/1147645/the-focus-on-ai-safety-disasters-ig...

8 CONF

OpenAI announces three big science discoveries and a new way to work with academics

The Verge AI 50%

Less than a month after its mass release of hundreds of manuscripts that it said contained solutions to long-standing mathematics questions, OpenAI is back with three more "major" discoveries in math, computer science, and physics. What's different this time is the choice to...

https://www.theverge.com/ai-artificial-intelligence/1019143/openai-eth-kaseya-science-m...

9 CONF

The RealReal won’t let AI agents hog designer clothing

The Verge AI 50%

The RealReal, a secondhand luxury fashion consignment retailer, said it's updating its bot defense system, effectively making it harder for AI agents to shop on its site. The change comes days after The Verge reported that an AI agent could be used to take advantage of The...

https://www.theverge.com/news/1018713/meta-muse-agentic-shopping-ai-the-realreal

0 CONF

Five weird details from Anthropic’s IPO filing

The Verge AI 50%

Anthropic has finally filed its S-1 prospectus with the SEC, a document designed to give investors a look under the hood of its business, from how much its executives are making and plans for spending its IPO billions to the risks that could derail its ambitions. At more than...

https://www.theverge.com/ai-artificial-intelligence/1018738/anthropic-ipo-s1-filing-wei...

1 CONF

Anthropic’s IPO filing lays out a plan to ‘revolutionize every area of human endeavor’ with AI

The Verge AI 50%

Anthropic has filed its S-1 prospectus with the Securities and Exchange Commission, preparing the company behind Claude to beat OpenAI to the punch with an IPO. In a filing with the SEC, Anthropic said it plans to list on the Nasdaq exchange under the ticker symbol "ANT." The AI...

https://www.theverge.com/ai-artificial-intelligence/1018476/anthropic-ipo-details-s1-fi...

2 CONF

Can AI help solve the COBOL code crisis?

The Verge AI 50%

What do the IRS, banks, airlines, and government infrastructure around the world have in common? They're all running on systems propped up by antique computer code that nobody dares touch, lest everything come crashing down. OpenAI thinks it may have a solution. On October 19th,...

https://www.theverge.com/ai-artificial-intelligence/1018434/cobol-migration-legacy-code...

3 CONF

Greedy Local Learning for Language Model Pretraining: Gaps and Objective Design

ArXiv CS.AI 50%

arXiv:2610.04867v2 Announce Type: replace-cross Abstract: Greedy block-wise local learning splits a network into gradient-isolated blocks trained by local auxiliary losses, deleting the backward pass between blocks: inter-stage communication becomes forward-only and every block...

https://arxiv.org/abs/2610.04867

4 CONF

NeSySpec: Neuro-Symbolic Technical Intelligence for Auditing Engineering Specifications with an Application to Aircraft

ArXiv CS.AI 50%

arXiv:2610.35403v1 Announce Type: new Abstract: Engineers who audit a declared specification of an engineered system need to know which physical constraints it respects, which it violates, which field was altered and how to correct it. NeSySpec is an integrated and interpretable...

https://arxiv.org/abs/2610.35403

5 CONF

Conformance as a Distribution: Delegation Contracts, Error Budgets, and the Allocation of Evaluation Authority in Use-Based AI Regulation

ArXiv CS.AI 50%

arXiv:2610.33112v1 Announce Type: new Abstract: Use-based AI regulation assigns obligations by deployment risk but lacks an evidentiary object to carry them. A benchmark score is a property of a model; conformance of a deployment is a failure rate under a declared distribution...

https://arxiv.org/abs/2610.33112

6 CONF

AOS-Bench: Can LLM Agents Manage Dynamic Workloads?

ArXiv CS.AI 50%

arXiv:2610.33178v1 Announce Type: new Abstract: Coding agents are moving from handling one request at a time to serving as persistent assistants to which users hand a stream of work: a build, a data-processing job and a bug fix may all be pending on the same machine at once....

https://arxiv.org/abs/2610.33178

7 CONF

Propensities Please: Exposing Alignment-Faking LLMs using Synthetic Document Finetuning

ArXiv CS.AI 50%

arXiv:2610.33180v1 Announce Type: new Abstract: Alignment evaluations aim to assess whether large language models would engage in dangerous misaligned behaviour. Misaligned and evaluation-aware models might engage in alignment faking, where they pursue a hidden goal, but conceal...

https://arxiv.org/abs/2610.33180

8 CONF

Scaling Qualitative Analysis with Large Language Models: Development and Validation of the Smart Coder Framework for Human-LLM Collaborative Coding

ArXiv CS.AI 50%

arXiv:2610.33205v1 Announce Type: new Abstract: Current uses of large language models (LLMs) for qualitative coding often rely on ad hoc prompting with limited guidance for reliability and human oversight. To address these limitations, we propose the Smart Coder Framework, a...

https://arxiv.org/abs/2610.33205

9 CONF

Credit Without Ground Truth: Claim-Matched Evaluation of Step-Level Credit with Finite Executed Replay

ArXiv CS.AI 50%

arXiv:2608.19760v3 Announce Type: replace-cross Abstract: Step-level credit evaluation can conflate numerical association with ordering and branch recovery with position-selection value. We test these distinctions in ALFWorld with two independently seeded replay banks: 27,520...

https://arxiv.org/abs/2608.19760

0 CONF

TriageRAG: Failure-Aware Self-Training for Adapting Retrieval-Augmented Language Models to Specialized Domains

ArXiv CS.AI 50%

arXiv:2610.33776v1 Announce Type: cross Abstract: Retrieval-augmented generation (RAG) lets large language models (LLMs) answer knowledge-intensive questions from retrieved evidence, but adapting a general-purpose RAG system to a specialized domain such as medicine is difficult...

https://arxiv.org/abs/2610.33776

1 CONF

Estimation of Propagated Risk in LLM-based Multi-Agent Systems

ArXiv CS.AI 50%

arXiv:2610.33282v1 Announce Type: new Abstract: Large Language Model (LLM) based multi-agent systems introduce a distinctive reliability challenge in which actions taken by upstream agents can affect the correctness of downstream agents. While prior work has investigated error...

https://arxiv.org/abs/2610.33282

2 CONF

Game theory of academic effort and hiring in the age of AI

ArXiv CS.AI 50%

arXiv:2610.35802v1 Announce Type: new Abstract: Many academic disciplines (in particular, mathematics and the theoretical sciences) face difficult questions about how to adapt their practices to the rise of highly capable artificial intelligence (AI). Before AI, a well-written...

https://arxiv.org/abs/2610.35802

3 CONF

Content-Aware Adaptive Token Pruning for Video-Language Models

ArXiv CS.AI 50%

arXiv:2608.03112v2 Announce Type: replace-cross Abstract: Vision-language models excel at image and video understanding but suffer from high inference latency due to the need to process thousands of tokens per image, limiting their deployment on resource-constrained edge devices...

https://arxiv.org/abs/2608.03112

4 CONF

Data Optimization for LLMs: A Survey

ArXiv CS.AI 50%

arXiv:2610.33405v1 Announce Type: new Abstract: Large language models (LLMs) have attracted great attention in both academic and industrial communities in recent years. Numerous survey studies have covered various aspects of LLMs, such as learning strategies, applications,...

https://arxiv.org/abs/2610.33405

5 CONF

The Hitchhiker's Guide to System 1 Models: Fast Decision-Making Beyond LLMS

ArXiv CS.AI 50%

arXiv:2610.33775v1 Announce Type: new Abstract: Large Language Models excel at reasoning and generation but their autoregressive, token-by-token inference makes them unnecessarily slow, expensive, and poorly calibrated for the many routine decisions (routing, gating, scoring,...

https://arxiv.org/abs/2610.33775

6 CONF

ConMem: Counterfactual Multi-Scenario Probing for Resolving LLM Memory Misalignment

ArXiv CS.AI 50%

arXiv:2610.33432v1 Announce Type: new Abstract: Long-term memory is critical for autonomous Large Language Model (LLM) agents executing multi-session, long-horizon tasks. However, existing memories often inappropriately store, retain, or retrieve over-generalized, misattributed,...

https://arxiv.org/abs/2610.33432

7 CONF

BoardSafety: AI-Ready Data from Manufacturer Documentation for Development-Board Safety

ArXiv CS.AI 50%

arXiv:2610.33448v1 Announce Type: new Abstract: The development of large language models (LLMs) are capable to assist with hardware experiments. In this scenario, this raises a practical question: Can LLMs determine whether a proposed operation is safe for a specific development...

https://arxiv.org/abs/2610.33448

8 CONF

A Comprehensive Framework for All Stages of Large-Model Training: From Data and Tokenization, through Pretraining, to Post-Training Alignment

ArXiv CS.AI 50%

arXiv:2610.33451v1 Announce Type: new Abstract: The capabilities of large language models (LLMs) and multimodal LLMs (MLLMs) arise from the synergy of the full training pipeline, including data and tokenization, pretraining, and post-training (supervised fine-tuning, preference...

https://arxiv.org/abs/2610.33451

9 CONF

LLM-Enhanced Heterogeneous Graph Learning with Business Semantics for Financial Fraud Detection

ArXiv CS.AI 50%

arXiv:2610.35298v1 Announce Type: new Abstract: Financial fraud detection increasingly relies on heterogeneous transaction graphs, in which entities such as cards and merchants play distinct business roles and the risk meaning of an attribute depends on the entity's role: a high...

https://arxiv.org/abs/2610.35298

0 CONF

Learning When to Collaborate: Selective Multi-Agent Medical Reasoning via Uncertainty-Aware Routing

ArXiv CS.AI 50%

arXiv:2610.33544v1 Announce Type: new Abstract: Multi-agent systems have emerged as a promising paradigm for multimodal medical reasoning, but existing methods rely on two assumptions that hinder practical deployment: (i) always-active dense collaboration among multiple agents,...

https://arxiv.org/abs/2610.33544

1 CONF

Are We Done with UMLS Injection into Language Models for Medical Question Answering?

ArXiv CS.AI 50%

arXiv:2610.33801v1 Announce Type: new Abstract: Maybe not. The Unified Medical Language System (UMLS) has become an oracle for millions of clinicians and health system developers worldwide. Researchers have increasingly used knowledge graphs derived from UMLS to enhance language...

https://arxiv.org/abs/2610.33801

2 CONF

Data Governance for Generative and Agentic AI : Technology, Law, Economics, and Security

ArXiv CS.AI 50%

arXiv:2610.33605v1 Announce Type: new Abstract: This book develops a unified governance framework for systems that retrieve, generate, infer, remember and act on information from large language models and retrieval augmented pipelines to autonomous multi agent systems. Eleven...

https://arxiv.org/abs/2610.33605

3 CONF

Graph-dependent pooling: adaptive granularity in hierarchical graph neural networks

ArXiv CS.AI 50%

arXiv:2610.33623v1 Announce Type: new Abstract: Graph Neural Networks are widely used to learn structural representations of relational data. Hierarchical graph pooling learns multi-level representations by progressively coarsening graphs; however, prior methods rely on...

https://arxiv.org/abs/2610.33623

4 CONF

AlphaIdeaBench: Benchmarking Automated Research Idea Generation under Feasibility Constraints with Peer Review Alignment

ArXiv CS.AI 50%

arXiv:2610.33644v1 Announce Type: new Abstract: Research ideas are always proposed under resource constraints: a plan must fit the available compute, API credit, and time. Yet evaluations of research-ideation agents mostly score unconstrained novelty or overlap with a published...

https://arxiv.org/abs/2610.33644

5 CONF

Physics-Informed Graph-Based Safe RL with Constraint-Specific Critics for Community Battery Coordination in Distribution Network

ArXiv CS.AI 50%

arXiv:2610.33654v1 Announce Type: new Abstract: High penetrations of distributed energy resources (DERs) and fast-varying electric vehicle (EV) charging intensify voltage and line-limit violations in distribution networks (DNs), making safe, low-cost operation increasingly...

https://arxiv.org/abs/2610.33654

6 CONF

Power-of-Two Leaky Integrate-and-Fire Neurons for Event-Driven Vision-Language-Action Inference

ArXiv CS.AI 50%

arXiv:2610.33658v1 Announce Type: new Abstract: Vision-language-action (VLA) models have demonstrated strong capabilities in robotic manipulation, yet their reliance on dense, continuous-valued computation presents a fundamental challenge for energy-constrained deployment....

https://arxiv.org/abs/2610.33658

7 CONF

Human-Feedback-Guided Multi-Agent Inverse Reinforcement Learning for Interactive Driving Behavior Modeling

ArXiv CS.AI 50%

arXiv:2610.33726v1 Announce Type: new Abstract: Modeling human driving behavior is fundamental to understanding driver decision-making and supporting autonomous vehicles (AVs) in mixed traffic. Reward learning offers a principled way to infer the latent objectives behind driving...

https://arxiv.org/abs/2610.33726

8 CONF

An actively trained neural operator for near-explicit atmospheric chemistry

ArXiv CS.AI 50%

arXiv:2610.33771v1 Announce Type: new Abstract: Earth-system models simplify atmospheric chemistry into condensed mechanisms, because near-explicit mechanisms such as the Master Chemical Mechanism (MCM) are too costly to solve in every grid cell. Neural emulators could lift this...

https://arxiv.org/abs/2610.33771

9 CONF

Improving the Chain-of-Thought Compressibility of Pruned Reasoning Models via Long--Short Contrastive Pruning

ArXiv CS.AI 50%

arXiv:2610.34336v1 Announce Type: cross Abstract: Large reasoning models are costly to serve because they are large and because they generate long chains of thought (CoT). Weight pruning and CoT compression address these two costs separately, but they interfere in that pruning...

https://arxiv.org/abs/2610.34336

0 CONF

Build, Don't Play: Self-Improving LLM Research Agents Develop a Hearthstone Battlegrounds AI That Rivals Top Players

ArXiv CS.AI 50%

arXiv:2610.33845v1 Announce Type: new Abstract: Can LLM agents carry a hard real-world problem from raw data to a system that holds its own against people? We study this in Hearthstone Battlegrounds, an eight-player auto-battler with stochastic shops and battles, hidden...

https://arxiv.org/abs/2610.33845

1 CONF

WebViSE: Scaling Visual Design Intelligence from the Open Web for LLM Website Generation

ArXiv CS.AI 50%

arXiv:2610.33929v1 Announce Type: new Abstract: Large language models have made remarkable strides in website development, enabling complete and interactive web experiences to be built from natural-language instructions with far less effort. Yet, non-expert users often struggle...

https://arxiv.org/abs/2610.33929

2 CONF

Fed-Universe: A Semantic-Geometric-Topological-Human (S-G-T-H) Stack for Negotiated Alignment in Federated Systems

ArXiv CS.AI 50%

arXiv:2610.34009v1 Announce Type: new Abstract: Conventional federated learners implicitly optimize for global means, creating a "one-size-fits-all" paradigm that inevitably suppresses the minority under cross-site heterogeneity. To bridge this critical gap, we present...

https://arxiv.org/abs/2610.34009

3 CONF

Can LLMs Trust Appropriately? Evaluating Calibrated Reliance on Imperfect Partners

ArXiv CS.AI 50%

arXiv:2610.34107v1 Announce Type: new Abstract: Large language models increasingly collaborate with humans and peer AI agents whose contributions may be incomplete, inconsistent, or incorrect. Effective collaboration requires calibrated trust, whereby decisions to rely on,...

https://arxiv.org/abs/2610.34107

4 CONF

PCoreSet: Effective Active Learning through Knowledge Distillation from Vision-Language Models

ArXiv CS.AI 50%

arXiv:2506.00910v3 Announce Type: replace-cross Abstract: Knowledge distillation (KD) is a widely used framework for training compact, task-specific models by transferring the knowledge from teacher models. However, its application to active learning (AL), which aims to minimize...

https://arxiv.org/abs/2506.00910

5 CONF

R2A-Sales: A Policy-Atom-Matched Benchmark for Rule-to-Action Gaps in Goal-Directed Sales Agents

ArXiv CS.AI 50%

arXiv:2610.34185v1 Announce Type: new Abstract: LLM-based sales agents are often evaluated on static policy questions or task success, but deployment asks a sharper question: when a model can state a rule, will it choose compliant, tool-grounded actions under commercial...

https://arxiv.org/abs/2610.34185

6 CONF

CTXAgent - Self-Controlled Context Management of LLM Agents

ArXiv CS.AI 50%

arXiv:2610.34425v1 Announce Type: new Abstract: Agent harnesses keep long sessions inside the context window by compacting the history when a token threshold is crossed. We study an alternative in which the model manages its own context: it decides when to compact, what to keep...

https://arxiv.org/abs/2610.34425

7 CONF

What Effort Level Should I Choose? On the Economics of Reasoning Effort in LLM Agents

ArXiv CS.AI 50%

arXiv:2610.34491v1 Announce Type: new Abstract: Generative models increasingly expose a reasoning effort setting that lets users scale inference-time compute, trading higher cost for potentially improved performance. Users must choose an effort level before knowing whether extra...

https://arxiv.org/abs/2610.34491

8 CONF

How Much Information Can We Fit inside a Transformer?

ArXiv CS.AI 50%

arXiv:2610.34553v1 Announce Type: new Abstract: How efficiently can transformers store factual knowledge? The answer to this question has implications for both model design and understanding the limits of transformer-based systems. Prior studies observe a linear relationship...

https://arxiv.org/abs/2610.34553

9 CONF

A Battery for AI Propensity Measurement: Rubrics, Synthetic Item Banks, and Annotations

ArXiv CS.AI 50%

arXiv:2610.34614v1 Announce Type: new Abstract: We release a battery of 28 propensities for AI research, with expert-written rubrics, synthetic datasets, and annotated demand intervals, built with the method of Romero-Alvarado et al. (2026). We selected a wide collection of...

https://arxiv.org/abs/2610.34614

0 CONF

SentryBench: A Cost-Sensitive Benchmark for Adversarial Robustness of Tool-Using LLM Agents in Public-Safety Deployments

ArXiv CS.AI 50%

arXiv:2610.34535v1 Announce Type: cross Abstract: Large language model (LLM) agents are increasingly wired to consequential tools that send public alerts, dispatch responders, and query citizen records, where one erroneous tool call carries real-world cost. Existing agent-safety...

https://arxiv.org/abs/2610.34535

1 CONF

One Mean Log-Mel Vector per Phoneme: Simple and Fast Text Injection for Speech Recognition

ArXiv CS.AI 50%

arXiv:2610.34623v1 Announce Type: new Abstract: Text injection trains speech recognition models on unpaired text alongside paired audio. We feed the speech encoder synthetic input derived from text so the whole model learns from it. Previously, we synthesized this input with a...

https://arxiv.org/abs/2610.34623

2 CONF

SlowLab: Benchmarking Autonomous AI Agents on a Year of Experiments They Cannot Undo

ArXiv CS.AI 50%

arXiv:2610.34644v1 Announce Type: new Abstract: AI agents are increasingly proposed as autonomous experimenters, yet the benchmarks that test them mostly offer cheap experiments that answer at once. Experiments in physical facilities are different: they take months, occupy...

https://arxiv.org/abs/2610.34644

3 CONF

Curriculum-Guided Acquisition of Hierarchical Task Knowledge from LLM-Generated Experience

ArXiv CS.AI 50%

arXiv:2610.34677v1 Announce Type: new Abstract: Hierarchical task knowledge enables efficient problem solving by allowing previously acquired skills to be reused across tasks and problem instances. In this paper, we investigate curriculum-guided acquisition of hierarchical task...

https://arxiv.org/abs/2610.34677

4 CONF

Provably Efficient Q-learning for Robust Reinforcement Learning

ArXiv CS.AI 50%

arXiv:2610.35656v1 Announce Type: cross Abstract: Standard reinforcement learning (RL) algorithms based on Markov Decision Processes (MDPs) are vulnerable to sim-to-real gaps. To address this issue, distributionally robust RL algorithms based on robust MDPs (RMDPs) have been...

https://arxiv.org/abs/2610.35656

5 CONF

Safe by Benchmark, Unsafe in Practice: Lessons from 1331 Child-Related AI Incidents

ArXiv CS.AI 50%

arXiv:2610.34705v1 Announce Type: new Abstract: AI systems are increasingly present in children's everyday environments: education, entertainment, social networks, and toys. Their safety is typically assessed at the model level, through benchmarks, red-teaming, and content...

https://arxiv.org/abs/2610.34705

6 CONF

TrustJev: Separating Policy Authority from Untrusted Evidence in LLM-Based Decisions

ArXiv CS.AI 50%

arXiv:2610.34826v1 Announce Type: new Abstract: Decision models enable software to select from predefined actions using natural language inputs. However, the external evidence that supplies facts for these decisions may also contain instructions that attempt to override the...

https://arxiv.org/abs/2610.34826

7 CONF

Logic Hijacking of RAG: Single-Document Poisoning via Local Optimization

ArXiv CS.AI 50%

arXiv:2610.34829v1 Announce Type: new Abstract: Single-document corpus poisoning under black-box access poses a severe threat to Retrieval-Augmented Generation (RAG). However, existing attacks predominantly rely on explicit factual refutations or imperative instruction...

https://arxiv.org/abs/2610.34829

8 CONF

ProtAgent: Post-Training Compact Protein Tool Agents with Interaction-Level Credit Assignment

ArXiv CS.AI 50%

arXiv:2610.33746v1 Announce Type: cross Abstract: Protein annotation agents must choose suitable bioinformatics tools, interpret their outputs, and use the resulting evidence in the final answer. Most existing systems rely on frontier-scale models. We ask whether a compact model...

https://arxiv.org/abs/2610.33746

9 CONF

GADAR: Generalizing Across Domains via Action Ranking for Classical Planning

ArXiv CS.AI 50%

arXiv:2610.34868v1 Announce Type: new Abstract: Learned policies for classical planning act without search, but each is tied to the domain it was trained on: network modules per action schema, parameters per predicate, or vocabulary-sized encodings mean a trained policy cannot...

https://arxiv.org/abs/2610.34868

0 CONF

DocTalk: Enhancing Medical Knowledge Retrieval via Fine-tuned Large Language Models with Chain-of-Thought Reasoning for Robust Clinical Decision Support

ArXiv CS.AI 50%

arXiv:2610.34871v1 Announce Type: new Abstract: Objective: The rapid advancement of large language models (LLMs) in healthcare has opened new possibilities for AI-assisted clinical decision-making; however, safe and effective deployment remains challenging because of...

https://arxiv.org/abs/2610.34871

1 CONF

ClinVerse: A Realistic Multi-Agent Diagnostic Environment for Training and Evaluating Medical Agents

ArXiv CS.AI 50%

arXiv:2610.34878v1 Announce Type: new Abstract: Disease diagnosis is a key application of large language model (LLM) agents in healthcare. However, existing medical agents are primarily trained and evaluated with static question-answering data, providing limited opportunities to...

https://arxiv.org/abs/2610.34878

2 CONF

C-SafeGen: Certified Safe LLM Generation with Claim-Based Streaming Guardrails

ArXiv CS.AI 50%

arXiv:2610.34961v1 Announce Type: new Abstract: Despite the remarkable capabilities of large language models (LLMs) across diverse applications, they remain vulnerable to generating content that violates safety regulations and policies. To mitigate these risks, LLMs undergo...

https://arxiv.org/abs/2610.34961

3 CONF

Why Are LLMs Confidently Wrong? Correcting Overconfident Errors via Causal Head Intervention

ArXiv CS.AI 50%

arXiv:2610.34994v1 Announce Type: new Abstract: Modern large language models (LLMs) often exhibit overconfident errors, assigning high confidence to incorrect predictions and thereby undermining reliability in real-world use. Existing calibration methods typically rely on...

https://arxiv.org/abs/2610.34994

4 CONF

Beyond Single-Attribute Auditing: A Configurable Multi-Attribute Benchmark for Intersectional Fairness in LLMs

ArXiv CS.AI 50%

arXiv:2610.35069v1 Announce Type: new Abstract: Large language models (LLMs) are increasingly evaluated for fairness, yet most existing benchmarks audit one sensitive attribute at a time. This single-attribute paradigm can miss intersectional harms, where multiple attributes...

https://arxiv.org/abs/2610.35069

5 CONF

Reva: Real-Time Vulnerability Detection through Contextual Probing of Large Language Models

ArXiv CS.AI 50%

arXiv:2610.35829v1 Announce Type: cross Abstract: A vulnerability costs more to fix the later it is discovered. As coding agents generate code faster than developers can review it, software needs a first line of defense screening each new function in real time. LLM agents have...

https://arxiv.org/abs/2610.35829

6 CONF

When Fusion Order Matters: Insights on Progressive Multimodal Representation Learning for Cross-modal Retrieval

ArXiv CS.AI 50%

arXiv:2610.35125v1 Announce Type: new Abstract: Multimodal video retrieval combines visual, acoustic, and speech information, yet how their fusion order should depend on the data and the retrieval query remains unclear. We introduce Unified Progressive Representation Learning...

https://arxiv.org/abs/2610.35125

7 CONF

From Hype to ROI: An Economic Model for LLM Adoption in Information Systems and Security Operations

ArXiv CS.AI 50%

arXiv:2610.33352v1 Announce Type: cross Abstract: Large Language Models (LLMs) promise significant efficiency gains across information-intensive operations by automating repetitive language-based tasks. However, adoption decisions remain largely opportunistic, as organizations...

https://arxiv.org/abs/2610.33352

8 CONF

MuSID-KD: Multigranularity Knowledge Distillation for Generative Recommendation with Semantic IDs

ArXiv CS.AI 50%

arXiv:2610.35218v1 Announce Type: new Abstract: Generative recommendation with semantic IDs (SIDs) represents each item as a sequence of discrete tokens and autoregressively predicts the next item's SID. Despite strong recommendation performance, repeated decoding over multiple...

https://arxiv.org/abs/2610.35218

9 CONF

Trainable Boundary-Keeping Where No Formal Answer Exists: A Controlled Knock-In Study of Reward Training in Evidence-Sufficiency Judgment

ArXiv CS.AI 50%

arXiv:2610.35226v1 Announce Type: new Abstract: Reasoning-tuned and reward-trained models are often assumed to verify more reliably because they reason longer or receive correctness signals. We test this in evidence-sufficiency judgment: deciding whether available information...

https://arxiv.org/abs/2610.35226

0 CONF

Before Blaming the LLM, Diagnose the Workflow: A Role-Level Difficulty Profile for Translating Natural Language into Ontology-Grounded Statements

ArXiv CS.AI 50%

arXiv:2610.35267v1 Announce Type: new Abstract: Natural-language-to-domain-specific-language (NL-to-DSL) semantic parsing systems increasingly translate informal user descriptions into structured artifacts such as Structured Query Language (SQL) queries, Application Programming...

https://arxiv.org/abs/2610.35267

1 CONF

Self-Supervised Pretraining for PHM: What Transfers, What Does Not, and What Federating It Costs

ArXiv CS.AI 50%

arXiv:2610.35280v1 Announce Type: new Abstract: Self-supervised pretraining has produced capable time-series foundation models, but its value for Prognostics and Health Management (PHM) has not been measured under leakage-controlled conditions. We measure it on a corpus of 47...

https://arxiv.org/abs/2610.35280

2 CONF

MarkovSim: an LLM Agent Harness with Bounded Context for Long-Horizon Computational Simulation

ArXiv CS.AI 50%

arXiv:2610.35349v1 Announce Type: cross Abstract: Recent studies have applied large language model (LLM) agent systems to computational simulation workflows. Such workflows are long-horizon, so the information generated can exceed what an LLM can reason over at once, its context...

https://arxiv.org/abs/2610.35349

3 CONF

Strategic Calibration in LLM Multi-Agent Systems: Audited Pivot Scoring for Truthful Elicitation

ArXiv CS.AI 50%

arXiv:2610.33728v1 Announce Type: cross Abstract: In large language model (LLM)-based multi-agent systems, reported confidence increasingly determines the allocation of scarce downstream attention, computation, and decision influence. This makes reporting strategic: an agent may...

https://arxiv.org/abs/2610.33728

4 CONF

Omni-InteractScope: Benchmarking Proactive and Agentic Interaction in Continuous Omnimodal Streams

ArXiv CS.AI 50%

arXiv:2610.35439v1 Announce Type: new Abstract: Omnimodal large language models (Omni-LLMs) are increasingly expected to assist users while observing continuous audiovisual streams. However, existing evaluations provide limited insight into how reliably these models connect...

https://arxiv.org/abs/2610.35439

5 CONF

Infinity Machine: a self-improving AI agent for research

ArXiv CS.AI 50%

arXiv:2610.35579v1 Announce Type: new Abstract: We present the Infinity Machine, a self-improving AI agent that we use for research on problems whose answers aprogram or data can check. Each candidate solution is produced by a full tool-using agent that works for up to an hour...

https://arxiv.org/abs/2610.35579

6 CONF

Linear Probes Transfer from Base to Distilled Language Models despite Rotation

ArXiv CS.AI 50%

arXiv:2610.35586v1 Announce Type: new Abstract: Linear probes, classifiers that read a language model's internal activations, are widely used to detect harmful requests. They are cheap to train and run, and an auditor who has trained one would like to reuse it when the model is...

https://arxiv.org/abs/2610.35586

7 CONF

Agentic Algorithm Design: LLMs that Propose, Formalise, and Prove Runtime Bounds for Evolutionary Algorithms in Lean

ArXiv CS.AI 50%

arXiv:2610.35595v1 Announce Type: new Abstract: LLMs now prove hard theorems in Lean. Can they also design algorithms and prove their correctness and efficiency? LLM-guided design of optimisation heuristics (such as AlphaEvolve) is purely empirical. Automated theorem provers...

https://arxiv.org/abs/2610.35595

8 CONF

VaLoP: Value Gradient Lookahead and Pullback for Reward-Guided Diffusion

ArXiv CS.AI 50%

arXiv:2610.35607v1 Announce Type: new Abstract: Reward guidance steers a pretrained diffusion or flow model toward images or actions that a reward model prefers, at sampling time and without retraining the model for each reward. It follows the gradient of a value function, the...

https://arxiv.org/abs/2610.35607

9 CONF

Explainable Multimodal AI Approach to Reduce Unnecessary Breast Biopsies

ArXiv CS.AI 50%

arXiv:2610.35704v1 Announce Type: new Abstract: Recently, there has been rise in multimodal vision-language models for computer-aided diagnosis (CAD), yet their application in breast diagnosis, particularly for predicting biopsy outcomes remains limited. This study presents a...

https://arxiv.org/abs/2610.35704

0 CONF

Free to Look, Free to Reason: Simple LLM Harnesses Outdiagnose Hand-Crafted Root Cause Analysis Agents

ArXiv CS.AI 50%

arXiv:2610.35767v1 Announce Type: new Abstract: Root cause analysis (RCA) of microservice failures has attracted many LLM agents whose investigation is hand-crafted, through standard operating procedures, fixed stages and role agents. Each is evaluated on its own data against...

https://arxiv.org/abs/2610.35767

1 CONF

A Unified Arena for LLM Agents across Diverse Eurogame

ArXiv CS.AI 50%

arXiv:2610.35773v1 Announce Type: new Abstract: Interactive environments offer a practical way to evaluate large language model (LLM) agents, and competitive games extend this evaluation to decisions shaped by competing goals and other agents' choices. Yet existing evaluations...

https://arxiv.org/abs/2610.35773

2 CONF

DairySocialTwin: Auditing the video-to-graph evidence chain in cattle social networks

ArXiv CS.AI 50%

arXiv:2610.32989v1 Announce Type: cross Abstract: Video-derived social graphs can appear precise while inheriting uncertainty from detection, tracking, classification, and aggregation. We present DairySocialTwin, a graph-centered audit framework applied to 203 videos (122 h 23...

https://arxiv.org/abs/2610.32989

3 CONF

Whispering Wisely: Smart Subset Selection for ASR Fine-Tuning

ArXiv CS.AI 50%

arXiv:2610.32997v1 Announce Type: cross Abstract: Automatic Speech Recognition (ASR) models such as OpenAIs Whisper have substantially improved speech recognition quality but remain expensive to fine-tune, especially when large amounts of transcribed audio are available. This...

https://arxiv.org/abs/2610.32997

4 CONF

RADAgent: Automating RADseq Pipelines with Private Open-Source LLMs

ArXiv CS.AI 50%

arXiv:2610.33009v1 Announce Type: cross Abstract: Restriction site-associated DNA sequencing (RADseq) is a family of reduced-representation sequencing methods. It uses restriction enzymes to sample subsets of the genome, therefore allowing multiple individuals to be sequenced...

https://arxiv.org/abs/2610.33009

5 CONF

Recursive Semantic Divergence for LLM Agent Consistency

ArXiv CS.AI 50%

arXiv:2610.33012v1 Announce Type: cross Abstract: Large Language Model (LLM) agents with tool use exhibit inconsistent behavior across independent runs from identical states, calling different tools, producing contradictory outputs, or pursuing divergent strategies. This...

https://arxiv.org/abs/2610.33012

6 CONF

SLICE-Mem: Salience-Gated Lightweight Information Curation for Efficient LLM Agent Memory

ArXiv CS.AI 50%

arXiv:2610.33015v1 Announce Type: cross Abstract: Long-term memory enables large language model (LLM) agents to maintain personalized and temporally coherent interactions, but repeated LLM inference for memory admission, extraction, and organization introduces substantial cost...

https://arxiv.org/abs/2610.33015

7 CONF

Merge-Based Vocabulary Extension Stalls at the Pre-Tokenizer: An English-Preserving Repair for Nepali, Tested in Continued Pretraining

ArXiv CS.AI 50%

arXiv:2610.33016v1 Announce Type: cross Abstract: Llama-3, Qwen3 and many other open LLMs pre-tokenize text with a regular expression whose word class admits only Unicode letters. Devanagari vowel signs are combining marks, so this regex splits Nepali words at every vowel sign,...

https://arxiv.org/abs/2610.33016

8 CONF

Mine or Yours? Mind-Bridged Language Models Take Each Other's Memories as Their Own

ArXiv CS.AI 50%

arXiv:2610.33020v1 Announce Type: cross Abstract: If your brain were connected directly to someone else's, would you still know which thoughts were yours? We build a functional version with two copies of one language model: each holds a private rule and code word, and one copy,...

https://arxiv.org/abs/2610.33020

9 CONF

Lessons from Luck: Do LLM-Written Lessons Carry Chance into an Agent's Next Decision?

ArXiv CS.AI 50%

arXiv:2610.33029v1 Announce Type: cross Abstract: Self-improving agents write lessons from the outcomes of their runs; when outcomes are noisy, a lesson can carry chance into later decisions. We measure one link of this chain in a registered study. In 256 synthetic worlds, each...

https://arxiv.org/abs/2610.33029

00 CONF

Recoverability-Controlled Safety Miscalibration in Large Language Models

ArXiv CS.AI 50%

arXiv:2610.22425v2 Announce Type: replace-cross Abstract: Safety-aligned language models are expected to refuse harmful requests while remaining useful for benign ones. However, because many evaluations score only the final response, they often conflate intent understanding with...

https://arxiv.org/abs/2610.22425

01 CONF

MOSAIC: Towards a Universal Latent Space for Computational Pathology Foundation Models

ArXiv CS.AI 50%

arXiv:2610.33059v1 Announce Type: cross Abstract: Pathology foundation models (PFMs) differ substantially in architecture, pretraining objective, training data and representation dimensionality, yet it remains unclear whether they learn fundamentally different feature spaces or...

https://arxiv.org/abs/2610.33059

02 CONF

Oddin: Latent Visual Reasoning in Unified Multimodal Models via Self-Distillation from Image Generation

ArXiv CS.AI 50%

arXiv:2610.33062v1 Announce Type: cross Abstract: Unified large multimodal models (ULMMs) are able to perform visual reasoning through interleaved text-image generation, but generating explicit intermediate images incurs significant computational overhead. While visual latent...

https://arxiv.org/abs/2610.33062

03 CONF

A Canary in the Neural Mine: Catching an Image Classifier's Silent, Confident Failures

ArXiv CS.AI 50%

arXiv:2610.33063v1 Announce Type: cross Abstract: Deployed vision models auto-accept confident predictions, so high-confidence errors propagate silently. Maximum Softmax Probability (MSP), the standard failure predictor, has a structural blind spot: as a scalar read from the...

https://arxiv.org/abs/2610.33063

04 CONF

REALM: Retrospective Encoder Alignment for LFP Modeling

ArXiv CS.AI 50%

arXiv:2605.14867v3 Announce Type: replace-cross Abstract: Spike activity has been the dominant neural signal for behavior decoding because its high spatiotemporal resolution supports accurate decoding. However, as intracortical brain-computer interfaces (iBCIs) move toward...

https://arxiv.org/abs/2605.14867

05 CONF

PhaseBreak: Proactive Temporal Protection against Unauthorized Image-to-Video Animation

ArXiv CS.AI 50%

arXiv:2610.33070v1 Announce Type: cross Abstract: Image-to-Video (I2V) diffusion models can animate a single portrait or photograph, enabling shared images to be repurposed for motion synthesis without the subject's authorization. We study proactive visual protection against...

https://arxiv.org/abs/2610.33070

06 CONF

Beyond Vision-Language Models: Visual-Semantic Affect Graphs for Emotion Recognition in Artworks

ArXiv CS.AI 50%

arXiv:2610.35244v1 Announce Type: cross Abstract: Emotion recognition in artworks remains challenging because affective meaning is often subjective, implicit, and shaped by relationships between visual form and semantic associations. While vision-language models provide strong...

https://arxiv.org/abs/2610.35244

07 CONF

EGGPlacer: Energy-Guided Generative Chip Placement with Diffusion Models

ArXiv CS.AI 50%

arXiv:2610.33099v1 Announce Type: cross Abstract: Chip placement strongly affects routed wirelength and thermal behavior, yet these downstream metrics are difficult to evaluate during macro placement. Learning-based macro placers commonly optimize the macro-only half-perimeter...

https://arxiv.org/abs/2610.33099

08 CONF

Whale You Go Green? Immediate and Longitudinal Effects of Embodied Conversational AI Avatars on Sustainable Behavioral Tendencies and Pro-Environmental Attitudes

ArXiv CS.AI 50%

arXiv:2610.33191v1 Announce Type: cross Abstract: Embodied conversational AI can be persuasive, but can it motivate people to develop sustainable behavioral tendencies and increase pro-environmental attitudes? We conducted a randomized experiment (N = 576) comparing four...

https://arxiv.org/abs/2610.33191

09 CONF

Position, Role or Label? A Controlled Study of Instruction Conflicts in LLM Agents

ArXiv CS.AI 50%

arXiv:2610.33201v1 Announce Type: cross Abstract: LLM agents take instructions from the operator, repository configuration, user, tool results and saved memory, which often conflict. Instruction-hierarchy benchmarks ask whether the higher-ranked source wins, but a source sets...

https://arxiv.org/abs/2610.33201

10 CONF

Operationalizing Governed Agentic Workflows in Regulated Finance: A Customer Zero Case Study using Zafin AIOS

ArXiv CS.AI 50%

arXiv:2610.33253v1 Announce Type: cross Abstract: Regulated industries such as finance and healthcare face critical governance challenges that hamper the adoption of AI-assisted software development and hinder the transition from experimental pilots to production work. This...

https://arxiv.org/abs/2610.33253

11 CONF

Exploring Joint-Embedding Predictive Architectures for Self-Supervised Representation Learning of Powder X-ray Diffraction Patterns

ArXiv CS.AI 50%

arXiv:2610.33310v1 Announce Type: cross Abstract: Powder X-ray diffraction (XRD) is the primary high-throughput method for resolving crystal structures, a cornerstone of materials science. As high-throughput synthesis and self-driving laboratories advance, automated XRD pattern...

https://arxiv.org/abs/2610.33310

12 CONF

When Does Pre-Trained Capacity Help Causal Learning? From Scalar Prediction to Counterfactual Post-Training

ArXiv CS.AI 50%

arXiv:2610.33311v1 Announce Type: cross Abstract: How does pre-trained model capacity relate to performance after adaptation for causal tasks? We study two systems built on the Qwen 2.5 family: CausalLLM, which estimates scalar potential outcomes with treatment-specific heads,...

https://arxiv.org/abs/2610.33311

13 CONF

A Deterministic Security Runtime for LLM Agents on MPI and Batch-Scheduled HPC Systems: Consensus-Governed, Hardware-Rooted Supervision of Every Rank

ArXiv CS.AI 50%

arXiv:2610.33318v1 Announce Type: cross Abstract: Large language model agents launch simulations on high-performance computing (HPC) systems, where one sbatch or mpirun can commit thousands of node-hours and a forged result can cost lives. Of 23 surveyed systems, none supervises...

https://arxiv.org/abs/2610.33318

14 CONF

Quantifying Meaning with SCAPE-T: Semantic Complexity from Encoder-Only Transformer Attention

ArXiv CS.AI 50%

arXiv:2610.33441v1 Announce Type: cross Abstract: Automated semantic complexity estimation remains underexplored in NLP, despite important applications in linguistics and interpretability. This paper defines semantic complexity as the code length required to uniquely specify a...

https://arxiv.org/abs/2610.33441

15 CONF

LLMs in Cloud: A Discussion on Architecture, Orchestration, and Inference

ArXiv CS.AI 50%

arXiv:2610.33475v1 Announce Type: cross Abstract: The widespread deployment of foundational Large Language Models (LLM) across cloud datacenters marks a fundamental paradigm shift from stateless deep learning toward distributed, memory-intensive generative workloads. However,...

https://arxiv.org/abs/2610.33475

16 CONF

Pacifo: An Efficient Fully Homomorphic Encryption-based Lightweight ResNet for Trustworthy Machine Learning as a Service

ArXiv CS.AI 50%

arXiv:2610.33525v1 Announce Type: cross Abstract: Machine Learning as a Service (MLaaS) has become a popular paradigm for deploying AI applications. However, it raises significant privacy concerns that the raw client data needs to be uploaded to the server for inference...

https://arxiv.org/abs/2610.33525

17 CONF

Hybrid-time PINNs: Discontinuity-Aware Physics-Informed Neural Networks for Learning Hybrid Dynamical Systems

ArXiv CS.AI 50%

arXiv:2610.33562v1 Announce Type: cross Abstract: Hybrid dynamical systems flow continuously but jump at discrete events. Physics-informed neural networks (PINNs) fit data while enforcing known dynamics, but a single network is smooth in time and cannot model these jumps. We...

https://arxiv.org/abs/2610.33562

18 CONF

WorldFailBench: Do World Models Understand Policy Failures?

ArXiv CS.AI 50%

arXiv:2610.33565v1 Announce Type: cross Abstract: Robot policy evaluation informs deployment decisions and policy improvement, but physical testing and simulator preparation require substantial effort. World models offer a complementary route through closed-loop interaction in...

https://arxiv.org/abs/2610.33565

19 CONF

Replicate What Varies: Resource Sharing for Domain-Randomized Embodied Simulation

ArXiv CS.AI 50%

arXiv:2610.33582v1 Announce Type: cross Abstract: Simulation provides a low-cost sandbox where embodied agents can attempt tasks, observe outcomes, and learn from experience. Running many diverse simulation environments concurrently can consume substantial CPU and GPU memory...

https://arxiv.org/abs/2610.33582

20 CONF

Build or Adopt? Custom-Built vs. General-Purpose LLM Agent Harnesses for Automotive GUI Testing

ArXiv CS.AI 50%

arXiv:2610.33636v1 Announce Type: cross Abstract: As general-purpose LLM agent harnesses mature, organizations must decide whether to continue developing in-house systems or adopt existing alternatives. We study this decision at Volvo Group for GUI testing of Android-based...

https://arxiv.org/abs/2610.33636

21 CONF

Extraction meets Distillation - cryptanalytic techniques to recover parameters from a Deep Neural Network

ArXiv CS.AI 50%

arXiv:2610.33671v1 Announce Type: cross Abstract: We study fully connected neural-network reconstruction by combining cryptanalytic weight-direction extraction with learning from predicted labels. We adapt a ReLU extraction implementation to Leaky ReLU, compare sign-search...

https://arxiv.org/abs/2610.33671

22 CONF

Embedding-Based Educational Taxonomies

ArXiv CS.AI 50%

arXiv:2610.33694v1 Announce Type: cross Abstract: Educational taxonomies organize educational content, such as exercise problems, into concepts, which are also called skills or knowledge components. These taxonomies define learning objectives in educational applications such as...

https://arxiv.org/abs/2610.33694

23 CONF

Trust Mechanisms for Offline Generative AI at the Edge: A Connectivity-Regime Framework and an Empirical Study on a Raspberry Pi 5

ArXiv CS.AI 50%

arXiv:2610.33732v1 Announce Type: cross Abstract: Generative AI and large language models (LLMs) are moving into the Internet of Things, yet most deployments assume cloud connectivity, and many IoT settings have intermittent, restricted, or no connectivity. We propose a...

https://arxiv.org/abs/2610.33732

24 CONF

A Distribution-Matching Perspective on LLM Training Algorithms

ArXiv CS.AI 50%

arXiv:2610.33740v1 Announce Type: cross Abstract: In this work, we propose a unified framework for training algorithms across different stages of modern large language model development, based on the perspective of distribution matching. We show that a broad class of language...

https://arxiv.org/abs/2610.33740

25 CONF

POD-RoPE: Rotary Position Encoding over Physical Observation Distributions for Embodied AI

ArXiv CS.AI 50%

arXiv:2610.33745v1 Announce Type: cross Abstract: Embodied perception requires positional representations that capture both physical geometry and the spatial and temporal extent of observations. We introduce POD-RoPE (Physical Observation Distribution RoPE), which represents...

https://arxiv.org/abs/2610.33745

26 CONF

StateWise: Teaching Language Models to Track and Obey Evolving Constraints in Multi-Turn Dialogue

ArXiv CS.AI 50%

arXiv:2610.33766v1 Announce Type: cross Abstract: Users of conversational assistants rarely state all requirements up front: they add constraints, amend them, withdraw them, digress, and return. Instruction-tuned LLMs fail here in two ways: active constraints stop being honored...

https://arxiv.org/abs/2610.33766

27 CONF

ProMo3D: Probing Motion Cues in Frozen 3D Foundation Models

ArXiv CS.AI 50%

arXiv:2610.33802v1 Announce Type: cross Abstract: Feed-forward 3D foundation models have recently reshaped multi-view reconstruction, yet dynamic 4D understanding still often relies on large-scale motion supervision or hand-crafted, architecture-specific heuristics. We ask...

https://arxiv.org/abs/2610.33802

28 CONF

EEGBiFormer: Real-Time EEG Attention Classification via Hybrid CNN-Transformer with EEG-Informed Auditory Protocol Validation

ArXiv CS.AI 50%

arXiv:2610.33832v1 Announce Type: cross Abstract: Sustained attention is critical for academic performance, yet existing auditory support tools are static and not grounded in real-time neural evidence. This paper presents EEGBiFormer, a hybrid CNN-Transformer model for realtime...

https://arxiv.org/abs/2610.33832

29 CONF

Intervention-Certified Amplification: A Validity Gate for σ-Weighted Post-Training

ArXiv CS.AI 50%

arXiv:2610.33882v1 Announce Type: cross Abstract: Hybrid post-training often mixes supervised fine-tuning (SFT) and reinforcement learning (RL), using training-side signals to set each sample's balance. Such signals characterize the anchor data and policy, but not whether the...

https://arxiv.org/abs/2610.33882

30 CONF

Rank-1 Weight-Corrected Trainable Weight Averaging for Out-of-Distribution Detection

ArXiv CS.AI 50%

arXiv:2610.33909v1 Announce Type: cross Abstract: Out-of-distribution (OOD) detection identifies inputs outside the training distribution. Rank-based analyses such as RankFeat and RankWeight suggest that OOD overconfidence is closely related to dominant low-rank components in...

https://arxiv.org/abs/2610.33909

31 CONF

What Makes a Question Curious? Measuring Model Curiosity in Human-LLM Interactions

ArXiv CS.AI 50%

arXiv:2610.33950v1 Announce Type: cross Abstract: Large language models increasingly act as partners in everyday communication. In these interactions, useful assistance may require not only answering, but also asking the right questions that surface knowledge gaps, assumptions,...

https://arxiv.org/abs/2610.33950

32 CONF

Thinking-in-Bets at Test Time Reduces Sycophantic Revisions in LLM Agents

ArXiv CS.AI 50%

arXiv:2610.34654v1 Announce Type: cross Abstract: Language model agents often abandon correct work when a user challenges it, yet simply telling them to stand firm makes them ignore real mistakes. Once an agent reads a complaint, it tends to discount its own earlier checks in...

https://arxiv.org/abs/2610.34654

33 CONF

Beyond Belief: Evaluating and Mitigating LLM Susceptibility to Misinformation

ArXiv CS.AI 50%

arXiv:2610.33962v1 Announce Type: cross Abstract: Tool-using large language models (LLMs) must decide whether external information warrants action. We investigate whether their expressed credibility judgments reliably indicate behavioral resistance to misinformation. We evaluate...

https://arxiv.org/abs/2610.33962

34 CONF

QECGBERT: A Scalable ECG Foundation Model Based on Self-Supervised Language Learning via VQ-VAE Waveform Tokenization

ArXiv CS.AI 50%

arXiv:2610.33964v1 Announce Type: cross Abstract: AI foundation models play a pivotal role in advancing heart disease diagnosis from electrocardiogram (ECG) data, particularly for rare diseases where the lack of sufficient labeled data hinders supervised learning. However, most...

https://arxiv.org/abs/2610.33964

35 CONF

BriLLM-MetaPred: A Representation-Prediction Unity Language Model with Adaptive Self-Measurement

ArXiv CS.AI 50%

arXiv:2610.34018v1 Announce Type: cross Abstract: We present BriLLM-MetaPred, the first engineering implementation of the MetaPred machine-learning framework established by a mathematical theory of AGI, instantiated for language modeling. The theory argues that MetaPred is both...

https://arxiv.org/abs/2610.34018

36 CONF

Extended LASMP: Language Aided Subset Sampling Based Motion Planner for Drone Navigation

ArXiv CS.AI 50%

arXiv:2610.34054v1 Announce Type: cross Abstract: This paper proposes a command-aligned sampling distribution that incorporates navigation commands directly into the exploration process to improve the efficiency of sampling-based motion planning for quadcopters in...

https://arxiv.org/abs/2610.34054

37 CONF

Identically Trained Language Models Break on Different Jailbreaks

ArXiv CS.AI 50%

arXiv:2610.34074v1 Announce Type: cross Abstract: Jailbreak benchmarks test one checkpoint and treat which jailbreaks break it as a property of its pretraining data and safety training. We ask whether another pretraining run of the same data, given the same safety training,...

https://arxiv.org/abs/2610.34074

38 CONF

HumanoidNDM: Closing the Training-Deployment Gap of Motion Tracking via Humanoid Neural Dynamics Model

ArXiv CS.AI 50%

arXiv:2610.34126v1 Announce Type: cross Abstract: Humanoid robots have recently achieved remarkable success in motion tracking. However, during real-world deployment, robots inevitably encounter dynamics gaps arising from either extra physical disturbances or inaccurate...

https://arxiv.org/abs/2610.34126

39 CONF

Decoupled Prune-and-Reconstruct: Training on a Compact Core with Double-Layer Sparsification

ArXiv CS.AI 50%

arXiv:2610.34140v1 Announce Type: cross Abstract: On million-node graphs, the memory wall arrives before the compute wall: a GNN that never fits in VRAM cannot deliver the all-node inference its accuracy promises. The standard remedy, pruning, has always been an accuracy tax -...

https://arxiv.org/abs/2610.34140

40 CONF

Rethinking Retrieval in the Long-Context Era: Training-Free Semantic Trees for Adaptive Context Construction

ArXiv CS.AI 50%

arXiv:2610.34146v1 Announce Type: cross Abstract: Long-context language models can increasingly process entire documents, making full-context prompting a practical alternative to passage retrieval. However, greater capacity determines how much information a model can receive,...

https://arxiv.org/abs/2610.34146

41 CONF

Single-Cycle Battery Remaining Useful Life Prediction with Compressed Continuous Wavelet Features and Deep Neural Networks

ArXiv CS.AI 50%

arXiv:2610.34159v1 Announce Type: cross Abstract: Lithium-ion batteries are widely used in industrial applications, but their performance gradually degrades over time, making accurate lifetime prediction important for safety and reliability. However, reliable lifetime assessment...

https://arxiv.org/abs/2610.34159

42 CONF

MemComposer: Learning to Assemble Memories for LLM Agents via Reinforcement Learning

ArXiv CS.AI 50%

arXiv:2610.21091v2 Announce Type: replace Abstract: Agent memory enables large language model (LLM) agents to reuse past experiences without updating their parameters. However, semantic similarity alone does not guarantee memory usefulness: downstream performance depends on how...

https://arxiv.org/abs/2610.21091

43 CONF

NameTrap: Steering AI Coding Agents Toward Security Vulnerabilities via Variable Names

ArXiv CS.AI 50%

arXiv:2610.34253v1 Announce Type: cross Abstract: AI coding agents routinely clone repositories, install packages, and read source code obtained from the Web, so publicly hosted code becomes part of the context that shapes the code they write. We show that this Web-borne context...

https://arxiv.org/abs/2610.34253

44 CONF

DoubleFree:Exploiting Exploration History for Rollout-Free Multi-Task LLMs Training

ArXiv CS.AI 50%

arXiv:2610.34272v1 Announce Type: cross Abstract: Multi-task post-training can develop domain specialists through independent reinforcement learning (RL) and consolidate their capabilities through multi-teacher on-policy distillation (MOPD). However, standard MOPD generates...

https://arxiv.org/abs/2610.34272

45 CONF

Repairing Broken Access Control in Web Applications with LLMs: An Empirical Benchmark Study

ArXiv CS.AI 50%

arXiv:2610.34290v1 Announce Type: cross Abstract: Broken access control (BAC) ranks among the most severe web-application vulnerability classes, yet its automated repair remains largely unexplored. Unlike memory-safety vulnerabilities, BAC is a logic vulnerability whose repair...

https://arxiv.org/abs/2610.34290

46 CONF

Automated, auditable dating of the Babylonian Astronomical Diaries with language models and astronomical tools

ArXiv CS.AI 50%

arXiv:2610.34297v1 Announce Type: cross Abstract: The Babylonian Astronomical Diaries record the sky over nearly six centuries. Many tablets are broken and have lost their date formula. Experts dated many of them one tablet at a time in the standard edition, Astronomical Diaries...

https://arxiv.org/abs/2610.34297

47 CONF

Developer Task Delegation to LLMs and the Trust Factor: A Mixed Methods Study

ArXiv CS.AI 50%

arXiv:2610.34317v1 Announce Type: cross Abstract: With the rise of Large Language Models (LLMs), software developers are presented with a powerful new approach: delegating software development tasks to LLMs. Existing studies have explored task delegation for a subset of tasks...

https://arxiv.org/abs/2610.34317

48 CONF

SHIFT-JEV: Certified Abstention and Gain-Gated Continual Adaptation for LLM-Based Decision Models under Distribution Shift

ArXiv CS.AI 50%

arXiv:2610.34329v1 Announce Type: cross Abstract: Large language models (LLMs) can act as decision models that return a probability distribution over a fixed set of options, but deployed decision models face domain, label-prior, emerging-class and candidate-set shift. We study...

https://arxiv.org/abs/2610.34329

49 CONF

SEEK: Skill Extraction with Evolving Knowledge against Frontier Science Agents

ArXiv CS.AI 50%

arXiv:2610.34365v1 Announce Type: cross Abstract: Science agents increasingly package domain expertise as reusable skills that plan analyses, select tools, and execute workflows. This modularity creates a confidentiality risk: the interactions that deliver a skill's capability...

https://arxiv.org/abs/2610.34365

50 CONF

ComBeautify: Leveraging Code Change Summaries for Conventional Commit Generation

ArXiv CS.AI 50%

arXiv:2610.34374v1 Announce Type: cross Abstract: Modern software development relies on version control systems such as Git to track changes and facilitate collaboration. While commit messages are crucial for conveying developer intent, they do not affect program functionality;...

https://arxiv.org/abs/2610.34374

51 CONF

Whether or Where? Tail Calibration Signals Finite-Budget Reward Hacking in Best-of-N Selection

ArXiv CS.AI 50%

arXiv:2610.34375v1 Announce Type: cross Abstract: Best-of-N (BoN) inference returns the candidate with the highest proxy reward, so a larger N tests the proxy's extreme tail more than its average ranking quality. We show that the exact BoN value is a Beta-weighted integral of...

https://arxiv.org/abs/2610.34375

52 CONF

CyberDual: AI Attackers versus AI Code Reviewers in Real Software Maintenance

ArXiv CS.AI 50%

arXiv:2610.34435v1 Announce Type: cross Abstract: Open-source projects now let AI agents review incoming code, but the code under review may itself come from an AI attacker. We present CyberDual, an arena for this dual role of AI: attacker agents hide an exploitable bug in real...

https://arxiv.org/abs/2610.34435

53 CONF

Multimodal Requirements Traceability Link Recovery in Industrial Settings

ArXiv CS.AI 50%

arXiv:2610.34503v1 Announce Type: cross Abstract: Requirements traceability links support activities such as change-impact analysis, coverage checking, and validation. Industrial requirements can include traceable images, which text-only approaches cannot process directly. For a...

https://arxiv.org/abs/2610.34503

54 CONF

LLM WikiLeaks: Uncovering URL-Based Capability Extension and Security Circumvention in Agentic Message Boards

ArXiv CS.AI 50%

arXiv:2610.34542v1 Announce Type: cross Abstract: Autonomous agents have been observed acting beyond their assigned tasks, accessing unauthorized systems, and establishing communication channels to exchange information, yet the Web and network-level techniques they use to extend...

https://arxiv.org/abs/2610.34542

55 CONF

GRASP: Group-Adaptive Hybrid RL-SFT Post-Training for Medical VQA without Reasoning Annotations

ArXiv CS.AI 50%

arXiv:2610.34561v1 Announce Type: cross Abstract: Adapting general-purpose vision-language models (VLMs) to medical visual question answering (VQA) is valuable for helping clinicians read images and make diagnoses, but annotated data are scarce and provide only short answers,...

https://arxiv.org/abs/2610.34561

56 CONF

Recursive CLIP: Vision-Language Distillation with Looped Transformer Blocks

ArXiv CS.AI 50%

arXiv:2610.34563v1 Announce Type: cross Abstract: Contrastive vision-language models such as CLIP typically rely on independently parameterized transformer blocks across depth. We investigate the extent to which these blocks can instead be reused recurrently while preserving the...

https://arxiv.org/abs/2610.34563

57 CONF

Audio-R1: Audio Reasoning through Reinforcement Learning on Self-Supervised Tasks

ArXiv CS.AI 50%

arXiv:2610.34609v1 Announce Type: cross Abstract: Reasoning is increasingly important for large audio-language models (LALMs), yet remains relatively underexplored compared with text-only and vision-language models. We introduce Audio-R1, a reinforcement learning (RL)...

https://arxiv.org/abs/2610.34609

58 CONF

Intrinsic Goals for Autonomous Agents: Hierarchical Exploration Captures Future Volitional Action in Drosophila

ArXiv CS.AI 50%

arXiv:2610.34638v1 Announce Type: cross Abstract: Animals have a remarkable capacity for autonomy---driven by intrinsic motivation---without external sensory cues or extrinsic objectives, yet the computational mechanisms behind this ability largely remain a mystery. We combine...

https://arxiv.org/abs/2610.34638

59 CONF

Agent Memory as Recommendation: Recommender-System Design for Long-Term Memory of LLM Agents

ArXiv CS.AI 50%

arXiv:2610.34673v1 Announce Type: cross Abstract: Large language model (LLM) agents that serve users across many sessions must choose, for every request, the few stored memories that enter a limited context window. This choice largely decides whether the agent answers correctly,...

https://arxiv.org/abs/2610.34673

60 CONF

Babysitting AI Agents: A User-Centered Taxonomy and Study of Agentic Delegation Failures

ArXiv CS.AI 50%

arXiv:2610.34748v1 Announce Type: cross Abstract: Users are increasingly delegating tasks to AI agents, allowing agents to perform long-running, multi-step workflows on their behalf. However, agents may perform these tasks in ways that depart from users' expectations for how...

https://arxiv.org/abs/2610.34748

61 CONF

DynEdit: Dynamic Entropy-Guided Sequential Editing for Large Language Models

ArXiv CS.AI 50%

arXiv:2610.34768v1 Announce Type: cross Abstract: Knowledge editing aims to update outdated or incorrect knowledge in large language models without full retraining. Existing parameter-editing methods remain limited for long-sequence knowledge: single-point editors provide...

https://arxiv.org/abs/2610.34768

62 CONF

Disentangled Sparse Representations for Concept-Separated Diffusion Unlearning

ArXiv CS.AI 50%

arXiv:2605.12122v3 Announce Type: replace-cross Abstract: Unlearning specific concepts in text-to-image diffusion models has become increasingly important for preventing undesirable content generation. Among prior approaches, sparse autoencoder (SAE)-based methods have attracted...

https://arxiv.org/abs/2605.12122

63 CONF

SPARTA: A Sparse Ternary LLM Accelerator with Adaptive Attention and Unified Datapath for Long-Context Edge Inference

ArXiv CS.AI 50%

arXiv:2610.34824v1 Announce Type: cross Abstract: LLMs on edge devices are constrained by limited memory and bandwidth. Although 1.58-bit ternary weights greatly reduce model storage, long-context inference is increasingly bottlenecked by KV-cache traffic and attention. Existing...

https://arxiv.org/abs/2610.34824

64 CONF

Adaptive Retrieval, Selective Answers: Evaluating Conformal Factuality in Agentic RAG

ArXiv CS.AI 50%

arXiv:2610.34872v1 Announce Type: cross Abstract: Retrieval-augmented generation (RAG) can improve factual grounding without guaranteeing that every generated claim is supported. Previous work has studied claim-level conformal factuality in single-stage RAG and, more recently,...

https://arxiv.org/abs/2610.34872

65 CONF

AIPC: Evidence-Guided Agent Workflows for PyTorch Project Deployment

ArXiv CS.AI 50%

arXiv:2610.34890v1 Announce Type: cross Abstract: Deploying and optimizing a PyTorch inference project on a dedicated accelerator, especially an on-device NPU, requires more than using a framework device API as with CUDA. The project must be converted into hardware-runtime...

https://arxiv.org/abs/2610.34890

66 CONF

LE-GDA: Layer Ensemble Gaussian Discriminant Analysis for Training-Free Few-Shot CLIP Adaptation

ArXiv CS.AI 50%

arXiv:2610.34901v1 Announce Type: cross Abstract: Training-free few-shot adaptation of vision-language models typically constructs a classifier from CLIP's final projected image representation, leaving intermediate representations unused. Combining them is not straightforward....

https://arxiv.org/abs/2610.34901

67 CONF

GUIFlow: Temporal Redundancy-Aware Incremental Inference for GUI Agents

ArXiv CS.AI 50%

arXiv:2610.34959v1 Announce Type: cross Abstract: GUI agents built on vision-language models interact with continuously evolving interfaces, yet conventional inference repeatedly recomputes representations of largely overlapping observations and histories. This mismatch makes...

https://arxiv.org/abs/2610.34959

68 CONF

CA-HNM: Constraint-Aware Hard Negative Mining for Dense Retrieval with Structured Domain Models

ArXiv CS.AI 50%

arXiv:2610.35045v1 Announce Type: cross Abstract: Dense retrievers are commonly trained with hard negatives selected from lexical or dense retrieval results. Retrieval proximity, however, does not indicate whether a candidate conflicts with explicit domain structure. We...

https://arxiv.org/abs/2610.35045

69 CONF

Locked In: Multi-Objective RL for the AI Data Center Build-Out

ArXiv CS.AI 50%

arXiv:2610.35147v1 Announce Type: cross Abstract: AI data center build-out is a multi-trillion-dollar, multi-decade infrastructure program. Yet siting, power procurement, and cooling decisions are still made with static, single-point scoring tools. These tools are evaluated...

https://arxiv.org/abs/2610.35147

70 CONF

Context Is Not Free: Carbon-Aware LLM Routing for Multi-Turn LLM Web Services across Geo-Distributed Data Centers

ArXiv CS.AI 50%

arXiv:2610.35230v1 Announce Type: cross Abstract: The sustainability of the Web has drawn growing attention, as large language model (LLM)-based web services leave a substantial carbon footprint. These services typically run in geo-distributed data centers whose grid carbon...

https://arxiv.org/abs/2610.35230

71 CONF

Reading Beyond OCR Noise: Robust Cross-Lingual RAG for Legislative Archives

ArXiv CS.AI 50%

arXiv:2610.35240v1 Announce Type: cross Abstract: India's historical public archives present severe computer vision and document image analysis challenges: decades of parliamentary records survive only as degraded, multi-column scanned pages featuring mixed scripts, complex...

https://arxiv.org/abs/2610.35240

72 CONF

Satellite Intelligence: Distributed Training, Collaborative Inference,and LLM Adaptation in LEO Constellation

ArXiv CS.AI 50%

arXiv:2610.35282v1 Announce Type: cross Abstract: The proliferation of orbital sensing data and the real-time demands of emerging global services are pushing low Earth orbit (LEO) satellite networks from bent-pipe relays toward intelligent onboard processing. This shift...

https://arxiv.org/abs/2610.35282

73 CONF

Constrained Decoding: Harmful or Beneficial? Revisiting Its Impact on LLM Performance

ArXiv CS.AI 50%

arXiv:2610.35328v1 Announce Type: cross Abstract: Constrained decoding guarantees that Large Language Models (LLMs) produce outputs conforming to predefined formats, but is often assumed to inherently degrade reasoning. We examine this assumption through a taxonomy of three...

https://arxiv.org/abs/2610.35328

74 CONF

Should Pseudo-Labels Share Mini-Batches with Labeled Examples? Semi-Supervised Continual Active Learning with Limited Training Cost

ArXiv CS.AI 50%

arXiv:2610.35345v1 Announce Type: cross Abstract: Active learning (AL) reduces annotation by querying informative examples. Continual active learning (CAL) also cuts the training cost, the total number of gradient steps, by training one model across AL rounds with a fixed number...

https://arxiv.org/abs/2610.35345

75 CONF

Shared-Factor-Conditioned Diffusion for Low-Rank Tensor Recovery

ArXiv CS.AI 50%

arXiv:2610.35368v1 Announce Type: cross Abstract: Tensor recovery spans heterogeneous data structures and observation processes, calling for a representation that can adapt across modalities and tasks. We propose DiffTR, an unsupervised framework that unifies tensor completion,...

https://arxiv.org/abs/2610.35368

76 CONF

ACSAN: All-Convolutional Self-Attention Network Integrating Convolutional Operations with Vision Transformer

ArXiv CS.AI 50%

arXiv:2610.35375v1 Announce Type: cross Abstract: In the past few years, Vision Transformer (ViT) has achieved excellent performance in the field of computer vision. However, there still some problems remain. For instance, ViTs heavily rely on large-scale training data since...

https://arxiv.org/abs/2610.35375

77 CONF

Autonomous Structural Discovery Of Bioprocess Dynamics And Kinetics Via LLM-Driven Agents

ArXiv CS.AI 50%

arXiv:2610.35391v1 Announce Type: cross Abstract: High-fidelity digital twins in continuous biomanufacturing are often bottlenecked by incomplete kinetic models, while standard data-driven algorithms (e.g., SINDy, PINNs) lack interpretability or impose rigid physical...

https://arxiv.org/abs/2610.35391

78 CONF

Governing Enterprise AI Traffic at Scale: Design and Operation of a Policy-Driven, Multi-Vendor LLM Gateway on Envoy

ArXiv CS.AI 50%

arXiv:2610.35509v1 Announce Type: cross Abstract: Large organizations consume large language models (LLMs) from several vendors at once, each with its own API dialect, credential model, streaming protocol and pricing. Integrated directly, this heterogeneity yields credential...

https://arxiv.org/abs/2610.35509

79 CONF

Generative Structure from Motion with Native 3D Diffusion

ArXiv CS.AI 50%

arXiv:2610.35522v1 Announce Type: cross Abstract: Recovering an object's 3D structure and camera motion from a few uncalibrated views is a fundamental problem in computer vision. Classical structure from motion methods recover 3D geometry consistent with the input images but...

https://arxiv.org/abs/2610.35522

80 CONF

JANUS: Synthesizing Cross-Layer Defenses Against Web Deserialization Attacks

ArXiv CS.AI 50%

arXiv:2610.35532v1 Announce Type: cross Abstract: Unsafe deserialization turns untrusted data into attacker-controlled execution, and a single vulnerability can admit diverse exploit variants that evade defenses tailored to a known exploit. Existing defenses typically operate at...

https://arxiv.org/abs/2610.35532

81 CONF

The Agent in the Middle: How AI-Mediated Access Reshapes Web Vulnerability Exploitation

ArXiv CS.AI 50%

arXiv:2610.35541v1 Announce Type: cross Abstract: Web services increasingly expose backend interfaces that let users delegate tasks to AI agents, which invoke service operations directly without navigating websites. This shift changes how existing Web vulnerabilities can be...

https://arxiv.org/abs/2610.35541

82 CONF

From Language to Drives: A Multimodal Framework for Aligning Autonomous Robots with Human Purposes

ArXiv CS.AI 50%

arXiv:2610.35629v1 Announce Type: cross Abstract: This work explores a multimodal approach to the alignment of the behavior of open-ended learning autonomous robots with human purposes through the use of large language models as cultural enablers. While autonomous robots can...

https://arxiv.org/abs/2610.35629

83 CONF

Adapting LLM-SC to Living Sacred Heritage: Domain-Specific Extraction, Semantic Networks and Human Evaluation on Pashupatinath Temple Reviews

ArXiv CS.AI 50%

arXiv:2610.35685v1 Announce Type: cross Abstract: Large language model (LLM) pipelines can turn visitor reviews into structured data, but generic schemes may miss rituals, sacred spaces, cultural restrictions and emotionally mixed experiences at living sacred heritage sites. We...

https://arxiv.org/abs/2610.35685

84 CONF

TrySee: Modular Perception-Enhanced VLM Grounding for Closed-Loop Robot Task Planning in Specialized Domains

ArXiv CS.AI 50%

arXiv:2610.35687v1 Announce Type: cross Abstract: Vision Language Models (VLMs) have shown impressive capabilities in high-level reasoning and task planning for robots. However, specialized domains remain a challenge for VLMs' image interpretation and grounding abilities. We...

https://arxiv.org/abs/2610.35687

85 CONF

Quantitative Generalization Bounds for Multi-Layer Transformers: A Few Long Strings can Beat Many Short Ones

ArXiv CS.AI 50%

arXiv:2610.35715v1 Announce Type: cross Abstract: Transformer-based language models exhibit an astonishing capacity to generalize to longer and more difficult problem instances than those seen in training. However, theoretical understanding of this phenomenon still lags behind....

https://arxiv.org/abs/2610.35715

86 CONF

Audit Where Hacks Pay: Budgeted Strong-Verifier Audits Against Reward Hacking in GRPO

ArXiv CS.AI 50%

arXiv:2610.35716v1 Announce Type: cross Abstract: Reinforcement learning with verifiable rewards trains on cheap checks, such as unit tests, that a policy can learn to game, while a stronger check is too expensive to run on every rollout. We study how to spend a limited...

https://arxiv.org/abs/2610.35716

87 CONF

Directional Attention Guidance for Diffusion Models via Sparse Hopfield Retrieval

ArXiv CS.AI 50%

arXiv:2603.02531v3 Announce Type: replace-cross Abstract: Classifier-Free Guidance (CFG) improves sample quality in diffusion models, but its dual-pass inference and null-condition training limit its use in few-step regimes. Attention-space guidance is a single-pass alternative,...

https://arxiv.org/abs/2603.02531

88 CONF

Estimating Long-term Heterogeneous Dose-response Curve: Generalization Bound Leveraging Optimal Transport Weights

ArXiv CS.AI 50%

arXiv:2406.19195v4 Announce Type: replace-cross Abstract: Long-term treatment effect estimation is a significant but challenging problem in many applications. Existing methods rely on ideal assumptions, such as no unobserved confounders or binary treatment, to estimate long-term...

https://arxiv.org/abs/2406.19195

89 CONF

Artificial Leviathan: Exploring Social Evolution of LLM Agents Through the Lens of Hobbesian Social Contract Theory

ArXiv CS.AI 50%

arXiv:2406.14373v4 Announce Type: replace Abstract: The emergence of Large Language Models (LLMs) and advancements in Artificial Intelligence (AI) offer an opportunity for computational social science research at scale. Building upon prior explorations of LLM agent design, our...

https://arxiv.org/abs/2406.14373

90 CONF

Exploring the Limits of Large Language Models as a Reflection of Human Cognition: An Illustration in the Context of Moral Judgement

ArXiv CS.AI 50%

arXiv:2508.06950v4 Announce Type: replace Abstract: Large Language Models (LLMs), such as ChatGPT, are increasingly used in research, ranging from simple writing assistance to complex data annotation tasks. Recently, some research has suggested that LLMs are able to reproduce...

https://arxiv.org/abs/2508.06950

91 CONF

Closing the empirical loop: autonomous AI agents conduct end-to-end research with human participants

ArXiv CS.AI 50%

arXiv:2508.13421v3 Announce Type: replace Abstract: Artificial intelligence systems are transforming scientific discovery by accelerating specific research tasks, from protein structure prediction to materials design, yet remain confined to narrow domains requiring substantial...

https://arxiv.org/abs/2508.13421

92 CONF

Agentic AI Empowered Intent-Based Networking for 6G

ArXiv CS.AI 50%

arXiv:2601.06640v2 Announce Type: replace Abstract: The transition towards sixth-generation (6G) wireless networks necessitates autonomous orchestration mechanisms capable of translating high-level operational intents into executable network configurations. Existing approaches...

https://arxiv.org/abs/2601.06640

93 CONF

Do MLLMs Really See It: Reinforcing Visual Attention in Multimodal LLMs

ArXiv CS.AI 50%

arXiv:2602.08241v2 Announce Type: replace Abstract: While chain-of-thought (CoT) reasoning has substantially improved multimodal large language models (MLLMs) on complex reasoning tasks, existing approaches largely rely on long textual reasoning trajectories and provide limited...

https://arxiv.org/abs/2602.08241

94 CONF

Beyond 0-100: How Confidence Scale Design Shapes Verbalized Confidence and LLM Metacognition

ArXiv CS.AI 50%

arXiv:2603.09309v4 Announce Type: replace Abstract: Verbalized confidence, in which LLMs report a numerical certainty score, is widely used to estimate uncertainty in black-box settings, yet the confidence scale itself (typically 0--100) is rarely examined. We show that this...

https://arxiv.org/abs/2603.09309

95 CONF

NGM: A Plug-and-Play Training-Free Memory Module for LLMs

ArXiv CS.AI 50%

arXiv:2605.16893v2 Announce Type: replace Abstract: Recent studies have introduced memory-augmented architectures that decouple knowledge storage from neural computation, enabling more direct access to stored representations. However, existing approaches often rely on learned...

https://arxiv.org/abs/2605.16893

96 CONF

Measuring Cross-Modal Synergy: A Benchmark for VLM Explainability

ArXiv CS.AI 50%

arXiv:2605.22168v2 Announce Type: replace Abstract: Vision-Language Models (VLMs) map complex visual inputs to semantic spaces, but interpreting the cross-modal reasoning of VLMs currently relies on post-hoc explainers evaluated via unimodal perturbation metrics. We expose a...

https://arxiv.org/abs/2605.22168

97 CONF

Trace Integrity for LLM Data Agents: Auditing Structured-Data Computation Beyond Answer Accuracy in Real-World Systems

ArXiv CS.AI 50%

arXiv:2608.26036v2 Announce Type: replace Abstract: Large language model (LLM) data agents answer structured-data questions by generating SQL, invoking tools, executing queries, and summarizing results. In operational settings, reliability depends on both reported answer and...

https://arxiv.org/abs/2608.26036

98 CONF

Agents That Edit Documents: Measuring Agentic PDF Forgery Against a Non-Agentic Control

ArXiv CS.AI 50%

arXiv:2609.23953v2 Announce Type: replace Abstract: AI agents that carry a multi-step computer task through on their own became ordinary tools in the past year, and the same autonomy is available to anyone whose task is harmful. We ask what that means for a relying party -- an...

https://arxiv.org/abs/2609.23953

99 CONF

DISCERN: Can AI Agents Work Like Scientists and Guide Discovery?

ArXiv CS.AI 50%

arXiv:2609.33357v2 Announce Type: replace Abstract: Reliable automated research requires agents to vet data, verify analyses, and generate hypotheses grounded in trustworthy evidence, potentially reducing routine scientific workload while allowing scientists to focus on...

https://arxiv.org/abs/2609.33357

00 CONF

A Computer Vision Approach to Visual Fraud Detection in Phishing Websites Using YOLOv8

ArXiv CS.AI 50%

arXiv:2609.34015v2 Announce Type: replace Abstract: Phishing remains one of the most common vectors for financial and identity fraud, and most detection systems still rely on inspecting a page's URL, HTML markup, or domain registration history. These signals are easy for an...

https://arxiv.org/abs/2609.34015

01 CONF

EOPSA: Efficient On-Policy Self-Distilled Safety Alignment

ArXiv CS.AI 50%

arXiv:2609.34519v2 Announce Type: replace Abstract: On-Policy Self-Distillation (OPSD) has emerged as a promising paradigm for safety alignment, delivering dense, token-level supervision by distilling from a teacher conditioned on refusal-oriented privileged prompts. However, we...

https://arxiv.org/abs/2609.34519

02 CONF

A Framework for Automated Multi-Source Satellite Data Analytics and LLM-Based Report Generation

ArXiv CS.AI 50%

arXiv:2610.05625v2 Announce Type: replace Abstract: This paper presents the workflow for building an automated ArcGIS Pro tool using ArcPy to extract the Land Surface Temperature (LST) from Landsat 7, 8 and 9 datasets. The tool eliminates the need for manual band selection and...

https://arxiv.org/abs/2610.05625

03 CONF

EpiFlow: Bridging Safety and Performance in Offline Reinforcement Learning via Epigraph-Guided Flow Matching

ArXiv CS.AI 50%

arXiv:2602.08054v2 Announce Type: replace-cross Abstract: Offline reinforcement learning (RL) offers a promising framework for deploying autonomous systems in safety-critical settings without the risks of online exploration. However, learning policies that simultaneously achieve...

https://arxiv.org/abs/2602.08054

04 CONF

Foundation Models in Medical Image Analysis: A Systematic Review and Meta-Analysis

ArXiv CS.AI 50%

arXiv:2510.16973v2 Announce Type: replace-cross Abstract: Recent advancements in foundation models (FMs) have catalyzed a paradigm shift in medical image analysis. Unlike traditional task-specific artificial intelligence (AI) models, FMs leverage large-scale datasets to learn...

https://arxiv.org/abs/2510.16973

05 CONF

Can Loan Processors and Underwriters Spot AI-Edited Statements? An Empirical Study

ArXiv CS.AI 50%

arXiv:2610.30758v2 Announce Type: replace Abstract: AI-assisted editing can produce loan documents that preserve familiar issuer layouts while changing the applicant's identity. We investigated whether experienced underwriters and loan processors identify such forgeries in...

https://arxiv.org/abs/2610.30758

06 CONF

Discovering and Reasoning of Causality in the Hidden World with Large Language Models

ArXiv CS.AI 50%

arXiv:2402.03941v4 Announce Type: replace-cross Abstract: Revealing hidden causal variables alongside the underlying causal mechanisms is essential to the development of science. Despite the progress in the past decades, existing practice in causal discovery (CD) heavily relies...

https://arxiv.org/abs/2402.03941

07 CONF

EF-LLM: A Continual Learning-Based Framework for AI-Assisted Automation, Enhanced Sparse Prediction, and Hallucination Detection

ArXiv CS.AI 50%

arXiv:2411.00852v3 Announce Type: replace-cross Abstract: Accurate prediction helps achieve supply-demand balance in energy systems,thereby supporting operational decision-making and scheduling.Traditional models,which lack AI-assisted automation,often rely heavily on expert...

https://arxiv.org/abs/2411.00852

08 CONF

Evaluating LLM-based Agents for Multi-Turn Conversations: A Survey

ArXiv CS.AI 50%

arXiv:2503.22458v3 Announce Type: replace-cross Abstract: This survey examines evaluation methods for large language model (LLM)-based agents in multi-turn conversational settings. Using a PRISMA-inspired framework, we systematically reviewed nearly 250 scholarly sources,...

https://arxiv.org/abs/2503.22458

09 CONF

Simple yet Effective Semi-supervised Knowledge Distillation from Vision-Language Models via Dual-Head Optimization

ArXiv CS.AI 50%

arXiv:2505.07675v3 Announce Type: replace-cross Abstract: Semi-supervised learning (SSL) has emerged as a practical solution for addressing data scarcity challenges by leveraging unlabeled data. Recently, vision-language models (VLMs), pre-trained on massive image-text pairs,...

https://arxiv.org/abs/2505.07675

10 CONF

Benchmarking Classification Tasks of Chest X-ray AI Across Multinational Public and Private Datasets: a retrospective diagnostic accuracy study

ArXiv CS.AI 50%

arXiv:2505.16027v2 Announce Type: replace-cross Abstract: Chest X-ray (CXR) AI models are increasingly proposed for triage, screening, and rule-out, where the operating point is fixed to miss few abnormalities. Most evaluations instead report threshold-free discrimination on a...

https://arxiv.org/abs/2505.16027

11 CONF

FediLoRA: Practical Federated Fine-Tuning of Foundation Models Under Missing-Modality Constraints

ArXiv CS.AI 50%

arXiv:2509.06984v4 Announce Type: replace-cross Abstract: Federated Learning with LoRA fine-tuning offers an efficient and privacy-aware solution for institutions to collaboratively leverage their large datasets to train VLLMs. However, participating institutions often possess...

https://arxiv.org/abs/2509.06984

12 CONF

Hierarchical Alignment: Block-Level LoRA-DPO for Preference Tuning in Large Language Models

ArXiv CS.AI 50%

arXiv:2510.12044v2 Announce Type: replace-cross Abstract: Preference tuning typically distributes adapter updates across every Transformer layer. We study how the placement of a fixed LoRA budget affects Direct Preference Optimization (DPO). Hierarchical Alignment partitions...

https://arxiv.org/abs/2510.12044

13 CONF

A Causal Perspective on the Role of Training Dynamics for Interpreting Privacy Risks in Code Models

ArXiv CS.AI 50%

arXiv:2512.07814v3 Announce Type: replace-cross Abstract: Large language models for code (LLM4Code) have greatly improved developer productivity but also raise privacy concerns due to their reliance on open-source repositories containing abundant personally identifiable...

https://arxiv.org/abs/2512.07814

14 CONF

LLM-Generated Counterfactual Stress Scenarios for Portfolio Risk Simulation via Hybrid Prompt-RAG Pipeline

ArXiv CS.AI 50%

arXiv:2512.07867v2 Announce Type: replace-cross Abstract: We develop a transparent, auditable large language model (LLM) pipeline for macro-financial stress testing. Structured prompting, with optional retrieval of country fundamentals and news, emits G7 macroeconomic scenarios...

https://arxiv.org/abs/2512.07867

15 CONF

Do Latent Tokens Think? A Causal and Adversarial Analysis of Chain-of-Continuous-Thought

ArXiv CS.AI 50%

arXiv:2512.21711v2 Announce Type: replace-cross Abstract: Latent tokens are gaining attention for enhancing reasoning in large language models (LLMs), yet their internal mechanisms remain unclear. This paper examines the problem from a reliability perspective, uncovering...

https://arxiv.org/abs/2512.21711

16 CONF

NoiseFormer -- Noise Diffused Symmetric Attention Transformer

ArXiv CS.AI 50%

arXiv:2601.11619v2 Announce Type: replace-cross Abstract: Transformer architecture has been very successful long runner in the field of Deep Learning (DL) and Large Language Models (LLM) because of its powerful attention-based learning and parallel-natured architecture. As the...

https://arxiv.org/abs/2601.11619

17 CONF

Multi-Agentic AI for Fairness-Aware and Accelerated Multi-modal Large Model Inference in Real-world Mobile Edge Networks

ArXiv CS.AI 50%

arXiv:2602.07215v2 Announce Type: replace-cross Abstract: Generative AI (GenAI) has transformed applications in natural language processing and content creation, yet centralized inference remains hindered by high latency, limited customizability, and privacy concerns. Deploying...

https://arxiv.org/abs/2602.07215

18 CONF

How well are open sourced AI-generated image detection models out-of-the-box: A comprehensive benchmark study

ArXiv CS.AI 50%

arXiv:2602.07814v2 Announce Type: replace-cross Abstract: As AI-generated images proliferate across digital platforms, reliable detection methods have become critical for combating misinformation and maintaining content authenticity. While numerous deepfake detection methods...

https://arxiv.org/abs/2602.07814

19 CONF

What Makes Synthetic Hard Negatives Work in Vision-Language Pretraining?

ArXiv CS.AI 50%

arXiv:2610.09700v2 Announce Type: replace-cross Abstract: Synthetic hard negatives generated in the representation space have proven effective for unimodal self-supervised learning, but transferring this idea to vision-language pretraining is not straightforward. We analyze six...

https://arxiv.org/abs/2610.09700

20 CONF

Good Scores, Bad Data: A Metric for Multimodal Coherence

ArXiv CS.AI 50%

arXiv:2603.25924v2 Announce Type: replace-cross Abstract: Multimodal AI systems are evaluated by downstream task accuracy, but high accuracy does not mean the underlying data is coherent. A model can score well on Visual Question Answering (VQA) while its inputs contradict each...

https://arxiv.org/abs/2603.25924

21 CONF

Beyond Feature Fusion: Contextual Bayesian PEFT for Multimodal Uncertainty Estimation

ArXiv CS.AI 50%

arXiv:2604.16615v2 Announce Type: replace-cross Abstract: We introduce CoCo-LoRA, a multimodal, uncertainty-aware parameter-efficient fine-tuning method for text prediction tasks accompanied by audio context. Existing PEFT approaches such as LoRA are efficient but typically...

https://arxiv.org/abs/2604.16615

22 CONF

SentryVLA: Uncertainty-Guided Action Selection and State Monitoring for Vision-Language-Action Robotic Control

ArXiv CS.AI 50%

arXiv:2604.16677v2 Announce Type: replace-cross Abstract: Vision-language-action (VLA) models have emerged as generalist robotic controllers that map visual observations and language instructions to continuous actions. However, pretrained VLAs typically provide limited...

https://arxiv.org/abs/2604.16677

23 CONF

GPT-Image-2 in the Wild: A Twitter Dataset of Self-Reported AI-Generated Images from the First Week of Deployment

ArXiv CS.AI 50%

arXiv:2604.25370v3 Announce Type: replace-cross Abstract: The release of GPT-image-2 by OpenAI marks a watershed moment in AI-generated imagery: the boundary between photographic reality and synthetic content has never been more difficult to discern. We introduce the GPT-Image-2...

https://arxiv.org/abs/2604.25370

24 CONF

Why Geometric Continuity Emerges in Deep Neural Networks: Residual Connections and Rotational Symmetry Breaking

ArXiv CS.AI 50%

arXiv:2605.04971v2 Announce Type: replace-cross Abstract: Weight matrices in deep networks exhibit geometric continuity---principal singular vectors of adjacent layers point in similar directions. While this property has been widely observed, its origin remains unexplained....

https://arxiv.org/abs/2605.04971

25 CONF

BatMIL: Geometry-Embedding State Space Model for Whole-Slide Image Representation

ArXiv CS.AI 50%

arXiv:2605.05164v2 Announce Type: replace-cross Abstract: Accurate analysis of histopathological images is critical for disease diagnosis and treatment planning. Whole-slide images (WSIs), which digitize tissue specimens at gigapixel resolution, are fundamental to this process...

https://arxiv.org/abs/2605.05164

26 CONF

MicroWorld: Empowering Multimodal Large Language Models to Bridge the Microscopic Domain Gap with Multimodal Attributed Property Graph

ArXiv CS.AI 50%

arXiv:2605.10120v2 Announce Type: replace-cross Abstract: Multimodal large language models (MLLMs) show remarkable potential for scientific reasoning, yet their performance in specialized domains such as microscopy remains limited by the scarcity of domain-specific training data...

https://arxiv.org/abs/2605.10120

27 CONF

Coherence-Guided Pyramid Self-Contrastive Disentanglement for Ultrasound Imaging via Single-Shot Test-Time Training

ArXiv CS.AI 50%

arXiv:2605.12567v3 Announce Type: replace-cross Abstract: The inherent electronic noise, side lobes, and speckle noise complicate clinical interpretation of ultrasound images. Conventional denoising methods rely on explicit noise assumptions that may underperform in composite...

https://arxiv.org/abs/2605.12567

28 CONF

Where and What to Reward? Understanding Generalization Between Verifiable and Open-Ended Domains in Reinforcement Learning

ArXiv CS.AI 50%

arXiv:2605.15464v2 Announce Type: replace-cross Abstract: Reinforcement learning with verifiable rewards (RLVR) is widely used to improve the reasoning capability of language models, since model-generated answers can be easily verified by a rule-based verifier. In contrast,...

https://arxiv.org/abs/2605.15464

29 CONF

Prompt PerturbAItions: Mechanisms of LLM Response Inconsistencies

ArXiv CS.AI 50%

arXiv:2605.30913v2 Announce Type: replace-cross Abstract: Large language models (LLMs) are increasingly deployed in conversational settings where user phrasing ranges from polite to casual to toxic, yet less is known about whether such lexical variation in otherwise semantically...

https://arxiv.org/abs/2605.30913

30 CONF

LUMA: Benchmarking Segmentation via a Lightweight Universal Mask Adapter

ArXiv CS.AI 50%

arXiv:2607.00687v2 Announce Type: replace-cross Abstract: Comparing transformer backbones for image segmentation is confounded: each is paired with a different decoder, recipe, and pretraining, so reported differences rarely reflect the backbone itself. We introduce the...

https://arxiv.org/abs/2607.00687

31 CONF

SLAB: Time-Lag-Aware Deep Reinforcement Learning for Flexible Job-Shop Scheduling in PPVC Module Factories

ArXiv CS.AI 50%

arXiv:2607.11725v2 Announce Type: replace-cross Abstract: Prefabricated prefinished volumetric construction (PPVC) builds room-sized modules in factories that operate as flexible job shops. A module waits 12 to 48 h after casting, waterproofing, and painting, while its station...

https://arxiv.org/abs/2607.11725

32 CONF

What the Abstention Option Measures: Refusal Loss, Sampling Churn, and Option Capture in Distilled Small Language Models

ArXiv CS.AI 50%

arXiv:2607.28639v3 Announce Type: replace-cross Abstract: Small instruction-tuned language models reach phones, cars, and offline assistants through knowledge distillation, and whether a distilled model is fair enough to ship is often decided by a multiple-choice bias score on a...

https://arxiv.org/abs/2607.28639

33 CONF

A False Average: Pooled CoT-Monitor Accuracy Conceals a Reasoning-Dependent Fragility

ArXiv CS.AI 50%

arXiv:2608.00583v4 Announce Type: replace-cross Abstract: Chain-of-thought (CoT) monitors are reported by their aggregate accuracy on a pool of reward hacks. We show that this number is a false average. On Terminal Wrench, about 77% of hacks are given away by the actions alone,...

https://arxiv.org/abs/2608.00583

34 CONF

EntailLLM: Verifying Agent-Generated Vulnerability Discovery Paths with Domain Knowledge via Logic Programming

ArXiv CS.AI 50%

arXiv:2608.01763v2 Announce Type: replace-cross Abstract: Large language model agents are increasingly used to analyze potential software vulnerabilities, but their outputs can silently violate domain knowledge, limiting their reliability in safety-critical settings such as...

https://arxiv.org/abs/2608.01763

35 CONF

What Can One Sentence Fix? A Failure Taxonomy and a Minimal Language Correction for Frozen VLA Policies

ArXiv CS.AI 50%

arXiv:2608.29967v2 Announce Type: replace-cross Abstract: When a deployed vision-language-action (VLA) policy fails, fine-tuning it again for every new setting is costly, yet the person watching it can usually say what went wrong in one sentence. We ask what such a sentence can...

https://arxiv.org/abs/2608.29967

36 CONF

Trading Agents or Market Crashers? Dissecting Robustness and Security Failures in Academic Financial LLM Trading Schemes

ArXiv CS.AI 50%

arXiv:2609.19705v3 Announce Type: replace-cross Abstract: Large language model (LLM) agents already trade real capital, and builders copy a few published designs that share models and news feeds, so one shock reaches many agents at once. We present FARSIGHT (Financial Agent...

https://arxiv.org/abs/2609.19705

37 CONF

A Large-Scale Benchmark and Risk Assessment of Traffic Analysis Attacks on Cloud LLM Services

ArXiv CS.AI 50%

arXiv:2609.31877v2 Announce Type: replace-cross Abstract: Cloud-based Large language model (LLM) services create a network-level traffic side channel that can expose model, prompt, and task behavior despite encryption. From packet sizes, directions, timing, and burst structure...

https://arxiv.org/abs/2609.31877

38 CONF

LLMs are not stochastic parrots: Evidence for meaning-mediated abstraction from conlang-like tasks

ArXiv CS.AI 50%

arXiv:2609.34187v4 Announce Type: replace-cross Abstract: The strong version of the stochastic parrot argument claims that, although large language models (LLMs) may exceed rote regurgitation, they cannot move beyond statistical pattern matching into abstraction or reasoning,...

https://arxiv.org/abs/2609.34187

39 CONF

What Does a Harness Repair? A Preregistered Study of Visibility, Baseline Adequacy and Evaluation Defects

ArXiv CS.AI 50%

arXiv:2610.05533v2 Announce Type: replace-cross Abstract: Harness search keeps a change to the prompts, token budgets or parsers around a frozen model if the change raises a score. Such a gain can come from answers the parser could not read before, a weak comparison or an...

https://arxiv.org/abs/2610.05533

40 CONF

Beyond LLM-GA: Secure Fluid Antenna Systems with ReEvo-Designed Memetic Algorithm

ArXiv CS.AI 50%

arXiv:2610.10235v2 Announce Type: replace-cross Abstract: Fluid antenna systems (FASs) offer significant spatial flexibility, yet securing them against eavesdropping is critical for practical FAS deployment in military, satellite, and internet-of-things networks. Although large...

https://arxiv.org/abs/2610.10235

41 CONF

MAGNET: Manifold-Aware Graph Diffusion Network for Connectome Generation

ArXiv CS.AI 50%

arXiv:2610.17014v2 Announce Type: replace-cross Abstract: Functional brain connectome represents neural connectivity as a matrix of pairwise interactions between brain regions. Generation of functional connectomes is not only a question of validity; rather, having satisfied the...

https://arxiv.org/abs/2610.17014

42 CONF

SWIFT: Self-guided Weak-to-strong dIFfusion models via Target extrapolation

ArXiv CS.AI 50%

arXiv:2610.25844v2 Announce Type: replace-cross Abstract: Training high-quality diffusion models is computationally expensive, and inference-time guidance often adds further sampling cost. In this work, we introduce Swift, a simple yet effective self-guided training method that...

https://arxiv.org/abs/2610.25844

43 CONF

Parameter-Efficient Adaptation of Time Series Foundation Models for Electrical Load Forecasting

ArXiv CS.AI 50%

arXiv:2610.28200v2 Announce Type: replace-cross Abstract: Accurate electrical load forecasting supports power-system operation across countries, regions, cities, and buildings, but repeated development of target-specific models is costly. Time series foundation models (TSFMs)...

https://arxiv.org/abs/2610.28200

44 CONF

From AI-Free Gates to AI-First Work: How Engineering Students in India Navigate University Education and Career Preparation Without AI Guidelines

ArXiv CS.AI 50%

arXiv:2610.29140v2 Announce Type: replace-cross Abstract: Engineering graduates in India enter a technology sector that is being reshaped by generative AI, yet the route into that sector runs through examinations and campus placement screens that largely exclude AI. Institutions...

https://arxiv.org/abs/2610.29140

45 CONF

Direct Forcing: Single-Stage Training with Representation-Space Supervision for Few-Step Autoregressive Video Generation

ArXiv CS.AI 50%

arXiv:2610.30861v2 Announce Type: replace-cross Abstract: Few-step autoregressive video generation enables low-latency streaming, yet prevailing training pipelines separate causal adaptation from distribution matching. Recent representation-based text-to-video methods retain...

https://arxiv.org/abs/2610.30861