NeoSignal

AI Intelligence Daily · Multi-Source · Cross-Verified

05 November 2026

136 verified · 241 raw · 5 sources

136	1	135	0
TOTAL	VERIFIED	CONFIRMED	EMERGING

Pipeline

Scrape HackerNews (Show/Top/New), Reddit (r/artificial, r/MachineLearning, r/singularity, r/LocalLLaMA), TechCrunch AI, VentureBeat AI, MIT Tech Review, The Verge AI, Wired AI, ArXiv CS.AI

Filter 40+ AI keyword patterns — LLMs, alignment, safety, research, tooling

Dedup Title-similarity (45% threshold) merges cross-source variants

Score Authenticity = 0.5 base + 0.3 per extra source + diversity + HN score bonuses

Rank Sorted by authenticity. Below 0.25 threshold are dropped.

Authenticity Tiers

VERIFIED	**0.80 - 1.00**	3+ independent sources including media
CONFIRMED	**0.50 - 0.79**	2 sources or 1 high-quality media
EMERGING	**0.25 - 0.49**	Single-source; passes AI keyword filter

Source Breakdown

Source	Articles	Category
ArXiv CS.AI	117	Media
TechCrunch AI	9	Media
The Verge AI	5	Media
HackerNews	5	Community
MIT Tech Review	1	Media

VERIFIED INTELLIGENCE — 1 story

VERI

Nothing teases a screen-free AI device ahead of launch next month

The Verge AI +1 source 80%

Nothing CEO Carl Pei said in a post on X that the company's launching its first AI device in December. That's just a little over a year after Pei announced plans for "AI-native" devices and an AI operating system. He didn't give away much about what exactly the new device is,...

https://www.theverge.com/ai-artificial-intelligence/1026586/nothing-screen-free-ai-devi...

CONFIRMED SIGNALS — 135 stories

CONF

Show HN: Kairos – Open-source risk and fraud decisioning platform

HackerNews 60% HN 104

Hi HN, I'm a long-time Superhuman user, and a few weeks ago I cloned it as a native Mac app. It's keyboard-first, with split inboxes, reminders, and a built-in calendar. It's local-first, including the AI features that use local models shipped in macOS 27, with the option to use...

https://github.com/finic-ai/kairos

CONF

Show HN: AI Gude – my feeds as a daily podcast, revived with AI

HackerNews 50% HN 7

Hi HN. AI Gude turns the news you follow into a short morning podcast, written and read by AI. It is open source and you host it yourself. “Ei Gude” is Hessian for good morning. Some background: I started this about a year and a half ago and never got back to it, because I...

https://github.com/alex0ptr/aigude

CONF

Show HN: Amigent – AI agent and remote desktop for Commodore Amiga over 50 KB/s

HackerNews 50% HN 3

No summary — see link for full article.

https://github.com/akustikrausch/amigent

CONF

Show HN: Homotope, vectorize images to SVG in the browser, MIT engine on NPM

HackerNews 50% HN 3

No summary — see link for full article.

https://github.com/stevenroman91/homotope-engine

CONF

Show HN: Using GPUs for other bs than AI – important prime record for the geeks

HackerNews 50% HN 2

Hey HN, if you want to be new prime gap hero here is the code that runs and gets you biggest merits (everyone knows what that is in prime world). I spent months to tune it, and in the last few days got couple of top 10 records just by testing. So get you precious GPU time, run...

https://github.com/djuk8/beacon/

CONF

Apple is reportedly cutting budgets for its internal AI usage

TechCrunch AI 50%

The company is said to be significantly dropping the amount it's spending on AI for its engineering teams.

https://techcrunch.com/2026/11/04/apple-is-reportedly-cutting-budgets-for-its-internal-...

CONF

OpenAI is already trying to shape the 2028 election

TechCrunch AI 50%

OpenAI’s Chris Lehane wants AI to define the 2028 election, proposing a new social contract that could benefit Americans while also advancing the AI giant’s own business interests.

https://techcrunch.com/2026/11/04/openai-is-already-trying-to-shape-the-2028-election/

CONF

Almost no one uses AI agents to book their Uber, company says

TechCrunch AI 50%

During its busiest week last quarter, Uber said only four people booked a ride using an AI agent.

https://techcrunch.com/2026/11/04/almost-no-one-uses-ai-agents-to-book-their-uber-compa...

0 CONF

Meta acquires AI coding startup Conductor

TechCrunch AI 50%

The startup allows software engineers to deploy teams of agents to work on various aspects of a project simultaneously.

https://techcrunch.com/2026/11/04/meta-acquires-ai-coding-startup-conductor/

1 CONF

Open source AI is booming. Hugging Face explains who profits.

TechCrunch AI 50%

Open source models are becoming more mainstream as enterprises, governments and startups alike rely on the promise of lower costs, more control, and more flexibility. But as increasingly capable models become more freely available, how do companies capture value? And with...

https://techcrunch.com/video/open-source-ai-is-booming-hugging-face-explains-who-profits/

2 CONF

DuckDuckGo pitches itself to schools amid backlash over Google’s AI policy

TechCrunch AI 50%

DuckDuckGo is courting schools frustrated by Google's AI restrictions, pointing educators to its search engine and options to disable AI features.

https://techcrunch.com/2026/11/04/duckduckgo-pitches-itself-to-schools-amid-backlash-ov...

3 CONF

Building on open source AI? Hugging Face’s policy chief on the costs, risks, and trade-offs

TechCrunch AI 50%

Open source models are becoming more mainstream as enterprises, governments and startups alike rely on the promise of lower costs, more control, and more flexibility. But as increasingly capable models become more freely available, how do companies capture value? And with...

https://techcrunch.com/podcast/building-on-open-source-ai-hugging-faces-policy-chief-on...

4 CONF

Google faces backlash for plan that would turn off Search for schools that restrict AI

TechCrunch AI 50%

Politicians on the left, like NYC mayor Zohran Mamdani and former FTC chair Lina Khan, accused Google of acting as an unregulated monopoly.

https://techcrunch.com/2026/11/04/google-faces-backlash-for-plan-that-would-turn-off-se...

5 CONF

Texas was the new epicenter of the AI data center buildout. That might be kaput.

MIT Tech Review 50%

Texas is the fastest-growing data center market in the US and ground zero for Big Tech’s largest AI buildouts. But the blue wave that swept the Lone Star State may jeopardize that position. Democrat James Talarico won the US Senate seat against Ken Paxton. Gina Hinojosa won the...

https://www.technologyreview.com/2026/11/04/1148043/texas-was-the-new-epicenter-of-the-...

6 CONF

Big Tech’s big AI buildout will not end until someone stops it

The Verge AI 50%

After four of the biggest competitors in the AI race held quarterly earnings calls late last month, two things are clear: Big Tech is not done spending record amounts on the AI data center buildout, and the AI data center buildout is not any closer to yielding its promised...

https://www.theverge.com/ai-artificial-intelligence/1027344/amazon-microsoft-alphabet-g...

7 CONF

Gemini for Home now plugs into AT&T’s smart security system

The Verge AI 50%

AT&T's smart home security system now integrates with Gemini for Home. The company says its Connected Life system is the "first" home security product to integrate with Google Home Premium, enabling a bunch of Gemini-powered features that can summarize home events, search your...

https://www.theverge.com/tech/1026787/att-gemini-for-home-connected-life-security-system

8 CONF

ChatGPT will now make you animated GIFs

The Verge AI 50%

ChatGPT can now create GIFs using a prompt or a photo uploaded to the AI chatbot. The update, which is rolling out alongside OpenAI's new Images 2.6 model, allows you to adjust the GIF's animation speed, edit individual frames, and share it with others. Before this change,...

https://www.theverge.com/ai-artificial-intelligence/1026474/openai-chatgpt-animated-gif...

9 CONF

Uber only fulfilled 4 rides from AI agents last quarter

The Verge AI 50%

AI agents won't be stealing Uber's lunch any time soon. Only four rides were arranged by third-party AI agents during the third quarter of 2026, Uber executives disclosed during an earnings call this morning. The volume of Uber Eats orders fulfilled by AI agents was a little...

https://www.theverge.com/transportation/1026306/uber-ai-agents-q3-earnings-disintermedi...

0 CONF

Who May Overrule the Agent? Evidence-Gated Authority in LLM Review Institutions

ArXiv CS.AI 50%

arXiv:2611.05390v1 Announce Type: new Abstract: Multi-agent language-model systems often add a reviewer to improve an upstream agent's decision. Yet a reviewer can both repair errors and replace decisions that were already correct. We study this as a problem of institutional...

https://arxiv.org/abs/2611.05390

1 CONF

Uncertainty-Aware Resource Orchestration for Agentic AI Systems

ArXiv CS.AI 50%

arXiv:2611.05392v1 Announce Type: new Abstract: Agentic artificial intelligence (AI) systems powered by large language models (LLMs) execute as autonomous multi-step pipelines that dynamically invoke tools, integrate observations into memory, and iterate until a task is resolved...

https://arxiv.org/abs/2611.05392

2 CONF

How Well Can Large Language Models Impersonate Individual Authors? A Comparative Evaluation

ArXiv CS.AI 50%

arXiv:2611.05771v1 Announce Type: cross Abstract: This study investigates the authorship impersonation capacity of six widely used large language models (LLMs): Gemini 3 Pro, Claude Opus, GPT 5.2, GPT 5.2 Pro, DeepSeek 3.2, and Grok 4.1 Fast. The evaluation was conducted using...

https://arxiv.org/abs/2611.05771

3 CONF

MASBench: Benchmarking LLM-based Multi-Agent Collaboration under Partial Observability

ArXiv CS.AI 50%

arXiv:2610.04672v2 Announce Type: replace Abstract: Large language models (LLMs) have progressively evolved into the core of autonomous agents. Building on this progress, LLM-based multi-agent systems (MAS) coordinate multiple agents into a synergistic team to accomplish complex...

https://arxiv.org/abs/2610.04672

4 CONF

Bidirectional Centered Score Matching for Discrete Diffusion Distillation

ArXiv CS.AI 50%

arXiv:2611.04279v2 Announce Type: replace Abstract: Discrete diffusion models with parallel decoding are often distilled using token-wise KL objectives. Although token-wise KL exactly matches the corresponding product-form block distributions, its probability-space optimization...

https://arxiv.org/abs/2611.04279

5 CONF

Reconciling Privacy with Provenance for Language Model Watermarks

ArXiv CS.AI 50%

arXiv:2611.06100v1 Announce Type: cross Abstract: Watermarking has great potential for regulating the use of language models (LMs), but its deployment has been marred by controversy, as many users argue that fine-grained tracking of LM usage violates their privacy. This tension...

https://arxiv.org/abs/2611.06100

6 CONF

Whose Transparency is it Anyway? Understanding Generative AI Transparency Needs and Gaps Across Stakeholders

ArXiv CS.AI 50%

arXiv:2611.05549v1 Announce Type: new Abstract: Transparency has become a central principle in AI governance, yet there is little agreement about what transparency means, for whom, and in what contexts. Drawing on interpretive flexibility within the Social Construction of...

https://arxiv.org/abs/2611.05549

7 CONF

AGORA: Reliability-Aware Aggregation of LLM Judges via Item-Response Theory

ArXiv CS.AI 50%

arXiv:2611.05556v1 Announce Type: new Abstract: LLM-as-a-judge is a common substitute for costly human evaluation, and aggregating several judges is the standard remedy for any single judge's bias. Judge pools mix judges of very different reliability, while the human-labeled...

https://arxiv.org/abs/2611.05556

8 CONF

BrainBind: Bridging Synchronized EEG and fMRI through Long-Form Naturalistic Video Representations

ArXiv CS.AI 50%

arXiv:2611.05581v1 Announce Type: new Abstract: Real-world vision is dynamic and scene-rich, while EEG and fMRI resolve it at complementary temporal and spatial scales. Most neural decoding systems, however, align one modality in isolation and mostly against static images or...

https://arxiv.org/abs/2611.05581

9 CONF

Verified but Vacuous: Three Failure Surfaces of LLM-Driven Verified Code Generation

ArXiv CS.AI 50%

arXiv:2611.05590v1 Announce Type: new Abstract: What should ``verified'' mean when evaluating LLM-generated code? Many existing pipelines report a single verification number. We argue this collapses three distinct properties measured in a conditional audit sequence: verifier...

https://arxiv.org/abs/2611.05590

0 CONF

ViBench 2.0: A Benchmark for Long-Horizon Software Building

ArXiv CS.AI 50%

arXiv:2611.05620v1 Announce Type: new Abstract: On Replit, our agentic engineering platform, users now build entire products without writing code, through successive feature requests to an agent. Yet existing coding benchmarks evaluate single, fully specified tasks, and many are...

https://arxiv.org/abs/2611.05620

1 CONF

Signals: Continuous Diagnostics of AI Agents

ArXiv CS.AI 50%

arXiv:2611.04507v2 Announce Type: replace Abstract: AI agents can make incorrect decisions or waste resources even when their tools and infrastructure work normally. Their nondeterministic behavior makes these problems difficult to reproduce and diagnose. We developed Signals to...

https://arxiv.org/abs/2611.04507

2 CONF

From Handcrafted to Generative: RLVR-Optimized Icon Generation for Alexa+

ArXiv CS.AI 50%

arXiv:2611.06123v1 Announce Type: new Abstract: Alexa+ devices with screens display animated icons during conversations to visually reinforce what is being said. The original set of these icons was hand-drawn by designers one at a time, a process that produces polished results...

https://arxiv.org/abs/2611.06123

3 CONF

Attributing Agent Failures to Their History Needs an Executed Counterfactual: Toward a Reporting Standard

ArXiv CS.AI 50%

arXiv:2611.05650v1 Announce Type: new Abstract: Published explanations of why long-horizon agents fail are commonly taxonomies labeled over completed trajectories: a judge reads a transcript and names the failure. We argue that a claim attributing an agent's behavior to its...

https://arxiv.org/abs/2611.05650

4 CONF

Quantifying Trait-State Hierarchy in EEG Foundation Model and Signal Processing Embeddings

ArXiv CS.AI 50%

arXiv:2611.05652v1 Announce Type: new Abstract: Electroencephalography (EEG) measures brain activity through recording electrical potentials of the scalp. EEG recordings capture two distinct sets of characteristics in the signal: (1) \emph{Traits} -- information that remains...

https://arxiv.org/abs/2611.05652

5 CONF

IDSISA: Influence-Distilled Sharded Training for Efficient Machine Unlearning

ArXiv CS.AI 50%

arXiv:2611.06519v1 Announce Type: cross Abstract: Machine unlearning removes the influence of specific training examples to comply with privacy regulations. Full retraining gives strong guarantees but is computationally prohibitive at scale; approximate methods are efficient but...

https://arxiv.org/abs/2611.06519

6 CONF

PAYWORTHY: Can LLMs Predict Which Freelance Jobs an AI Agent Will Complete?

ArXiv CS.AI 50%

arXiv:2611.05759v1 Announce Type: new Abstract: Before an AI agent takes on a paid job, someone has to judge whether it is likely to succeed. We test whether frontier language models can make this call before any attempt. We propose PAYWORTHY which contains 1{,}494 real...

https://arxiv.org/abs/2611.05759

7 CONF

FlyHash for LLM Agent Memory: An Empirical Validation and a New Application

ArXiv CS.AI 50%

arXiv:2611.05791v1 Announce Type: new Abstract: Large language model (LLM) agents increasingly rely on external memory stores to retain information across conversations, typically via dense vector embeddings searched with approximate nearest-neighbor methods. The fruit fly's...

https://arxiv.org/abs/2611.05791

8 CONF

Vision Paper: Why Distributed LLM Post-Training and Decentralized Alignment Are Now Essential

ArXiv CS.AI 50%

arXiv:2611.05841v1 Announce Type: new Abstract: This vision paper argues that distributed learning, including federated and decentralized, should become a first-class paradigm in LLM post-training, not merely a privacy add-on. Current methods like RLHF and DPO rest on an...

https://arxiv.org/abs/2611.05841

9 CONF

Auditable Unlearning for LLMs

ArXiv CS.AI 50%

arXiv:2611.05866v1 Announce Type: new Abstract: Machine unlearning aims to remove a user's data from a trained model after a deletion request. Existing methods are majorly evaluated on three axes - how well they forget, how much capability they retain, and how robust they are to...

https://arxiv.org/abs/2611.05866

0 CONF

D3D: Decoupled Diffusion Drafting for Asynchronous Speculative Decoding

ArXiv CS.AI 50%

arXiv:2611.06425v1 Announce Type: new Abstract: Target-conditioned diffusion drafters predict candidate tokens in parallel, yet each new block must wait for updated target-model hidden states. This forces drafting and verification to run sequentially, limiting speedup even...

https://arxiv.org/abs/2611.06425

1 CONF

Autonomous agents for electricity-system planning

ArXiv CS.AI 50%

arXiv:2611.05911v1 Announce Type: new Abstract: Planning electricity systems requires repeated decisions that connect regional resources, investment, operation and network constraints. Here we present an autonomous planning framework that carries prepared regional scenarios...

https://arxiv.org/abs/2611.05911

2 CONF

Beyond Risk Cues: A Mechanism-Level Benchmark and Verifiable Debate for Scam Video Detection on the Web

ArXiv CS.AI 50%

arXiv:2611.06110v1 Announce Type: new Abstract: Online scam videos pose significant threats to individuals and society. Existing studies primarily focus on binary classification, offering limited insights into how scams work and the reasoning behind detection decisions....

https://arxiv.org/abs/2611.06110

3 CONF

SafeStart: Mitigating the Overlooked Safety Cost of Active Expert Reduction in MoE LLM Serving

ArXiv CS.AI 50%

arXiv:2611.06147v1 Announce Type: new Abstract: Mixture-of-Experts (MoE) language models send each token through a few of many expert subnetworks. Serving methods make these models faster by running even fewer experts per token, and they are validated on accuracy benchmarks. We...

https://arxiv.org/abs/2611.06147

4 CONF

A Training-Free Socket for DNA in Frozen Language Models

ArXiv CS.AI 50%

arXiv:2611.05401v1 Announce Type: cross Abstract: Connecting a DNA foundation model to a large language model (LLM) usually requires training a connector, and often the LLM too. We show that a closed-form map is enough for gene-level questions. Our method, the socket, fits one...

https://arxiv.org/abs/2611.05401

5 CONF

STEK: Structural-Terrain-Enhanced Koopman-Inspired Cross-Channel Regularization for Retrospective Battery Multi-Target Estimation

ArXiv CS.AI 50%

arXiv:2611.06299v1 Announce Type: new Abstract: Retrospective battery characterization across heterogeneous cells is challenging because cycling records vary substantially with chemistry, protocol, cell format, capacity, and operating conditions. An endpoint charge--discharge...

https://arxiv.org/abs/2611.06299

6 CONF

Corralling AI Risks with Grassroots Contracts

ArXiv CS.AI 50%

arXiv:2611.06374v1 Announce Type: new Abstract: AI's inflection point is a fountain of alarming risks. Digital platforms in particular are vulnerable to three major AI risks, as those in control of global platforms cede control to AI, of the algorithms, of membership and of...

https://arxiv.org/abs/2611.06374

7 CONF

Does Context Compression Preserve Refusal Alignment?

ArXiv CS.AI 50%

arXiv:2611.06380v1 Announce Type: new Abstract: Context compression reduces inference cost by encoding inputs into compact representations while preserving semantic content. An open question is whether semantic preservation alone is sufficient to maintain downstream behaviours...

https://arxiv.org/abs/2611.06380

8 CONF

FPTR: Feasibility-Preserving Transactional Refinement for Deadline-Constrained Joint Beam and Resource Allocation

ArXiv CS.AI 50%

arXiv:2611.06491v1 Announce Type: new Abstract: Short-interval wireless schedulers must return a complete, feasible beam-and-resource allocation before each execution deadline, and a late or invalid output gives no utility in the current control cycle. Searching longer can...

https://arxiv.org/abs/2611.06491

9 CONF

UpSafe$^\circ$C: Upcycling for Controllable Safety in Large Language Models

ArXiv CS.AI 50%

arXiv:2510.02194v2 Announce Type: replace Abstract: Large Language Models (LLMs) have achieved remarkable progress across a wide range of tasks, but remain vulnerable to safety risks such as harmful content generation and jailbreak attacks. Existing safety techniques---including...

https://arxiv.org/abs/2510.02194

0 CONF

RAGAM-AI: A Hybrid Reinforcement Learning based Agentic approach powered by Generative AI for Edge Resource Management

ArXiv CS.AI 50%

arXiv:2611.06554v1 Announce Type: new Abstract: Resource management for distributed edge networks hosting different services has been handled efficiently using Reinforcement learning (RL). Reward generalization in complex workflows is still an open issue within the RL realm....

https://arxiv.org/abs/2611.06554

1 CONF

One Scale, Four Texts: Gaps, Overlaps and Silent Repairs in Published Grading Schemes

ArXiv CS.AI 50%

arXiv:2611.06569v1 Announce Type: new Abstract: Category systems often reach their users as text, and readers apply them with tacit repairs. We ask what such a text determines once the repairs are removed, using diabetic retinopathy grading, where one international scale...

https://arxiv.org/abs/2611.06569

2 CONF

From Certified Pairs to a Certified Hub: Reasoner-Checked Crosswalks for the Information Exchange Standard

ArXiv CS.AI 50%

arXiv:2611.06628v1 Announce Type: new Abstract: A national digital twin was conceived as an ecosystem of twins connected through mappings to a common foundation, but the roadmap that promises to confirm those mappings are valid never says what valid means. We give the test a...

https://arxiv.org/abs/2611.06628

3 CONF

Where Search Agents Go Off the Rails: Graph-Grounded Step-Level Verification

ArXiv CS.AI 50%

arXiv:2611.06662v1 Announce Type: new Abstract: Search-augmented reasoning systems can reach correct answers while making errors in intermediate steps. Reliable auditing must check whether evidence D covers the question Q, reasoning T addresses Q, and T is grounded in...

https://arxiv.org/abs/2611.06662

4 CONF

LifeVal: Evaluating AI Model Performance on Personalized Agentic Consumer Tasks

ArXiv CS.AI 50%

arXiv:2611.06687v1 Announce Type: new Abstract: We introduce LifeVal, a benchmark evaluating AI model performance on personalized agentic consumer tasks. LifeVal covers the most valuable life tasks that a personalized digital superassistant could automate, representing over 1.7...

https://arxiv.org/abs/2611.06687

5 CONF

Unsupervised Neural Operators for Rotating Relativistic Stars: High-Fidelity Equation of State Inference and Pulsar Spin Evolution at Scale

ArXiv CS.AI 50%

arXiv:2611.05790v1 Announce Type: cross Abstract: Observed populations of rotating neutron stars~(NSs) with millisecond-order periods are powerful probes of fundamental physics. Mass-radius measurements inform poorly constrained properties of degenerate nuclear matter, whereas...

https://arxiv.org/abs/2611.05790

6 CONF

CompressionBench: How can your coding agents compress LLMs?

ArXiv CS.AI 50%

arXiv:2611.06858v1 Announce Type: new Abstract: Recursive self-improvement requires AI systems to improve the models and infrastructure used by subsequent systems. LLM compression offers a concrete efficiency objective: retain a pretrained model's behavior within a smaller model...

https://arxiv.org/abs/2611.06858

7 CONF

On Turn-Level Best-of-N Sampling for Scaling LLM Agents

ArXiv CS.AI 50%

arXiv:2611.06868v1 Announce Type: new Abstract: In this work, we evaluate a simple inference-time scaling approach based on \textit{turn-level best-of-N sampling} (\method) for large language model (LLM) agents on complex, long-horizon tasks. We instantiate this using a...

https://arxiv.org/abs/2611.06868

8 CONF

Dull, Dirty, Dangerous? Working Conditions and the Limits of Occupational AI Exposure

ArXiv CS.AI 50%

arXiv:2611.06888v1 Announce Type: new Abstract: Automation is often justified as relief from dull, dirty, and dangerous work. I examine how these conditions relate to occupational AI exposure using 21 O*NET items across 754 US occupations. Higher exposure is associated with...

https://arxiv.org/abs/2611.06888

9 CONF

BEAKER: An Expert-Curated Benchmark for Embodied Brains in Self-Driving Chemical Laboratories

ArXiv CS.AI 50%

arXiv:2611.06911v1 Announce Type: new Abstract: Self-Driving Chemical Laboratories (SDCLs) are moving chemical experimentation toward embodied automation, where machines must perceive laboratory scenes, reason about experimental states, and act under physical, procedural, and...

https://arxiv.org/abs/2611.06911

0 CONF

MvBody: Multi-View-Based Hybrid Transformer Using Optical 3D Body Scan for Explainable Cesarean Section Prediction

ArXiv CS.AI 50%

arXiv:2511.03212v1 Announce Type: cross Abstract: Accurately assessing the risk of cesarean section (CS) delivery is critical, especially in settings with limited medical resources, where access to healthcare is often restricted. Early and reliable risk prediction allows...

https://arxiv.org/abs/2511.03212

1 CONF

Hallucination Is Linearly Decodable from Mid-Layer Hidden States in Quantized LLMs

ArXiv CS.AI 50%

arXiv:2606.02628v1 Announce Type: cross Abstract: We investigate whether open-source LLMs encode a linearly separable truthfulness signal in their hidden states, and at which network depth this signal is strongest. Across three 7B--8B instruction-tuned models (Llama-3.1-8B,...

https://arxiv.org/abs/2606.02628

2 CONF

AI-Driven SaaS Platform Integration: AI Express for Unified Generative AI Services with Pay-As-You-Go Credit Systems and Scalable Microservice Architecture

ArXiv CS.AI 50%

arXiv:2611.05389v1 Announce Type: cross Abstract: Generative AI is currently fragmented, with separate subscriptions required for text generation, image synthesis, and code assistance, leading to high costs and workflow disruptions. AI Express is a SaaS platform that...

https://arxiv.org/abs/2611.05389

3 CONF

PEEK: One-Step-Look-Ahead Exposure Bias Correction for Diffusion Sampling

ArXiv CS.AI 50%

arXiv:2611.05407v1 Announce Type: cross Abstract: Diffusion image generators are successful due to accurate denoising networks and strong numerical samplers. Yet the two components are not fully matched. During sampling, the denoiser predicts for the next step based on flawed...

https://arxiv.org/abs/2611.05407

4 CONF

LLM-based NPCs in Video Games: A Systematic Mapping Study

ArXiv CS.AI 50%

arXiv:2611.05427v1 Announce Type: cross Abstract: For decades, non-player characters (NPCs) have been driven by rule-based systems whose personalities and interactions are scripted before release. Large language models (LLMs) promise NPCs that converse openly, adapt to new...

https://arxiv.org/abs/2611.05427

5 CONF

Clips2Story: Training-Free Video Storyboarding and Editing Using Multimodal Retrieval-Embedded Generation

ArXiv CS.AI 50%

arXiv:2611.05992v1 Announce Type: cross Abstract: Video editors often work with large collections of fragmented footage and must select and organize relevant clips into a coherent story. However, many existing video editing and summarization methods assume a temporally ordered...

https://arxiv.org/abs/2611.05992

6 CONF

Hybrid LLM-Deterministic Pipelines for Enterprise-Scale Code Migration at Bloomberg

ArXiv CS.AI 50%

arXiv:2611.05457v1 Announce Type: cross Abstract: Large organizations carry a long tail of infrastructure migrations that are worth doing but rarely urgent enough to compete with feature work, because each one costs significant engineering time per repository. This paper reports...

https://arxiv.org/abs/2611.05457

7 CONF

TSFM Meets LLM: Context as Covariates

ArXiv CS.AI 50%

arXiv:2611.05515v1 Announce Type: cross Abstract: Textual information is ubiquitous in real-world time series datasets, appearing as news, reports, and annotations that hold key predictive signals not deducible from numerical history alone. Large Language Models (LLMs) interpret...

https://arxiv.org/abs/2611.05515

8 CONF

How Finely Can Vision-Language Models See Biology? Bio-VQA: A Benchmark for Taxonomic, Morphological, Spatial and Individual Reasoning

ArXiv CS.AI 50%

arXiv:2611.05559v1 Announce Type: cross Abstract: In biodiversity monitoring, a single image raises several questions: which taxon, how many, which trait, and whether the individual has been seen before. Vision-language models (VLMs) could answer all of these from the image, but...

https://arxiv.org/abs/2611.05559

9 CONF

Discarded or Unreachable? Diagnosing and Repairing Spatial Relation Failure in Vision--Language Similarity

ArXiv CS.AI 50%

arXiv:2611.05565v1 Announce Type: cross Abstract: Contrastive vision-language models score left/right relations near chance, but the symptom has two causes that call for opposite repairs: the relation may survive in the pooled embedding yet be unreachable by native image-text...

https://arxiv.org/abs/2611.05565

0 CONF

Just the Facts: How Dialogues with AI Reduce Conspiracy Beliefs

ArXiv CS.AI 50%

arXiv:2611.05572v1 Announce Type: cross Abstract: Personalized dialogues with AI substantially reduce conspiracy beliefs. To identify why, we conducted two experiments with U.S. adults recruited online (total N = 2,193) that removed possible causes one at a time. The debunking...

https://arxiv.org/abs/2611.05572

1 CONF

Multimodal Fusion Methods for Whole-Slide Images and Clinical Data in Inflammatory Bowel Disease

ArXiv CS.AI 50%

arXiv:2611.05579v1 Announce Type: cross Abstract: Multimodal AI fusion of histology and clinical data is an active research area in computational pathology, but fusion methods are rarely compared under identical conditions and across different clinical endpoints. We present a...

https://arxiv.org/abs/2611.05579

2 CONF

TACT: A Pragmatic Approach to Qualitative Analysis with an AI Council

ArXiv CS.AI 50%

arXiv:2611.05702v1 Announce Type: cross Abstract: Qualitative data can tell software teams what their stakeholders need, yet rigorously analyzing thousands of qualitative responses takes expert time that few teams can afford or justify. Teams therefore often face a false choice:...

https://arxiv.org/abs/2611.05702

3 CONF

Do Language Models Know Which Constraints Belong Together?

ArXiv CS.AI 50%

arXiv:2611.05726v1 Announce Type: cross Abstract: Dantzig-Wolfe reformulation strengthens the linear relaxation of a mixed-integer program by convexifying a chosen set of constraint families, and every candidate costs a column generation run. The choice is hard when families...

https://arxiv.org/abs/2611.05726

4 CONF

Learning Hierarchical Patch Splitting Policies for Faster Vision Transformers

ArXiv CS.AI 50%

arXiv:2611.05742v1 Announce Type: cross Abstract: Vision Transformers (ViTs) are bottlenecked by the length of their input token sequence. Standard patchification fixes this length with a single patch size applied across the image, regardless of regional content. Prior work on...

https://arxiv.org/abs/2611.05742

5 CONF

Text-Only Domain Adaptation for Low-Resource Medical ASR: Shallow Fusion and Density Ratio Rescoring for Hausa

ArXiv CS.AI 50%

arXiv:2611.05744v1 Announce Type: cross Abstract: Automatic speech recognition (ASR) systems often struggle with specialised medical vocabulary, especially in low-resource languages where domain-specific speech data are scarce. We investigate text-only domain adaptation for...

https://arxiv.org/abs/2611.05744

6 CONF

Pairing Expert Critique with LLM-Driven Refinement: An Iterative Approach for Curating Quality AI Evaluation Datasets

ArXiv CS.AI 50%

arXiv:2611.05749v1 Announce Type: cross Abstract: Evaluating and improving domain-specific LLM responses requires high-quality reference data and rubrics, yet expert authoring is slow and expensive, and synthetic datasets may fail to capture practitioner preferences unless they...

https://arxiv.org/abs/2611.05749

7 CONF

DARE: Difficulty Analysis for Repair Enhancement

ArXiv CS.AI 50%

arXiv:2611.05797v1 Announce Type: cross Abstract: LLM-based coding agents are primarily evaluated by their solve rates. However, this metric does not reveal how difficult the solved tasks are or account for differences intask difficulty across benchmarks. Determining the...

https://arxiv.org/abs/2611.05797

8 CONF

Matching the Training-Time Weight Rule in Inference-Time Expert Deactivation

ArXiv CS.AI 50%

arXiv:2611.05804v1 Announce Type: cross Abstract: Expert deactivation at inference time changes both token routing and expert weights, and its effect on performance depends on how these changes are implemented. We compare removal, freezing, remaining-expert renormalization, and...

https://arxiv.org/abs/2611.05804

9 CONF

When File Ingestion Misleads Language Models: Measuring and Mitigating Divergences in Office Documents

ArXiv CS.AI 50%

arXiv:2611.05806v1 Announce Type: cross Abstract: Office documents are now routinely uploaded to large language models (LLMs) for processing. However, an LLM does not always see what a human sees, because it typically reads the text that an extractor produces from the document's...

https://arxiv.org/abs/2611.05806

0 CONF

Generative AI for urban planning: synthesizing land-use maps with satellite imagery via DualControlNet

ArXiv CS.AI 50%

arXiv:2611.05859v1 Announce Type: cross Abstract: Various land-use models characterize land-use patterns and facilitate scenario planning. However, existing models are largely predictive, forecasting transitions based on numeric data without synthesizing new land-use maps from...

https://arxiv.org/abs/2611.05859

1 CONF

AI-Assisted Classification under Correlation Neglect and Trust

ArXiv CS.AI 50%

arXiv:2611.05870v1 Announce Type: cross Abstract: When does an AI assistant improve human decision-making, and when does it hurt? We argue that the answer depends not just on AI accuracy but also on how much its evidence overlaps with the human's and how the human combines the...

https://arxiv.org/abs/2611.05870

2 CONF

Do Saliency Maps Fail Dense Breasts? A Chance-Normalized, Confounder-Aware Audit of Explanation Fairness in Mammography

ArXiv CS.AI 50%

arXiv:2611.05909v1 Announce Type: cross Abstract: Saliency maps, offered as evidence that mammography classifiers use the cancer, are untested in dense breasts, which can hide cancers. We tested whether they locate cancers worse there, and why. On 1,345 malignant CBIS-DDSM...

https://arxiv.org/abs/2611.05909

3 CONF

Intra-Architecture Source Attribution for Stable Diffusion: A New Task with Benchmark Dataset and Baseline

ArXiv CS.AI 50%

arXiv:2611.06000v1 Announce Type: cross Abstract: Designing systems that attribute a generated image to its source model has lately become a major area of AI research. Most existing datasets and systems related to image source attribution are targeted towards the detection of...

https://arxiv.org/abs/2611.06000

4 CONF

Practical Estimation of the Bayes Optimal Fairness-Accuracy Tradeoff with Soft Labels

ArXiv CS.AI 50%

arXiv:2611.06027v1 Announce Type: cross Abstract: There is a fundamental limit to the best prediction error of any model for a given data distribution. In classification, the Bayes error quantifies this limit and serves as both a benchmark for trained models and a criterion to...

https://arxiv.org/abs/2611.06027

5 CONF

EvoWorld: Towards Out-of-Sight World Evolution in Generative Video Models

ArXiv CS.AI 50%

arXiv:2611.06748v1 Announce Type: cross Abstract: Generative video models are increasingly used as visual world simulators, yet their primary strength lies in rendering plausible pixels rather than reasoning about how a world should evolve. Coherent world evolution requires more...

https://arxiv.org/abs/2611.06748

6 CONF

The Harness Is the Model: Agent Scaffolding Changes a Fixed LLM's Compliance with Banking Conduct Rules, Unevenly Across Models

ArXiv CS.AI 50%

arXiv:2611.06116v1 Announce Type: cross Abstract: Banks validate language models, but they deploy systems: a model inside a harness that decides what it sees, which tools it can call and who checks its actions. We test whether that harness changes compliance with banking conduct...

https://arxiv.org/abs/2611.06116

7 CONF

Grounding in the Dark: Text-Only Training for Zero-Shot Video Temporal Localization

ArXiv CS.AI 50%

arXiv:2611.06149v1 Announce Type: cross Abstract: Video-LLMs are typically trained on paired video-text data for temporal tasks, yet collecting precise temporal annotations and training on large-scale video datasets are both costly. We investigate whether a text-only LLM, with...

https://arxiv.org/abs/2611.06149

8 CONF

Reasoning-Aware Cross-Modal Association Unlearning in Multimodal Large Reasoning Models

ArXiv CS.AI 50%

arXiv:2611.06168v1 Announce Type: cross Abstract: Multimodal large reasoning models (MLRMs) integrate visual inputs with textual and parametric knowledge, allowing sensitive information to be accessed through cross-modal associations during reasoning. This raises privacy...

https://arxiv.org/abs/2611.06168

9 CONF

AssameseLM-E2B: An Open-Weight Assamese Language Model for Efficient Consumer-Grade Deployment

ArXiv CS.AI 50%

arXiv:2611.06170v1 Announce Type: cross Abstract: While Large Language Models (LLMs) excel in high-resource languages, their performance often degrades in low-resource languages due to tokenization inefficiencies, generative instability, and linguistic interference. This paper...

https://arxiv.org/abs/2611.06170

0 CONF

Editing Large Language Models with Geometry-Aware Regularization

ArXiv CS.AI 50%

arXiv:2611.06185v1 Announce Type: cross Abstract: Knowledge editing has become a key technique for updating knowledge in large language models. However, existing methods neglect the geometric structure of the knowledge representation space, leading to two key limitations: they...

https://arxiv.org/abs/2611.06185

1 CONF

Rubric Curriculum RL: Exploiting the Generation-Verification Gap in Non-Verifiable Domains

ArXiv CS.AI 50%

arXiv:2611.06244v1 Announce Type: cross Abstract: Reinforcement learning with verifiable rewards (RLVR) has led to significant gains in math and code generation. Extending these gains to open-ended domains remains challenging: ground-truth verification is unavailable, human...

https://arxiv.org/abs/2611.06244

2 CONF

NextGAN: Towards Simple and Scalable GAN Training for Image and Video Generation

ArXiv CS.AI 50%

arXiv:2611.06270v1 Announce Type: cross Abstract: Adversarial training enables one- and few-step image and video generation, but conventional adversarial losses supervise critic values without directly specifying the gradients used to update the generator. We present NEXTGAN,...

https://arxiv.org/abs/2611.06270

3 CONF

A.S.E 2.0: Understanding How Repository Context Impacts the Security of LLM-Generated Code

ArXiv CS.AI 50%

arXiv:2611.06294v1 Announce Type: cross Abstract: Large language model (LLM) coding assistants and agents construct a repository context before generating code. For security-sensitive code, this context decides whether the model sees the project's safeguards (e.g., authorization...

https://arxiv.org/abs/2611.06294

4 CONF

Prism: Bidirectional Self-Teaching for Unified Multimodal Models

ArXiv CS.AI 50%

arXiv:2611.06343v1 Announce Type: cross Abstract: Unified Multimodal Models (UMMs) aim to support visual understanding and generation within a single model. Existing post-training methods either use understanding as a fixed teacher to improve generation, or rely on ground-truth...

https://arxiv.org/abs/2611.06343

5 CONF

Today's Work, Tomorrow's Workforce: Incentivizing Learning in AI-Assisted Work

ArXiv CS.AI 50%

arXiv:2611.06346v1 Announce Type: cross Abstract: At work, people develop skills through experience: novices acquire skills by completing tasks, while experts retain theirs through continued practice. AI assistance changes this process by offering workers a choice between...

https://arxiv.org/abs/2611.06346

6 CONF

Traffic-VL: Reinforced Vision-Language Models for Generalizable End-to-End Traffic Signal Control

ArXiv CS.AI 50%

arXiv:2611.06353v1 Announce Type: cross Abstract: Existing traffic signal control methods typically rely on structured traffic states produced by dedicated perception pipelines. Emerging visual approaches often use idealized simulator images and convert them into intermediate...

https://arxiv.org/abs/2611.06353

7 CONF

A Bag of Tricks for Defending Deep Learning-based Face Recognition Pipelines against Backdoor Attacks

ArXiv CS.AI 50%

arXiv:2611.06412v1 Announce Type: cross Abstract: The protection of mature, multi-stage Face Recognition Systems is crucial to system providers. This technical report targets the threat model most relevant to industry practitioners: although model training is performed in-house,...

https://arxiv.org/abs/2611.06412

8 CONF

Mind Control of LLMs via Closed-Loop Stimulation: Out-of-Distribution Robustness

ArXiv CS.AI 50%

arXiv:2611.06418v1 Announce Type: cross Abstract: Brains and large language models need not be considered equivalent systems to reveal an interpretable commonality: both exhibit high-dimensional, nonlinear dynamics whose behavior can only be modeled imperfectly. This creates a...

https://arxiv.org/abs/2611.06418

9 CONF

Within-Family EfficientNet Scaling under Fixed-Resolution Transfer Learning for Spectral-Domain Optical Coherence Tomography Retinal Disease Classification

ArXiv CS.AI 50%

arXiv:2611.06452v1 Announce Type: cross Abstract: Background/Objectives: CNNs perform well in retinal disease classification from spectral-domain optical coherence tomography (SD-OCT), but whether EfficientNet's natural-image scaling hierarchy transfers to medical-image transfer...

https://arxiv.org/abs/2611.06452

00 CONF

Microdroplet Generator Surrogate Enhancement and Inverse Design Through an Uncertainty-Guided Framework

ArXiv CS.AI 50%

arXiv:2611.06464v1 Announce Type: cross Abstract: Microfluidic droplet generation is critical for many applications, yet efficient device design in high-dimensional parameter spaces remains computationally prohibitive due to complex multiphase flow dynamics. This work presents a...

https://arxiv.org/abs/2611.06464

01 CONF

ArcRAG: Agentic Compositional Retrieval over Knowledge Graphs

ArXiv CS.AI 50%

arXiv:2611.06481v1 Announce Type: cross Abstract: Graph-based retrieval-augmented generation (GraphRAG) answers questions by retrieving over a knowledge graph built from a corpus. Yet retrieval strategies are designed around the structure their own construction pipeline...

https://arxiv.org/abs/2611.06481

02 CONF

NaijaFacts: A Framework for Disentangling Language from Knowledge in Facts about Nigeria

ArXiv CS.AI 50%

arXiv:2611.06504v1 Announce Type: cross Abstract: Multilingual benchmarks for African languages mostly translate generic questions into local languages, measuring linguistic competence but not whether models know facts about Africa. We introduce \naija, a benchmark of factual...

https://arxiv.org/abs/2611.06504

03 CONF

Signed, Sealed, Leaked: Trust Beyond Signature for MCP Toolchains

ArXiv CS.AI 50%

arXiv:2611.06585v1 Announce Type: cross Abstract: AI agents rely on third-party tools, such as Model Context Protocol (MCP) servers, whose publishers can be authenticated and whose execution can be restricted by runtime security policies. Yet authenticity does not guarantee...

https://arxiv.org/abs/2611.06585

04 CONF

Does Frontier AI Strengthen Cyber Defense? Evidence from Project Glasswing

ArXiv CS.AI 50%

arXiv:2611.06629v1 Announce Type: cross Abstract: Defenders and attackers race to find the same software flaws. I use the launch of Project Glasswing, which gave a small group of large US tech firms early access to Claude Mythos Preview, a model for finding software...

https://arxiv.org/abs/2611.06629

05 CONF

Noise-regularized training of quantum circuit Born machines under a finite shot budget

ArXiv CS.AI 50%

arXiv:2611.06630v1 Announce Type: cross Abstract: Quantum circuit Born machines (QCBMs) can learn distributions that are hard to sample from classically, but their value depends on generalizing from finite training data. Here, we study how hardware noise affects this...

https://arxiv.org/abs/2611.06630

06 CONF

AgentTrustCI: Build-Time Policy, Grounding, and Trajectory Tests for Java LLM Agents

ArXiv CS.AI 50%

arXiv:2611.06674v1 Announce Type: cross Abstract: Infrastructure security does not establish whether an LLM agent selected an allowed tool, grounded its answer in tool output, or followed a required workflow. AgentTrustCI is a Java/JUnit 5 framework that expresses these...

https://arxiv.org/abs/2611.06674

07 CONF

Reasoning over Non-Parametric Memory in Multilingual LLMs: A Cognitive-Based Analysis of Relational Knowledge

ArXiv CS.AI 50%

arXiv:2611.06715v1 Announce Type: cross Abstract: Retrieval-augmented generation (RAG) assumes that language models can reason over relations stated in their context rather than fall back on associations memorized during pretraining. We study this evidence-conditioned relational...

https://arxiv.org/abs/2611.06715

08 CONF

Smartphone-Aided VR Exergames for HIIT: Gamifying Fitness with Real-Time Pose Estimation

ArXiv CS.AI 50%

arXiv:2611.06734v1 Announce Type: cross Abstract: This paper introduces an innovative virtual reality (VR) exergame system designed to gamify High-Intensity Interval Training (HIIT) workouts through 11 predefined combo moves, including squat-punch combinations, bilateral arm...

https://arxiv.org/abs/2611.06734

09 CONF

Context--Answer Metamodels: Forecasting LLM Activations Before the First Answer Token

ArXiv CS.AI 50%

arXiv:2611.06837v1 Announce Type: cross Abstract: Under stochastic decoding, a single context can elicit many different answers from a language model. Characterizing this variation typically requires many expensive autoregressive rollouts, each revealing only one realization. We...

https://arxiv.org/abs/2611.06837

10 CONF

Stabilizing Off-policy LLM Optimization with Prefix Importance Ratio

ArXiv CS.AI 50%

arXiv:2601.22718v2 Announce Type: replace Abstract: Reinforcement learning (RL) post-training has increasingly demonstrated strong ability to elicit reasoning behaviors in large language models (LLMs). For training efficiency, rollouts are typically generated in an off-policy...

https://arxiv.org/abs/2601.22718

11 CONF

Ordinal Credit Assignment: Multi-Step Reinforcement Learning from Comparisons Alone

ArXiv CS.AI 50%

arXiv:2605.06375v2 Announce Type: replace-cross Abstract: Policy-gradient methods need a scalar reward, yet for many tasks no formula captures what we want, while judging which of two outcomes is better is easy. We study reinforcement learning that receives only such comparisons...

https://arxiv.org/abs/2605.06375

12 CONF

Reinforcement Learning in GUI Agents: A Survey Toward Digital Inhabitants

ArXiv CS.AI 50%

arXiv:2604.27955v2 Announce Type: replace Abstract: Graphical User Interface (GUI) agents have emerged as a promising paradigm for intelligent systems that perceive and interact with graphical interfaces visually. Yet supervised fine-tuning alone cannot handle long-horizon...

https://arxiv.org/abs/2604.27955

13 CONF

Learning to Learn from Multimodal Experience

ArXiv CS.AI 50%

arXiv:2605.16857v2 Announce Type: replace Abstract: Multimodal agents increasingly interact with real-world-like environments and generate rich trajectories spanning perception, reasoning, and action. A key research focus is how to transform these raw multimodal trajectories...

https://arxiv.org/abs/2605.16857

14 CONF

A Roadmap to Impactful Pluralistic Alignment Research

ArXiv CS.AI 50%

arXiv:2607.22305v2 Announce Type: replace Abstract: Much research on responsible AI aims to change the behavior of AI systems that billions of people use. Pluralistic value alignment---the goal of building AI systems that represent and serve diverse human values and...

https://arxiv.org/abs/2607.22305

15 CONF

Using profiles of cognitive capability to assess AI suitability for workplace tasks

ArXiv CS.AI 50%

arXiv:2608.25623v2 Announce Type: replace Abstract: Organisations deploying AI face a scoping problem: which tasks can be automated, which should remain with humans, and which are best shared between the two. Aggregate benchmark scores provide little insight into where systems...

https://arxiv.org/abs/2608.25623

16 CONF

Reliability Engineering for AI Systems: Challenges, Methods, and Directions

ArXiv CS.AI 50%

arXiv:2609.35316v3 Announce Type: replace Abstract: AI reliability concerns whether an AI system performs its intended function dependably over a stated period and under stated operating conditions, with clear evidence. As these systems become more autonomous, this function...

https://arxiv.org/abs/2609.35316

17 CONF

AutoBCI: Forecast-Guided Agentic Neural Architecture Discovery for EEG-Based Brain-Computer Interfaces

ArXiv CS.AI 50%

arXiv:2609.35456v2 Announce Type: replace Abstract: EEG-based brain-computer interfaces support a broad range of applications, yet designing decoding architectures that perform well across diverse tasks remains challenging. We introduce AutoBCI, an agentic framework in which a...

https://arxiv.org/abs/2609.35456

18 CONF

MEA: A Reward-Driven Multi-Agent System for Faithful Model Explanations

ArXiv CS.AI 50%

arXiv:2610.02480v2 Announce Type: replace Abstract: Recent years have seen the employment of a plethora of machine learning (ML) models in high-stakes domains, but they remain largely opaque to the practitioners who act on their predictions. While post-hoc explanation methods...

https://arxiv.org/abs/2610.02480

19 CONF

RetryBench: Separating Retry Permission from Recovery Outcomes in GPT Tool Use

ArXiv CS.AI 50%

arXiv:2610.15884v4 Announce Type: replace Abstract: LLM agents that call external tools must decide what to do when a tool call times out: retrying risks duplicate side effects, while not retrying can leave the task incomplete. We present RetryBench, an evaluation protocol for...

https://arxiv.org/abs/2610.15884

20 CONF

Population-Aligned Persona Generation for LLM-based Social Simulation

ArXiv CS.AI 50%

arXiv:2509.10127v3 Announce Type: replace-cross Abstract: Recent advances in large language models (LLMs) have enabled large-scale, high-fidelity social simulations, creating new opportunities for computational social science. However, constructing persona sets that faithfully...

https://arxiv.org/abs/2509.10127

21 CONF

AutoEnergyDesign: An LLM-Driven Agent for Automatic Energy System Design Optimization

ArXiv CS.AI 50%

arXiv:2610.35160v2 Announce Type: replace Abstract: The design optimization of building energy systems represents a complex, high-dimensional search problem involving discrete technology selections, continuous sizing parameters, and intricate economic trade-offs in which every...

https://arxiv.org/abs/2610.35160

22 CONF

VCoT-SQL: Enhancing Text-to-SQL Generation through Dual Chain-of-Thought Validation Training

ArXiv CS.AI 50%

arXiv:2611.00815v2 Announce Type: replace Abstract: Open-source LLMs provide a flexible alternative to closed-source models for Text-to-SQL but still struggle with complex queries. Finetuning with chain-of-thought (CoT) data improves reasoning, yet standard datasets lack CoT...

https://arxiv.org/abs/2611.00815

23 CONF

ScamWorld: An Interactive Environment for Evaluating Web Agents Against Online Fraud

ArXiv CS.AI 50%

arXiv:2611.04879v2 Announce Type: replace Abstract: Web agents increasingly browse, shop, and manage accounts on behalf of their users. Prior work has shown that they can be misled by adversarial content, including prompt injections and deceptive interfaces. Beyond such attacks,...

https://arxiv.org/abs/2611.04879

24 CONF

ENACT-Heart -- ENsemble-based Assessment Using CNN and Transformer on Heart Sounds

ArXiv CS.AI 50%

arXiv:2502.16914v2 Announce Type: replace-cross Abstract: ENACT-Heart is a two-branch ensemble for heart-sound classification. A Vision Transformer (ViT) processes spectrogram images, while a convolutional neural network (CNN) processes visualizations of waveforms and spectral...

https://arxiv.org/abs/2502.16914

25 CONF

Benchmarking Pretrained Molecular Embedding Models For Molecular Representation Learning

ArXiv CS.AI 50%

arXiv:2508.06199v5 Announce Type: replace-cross Abstract: Pretrained molecular embedding models spanning graph neural networks, text transformers, and hybrid architectures have been widely used in drug discovery, yet no large-scale, rigorous comparison of their representations...

https://arxiv.org/abs/2508.06199

26 CONF

A Financial Brain Scan of the LLM

ArXiv CS.AI 50%

arXiv:2508.21285v4 Announce Type: replace-cross Abstract: We use sparse autoencoders (SAEs) to "brain scan" large language models, mapping internal activations to interpretable concepts and enabling targeted "steering" of specific features. We show that prompting models to adopt...

https://arxiv.org/abs/2508.21285

27 CONF

PointRAFT: 3D deep learning for high-throughput prediction of potato tuber weight from partial point clouds

ArXiv CS.AI 50%

arXiv:2512.24193v2 Announce Type: replace-cross Abstract: Key phenotypic traits, such as crop size, volume, and weight, can be estimated non-destructively from point clouds. In practice, however, point clouds of crops are typically incomplete due to self-occlusion, leading to...

https://arxiv.org/abs/2512.24193

28 CONF

Disentangle to create: Architecture-level disentangled diffusion model for few-shot Chinese font generation

ArXiv CS.AI 50%

arXiv:2602.18874v2 Announce Type: replace-cross Abstract: Few-shot Chinese font generation aims to synthesize new characters in a target style using only a handful of reference images. Achieving accurate content rendering and faithful style transfer requires effective...

https://arxiv.org/abs/2602.18874

29 CONF

Manifold-Matching Autoencoders: Embedding Data into Chosen Geometric Structures

ArXiv CS.AI 50%

arXiv:2603.16568v2 Announce Type: replace-cross Abstract: Manipulating the structure of data in the latent space of an autoencoder is a long-standing goal of representation learning. We propose a simple regularization, Manifold Matching, that imposes an explicit geometric...

https://arxiv.org/abs/2603.16568

30 CONF

Emergent Transfer of a Physics Foundation Model from Simulation to Laboratory Turbulence

ArXiv CS.AI 50%

arXiv:2606.01470v2 Announce Type: replace-cross Abstract: Whether physics foundation models can be usefully deployed on laboratory experiments remains an open question for scientific machine learning (ML). We test this question on the Rayleigh-Taylor instability (RTI), a...

https://arxiv.org/abs/2606.01470

31 CONF

MAIA: Multi-Agent Intent Articulation for Requirement Discovery in Art Commissions

ArXiv CS.AI 50%

arXiv:2609.12097v2 Announce Type: replace-cross Abstract: In bespoke art commissions, laypeople know what they feel but lack the words to specify it: one participant wanted a laid-off truck driver depicted as "a ghost in his own machine" but left the medium, scale, and palette...

https://arxiv.org/abs/2609.12097

32 CONF

SimForcing: Distilling Simulation Motion Priors into Real-Domain Robot World Models

ArXiv CS.AI 50%

arXiv:2610.06598v2 Announce Type: replace-cross Abstract: Action-conditioned robot world models must respond precisely to robot trajectories while preserving realistic visual dynamics, yet learning both from heterogeneous robot videos remains challenging. Simulation offers...

https://arxiv.org/abs/2610.06598

33 CONF

Revisiting the 16x Pretraining Compute Gap between Diffusion and Autoregressive Language Models

ArXiv CS.AI 50%

arXiv:2610.17590v2 Announce Type: replace-cross Abstract: Diffusion language models (DLMs) have attracted wide attention for their fast, parallel decoding, but they are believed to be far less efficient to train than autoregressive language models (ARMs): a widely cited study...

https://arxiv.org/abs/2610.17590

34 CONF

RefCraft: Benchmarking Visual-Reference-Conditioned Agentic Generation through Web Development

ArXiv CS.AI 50%

arXiv:2610.26826v2 Announce Type: replace-cross Abstract: Beyond basic tool use and coding skills, effectively incorporating visual references into agentic generation is equally essential for applying AI agents to real-world productivity tasks. We study this crucial capability...

https://arxiv.org/abs/2610.26826

35 CONF

Testing concept-level audits of melanoma classifiers with diffusion counterfactuals and matched positive controls

ArXiv CS.AI 50%

arXiv:2611.02119v2 Announce Type: replace-cross Abstract: Concept-based explanations of melanoma classifiers are rarely tested against controls with known targets. We built a counterfactual audit from a diffusion model conditioned on seven 7-point checklist concepts and tested...

https://arxiv.org/abs/2611.02119

36 CONF

Aligning Training and Rollout Policies for End-to-End FP4 RL on Mixture-of-Experts LLMs

ArXiv CS.AI 50%

arXiv:2611.04376v2 Announce Type: replace-cross Abstract: In this work, we study end-to-end FP4 reinforcement learning (RL) post-training for Mixture-of-Experts (MoE) large language models. Recent work on FP4 RL has introduced effective recovery recipes that restore accuracy...

https://arxiv.org/abs/2611.04376